# Gauge-Covariant Variational Free Energy and Exact Coarse-Graining

Agents as Local Sections at One Base Point: Collective Evidence, Attention, Inference Histories, and Formation

**Robert C. Dennis**

Independent Researcher

**Date: September 2026**

# Contents


**I Single-Point Foundations** **1**

**1 Program, Scope, and Standing Conventions** **2**
1.1 Program aim . . . 2
1.2 Objects and the local-first architecture . . . 3
1.3 Results and reading map . . . 3
1.4 Standing conventions and scope . . . 4
1.5 Status of statements . . . 6

**2 Geometry at One Base Point** **7**
2.1 The base point, the principal torsor, and associated law fibers . . . 7
2.2 Frames, passive re-expression, and the active gauge group at $c_*$ . . . 8
2.2.1 Pointwise frame comparisons and the active gauge group . . . 8
2.2.2 The common frame is a gauge choice . . . 9
2.3 Agents and cross-associated-bundle morphisms . . . 10
2.4 Independent links at $c_*$: vertex law, holonomy, edge twists . . . 11
2.5 Running example: three agents on one loop . . . 15

**3 Probability Semantics at $c_*$** **16**
3.1 Latent and record spaces at $c_*$ . . . 16
3.2 Structural data and kernel signatures . . . 17
3.3 Reference measures . . . 18
3.4 Kernels first, densities second . . . 18
3.5 Absolute continuity . . . 19
3.6 The relative-entropy chain rule and data processing . . . 20
3.7 Recognition marginals, couplings, and selection . . . 21
3.8 Levels of description . . . 23

**4 The Fixed Normalized Generative Law** **26**
4.1 Parameters, the directed latent law, and the typing prohibition . . . 26
4.2 Normalized record extension and the undirected alternative . . . 28
4.3 Covariance under passive frame re-expression . . . 30

**5 One Evidence Bound and Local-Collective Free Energy** **32**
5.1 Recognition coupling and total correlation . . . 32
5.2 Evidence domain . . . 33

5.3 The extended evidence bound and its record presentation . . . 34
5.4 Restriction principle . . . 37
5.5 The factorwise decomposition . . . 39
5.6 Observations as interaction records . . . 40
5.7 Local conditional free energy and the potential identity . . . 41
5.8 Additive accounting . . . 45
5.9 Model coordinates and accepted updates . . . 46
5.10 What remains open . . . 48

**6 Attention and Selective Communication 50**
6.1 Attention as a latent source label . . . 50
6.2 Communication as a normalized channel . . . 56
6.3 Factored links and loops . . . 57
6.4 Comparison, content, and the transformer correspondence . . . 59
6.4.1 Exact receiver update and transformer attention . . . 60
6.5 Compositions as records . . . 63
6.6 Information loss and descent . . . 64

**7 Fiber Information Geometry and Inference Histories 67**
7.1 Fiber statistical tensors . . . 68
7.2 Divergence jets and the orbit Fisher form . . . 69
7.3 Markov coarse morphisms and the vertical Fisher defect . . . 70
7.4 Configuration manifolds and the exact free-energy lift . . . 72
7.4.1 Canonicity boundary of the lift . . . 75
7.5 A finite correlated lift: the parity witness . . . 76
7.6 Natural-gradient orbits . . . 78
7.7 Fisher duration . . . 80
7.8 Orthogonal clocks . . . 82
7.9 Record clocks . . . 83
7.10 The coarse configuration map and oriented semiconjugacy . . . 84
7.11 Semiconjugacy on symmetry quotients . . . 87
7.12 Coarse-graining depth, orbit parameter, and duration . . . 87

**8 Two-Timescale Dynamics 89**
8.1 Block descent at fixed parameter . . . 89
8.2 Fast recognition and slow model coordinates . . . 90
8.3 Parent-sensitive inference and the prepared parent . . . 96

**II Exact Coarse-Graining at One Base Point 99**

**9 Coarse Channels 100**
9.1 Markov coarse maps . . . 100

9.2 KL contraction, equality, and recovery . . . 101
9.3 The common-channel VFE identity and the pointwise parent datum . . . 102
9.4 Pushforward, precomposition, and Galerkin restriction . . . 104
9.4.1 Galerkin operators and message passing . . . 105
9.5 Gauge and permutation equivariance . . . 106
9.6 Scale diagrams . . . 108
9.7 Labeled families and retained label information . . . 108
9.8 Retained-label Fisher asymptotics . . . 111

**10 Gauge-Retained Network State and Fixed Sectors 114**
10.1 Levelwise fiber state and scale intertwiners at $c_*$ . . . 114
10.2 Rooted trees, based holonomy, and dressed edges . . . 116
10.3 Linear feature tier and tree dependence . . . 118
10.4 Holonomy-conditioned marginal-law modes . . . 119
10.5 Structural fixed sectors . . . 122
10.6 Cut links and rectangular sheaf closure . . . 124
10.7 The frame sector . . . 129
10.8 An operational record at $c_*$ sensitive to loop holonomy . . . 130

**11 Exact Agent-Network Coarse-Graining 132**
11.1 Exact effective action . . . 132
11.2 Bounded action calculus and contraction . . . 134
11.3 Finite-network interaction coordinates . . . 135
11.4 The exact interaction step, its residual, and finite-step multipliers . . . 137
11.5 Exact closure generically requires hyperedges . . . 142
11.6 Fine–meta bridge kernels . . . 143
11.7 Exact attention between meta-agents . . . 144
11.8 Path laws and memory . . . 147
11.8.1 The hidden history and its filter . . . 149
11.8.2 Finite-memory projection . . . 150
11.9 The exact retained state . . . 151

**12 The Multiscale Graph VFE and Closing the Tower 154**
12.1 The full answer before decomposition . . . 154
12.2 A normalized hierarchical graph model . . . 155
12.3 The scale-local two-channel graph energy . . . 158
12.4 Partition variables and event pushforwards . . . 159
12.4.1 A random partition catalog . . . 159
12.4.2 Membership kernels and event pushforwards . . . 161
12.5 Coarse links and what the parent retains . . . 162
12.6 How a parent influences its children . . . 163
12.7 Closing the tower . . . 164

12.7.1 A declared root . . . 164
12.7.2 Three loci of participation . . . 165
12.7.3 The remaining hierarchy question . . . 169

**III Formation and Interpretation 171**

**13 Formation Witnesses 172**
13.1 The pointwise carrier and its restoring flag . . . 174
13.2 One normalized experiment with a nonlinear source record . . . 175
13.3 Physical posterior, restricted optimizer, and recognition . . . 177
13.4 Exact marked parents and their inherited inference . . . 178
13.5 Complete model scores and dissipation . . . 179
13.6 Preparation and successive activation . . . 180
13.7 The finite-depth formation theorem . . . 183
13.8 A two-triangle instance and the remaining scope . . . 186
13.9 An endogenous four-cycle witness . . . 187

**14 Interpretation and Its Limits 194**
14.1 Scaffolding, idle posits, and structural readings . . . 194
14.2 Participation, formation, and testability . . . 195

**15 Related Work 201**
15.1 Gauge theory, free energy, and multi-agent active inference . . . 201
15.2 Connection Laplacians, sheaves, and equivariant networks . . . 202
15.3 Averaging and path-space coarse-graining . . . 202
15.4 Information-geometric and Bayesian renormalization . . . 203
15.5 Network renormalization, graph coarsening, and multigrid . . . 204

**A Measure-Theoretic Lemmas and Integrability Witnesses 206**
A.1 Kernel measurability and proofs for the probability chapter . . . 206
A.2 The relative-log representation and integrability witnesses . . . 209

**B Deferred Proofs 211**
B.1 Proofs for Attention and Selective Communication . . . 211
B.2 Proofs for Fiber Information Geometry and Inference Histories . . . 214
B.3 Proofs for Coarse Channels . . . 218
B.4 Proofs for Exact Agent-Network Coarse-Graining . . . 222

**C Proofs for Finite-Depth Formation 225**
C.1 Exact marginals, rates, and parameter derivatives . . . 225
C.2 Prepared laws and coupled continuation . . . 228
C.3 Entropy estimates . . . 231
C.4 Exact fast marginals and complete source histories . . . 236

C.5 Coherence, source movement, and continuous paths . . . . . 240
C.6 The Gaussian core across an early interface . . . . . 245
C.7 The endogenous four-cycle . . . . . 247
C.8 Depth one under exclusive source gating . . . . . 253
C.9 Proof of Pairing Selection . . . . . 255
C.10 Proof of Joint Selection Along a Schedule . . . . . 266
C.11 Coherence and Persistence of the Selected Pairing . . . . . 271
C.11.1 The first contrast on the original episode window . . . . . 274
C.12 Activation from the Actual Conditioned Recognition Law . . . . . 276

**D Notation** **284**

**E Open Problems** **287**

**F Numerical Index** **293**

# List of Tables


1.1 Status of statements. . . . . 6

6.1 Correspondence between transformer attention and the pointwise attention network at $c_*$. . . . . 62

11.1 The retained state of exact agent-network coarse-graining at $c_*$. . . . . 152

13.1 Chapter-local symbols of the formation families, shared with the proof appendix. . 173

C.1 Lemma dependencies of the finite-depth formation theorem. . . . . 225

D.1 Notation of the main manuscript. . . . . 284

E.1 Open problems of the main manuscript. . . . . 287

F.1 Numerical index of the main manuscript. . . . . 293

# Part I

# Single-Point Foundations

# Chapter 1

# Program, Scope, and Standing Conventions

This manuscript asks what collective inference is available to agents that are local sections of a belief and a model associated bundle of one principal bundle, placed at one base point and interacting through independent links under one normalized joint law, and which parts of that inference survive exact coarse-graining of the agents into meta-agents. The answer has four main parts. The whole population has one evidence bound, an identity in the extended reals whose gap vanishes exactly at the posterior as a measure (Corollary 5.7 and theorem 5.8). Each agent's local objective is a conditional free energy whose differential is that of the collective free energy (Theorem 5.23). Coarse channels of the agent network preserve the evidence and decompose the free energy exactly (Theorem 9.6), with an effective action for the retained coordinates (Theorem 11.3). A finite-depth family forms successive meta-agents under free-energy descent with learning, in a stated scope (Theorem 13.17). Section 1.3 maps these results to the chapters, and Section 1.4 fixes, once for the whole manuscript, the defaults and the claims it declines.

## 1.1 Program aim

The program is a participatory multi-agent theory. Agents are local sections of associated bundles, collectives form through the agents' coupled inference, and parents act back on their children through normalized laws rather than through quantities that read the children's beliefs. Three requirements follow from that aim, and the manuscript is organized around them. Every collective object must be a complete probability object, constructed from the agents' mechanisms and interaction records or selected explicitly, never an additional agent and never a sum of local objectives that happens to look like one. Every statement about a parent must hold for the normalized parent law that coarse-graining actually produces. And formation of a parent must be a theorem about a declared family, with its scope stated, rather than a reading of a formula.

The present phase fixes one base point $c_*$ and settles the theory there before any transport across the base is considered. At $c_*$ the exact normalized parent law exists for every coarse channel of the agent network (Corollary 9.10), and Theorem 13.17 gives a formation witness under a supplied hierarchy of interaction strengths. On a homogeneous four-cycle a random initial source assignment already makes a collective form from every Gaussian initial law, with source learning idle (Proposition 14.4). At each sufficiently small fixed temperature, slow source learning selects

one matching in the assignment law for almost every Gaussian initial mean, while the frozen-prior pairing difference vanishes at long times (Theorem 14.5). How a parent's own dynamics are determined by the network it summarizes is Open Problem 11.38, the defining open problem of the phase.

## 1.2 Objects and the local-first architecture

Agent $i$ is the object $\mathcal{A}^i = (\mathcal{C}_i; q_i^b, q_i^m, u_i)$ of Definition 2.8: a belief section and a model section of the associated bundles $\mathcal{E}_b$ and $\mathcal{E}_m$ of one principal $G$-bundle, expressed in one common frame $u_i$, a gauge choice by Proposition 2.7. The two channels may carry inequivalent representations, and cross morphisms between them are additional structure (Definition 2.9). Interactions are normalized record kernels $K_a$ on the scopes of an interaction multigraph, and the links $\Theta_e^x$ on its edge copies are independent data with loop holonomy and per-edge invariant classes (Proposition 2.12). The primitive descriptive data are the agents, their local recognition laws, their evaluated mechanisms, and the record kernels. Population laws enter only by construction (Proposition 4.4) or by an explicit selection (Definition 3.15). The architecture is

$$\{\mathcal{A}^i, G_{\theta,i}, \mathbb{Q}_{i,o,X}, K_a\}_{i\in V, a\in\mathfrak{A}} \longrightarrow \begin{cases} \mathbb{P}_{\theta,V} & \text{normalized construction,} \\ \mathbb{Q}_{V,o,X} \in \mathrm{Cpl}(\mathbb{Q}_{\bullet,o,X}) & \text{measurable selection,} \end{cases} \longrightarrow \mathcal{L}_V, \tag{1.1}$$

so the constructed joint and the selected coupling precede every local and collective objective. On the evidence domain of Hypothesis 5.3 there is one identity,

$$\log p_{\theta,V}(o \mid X) = \mathcal{L}_V(\mathbb{Q}_{V,o,X}; X, o) + D_{\mathrm{KL}}\big(\mathbb{Q}_{V,o,X} \big\| \mathbf{\Pi}_{\theta,V,o,X}\big), \tag{1.2}$$

with equality of the bound and the log evidence exactly when the recognition coupling is the posterior as a measure, which matching coordinate marginals does not achieve (Theorem 5.8). The identity is invariant under passive re-expression of the frames when the kernels are covariant (Proposition 4.9 and corollary 4.10). Excluding the posterior from a recognition family makes every pointwise gap positive, but it does not make the optimized gap positive or attained (Section 5.4).

## 1.3 Results and reading map

Part I builds the single-point foundations. Chapter 2 fixes the geometry at $c_*$: the active gauge group is one copy of $G$, links carry per-edge twist classes, and three notions of agreement are separated. Chapter 3 fixes kernels before densities, proves the chain rule used everywhere, and distinguishes sample, law-point, and configuration levels. Chapter 4 builds the normalized joint on arbitrary record cycles and states the typing prohibition. Chapter 5 derives the evidence bound and the local-collective potential identity, with observations entering as interaction records (Theorem 5.17). Chapter 6 places attention inside the fixed joint as a latent source label (Proposition 6.2), builds selective communication, and records its correspondence with transformer attention (Table 6.1).

Chapter 7 develops fiber information geometry, Fisher contraction under Markov morphisms (Theorem 7.8), and inference histories measured by a Fisher duration that is invariant under reparameterization (Theorem 7.36). Chapter 8 types fast beliefs and slow models (Definition 8.2) and proves the reduced slow flow (Proposition 8.4).

Part II coarse-grains at one base point. Chapter 9 separates Markov coarse channels, energy restriction, and parameter aggregation (Section 9.4) and proves the exact coarse free energy and the pointwise parent datum (Theorem 9.6 and corollary 9.10). Chapter 10 gives the gauge-retained network state, the kernel–holonomy fixed sectors (Proposition 10.17), and the holonomy-conditioned marginal law (Theorem 10.9). Chapter 11 is exact coarse-graining of the agent network: effective action, residual, finite-step multipliers, and the native growth bound (Theorems 11.3 and 11.4 and proposition 11.19), with the retained state collected in Table 11.1 and a parent-sensitive mechanism with path laws and an exact finite-memory cost (Section 8.3). Chapter 12 assembles the multiscale free energy and closes a tower of priors (Section 12.7).

Part III turns to formation and interpretation. Chapter 13 constructs, for every finite depth, a family in which successive coherence episodes come from linear relaxation of a prepared mean through a supplied hierarchy of interaction strengths; source jumps and learning are perturbations the proof controls, not the cause of formation, and a separate witness produces a four-cycle endogenously (Proposition 13.22). Chapter 14 states what the mathematics does and does not support, and Chapter 15 positions the construction. The appendices hold measure-theoretic lemmas (Appendix A), the formation proofs (Appendix C), the notation (Appendix D), the open problems (Appendix E), and the numerical index (Appendix F).

Two supplements accompany the manuscript, and no result here depends on them. Supplement G develops the multivariate-Gaussian realization. Supplement B collects base transport, covers and connection descent, multi-point designs, infinite networks with their asymptotic apparatus, and the lattice cluster.

A first reading can take Chapters 2 to 4 for the objects and the joint, then Chapter 5 for the bound and Chapter 6 for attention, then Chapters 9 and 11 for exact coarse-graining and Chapter 13 for formation; proofs, witnesses, and the ledger can be recovered through their cross-references on a second reading.

## 1.4 Standing conventions and scope

Everything takes place at one base point. The design is $\{c_*\}$, design and incidence indices are dropped, and every agent's frame, beliefs, and models are points of the fibers over $c_*$. Base maps and base curves, neighborhoods, connection descent, transport along the base, section spaces over $\mathcal{C}$, and continuum law theory are deferred to the base supplement named in Section 1.3.

The following defaults hold throughout, and a statement mentions only departures from them. Measurable spaces are nonempty and standard Borel, and $\mathcal{P}(\mathsf{S})$ denotes the probability laws on a measurable space $\mathsf{S}$. Laws and kernels are normalized and measurable. A coarse channel

is a Markov kernel on latent coordinates that does not depend on the recognition law, leaves observations unchanged, and is vertical over the identity of $c_*$. An admitted observation has finite positive evidence. Structural data are parameters unless a configuration prior or a structural channel is declared, and the declaration names its law. The interaction graph at $c_*$ is the support of the source prior, a structural parameter that is complete unless masked. Gibbs attention weights are strictly positive on positive-prior edges when their costs are finite; the extended-cost rows of Proposition 12.9 can assign zero weight to infinite-cost edges. Dynamic structure is carried by attention weights and link values. Discrete creation and deletion of edges is not part of the model; it is Open Problem 6.9. Coarse multigraphs are induced structure. Absolute continuity of a recognition law with respect to the posterior is not a default: it is clause (H3) of Hypothesis 5.3, and its failure is the singular branch of the extended bound.

Three transports must remain distinct. A pointwise comparison $T_{ij}$ relates two frames in the same principal fiber and involves no curve (Proposition 2.5). A link $\Theta_e^x$ is an independent variable on an edge copy of the interaction multigraph (Definition 2.11). Parallel transport along a base curve is not used. Graph links are not base transport, and graph-link holonomy is neither the holonomy of a base connection nor a property of the topology of the principal bundle.

The word coarse has two uses that must not be merged. A Markov channel discards sample information while preserving normalization; a hard restriction substitutes a lower-dimensional configuration into an energy. They can agree in a special realization, but a shared word transfers no theorem, and parameter aggregation inside a closed family is a third operation (Section 9.4). Likewise, an inference history is an oriented unparameterized orbit in a configuration space, an ordered interaction-record path retained by exact coarse-graining is a sample coordinate, coarse-graining depth is a scale index, and a Fisher duration is a statistical length.

**Scope.** The manuscript declines the following claims. None of the inference index, flow parameter, transition index, ordered path coordinate, coarse-graining depth, or Fisher duration is identified with physical time; such an identification would need a named target clock, a calibration map with an uncertainty model, causal and compositional consistency, and discriminating evidence against alternative parameterizations. The contextual base is not spacetime and carries no Lorentzian structure. Sections, laws, and meta-agents are mathematical objects: calling a boundary node an agent is a naming decision, no claim of autonomous agency or of what exists is made, and population laws, population samples, and blocks are not agents. No canonical coupling selector, partition, or membership selector is asserted. No renormalization semigroup with rescaling, fixed points, or beta functions is constructed; Supplement B holds that asymptotic apparatus. The identification of any limiting object with a physical law is Open Problem 14.7. Declining a claim is not refuting it, and a refuted statement appears instead as an established negative result. [NOT-CLAIMED]

## 1.5 Status of statements

Every definition, standing hypothesis, lemma, proposition, theorem, corollary, conjecture, open problem, and claim-bearing remark is a numbered statement. Definitions, standing hypotheses, conjectures, and open problems carry their status in their heading type. Every lemma, proposition, theorem, corollary, and claim-bearing remark carries exactly one of three tags, one for all of its conclusions. Text that motivates, summarizes, or points carries no tag and makes no new claim. The fourth tag of Table 1.1, for declined claims, appears once in each scope paragraph and nowhere else. Open problems are collected in Table E.1, and numerical statements with their checks in Table F.1.

**Table 1.1:** Status of statements.

| Tag or heading | Meaning |
|---|---|
| [ESTABLISHED] | Proved in this manuscript. A refutation is an established negative statement. |
| [STANDARD] | A known result, cited to a checked primary source with a theorem or section locator. |
| [NUMERICAL] | Supported only by a named computation, indexed with its control and seed; computation is not proof. |
| Definition | A declared type, construction, or convention; nothing is proved. |
| Hypothesis | A standing restriction adopted for a chapter or part; it states what it excludes. |
| Conjecture | A precise unproved statement with a named observation that would refute it. |
| Open problem | An unsettled question with what holds, what is unproved, and what would settle it. |
| [NOT-CLAIMED] | A claim declined in a scope paragraph; declining is not refuting. |

## Physicist's summary

The manuscript studies agents that are local sections of belief and model bundles, all at one base point, coupled through independent links and interaction records under a single normalized joint law. There is one evidence bound for the whole population, the local objectives are exact conditional pieces of it, and coarse-graining the agent network into meta-agents is done with Markov channels that preserve the evidence and split the free energy exactly. Formation of meta-agents is proved for a declared finite-depth family in a stated regime and conjectured more generally. Physical time, spacetime, agency, and transport across the base are not claimed; they are either deferred to the base supplement or declined, and the defaults stated above are not repeated in later chapters.

# Chapter 2

# Geometry at One Base Point

This chapter answers what geometric data an agent carries at one base point $c_*$ and which of those data are invariant. The active gauge group at $c_*$ is one copy of $G$ inside the passive vertex group $\prod_{i\in V} G$ (Proposition 2.6); every independent link carries a gauge-invariant conjugacy class, and loop holonomy is the telescoped product of these per-edge twists (Proposition 2.12); frame agreement, agreement of an equivariant statistic, and law coherence up to holonomy stabilizers are distinct (Proposition 2.18). Chapter 4 uses passive re-expression for covariance of the generative law, Chapter 6 builds attention on link-transported comparisons, and Chapter 10 turns holonomy and the trivializing criterion (Proposition 2.13) into fixed sectors of coarse channels.

**Scope.** Normalizability of a law on links built from the loop penalties of Remark 2.16 is not asserted for a noncompact or general linear $G$ (the compact unitary case is in the remark), and whether law coherence bears on the selection of a partition into meta-agents is part of Open Problem 9.19. [NOT-CLAIMED]

## 2.1 The base point, the principal torsor, and associated law fibers

**Definition 2.1 (Contextual base and base point).** Let $\mathcal{C}$ be a finite-dimensional, second countable, Hausdorff smooth manifold whose points $c$ are called contexts, and fix one base point $c_* \in \mathcal{C}$.

**Definition 2.2 (Common principal gauge bundle and agent frames).** Let $\pi$ be a smooth principal right $G$-bundle for a Lie group $G$ and, for each agent $i \in V$, let $u_i$ be one local frame on an open support $\mathcal{C}_i \ni c_*$:

$$\pi : \mathscr{P}_G \longrightarrow \mathcal{C}, \tag{2.1}$$

$$u_i : \mathcal{C}_i \longrightarrow \mathscr{P}_G, \qquad \pi \circ u_i = \mathrm{id}_{\mathcal{C}_i}\,. \tag{2.2}$$

The fiber $(\mathscr{P}_G)_{c_*}$ is a right $G$-torsor and the values $u_i := u_i(c_*)$ are points of it. For any reference point $\sigma \in (\mathscr{P}_G)_{c_*}$ the frames are $u_i = \sigma \cdot U_i^{-1}$ for unique $U_i \in G$.

The same frame $u_i$ serves the belief and the model channel; Proposition 2.7 shows that this is a gauge choice. On the belief and model sample fibers $(\mathsf{K}, \mathscr{K})$ and $(\mathsf{M}, \mathscr{M})$ declare law fibers, measurable representations, which need not be equivalent, faithful, or of equal dimension, and their

pushforward actions on laws,

$$\mathcal{B}_b \subseteq \mathcal{P}(\mathsf{K}), \qquad \mathcal{B}_m \subseteq \mathcal{P}(\mathsf{M}), \tag{2.3}$$

$$\rho_b : G \longrightarrow \mathrm{Aut}(\mathsf{K}), \qquad \rho_m : G \longrightarrow \mathrm{Aut}(\mathsf{M}), \tag{2.4}$$

$$\widehat{\rho}_b(g)q = (\rho_b(g))_{\#}q, \qquad \widehat{\rho}_m(g)q = (\rho_m(g))_{\#}q, \tag{2.5}$$

so that $(\widehat{\rho}_x(g)q)(A) = q(\rho_x(g)^{-1}A)$.

**Hypothesis 2.3 (Selected smooth tier).** Whenever a differential statement is made, $\mathcal{B}_b$ and $\mathcal{B}_m$ are finite-dimensional smooth parametrized-measure models, each closed under its $G$-action, and the actions $(g, q) \mapsto (\rho_x(g))_{\#}q$ are smooth. The hypothesis excludes stratified and infinite-dimensional law fibers; a statistical reading of the tier further requires differentiability in quadratic mean and the integrability each tensor uses, and a Fisher tensor is a metric only where it is nondegenerate.

**Definition 2.4 (Belief and model associated bundles).** On invariant law fibers define the bundles below, with projections $\varpi_x : \mathcal{E}_x \to \mathcal{C}$ and the quotient convention of the second line,

$$\mathcal{E}_b = \mathscr{P}_G \times_{\widehat{\rho}_b} \mathcal{B}_b, \qquad \mathcal{E}_m = \mathscr{P}_G \times_{\widehat{\rho}_m} \mathcal{B}_m, \tag{2.6}$$

$$(u \cdot g, \beta) \sim (u, \widehat{\rho}_x(g)\beta). \tag{2.7}$$

At $c_*$ the frame $u_i$ identifies $\mathcal{B}_x$ with the fiber $(\mathcal{E}_x)_{c_*}$ by $\beta \mapsto [u_i, \beta]$. This is the standard associated-bundle construction [52, Section I.5], smooth under Hypothesis 2.3.

## 2.2 Frames, passive re-expression, and the active gauge group at $c_*$

A passive re-expression of agent $i$ is a choice $g_i \in G$ that replaces the frame as in the first line below. Since $[u_i, \beta] = [u_i' \cdot g_i, \beta] = [u_i', \widehat{\rho}_x(g_i)\beta]$, associated coordinates change by $\beta' = \widehat{\rho}_x(g_i)\beta$, and the represented sample-coordinate maps are those of the second line:

$$u_i' = u_i \cdot g_i^{-1}, \tag{2.8}$$

$$R_i^b = \rho_b(g_i), \qquad R_i^m = \rho_m(g_i), \qquad k_i' = R_i^b k_i, \quad m_i' = R_i^m m_i. \tag{2.9}$$

The group element, not a represented operator, multiplies the frame, and the tuple $(g_i)_{i \in V}$ ranges over the passive vertex group $\prod_{i \in V} G$.

### 2.2.1 Pointwise frame comparisons and the active gauge group

**Proposition 2.5 (Frame comparisons at one base point).** For $i, j \in V$ there is a unique $T_{ij} \in G$ with the first relation below, for every reference point $\sigma$. The comparisons satisfy $T_{ii} = e_G$, $T_{ji} = T_{ij}^{-1}$,

and $T_{ij}T_{jl} = T_{il}$, and under passive re-expressions they transform by the second line:

$$u_j = u_i \cdot T_{ij}, \qquad T_{ij} = U_i U_j^{-1}, \tag{2.10}$$

$$T'_{ij} = g_i T_{ij} g_j^{-1}, \qquad U'_i = g_i U_i. \tag{2.11}$$

[ESTABLISHED]

*Proof.* Freeness and transitivity of the torsor give uniqueness of $T_{ij}$, and $u_j = \sigma \cdot U_j^{-1} = u_i \cdot U_i U_j^{-1}$ gives the formula; the identities follow from uniqueness. From $u'_j = u_j \cdot g_j^{-1} = u'_i \cdot g_i T_{ij} g_j^{-1}$ and $u'_i = \sigma \cdot U_i^{-1} g_i^{-1}$ follow both laws in (2.11). □

**Proposition 2.6 (Active gauge image at one base point).** Let $F \in \mathrm{Aut}_G((\mathscr{P}_G)_{c_*})$, the group of right-$G$-equivariant bijections of the principal fiber over $c_*$, have local functions $F(u_i) = u_i \cdot \varphi_i$. Then

$$\varphi_j = T_{ij}^{-1} \varphi_i T_{ij}, \qquad i, j \in V, \tag{2.12}$$

and $F \mapsto (\varphi_i)_{i \in V}$ is a group isomorphism onto the subgroup $\mathcal{A}_{c_*} \subseteq \prod_{i \in V} G$ cut out by (2.12). For each $i_0 \in V$ the projection $(\varphi_i)_{i \in V} \mapsto \varphi_{i_0}$ is an isomorphism $\mathcal{A}_{c_*} \cong G$. In the frame $u_i$, $F$ changes the coordinates of both associated fibers by $\beta \mapsto \widehat{\rho}_x(\varphi_i)\beta$, which is the coordinate change of the passive re-expression $g_i = \varphi_i$. Hence the active transformations at $c_*$ form one copy $\mathcal{A}_{c_*} \cong G$ of the passive vertex group, a proper subgroup of $\prod_{i \in V} G$ when $G$ is nontrivial and $|V| \geq 2$, of dimension $\dim G$ against $|V| \dim G$ for a Lie group. The local functions at $c_*$ of every gauge automorphism of $\mathscr{P}_G$ form one element of $\mathcal{A}_{c_*}$. [ESTABLISHED]

*Proof.* Equivariance gives $F(u_j) = F(u_i \cdot T_{ij}) = u_i \cdot \varphi_i T_{ij} = u_j \cdot T_{ij}^{-1} \varphi_i T_{ij}$, and uniqueness in the torsor gives (2.12). Since $F(F'(u)) = F(u \cdot \varphi') = u \cdot \varphi\varphi'$, the map is a homomorphism. The value $\varphi_{i_0}$ determines $F$ on the whole torsor, so the map is injective, and for $\varphi_{i_0} \in G$ the map $u_{i_0} \cdot g \mapsto u_{i_0} \cdot \varphi_{i_0} g$ is an equivariant bijection, which proves $\mathcal{A}_{c_*} \cong G$. The coordinate statement is $F[u_i, \beta] = [u_i \cdot \varphi_i, \beta] = [u_i, \widehat{\rho}_x(\varphi_i)\beta]$, the passive coordinate change with $g_i = \varphi_i$. If $G \neq \{e_G\}$ and $i \neq j$, the tuple with $\varphi_i = a \neq e_G$ and all other entries $e_G$ violates (2.12), so the subgroup is proper. A gauge automorphism restricts to an equivariant bijection of each fiber. □

### 2.2.2 The common frame is a gauge choice

**Proposition 2.7 (Two atlases reduce to one by a passive re-expression).** Suppose belief and model coordinates were given in two frames $u_i^b, u_i^m$ of the same fiber. Then there is a unique $h_i \in G$ with the first relation below, and under independent re-expressions $u_i^{b\prime} = u_i^b \cdot (g_i^b)^{-1}$ and $u_i^{m\prime} = u_i^m \cdot (g_i^m)^{-1}$ it transforms by the second:

$$u_i^m = u_i^b \cdot h_i, \tag{2.13}$$

$$h'_i = g_i^b h_i (g_i^m)^{-1}. \tag{2.14}$$

In particular the re-expression $g_i^b = e_G$, $g_i^m = h_i$ of the model atlas alone makes the two frames equal. In two atlases the belief and model comparisons, the local functions $F(u_i^x) = u_i^x \cdot \varphi_i^x$ of a

fiber automorphism $F$, and the representatives of a cross morphism are related by

$$T_{ij}^{m} = h_i^{-1} T_{ij}^{b} h_j, \tag{2.15}$$

$$\varphi_i^{m} = h_i^{-1} \varphi_i^{b} h_i, \tag{2.16}$$

$$\Phi_i^{\mathrm{conv}} = \widehat{\rho}_m(h_i) \circ \Phi_i^{bm}, \tag{2.17}$$

where $\Phi_i^{bm}$ is the representative of $\Phi$ with source frame $u_i^b$ and target frame $u_i^m$, and $\Phi_i^{\mathrm{conv}}$ is its representative with both frames equal to $u_i^b$. All three relations become identities when $h_i = e_G$. [ESTABLISHED]

*Proof.* Uniqueness of $h_i$ is the torsor property. Substituting the re-expressions gives $u_i^{m\prime} = u_i^b \cdot h_i (g_i^m)^{-1} = u_i^{b\prime} \cdot g_i^b h_i (g_i^m)^{-1}$, which is (2.14). The two expressions $u_j^m = u_j^b \cdot h_j = u_i^b \cdot T_{ij}^b h_j$ and $u_j^m = u_i^m \cdot T_{ij}^m = u_i^b \cdot h_i T_{ij}^m$ give (2.15). Equivariance gives $F(u_i^m) = F(u_i^b) \cdot h_i = u_i^b \cdot \varphi_i^b h_i = u_i^m \cdot h_i^{-1} \varphi_i^b h_i$. Finally, $\Phi[u_i^b, \beta] = [u_i^m, \Phi_i^{bm}\beta] = [u_i^b, \widehat{\rho}_m(h_i)\Phi_i^{bm}\beta]$ by (2.7). □

No energy, kernel, or evaluator reads $h_i$, so the common frame

$$u_i^b = u_i^m = u_i, \qquad h_i = e_G, \qquad T_{ij}^b = T_{ij}^m = T_{ij} \tag{2.18}$$

is adopted throughout.

## 2.3 Agents and cross-associated-bundle morphisms

**Definition 2.8 (Agent).** Agent $i$ is the section-bearing object $\mathcal{A}^i$ of the first line below, with the sections of the second line, and a population sample at $c_*$ is a point $y$ of a product of sample fibers:

$$\mathcal{A}^i = \big(\mathcal{C}_i; q_i^b, q_i^m, u_i\big), \tag{2.19}$$

$$q_i^b \in \Gamma(\mathcal{C}_i, \mathcal{E}_b|_{\mathcal{C}_i}), \qquad q_i^m \in \Gamma(\mathcal{C}_i, \mathcal{E}_m|_{\mathcal{C}_i}), \tag{2.20}$$

$$y = \big((k_1, m_1), \ldots, (k_N, m_N)\big) \in \prod_{i \in V} (\mathsf{K} \times \mathsf{M}). \tag{2.21}$$

The values of $q_i^b, q_i^m$ at $c_*$ are points of $(\mathcal{E}_b)_{c_*}$ and $(\mathcal{E}_m)_{c_*}$, whose coordinates in the frame $u_i$ are laws in $\mathcal{B}_b$ and $\mathcal{B}_m$; the same letters denote these coordinate laws, which Definition 3.13 identifies with the local recognition marginals. Agents need not have equal supports.

**Definition 2.9 (Cross morphisms).** Cross morphisms are measurable fiber-preserving maps $\Phi : \mathcal{E}_b \to \mathcal{E}_m$ and $\widetilde{\Phi} : \mathcal{E}_m \to \mathcal{E}_b$, $\varpi_m \circ \Phi = \varpi_b$ and $\varpi_b \circ \widetilde{\Phi} = \varpi_m$. At $c_*$ they are maps between the fibers $(\mathcal{E}_b)_{c_*}$ and $(\mathcal{E}_m)_{c_*}$. They are not principal-bundle morphisms, gauge automorphisms, same-channel transports, or assumed inverses, and their existence is additional structure. Under Hypothesis 2.3 a differential statement also assumes them smooth. In the frame $u_i$ their representatives $\Phi_i : \mathcal{B}_b \to \mathcal{B}_m$ and $\widetilde{\Phi}_i : \mathcal{B}_m \to \mathcal{B}_b$ are defined by $\Phi[u_i, \beta] = [u_i, \Phi_i \beta]$ and $\widetilde{\Phi}[u_i, \beta] = [u_i, \widetilde{\Phi}_i \beta]$.

**Proposition 2.10 (Representatives and induced cross morphisms).** The representatives transform under passive re-expression by the first two lines below, and those of different agents are related by the third:

$$\Phi_i' = \widehat{\rho}_m(g_i) \circ \Phi_i \circ \widehat{\rho}_b(g_i)^{-1}, \tag{2.22}$$

$$\widetilde{\Phi}_i' = \widehat{\rho}_b(g_i) \circ \widetilde{\Phi}_i \circ \widehat{\rho}_m(g_i)^{-1}, \tag{2.23}$$

$$\Phi_i \circ \widehat{\rho}_b(T_{ij}) = \widehat{\rho}_m(T_{ij}) \circ \Phi_j, \qquad \widetilde{\Phi}_i \circ \widehat{\rho}_m(T_{ij}) = \widehat{\rho}_b(T_{ij}) \circ \widetilde{\Phi}_j. \tag{2.24}$$

If maps $\Phi_0 : \mathcal{B}_b \to \mathcal{B}_m$ and $\widetilde{\Phi}_0 : \mathcal{B}_m \to \mathcal{B}_b$ satisfy the first line below for every $g \in G$, then the maps of the second line

$$\Phi_0 \circ \widehat{\rho}_b(g) = \widehat{\rho}_m(g) \circ \Phi_0, \qquad \widetilde{\Phi}_0 \circ \widehat{\rho}_m(g) = \widehat{\rho}_b(g) \circ \widetilde{\Phi}_0, \tag{2.25}$$

$$[p, \beta] \longmapsto [p, \Phi_0 \beta], \qquad [p, \beta] \longmapsto [p, \widetilde{\Phi}_0 \beta] \tag{2.26}$$

are well-defined cross morphisms, whose representatives in every frame are $\Phi_0$ and $\widetilde{\Phi}_0$; conversely, $[p, \beta] \mapsto [p, \Phi_0 \beta]$ is well defined only if $\Phi_0$ satisfies (2.25). Not every cross morphism is induced in this way, and for a linear specialization with irreducible, inequivalent representations every equivariant linear map between the sample fibers vanishes, while cross morphisms still exist. [ESTABLISHED]

*Proof.* From $[u_i, \beta] = [u_i', \widehat{\rho}_b(g_i)\beta]$ we get $\Phi[u_i', \beta'] = [u_i, \Phi_i \widehat{\rho}_b(g_i)^{-1}\beta'] = [u_i', \widehat{\rho}_m(g_i)\Phi_i \widehat{\rho}_b(g_i)^{-1}\beta']$, which is (2.22); interchanging the channels gives (2.23). Evaluating $\Phi[u_j, \beta] = \Phi[u_i, \widehat{\rho}_b(T_{ij})\beta]$ in both frames gives (2.24). The representatives $(p \cdot g, \beta)$ and $(p, \widehat{\rho}_b(g)\beta)$ of one class have the same image under (2.26) exactly when $\widehat{\rho}_m(g)\Phi_0\beta = \Phi_0 \widehat{\rho}_b(g)\beta$, which proves both directions of the induced statement. For a cross morphism that is not induced, let $G = SO(2)$ act on the laws on $\mathbb{R}^2$ in both channels by the defining representation, fix $\sigma \in (\mathscr{P}_G)_{c_*}$, and put $\Phi[\sigma, \beta] = [\sigma, D_\# \beta]$ with $D = \mathrm{diag}(2, 1)$, which is well defined because every class has a unique representative with principal component $\sigma$. By (2.22) its representative in the frame $\sigma \cdot g^{-1}$ is $\widehat{\rho}_m(g) D_\# \widehat{\rho}_b(g)^{-1}$, and for the rotation $R_{\pi/2}$ one has $D_\#(R_{\pi/2})_\# \delta_{(1,0)} = \delta_{(0,1)} \neq \delta_{(0,2)} = (R_{\pi/2})_\# D_\# \delta_{(1,0)}$, so the representative depends on the frame. In the linear specialization, the kernel and image of an equivariant linear map between irreducible representation spaces are invariant subspaces, so the map is zero or an isomorphism, and inequivalence excludes an isomorphism. □

Law-level equivariance is weaker than linear equivariance: an equivariant Markov kernel induces an equivariant map on laws without supplying a linear map between the sample fibers.

## 2.4 Independent links at $c_*$: vertex law, holonomy, edge twists

**Definition 2.11 (Typed graph links).** Let $\mathfrak{I} = (V, E, {}^-)$ be a finite interaction multigraph at $c_*$: a finite set $E$ of oriented edge copies $e : j \to i$ with a fixed-point-free reversal $e \mapsto \bar{e}$, parallel copies allowed. For each channel $x \in \{b, m\}$ and each $e : j \to i$ declare a link $\Theta_e^x$ as in the first line below, with link map $\mathfrak{t}_e^x(u_j \cdot g) = u_i \cdot \Theta_e^x g$ of the fiber; endpoint notation $\Theta_{ij}^x$ is used only when the copy $j \to i$ is unique. Under passive re-expression the links obey the vertex law of the second and third lines,

with $R_i^x = \rho_x(g_i)$, because $\mathrm{t}_e^x(u'_j \cdot g) = u_i \cdot \Theta_e^x g_j^{-1} g = u'_i \cdot g_i \Theta_e^x g_j^{-1} g$. For a walk $p = (e_0, \dots, e_{r-1})$ from $j$ to $i$, that is, with $e_a : i_{a+1} \to i_a$, $i_0 = i$, and $i_r = j$, the path product and, when $i = j$, the holonomy are those of the fourth line, and the common-link specialization of the fifth line declares that the two link maps coincide in the common frame:

$$\Theta_e^x \in G, \qquad \Theta_{\bar{e}}^x = (\Theta_e^x)^{-1}, \tag{2.27}$$

$$\Theta_e^{x\prime} = g_i \Theta_e^x g_j^{-1}, \tag{2.28}$$

$$\rho_x(\Theta_e^{x\prime}) = R_i^x \rho_x(\Theta_e^x)(R_j^x)^{-1}, \tag{2.29}$$

$$\Theta_p^x = \Theta_{e_0}^x \Theta_{e_1}^x \cdots \Theta_{e_{r-1}}^x, \qquad \mathrm{Hol}^x(p) = \Theta_p^x, \tag{2.30}$$

$$\Theta_e^b = \Theta_e^m = \Theta_e. \tag{2.31}$$

**Proposition 2.12 (Per-edge twist invariance at one base point).** Fix a reference point $\sigma$ with $u_i = \sigma \cdot U_i^{-1}$, and define the per-edge twist $\mathring{\Theta}_e^x$ of $e : j \to i$ by the first line below; a closed walk $p$ based at $i_0$ then has the holonomy of the second line:

$$\mathrm{t}_e^x(\sigma) = \sigma \cdot \mathring{\Theta}_e^x, \qquad \mathring{\Theta}_e^x = U_i^{-1} \Theta_e^x U_j, \tag{2.32}$$

$$\mathrm{Hol}^x(p) = U_{i_0} \mathring{\Theta}_{e_0}^x \mathring{\Theta}_{e_1}^x \cdots \mathring{\Theta}_{e_{r-1}}^x U_{i_0}^{-1}, \qquad \mathrm{Hol}^x(p)' = g_{i_0} \mathrm{Hol}^x(p) g_{i_0}^{-1}. \tag{2.33}$$

Then $\mathrm{t}_e^x$ is a right-$G$-equivariant bijection that does not depend on the frames, and $\mathrm{t}_{\bar{e}}^x = (\mathrm{t}_e^x)^{-1}$. Every twist is unchanged by arbitrary passive re-expressions; a change of reference point $\sigma' = \sigma \cdot g_0$ conjugates every twist by the same $g_0$; and an active $F$ with $F(\sigma) = \sigma \cdot \kappa$ sends $\mathrm{t}_e^x$ to $F \circ \mathrm{t}_e^x \circ F^{-1}$ and every twist to $\kappa \mathring{\Theta}_e^x \kappa^{-1}$. Each edge copy therefore carries the conjugacy class $[\mathring{\Theta}_e^x]$, invariant under all passive and active changes, and two link data sets on the same multigraph are related by one active automorphism exactly when their twist tuples, computed from one reference point, are related by one simultaneous conjugation. The conjugacy class of $\mathrm{Hol}^x(p)$ and the condition $\mathrm{Hol}^x(p) = e_G$ are invariant. [ESTABLISHED]

*Proof.* Every point of the fiber is $u_j \cdot g$ for a unique $g$, so $\mathrm{t}_e^x$ is a well-defined equivariant bijection whose inverse $u_i \cdot g \mapsto u_j \cdot (\Theta_e^x)^{-1} g$ is $\mathrm{t}_{\bar{e}}^x$, and by the computation in Definition 2.11 re-expressed frames and links define the same map. Since $\sigma = u_j \cdot U_j$, $\mathrm{t}_e^x(\sigma) = u_i \cdot \Theta_e^x U_j = \sigma \cdot U_i^{-1} \Theta_e^x U_j$; the left side depends only on $\mathrm{t}_e^x$ and $\sigma$, which proves passive invariance, and $\mathrm{t}_e^x(\sigma \cdot g_0) = (\sigma \cdot g_0) \cdot g_0^{-1} \mathring{\Theta}_e^x g_0$. For the active statement, $F \mathrm{t}_e^x F^{-1}(\sigma) = F(\sigma \cdot \mathring{\Theta}_e^x \kappa^{-1}) = \sigma \cdot \kappa \mathring{\Theta}_e^x \kappa^{-1}$; an equivariant self-map of the torsor is determined by its value at $\sigma$, and $F \mapsto \kappa$ is a bijection onto $G$, which gives the criterion for two link data sets. Finally $\Theta_{e_a}^x = U_{i_a} \mathring{\Theta}_{e_a}^x U_{i_{a+1}}^{-1}$ telescopes, and $U'_{i_0} = g_{i_0} U_{i_0}$ by (2.11). □

**Proposition 2.13 (Trivializing criterion).** Fix a subfamily $I \subseteq V$ and elements $\Xi_e \in G$ on the internal oriented edge copies of $I$. The following are equivalent:

$$\begin{aligned} &\exists V_i \in G \ (i \in I) \quad \Xi_e = V_i V_j^{-1} \quad \text{for every internal } e : j \to i, \\ &\Xi_{e_0} \Xi_{e_1} \cdots \Xi_{e_{r-1}} = e_G \quad \text{for every closed internal walk } (e_0, \dots, e_{r-1}). \end{aligned} \tag{2.34}$$

When they hold, the potentials $V_i$ are unique up to right multiplication by one element of $G$ on each connected component of the internal multigraph. The statement and proof use only the group structure, so they hold for any group in place of $G$, in particular for a subgroup of $\mathrm{GL}(K)$ containing represented links. [ESTABLISHED]

*Proof.* The first condition telescopes around every closed walk. Conversely, backtracks $(e, \bar{e})$ are closed walks, so $\Xi_{\bar{e}} = \Xi_e^{-1}$ and reversed walks have inverse products. Choose a base vertex in each connected component and set $V_i = \Xi_p$ for an internal walk $p$ from that vertex to $i$, with products ordered as in (2.30); two choices differ by a closed walk and agree, and extending a walk to $j$ by $e : j \to i$ gives $V_i = \Xi_e V_j$. If also $\Xi_e = W_i W_j^{-1}$, then $W_i^{-1} V_i = W_j^{-1} V_j$ along every internal edge, so $W_i^{-1} V_i$ is constant on each component. □

**Definition 2.14 (Graph-flat, jointly trivializing, and frame-aligned links).** The links of channel $x$ are graph flat on $I \subseteq V$ when Proposition 2.13 holds for $\Xi_e = \Theta_e^x$, and $I$ is jointly trivializing when this holds in both channels. The links are frame aligned when $\Theta_e^x = T_{ij} = U_i U_j^{-1}$ for every $e : j \to i$, equivalently when every twist $\mathring{\Theta}_e^x$ equals $e_G$ and every link map $\mathfrak{t}_e^x$ is the identity of the fiber.

**Remark 2.15 (Frame alignment, graph flatness, and per-edge classes).** Frame alignment implies graph flatness, because $T_{ij} = U_i U_j^{-1}$ is a coboundary, but not conversely: on two vertices joined by one edge copy $e : 2 \to 1$ and its reversal, every closed walk is a backtrack, so any $\Theta_e^x \neq T_{12}$ is graph flat with potentials $V_1 = \Theta_e^x$, $V_2 = e_G$, while $[\mathring{\Theta}_e^x] \neq \{e_G\}$. Graph flatness of the links is graph flatness of the twists, since $\Theta_e^x = V_i V_j^{-1}$ exactly when $\mathring{\Theta}_e^x = W_i W_j^{-1}$ with $W_i = U_i^{-1} V_i$. A single link coordinate between distinct vertices carries no invariant: the passive re-expression $g_i = e_G$, $g_j = \Theta_e^x$ sends $\Theta_e^x$ to $e_G$. On a tree every loop product is trivial, yet the per-edge classes can all differ from $\{e_G\}$. Graph flatness is a property of the links alone: the one fiber over $c_*$ is always a trivial torsor, links on a triangle with nontrivial loop product are not graph flat, and all-identity links are graph flat whatever the topology of $\mathscr{P}_G$. [ESTABLISHED]

**Remark 2.16 (Invariant loop penalties).** Let $p$ be a closed walk based at $i_0$ and $\mathsf{U} = \rho_x(\mathrm{Hol}^x(p))$. If $\rho_x$ is unitary of dimension $n$, then $\frac{1}{2n} \operatorname{Re} \operatorname{Tr}(2I_n - \mathsf{U} - \mathsf{U}^*) = \frac{1}{2n} \operatorname{Tr}((I_n - \mathsf{U})^*(I_n - \mathsf{U})) \geq 0$, with equality exactly when $\mathsf{U} = I_n$, and the value is invariant under passive re-expression, which conjugates $\mathsf{U}$ by the unitary $R_{i_0}^x = \rho_x(g_{i_0})$. For $n \geq 2$ no norm on $\mathfrak{gl}(n, \mathbb{R})$ is invariant under conjugation by all of $\mathrm{GL}(n, \mathbb{R})$: conjugation by $\mathrm{diag}(2, 1, \ldots, 1)$ doubles the matrix unit $E_{12}$. A positive invariant penalty for general linear links therefore needs further data: if a symmetric positive definite fiber metric $\mathsf{G}^{\mathrm{fib}}$ at $i_0$ is declared to transform by $\mathsf{G}^{\mathrm{fib}} \mapsto (R_{i_0}^x)^{-\top} \mathsf{G}^{\mathrm{fib}} (R_{i_0}^x)^{-1}$, then $\operatorname{Tr}\big((\mathsf{G}^{\mathrm{fib}})^{-1} (I_n - \mathsf{U})^\top \mathsf{G}^{\mathrm{fib}} (I_n - \mathsf{U})\big)$ is invariant, nonnegative, and zero exactly when $\mathsf{U} = I_n$. For compact $G$ with $\rho_x$ unitary, the penalty $\frac{1}{2n} \operatorname{Tr}((I_n - \mathsf{U})^*(I_n - \mathsf{U}))$ of the first sentence lies in $[0, 2]$, so for finitely many closed walks $p$ with holonomies $\mathsf{U}_p = \rho_x(\mathrm{Hol}^x(p))$ and coefficients $\mathrm{c}_p \geq 0$ the law on links with density $\exp\big(-\sum_p \mathrm{c}_p \frac{1}{2n} \operatorname{Tr}((I_n - \mathsf{U}_p)^*(I_n - \mathsf{U}_p))\big)$ against the product, over one orientation of each edge copy, of the Haar probability measure of $G$ has normalizer in $[e^{-2\sum_p \mathrm{c}_p}, 1]$. [ESTABLISHED]

*Proof.* Unitarity gives $(I_n - \mathsf{U})^*(I_n - \mathsf{U}) = 2I_n - \mathsf{U} - \mathsf{U}^*$, whose trace is real and equals the squared Frobenius norm of $I_n - \mathsf{U}$. With $\Delta = I_n - \mathsf{U}$ and $\mathsf{G} = \mathsf{G}^{\mathrm{fib}}$, passive re-expression sends $\Delta$ to $R^x_{i_0}\Delta(R^x_{i_0})^{-1}$, cyclicity of the trace gives invariance, and $\mathrm{Tr}(\mathsf{G}^{-1}\Delta^\top\mathsf{G}\Delta) = \|\mathsf{G}^{1/2}\Delta\mathsf{G}^{-1/2}\|_{\mathrm{F}}^2$. A norm $\|\cdot\|_\star$ with $\|E_{12}\|_\star > 0$ cannot satisfy $\|2E_{12}\|_\star = \|E_{12}\|_\star$. For the compact case, $\mathrm{Tr}(2I_n - \mathsf{U} - \mathsf{U}^*) = 2n - 2\,\mathrm{Re}\,\mathrm{Tr}\,\mathsf{U}$ and $|\mathrm{Tr}\,\mathsf{U}| \le n$ for unitary $\mathsf{U}$, so that penalty lies in $[0, 2]$ and the exponent in $[-2\sum_p \mathrm{c}_p, 0]$; each $\mathsf{U}_p$ is a continuous function of the links, and integrating against a probability measure gives the two bounds on the normalizer. □

**Definition 2.17 (Law coherence).** Let $I \subseteq V$ induce a connected internal multigraph. Laws $q_i \in \mathcal{B}_x$, $i \in I$, are law coherent on $I$ when

$$q_i = \widehat{\rho}_x(\Theta^x_e)q_j \qquad \text{on each internal edge copy } e : j \to i. \tag{2.35}$$

For a $G$-set $S$ and a map $s : \mathcal{B}_x \to S$ with $s(\widehat{\rho}_x(g)q) = g \cdot s(q)$, the laws are $s$-agreeing when $s(q_i) = \Theta^x_e \cdot s(q_j)$ on each internal edge copy; the mean of laws on a linear representation space is the basic instance. The links have frame agreement on $I$ when they are graph flat on $I$.

**Proposition 2.18 (Three agreements).** Let $I$ be as in Definition 2.17 and fix $i_0 \in I$. (i) The laws are law coherent exactly when $q_{i_0}$ is fixed by $\widehat{\rho}_x(\mathrm{Hol}^x(p))$ for every closed internal walk $p$ based at $i_0$ and $q_i = \widehat{\rho}_x(\Theta^x_{p_i})q_{i_0}$ for some, hence every, internal walk $p_i$ from $i_0$ to $i$; for an equivariant statistic the same holds with $s(q_{i_0})$ fixed by the holonomy, so that frame agreement and mean agreement on edges, the two conditions of the first line below, are conditions on the links and on the means, and mean agreement holds exactly when the base mean satisfies the second line and the other means are its transports:

$$\mathrm{Hol}^x(p) = e_G \text{ for every closed internal walk } p, \quad \mu_i = \rho_x(\Theta^x_e)\mu_j \text{ on each internal } e : j \to i, \tag{2.36}$$

$$\mu_{i_0} \in \bigcap_{p \text{ closed at } i_0} \ker\big(\rho_x(\mathrm{Hol}^x(p)) - \mathrm{id}\big). \tag{2.37}$$

(ii) Law coherence implies $s$-agreement for every equivariant statistic $s$, and for the mean the converse fails. (iii) Frame agreement neither implies nor is implied by law coherence or by mean agreement. (iv) With several independent cycles, trivial holonomy or a fixed base mean for one loop is necessary but not sufficient. In particular, for $G = \mathrm{GL}(2, \mathbb{R})$ acting on laws on $\mathbb{R}^2$, three agents with the common law $\mathcal{N}(0, \Sigma)$, $\Sigma = \left(\begin{smallmatrix} 4 & 1 \\ 1 & 1 \end{smallmatrix}\right)$, and links $\Theta_{12} = \Sigma^{1/2}R(0.7)\Sigma^{-1/2}$ and $\Theta_{23} = \Theta_{31} = I_2$, with $R(\phi)$ the rotation by $\phi$, are law coherent while the triangle holonomy $\Theta_{12}$ is not the identity, is not orthogonal, and fixes no nonzero vector. [ESTABLISHED]

*Proof.* (i) Composing (2.35) along an internal walk $p$ from $j$ to $i$ gives $q_i = \widehat{\rho}_x(\Theta^x_p)q_j$, which yields both conditions. Conversely, let the conditions hold; for $e : j \to i$ and walks $p_i, p_j$ from $i_0$, the walk that follows $p_j$, then $e$, then $p_i$ backward is closed at $i_0$ with holonomy $(\Theta^x_{p_i})^{-1}\Theta^x_e\Theta^x_{p_j}$, so $\widehat{\rho}_x(\Theta^x_e)q_j = \widehat{\rho}_x(\Theta^x_{p_i})\widehat{\rho}_x\big((\Theta^x_{p_i})^{-1}\Theta^x_e\Theta^x_{p_j}\big)q_{i_0} = q_i$. The statistic version is the same argument in $S$, and the vectors left unchanged by the linear map $\rho_x(\mathrm{Hol}^x(p))$ form $\ker(\rho_x(\mathrm{Hol}^x(p)) - \mathrm{id})$. (ii) Apply $s$ to (2.35). For the converse take two agents joined by one edge copy with identity link and the

laws $\mathcal{N}(0, I_2)$ and $\mathcal{N}(0, 2I_2)$: the means agree and the laws do not. (iii) The same two agents with the laws $\mathcal{N}(0, I_2)$ and $\mathcal{N}(v, I_2)$, $v \neq 0$, have frame agreement without mean agreement or law coherence. In the displayed witness, $\Theta_{12}\Sigma\Theta_{12}^{\top} = \Sigma^{1/2}R(0.7)R(0.7)^{\top}\Sigma^{1/2} = \Sigma$, so every transported law equals $\mathcal{N}(0, \Sigma)$ and all means vanish, while $\Theta_{12}$ is similar to a rotation by $0.7$, hence has no eigenvalue $1$, and $\Theta_{12}\Theta_{12}^{\top} \neq I_2$ because $\Sigma$ is not a multiple of $I_2$. (iv) Two vertices joined by three parallel copies with $\Theta_e = \Theta_f \neq \Theta_g$ have trivial holonomy on the walk through $e$ and $\bar{f}$ and nontrivial holonomy on the walk through $e$ and $\bar{g}$. □

## 2.5 Running example: three agents on one loop

Take $G = SO(3)$, $\rho_b$ the defining representation on $\mathsf{K} = \mathbb{R}^3$, and $\rho_m$ the trivial representation on $\mathsf{M} = \mathbb{R}$, so the channels carry inequivalent representations. Let $V = \{1, 2, 3\}$, let $\mathfrak{I}$ be the triangle with edge copies $e_1 : 2 \to 1$, $e_2 : 3 \to 2$, $e_3 : 1 \to 3$ and their reversals, and, with $R_z(\phi)$ and $R_x(\phi)$ the rotations by $\phi$ about the third and first coordinate axes, declare the belief links of the first line below; model links act trivially. Choose $\sigma = u_1$ and, for display, $u_2 = u_3 = u_1$, so $U_i = e_G$ and each twist equals its link. The per-edge classes are the angles $\pi/2$, $\pi/2$, and $0$, and the loop $p = (e_1, e_2, e_3)$ based at agent $1$ has the holonomy of the second line,

$$\Theta^b_{e_1} = R_z(\tfrac{\pi}{2}), \qquad \Theta^b_{e_2} = R_x(\tfrac{\pi}{2}), \qquad \Theta^b_{e_3} = I_3, \tag{2.38}$$

$$\mathrm{Hol}^b(p) = R_z(\tfrac{\pi}{2})R_x(\tfrac{\pi}{2}), \qquad \mathrm{Tr}\,\mathrm{Hol}^b(p) = 0, \tag{2.39}$$

a rotation by $2\pi/3$; two quarter turns about one axis would give a rotation by $\pi$ with the same classes, so the class tuple does not determine the holonomy, while the tuple modulo one simultaneous conjugation does (Proposition 2.12). A passive re-expression $g_2$ changes $\Theta^b_{e_1}$ to $\Theta^b_{e_1}g_2^{-1}$ and $\Theta^b_{e_2}$ to $g_2\Theta^b_{e_2}$ but changes no twist, recomputed with $U_2' = g_2$, and not the holonomy. The loop is not graph flat, and the isotropic laws $q_i^b = \mathcal{N}(0, s^2 I_3)$ are law coherent on it, because every rotation fixes them (Proposition 2.18). One record reads the transported discrepancy of agents $1$ and $2$: the Markov kernel from $\mathsf{K} \times \mathsf{K}$ to $\{0, 1\}$ with

$$K_{a_1}\big(k_1, k_2, \{1\}\big) = \exp\Big(-\tfrac{1}{2}\big|k_1 - \rho_b(\Theta^b_{e_1})k_2\big|^2\Big), \tag{2.40}$$

which is gauge inert, since passive re-expressions rotate $k_1 - \rho_b(\Theta^b_{e_1})k_2$ by $\rho_b(g_1)$.

## Physicist's summary

At one base point all frames lie in one principal fiber, so they are compared directly by $T_{ij} = U_i U_j^{-1}$, the only active symmetry is one global rotation of that fiber, and per-agent frame changes are bookkeeping. Each independent link carries an invariant, its conjugacy class after dressing by the frames, whereas in lattice gauge theory only closed-loop products are invariant [85]. Agreement comes in three kinds, flat links, agreement of a statistic such as the mean, and law coherence, which survives nontrivial holonomy when the base law has a nontrivial stabilizer.

# Chapter 3

# Probability Semantics at $c_*$

On which measurable objects, and in which order, does every probabilistic statement at $c_*$ rest? Laws and kernels come first, densities follow as Radon–Nikodym derivatives against references fixed for a whole family, and logarithms of density ratios enter only under Definition 3.10; this order makes the evidence bound of Chapter 5 an identity in the extended reals. The main results are common domination of a mixed family exactly when its atoms are countable (Theorem 3.5), the relative-entropy chain rule in $[0, +\infty]$ (Lemma 3.12) that every later decomposition cites, and three levels of description, at which a law-point energy is not a sample energy and point-mass recognition of law points has infinite cost (Propositions 3.20 and 3.21). Chapter 4 builds the joint on these spaces, Chapter 5 applies the chain rule, and Chapters 6 and 13 use the levels.

**Scope.** Injectivity of a model evaluator is not assumed, and no regularity of the quotient of presentations with the same evaluated kernel is supplied; presentations are retained. The coupling selector is applied once, and its equivariance under frame re-expression (rerunning it on transformed marginals returns the transformed coupling) is neither assumed nor claimed. Open are a normalized law-point prior (Open Problem 3.23) and likelihoods for record families without a common dominating measure (Open Problem 3.9). [NOT-CLAIMED]

## 3.1 Latent and record spaces at $c_*$

**Definition 3.1 (Sample spaces and product sigma-algebras).** For each agent $i \in V$ let $(\mathsf{K}_i, \mathscr{K}_i)$ and $(\mathsf{M}_i, \mathscr{M}_i)$ be its belief and model sample fibers at $c_*$, identified with $\mathsf{K}$ and $\mathsf{M}$ by the frame $u_i$, and put $\mathsf{Y}_i = \mathsf{K}_i \times \mathsf{M}_i$. Let $\mathfrak{A}$ be a finite set of records; each $a \in \mathfrak{A}$ has a scope $\partial a \subseteq V$ and a record space $(\mathsf{O}_a, \mathscr{O}_a)$. The latent and observation spaces are

$$\mathsf{Y} = \prod_{i \in V} \mathsf{Y}_i, \qquad \mathscr{Y} = \bigotimes_{i \in V} (\mathscr{K}_i \otimes \mathscr{M}_i), \tag{3.1}$$

$$\mathsf{O} = \prod_{a \in \mathfrak{A}} \mathsf{O}_a, \qquad \mathscr{O} = \bigotimes_{a \in \mathfrak{A}} \mathscr{O}_a, \tag{3.2}$$

with points $y = ((k_i, m_i))_{i \in V}$ and $o = (o_a)_{a \in \mathfrak{A}}$, and $\mathsf{Y}_{\partial a} = \prod_{i \in \partial a} \mathsf{Y}_i$. An observation of a single agent is a record with singleton scope.

A population sample is a point of Y, and evaluating a sample field at $c_*$ returns

$$\mathrm{Ev}(y) = \big((k_i(c_*), m_i(c_*))\big)_{i\in V} \in \mathsf{Y}, \tag{3.3}$$

whereas the section values $q_i^b, q_i^m$ at $c_*$ are laws on $\mathsf{K}_i$ and $\mathsf{M}_i$, which become latent coordinates only at the law-point level of Section 3.8. Missing or aggregated observations are represented by an observation kernel into a coarser record space, which changes O and leaves Y unchanged.

## 3.2 Structural data and kernel signatures

**Definition 3.2 (Structural configuration and kernel signatures).** Fix a measurable configuration space $(\mathsf{X}, \mathscr{X})$. Every $X \in \mathsf{X}$ uses the same agent set $V$, the same records and scopes, and the same fiber identifications, so $(\mathsf{Y}, \mathscr{Y})$ and $(\mathsf{O}, \mathscr{O})$ do not vary with $X$. A configuration carries the interaction multigraph $\mathfrak{I}$ of Definition 2.11, its links $\Theta_e^x$, the frames $u_i$, and the parameters of the mechanisms. With $(\mathsf{I}_i, \mathscr{I}_i)$ an access space for each $i \in V$, the access maps, local recognition kernels, record kernels, and population laws have the signatures

$$\mathrm{Acc}_i : \mathsf{O}\times\mathsf{X} \longrightarrow \mathsf{I}_i, \quad \mathbb{Q}_i : \mathsf{I}_i \rightsquigarrow \mathsf{Y}_i, \qquad \mathbb{Q}_{i,o,X} := \mathbb{Q}_i(\mathrm{Acc}_i(o,X),\cdot), \tag{3.4}$$

$$K_a : (\mathsf{X}\times\mathsf{Y}_{\partial a}) \rightsquigarrow \mathsf{O}_a, \qquad a \in \mathfrak{A}, \tag{3.5}$$

$$\mathbb{P}_{\theta,V} : (\mathsf{X},\mathscr{X}) \longrightarrow \mathcal{P}\big(\mathsf{Y}\times\mathsf{O}, \mathscr{Y}\otimes\mathscr{O}\big), \qquad X \longmapsto \mathbb{P}_{\theta,V}(\cdot \mid X), \tag{3.6}$$

$$\mathbb{Q}_V : (\mathsf{O}\times\mathsf{X}, \mathscr{O}\otimes\mathscr{X}) \longrightarrow \mathcal{P}\big(\mathsf{Y},\mathscr{Y}\big), \qquad (o,X) \longmapsto \mathbb{Q}_{V,o,X}, \tag{3.7}$$

where each record coordinate occurs once, $\mathbb{P}_{\theta,V}$ is constructed in Chapter 4, and $\mathbb{Q}_{V,o,X}$ is the coupling selected in Section 3.7; neither is primitive data. The identity access map is the unrestricted interface; a stricter one restricts what the local law may depend on, not the population law (Proposition 3.18). An agent's generative-model structure is the pair $(q_i^m, \mathrm{ev}_i)$ of its model law and the evaluator of Definition 3.3. The generative kernel reads only $\theta$, $X$, and latent coordinates (Hypothesis 4.3), and a configuration prior, when declared, is a law on $(\mathsf{X}, \mathscr{X})$.

**Definition 3.3 (Model presentations and evaluated kernels).** For each $i \in V$ and structural value $X$, let $(\mathsf{S}_i^X, \mathscr{S}_i^X)$ and $(\mathsf{T}_i^X, \mathscr{T}_i^X)$ contain every source and output coordinate of agent $i$'s generative factor, and let $(\mathsf{M}_i, \mathscr{M}_i)$ be its presentation space, so that a model sample is $m_i \in \mathsf{M}_i$ and the model law is $q_i^m \in \mathcal{P}(\mathsf{M}_i)$. The evaluator is a map

$$\mathrm{ev}_i : \mathsf{M}_i \longrightarrow \mathrm{Kern}(\mathsf{S}_i^X, \mathsf{T}_i^X), \qquad m_i \longmapsto K_{i,m_i}^X, \tag{3.8}$$

with $(m_i, z) \mapsto K_{i,m_i}^X(z, B)$ measurable for every $B \in \mathscr{T}_i^X$ and $K_{i,m_i}^X(z, \mathsf{T}_i^X) = 1$. Each source space retains every observation, structural, or parent coordinate its factor requires. In the definite case $q_i^m = \delta_{m_i}$. Neither a model law, a recognition parameter, nor a posterior is an argument of an evaluated kernel.

## 3.3 Reference measures

**Definition 3.4 (Reference measures).** Each Euclidean coordinate carries its Borel sigma-algebra and Lebesgue measure, and each countable coordinate its power set and counting measure. A mixed real coordinate carries one parameter-independent reference for the whole family of laws placed on it, generative and recognition alike: when the union $A_*$ of the atom sets of that family is countable, it carries $\nu_{\mathrm{mix}}$ of (3.10); when $A_*$ is uncountable no density model is declared on that coordinate (Theorem 3.5). Any other fiber carries an explicitly declared sigma-finite measure. With $\nu_i^k, \nu_i^m, \nu_a$ the component measures, the product references are

$$\nu^{\mathsf{Y}} = \bigotimes_{i\in V} (\nu_i^k \otimes \nu_i^m), \qquad \nu^{\mathsf{O}} = \bigotimes_{a\in\mathfrak{A}} \nu_a. \tag{3.9}$$

**Theorem 3.5 (Common domination for a mixed family).** Let $(\mu_t)_{t\in\mathbb{T}}$ be any family of Borel probability measures on $\mathbb{R}$ with no singular-continuous part, written $\mu_t = f_t \,\mathrm{Leb} + \sum_{a\in A_t} p_t(a)\delta_a$ with $A_t := \{a \in \mathbb{R} : \mu_t(\{a\}) > 0\}$. A single sigma-finite Borel measure dominates every $\mu_t$ if and only if $A_* := \bigcup_{t\in\mathbb{T}} A_t$ is countable. When $A_* = \{a_1, a_2, \ldots\}$, one such parameter-independent measure is

$$\nu_{\mathrm{mix}} := \mathrm{Leb} + \sum_{n\geq 1} 2^{-n}\delta_{a_n}. \tag{3.10}$$

In particular, the moving-atom family $\mu_t = \frac{1}{2}\mathcal{N}(0,1) + \frac{1}{2}\delta_t$, $t \in \mathbb{R}$, has no common sigma-finite dominating measure. [ESTABLISHED]

The proof is in Appendix A.

**Remark 3.6 (Dirac families are not dominated).** On an uncountable standard Borel space, an uncountable family of Dirac laws admits no common sigma-finite dominating measure, so density charts cannot be inferred from a general law fiber. [ESTABLISHED]

*Proof.* A dominating measure must give positive mass to every atom of the family, and a sigma-finite measure has at most countably many points of positive mass: on each finite-measure member of a countable cover, the points of mass at least $1/n$ are finitely many. □

## 3.4 Kernels first, densities second

**Definition 3.7 (Normalized kernels; densities as derivatives).** The objects (3.6) and (3.7) are probability kernels. When $\mathbb{P}_{\theta,V}(\cdot \mid X) \ll \nu^{\mathsf{Y}} \otimes \nu^{\mathsf{O}}$, its density, the Radon–Nikodym derivative $p_{\theta,V}(o, y \mid X)$, and the evidence, its marginal density, are determined by

$$\mathbb{P}_{\theta,V}(dy, do \mid X) = p_{\theta,V}(o, y \mid X)\nu^{\mathsf{Y}}(dy)\nu^{\mathsf{O}}(do), \tag{3.11}$$

$$p_{\theta,V}(o \mid X) = \int_{\mathsf{Y}} p_{\theta,V}(o, y \mid X)\nu^{\mathsf{Y}}(dy). \tag{3.12}$$

When $\mathbb{Q}_{V,o,X} \ll \nu^{\mathsf{Y}}$ its density is written $q_{V,o,X}(y)$. Densities are determined up to null sets of the reference, and versions that agree outside such a set denote the same density. Measurability in $(o, X)$ of the evidence and of recognition expectations is Proposition A.3, and jointly measurable densities are Corollary A.2; an unnormalized energy in a later chapter comes with a finite positive normalizer.

**Hypothesis 3.8 (Regular-conditional convention).** For fixed $(\theta, X)$ the observation marginal and the posterior, a regular conditional probability, are

$$\mathbb{P}^{\mathsf{O}}_{\theta,V,X}(A) := \mathbb{P}_{\theta,V}(\mathsf{Y} \times A \mid X), \qquad A \in \mathscr{O}, \tag{3.13}$$

$$\mathbf{\Pi}_{\theta,V,o,X}(B) = \mathbb{P}_{\theta,V}(B \mid o, X), \qquad B \in \mathscr{Y}, \tag{3.14}$$

the posterior being unique for $\mathbb{P}^{\mathsf{O}}_{\theta,V,X}$-almost every $o$. When the domination of (3.11) holds, (3.12) is a density of $\mathbb{P}^{\mathsf{O}}_{\theta,V,X}$, and after a change on a $\nu^{\mathsf{O}}$-null set there is a measurable set $\mathsf{O}^{\mathrm{reg}}_{\theta,V,X}$ with

$$\mathbb{P}^{\mathsf{O}}_{\theta,V,X}\big(\mathsf{O}^{\mathrm{reg}}_{\theta,V,X}\big) = 1, \qquad 0 < p_{\theta,V}(o \mid X) < \infty \quad \text{on } \mathsf{O}^{\mathrm{reg}}_{\theta,V,X}, \tag{3.15}$$

$$\mathbf{\Pi}_{\theta,V,o,X}(B) = \frac{\int_B p_{\theta,V}(o, y \mid X)\nu^{\mathsf{Y}}(dy)}{p_{\theta,V}(o \mid X)}, \qquad o \in \mathsf{O}^{\mathrm{reg}}_{\theta,V,X}. \tag{3.16}$$

Every fixed-observation posterior, evidence-bound, or free-energy statement is asserted for $\mathbb{P}^{\mathsf{O}}_{\theta,V,X}$-almost every $o$, represented by $o \in \mathsf{O}^{\mathrm{reg}}_{\theta,V,X}$. The hypothesis excludes everywhere-pointwise statements unless a jointly measurable density, an evidence representative, and a conditional version are declared as model data; the joint measure alone does not determine a posterior at an exceptional observation (Remark A.5).

**Open Problem 3.9 (Nondominated likelihood families).** The density formulas here use, for every record kernel $K_a$, one sigma-finite $\nu_a$ dominating $\{K_a(X, y_{\partial a}, \cdot)\}$ and a jointly measurable density. For a family without one, such as $K_a(X, y_{\partial a}, \cdot) = \delta_{h(y_{\partial a})}$ with $h$ taking uncountably many values (Remark 3.6), the joint law and its regular conditional posterior exist, but the likelihood $dK_a(X, y_{\partial a}, \cdot)/d\nu_a(o_a)$ is undefined and the evidence density can fail to exist. Open is a treatment of the evidence bound, its local disintegrations, and the coarse channels at the level of joint measures and regular conditional laws, or hypotheses restoring a jointly measurable likelihood, with a rule for propagating record-null choices of versions. It is settled by such a treatment with proofs that the exact identities of Chapters 5 and 9 hold in it, or by a record family for which they provably have no measure-level counterpart.

## 3.5 Absolute continuity

**Definition 3.10 (Logarithmic density ratios and relative entropy).** Wherever a logarithmic density ratio or its expectation is written, the numerator measure is absolutely continuous with

respect to the denominator measure. Relative entropy is

$$D_{\mathrm{KL}}(Q\|P) = \begin{cases} \int \log\Big(\frac{dQ}{dP}(y)\Big) Q(dy), & Q \ll P, \\ +\infty, & \text{otherwise}, \end{cases} \tag{3.17}$$

the first branch being the divergence of Kullback and Leibler [53] and the value $+\infty$ off absolute continuity the convention of Csiszár [24]. The negative part of the integrand is $Q$-integrable, since $\int_{\{dQ/dP<1\}} -\log(dQ/dP)dQ \le 1/e$, so the first branch lies in $[0,+\infty]$.

**Proposition 3.11 (Density criterion for absolute continuity).** Let $Q$ and $P$ be probability measures with $Q \ll \nu$ and $P \ll \nu$ for a sigma-finite $\nu$, with densities $q$ and $p$, and put $Z = \{p = 0\}$. Then

$$Q \ll P \iff Q(Z) = 0 \iff q = 0 \quad \nu\text{-almost everywhere on } Z, \tag{3.18}$$

$$Q(S) = 1, \qquad P(S) = 0, \tag{3.19}$$

where the first line holds, and whether or not $Q$ is dominated by $\nu$, a measurable $S$ satisfying the second gives $Q \not\ll P$ and $D_{\mathrm{KL}}(Q\|P) = +\infty$. [ESTABLISHED]

*Proof.* First $P(Z) = \int_Z p(y)\nu(dy) = 0$, so $Q \ll P$ gives $Q(Z) = 0$. Conversely, if $Q(Z) = 0$ and $P(A) = 0$, then $\nu(A \cap \{p > 0\}) = 0$, so $Q(A \cap \{p > 0\}) = 0$, while $Q(A \cap Z) \le Q(Z) = 0$. The second equivalence follows from $Q(Z) = \int_Z q(y)\nu(dy)$. The last claim is the definition of absolute continuity together with (3.17). □

A law carried by a Lebesgue-null set, such as a proper affine subspace, has no Lebesgue density, and (3.19) shows its singularity (Proposition 5.4).

## 3.6 The relative-entropy chain rule and data processing

**Lemma 3.12 (Relative-entropy chain rule, data processing, and bijection invariance).** Let $P, Q$ be laws on $\mathsf{Y}_1 \times \mathsf{Y}_2$ with first marginals $P_1, Q_1$ and conditional kernels $P_{2|1}, Q_{2|1}$. Then:

(a) the map $y_1 \mapsto D_{\mathrm{KL}}(Q_{2|1}(y_1,\cdot)\|P_{2|1}(y_1,\cdot))$ is measurable and

$$D_{\mathrm{KL}}(Q\|P) = D_{\mathrm{KL}}(Q_1\|P_1) + \int D_{\mathrm{KL}}\big(Q_{2|1}(y_1,\cdot)\big\|P_{2|1}(y_1,\cdot)\big) Q_1(dy_1) \tag{3.20}$$

in $[0,+\infty]$, with a right side that does not depend on the chosen conditional versions;

(b) on a finite product $\mathsf{Y} = \prod_{b=1}^{B} \mathsf{Y}_b$, for any law $Q$ with marginals $Q_b$ and any laws $\rho_b$ on $\mathsf{Y}_b$,

$$D_{\mathrm{KL}}\Big(Q\Big\|\bigotimes_b \rho_b\Big) = D_{\mathrm{KL}}\Big(Q\Big\|\bigotimes_b Q_b\Big) + \sum_b D_{\mathrm{KL}}(Q_b\|\rho_b) \tag{3.21}$$

in $[0,+\infty]$;

(c) for a law $\mu$ on $\mathsf{Y}_1$, the minimum of $D_{\mathrm{KL}}(Q\|P)$ over all laws $Q$ with $Q_1 = \mu$ is $D_{\mathrm{KL}}(\mu\|P_1)$, attained by $\mu(dy_1)P_{2|1}(y_1, dy_2)$, uniquely when $D_{\mathrm{KL}}(\mu\|P_1) < \infty$;

(d) for laws $P, Q$ on $\mathsf{Y}$ and every Markov kernel $K : \mathsf{Y} \rightsquigarrow \mathsf{Z}$,

$$D_{\mathrm{KL}}(QK\|PK) \le D_{\mathrm{KL}}(Q\|P) \qquad \text{in } [0, +\infty]; \tag{3.22}$$

(e) for a bimeasurable bijection $A : \mathsf{Y} \to \mathsf{Y}'$,

$$\frac{d(A_\# Q)}{d(A_\# P)} = \frac{dQ}{dP} \circ A^{-1} \quad A_\# P\text{-almost surely, if } Q \ll P, \tag{3.23}$$

$$D_{\mathrm{KL}}(A_\# Q\|A_\# P) = D_{\mathrm{KL}}(Q\|P) \quad \text{in } [0, +\infty]. \tag{3.24}$$

[ESTABLISHED]

The proof is in Appendix A. Clauses (a), (b), and (d) need no absolute-continuity or finiteness hypothesis.

## 3.7 Recognition marginals, couplings, and selection

**Definition 3.13 (Local recognition laws and their coupling class).** With $\mathrm{pr}_i : \mathsf{Y} \to \mathsf{Y}_i$ the agent-block projection and $\mathrm{pr}_i^k, \mathrm{pr}_i^m$ the coordinate projections of $\mathsf{Y}_i$, for each admitted $(o, X)$ put

$$\mathrm{Cpl}(\mathbb{Q}_{\bullet,o,X}) := \Big\{\mathbb{J} \in \mathcal{P}(\mathsf{Y}) : (\mathrm{pr}_i)_\# \mathbb{J} = \mathbb{Q}_{i,o,X} \text{ for every } i \in V\Big\}, \tag{3.25}$$

$$q_i^{b;o,X} := (\mathrm{pr}_i^k)_\# \mathbb{Q}_{i,o,X}, \qquad q_i^{m;o,X} := (\mathrm{pr}_i^m)_\# \mathbb{Q}_{i,o,X}. \tag{3.26}$$

The belief and model recognition marginals $q_i^{b;o,X}, q_i^{m;o,X}$ are the values at $c_*$, in the frame $u_i$, of agent $i$'s sections $q_i^b, q_i^m$ at the admitted $(o, X)$. They are local data; a member of (3.25) is population-level dependence data.

**Proposition 3.14 (Existence and convexity of the coupling class).** For every admitted $(o, X)$ the class $\mathrm{Cpl}(\mathbb{Q}_{\bullet,o,X})$ is nonempty and convex, and

$$(o, X) \longmapsto \bigotimes_{i \in V} \mathbb{Q}_{i,o,X} \tag{3.27}$$

is a measurable product-coupling kernel. [ESTABLISHED]

*Proof.* On measurable rectangles the product is a product of measurable functions, and a monotone-class argument extends measurability; its $i$-th marginal is $\mathbb{Q}_{i,o,X}$ because the other factors have unit mass. Marginalization is affine, so convex combinations of couplings are couplings. □

**Definition 3.15 (Selected population recognition kernel).** A selection rule $\mathsf{S}$ is declared, with law spaces carrying their evaluation sigma-algebras, such that the population recognition kernel

$$\mathbb{Q}_V : (\mathsf{O} \times \mathsf{X}, \mathscr{O} \otimes \mathscr{X}) \longrightarrow \mathcal{P}(\mathsf{Y}, \mathscr{Y}), \qquad (o, X) \longmapsto \mathbb{Q}_{V,o,X}, \tag{3.28}$$

$$\mathbb{Q}_{V,o,X} = \mathsf{S}_{o,X}\big((\mathbb{Q}_{i,o,X})_{i\in V}\big) \in \mathrm{Cpl}(\mathbb{Q}_{\bullet,o,X}), \tag{3.29}$$

has $(o, X) \mapsto \mathbb{Q}_{V,o,X}(B)$ measurable for every $B \in \mathscr{Y}$. The law may correlate agents and channels; the product coupling is one admissible selection, not the default, and copulas, common latent variables, ordered conditionals, or graphical families supply others.

**Hypothesis 3.16 (Constrained coupling selection).** When the selection must lie in an imposed subfamily $\mathrm{Cpl}^{\mathrm{res}}(\mathbb{Q}_{\bullet,o,X}) \subseteq \mathrm{Cpl}(\mathbb{Q}_{\bullet,o,X})$, every fiber $\mathrm{Cpl}^{\mathrm{res}}(\mathbb{Q}_{\bullet,o,X})$ is assumed nonempty, and a selector $\mathsf{S}$ from these fibers is assumed with $(o, X) \mapsto \mathsf{S}_{o,X}((\mathbb{Q}_{i,o,X})_{i\in V})(B)$ measurable for every $B \in \mathscr{Y}$. The hypothesis excludes restricted families whose fibers are empty or admit no measurable selection; the two conditions are independent, and the product coupling proves neither.

**Proposition 3.17 (Coordinate marginals do not determine a sufficiently rich joint).** For a finite family of standard Borel blocks, the class $\mathrm{Cpl}((\mu_i)_{i\in V})$ is a singleton if and only if at most one marginal $\mu_i$ is non-Dirac. Hence, when two local recognition laws are non-Dirac, the local marginals do not determine $\mathbb{Q}_{V,o,X}$. [ESTABLISHED]

*Proof.* Let $\mu_1, \mu_2$ be non-Dirac. A non-Dirac law charges some measurable set with probability in $(0, 1)$, since otherwise the members or complements of probability one of a countable separating generator would meet in a set of probability one with at most one point. Choose such sets $A, B$ and, for $0 < |\varepsilon| < 1$, put

$$d\mathbb{J}_\varepsilon(x, x') = \big[1 + \varepsilon(\mathbf{1}_A(x) - \mu_1(A))(\mathbf{1}_B(x') - \mu_2(B))\big] d\mu_1(x) d\mu_2(x'). \tag{3.30}$$

The bracket is at least $1 - |\varepsilon| > 0$ and integrating in either variable removes the perturbation, so $\mathbb{J}_\varepsilon$ couples $\mu_1, \mu_2$, while $\mathbb{J}_\varepsilon(A \times B) = \mu_1(A)\mu_2(B) + \varepsilon\mu_1(A)(1 - \mu_1(A))\mu_2(B)(1 - \mu_2(B))$ varies with $\varepsilon$; tensoring with the other marginals gives distinct couplings of the family. Conversely, if at most one marginal is non-Dirac, every coupling is carried by the Dirac atoms in the other coordinates, and its law there is the remaining marginal. □

**Proposition 3.18 (Limits of access restriction).** A restriction of the access maps constrains the local laws and not the population law, and it has an evidence price that no selector removes.

(i) Let $V = \{1, 2\}$ with binary blocks, let both local laws be $\mathrm{Ber}(\frac{1}{2})$ at every observation, so that every access map is respected, and fix $0 < r < 1$. The selector that returns the coupling with $\mathbb{Q}_{V,o,X}(y_1 = y_2) = (1 + r)/2$ when a record $o_2$ of agent $2$ equals $1$ and $(1 - r)/2$ otherwise satisfies (3.29), and the population law depends on $o_2$ while no local law does. A decentralized information claim therefore also requires a selector indexed by $((\mathrm{Acc}_i(o, X))_{i\in V}, X)$ alone.

(ii) Let $V = \{1, 2\}$ with binary latents, the prior with mass $0.45$ on each of $(0, 0)$ and $(1, 1)$ and $0.05$ on each of $(0, 1)$ and $(1, 0)$, and singleton records $o_1, o_2$ that report $y_1, y_2$ correctly with probability $0.8$, independently given $y$. Let $\mathrm{Acc}_1(o, X) = o_1$, so that $\mathbb{Q}_{1,o,X} = \mathrm{Ber}(t)$ is the same at $o = (1, 0)$ and $o = (1, 1)$. The posterior marginals of $y_1$ at these observations, written as the probabilities of $y_1 = 0$ and $y_1 = 1$, are $\Pi_1^{(0)} = (37/89, 52/89)$ and $\Pi_1^{(1)} = (13/161, 148/161)$, and every recognition law with this access restriction satisfies

$$\begin{aligned}\max_{o_2\in\{0,1\}} D_{\mathrm{KL}}\big(\mathbb{Q}_{V,(1,o_2),X}\big\|\mathbf{\Pi}_{\theta,V,(1,o_2),X}\big) &\geq \max_{o_2\in\{0,1\}} D_{\mathrm{KL}}\big(\mathrm{Ber}(t)\big\|\Pi_1^{(o_2)}\big)\\ &\geq \sup_{\xi\in[0,1]}\big(-\log Z_\xi\big) = 0.08849\ldots,\end{aligned} \tag{3.31}$$

where $Z_\xi = \sum_{y_1} \Pi_1^{(0)}(y_1)^\xi \Pi_1^{(1)}(y_1)^{1-\xi}$. The last bound is sharp: the supremum is attained at $\xi \approx 0.548$, and at $t \approx 0.783$ the middle term equals it. Hence the equality condition of Theorem 5.8 lies outside the access-restricted class, and the exact E-coordinate of Corollary 5.22, which optimizes over every law with a prescribed complement marginal, need not be access restricted.

[ESTABLISHED]

*Proof.* (i) Both couplings have $\mathrm{Ber}(\frac{1}{2})$ marginals. (ii) Bayes' rule gives the posterior marginals, and projection onto $y_1$ with Lemma 3.12(d) gives the first inequality. For $\xi \in [0, 1]$ let $\Pi^{[\xi]} = Z_\xi^{-1}(\Pi_1^{(0)})^\xi(\Pi_1^{(1)})^{1-\xi}$. Expanding logarithms gives $\xi D_{\mathrm{KL}}(\mathrm{Ber}(t)\|\Pi_1^{(0)}) + (1-\xi)D_{\mathrm{KL}}(\mathrm{Ber}(t)\|\Pi_1^{(1)}) = D_{\mathrm{KL}}(\mathrm{Ber}(t)\|\Pi^{[\xi]}) - \log Z_\xi \geq -\log Z_\xi$, and a maximum dominates every convex combination. The concave function $\xi \mapsto -\log Z_\xi$ has derivative $\mathbb{E}_{\Pi^{[\xi]}} \log(\Pi_1^{(1)}/\Pi_1^{(0)})$, which vanishes at its interior maximizer $\xi_\star$, where the two divergences from $\Pi^{[\xi_\star]}$ are therefore equal; the identity with $\mathrm{Ber}(t) = \Pi^{[\xi_\star]}$ shows that both equal $-\log Z_{\xi_\star}$, the Chernoff information [21, Section 11.9]. The value and maximizers follow from the closed forms, and every access-restricted law has a positive gap at one of the two observations. □

## 3.8 Levels of description

**Definition 3.19 (Sample, law-point, and configuration levels).** At the sample level the agent block is $\mathsf{Y}_i$ of Definition 3.1, the population recognition law $\mathbb{Q}_{V,o,X}$ is a law of samples, and the belief of agent $i$ is the recognition marginal $q_i^{b;o,X}$ of (3.26). Generative kernels and the record kernels (3.5) read sample coordinates. At the law-point level, law points are latent coordinates, and the law-point blocks are

$$\mathsf{Y}_i^{\mathrm{law}} = (\mathcal{E}_b)_{c_*} \times (\mathcal{E}_m)_{c_*} \ni \big(\mathfrak{q}_i^b, \mathfrak{q}_i^m\big), \qquad \mathsf{Y}^{\mathrm{law}} = \prod_{i\in V} \mathsf{Y}_i^{\mathrm{law}}, \tag{3.32}$$

possibly augmented by further latent coordinates; through the frame $u_i$, $\mathfrak{q}_i^b$ and $\mathfrak{q}_i^m$ are laws on $\mathsf{K}_i$ and $\mathsf{M}_i$. A generative law at this level is a law on $\mathsf{Y}^{\mathrm{law}} \times \mathsf{O}$ whose kernels may read law points, and a recognition law is a law over law points. At the configuration level the latent state is $(w, y) \in \mathsf{Y}^{\mathrm{cfg}} \times \mathsf{Y}$, where the configuration space $\mathsf{Y}^{\mathrm{cfg}}$ carries a measurable presentation map $\mathrm{pres} : \mathsf{Y}^{\mathrm{cfg}} \to \mathsf{Y}^{\mathrm{law}}$

and may carry further latent coordinates such as random links. A lifted generative law is a law $\mathbb{P}^{\mathrm{cfg}}_{\theta,V}(dw, dy, do \mid X)$, and a lifted recognition law has the form $\mathbb{J}(dw, dy) = \mathbb{J}^{\mathrm{cfg}}(dw)\kappa^{\mathrm{pres}}(w, dy)$ with a presentation kernel $\kappa^{\mathrm{pres}} : \mathsf{Y}^{\mathrm{cfg}} \rightsquigarrow \mathsf{Y}$ whose marginal at each agent is the law presented by the corresponding component of $\mathrm{pres}(w)$.

The levels define different optimization problems, so an energy enters a free energy only after its level is named. The source energies of (6.2) and every record kernel are sample-level. The transported divergences (12.12) and the role divergences (6.35) are measurable $[0, +\infty]$-valued functions of law points (Corollary A.4), admissible as generative energies when their arguments are latent law points; read at the sample level they would make a generative factor read recognition marginals, which Hypothesis 4.3 excludes. A block admitted by Hypothesis 5.5 with law points is a law-point or configuration block, and the formation model (13.1) is configuration-level, with records that read the configuration.

**Proposition 3.20 (A law-point energy is not a sample energy).** Let $q_i$ be a receiver law and $p_1, p_2$ source laws in $\mathcal{B}_x$ with $D_{\mathrm{KL}}(q_i \| \widehat{\rho}_x(\Theta^x_e)p_k) < \infty$ for $k = 1, 2$, whose transported laws have densities against a common dominating measure that differ on a set of positive $q_i$-measure. Then for $0 < s < 1$

$$D_{\mathrm{KL}}\big(q_i \| \widehat{\rho}_x(\Theta^x_e)(sp_1 + (1-s)p_2)\big) < sD_{\mathrm{KL}}\big(q_i \| \widehat{\rho}_x(\Theta^x_e)p_1\big) + (1-s)D_{\mathrm{KL}}\big(q_i \| \widehat{\rho}_x(\Theta^x_e)p_2\big),$$

and no recognition-independent sample energy $D$ satisfies $\mathbb{E}_{q_i \otimes q} D = D_{\mathrm{KL}}(q_i \| \widehat{\rho}_x(\Theta^x_e)q)$ at $q = p_1$, at $q = p_2$, and at one strict mixture of them. [ESTABLISHED]

*Proof.* Pushforward preserves mixtures. Let $\alpha, \beta$ be the $q_i$-densities of the absolutely continuous parts of the two transported laws. Finite divergences give $q_i \ll$ each transported law, so $\alpha, \beta > 0$ $q_i$-almost everywhere and the divergences are $\int -\log \alpha(y) q_i(dy)$ and $\int -\log \beta(y) q_i(dy)$, well defined because $\int (\log \alpha(y))^+ q_i(dy) \le \int \alpha(y) q_i(dy) \le 1$. The mixture has relative density $s\alpha + (1-s)\beta$, and $-\log(s\alpha + (1-s)\beta) \le -s \log \alpha - (1-s) \log \beta$, strictly on the set $\{\alpha \neq \beta\}$ of positive $q_i$-measure; integrating gives the strict inequality. An energy $D$ contributes $\mathbb{E}_{q_i \otimes q} D$, affine in $q$ along mixtures, so matching the divergence at $p_1$ and $p_2$ forces the affine value at the mixture, which the strict inequality excludes. □

**Proposition 3.21 (Point-mass recognition of law points has infinite cost).** At the law-point level let $\mathbb{P}^{\mathrm{law}}_{\theta,V}(\cdot \mid X)$ be a generative law with latent marginal $\mathbb{P}^{\mathrm{law},\mathsf{Y}}_{\theta,V}$ and posterior $\mathbf{\Pi}^{\mathrm{law}}_{\theta,V,o,X} \ll \mathbb{P}^{\mathrm{law},\mathsf{Y}}_{\theta,V}$ at an observation with finite positive evidence; the domination holds when every record kernel has a density against a fixed reference. If $\bar{\mathsf{q}}$ is not an atom of the prior, then

$$D_{\mathrm{KL}}\big(\delta_{\bar{\mathsf{q}}} \big\| \mathbf{\Pi}^{\mathrm{law}}_{\theta,V,o,X}\big) = +\infty, \tag{3.33}$$

and the extended evidence bound of $\delta_{\bar{\mathsf{q}}}$ is $-\infty$. Consequently, against a non-atomic law-point prior, a row or an update computed at fixed law points is not the recognition optimum of a finite evidence bound at the law-point level: with the law points latent it is a generative posterior quantity conditional

on them, and with the law points declared as parameters in $(\theta, X)$ it is a function of parameters whose change is an M-coordinate subject to Definition 5.28. [ESTABLISHED]

*Proof.* Singletons are measurable, and domination gives $\mathbf{\Pi}^{\mathrm{law}}_{\theta,V,o,X}(\{\bar{\mathfrak{q}}\}) = 0$, so (3.19) holds for $\delta_{\bar{\mathfrak{q}}}$ with $S = \{\bar{\mathfrak{q}}\}$. The extended bound (5.15) subtracts this divergence from a finite log evidence, and a recognition optimum with finite bound has finite divergence. □

**Proposition 3.22 (Configuration-lift identity).** Let $\kappa^{\mathrm{cfg}} : \mathsf{Y} \rightsquigarrow \mathsf{Y}^{\mathrm{cfg}}$ be a Markov kernel, and let the lifted generative law be $\mathbb{P}_{\theta,V}(dy, do \mid X)\kappa^{\mathrm{cfg}}(y, dw)$, so that its $(y, o)$ marginal is the sample-level law and its records read only sample coordinates. Then at every $o \in \mathsf{O}^{\mathrm{reg}}_{\theta,V,X}$ the evidence is unchanged, the lifted posterior may be selected as $\mathbf{\Pi}_{\theta,V,o,X}(dy)\kappa^{\mathrm{cfg}}(y, dw)$, and every lifted recognition law $\mathbb{J}$ with sample marginal $\bar{\mathbb{Q}}$ and conditional kernel $\mathbb{J}(dw \mid y)$ satisfies

$$\mathcal{L}^{\mathrm{ext,cfg}}_V(\mathbb{J}; X, o) = \mathcal{L}^{\mathrm{ext}}_V(\bar{\mathbb{Q}}; X, o) - \mathbb{E}_{\bar{\mathbb{Q}}}\Big[D_{\mathrm{KL}}\big(\mathbb{J}(\cdot \mid Y)\big\|\kappa^{\mathrm{cfg}}(Y, \cdot)\big)\Big] \tag{3.34}$$

in $[-\infty, \log p_{\theta,V}(o \mid X)]$, where $\mathcal{L}^{\mathrm{ext,cfg}}_V$ is (5.15) for the lifted law. The lift never raises the bound of the induced sample law, and when that bound is finite it leaves it unchanged exactly when $\mathbb{J}(\cdot \mid y) = \kappa^{\mathrm{cfg}}(y, \cdot)$ for $\bar{\mathbb{Q}}$-almost every $y$. [ESTABLISHED]

*Proof.* The $(y, o)$ marginal of the lifted law is $\mathbb{P}_{\theta,V}$, so the evidence is unchanged, and disintegrating along $o$ gives the stated posterior. Lemma 3.12(a), with $y$ as the first coordinate, gives the lifted divergence as $D_{\mathrm{KL}}(\bar{\mathbb{Q}}\|\mathbf{\Pi}_{\theta,V,o,X})$ plus the conditional term in $[0, +\infty]$, and the conditional term vanishes exactly under the stated almost-everywhere equality. Substitution in (5.15) subtracts two terms in $[0, +\infty]$ from a finite log evidence. □

**Open Problem 3.23 (A normalized law-point prior).** No chapter constructs a normalized generative law on $\mathsf{Y}^{\mathrm{law}}$ for a general class of law points, which the law-point energies (12.12) and (6.35) need to be evidence-bound terms rather than a separately declared potential. The law fibers supply no canonical sigma-finite reference ($\mathcal{P}(\mathsf{K})$ is in general infinite dimensional), and the Haar measure of a noncompact $G$ is infinite, so a Haar prior on a free orbit is not normalizable; the formation model supplies one normalized configuration-level prior, (13.7), for finite-dimensional Gaussian presentations with a proper, non-invariant reference. It is settled for a class of law points by a sigma-finite reference on the law-point or presentation space, a finite normalizer for the exponentiated energies, the transformation law of the prior under the active gauge action, record kernels that read law points, and a compatibility relation between the law-point posterior and the sample-level recognition marginals.

## Physicist's summary

Everything probabilistic at $c_*$ is a law or a kernel, with densities against references fixed per family, and one chain rule in $[0, +\infty]$ does all later bookkeeping. Local recognition laws do not determine their coupling, restricted access has an exact evidence price, and attention's transported divergence is a law-point energy.

# Chapter 4

# The Fixed Normalized Generative Law

This chapter answers how one normalized joint law of latent samples and records is built at $c_*$ and what passive re-expression of the frames does to it. The joint is normalized for every finite record hypergraph, cyclic ones included (Proposition 4.4), local normalization of reciprocal latent conditionals supplies no joint (Proposition 4.6), and under kernel covariance the evidence and the bound of a transported coupling are invariant (Proposition 4.9 and corollary 4.10). Its evidence is bounded in Chapter 5, and attention and coarse channels act on it in Chapters 6 and 9.

**Scope.** The cross morphisms supply no generative factor, and a model with factors in both directions needs its own normalized joint. No quotient of the parameter space by passive re-expression is formed, and no parameter count subtracts the dimension of the vertex group, whose action can have nontrivial stabilizers. [NOT-CLAIMED]

## 4.1 Parameters, the directed latent law, and the typing prohibition

**Definition 4.1 (Parameters, structural data, and generative factors).** Fix the agent set $V = \{1, \dots, N\}$ and the spaces and references of Chapter 3. Structural data $X \in \mathsf{X}$ contain a finite directed acyclic graph $\Gamma$ on $V$ with parent sets $\mathrm{pa}(i)$, the interaction multigraph $\mathfrak{I}$ with its links, and the frames; the numerical parameters that a learning coordinate may move are denoted $\theta$. The split between $X$ and $\theta$ is declared, not derived, and every typing statement below applies to both. For each agent $i$ declare the Markov kernels of the first line below, its model and state factors, which are its model and belief priors and are generative factors, not components of the agent of Definition 2.8; at a root, where $\mathrm{pa}(i) = \varnothing$, the first is a proper model law and the second a bridge from the agent's own model coordinate to its belief coordinate. The state factor is the evaluated kernel of Definition 3.3 with source $\mathsf{S}_i^X = \prod_{j \in \mathrm{pa}(i)} \mathsf{K}_j$, the empty product being a singleton $\{*\}$, and target $\mathsf{T}_i^X = \mathsf{K}_i$ (second line):

$$P^m_{\theta,i}(dm_i \mid m_{\mathrm{pa}(i)}, X), \qquad P^k_{\theta,i}(dk_i \mid k_{\mathrm{pa}(i)}, m_i, X), \tag{4.1}$$

$$\mathrm{ev}_i(m_i) = K^X_{i,m_i}, \qquad K^X_{i,m_i}(k_{\mathrm{pa}(i)}, B) := P^k_{\theta,i}(B \mid k_{\mathrm{pa}(i)}, m_i, X). \tag{4.2}$$

With $\mathsf{Y}_{\mathrm{pa}(i)} = \prod_{j\in\mathrm{pa}(i)} \mathsf{Y}_j$, the agent mechanism is the ordered composite below, which reads the whole parent block $y_{\mathrm{pa}(i)}$ and has unit mass:

$$G_{\theta,i} : (\mathsf{X} \times \mathsf{Y}_{\mathrm{pa}(i)}) \rightsquigarrow \mathsf{Y}_i, \tag{4.3}$$

$$G_{\theta,i}\big(d(k_i, m_i) \mid y_{\mathrm{pa}(i)}, X\big) := P^m_{\theta,i}(dm_i \mid m_{\mathrm{pa}(i)}, X) K^X_{i,m_i}(k_{\mathrm{pa}(i)}, dk_i). \tag{4.4}$$

Definition 8.2 divides these coordinates into fast recognition and slow model coordinates. At a root the state factor induces, for a model law $s$ on $\mathsf{M}_i$, the belief law

$$\widetilde{\Phi}^{\mathrm{gen}}_i(s)(B) = \int_{\mathsf{M}_i} P^k_{\theta,i}(B \mid m, X) s(dm), \qquad B \in \mathscr{K}_i, \tag{4.5}$$

which can serve as the representative at $c_*$ of the cross morphism $\widetilde{\Phi}$ of Definition 2.9 when that identification is declared; at a nonroot the evaluated kernel is indexed by the parent belief tuple and is a kernel-valued bridge, not a map of law fibers.

**Definition 4.2 (Finite directed latent law).** Choose a topological ordering $v_1, \dots, v_N$ of $\Gamma$. The latent law is the ordered composition of the first line below, where the arrow denotes ordered application of kernels as in the second line, for bounded measurable $f$; every conditioning coordinate is introduced before its kernel is applied:

$$\mathbb{P}^{\mathsf{Y}}_{\theta,V}(dy \mid X) := \prod_{l=1}^{\overrightarrow{N}} G_{\theta,v_l}\big(dy_{v_l} \mid y_{\mathrm{pa}(v_l)}, X\big), \tag{4.6}$$

$$\int f(y) \mathbb{P}^{\mathsf{Y}}_{\theta,V}(dy \mid X) = \int G_{\theta,v_1}(dy_{v_1} \mid X) \cdots \int G_{\theta,v_N}\big(dy_{v_N} \mid y_{\mathrm{pa}(v_N)}, X\big) f(y). \tag{4.7}$$

When the factors have densities against the component references, the densities are

$$g_{\theta,i}(k_i, m_i \mid y_{\mathrm{pa}(i)}, X) = p^m_{\theta,i}(m_i \mid m_{\mathrm{pa}(i)}, X) p^k_{\theta,i}(k_i \mid k_{\mathrm{pa}(i)}, m_i, X), \tag{4.8}$$

$$p^{\mathsf{Y}}_{\theta,V}(y \mid X) = \prod_{i\in V} g_{\theta,i}(k_i, m_i \mid y_{\mathrm{pa}(i)}, X), \qquad \mathbb{P}^{\mathsf{Y}}_{\theta,V}(dy \mid X) = p^{\mathsf{Y}}_{\theta,V}(y \mid X) \nu^{\mathsf{Y}}(dy). \tag{4.9}$$

**Hypothesis 4.3 (Typing prohibition).** Each factor in (4.1) and each record kernel is a function only of $\theta$, $X$, and its displayed latent coordinates; no generative kernel takes as an argument the recognition law $\mathbb{Q}_{V,o,X}$, one of its marginals or parameters, or the posterior $\mathbf{\Pi}_{\theta,V,o,X}$. In the two-argument form $\mathcal{F}_V(\mathbb{Q};\theta)$ of the free energy [64], every evaluation holds the generative argument $(\theta, X)$ fixed while the recognition argument varies, which excludes the composite $\mathbb{Q} \mapsto \mathcal{F}_V(\mathbb{Q};\theta(\mathbb{Q}))$ inside one evaluation: a generative factor that reads the recognition law being scored, such as a transition built from a live peer recognition marginal, is a map from recognition laws to kernels rather than a factor of one joint. A parameter may depend on the observation or on an earlier recognition law if it is held fixed while the bound is evaluated, so slow evolution of an agent's model or prior is admissible as recognition dynamics of the model channel under one fixed joint or as accepted

M-coordinates on $\theta$ (Propositions 5.27 and 5.30 and definition 5.28), each accepted value defining a new fixed joint to which the evidence identity applies afresh.

Coupling through latents and loops of inference updates under a fixed joint remain admissible (Section 12.7), and a family of joints indexed by recognition laws designates no target (Proposition 5.11).

## 4.2 Normalized record extension and the undirected alternative

**Proposition 4.4 (Normalization by latent composition and record attachment).** The iterated integral (4.7) defines a probability kernel $X \mapsto \mathbb{P}^{\mathsf{Y}}_{\theta,V}(\cdot \mid X)$. If the record coordinates are conditionally independent given $(X, y)$, the population law is the kernel of the first line below, and the second line holds for every finite record hypergraph, cyclic or not, and every $(\theta, X)$:

$$\mathbb{P}_{\theta,V}(dy, do \mid X) = \mathbb{P}^{\mathsf{Y}}_{\theta,V}(dy \mid X) \prod_{a \in \mathfrak{A}} K_a(X, y_{\partial a}, do_a), \tag{4.10}$$

$$\mathbb{P}_{\theta,V}(\mathsf{Y} \times \mathsf{O} \mid X) = 1. \tag{4.11}$$

If $K_a(X, y_{\partial a}, do_a) = \exp\big(-E_{a,o}(y_{\partial a})\big)\nu_a(do_a)$ with a jointly measurable negative log density $E_{a,o}$ and $\nu^{\mathsf{O}} = \bigotimes_a \nu_a$, then

$$p_{\theta,V}(o, y \mid X) = p^{\mathsf{Y}}_{\theta,V}(y \mid X) \exp\Big(-\sum_{a \in \mathfrak{A}} E_{a,o}(y_{\partial a})\Big) \tag{4.12}$$

is its density. [ESTABLISHED]

*Proof.* Integrating a bounded measurable function against a probability kernel preserves boundedness and measurability in the remaining variables, so induction along the topological order defines a measurable kernel in $X$, and $f \equiv 1$ gives unit mass; finite kernel composition gives measurability of (4.10). By Tonelli's theorem, integrate the record coordinates first, in any order: each $K_a$ has unit mass in $o_a$, so the finite product disappears and leaves $\mathbb{P}^{\mathsf{Y}}_{\theta,V}(dy \mid X)$ whatever the overlaps of the scopes, since the records are conditionally independent children of one latent configuration. The same computation with the Radon–Nikodym derivatives proves (4.12). □

Each record kernel is owned once by the joint, while the coordinate-local free energies of Chapter 5 include every factor incident to a coordinate.

**Running example.** The agents of Section 2.5 carry the directed joint with $\mathrm{pa}(1) = \varnothing$, $\mathrm{pa}(2) = \mathrm{pa}(3) = \{1\}$, independent $\mathcal{N}(0,1)$ model coordinates, $k_1 \mid m_1 \sim \mathcal{N}(0, e^{-m_1} I_3)$, and children transported from agent 1, $k_2 \mid k_1, m_2 \sim \mathcal{N}\big(\rho_b(\Theta^b_{\bar{e}_1})k_1, e^{-m_2} I_3\big)$ and $k_3 \mid k_1, m_3 \sim \mathcal{N}\big(\rho_b(\Theta^b_{e_3})k_1, e^{-m_3} I_3\big)$. Two binary records are attached: $a_1$ with scope $\{1,2\}$ as in (2.40), and $a_2$ with scope $\{2,3\}$ and success probability $\exp(-\frac{1}{2}|k_2 - \rho_b(\Theta^b_{e_2})k_3|^2)$. The records close the latent tree into the triangle, and Proposition 4.4 normalizes the joint. At the conditional means of the children the residual read by $a_2$ is $\rho_b(\Theta^b_{\bar{e}_1})k_1 - \rho_b(\Theta^b_{e_2}\Theta^b_{e_3})k_1 = \rho_b(\Theta^b_{e_1})^{-1}(I_3 - \rho_b(\mathrm{Hol}^b(p)))k_1$, whose length is $\sqrt{3}$ times that of the component of $k_1$ orthogonal to the rotation axis of the holonomy (2.39), so the records see the loop holonomy through an invariant.

**Definition 4.5 (Normalized undirected model).** For nonnegative potentials on the agents, on the edges of an undirected graph with edge set $E^{\mathrm{u}}$, and on the records, put $\widetilde{p}_X$ and $Z_X$ as in the first two lines below. When $0 < Z_X < \infty$, a condition that can depend on the parameters, the undirected model is the law with the density of the third line:

$$\widetilde{p}_X(o,y) = \prod_{i\in V} \psi_i(k_i, m_i; X) \prod_{\{i,j\}\in E^{\mathrm{u}}} \psi_{ij}(y_i, y_j; X) \prod_{a\in\mathfrak{A}} \psi_a(o_a, y_{\partial a}; X), \tag{4.13}$$

$$Z_X = \int \widetilde{p}_X(o,y)\nu^{\mathsf{O}}(do)\nu^{\mathsf{Y}}(dy), \tag{4.14}$$

$$p^{\mathrm{Gibbs}}_{\theta,V}(o,y \mid X) = Z_X^{-1}\widetilde{p}_X(o,y). \tag{4.15}$$

**Proposition 4.6 (Local normalization is insufficient).** Normalized agent potentials and an everywhere positive, finite, continuous edge potential need not give a finite $Z_X$. [ESTABLISHED]

*Proof.* Take two agents with scalar belief coordinates $k_1, k_2 \in \mathbb{R}$ under Lebesgue measure, singleton model and record spaces with unit point mass, and potentials with product exponent $-\frac{1}{2}k^{\top}\mathsf{W}_{\mathrm{c}}k$,

$$\psi_i(k_i, m_i; X) = (2\pi)^{-1/2}e^{-k_i^2/2}, \qquad \psi_{12}(y_1, y_2; X) = e^{\mathrm{c}k_1k_2}, \qquad \mathrm{c} \geq 1, \tag{4.16}$$

$$\mathsf{W}_{\mathrm{c}} = \left(\begin{smallmatrix} 1 & -\mathrm{c} \\ -\mathrm{c} & 1 \end{smallmatrix}\right), \qquad \operatorname{spec}\mathsf{W}_{\mathrm{c}} = \{1-\mathrm{c}, 1+\mathrm{c}\}; \tag{4.17}$$

in orthonormal sum and difference coordinates the integral factorizes, and the factor in the direction of the eigenvalue $1 - \mathrm{c} \leq 0$ diverges. □

**Remark 4.7 (Three uses of a reciprocal energy).** Let $U$ be a measurable energy on Y, bounded below, that couples coordinates reciprocally. If $U$ enters a potential (4.13) with $Z_X = +\infty$, no law is defined and there is no evidence to bound. If it enters with $0 < Z_X < \infty$, the undirected model is a legitimate joint to which Chapter 5 applies. If $U$ is appended to the free energy of a joint $\mathbb{P}_{\theta,V}$ that does not contain it, the evidence of that joint is unaffected, and for every recognition law $Q$ in the classical domain of Hypothesis 5.3 with $\mathbb{E}_Q|U| < \infty$,

$$\mathcal{F}_V[Q; X, o] + \mathbb{E}_Q U = \widetilde{\mathcal{F}}_V[Q; X, o] - \log Z_U, \qquad Z_U = \int e^{-U(y)}\mathbb{P}_{\theta,V}(dy, do \mid X) \in (0, \infty), \tag{4.18}$$

where $\widetilde{\mathcal{F}}_V$ is the free energy of the tilted joint $Z_U^{-1}e^{-U(y)}\mathbb{P}_{\theta,V}(dy, do \mid X)$. The appended objective is therefore an evidence bound of that tilted joint only after $\log Z_U$, which depends on $(\theta, X)$, is restored. [ESTABLISHED]

*Proof.* Since $U$ is bounded below, $0 < Z_U \leq \sup e^{-U} < \infty$, and the tilted joint has density $p_{\theta,V}(o, y \mid X)e^{-U(y)}/Z_U$. With $\mathcal{F}_V[Q; X, o] = \mathbb{E}_Q[\log q_{V,o,X}(Y) - \log p_{\theta,V}(o, Y \mid X)]$, adding $\mathbb{E}_Q U$ and splitting $\log Z_U$ off the tilted log density gives (4.18); all expectations are finite under the stated integrability. □

The directed construction is used because its normalization does not depend on the parameters.

## 4.3 Covariance under passive frame re-expression

A passive re-expression $(g_i)_{i\in V}$ acts on samples through $R_i^x = \rho_x(g_i)$ of (2.9), and on the latent and joint spaces by the first two lines below, so records are gauge inert. The structural data transform to $X'$ with the frames $u_i \cdot g_i^{-1}$ and, by the vertex law (2.28), the links of the third line; a structural constraint tying parallel copies $e, e' : j \to i$ to one link is preserved, because both copies transform by the same pair (fourth line). Active invariance is invariance under the one copy $\mathcal{A}_{c_*} \cong G$ of Proposition 2.6.

$$R = \bigoplus_{i\in V} (R_i^b \oplus R_i^m) : \mathsf{Y} \longrightarrow \mathsf{Y}, \tag{4.19}$$

$$T(y, o) = (Ry, o), \tag{4.20}$$

$$\Theta_e^{x\prime} = g_i \Theta_e^x g_j^{-1}, \qquad e : j \to i,\ x \in \{b, m\}, \tag{4.21}$$

$$\Theta_e^x = \Theta_{e'}^x \quad \Longrightarrow \quad g_i \Theta_e^x g_j^{-1} = g_i \Theta_{e'}^x g_j^{-1}. \tag{4.22}$$

**Hypothesis 4.8 (Kernel covariance).** The transformed parameters $(\theta', X')$ are such that every factor obeys the first two lines below, with the empty parent tuple at roots, and every record kernel obeys the third, where $R_{\partial a}$ is the restriction of $R$ to the incident blocks; for a record with scope $\{i\}$ the third line reads as the fourth:

$$P^m_{\theta',i}\big(\cdot \mid R^m_{\mathrm{pa}(i)} m_{\mathrm{pa}(i)}, X'\big) = (R_i^m)_{\#} P^m_{\theta,i}\big(\cdot \mid m_{\mathrm{pa}(i)}, X\big), \tag{4.23}$$

$$P^k_{\theta',i}\big(\cdot \mid R^b_{\mathrm{pa}(i)} k_{\mathrm{pa}(i)}, R_i^m m_i, X'\big) = (R_i^b)_{\#} P^k_{\theta,i}\big(\cdot \mid k_{\mathrm{pa}(i)}, m_i, X\big), \tag{4.24}$$

$$K_a\big(X', R_{\partial a} y_{\partial a}, \cdot\big) = K_a\big(X, y_{\partial a}, \cdot\big), \qquad a \in \mathfrak{A}, \tag{4.25}$$

$$K_a\big(X', (R_i^b k_i, R_i^m m_i), \cdot\big) = K_a\big(X, (k_i, m_i), \cdot\big). \tag{4.26}$$

The hypothesis excludes mechanisms whose coordinate form is compatible with no re-expression of the parameters.

**Proposition 4.9 (Joint measure covariance and evidence invariance).** Under Hypothesis 4.8, with $R$ and $T$ as in (4.19)–(4.20), the first line below holds as probability measures, without a density or reference hypothesis, and the observation marginals agree, so the second holds:

$$\mathbb{P}_{\theta',V}(\cdot \mid X') = T_{\#} \mathbb{P}_{\theta,V}(\cdot \mid X), \tag{4.27}$$

$$p_{\theta',V}(o \mid X') = p_{\theta,V}(o \mid X) \quad \text{for } \nu^{\mathsf{O}}\text{-almost every } o. \tag{4.28}$$

If latent densities exist and $R_{\#}\nu^{\mathsf{Y}} \ll \nu^{\mathsf{Y}}$, then with $j_R$ of the first line below the second holds $(\nu^{\mathsf{Y}} \otimes \nu^{\mathsf{O}})$-almost everywhere, so composing a density with $R^{-1}$ represents the pushforward for every density exactly when $j_R = 1$ almost everywhere. If every represented block is an invertible linear map between Euclidean spaces of equal dimension, with Lebesgue reference (equal within

each channel, with $\dim \mathsf{K} \neq \dim \mathsf{M}$ allowed), then $j_R = |\det R|^{-1}$ and the third line holds:

$$j_R(y') := \big(d(R_{\#}\nu^{\mathsf{Y}})/d\nu^{\mathsf{Y}}\big)(y'), \tag{4.29}$$

$$p_{\theta',V}(o, y' \mid X') = p_{\theta,V}(o, R^{-1}y' \mid X) j_R(y'), \tag{4.30}$$

$$p_{\theta',V}(o, y' \mid X') = |\det R|^{-1} p_{\theta,V}(o, R^{-1}y' \mid X). \tag{4.31}$$

Finally, the root law bridges (4.5) of $\theta$ and $\theta'$ satisfy $\widetilde{\Phi}_i^{\mathrm{gen}\prime} = \widehat{\rho}_b(g_i) \circ \widetilde{\Phi}_i^{\mathrm{gen}} \circ \widehat{\rho}_m(g_i)^{-1}$, the law (2.23) of a cross-morphism representative. [ESTABLISHED]

*Proof.* Apply (4.23) and (4.24) along the topological ordering: finite kernel induction gives $\mathbb{P}^{\mathsf{Y}}_{\theta',V}(\cdot \mid X') = R_{\#}\mathbb{P}^{\mathsf{Y}}_{\theta,V}(\cdot \mid X)$. For bounded measurable $F$ on $\mathsf{Y} \times \mathsf{O}$, write $\mathbb{P}_{\theta',V}(\cdot \mid X')$ by (4.10), substitute $y' = Ry$ in the latent integral, and replace the record kernels at $(X', R_{\partial a} y_{\partial a})$ by those at $(X, y_{\partial a})$ by (4.25); this gives $\int F(y,o)\mathbb{P}_{\theta',V}(dy, do \mid X') = \int F(T(y,o))\mathbb{P}_{\theta,V}(dy, do \mid X)$, which is (4.27), and projecting to $o$ gives (4.28). For $Q = q\nu^{\mathsf{Y}}$ and measurable $B$, $R_{\#}Q(B) = \int_{R^{-1}B} q(y)\nu^{\mathsf{Y}}(dy) = \int_B q(R^{-1}y')(R_{\#}\nu^{\mathsf{Y}})(dy')$, so $R_{\#}Q = (q \circ R^{-1}) j_R \nu^{\mathsf{Y}}$; applied to the joint density this is (4.30), and linear change of variables gives $j_R = |\det R|^{-1}$. For the bridge, (4.24) at a root gives $\widetilde{\Phi}_i^{\mathrm{gen}\prime}(\widehat{\rho}_m(g_i)s) = \int P^k_{\theta',i}(\cdot \mid R_i^m m, X')s(dm) = \widehat{\rho}_b(g_i)\widetilde{\Phi}_i^{\mathrm{gen}}(s)$. □

**Corollary 4.10 (Independent channel-frame invariance).** Under Hypothesis 4.8, let $\mathbb{Q}_{V,o,X}$ be selected before the re-expression and transported as a whole, $\mathbb{Q}'_{V,o,X'} = R_{\#}\mathbb{Q}_{V,o,X}$. Then for $\mathbb{P}^{\mathsf{O}}_{\theta,V,X}$-almost every $o$ the posterior transforms as $\mathbf{\Pi}_{\theta',V,o,X'} = R_{\#}\mathbf{\Pi}_{\theta,V,o,X}$, and

$$D_{\mathrm{KL}}(\mathbb{Q}'_{V,o,X'} \| \mathbf{\Pi}_{\theta',V,o,X'}) = D_{\mathrm{KL}}(\mathbb{Q}_{V,o,X} \| \mathbf{\Pi}_{\theta,V,o,X}), \quad \mathcal{L}'^{\mathrm{ext}}_V(\mathbb{Q}'_{V,o,X'}; X', o) = \mathcal{L}^{\mathrm{ext}}_V(\mathbb{Q}_{V,o,X}; X, o) \tag{4.32}$$

in the extended reals, the singular branch included, with $\mathcal{L}^{\mathrm{ext}}_V$ as in (5.15); on the classical domain of both sides the same holds for $\mathcal{L}_V$. [ESTABLISHED]

*Proof.* By (4.27), for measurable $B \subseteq \mathsf{Y}$ and $A \subseteq \mathsf{O}$, $\mathbb{P}_{\theta',V}(B \times A \mid X') = \mathbb{P}_{\theta,V}(R^{-1}B \times A \mid X) = \int_A (R_{\#}\mathbf{\Pi}_{\theta,V,o,X})(B)\mathbb{P}^{\mathsf{O}}_{\theta,V,X}(do)$, and the observation marginals agree, so $R_{\#}\mathbf{\Pi}_{\theta,V,o,X}$ is a version of the new posterior. Lemma 3.12(e) gives equality of the divergences in $[0, +\infty]$. The evidence agrees by (4.28), and substitution in (5.15) subtracts equal quantities from equal finite numbers; the classical functional equals the extended one on its domain. □

The running example satisfies Hypothesis 4.8 with $\theta' = \theta$, since $\rho_m$ is trivial and its isotropic state covariances, link-transported means, and record lengths are compatible with (4.21).

## Physicist's summary

The generative model is one normalized joint of models, beliefs given parents, and records, which may close arbitrary cycles; reciprocal latent conditionals define no joint. No factor may read the recognition law it is scored against, and passive frame changes preserve evidence and bound.

# Chapter 5

# One Evidence Bound and Local-Collective Free Energy

Chapter 4 fixed one normalized joint law and chapter 3 one correlated recognition coupling of the whole population. This chapter asks what exact relation binds the two, and how the objective of one agent, or of a block of agents, is related to that of the population. The answer is one extended evidence bound, defined for every probability law on the latent space (definition 5.6 and corollary 5.7), and the local–global potential identity (theorem 5.23): with the outside marginal held fixed, the change of the conditional local free energy of a block equals the change of the collective free energy, so the agents' local objectives form an exact potential game (corollary 5.24). Every local object is a disintegration of the one bound through lemma 3.12. The bound is the objective of chapters 6 to 8, and chapters 9 and 12 transport it along evidence-preserving channels.

**Scope.** The chapter does not claim that a maximizer of the bound over a restricted recognition family exists, that the bound is tight for any particular restricted family, that the alternating scheme of exact coordinates converges (open problem 5.32), or that a block of agents together with its local free energy is a meta-agent. No criterion internal to one bound that selects a recognition structure is known (open problem 5.14). [NOT-CLAIMED]

## 5.1 Recognition coupling and total correlation

The population recognition law is the coupling $\mathbb{Q}_{V,o,X} \in \mathrm{Cpl}(\mathbb{Q}_{\bullet,o,X})$ selected by definition 3.15, and every restriction used later is a subfamily $\mathcal{Q}'$ of proposition 5.9. Its marginals do not determine it once two of them are non-Dirac (proposition 3.17), so they cannot replace it in an exact entropy or evidence calculation; the next two statements give the error its exact value and sign.

**Corollary 5.1 (Extended total-correlation chain identity).** Let $Y = (Y_j)_{j\in\mathcal{J}}$ take values in a finite product of standard Borel spaces, let $Q$ be a probability law with marginals $Q_j$, and let $(\varrho_j)_{j\in\mathcal{J}}$ be arbitrary probability laws. With the total correlation

$$\mathrm{TC}(Q) := D_{\mathrm{KL}}\Big(Q \,\Big\|\, \bigotimes_j Q_j\Big) \in [0,+\infty], \tag{5.1}$$

the identity

$$D_{\mathrm{KL}}\Big(Q \,\Big\|\, \bigotimes_j \varrho_j\Big) = \mathrm{TC}(Q) + \sum_j D_{\mathrm{KL}}(Q_j \| \varrho_j) \tag{5.2}$$

holds in $[0, +\infty]$, including every singular branch. The infinite branch occurs for a continuous diagonal: if $U \sim \mathrm{Unif}[0,1]$ and $Q = \mathrm{law}(U,U)$, both marginals are uniform and $\mathrm{TC}(Q) = +\infty$. [ESTABLISHED]

*Proof.* The identity is (3.21) of lemma 3.12 with its first term named. The diagonal is null for the uniform law on the square and carries all the mass of $Q$, so $Q$ is singular with respect to $Q_1 \otimes Q_2$. □

**Proposition 5.2 (Finite-density total correlation and the two signs).** Partition the latent coordinates into blocks $j$, let $\mathbb{Q}_{V,o,X}$ have density $q_{V,o,X}$ with respect to $\nu^{\mathsf{Y}}$ and block marginal densities $q_j$, and suppose the joint and block differential entropies are finite. Then the total correlation $\mathrm{TC}(\mathbb{Q}_{V,o,X}) \geq 0$ of the blocks, (5.1), satisfies

$$\mathbb{E}_{\mathbb{Q}_{V,o,X}}\big[\log q_{V,o,X}(Y)\big] = \sum_j \mathbb{E}_{\mathbb{Q}_{V,o,X}}\big[\log q_j(Y_j)\big] + \mathrm{TC}(\mathbb{Q}_{V,o,X}), \tag{5.3}$$

with equality of the two log-density expectations exactly when the blocks are independent. If in addition $\mathbb{E}_{\mathbb{Q}_{V,o,X}}|\log p_{\theta,V}(o, Y \mid X)| < \infty$, the substituted free energy and its pseudo-ELBO,

$$\widetilde{\mathcal{F}} := \mathbb{E}_{\mathbb{Q}_{V,o,X}}\Big[-\log p_{\theta,V}(o, Y \mid X) + \sum_j \log q_j(Y_j)\Big], \qquad \widetilde{\mathcal{L}} := -\widetilde{\mathcal{F}}, \tag{5.4}$$

satisfy

$$\widetilde{\mathcal{F}} = \mathcal{F}_V - \mathrm{TC}(\mathbb{Q}_{V,o,X}), \qquad \widetilde{\mathcal{L}} = \mathcal{L}_V(\mathbb{Q}_{V,o,X}; X, o) + \mathrm{TC}(\mathbb{Q}_{V,o,X}), \tag{5.5}$$

where $\mathcal{L}_V$ and $\mathcal{F}_V$ are the classical functionals of theorem 5.8. The pseudo-ELBO exceeds the ELBO by the total correlation and need not bound the log evidence. [ESTABLISHED]

*Proof.* Take $\varrho_j = \mathbb{Q}_{V,o,X,j}$ in corollary 5.1; the entropy hypotheses make every term finite and allow the logarithm of $q_{V,o,X} / \prod_j q_j$ to be split, which is (5.3). Subtracting (5.4) from $\mathcal{F}_V = \mathbb{E}_{\mathbb{Q}_{V,o,X}}[-\log p_{\theta,V} + \log q_{V,o,X}]$ leaves $\mathrm{TC}$, and negation gives the second identity. □

## 5.2 Evidence domain

**Hypothesis 5.3 (Evidence-domain hypotheses).** The extended bound of section 5.3 uses (H1)–(H2); its classical split into two expected log densities also uses (H3)–(H4).

**(H1) Types and version.** The joint $\mathbb{P}_{\theta,V}(\cdot \mid X)$ of chapter 4 has the density $p_{\theta,V}(o, y \mid X)$ relative to its reference, and $\mathbf{\Pi}_{\theta,V,o,X}$ is the version (3.16) on the regular set $\mathsf{O}^{\mathrm{reg}}_{\theta,V,X}$ of hypothesis 3.8, so every posterior and bound below holds for $\mathbb{P}^{\mathsf{O}}_{\theta,V,X}$-almost every $o$. A comparison of finitely many parameter values uses an observation in the intersection of their regular sets. No density of the recognition law is required.

**(H2) Positive finite evidence.** $0 < p_{\theta,V}(o \mid X) < \infty$; this excludes the observations at which the posterior formula would divide by zero or by infinity.

**(H3) Absolute continuity.** $\mathbb{Q}_{V,o,X} \ll \mathbf{\Pi}_{\theta,V,o,X}$. Since the posterior is dominated by $\nu^{\mathsf{Y}}$, the recognition law then has a density $q_{V,o,X} = d\mathbb{Q}_{V,o,X}/d\nu^{\mathsf{Y}}$, with a jointly measurable version for parameterized kernels under theorem A.1 and corollary A.2. Clause (H3) excludes every recognition law that charges a posterior-null set; such a law lies on the singular branch of the extended bound, where the bound equals $-\infty$ (proposition 5.4).

**(H4) Log-integrability.** The generative and recognition halves

$$\mathbb{E}_{\mathbb{Q}_{V,o,X}} \left| \log p_{\theta,V}(o, Y \mid X) \right| < \infty, \tag{5.6}$$

$$\mathbb{E}_{\mathbb{Q}_{V,o,X}} \left| \log q_{V,o,X}(Y) \right| < \infty, \tag{5.7}$$

both hold. Clause (H4) excludes recognition laws whose tails make the two classical expectations individually infinite. It is not implied by (H1)–(H3), and the extended bound is well defined without it; proposition A.7 gives a Cauchy witness whose formal difference is $-\infty$ and a countable witness whose two expectations are both $-\infty$.

**Proposition 5.4 (Posterior-null support has infinite cost).** Let $Q$ be a probability law on the latent space with $Q(S) = 1$ and $\mathbf{\Pi}_{\theta,V,o,X}(S) = 0$ for some measurable $S$. Then (H3) fails for $Q$, $D_{\mathrm{KL}}(Q \| \mathbf{\Pi}_{\theta,V,o,X}) = +\infty$, and $\mathcal{L}_V^{\mathrm{ext}}(Q; X, o) = -\infty$. In particular this holds when the posterior is dominated by Lebesgue measure on $\mathbb{R}^d$ and $Q$ is supported on a proper affine subspace. A proper closed support alone does not suffice: the uniform law on $[0, \frac{1}{2}]$ has divergence $\log 2$ from the uniform law on $[0, 1]$. [ESTABLISHED]

*Proof.* The first claims are the null-set clause of proposition 3.11 and (5.15). A proper affine subspace is Lebesgue null, and the uniform law on $[0, \frac{1}{2}]$ has density $2$ on its support relative to the uniform law on $[0, 1]$. □

## 5.3 The extended evidence bound and its record presentation

**Hypothesis 5.5 (Agent state and interaction records).** Agent $i$ has a block space $(\mathsf{Y}_i, \mathscr{Y}_i)$, and $(\mathsf{Y}, \mathscr{Y}) = \prod_{i \in V} (\mathsf{Y}_i, \mathscr{Y}_i)$. The block may contain auxiliary latent variables and law points of $(\mathcal{E}_b)_{c_*}$ and $(\mathcal{E}_m)_{c_*}$, which are then latent coordinates at the law-point or configuration level of definition 3.19, not the recognition marginal of agent $i$. The baseline is the latent population kernel $\mathbb{P}^{\mathsf{Y}}_{\theta,V} : (\mathsf{X}, \mathscr{X}) \rightsquigarrow (\mathsf{Y}, \mathscr{Y})$ of definition 4.2. Each record $a \in \mathfrak{A}$ occurs once and has a finite scope $\partial a \subseteq V$, a record space $(\mathsf{O}_a, \mathscr{O}_a)$, and a record kernel $K_a$ of the type (3.5), all of whose values $K_a(X, y_{\partial a}, \cdot)$ are dominated by one $\sigma$-finite measure $\nu_a$ with a jointly measurable negative log density $E_{a,o}$ in $(-\infty, +\infty]$,

$$K_a(X, y_{\partial a}, do_a) = e^{-E_{a,o}(X, y_{\partial a})} \nu_a(do_a), \tag{5.8}$$

where $E_{a,o}$ depends on $o$ only through $o_a$. A record shared by agents $i$ and $j$ is one element of $\mathfrak{A}$, and writing both $E_{ij}$ and $E_{ji}$ asserts two directed records. Binary success records $K_a(y_{\partial a}, \{1\}) = e^{-E(y_{\partial a})}$ with measurable $E : \mathsf{Y}_{\partial a} \to [0, +\infty]$ are admitted. The kernels and $\mathbb{P}^{\mathsf{Y}}_{\theta,V}$ read neither a recognition law, a variational parameter, nor a posterior (hypothesis 4.3), and under a common vertex gauge the coordinate presentation of each $K_a$ is equivariant under the simultaneous change of the incident frames, as in (4.25), so the kernel itself is frame independent. The hypothesis excludes record families without a common dominating measure (open problem 3.9) and every generative factor that reads recognition.

Correlated baseline states are permitted; a graph-local Markov blanket needs in addition the corresponding locality of the regular conditionals of $\mathbb{P}^{\mathsf{Y}}_{\theta,V}$. Frame independence can be realized by giving record $a$ its own frame and a link $\Theta^x_{a\leftarrow i}$ on each leg, so that the density reads $\rho_x(\Theta^x_{a\leftarrow i})y^x_i$; comparing the belief and model channels still requires the cross morphisms of section 2.3. From here on one $X \in \mathsf{X}$ is fixed and suppressed from $K_a$, $E_{a,o}$ and $\mathbb{P}^{\mathsf{Y}}_{\theta,V}$.

**Definition 5.6 (Record likelihood, slice measure, and extended bound).** Assume hypothesis 5.5 and (H1)–(H2), and fix $o \in \mathsf{O}^{\mathrm{reg}}_{\theta,V,X}$. The record likelihood and the interaction energy are

$$L_o(y) := \prod_{a\in\mathfrak{A}} e^{-E_{a,o}(y_{\partial a})} \in [0, +\infty), \tag{5.9}$$

$$H_o(y) := -\log L_o(y) = \sum_{a\in\mathfrak{A}} E_{a,o}(y_{\partial a}) \in (-\infty, +\infty]; \tag{5.10}$$

$H_o$ has one term for each record, never equals $-\infty$ because each factor density is finite, and can be negative. With $\nu^{\mathsf{O}} = \bigotimes_a \nu_a$ the joint, normalized on every finite record hypergraph, cyclic or not, by proposition 4.4, is

$$\mathbb{P}_{\theta,V}(dy, do \mid X) = \mathbb{P}^{\mathsf{Y}}_{\theta,V}(dy) \prod_{a\in\mathfrak{A}} K_a(y_{\partial a}, do_a) = \mathbb{P}^{\mathsf{Y}}_{\theta,V}(dy) L_o(y) \nu^{\mathsf{O}}(do). \tag{5.11}$$

The evidence and the posterior slice are

$$p_{\theta,V}(o \mid X) = \int_{\mathsf{Y}} L_o(y) \mathbb{P}^{\mathsf{Y}}_{\theta,V}(dy) \in (0, +\infty), \tag{5.12}$$

$$\mathbf{\Pi}_{\theta,V,o,X}(dy) = p_{\theta,V}(o \mid X)^{-1} L_o(y) \mathbb{P}^{\mathsf{Y}}_{\theta,V}(dy), \tag{5.13}$$

and the slice measure is

$$M_{\theta,V,o,X}(S) := \int_S p_{\theta,V}(o, y \mid X) \nu^{\mathsf{Y}}(dy) = \int_S L_o(y) \mathbb{P}^{\mathsf{Y}}_{\theta,V}(dy) = p_{\theta,V}(o \mid X) \mathbf{\Pi}_{\theta,V,o,X}(S). \tag{5.14}$$

For every probability law $Q$ on $(\mathsf{Y}, \mathscr{Y})$ the extended ELBO and free energy, which are well-defined extended reals because $\log p_{\theta,V}(o \mid X)$ is finite and relative entropy lies in $[0, +\infty]$, are

$$\mathcal{L}_V^{\mathrm{ext}}(Q; X, o) := \log p_{\theta,V}(o \mid X) - D_{\mathrm{KL}}\big(Q \| \mathbf{\Pi}_{\theta,V,o,X}\big) \in \big[-\infty, \log p_{\theta,V}(o \mid X)\big], \tag{5.15}$$

$$\mathcal{F}_V^{\mathrm{ext}}[Q; X, o] := -\mathcal{L}_V^{\mathrm{ext}}(Q; X, o) = -\log p_{\theta,V}(o \mid X) + D_{\mathrm{KL}}\big(Q \| \mathbf{\Pi}_{\theta,V,o,X}\big), \tag{5.16}$$

and with $\theta$, $V$ and $X$ fixed the collective free energy is written

$$\mathcal{F}_o^{\mathrm{ext}}(Q) := -\mathcal{L}_o^{\mathrm{ext}}(Q) := \mathcal{F}_V^{\mathrm{ext}}[Q; X, o]. \tag{5.17}$$

No domination, density, or mean-field condition on $Q$ enters.

**Corollary 5.7 (Extended gap identity and energy form).** For every probability law $Q$ on the latent space,

$$\log p_{\theta,V}(o \mid X) - \mathcal{L}_V^{\mathrm{ext}}(Q; X, o) = \mathcal{F}_V^{\mathrm{ext}}[Q; X, o] + \log p_{\theta,V}(o \mid X) = D_{\mathrm{KL}}\big(Q \| \mathbf{\Pi}_{\theta,V,o,X}\big), \tag{5.18}$$

so $\mathcal{L}_V^{\mathrm{ext}} \le \log p_{\theta,V}(o \mid X)$ with equality exactly when $Q = \mathbf{\Pi}_{\theta,V,o,X}$ as measures. If moreover $Q \ll \mathbb{P}^{\mathsf{Y}}_{\theta,V}$, $D_{\mathrm{KL}}(Q \| \mathbb{P}^{\mathsf{Y}}_{\theta,V}) < \infty$, and the negative part of $H_o$ is $Q$-integrable, then

$$\mathcal{F}_o^{\mathrm{ext}}(Q) = D_{\mathrm{KL}}\big(Q \| \mathbb{P}^{\mathsf{Y}}_{\theta,V}\big) + \mathbb{E}_Q H_o \quad \text{in } (-\infty, +\infty]. \tag{5.19}$$

No formula $\log p_{\theta,V} = \mathcal{L}_V^{\mathrm{ext}} + D_{\mathrm{KL}}$ is asserted on the singular branch, where it would read $-\infty + \infty$. [ESTABLISHED]

*Proof.* The gap identity adds or subtracts the finite number $\log p_{\theta,V}(o \mid X)$ in (5.15)–(5.16), and the equality case is the zero case of relative entropy. For the energy form let $r' = dQ/d\mathbb{P}^{\mathsf{Y}}_{\theta,V}$. If $Q \ll M_{\theta,V,o,X}$, then $dQ/dM_{\theta,V,o,X} = r' e^{H_o}$ $Q$-almost surely, and proposition A.6 gives $\mathcal{F}_o^{\mathrm{ext}}(Q) = \int (\log r' + H_o) dQ$; the first term integrates to the finite $D_{\mathrm{KL}}(Q \| \mathbb{P}^{\mathsf{Y}}_{\theta,V})$ and the second is quasi-integrable, so the integral splits. If $Q \not\ll M_{\theta,V,o,X}$, some set null for $M_{\theta,V,o,X}$ has positive $Q$-mass; since $M_{\theta,V,o,X}$ and $\mathbb{P}^{\mathsf{Y}}_{\theta,V}$ agree in null sets on $\{L_o > 0\}$ and $Q \ll \mathbb{P}^{\mathsf{Y}}_{\theta,V}$, that mass lies in $\{H_o = +\infty\}$, and both sides equal $+\infty$. □

**Theorem 5.8 (Classical finite evidence identity).** Under hypothesis 5.3, the extended bound has the separately integrable representation

$$\mathcal{L}_V(\mathbb{Q}_{V,o,X}; X, o) := \mathbb{E}_{\mathbb{Q}_{V,o,X}}\big[\log p_{\theta,V}(o, Y \mid X) - \log q_{V,o,X}(Y)\big] = \mathcal{L}_V^{\mathrm{ext}}(\mathbb{Q}_{V,o,X}; X, o), \tag{5.20}$$

$$\mathcal{F}_V[\mathbb{Q}_{V,o,X}; X, o] := -\mathcal{L}_V(\mathbb{Q}_{V,o,X}; X, o), \tag{5.21}$$

and for $\mathbb{P}^{\mathsf{O}}_{\theta,V,X}$-almost every observation

$$\log p_{\theta,V}(o \mid X) = \mathcal{L}_V(\mathbb{Q}_{V,o,X}; X, o) + D_{\mathrm{KL}}\big(\mathbb{Q}_{V,o,X} \| \mathbf{\Pi}_{\theta,V,o,X}\big). \tag{5.22}$$

Consequently $\mathcal{L}_V \le \log p_{\theta,V}(o \mid X)$, with equality if and only if $\mathbb{Q}_{V,o,X} = \mathbf{\Pi}_{\theta,V,o,X}$ as probability measures on $(\mathsf{Y}, \mathscr{Y})$. Equality of all coordinate marginals is not sufficient. [ESTABLISHED]

*Proof.* The posterior and the slice measure have the same null sets, so (H3) gives $\mathbb{Q}_{V,o,X} \ll M_{\theta,V,o,X}$, and since the two measures have the densities $q_{V,o,X}$ and $p_{\theta,V}(o,\cdot \mid X)$ relative to $\nu^{\mathsf{Y}}$, $d\mathbb{Q}_{V,o,X}/dM_{\theta,V,o,X} = q_{V,o,X}/p_{\theta,V}(o,\cdot \mid X)$ $\mathbb{Q}_{V,o,X}$-almost surely. Proposition A.6 gives $\mathcal{L}_V^{\mathrm{ext}} = -\int \log(q_{V,o,X}/p_{\theta,V}(o,\cdot \mid X)) d\mathbb{Q}_{V,o,X}$, and (H4) splits the integral, which is (5.20). The bound is then finite, so (5.15) rearranges to (5.22), and the equality case is that of corollary 5.7. Distinct measures with equal coordinate marginals are supplied by proposition 3.17. □

## 5.4 Restriction principle

**Proposition 5.9 (Restriction principle).** Fix $(\theta, X)$ and $o$ as in definition 5.6, and write $P^\star = \mathbf{\Pi}_{\theta,V,o,X}$, so that for every law $Q$

$$\mathcal{L}_V^{\mathrm{ext}}(Q; X, o) = \log p_{\theta,V}(o \mid X) - D_{\mathrm{KL}}(Q \| P^\star). \tag{5.23}$$

For arbitrary sets of probability laws $\mathcal{Q}' \subseteq \mathcal{Q}$,

$$\sup_{Q \in \mathcal{Q}'} \mathcal{L}_V^{\mathrm{ext}}(Q; X, o) = \log p_{\theta,V}(o \mid X) - \inf_{Q \in \mathcal{Q}'} D_{\mathrm{KL}}(Q \| P^\star), \tag{5.24}$$

the supremum over $\mathcal{Q}'$ is at most the supremum over $\mathcal{Q}$, and when both infima are finite

$$\sup_{\mathcal{Q}} \mathcal{L}_V^{\mathrm{ext}} - \sup_{\mathcal{Q}'} \mathcal{L}_V^{\mathrm{ext}} = \inf_{\mathcal{Q}'} D_{\mathrm{KL}}(Q \| P^\star) - \inf_{\mathcal{Q}} D_{\mathrm{KL}}(Q \| P^\star) \ge 0. \tag{5.25}$$

No attainment of the suprema or infima is required. [ESTABLISHED]

*Proof.* Identity (5.23) is (5.15). Taking suprema turns a finite constant minus a nonnegative extended divergence into that constant minus its infimum, including the value $-\infty$. Inclusion gives monotonicity, and subtraction is legitimate only when both infima are finite. □

The evidence in (5.23) is constant because the generative kernel does not read $Q$ (hypothesis 4.3); the next two propositions show what fails when a model moves with $Q$.

**Remark 5.10 (Orientation, separation, and attainment).** The infimum in (5.24) is a reverse information projection, since $Q$ is the first argument of the divergence. Excluding $P^\star$ from a family does not make the optimum positive or attained: for $P^\star = \mathcal{N}(0,1)$ the family $\{\mathcal{N}(\mu,1) : \mu > 0\}$ has gaps $\mu^2/2$ with unattained infimum zero. A positive optimized gap is the condition $\inf_{Q \in \mathcal{Q}} D_{\mathrm{KL}}(Q \| P^\star) > 0$. A family constrained to product laws need not be convex, because mixtures of product laws need not be product laws. Strict convexity of relative entropy on the ambient law space therefore does not establish uniqueness on that family; on a convex feasible family it does give at most one finite minimizer. [ESTABLISHED]

**Proposition 5.11 (No distinguished target).** Let $\mathcal{Q}$ be a nonempty family of recognition laws, and replace one factor of (4.6) by a nonconstant $Q$-indexed family, giving a normalized joint $P_{\theta,Q}$ for each $Q \in \mathcal{Q}$. Assume at the observation that $0 < p_{\theta,Q}(o \mid X) < \infty$, that $Q \ll P_{\theta,Q}(\cdot \mid o, X)$, and that the log densities defining the ELBO are integrable. Then for each fixed $Q$

$$\log p_{\theta,Q}(o \mid X) = \mathcal{L}(Q; X) + D_{\mathrm{KL}}\big(Q \| P_{\theta,Q}(\cdot \mid o, X)\big), \tag{5.26}$$

and with $\mathsf{e}(Q) = \log p_{\theta,Q}(o \mid X)$ two members satisfy

$$\mathcal{L}(Q'; X) - \mathcal{L}(Q; X) = \mathsf{e}(Q') - \mathsf{e}(Q) - \big[D_{\mathrm{KL}}\big(Q' \| P_{\theta,Q'}(\cdot \mid o, X)\big) - D_{\mathrm{KL}}\big(Q \| P_{\theta,Q}(\cdot \mid o, X)\big)\big]. \tag{5.27}$$

On a level set of $\mathsf{e}$ the change of the bound is the negative change of the divergence; off it the two changes differ by $\mathsf{e}(Q') - \mathsf{e}(Q)$. The family selects no joint as the model. [ESTABLISHED]

*Proof.* Apply theorem 5.8 to each fixed pair $(Q, P_{\theta,Q})$ and subtract. No member is distinguished because no selection rule was declared. □

**Proposition 5.12 (The stronger consequences fail).** Nonconstant dependence of $P_{\theta,Q}$ on $Q$ does not imply that $\mathsf{e}$ varies, that no ascent improves a fixed quantity, or that no single member's evidence bounds the family of bounds. Explicitly, for one binary latent $y$, one binary observation, and $Q_{\mathrm{p}} = \mathrm{Ber}(\mathrm{p})$ with $\mathrm{p} \in (0, 1)$, the joint

$$P_{\theta,Q_{\mathrm{p}}}(o, y \mid X) = \tfrac{1}{2}\,\mathrm{Ber}\big(\tfrac{1}{2}\mathrm{p} + \tfrac{1}{4}\big)(y) \tag{5.28}$$

is normalized and injective in $\mathrm{p}$, has $\mathsf{e}(Q_{\mathrm{p}}) = \log \frac{1}{2}$ for every $\mathrm{p}$, and has gaps

$$\tfrac{9}{10}\log\tfrac{9}{7} + \tfrac{1}{10}\log\tfrac{1}{3} = 0.1163218\ldots, \qquad \tfrac{7}{10}\log\tfrac{7}{6} + \tfrac{3}{10}\log\tfrac{3}{4} = 0.0216009\ldots, \qquad 0 \tag{5.29}$$

at $\mathrm{p} = \frac{9}{10}, \frac{7}{10}, \frac{1}{2}$. [ESTABLISHED]

*Proof.* The observation marginal is $\frac{1}{2}$ for every $\mathrm{p}$, and the posterior is $\mathrm{Ber}(\frac{1}{2}\mathrm{p} + \frac{1}{4})$, so the gaps are the Bernoulli relative entropies displayed. The bound increases to the fixed evidence along the three values, and the member at $\mathrm{p} = \frac{1}{2}$ has evidence equal to the supremum of the bounds. □

**Proposition 5.13 (ELBO values do not order unrelated joints).** Let $P_\theta$ and $P_{\theta'}$ be normalized joints over the same observation, possibly with different latent inventories, with recognition families $\mathcal{Q}$ and $\mathcal{Q}'$ on their latent spaces, and let $o$ lie in the regular set of each. Suppose both infima of the posterior relative entropy over the respective recognition families are finite. If no evidence-preserving channel relates the two joints, their optimized bounds satisfy

$$\mathcal{L}^\star - \mathcal{L}'^\star = \big[\log p_\theta(o \mid X) - \log p_{\theta'}(o \mid X)\big] - \Big[\inf_{\mathcal{Q}} D_{\mathrm{KL}} - \inf_{\mathcal{Q}'} D_{\mathrm{KL}}\Big], \tag{5.30}$$

which combines an evidence difference, a model-comparison quantity, with a tightness difference, a recognition-quality quantity. The sign of the left side orders neither the models nor the recognition families: with binary $O$ and $Y$, constants $\mathrm{c}_1 \le 0 \le \mathrm{c}_2$, and the normalized joints $P(o, 1) = e^{\mathrm{c}_1 - \mathrm{c}_2}$,

$P(o,0) = e^{c_1}(1-e^{-c_2})$, $P(\bar{o},0) = 1-e^{c_1}$, $P(\bar{o},1) = 0$ under the singleton family $\{\delta_1\}$, the pairs $(c_1,c_2) = (-1,0),(0,2)$ give bounds $-1 > -2$ with the reverse evidence order, and the pairs $(0,1),(-2,0)$ give bounds $-1 > -2$ with the reverse tightness order. [ESTABLISHED]

*Proof.* Apply (5.24) to each model and subtract. In the witness, $\log P(O=o) = c_1$ and $P(Y = 1 \mid o) = e^{-c_2}$, so the minimized gap of $\{\delta_1\}$ is $c_2$ and the optimized bound is $c_1 - c_2$. An unused deterministic coordinate appended to either model, with its point mass appended to the recognition law, makes the latent inventories literally different and changes no evidence, divergence, or bound. □

**Open Problem 5.14 (An intrinsic criterion).** No criterion internal to one evidence bound at a fixed model is known that selects a nondegenerate recognition structure. Along nested recognition families the optimized bound is monotone by inclusion (proposition 5.9), so minimizing the gap selects the largest family; an arbitrary replacement joint supplies no common evidence (proposition 5.13); and evidence-preserving channels, which do order bounds across latent spaces, do so without selecting a nondegenerate partition (section 9.3). The problem is settled positively by a functional on admissible channels or partitions, expressible relative to one fixed joint and not monotone toward an endpoint, whose optimizer is nondegenerate, and negatively by a proof of endpoint degeneracy for a precisely stated class of such functionals.

## 5.5 The factorwise decomposition

**Proposition 5.15 (Exact factorwise decomposition).** Assume (H1)–(H3) of hypothesis 5.3 and factorwise log-integrability: the absolute logarithm of every component density of (4.8) and every record energy $E_{a,o}$ is $\mathbb{Q}_{V,o,X}$-integrable. Then (i) the generative half (5.6) of (H4) holds, and (H4) holds exactly when the recognition half (5.7) does; (ii) under (H4) the free energy has exactly one decomposition into factor terms,

$$\begin{aligned}\mathcal{F}_V[\mathbb{Q}_{V,o,X}; X, o] = &\sum_{i\in V} \mathbb{E}_{\mathbb{Q}_{V,o,X}}\left[-\log p^m_{\theta,i}(m_i \mid m_{\mathrm{pa}(i)}, X)\right] + \sum_{i\in V} \mathbb{E}_{\mathbb{Q}_{V,o,X}}\left[-\log p^k_{\theta,i}(k_i \mid k_{\mathrm{pa}(i)}, m_i, X)\right] \\ &+ \sum_{a\in\mathfrak{A}} \mathbb{E}_{\mathbb{Q}_{V,o,X}}\left[E_{a,o}(Y_{\partial a})\right] + \mathbb{E}_{\mathbb{Q}_{V,o,X}}\left[\log q_{V,o,X}(Y)\right], \end{aligned} \tag{5.31}$$

with $\mathrm{pa}(i) = \varnothing$ at the roots, $2N + |\mathfrak{A}|$ generative factors, and a final term that is one full-joint recognition log density, not a sum of marginal entropies; (iii) factorwise integrability does not imply the recognition half, and the generative half does not imply factorwise integrability. [ESTABLISHED]

*Proof.* By (4.8) and hypothesis 5.5,

$$\log p_{\theta,V}(o, y \mid X) = \sum_{i\in V}\left[\log p^m_{\theta,i}(m_i \mid m_{\mathrm{pa}(i)}, X) + \log p^k_{\theta,i}(k_i \mid k_{\mathrm{pa}(i)}, m_i, X)\right] - \sum_{a\in\mathfrak{A}} E_{a,o}(y_{\partial a}), \tag{5.32}$$

a finite sum, so the triangle inequality gives (i). Under (H4), taking $\mathbb{Q}_{V,o,X}$-expectations of the negative of (5.32) term by term and adding $\mathbb{E}_{\mathbb{Q}_{V,o,X}}[\log q_{V,o,X}]$ gives (5.31) from (5.20). For (iii),

appendix A.2 gives a model with bounded factor logarithms whose recognition density has an infinite mean logarithm, and a root law $\mathrm{c}/(m(\log m)^2)$ whose state bridge cancels its logarithm, so that the complete log joint is integrable while one factor is not. □

## 5.6 Observations as interaction records

**Definition 5.16 (Operational environment node).** An operational environment node is a standard Borel state space with a normalized state law and a measurable message kernel.

**Theorem 5.17 (Observation–interaction kernel equivalence).** Every normalized observation kernel $K(do \mid y)$ from a measurable space to a standard Borel space is the message of an operational environment node: there are $U \sim \mathrm{Unif}[0,1]$ and a measurable map $F$ such that

$$O = F(Y,U), \qquad \Pr\big(F(y,U) \in A\big) = K(A \mid y). \tag{5.33}$$

Conversely, marginalizing any environment state and message policy produces a normalized observation kernel. [ESTABLISHED]

*Proof.* The forward direction is the randomization lemma for kernels [49]. By Kuratowski's theorem a standard Borel space admits a Borel isomorphism $\phi$ onto a Borel subset of $\mathbb{R}$. The function $G(y,x) = K(\phi^{-1}(-\infty,x] \mid y)$ is measurable in $y$ and a distribution function in $x$, so the quantile $F_0(y,u) = \inf\{x : G(y,x) \geq u\}$ on $(0,1)$ satisfies $\{F_0 \leq x\} = \{(y,u) : u \leq G(y,x)\}$; hence $F_0$ is jointly measurable and $\Pr(F_0(y,U) \leq x) = G(y,x)$. The value $F_0(y,U)$ lies in the image of $\phi$ almost surely; setting $F = \phi^{-1} \circ F_0$ there and a fixed point elsewhere gives (5.33). Conversely, a state law $\varrho_{\mathrm{env}}(du)$ and a message policy $\kappa_{\mathrm{env}}(y,u,A)$ give $K(A \mid y) = \int \kappa_{\mathrm{env}}(y,u,A)\varrho_{\mathrm{env}}(du)$, which is measurable in $y$, countably additive in $A$, and normalized. □

**Proposition 5.18 (Conditioning information and retained statistics).** Let $(Y,O)$ have joint law $P$ on standard Borel spaces with marginals $P^Y$ and $P^O$, let $I(Y;O) = D_{\mathrm{KL}}(P \| P^Y \otimes P^O)$, and for a measurable statistic $S = T(O)$ let $I(Y;O \mid S) = \int D_{\mathrm{KL}}(P_{Y,O|s} \| P_{Y|s} \otimes P_{O|s}) P^S(ds)$. The information carried by the record is

$$I(Y;O) = \int D_{\mathrm{KL}}\big(P(dY \mid o) \| P^Y\big) P^O(do) \in [0,+\infty]. \tag{5.34}$$

Moreover

$$I(Y;O) = I(Y;S) + I(Y;O \mid S) \quad \text{in } [0,+\infty], \tag{5.35}$$

so replacing $O$ by $S$ leaves the posterior unchanged exactly when $I(Y;O \mid S) = 0$, up to the choice of conditional versions. Deleting the record makes $S$ constant and returns the prior; for $Y \sim \mathrm{Ber}(1/2)$ and $O = Y$, $I(Y;O) = \log 2$ and $P(dY \mid O = o) = \delta_o$. A kernel that does not depend on $y$ carries zero information. [ESTABLISHED]

*Proof.* Disintegrating $D_{\mathrm{KL}}(P\|P^Y \otimes P^O)$ over $O$ with lemma 3.12, whose $O$-marginals agree, gives (5.34). Since $o \mapsto (T(o), o)$ is a bimeasurable bijection onto its graph, relative entropy is unchanged by passing to $(Y, S, O)$ (3.24). Disintegrating $D_{\mathrm{KL}}(P_{Y,S,O}\|P^Y \otimes P_{S,O})$ over $(Y, S)$ gives $I(Y;S) + \mathbb{E}_{P_{Y,S}} D_{\mathrm{KL}}(P_{O|Y,S}\|P_{O|S})$, because under the reference $O$ is independent of $Y$ given $S$ with law $P_{O|S}$; for each $s$, disintegrating $D_{\mathrm{KL}}(P_{Y,O|s}\|P_{Y|s} \otimes P_{O|s})$ over $Y$ and integrating in $s$ identifies the second term with $I(Y;O \mid S)$, which proves (5.35); the discrete case is Cover and Thomas [21, Sec. 2.5]. Posterior sufficiency is the vanishing of the conditional term, the Bernoulli values are direct, and a $y$-independent kernel makes $P = P^Y \otimes P^O$. □

**Remark 5.19 (Environment nodes in the variational problem).** Theorem 5.17 is exact for the joint law of $(Y, O)$, and it keeps the record variable $O$ and hence the conditioning $\sigma$-algebra $\sigma(O)$, but it is not exact for a variational problem over the enlarged state $(Y, U)$. For $P^O$-almost every $o$ the posterior of $(Y, U)$ is concentrated on the level set $\{F(y, u) = o\}$. For $Y \sim \mathcal{N}(0, 1)$ and $F(y, u) = y + F_{\mathcal{N}}^{-1}(u)$, with $F_{\mathcal{N}}$ the standard normal distribution function, which realizes $K(do \mid y) = \mathcal{N}(y, 1)$, that level set is the Lebesgue-null curve $\{u = F_{\mathcal{N}}(o - y)\}$, so every recognition law on $(Y, U)$ with a Lebesgue density has relative entropy $+\infty$ by proposition 5.4, whereas the fiber-restricted law $q(dy)\delta_{F_{\mathcal{N}}(o-y)}(du)$, the image of $q$ under the injective map $y \mapsto (y, F_{\mathcal{N}}(o - y))$ that also carries the posterior of $Y$ to that of $(Y, U)$, has relative entropy $D_{\mathrm{KL}}(q\|P(dY \mid o))$. A finite variational problem therefore integrates $U$ out or restricts recognition to that fiber. [ESTABLISHED]

## 5.7 Local conditional free energy and the potential identity

Fix a nonempty block $B \subseteq V$, one agent or any aggregate, and write $y = (y_B, y_{-B})$; the probability model does not change with $B$.

**Lemma 5.20 (Block disintegration of the posterior).** Let $\mathfrak{A}_B = \{a \in \mathfrak{A} : \partial a \cap B \neq \varnothing\}$ and define the incident and outside likelihoods and energies

$$L_{B,o}(y_B; y_{-B}) = \prod_{a \in \mathfrak{A}_B} e^{-E_{a,o}(y_{\partial a})}, \qquad L_{\bar{B},o}(y_{-B}) = \prod_{a \notin \mathfrak{A}_B} e^{-E_{a,o}(y_{\partial a})}, \tag{5.36}$$

$$H_{B,o}(y_B; y_{-B}) = -\log L_{B,o} = \sum_{a \in \mathfrak{A}_B} E_{a,o}(y_{\partial a}), \qquad H_{\bar{B},o}(y_{-B}) = -\log L_{\bar{B},o}(y_{-B}), \tag{5.37}$$

so that $L_o = L_{B,o} L_{\bar{B},o}$ and $H_o = H_{B,o} + H_{\bar{B},o}$ in $(-\infty, +\infty]$. Fix a measurable version of the baseline disintegration

$$\mathbb{P}^{\mathsf{Y}}_{\theta,V}(dy_B, dy_{-B}) = \mathbb{P}^{\mathsf{Y}}_{\theta,V,-B}(dy_{-B}) \mathbb{P}^{\mathsf{Y}}_{\theta,V,B}(dy_B \mid y_{-B}), \tag{5.38}$$

and define the local evidence and the likelihood of the outside state,

$$Z_B(y_{-B}) = \int L_{B,o}(y_B; y_{-B}) \mathbb{P}^{\mathsf{Y}}_{\theta,V,B}(dy_B \mid y_{-B}), \tag{5.39}$$

$$\widetilde{L}_{B,o}(y_{-B}) = \int L_o(y_B, y_{-B}) \mathbb{P}^{\mathsf{Y}}_{\theta,V,B}(dy_B \mid y_{-B}). \tag{5.40}$$

Then (i) the evidence and the outside posterior marginal are

$$p_{\theta,V}(o \mid X) = \int \widetilde{L}_{B,o} d\mathbb{P}^{\mathsf{Y}}_{\theta,V,-B}, \qquad \mathbf{\Pi}_{\theta,V,o,X,-B}(dy_{-B}) = p_{\theta,V}(o \mid X)^{-1} \widetilde{L}_{B,o}(y_{-B}) \mathbb{P}^{\mathsf{Y}}_{\theta,V,-B}(dy_{-B}); \tag{5.41}$$

(ii) the set $\Omega_{B,o}$ of outside states at which $0 < \widetilde{L}_{B,o} < \infty$, $0 < L_{\bar{B},o} < \infty$, $0 < Z_B < \infty$, and $\widetilde{L}_{B,o} = L_{\bar{B},o} Z_B$ has $\mathbf{\Pi}_{\theta,V,o,X,-B}$-measure one, though it need not be $\mathbb{P}^{\mathsf{Y}}_{\theta,V,-B}$-full; (iii) the kernel

$$\mathbf{\Pi}_{\theta,V,o,X,B}(dy_B \mid y_{-B}) = Z_B(y_{-B})^{-1} L_{B,o}(y_B; y_{-B}) \mathbb{P}^{\mathsf{Y}}_{\theta,V,B}(dy_B \mid y_{-B}) \qquad (y_{-B} \in \Omega_{B,o}), \tag{5.42}$$

extended by any measurable probability kernel off $\Omega_{B,o}$, satisfies

$$\mathbf{\Pi}_{\theta,V,o,X}(dy_B, dy_{-B}) = \mathbf{\Pi}_{\theta,V,o,X,-B}(dy_{-B}) \mathbf{\Pi}_{\theta,V,o,X,B}(dy_B \mid y_{-B}). \tag{5.43}$$

The posterior determines this kernel, and every pointwise local quantity built from it, only $\mathbf{\Pi}_{\theta,V,o,X,-B}$-almost everywhere; two baseline versions give local posteriors that agree there. [ESTABLISHED]

*Proof.* Tonelli's theorem applied to (5.12)–(5.13) with (5.38) gives (i); defining $\widetilde{L}_{B,o}$ directly avoids the product $L_{\bar{B},o} Z_B$, which can be $0 \cdot \infty$ off $\Omega_{B,o}$. For (ii), finite evidence makes $\widetilde{L}_{B,o} < \infty$ off a $\mathbb{P}^{\mathsf{Y}}_{\theta,V,-B}$-null set, and by (i) the outside posterior charges only $\{\widetilde{L}_{B,o} > 0\}$. There $L_{\bar{B},o}$ is finite because each factor density is, positive because $L_{\bar{B},o} = 0$ would force $\widetilde{L}_{B,o} = 0$, and independent of $y_B$, so it leaves the conditional integral and $Z_B = \widetilde{L}_{B,o} / L_{\bar{B},o} \in (0, \infty)$. For (iii), a bounded measurable $f$ has $\int f d\mathbf{\Pi}_{\theta,V,o,X} = p_{\theta,V}(o \mid X)^{-1} \int L_{\bar{B},o} \int f L_{B,o} d\mathbb{P}^{\mathsf{Y}}_{\theta,V,B} d\mathbb{P}^{\mathsf{Y}}_{\theta,V,-B}$, and on $\Omega_{B,o}$ the inner integral is $Z_B$ times the integral of $f$ against (5.42); (i) and (ii) turn the outer weight into $\mathbf{\Pi}_{\theta,V,o,X,-B}$. Almost-everywhere uniqueness is uniqueness of disintegrations; two baseline versions agree $\mathbb{P}^{\mathsf{Y}}_{\theta,V,-B}$-almost everywhere, hence $\mathbf{\Pi}_{\theta,V,o,X,-B}$-almost everywhere by (i). □

**Definition 5.21 (Local conditional free energy).** For $y_{-B} \in \Omega_{B,o}$ and every probability law $r_B(\cdot \mid y_{-B})$ on $\mathsf{Y}_B$, the local conditional ELBO and free energy of the block are

$$\mathcal{F}^{\mathrm{ext}}_{B,o}(r_B; y_{-B}) := -\mathcal{L}^{\mathrm{ext}}_{B,o}(r_B; y_{-B}) := -\log Z_B(y_{-B}) + D_{\mathrm{KL}}\big(r_B(\cdot \mid y_{-B}) \| \mathbf{\Pi}_{\theta,V,o,X,B}(\cdot \mid y_{-B})\big). \tag{5.44}$$

For $B = \{i\}$ this is the local free energy of agent $i$. It contains every record incident to the block and no other.

**Corollary 5.22 (Local Gibbs bound and the exact block E-coordinate).** (i) $\mathcal{L}^{\mathrm{ext}}_{B,o}(r_B; y_{-B}) \le \log Z_B(y_{-B})$, with equality exactly at $r_B = \mathbf{\Pi}_{\theta,V,o,X,B}(\cdot \mid y_{-B})$. If $r_B \ll \mathbb{P}^{\mathsf{Y}}_{\theta,V,B}(\cdot \mid y_{-B})$ with finite

relative entropy and the negative part of $H_{B,o}(\cdot; y_{-B})$ is $r_B$-integrable,

$$\mathcal{F}^{\mathrm{ext}}_{B,o}(r_B; y_{-B}) = D_{\mathrm{KL}}\big(r_B(\cdot \mid y_{-B}) \| \mathbb{P}^{\mathsf{Y}}_{\theta,V,B}(\cdot \mid y_{-B})\big) + \mathbb{E}_{r_B(\cdot|y_{-B})} H_{B,o}(Y_B; y_{-B}). \tag{5.45}$$

(ii) If $Q(dy_B, dy_{-B}) = Q_{-B}(dy_{-B}) r_B(dy_B \mid y_{-B})$ and $Q_{-B} \ll \mathbf{\Pi}_{\theta,V,o,X,-B}$, then in $[0, +\infty]$, with the baseline conditional of (5.38) in the second identity,

$$D_{\mathrm{KL}}\big(Q \| \mathbf{\Pi}_{\theta,V,o,X}\big) = D_{\mathrm{KL}}\big(Q_{-B} \| \mathbf{\Pi}_{\theta,V,o,X,-B}\big) + \mathbb{E}_{Q_{-B}} D_{\mathrm{KL}}\big(r_B(\cdot \mid Y_{-B}) \| \mathbf{\Pi}_{\theta,V,o,X,B}(\cdot \mid Y_{-B})\big), \tag{5.46}$$

$$D_{\mathrm{KL}}\big(Q \| \mathbb{P}^{\mathsf{Y}}_{\theta,V}\big) = D_{\mathrm{KL}}\big(Q_{-B} \| \mathbb{P}^{\mathsf{Y}}_{\theta,V,-B}\big) + \mathbb{E}_{Q_{-B}} D_{\mathrm{KL}}\big(r_B(\cdot \mid Y_{-B}) \| \mathbb{P}^{\mathsf{Y}}_{\theta,V,B}(\cdot \mid Y_{-B})\big). \tag{5.47}$$

(iii) Let $\mathcal{B}$ be any block of latent coordinates with complement $-\mathcal{B}$, let $P_-$ be the $-\mathcal{B}$ marginal of $P = \mathbf{\Pi}_{\theta,V,o,X}$, let $R$ be any law on the $-\mathcal{B}$ coordinates, and let $\kappa_{\mathcal{B}}(dY_{\mathcal{B}} \mid Y_{-\mathcal{B}})$ be a Markov kernel that agrees $P_-$-almost everywhere with a regular conditional law of $P$ (for an agent block, (5.42)). Among all laws with $-\mathcal{B}$ marginal $R$, the divergence from $P$ has minimum $D_{\mathrm{KL}}(R \| P_-) \in [0, +\infty]$, attained by

$$\mathbb{Q}^{\mathrm{new}}_{V,o,X}(dY_{\mathcal{B}}, dY_{-\mathcal{B}}) = \kappa_{\mathcal{B}}(dY_{\mathcal{B}} \mid Y_{-\mathcal{B}}) R(dY_{-\mathcal{B}}), \tag{5.48}$$

which therefore maximizes $\mathcal{L}^{\mathrm{ext}}_V$ on that class. The minimizer is unique when $D_{\mathrm{KL}}(R \| P_-) < \infty$; otherwise every law in the class has extended bound $-\infty$. If the minimum is finite and

$$\mathbb{E}_{\mathbb{Q}^{\mathrm{new}}_{V,o,X}}\big[|\log p_{\theta,V}(o, Y \mid X)| + |\log q^{\mathrm{new}}_{V,o,X}(Y)|\big] < \infty, \tag{5.49}$$

the new law satisfies (H3)–(H4) and maximizes the classical $\mathcal{L}_V$ among the laws with marginal $R$ that satisfy hypothesis 5.3. Finite complement relative entropy does not imply (5.49).

(iv) If $\mathcal{F}^{\mathrm{ext}}_o(Q) < \infty$ for $Q = Q_{-B} r_B$, then for $\mathrm{w} \in [0, 1]$ the damped kernel $r^{\mathrm{w}}_B = (1 - \mathrm{w}) r_B + \mathrm{w} \mathbf{\Pi}_{\theta,V,o,X,B}$ satisfies

$$\mathcal{F}^{\mathrm{ext}}_o(Q_{-B} r^{\mathrm{w}}_B) \le (1 - \mathrm{w}) \mathcal{F}^{\mathrm{ext}}_o(Q) + \mathrm{w} \mathcal{F}^{\mathrm{ext}}_o(Q_{-B} \mathbf{\Pi}_{\theta,V,o,X,B}) \le \mathcal{F}^{\mathrm{ext}}_o(Q), \tag{5.50}$$

so replacing $r_B$ by the local posterior with $Q_{-B}$ fixed does not increase $\mathcal{F}^{\mathrm{ext}}_o$. [ESTABLISHED]

*Proof.* For (i), the conditional model with reference $\mathbb{P}^{\mathsf{Y}}_{\theta,V,B}(\cdot \mid y_{-B})$ and likelihood $L_{B,o}(\cdot; y_{-B})$ has evidence $Z_B(y_{-B}) \in (0, \infty)$ and posterior (5.42), and corollary 5.7 applies to it. Part (ii) is lemma 3.12 with the disintegrations (5.43) and (5.38). For (iii), the fixed-marginal clause of lemma 3.12, with the $-\mathcal{B}$ coordinates first, gives the minimum $D_{\mathrm{KL}}(R \| P_-)$, the minimizer (5.48), and uniqueness when the minimum is finite; then $R \ll P_-$, the completion of $\kappa_{\mathcal{B}}$ off a $P_-$-null set is irrelevant, and $d\mathbb{Q}^{\mathrm{new}}_{V,o,X}/dP = (dR/dP_-)(Y_{-\mathcal{B}})$. When the minimum is infinite, the data-processing clause (3.22) applied to the projection makes every divergence in the class infinite. Definition 5.6 converts divergence into the extended bound, and theorem 5.8 gives the classical statement. The countable witness of appendix A.2 has $R = P_-$ and fails (5.49). For (iv), finiteness makes $Q_{-B} \ll \mathbf{\Pi}_{\theta,V,o,X,-B}$ by data processing, and (5.46) leaves the outside term unchanged; the exact

replacement sets the conditional term to zero, and convexity of relative entropy in its first argument, applied for each $y_{-B}$, gives (5.50). □

The exact full E-step is case (iii) with $\mathcal{B}$ equal to all coordinates; a coordinate that optimizes inside a product family does not have the form (5.48) and leaves the gap of proposition 5.9. By (iv) a sequential sweep of exact or damped block updates does not increase the collective free energy; replacing the conditionals of several correlated blocks at once is not covered, because each replacement changes the outside marginal of the others.

**Theorem 5.23 (Exact local–global potential identity).** Let $Q = Q_{-B} r_B$ and $Q' = Q_{-B} r'_B$ have the same outside marginal with $Q_{-B} \ll \mathbf{\Pi}_{\theta,V,o,X,-B}$, and suppose both have finite divergence from $\mathbf{\Pi}_{\theta,V,o,X}$. Then

$$\mathcal{F}^{\mathrm{ext}}_o(Q') - \mathcal{F}^{\mathrm{ext}}_o(Q) = \mathbb{E}_{Q_{-B}}\big[\mathcal{F}^{\mathrm{ext}}_{B,o}(r'_B; Y_{-B}) - \mathcal{F}^{\mathrm{ext}}_{B,o}(r_B; Y_{-B})\big]. \tag{5.51}$$

When in addition every baseline relative entropy and the positive and negative parts of the energies are integrable,

$$\mathcal{F}^{\mathrm{ext}}_o(Q) = D_{\mathrm{KL}}\big(Q_{-B} \| \mathbb{P}^{\mathsf{Y}}_{\theta,V,-B}\big) + \mathbb{E}_{Q_{-B}}\big[H_{\bar{B},o}(Y_{-B}) + \mathcal{F}^{\mathrm{ext}}_{B,o}(r_B; Y_{-B})\big]. \tag{5.52}$$

Every exact local update is thus a coordinate update of the collective free energy. The fixed outside marginal is load bearing. [ESTABLISHED]

*Proof.* Apply (5.46) to both laws. The finite divergences make both conditional terms integrable, the outside terms agree, the constant $-\log p_{\theta,V}(o \mid X)$ cancels, and so does $-\log Z_B(Y_{-B})$ inside the two local functionals. The decomposition follows from (5.47), (5.45), and (5.19). □

**Corollary 5.24 (Exact potential game).** Let the blocks be singletons and choose a nonempty strategy set $\mathcal{R}_i$ for each agent. Suppose every profile $(r_i) \in \prod_{i \in V} \mathcal{R}_i$ defines a product recognition law $Q = \bigotimes_{i \in V} r_i$ with $D_{\mathrm{KL}}(Q \| \mathbf{\Pi}_{\theta,V,o,X}) < \infty$ and $\mathbb{E}_{Q_{-i}} |\log Z_{\{i\}}(Y_{-i})| < \infty$ for every $i$, where $Q_{-i} = \bigotimes_{j \neq i} r_j$. The outside-averaged local free energy

$$\overline{\mathcal{F}}_i(r_i; r_{-i}) := \mathbb{E}_{Q_{-i}} \mathcal{F}^{\mathrm{ext}}_{\{i\},o}(r_i; Y_{-i}) \tag{5.53}$$

is finite, and between two such profiles that differ only in $r_i$ it changes by exactly the change of $\mathcal{F}^{\mathrm{ext}}_o$. The players $V$, strategy sets $\mathcal{R}_i$, and costs $\overline{\mathcal{F}}_i$ therefore form an exact potential game with potential $\mathcal{F}^{\mathrm{ext}}_o$ in the sense of Monderer and Shapley [62, Sec. 2], written for costs instead of payoffs. [ESTABLISHED]

*Proof.* Finite divergence gives $Q_{-i} \ll \mathbf{\Pi}_{\theta,V,o,X,-i}$ by data processing, and (5.46) makes the averaged conditional divergence finite; with the hypothesis on $\log Z_{\{i\}}$ this makes (5.53) finite. A unilateral change leaves $Q_{-i}$ fixed, and theorem 5.23 with $B = \{i\}$ gives the equality of changes. □

Figure 5.1 shows which records enter a block potential; block-gradient dynamics are in section 8.1.

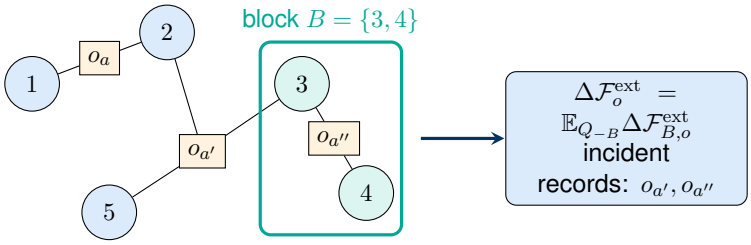



**Figure 5.1:** A block-local free energy contains every record incident to the block: $o_{a'}$ crosses the boundary and enters the block conditional, while $o_a$ lies outside and cancels.

## 5.8 Additive accounting

The local free energies of definition 5.21 are coordinate potentials of the collective free energy, not shares of it; an additive single-count ledger is a different construction.

**Corollary 5.25 (Correlated ledger and incident overcounting).** Suppose $\mathbb{P}_{\theta,V}^{\mathsf{Y}} = \bigotimes_{i\in V} \varrho_i$ and let $Q$ be any law on the latent space. Then (i) in $[0, +\infty]$,

$$D_{\mathrm{KL}}\big(Q\|\mathbb{P}_{\theta,V}^{\mathsf{Y}}\big) = \mathrm{TC}(Q) + \sum_i D_{\mathrm{KL}}(Q_i\|\varrho_i); \tag{5.54}$$

(ii) whenever (5.19) holds,

$$\mathcal{F}_o^{\mathrm{ext}}(Q) = \mathrm{TC}(Q) + \sum_i D_{\mathrm{KL}}(Q_i\|\varrho_i) + \mathbb{E}_Q\Big[\sum_{a\in\mathfrak{A}} E_{a,o}(Y_{\partial a})\Big], \tag{5.55}$$

in which each record appears once and $\mathrm{TC}(Q) = 0$ exactly when $Q$ is a product law, so mean field is a special case of the ledger and not a premise of it; (iii) if every record has a nonempty scope, the singleton incident energies satisfy, pointwise in $(-\infty, +\infty]$,

$$H_{\{i\},o}(y_i; y_{-i}) = \sum_{a: i\in\partial a} E_{a,o}(y_{\partial a}), \qquad \sum_{i\in V} H_{\{i\},o}(y_i; y_{-i}) = \sum_{a\in\mathfrak{A}} |\partial a| E_{a,o}(y_{\partial a}), \tag{5.56}$$

so the singleton local free energies count each record $|\partial a|$ times and are not the single-count ledger; (iv) this is a statement about these local free energies, not a bar to additive decompositions: iterating lemma 3.12 along any ordering of $V$ writes $D_{\mathrm{KL}}(Q\|\mathbf{\Pi}_{\theta,V,o,X})$ as one term per agent, each conditioned on the agents before it. [ESTABLISHED]

*Proof.* Part (i) is corollary 5.1 with $\varrho_j = \varrho_i$, and (ii) substitutes it into (5.19). For (iii), every $E_{a,o}$ lies in $(-\infty, +\infty]$ and both index sets are finite, so the rearrangement is valid pointwise. Part (iv) is repeated application of the chain rule, peeling one agent at a time. □

The same chain rule accounts for link values once they are latent; the interaction graph itself stays structural (section 1.4).

**Proposition 5.26 (Collective free energy with random link values).** Let the link family $\Theta = (\Theta_e^x)$ on the edge copies of $\mathfrak{I}$ be a latent variable with values in a standard Borel space $\mathscr{T}$. Let $\nu^{\mathscr{T}}$ be a probability law on $\mathscr{T}$, let $\mathscr{S} : \mathscr{T} \to (-\infty, +\infty]$ be a measurable action with $0 < Z_{\mathscr{S}} = \int e^{-\mathscr{S}} d\nu^{\mathscr{T}} < \infty$, and let the link prior be $Z_{\mathscr{S}}^{-1} e^{-\mathscr{S}} \nu^{\mathscr{T}}$. For each $\Theta$ let the record model of hypothesis 5.5 with links $\Theta$ have evidence $p(o \mid \Theta) \in (0, \infty)$, measurable in $\Theta$, posterior $\mathbf{\Pi}_\Theta$, and free energy $\mathcal{F}_{\mathrm{matter}}(Q_\Theta; o, \Theta) = -\log p(o \mid \Theta) + D_{\mathrm{KL}}(Q_\Theta \| \mathbf{\Pi}_\Theta)$, and assume the joint evidence $p(o) = \int p(o \mid \Theta) Z_{\mathscr{S}}^{-1} e^{-\mathscr{S}(\Theta)} \nu^{\mathscr{T}}(d\Theta)$ is finite. For a recognition law $Q(d\Theta, dy) = R(d\Theta) Q_\Theta(dy)$ with $D_{\mathrm{KL}}(R \| \nu^{\mathscr{T}}) < \infty$ and $\mathbb{E}_R |\mathscr{S}| < \infty$, the extended free energy of the joint model is

$$\mathcal{F} = \mathbb{E}_R \mathcal{F}_{\mathrm{matter}}(Q_\Theta; o, \Theta) + \mathbb{E}_R \mathscr{S} + D_{\mathrm{KL}}(R \| \nu^{\mathscr{T}}) + \log Z_{\mathscr{S}} \quad \text{in } (-\infty, +\infty]. \tag{5.57}$$

Omitting the link entropy term or the normalizer changes the free energy; when the links are structural parameters instead of latent variables, only $\mathcal{F}_{\mathrm{matter}}$ appears. [ESTABLISHED]

*Proof.* The posterior of the joint model has conditional law $\mathbf{\Pi}_\Theta$ given $\Theta$ and link marginal $\mathbf{\Pi}^{\mathscr{T}}(d\Theta) = p(o)^{-1} p(o \mid \Theta) Z_{\mathscr{S}}^{-1} e^{-\mathscr{S}(\Theta)} \nu^{\mathscr{T}}(d\Theta)$, so the chain rule (3.20) gives $D_{\mathrm{KL}}(Q \| \mathbf{\Pi}) = D_{\mathrm{KL}}(R \| \mathbf{\Pi}^{\mathscr{T}}) + \mathbb{E}_R D_{\mathrm{KL}}(Q_\Theta \| \mathbf{\Pi}_\Theta)$ in $[0, +\infty]$, with a measurable integrand. Since $\mathbb{E}_R |\mathscr{S}| < \infty$, $R$ charges no set where $\mathscr{S} = +\infty$, so $R \ll \mathbf{\Pi}^{\mathscr{T}}$ and $\log(dR/d\mathbf{\Pi}^{\mathscr{T}}) = \log(dR/d\nu^{\mathscr{T}}) + \mathscr{S} + \log Z_{\mathscr{S}} - \log p(o \mid \Theta) + \log p(o)$. The first two terms are $R$-integrable, and so is the positive part $f$ of $\log p(o \mid \cdot)$: for $f_n = \min(f, n)$, the nonnegative divergence from the law proportional to $e^{f_n}$ times the prior gives $\mathbb{E}_R f_n \le D_{\mathrm{KL}}(R \| Z_{\mathscr{S}}^{-1} e^{-\mathscr{S}} \nu^{\mathscr{T}}) + \log(1 + p(o))$, and that divergence equals $D_{\mathrm{KL}}(R \| \nu^{\mathscr{T}}) + \mathbb{E}_R \mathscr{S} + \log Z_{\mathscr{S}}$. Integrating term by term and adding the conditional term, whose integrand is nonnegative, turns $\mathcal{F} = -\log p(o) + D_{\mathrm{KL}}(Q \| \mathbf{\Pi})$ into (5.57); every summand has an integrable negative part, so no difference of infinities occurs. □

## 5.9 Model coordinates and accepted updates

In the two-argument form $\mathcal{F}_V(\mathbb{Q}; \theta)$ of hypothesis 4.3, block coordinates move $\mathbb{Q}$ at fixed $\theta$ and a model coordinate moves $\theta$ at fixed $\mathbb{Q}$; each accepted value defines a new fixed joint.

**Proposition 5.27 (Exact M-coordinate).** Hold $\mathbb{Q}_{V,o,X}$ fixed, and suppose the parameterized family supplies a common regular-observation set or jointly measurable pointwise versions for the candidates compared, with $\mathbb{E}_{\mathbb{Q}_{V,o,X}} |\log q_{V,o,X}(Y)| < \infty$ and $\mathbb{E}_{\mathbb{Q}_{V,o,X}} |\log p_{\theta',V}(o, Y \mid X)| < \infty$ for every candidate $\theta'$. If a maximizer exists in the declared parameter space, then every

$$\theta^{\mathrm{new}} \in \operatorname*{argmax}_{\theta'} \mathbb{E}_{\mathbb{Q}_{V,o,X}} \left[ \log p_{\theta',V}(o, Y \mid X) \right] \tag{5.58}$$

maximizes $\mathcal{L}_{V,\theta'}(\mathbb{Q}_{V,o,X}; X, o)$ over the same space. The finiteness hypotheses cannot be dropped. [ESTABLISHED]

*Proof.* The recognition term $\mathbb{E}_{\mathbb{Q}_{V,o,X}} [\log q_{V,o,X}]$ is finite and does not depend on $\theta'$. For the last sentence, on $\{2, 3, \ldots\}$ let $q_n = \mathrm{c}/(n(\log n)^2)$ with $\mathrm{c}$ normalizing, $h_n = \mathbf{1}_{\{n=2\}} - q_2$, and $r_n = 1 + \mathrm{c}_2 h_n$ with $0 < \mathrm{c}_2 < 1/q_2$, so that $r$ is bounded, strictly positive, nonconstant, and $\sum_n q_n r_n = 1$. With a

trivial record, let the two candidates have latent laws $p_0 = q$ and $p_1 = qr$, and let the recognition law be $q$. The entropy of $q$ is infinite (appendix A.2), so $\mathbb{E}_q \log p_0 = \mathbb{E}_q \log p_1 = -\infty$ and both candidates maximize the expected complete log-likelihood, whereas their extended bounds are $\mathbb{E}_q \log(p_0/q) = 0$ and $\mathbb{E}_q \log r = -D_{\mathrm{KL}}(q\|p_1) < 0$. □

**Definition 5.28 (Generalized M-coordinate acceptance).** A generalized M-coordinate need not attain a maximum, but it must pass the acceptance condition

$$\mathcal{L}^{\mathrm{ext}}_{V,\theta^{\mathrm{new}}}(\mathbb{Q}_{V,o,X}; X, o) \geq \mathcal{L}^{\mathrm{ext}}_{V,\theta^{\mathrm{old}}}(\mathbb{Q}_{V,o,X}; X, o), \tag{5.59}$$

evaluated at an observation in the regular sets of both candidates, or with jointly measurable versions declared as part of the parameterized model. No recognition density and no (H4) is required; on the classical domain the condition is the same inequality for (5.20), and (5.58) is a sufficient condition.

**Corollary 5.29 (Evidence-free acceptance test).** Suppose $\mathbb{Q}_{V,o,X} \ll M_{\theta^{\mathrm{old}},V,o,X}$, the two slice measures (5.14) are mutually absolutely continuous, and both extended values are finite. Then

$$\mathcal{L}^{\mathrm{ext}}_{V,\theta^{\mathrm{new}}}(\mathbb{Q}_{V,o,X}; X, o) - \mathcal{L}^{\mathrm{ext}}_{V,\theta^{\mathrm{old}}}(\mathbb{Q}_{V,o,X}; X, o) = \mathbb{E}_{\mathbb{Q}_{V,o,X}}\left[\log \frac{p_{\theta^{\mathrm{new}},V}(o, Y \mid X)}{p_{\theta^{\mathrm{old}},V}(o, Y \mid X)}\right], \tag{5.60}$$

so the test (5.59) needs no evidence value. [ESTABLISHED]

*Proof.* With $r_x = d\mathbb{Q}_{V,o,X}/dM_{\theta^x,V,o,X}$ for $x \in \{\mathrm{old}, \mathrm{new}\}$, the Radon–Nikodym chain rule gives $\log r_{\mathrm{old}} - \log r_{\mathrm{new}} = \log(dM_{\theta^{\mathrm{new}},V,o,X}/dM_{\theta^{\mathrm{old}},V,o,X})$ almost surely, and this derivative is the ratio of the two joint densities at $o$. By proposition A.6 the negative part of each $\log r_x$ is integrable, so finiteness makes each $\log r_x$ integrable, and subtracting the two representations (A.4) proves the identity. □

**Proposition 5.30 (Evidence monotonicity).** Suppose (H1)–(H2) hold for the fixed $\mathbb{Q}_{V,o,X}$ at both $\theta^{\mathrm{old}}$ and $\theta^{\mathrm{new}}$. If the E-coordinate is exact at the old parameter, $\mathbb{Q}_{V,o,X} = \mathbf{\Pi}_{\theta^{\mathrm{old}},V,o,X}$, and the M-coordinate satisfies (5.59), then

$$\log p_{\theta^{\mathrm{new}},V}(o \mid X) \geq \mathcal{L}^{\mathrm{ext}}_{V,\theta^{\mathrm{new}}}(\mathbb{Q}_{V,o,X}; X, o) \geq \mathcal{L}^{\mathrm{ext}}_{V,\theta^{\mathrm{old}}}(\mathbb{Q}_{V,o,X}; X, o) = \log p_{\theta^{\mathrm{old}},V}(o \mid X). \tag{5.61}$$

Consequently a schedule that places a full exact E phase before each accepted M phase has nondecreasing evidence values, which converge in $(-\infty, +\infty]$ and to a finite limit when bounded above. Block E-coordinates and accepted M-coordinates make the compared bound values nondecreasing on a common functional; the transfer to the evidence in (5.61) uses the full E phase. [ESTABLISHED]

*Proof.* The first inequality is corollary 5.7 at $\theta^{\mathrm{new}}$, the second is (5.59), and the equality is the zero-gap case at $\theta^{\mathrm{old}}$; the singular-branch sum $\mathcal{L}^{\mathrm{ext}} + D_{\mathrm{KL}}$ is never written. A nondecreasing sequence of extended reals converges. Each block coordinate does not decrease the bound by corollary 5.22, and each accepted step does not decrease it by (5.59); only the final equality of

(5.61) uses the exact E phase. Under (H3)–(H4) at both candidates this is the classical conclusion of Dempster et al. [26] in the reading of Neal and Hinton [64]. □

Convergence of the parameter sequence needs regularity conditions of the kind identified for the classical algorithm by Wu [87]. A finite step along a natural-gradient or other optimizer direction is a proposal, not an accepted coordinate, and it can decrease the bound.

**Proposition 5.31 (Finite natural-gradient steps need not be monotone).** Let $\mathsf{G} \succ 0$, write $\nabla\mathcal{L}_V$ for the gradient of $\theta' \mapsto \mathcal{L}_{V,\theta'}(\mathbb{Q}_{V,o,X}; X, o)$, and consider the step $\theta^+ = \theta' + \mathrm{h}\mathsf{G}^{-1}\nabla\mathcal{L}_V$ of length $\mathrm{h} > 0$. At a nonstationary point the direction is an ascent direction,

$$\frac{d}{d\mathrm{h}}\mathcal{L}_{V,\theta^+}\Big|_{\mathrm{h}=0} = \nabla\mathcal{L}_V^{\top}\mathsf{G}^{-1}\nabla\mathcal{L}_V > 0. \tag{5.62}$$

If $\mathcal{L}_{V,\theta''} = \mathcal{L}^\star - \frac{1}{2}(\theta'' - \theta^\star)^{\top}\mathsf{H}(\theta'' - \theta^\star)$ with $\mathsf{H} \succ 0$ for $\theta''$ in a neighborhood $\mathcal{U}$, $d_{\max}$ is the largest generalized eigenvalue of $(\mathsf{H}, \mathsf{G})$, and $\mathbf{z}$ is a corresponding eigenvector with $\mathbf{z}^{\top}\mathsf{G}\mathbf{z} = 1$, then at every $\theta' = \theta^\star + \mathrm{c}_1\mathbf{z} \in \mathcal{U}$ with $\mathrm{c}_1 \neq 0$ and every step whose endpoint lies in $\mathcal{U}$,

$$\mathcal{L}_{V,\theta^+} - \mathcal{L}_{V,\theta'} = \tfrac{1}{2}d_{\max}\mathrm{c}_1^2\big[1 - (1 - \mathrm{h}d_{\max})^2\big] < 0 \qquad \text{whenever} \qquad \mathrm{h} > \frac{2}{d_{\max}}. \tag{5.63}$$

[ESTABLISHED]

*Proof.* The first display is the chain rule. With $\Delta\theta = \theta' - \theta^\star = \mathrm{c}_1\mathbf{z}$, $\nabla\mathcal{L}_V = -\mathsf{H}\Delta\theta$ and $\mathsf{H}\mathbf{z} = d_{\max}\mathsf{G}\mathbf{z}$ give $\theta^+ - \theta^\star = (1 - \mathrm{h}d_{\max})\mathrm{c}_1\mathbf{z}$; substituting in the quadratic gives the second display, whose bracket is negative exactly when $\mathrm{h}d_{\max} > 2$. □

Monotonicity at a finite step therefore needs a line search, a trust region, or another rule that verifies (5.59) on the same functional; a better preconditioner does not remove the requirement, because (5.63) is stated in the preconditioned spectrum.

## 5.10 What remains open

**Open Problem 5.32 (Convergence of the alternating scheme).** Proposition 5.30 gives monotone convergence of the evidence values under a full exact E phase and nothing about the parameters, and since (5.59) is a weak inequality the null step is an accepted M-coordinate, so the general scheme of block E-coordinates (5.48) and accepted M-coordinates need not move its parameter at all. The problem concerns models of hypothesis 5.5 with finitely many latent states, parameters in an open cone of a Euclidean space, and joint densities twice continuously differentiable in the parameters, under the scheme that alternates a full exact E phase with the exact M-coordinate (5.58) whose maximizer is unique. The regularity conditions of Wu [87] concern the limit points of the parameter sequence and its convergence when the stationary points at its limit value are isolated; isolation is not among the hypotheses of the class, in which uniqueness of the maximizer leaves open that the evidence is constant along a curve of parameters on which the complete-data likelihood is not. It is settled by conditions checkable on the model, including

compactness of the upper level sets of the evidence in the cone, that are met by some model of the class with parameter dimension at least two whose stationary set at the limit value of an orbit is not discrete, under which every orbit of the scheme converges to one stationary point of the extended bound, or by a model of that class with compact upper level sets whose orbit has more than one limit point.

## Physicist's summary

The population carries one generative law and one recognition law, and the only evidence bound is $\log p(o) - D_{\mathrm{KL}}(Q\|\mathbf{\Pi})$, defined for every $Q$ and equal to $-\infty$ when $Q$ charges what the posterior cannot see; everything else is bookkeeping of that relative entropy through the chain rule. Substituting marginal entropies overstates the bound by exactly the total correlation. Observations are records attached to latent states, and an environment can be rewritten as boundary nodes emitting them. Each block of agents has a local free energy built from the records that touch it; with the rest held fixed, lowering it lowers the collective free energy by the same amount, so the population is an exact potential game. Summed, the local free energies count every shared record once per participant. Parameters move only between evaluations, by accepted steps, and the evidence rises monotonically when each step follows an exact posterior refresh; convergence is open.

# Chapter 6

# Attention and Selective Communication

Which agent does a receiver listen to, and what does it hear? Attention enters the joint of chapter 4 as a latent source label: the softmax row is its generative posterior at a latent state, and the optimal recognition row held constant in the latent state is a different softmax, over averaged energies (proposition 6.2), at the exact cost of proposition 6.3. Communication enters as normalized encoder, transmission, and decoder kernels, and the exact receiver update of a comparison record fixes the value map by the query and key maps (proposition 6.19). The row calculus and the reduced curvature of proposition 6.8 feed chapter 8, the event law of definition 6.4 is the attention state that chapter 11 coarse-grains, and the comparison records supply the endpoint energies of chapter 10. Rows, events, links, and divergences are indexed by oriented edge copies $e : j \to i$; the pair $ij$ names a unique copy, a pair quantity is the sum over parallel copies, and the construction serves both channels, with a model-channel row written $\beta^m$ where the channel must be named. Longer proofs are deferred to appendix B.

This chapter does not claim that transformer attention as a whole is derived (table 6.1; feed-forward blocks and layer normalization have no counterpart), that declared per-head rows follow from a label model, or that an interface map, its rank, or its coding establishes a channel capacity, a reconstruction of the sender's belief, or shared meaning. Capacity needs an admissible input family and, for a nonzero continuous Gaussian channel, an input resource constraint; a coding theorem also needs repeated uses, errors, and rates [72]; a task-dependent distortion or sufficient statistic is further model data. Open are the label-level multi-head model (open problem 6.22), discrete creation and deletion of interaction edges (open problem 6.9), and formation with evolving selective interfaces (open problem 6.27). [NOT-CLAIMED]

## 6.1 Attention as a latent source label

**Definition 6.1 (Per-receiver source-label model).** Let $\mathcal{I}_{\mathrm{att}}$ be a finite set of receiving rows, each an agent $i \in V$ or an agent–head pair $(i, h)$ that inherits its agent's latent state (written $i$ when no head is needed). Row $i$ carries a source label $J_i$ valued in the incoming edge copies of $i$, with a structural prior row $\pi_i$ independent of recognition whose support is the source set $\mathcal{N}_i$; source sums and indices range over $\mathcal{N}_i$. By section 1.4 the interaction multigraph at $c_*$ is the union of these supports, complete unless a mask sets a prior weight to zero, and whether $\mathcal{N}_i$ contains the self

edge is part of the model. The labels augment the latent baseline independently,

$$\mathbb{P}^{Y,\mathrm{aug}}_{\theta,V}(dy, dj) = \mathbb{P}^{Y}_{\theta,V}(dy) \prod_{i \in \mathcal{I}_{\mathrm{att}}} \pi_i(dj_i). \tag{6.1}$$

Each row has a temperature $\tau_i > 0$ and finite measurable energies $D_{ij} : \mathsf{Y} \to \mathbb{R}$ that contain no recognition law and are gauge-invariant functions of sample-level latent states (definition 3.19) and of the links, and at the selected record row $i$ has the record density

$$L_i(o_i \mid y, J_i = j) = \mathrm{c}_i(o_i, y) \exp[-D_{ij}(y)/\tau_i], \tag{6.2}$$

with $\mathrm{c}_i$ independent of $j$; the kernel off the selected slice is part of the model. When $\mathrm{c}_i(o_i, \cdot)$ depends only on $y_i$ and each $D_{ij}$ only on $(y_i, y_j)$, the record is one hyperedge factor of hypothesis 5.5 with scope $\{i\} \cup \mathcal{N}_i$, and through (5.36) it enters the local free energy of the receiver and of every candidate source. Each row carries exactly one label, so receiver occupancy is fixed; inferring it needs a null label in the source alphabet or a different joint label model with its own likelihood, such as definition 6.5.

The slice form (6.2) is inhabited: a binary success record with success probability $\exp[-D_{ij}/\tau_i]$ and every $D_{ij} \geq 0$ realizes it with $\mathrm{c}_i = 1$, as in (5.10), and so does the comparison record of proposition 6.19 at every record value.

**Proposition 6.2 (Exact posterior and recognition attention).** Assume definition 6.1 and label exclusivity at the complete selected record $o$: the augmented likelihood factors as

$$L^{\mathrm{aug}}_o(y, j) = L^{Y}_o(y) \prod_{i \in \mathcal{I}_{\mathrm{att}}} \mathrm{c}_i(o_i, y) \exp[-D_{ij_i}(y)/\tau_i], \tag{6.3}$$

with $L^{Y}_o : \mathsf{Y} \to [0, +\infty)$ finite, measurable, and free of every label, and no other record or factor of the generative law reads any $J_i$. (a) On $\{L^{Y}_o \prod_k \mathrm{c}_k(o_k, \cdot) > 0\}$ the generative posterior row given $y$ and the record is

$$\beta^{P}_{ij}(y) = \pi_{ij} e^{-D_{ij}(y)/\tau_i} \Big/ \sum_k \pi_{ik} e^{-D_{ik}(y)/\tau_i}, \tag{6.4}$$

a version that may also be chosen off that set; it is strictly positive exactly on $\mathcal{N}_i$, and if row $i$ has nonnegative energies, an included self edge with $D_{ii} = 0$, and equal prior weights, no source has a larger posterior weight than the self edge. (b) Restrict the augmented recognition law to the full factorization

$$\mathbb{Q}(dy, dj) = \mathbb{Q}_Y(dy) \prod_{i \in \mathcal{I}_{\mathrm{att}}} \beta^{Q}_i(dj_i), \qquad \beta^{Q}_i(dj_i) \text{ independent of } y, \tag{6.5}$$

which makes the labels conditionally independent and every row independent of the latent state, and assume $\mathbb{E}_{\mathbb{Q}_Y} |D_{ij}| < +\infty$ for every source and $\mathbb{E}_{\mathbb{Q}_Y} |\log \mathrm{c}_i(o_i, Y)| < +\infty$. Then row $i$ contributes

(6.6) to the collective VFE, whose first term does not depend on $\beta_i^Q$, with (6.7):

$$-\mathbb{E}_{\mathbb{Q}_Y}\log \mathrm{c}_i(o_i,Y)+\mathcal{F}_i^{\mathrm{att}}(\beta_i^Q), \tag{6.6}$$

$$\mathcal{F}_i^{\mathrm{att}}(\beta_i^Q)=D_{\mathrm{KL}}(\beta_i^Q\|\pi_i)+\frac{1}{\tau_i}\sum_j \beta_{ij}^Q\mathbb{E}_{\mathbb{Q}_Y}D_{ij}. \tag{6.7}$$

Its unique minimizer on the open simplex over $\mathcal{N}_i$ is

$$\beta_{ij}^{Q\star}=\pi_{ij}e^{-\mathbb{E}_{\mathbb{Q}_Y}D_{ij}/\tau_i}\Big/\sum_k \pi_{ik}e^{-\mathbb{E}_{\mathbb{Q}_Y}D_{ik}/\tau_i}, \tag{6.8}$$

and the row functional $\sum_j \beta_{ij}\mathbb{E}_{\mathbb{Q}_Y}D_{ij}+\tau_i D_{\mathrm{KL}}(\beta_i\|\pi_i)$ equals $\tau_i\mathcal{F}_i^{\mathrm{att}}$ with the same minimizer, while the collective VFE contains $\mathcal{F}_i^{\mathrm{att}}$ itself. (c) For an arbitrary recognition law $\mathbb{Q}_Y(dy)\mathbb{Q}_{J|Y}(dj\mid y)$ with row marginals $\mathbb{Q}_{J_i|Y}$, and with $-\mathbb{E}_{\mathbb{Q}_Y}\log L_o^Y(Y)$ counted in the state sector, the attention-label contribution is

$$\begin{aligned}&-\sum_{i\in\mathcal{I}_{\mathrm{att}}}\mathbb{E}_{\mathbb{Q}_Y}\log \mathrm{c}_i(o_i,Y)+\mathbb{E}_{\mathbb{Q}_Y}\,\mathrm{TC}(\mathbb{Q}_{J|Y})\\&\quad+\sum_{i\in\mathcal{I}_{\mathrm{att}}}\mathbb{E}_{\mathbb{Q}_Y}\left[D_{\mathrm{KL}}(\mathbb{Q}_{J_i|Y}\|\pi_i)+\frac{1}{\tau_i}\mathbb{E}_{\mathbb{Q}_{J_i|Y}}D_{iJ_i}(Y)\right],\end{aligned} \tag{6.9}$$

where $\mathrm{TC}(\mathbb{Q}_{J|Y=y})=D_{\mathrm{KL}}\big(\mathbb{Q}_{J|Y=y}\|\prod_i\mathbb{Q}_{J_i|Y=y}\big)$. Its minimum over unrestricted conditional laws is attained by the product of the posterior rows (6.4), and row marginals independent of $y$ do not by themselves give (6.7), since a correlated conditional law keeps the nonnegative total-correlation term. [ESTABLISHED]

**Proposition 6.3 (Cost of the constant-row restriction).** Assume the label exclusivity of proposition 6.2, fix $\mathbb{Q}_Y$ with $\mathbb{E}_{\mathbb{Q}_Y}|D_{ij}|<+\infty$ for every row $i$ and $j\in\mathcal{N}_i$, write $d_{ij}=\mathbb{E}_{\mathbb{Q}_Y}D_{ij}$, and define

$$R_i^Q:=-\log\sum_j \pi_{ij}e^{-d_{ij}/\tau_i},\qquad R_i^P(y):=-\log\sum_j \pi_{ij}e^{-D_{ij}(y)/\tau_i}. \tag{6.10}$$

Over constant rows (6.5) the minimum of $\mathcal{F}_i^{\mathrm{att}}$ is $R_i^Q$, attained at (6.8); over all conditional label laws $\mathbb{Q}_{J|Y}$ the minimum of the label terms of (6.9) other than $-\sum_i\mathbb{E}_{\mathbb{Q}_Y}\log \mathrm{c}_i(o_i,Y)$ is $\sum_i\mathbb{E}_{\mathbb{Q}_Y}R_i^P(Y)$, attained by the product of the posterior rows (6.4). For each row the difference is exactly

$$R_i^Q-\mathbb{E}_{\mathbb{Q}_Y}R_i^P(Y)=\mathbb{E}_{\mathbb{Q}_Y}D_{\mathrm{KL}}\big(\beta_i^{Q\star}\|\beta_i^P(Y)\big)\geq 0, \tag{6.11}$$

which vanishes exactly when every difference $D_{ij}(Y)-D_{ik}(Y)$ is $\mathbb{Q}_Y$-almost surely constant. If further $D_{ij}=\widetilde{D}_{ij}\circ\psi_i$ for one measurable map $\psi_i:\mathsf{Y}\to\mathsf{Z}_i$ into a finite-dimensional real vector space with finite $\mathbb{Q}_Y$-mean $\bar{\psi}_i$, and every $\widetilde{D}_{ij}-\widetilde{D}_{ik}$ is affine, the recognition row is the posterior row at the recognition mean,

$$\beta_{ij}^{Q\star}=\pi_{ij}e^{-\widetilde{D}_{ij}(\bar{\psi}_i)/\tau_i}\Big/\sum_k \pi_{ik}e^{-\widetilde{D}_{ik}(\bar{\psi}_i)/\tau_i}. \tag{6.12}$$

In a linear sample fiber, the energies $\frac{1}{2}(y_i-\mathring{y}_{ij})^\top W_i(y_i-\mathring{y}_{ij})$ with $W_i \succ 0$ and structural templates $\mathring{y}_{ij}$ have affine differences. If instead $D_{ij}(y) = \frac{1}{2}r_{ij}^\top W_i r_{ij}$ with $r_{ij} = y_i - \rho_b(\Theta^b_{ij})y_j$ for a unique edge copy, and $Y_i$ and $Y_j$ are uncorrelated under $\mathbb{Q}_Y$ with finite second moments, then $\beta_i^{Q\star}$ is the posterior row at the recognition mean with the logit of source $j$ lowered by $\operatorname{Tr}(W_i\rho_b(\Theta^b_{ij})\operatorname{Cov}(Y_j)\rho_b(\Theta^b_{ij})^\top)/(2\tau_i)$. [ESTABLISHED]

The proofs of propositions 6.2 and 6.3 are in appendix B.1. Coarse-graining a row, a conditional law of a source given a receiver, needs a receiver law, which the next two definitions supply.

**Definition 6.4 (Receiver occupancy and typed event law).** For a role mark $x \in \{P, \Pi, Q\}$, a receiver occupancy $\alpha_i^x$, a probability on the rows, and a row $\beta_i^x$ on $\mathcal{N}_i$ give the marked edge-event law

$$\eta^x_{ij} = \alpha_i^x \beta^x_{ij}, \qquad \sum_{i,j} \eta^x_{ij} = 1. \tag{6.13}$$

The mark names the law: $P$ is the generative joint conditioned on the latent state and the record, so $\eta^P$, $\alpha^P$, and $\beta^P$ are functions of $y$; $\Pi$ is the posterior slice, with $\eta^\Pi_{ij} = \mathbb{E}_{\mathbf{\Pi}}\eta^P_{ij}(Y)$; and $Q$ is the selected recognition coupling. The event law carries no head index. The occupancy is not supplied by definition 6.1, in which every row carries its own label, and it enters in one of two typed ways, each use naming which: a structural receiver-selection law, fixed with the coarse channel before observation and independent of recognition, or the receiver marginal of a single joint edge label as in definition 6.5, with the label prior as its prior and its posterior from Bayes' formula. On $\{\alpha_i^x > 0\}$ the row is recovered by disintegration, $\beta^x_{ij} = \eta^x_{ij}/\alpha_i^x$; a conditional row alone is not recursive state.

**Definition 6.5 (Single-edge joint-label attention model).** This generative model differs from definition 6.1: the baseline is augmented by one joint edge label $(i, j)$, a single receiver–source pair for the selected record, with recognition-independent prior

$$\mathbb{P}^{\mathrm{att}}_{\theta,V}(dy, di, dj) = a_i\pi_{ij}\mathbb{P}^Y_{\theta,V}(dy), \qquad \sum_i a_i = 1, \quad \sum_j \pi_{ij} = 1, \tag{6.14}$$

and the augmented likelihood at the complete selected record $o$ factors as

$$L^{\mathrm{edge}}_o(y, i, j) = L^Y_o(y)\widetilde{\mathrm{c}}_i(o, y)\exp[-D_{ij}(y)/\tau_i], \qquad \tau_i > 0, \tag{6.15}$$

with $L^Y_o$ and $\widetilde{\mathrm{c}}_i$ finite, nonnegative, and measurable, $L^Y_o$ free of $(i, j)$, $\widetilde{\mathrm{c}}_i$ free of $j$, each $D_{ij}$ finite and measurable, and no other record or factor of the generative law reading the label; this exclusivity excludes every model in which several receivers attend at once. In general $\widetilde{\mathrm{c}}_i$ is not the $\mathrm{c}_i$ of (6.2): if the selected receiver's record has slice density $\mathrm{c}_i(o_i, y)\exp[-D_{ij}(y)/\tau_i]$ while every other receiver's record keeps a label-free kernel $g_k(o_k \mid y)$, then where every $g_k(o_k \mid y) > 0$ the factorization holds with $L^Y_o = \prod_k g_k(o_k \mid y)$ and $\widetilde{\mathrm{c}}_i = \mathrm{c}_i/g_i(o_i \mid y)$, a likelihood ratio against the non-event kernel. The evidence-weighted occupancy of this model is proposition 11.30; the per-receiver model collapses with its declared selection law instead (proposition 11.31).

**Corollary 6.6 (Self edges and blocking).** Let $N \geq 2$ agents carry one row each with uniform receiver occupancy $\alpha_i = 1/N$, partition them into blocks of equal size $n$ dividing $N$, and push an event law with rows independent of the latent state along the partition by block sums, $\eta^{\mathrm{c}}_{IJ} = \sum_{i\in I, j\in J} \eta_{ij}$, with $\alpha^{\mathrm{c}}_I = \sum_J \eta^{\mathrm{c}}_{IJ}$ and $\beta^{\mathrm{c}}_{IJ} = \eta^{\mathrm{c}}_{IJ}/\alpha^{\mathrm{c}}_I$. If every fine row is uniform over $j \neq i$, the block event masses are $n^2/(N(N-1))$ from $I$ to $J \neq I$ and $n(n-1)/(N(N-1))$ from $I$ to itself, the block occupancy is $n/N$, and

$$\beta^{\mathrm{c}}_{II} = \frac{n-1}{N-1}, \qquad \beta^{\mathrm{c}}_{IJ} = \frac{n}{N-1} \quad (J \neq I). \tag{6.16}$$

If every fine row is uniform over all $N$ sources, including the self edge, every coarse row is uniform with weight $n/N$. A fine no-self convention is therefore not preserved: the coarse source set of $I$ contains $I$. [ESTABLISHED]

*Proof.* Without self edges each fine event has mass $1/(N(N-1))$, a block pair $I \neq J$ contains $n^2$ of them, $I \times I$ contains $n(n-1)$, and the block occupancy is $n/N$; with self edges each event has mass $1/N^2$ and every block pair contains $n^2$ events. Dividing gives both cases. □

**Corollary 6.7 (Row flows, dissipation, and parameter paths).** Hold $\mathbb{Q}_Y$, and hence every $\mathbb{E}_{\mathbb{Q}_Y} D_{ij}$, fixed. With row relaxation rate $\mathsf{r}_i \geq 0$, the Fisher natural gradient of (6.7) on the open simplex over $\mathcal{N}_i$ is the replicator equation (6.17), with partial row-free-energy rate (6.18):

$$\dot\beta^Q_{ij} = -\mathsf{r}_i \beta^Q_{ij} \left( \mathsf{c}_{ij} - \sum_k \beta^Q_{ik} \mathsf{c}_{ik} \right), \qquad \mathsf{c}_{ij} = \log \frac{\beta^Q_{ij}}{\pi_{ij}} + \frac{\mathbb{E}_{\mathbb{Q}_Y} D_{ij}}{\tau_i}, \tag{6.17}$$

$$\frac{d}{dt} \mathcal{F}^{\mathrm{att}}_i \Big|_{\beta_i} = -\mathsf{r}_i \operatorname{Var}_{\beta^Q_i}(\mathsf{c}_{ij}) \leq 0. \tag{6.18}$$

If instead $\mathbb{Q}_{Y,\lambda(t)}$ moves along a $C^1$ curve $\lambda$ in an open coordinate domain, with every $d_{ij}(\lambda) = \mathbb{E}_{\mathbb{Q}_{Y,\lambda}} D_{ij}$ of class $C^1$, and the row follows (6.17) with $d_{ij}(\lambda(t))$ in place of $\mathbb{E}_{\mathbb{Q}_Y} D_{ij}$, then $\mathcal{F}^{\mathrm{att}}_i(\beta_i; \lambda) = D_{\mathrm{KL}}(\beta_i \| \pi_i) + \tau_i^{-1} \sum_j \beta_{ij} d_{ij}(\lambda)$ satisfies

$$\frac{d}{dt} \mathcal{F}^{\mathrm{att}}_i(\beta^Q_i(t); \lambda(t)) = -\mathsf{r}_i \operatorname{Var}_{\beta^Q_i}(\mathsf{c}_{ij}) + \frac{1}{\tau_i} \sum_j \beta^Q_{ij} \nabla_\lambda d_{ij}(\lambda) \cdot \dot\lambda, \tag{6.19}$$

whose second term has no sign in general. Along any differentiable path on $\{\alpha_i > 0\}$ an event law and its row satisfy (6.20), and a softmax row $\beta_j = \pi_j e^{-E_j/\tau} / \sum_k \pi_k e^{-E_k/\tau}$ on a strictly positive prior support, with $(\pi, E, \tau)$ differentiable and $\tau > 0$, satisfies (6.21):

$$\dot\beta_{ij} = \frac{\dot\eta_{ij} - \beta_{ij} \dot\alpha_i}{\alpha_i}, \tag{6.20}$$

$$\dot\beta_j = \beta_j \left( \chi_j - \sum_k \beta_k \chi_k \right), \qquad \chi_j = \partial_t \log \pi_j - \frac{\dot E_j}{\tau} + \frac{E_j \dot\tau}{\tau^2}. \tag{6.21}$$

[ESTABLISHED]

The proof is in appendix B.1. No scale map determines the path in (6.21), and on $\{\alpha_i = 0\}$ the flow variable is $\eta$, not the quotient row.

**Proposition 6.8 (Attention-reduced potential and its curvature).** Fix a row $i$, an open set $U$ of a finite-dimensional real vector space, and $E_{ij} \in C^2(U)$ for $j \in \mathcal{N}_i$, and define

$$R_i(x) = -\log \sum_j \pi_{ij} e^{-E_{ij}(x)/\tau_i}, \qquad \beta_{ij}(x) = \pi_{ij} e^{-E_{ij}(x)/\tau_i} \Big/ \sum_k \pi_{ik} e^{-E_{ik}(x)/\tau_i}. \tag{6.22}$$

Then $R_i \in C^2(U)$, and

$$\nabla R_i = \frac{1}{\tau_i} \sum_j \beta_{ij} \nabla E_{ij}, \qquad \nabla^2 R_i = \frac{1}{\tau_i} \sum_j \beta_{ij} \nabla^2 E_{ij} - \frac{1}{\tau_i^2} \operatorname{Cov}_{\beta_i}(\nabla E_i), \tag{6.23}$$

where $\operatorname{Cov}_{\beta_i}(\nabla E_i) = \sum_j \beta_{ij} (\nabla E_{ij} - \bar{g})(\nabla E_{ij} - \bar{g})^\top$ and $\bar{g} = \sum_j \beta_{ij} \nabla E_{ij}$. With $\mathscr{J}_i(x, \beta) = D_{\mathrm{KL}}(\beta \| \pi_i) + \tau_i^{-1} \sum_j \beta_j E_{ij}(x)$, the gradient is the frozen-weight gradient $\partial_x \mathscr{J}_i(x, \beta_i(x))$, and the Hessian is the Schur complement, over the tangent space $\{w : \sum_j w_j = 0\}$ of the source simplex, of the joint Hessian of $\mathscr{J}_i$ at $(x, \beta_i(x))$. The potential arises under both label treatments: under the constant-row restriction, with $x = \lambda$ the coordinate of a recognition family $\mathbb{Q}_{Y,\lambda}$ whose averaged energies $E_{ij}(\lambda) = \mathbb{E}_{\mathbb{Q}_{Y,\lambda}} D_{ij}$ are $C^2$, $R_i(\lambda)$ is the row-optimized value $R_i^Q$ of proposition 6.3; under exact label integration, with $x = y$ and $E_{ij} = D_{ij}$ on an open set of latent coordinates where every $D_{ij}$ is $C^2$, summing over $J_i$ in (6.3) replaces the label factor by $\mathrm{c}_i(o_i, y) e^{-R_i(y)}$, so $R_i = R_i^P$ enters the negative log posterior density of the latent state relative to its baseline. In both cases the correction along a direction $v$ is $-\tau_i^{-2} \operatorname{Var}_{\beta_i}(\nabla E_{ij} \cdot v)$, so positivity of the frozen-weight curvature $\sum_j \beta_{ij} \nabla^2 E_{ij}$ does not imply positivity of $\nabla^2 R_i$, and in the following normalized witness the correction reverses the sign. Let the latent $X \in \mathbb{R}$ have baseline $\mathcal{N}(0, \tau/\omega)$ with $\omega, \tau > 0$, let one row have source alphabet $\{-1, 1\}$ with uniform prior, and let the record $O \in \mathbb{R}$ have kernel $\mathcal{N}(x - j\delta, \tau)$ given $J = j$, with $\delta > 0$. This kernel is $\mathrm{c} \exp[-D_j^o(x)/\tau]$ with $\mathrm{c} = (2\pi\tau)^{-1/2}$ and $D_j^o(x) = (o - x + j\delta)^2/2$, so (6.2) holds at every record value, and the evidence at $o = 0$, $\exp[-\delta^2/(2\tau(1 + \omega^{-1}))]/\sqrt{2\pi\tau(1 + \omega^{-1})}$, is finite and positive. Exact label integration gives the negative log posterior density

$$\frac{(\omega + 1)x^2}{2\tau} - \log\cosh\frac{\delta x}{\tau} + \text{const}, \qquad \text{with second derivative } \frac{1}{\tau}\left(\omega + 1 - \frac{\delta^2}{\tau}\right) \text{ at } x = 0, \tag{6.24}$$

whereas the prior curvature plus the frozen-weight curvature is $(\omega + 1)/\tau$. Under the constant-row restriction with Gaussian recognition $\mathcal{N}(\mathsf{m}, \mathsf{v})$, $\mathsf{v} > 0$, the averaged energies $((\mathsf{m} - j\delta)^2 + \mathsf{v})/2$ have differences affine in $\mathsf{m}$, and the $\mathsf{m}$-curvature of the row-optimized collective VFE at $\mathsf{m} = 0$ is again $(\omega + 1 - \delta^2/\tau)/\tau$. For $(\omega, \delta, \tau) = (1, 2, 1)$ the frozen-weight value is $2$ and the reduced value $-2$, although every source energy and the prior are convex, so the symmetric point is a strict local maximum of the reduced functional along the mean. [ESTABLISHED]

The proof is in appendix B.1.

**Open Problem 6.9 (Discrete creation and deletion of interaction edges).** At $c_*$ the interaction multigraph is the support of the source prior and belongs to the structural data. Every row of this chapter is strictly positive on it (proposition 6.2), and the row flows change no support at finite time: along (6.17), $\beta^Q_{ij}(t) = \beta^Q_{ij}(0)\exp[-\int_0^t \mathsf{r}_i(\mathsf{c}_{ij} - \bar{\mathsf{c}}_i)ds]$ with $\bar{\mathsf{c}}_i = \sum_k \beta^Q_{ik}\mathsf{c}_{ik}$ whenever the costs stay finite. Kinetics that change the support lie outside the model of this manuscript, and the problem fixes one rule so that the question is which normalized free energy, if any, such kinetics descend. Let the edge set be a multigraph on $V$ with edge multiplicity at most a fixed bound, let a comparison energy $D_{ij}$ be declared for every candidate copy $j \to i$, present or absent, and let the rows on the current support follow (6.17). The threshold rule deletes a present copy at a rate that is a continuous nonincreasing function of its recognition row weight $\beta^Q_{ij}$, positive below a threshold in $(0, 1)$ and zero above it, and creates an absent copy at a rate that is a continuous nonincreasing function of its averaged energy $\mathbb{E}_{\mathbb{Q}_Y} D_{ij}$, positive below a finite threshold and zero above it; at a deletion the receiver's row is restricted to the remaining copies and renormalized, and at a creation the new copy enters the row at the deletion threshold and the row is renormalized. The rule reads the recognition law and the averaged energies and no law over multigraphs. The problem is settled, for some pair of such rate functions, by a normalized joint at $c_*$ in which the edge set is a structural coordinate, declared as a configuration prior or structural channel with a named law and coupled to the link, interface, and label variables, such that the collective VFE decomposition of definition 5.6 holds with the edge set as a coordinate, the event law of definition 6.4 and its pushforward in section 11.7 extend to the random support with the zero-occupancy convention used there, and the recognition law over supports, evolving by the master equation of the rule with the rows on each support following (6.17), has nonincreasing collective VFE from every initial support and recognition law; negatively, by a proof that for every pair of such rate functions no joint of that class has this property.

## 6.2 Communication as a normalized channel

**Definition 6.10 (Directed communication interface).** For an oriented edge copy $e : j \to i$, whose independent link is that of definition 2.11, and a channel $x \in \{b, m\}$, let $\mathsf{S}^x_j$ be the source sample space. Choose transmitted and received message spaces $\mathsf{H}^-_e$ and $\mathsf{H}^+_e$, a record space $\mathsf{O}_e$, receiver side information $v$ in a measurable space, and encoding, transmission, and decoding kernels, in the sense of the communication model of Shannon [72], jointly measurable in $v$, whose composition is the composed value kernel $K^v_e$:

$$\mathsf{E}^v_e : \mathsf{S}^x_j \rightsquigarrow \mathsf{H}^-_e, \qquad \mathsf{N}^v_e : \mathsf{H}^-_e \rightsquigarrow \mathsf{H}^+_e, \qquad \mathsf{D}^v_e : \mathsf{H}^+_e \rightsquigarrow \mathsf{O}_e, \tag{6.25}$$

$$K^v_e(s, do) = \int \mathsf{E}^v_e(s, d\xi^-)\mathsf{N}^v_e(\xi^-, d\xi^+)\mathsf{D}^v_e(\xi^+, do). \tag{6.26}$$

**Corollary 6.11 (Normalization and the kernel-level gauge law).** (a) $K^v_e$ is a Markov kernel jointly measurable in $v$, and a source law $q^x_j$ gives the law $q^x_j K^v_e$ on $\mathsf{O}_e$, a candidate law when $\mathsf{O}_e$ is a receiver sample space, whose membership in a restricted receiver family requires a separate

closure condition. (b) Let $\mathsf{S}_j^x$ be the channel sample fiber on which $G$ acts measurably by $\rho_x$ as in (2.4), $\mathsf{O}_e$ the receiver's channel sample fiber, the side information gauge inert, and $\phi_e^{\pm}$ separately chosen bimeasurable recoordinatizations of the message spaces, and write $\mathsf{g}_i$ for the old-to-new bijection of the channel-$x$ sample fiber induced at agent $i$ by the vertex re-expression of (2.28). The re-expressed kernels

$$\begin{gathered}\mathsf{E}_e^{v\prime}(s',\cdot)=(\phi_e^-)_{\#}\mathsf{E}_e^v(\mathsf{g}_j^{-1}s',\cdot),\qquad \mathsf{N}_e^{v\prime}(\xi',\cdot)=(\phi_e^+)_{\#}\mathsf{N}_e^v((\phi_e^-)^{-1}\xi',\cdot),\\ \mathsf{D}_e^{v\prime}(\xi',\cdot)=(\mathsf{g}_i)_{\#}\mathsf{D}_e^v((\phi_e^+)^{-1}\xi',\cdot)\end{gathered}\tag{6.27}$$

compose to $K_e^{v\prime}(s',\cdot)=(\mathsf{g}_i)_{\#}K_e^v(\mathsf{g}_j^{-1}s',\cdot)$, so $q_j'K_e^{v\prime}=(\mathsf{g}_i)_{\#}(q_jK_e^v)$ for $q_j'=(\mathsf{g}_j)_{\#}q_j$; for a gauge-inert record space the pushforward by $\mathsf{g}_i$ is omitted. (c) For the represented link $\rho_x(\Theta_e^x)$, a bimeasurable bijection, $F_e^v(s,\cdot):=K_e^v(\rho_x(\Theta_e^x)^{-1}s,\cdot)$ satisfies $K_e^v(s,\cdot)=F_e^v(\rho_x(\Theta_e^x)s,\cdot)$, and the link law gives $F_e^{v\prime}(s',\cdot)=(\mathsf{g}_i)_{\#}F_e^v(\mathsf{g}_i^{-1}s',\cdot)$; this separation uses the invertibility of the link and does not determine the link from $K_e^v$. [ESTABLISHED]

The proof is in appendix B.1. A kernel may encode a belief-law or model-law point only when that point is a latent coordinate of hypothesis 5.5 at the law-point level of definition 3.19; by hypothesis 4.3 it never reads the live recognition law over that point, and on the model channel a compatible evaluator is additional data.

## 6.3 Factored links and loops

**Definition 6.12 (Factored matrix communication).** On the finite-dimensional vector tier, write $E_i$ for the sample fiber of agent $i$ in channel $x$, distinct from the law fiber $(\mathcal{E}_x)_{c_*}$, and suppress $x$. For each edge copy $e:j\to i$ choose linear maps, possibly depending on the copy and with message spaces of possibly different dimensions,

$$\mathsf{L}_e^{\mathrm{enc}}:E_j\to H_e^-,\qquad \mathsf{L}_e^{\mathrm{tr}}:H_e^-\to H_e^+,\qquad \mathsf{L}_e^{\mathrm{dec}}:H_e^+\to E_i,\qquad \mathsf{L}_e=\mathsf{L}_e^{\mathrm{dec}}\mathsf{L}_e^{\mathrm{tr}}\mathsf{L}_e^{\mathrm{enc}},\tag{6.28}$$

separately in the belief and model representations, crossing channels only through the morphisms of section 2.3. The Markov kernel of a linear realization is the composed value kernel $K_e^v$, not $\mathsf{L}_e$. For represented frames $U_j:\mathbb{R}^{\dim E}\to E_j$, a value extraction $\mathsf{M}^{\mathrm{val}}:\mathbb{R}^{\dim E}\to H$ and an output placement $\mathsf{M}^{\mathrm{out}}:H\to\mathbb{R}^{\dim E}$ give the shared-head specialization

$$\mathsf{L}_e^{\mathrm{enc}}=\mathsf{M}^{\mathrm{val}}U_j^{-1},\qquad \mathsf{L}_e^{\mathrm{tr}}=I_H,\qquad \mathsf{L}_e^{\mathrm{dec}}=U_i\mathsf{M}^{\mathrm{out}},\qquad \mathsf{L}_e=U_i\mathsf{M}^{\mathrm{out}}\mathsf{M}^{\mathrm{val}}U_j^{-1}.\tag{6.29}$$

With $\mathsf{G}_j=U_j^{-\top}U_j^{-1}$, sending uses $U_j\,\mathrm{im}(\mathsf{M}^{\mathrm{val}})^{\top}$ and placement uses $U_i\,\mathrm{im}\,\mathsf{M}^{\mathrm{out}}$; these and the comparison directions can overlap and need not partition a frame, and $\mathsf{M}^{\mathrm{out}}\mathsf{M}^{\mathrm{val}}$ need be neither orthogonal nor idempotent.

**Proposition 6.13 (Endpoint covariance and rank).** (a) Under the old-to-new maps $\mathsf{g}_i$ of corollary 6.11, here linear, and independent message coordinate changes $\phi_e^-\in\mathrm{GL}(H_e^-)$ and

$\phi_e^+ \in \mathrm{GL}(H_e^+)$,

$$\mathsf{L}_e^{\mathrm{enc}\prime} = \phi_e^- \mathsf{L}_e^{\mathrm{enc}} \mathsf{g}_j^{-1}, \qquad \mathsf{L}_e^{\mathrm{tr}\prime} = \phi_e^+ \mathsf{L}_e^{\mathrm{tr}} (\phi_e^-)^{-1}, \qquad \mathsf{L}_e^{\mathrm{dec}\prime} = \mathsf{g}_i \mathsf{L}_e^{\mathrm{dec}} (\phi_e^+)^{-1},$$
$$\mathsf{L}'_e = \mathsf{g}_i \mathsf{L}_e \mathsf{g}_j^{-1}, \tag{6.30}$$

and the message-basis freedom does not enlarge the structure group. (b) Direct composition gives

$$\ker \mathsf{L}_e^{\mathrm{enc}} \subseteq \ker \mathsf{L}_e, \qquad \mathrm{im}\, \mathsf{L}_e \subseteq \mathrm{im}\, \mathsf{L}_e^{\mathrm{dec}}, \qquad \mathrm{rank}\, \mathsf{L}_e \le \min(\mathrm{rank}\, \mathsf{L}_e^{\mathrm{enc}}, \mathrm{rank}\, \mathsf{L}_e^{\mathrm{tr}}, \mathrm{rank}\, \mathsf{L}_e^{\mathrm{dec}}), \tag{6.31}$$

so the encoder accesses the quotient $E_j/\ker \mathsf{L}_e^{\mathrm{enc}}$, the potential placement space is $\mathrm{im}\, \mathsf{L}_e^{\mathrm{dec}}$, and the reachable linear signal is the possibly smaller $\mathrm{im}\, \mathsf{L}_e$. An SPD source metric $\mathsf{G}_j$ identifies the quotient with $(\ker \mathsf{L}_e^{\mathrm{enc}})^{\perp_{\mathsf{G}_j}} = \mathsf{G}_j^{-1} \mathrm{im}(\mathsf{L}_e^{\mathrm{enc}})^\top$, and $\mathsf{G}'_j = \mathsf{g}_j^{-\top} \mathsf{G}_j \mathsf{g}_j^{-1}$, so a distinguished sending subspace requires a metric while the kernel and quotient do not. Every finite-dimensional $\mathsf{L}_e$ factors through its rank. (c) A weighted sum $\sum_h w_h \mathsf{L}_h$ of maps $E_j \to E_i$, each factoring through a space $H_h$, has rank at most $\min(\dim E_i, \sum_h \dim H_h)$, attained with rank-one factors: for $\dim E_i = \dim E_j = 2$, identity frames, $\mathsf{L}_h = \mathbf{e}_h \mathbf{e}_h^\top$, and positive weights, the sum is $\mathrm{diag}(w_1, w_2)$. [ESTABLISHED]

**Proposition 6.14 (Connection data and communication loops).** Keep the represented invertible link $\Theta_e$ of $e : j \to i$ on the vector tier, with $\Theta_{\bar{e}} = \Theta_e^{-1}$. (a) For any $\mathsf{L}_e$ and a chosen link, $F_e := \mathsf{L}_e \Theta_e^{-1}$ gives

$$\mathsf{L}_e = F_e \Theta_e, \qquad F'_e = \mathsf{g}_i F_e \mathsf{g}_i^{-1}, \tag{6.32}$$

and if $\mathsf{L}_e^{\mathrm{enc}} = \widetilde{\mathsf{L}}_e^{\mathrm{enc}} \Theta_e$ for some $\widetilde{\mathsf{L}}_e^{\mathrm{enc}} : E_i \to H_e^-$, then $F_e = \mathsf{L}_e^{\mathrm{dec}} \mathsf{L}_e^{\mathrm{tr}} \widetilde{\mathsf{L}}_e^{\mathrm{enc}}$. The separation does not recover $\Theta_e$ from communication, and a receiver-only processing map $F_e = F_i$ exists exactly when $\mathsf{L}_e \Theta_e^{-1}$ agrees on every incoming edge copy at $i$. (b) Communication on the reverse copy is separately specified and need not be an inverse, an adjoint, or of the same rank. For a closed walk $p = (e_0, \dots, e_{r-1})$ with $e_a : i_{a+1} \to i_a$ and $i_r = i_0$, the communication round trip

$$\mathsf{L}_p = \mathsf{L}_{e_0} \cdots \mathsf{L}_{e_{r-1}} \qquad \text{satisfies} \qquad \mathsf{L}'_p = \mathsf{g}_{i_0} \mathsf{L}_p \mathsf{g}_{i_0}^{-1}, \tag{6.33}$$

even when the maps are singular, so its rank, characteristic polynomial, and trace are invariant; a singular round trip records processing and possible loss of distinguishability and is not the link holonomy $\mathrm{Hol}^x(p)$ of (2.30). (c) For two sites with identity frames and identity links, $e : 1 \to 2$ reading coordinate one and writing coordinate two, and $f : 2 \to 1$ the reverse,

$$\mathsf{L}_e = \left(\begin{smallmatrix} 0 & 0 \\ 1 & 0 \end{smallmatrix}\right), \qquad \mathsf{L}_f = \left(\begin{smallmatrix} 0 & 1 \\ 0 & 0 \end{smallmatrix}\right), \qquad \mathsf{L}_f \mathsf{L}_e = \left(\begin{smallmatrix} 1 & 0 \\ 0 & 0 \end{smallmatrix}\right), \qquad \mathsf{L}_e \mathsf{L}_f = \left(\begin{smallmatrix} 0 & 0 \\ 0 & 1 \end{smallmatrix}\right) \tag{6.34}$$

give round trips that preserve one direction and discard the other despite trivial link holonomy, and $2\mathsf{L}_f$ in place of $\mathsf{L}_f$ adds gain at the same rank. The shared-head family (6.29) has length-$r$ product $U_{i_0}(\mathsf{M}^{\mathrm{out}} \mathsf{M}^{\mathrm{val}})^r U_{i_r}^{-1}$, which depends on a path only through its length and endpoints; independent edge maps remove that restriction. (d) Link holonomy and communication rank are independent: for $G = \mathrm{GL}(2)$ in its defining representation and two parallel copies $a, b : 1 \to 2$ with $\Theta_a = I_2$ and

$\Theta_b = \left(\begin{smallmatrix} 1 & 1 \\ 0 & 1 \end{smallmatrix}\right)$, the loop through $b$ and $\bar{a}$ has nonidentity holonomy, while arbitrary singular $\mathsf{L}_a, \mathsf{L}_b$ are admissible through $F_a = \mathsf{L}_a \Theta_a^{-1}$ and $F_b = \mathsf{L}_b \Theta_b^{-1}$. [ESTABLISHED]

The proofs of propositions 6.13 and 6.14 are in appendix B.1. General matrix link fields also occur in linear lattice gauge theory [83], whose compact-group action and continuum conclusions do not follow from this endpoint algebra.

## 6.4 Comparison, content, and the transformer correspondence

**Definition 6.15 (Role comparison, quadratic disagreement, and mean operator).** For an edge copy $e : j \to i$, comparison uses data separate from content transmission: a comparison space $H_e^{\mathrm{cmp}}$ and query and key maps $S_e^Q : E_i \to H_e^{\mathrm{cmp}}$ and $S_e^K : E_j \to H_e^{\mathrm{cmp}}$, or, in the measurable setting, measurement kernels with a common output space. For law points $q_i, q_j$ the role divergence is

$$d_e = D_{\mathrm{KL}}\big((S_e^Q)_{\#} q_i \| (S_e^K)_{\#} q_j\big), \tag{6.35}$$

a law-point-level energy (definition 3.19) admitted where it is finite; this cospan compares observables and chooses no map back into $E_i$. Reference maps with $S_e^Q = S_{\mathrm{ref}}^Q U_i^{-1}$ and $S_e^K = S_{\mathrm{ref}}^K U_j^{-1}$ reproduce projected query and key comparison; with independent links the sender can instead be measured by $S_e^K = S_{\mathrm{ref}}^K U_i^{-1} \Theta_e$, which recovers the reference formula when $\Theta_e = U_i U_j^{-1}$ without imposing it. For an SPD state metric $\mathsf{G}_i$ and an SPD comparison precision $W_e$, the quadratic disagreement is $V_e(k_i, k_j) = \frac{1}{2}(S_e^Q k_i - S_e^K k_j)^\top W_e (S_e^Q k_i - S_e^K k_j)$. Declared per-head rows $\beta_{ij}^{(h)}$, value extractions $\mathsf{M}_h^{\mathrm{val}} : \mathbb{R}^{\dim E} \to H_h$, and output placements $\mathsf{M}_h^{\mathrm{out}} : H_h \to \mathbb{R}^{\dim E}$ define the mean operator

$$\mu_i^{\mathrm{new}} = \mu_i + \sum_{j,h} \beta_{ij}^{(h)} U_i \mathsf{M}_h^{\mathrm{out}} \mathsf{M}_h^{\mathrm{val}} U_j^{-1} \mu_j, \tag{6.36}$$

which in reference coordinates is value extraction, weighted aggregation, output placement, and a residual, the operator form of multi-head attention [80]. The operator is declared: definition 6.1 constructs one label per row, and nothing here selects the content maps.

**Proposition 6.16 (Natural-gradient response to a quadratic disagreement).** Fix the maps, an SPD state metric $\mathsf{G}_i$, and an SPD precision $W_e$. The negative natural gradient of $V_e$ in $k_i$ is

$$-\operatorname{grad}_{\mathsf{G}_i} V_e = -\mathsf{G}_i^{-1} S_e^{Q\top} W_e S_e^Q k_i + \mathsf{G}_i^{-1} S_e^{Q\top} W_e S_e^K k_j. \tag{6.37}$$

Descent of this energy thus determines one receiving map $\mathsf{G}_i^{-1} S_e^{Q\top} W_e S_e^K$ and one self term, so a declared insertion such as (6.36) is a natural-gradient step of a quadratic disagreement only when its maps have this form; identifying $\mathsf{G}_i$ with a Fisher metric requires a regular statistical model that supplies it (chapter 7). [ESTABLISHED]

*Proof.* Differentiate $V_e$ in $k_i$ and multiply by $\mathsf{G}_i^{-1}$. □

**Corollary 6.17 (Softmax over role divergences).** (a) When the latent coordinates include the law points compared in (6.35), at the law-point level of definition 3.19, a latent configuration $y$ gives the source energies $D_{ij}(y) = d_{ij}$. On a finite positive-prior source set with finite costs and $\tau_i > 0$, the categorical row functional satisfies

$$\sum_j \beta_{ij} d_{ij} + \tau_i D_{\mathrm{KL}}(\beta_i \| \pi_i) = \tau_i D_{\mathrm{KL}}\big(\beta_i \| \beta_i^P(y)\big) - \tau_i \log \sum_j \pi_{ij} e^{-d_{ij}/\tau_i}, \tag{6.38}$$

so the posterior row (6.4) is its unique minimizer: softmax attention over divergences between current law points is the posterior row at a latent configuration. The recognition optimum (6.8) reduces to it for $\mathbb{Q}_Y = \delta_y$, and when the latent baseline gives $\{y\}$ zero mass that point mass is not absolutely continuous with respect to the posterior slice, so its relative entropy and VFE are infinite (proposition 5.4); the law-point prior that such a latent requires is open problem 3.23. (b) Let the projected key covariance be the same SPD matrix $\Sigma^K$ for every source, and put $\mu_i^Q = S_{\mathrm{ref}}^Q U_i^{-1} \mu_i$ and $\mu_j^K = S_{\mathrm{ref}}^K U_j^{-1} \mu_j$. For Gaussian law points

$$d_{ij} = d_i^\circ - (\mu_i^Q)^\top (\Sigma^K)^{-1} \mu_j^K + \tfrac{1}{2} (\mu_j^K)^\top (\Sigma^K)^{-1} \mu_j^K, \tag{6.39}$$

where $d_i^\circ$ collects every term independent of $j$, including the query covariance terms, so whitening gives a dot-product logit together with a key-norm bias and $\log \pi_{ij}$. The bias disappears when the whitened key norms are constant, or when a source prior cancels it at these comparison laws, and such compensation in a population model requires its own source law. Unequal projected key covariances generally retain additional source-dependent terms. [ESTABLISHED]

The proof is in appendix B.1.

**Corollary 6.18 (Content changes routing).** Let $p(j, y, v) = \pi_j p_j(y) r_j(v \mid y)$ and $q(j, y, v) = \beta_j \widetilde{q}(y) s_j(v \mid y)$ be normalized laws on a source label, a comparison state, and a content variable, with $q \ll p$ and finite terms. Then

$$D_{\mathrm{KL}}(q \| p) = D_{\mathrm{KL}}(\beta \| \pi) + \sum_j \beta_j \left[ D_{\mathrm{KL}}(\widetilde{q} \| p_j) + \mathbb{E}_{\widetilde{q}} D_{\mathrm{KL}}(s_j(\cdot \mid Y) \| r_j(\cdot \mid Y)) \right], \tag{6.40}$$

so the conditional content term generally changes the optimal routing, and matching $s_j = r_j$ removes it. [ESTABLISHED]

*Proof.* Apply lemma 3.12 to the disintegration over $j$, then over $y$ given $j$. □

### 6.4.1 Exact receiver update and transformer attention

**Proposition 6.19 (Exact receiver update in the comparison-record model).** Fix a receiver $i$ with a finite source set on which $\pi_{ij} > 0$, each source $j$ reaching $i$ through one edge copy $e : j \to i$, and suppose that these copies share the comparison space $H_i^{\mathrm{cmp}}$, the query map $S_i^Q := S_e^Q$, and an SPD precision $W_i$, while the key maps $S_e^K$ may depend on the copy. Let the attention record of

definition 6.1 be the comparison record

$$o_i \mid k_i, J_i = j, y_{-i} \sim \mathcal{N}\big(S_i^Q k_i - S_e^K k_j, \tau_i W_i^{-1}\big), \qquad \tau_i > 0, \tag{6.41}$$

observed at $o_i = 0$, where $y_{-i}$ collects every latent coordinate other than $k_i$ and $J_i$ and so fixes the source states $k_j$. At $o_i = 0$ its density is $\mathrm{c}_i \exp[-V_e(k_i, k_j)/\tau_i]$, with $V_e$ of definition 6.15 and $\mathrm{c}_i = (2\pi\tau_i)^{-\dim H_i^{\mathrm{cmp}}/2}(\det W_i)^{1/2}$, so it realizes (6.2). Suppose that it is the only factor reading $J_i$, as in (6.3), and that, given $y_{-i}$ and every record other than $o_i$, the law of $k_i$ is $\mathcal{N}(\mu_i, \Sigma_i)$ with $\Sigma_i \succ 0$. Write $\widetilde{W}_i = W_i/\tau_i$ and

$$\Sigma_i^+ = \big(\Sigma_i^{-1} + S_i^{Q\top}\widetilde{W}_i S_i^Q\big)^{-1}, \qquad \Sigma_i^{\mathrm{pred}} = \widetilde{W}_i^{-1} + S_i^Q \Sigma_i S_i^{Q\top}, \qquad r_{ij} = S_i^Q \mu_i - S_e^K k_j. \tag{6.42}$$

(a) Given $y_{-i}$ and all records, $J_i$ has the posterior weights $w_{ij}$, normalized over the source set, $k_i$ has the conditional components, and the conditional mean of $k_i$ is

$$w_{ij} \propto \pi_{ij} \exp\big[-\tfrac{1}{2} r_{ij}^\top (\Sigma_i^{\mathrm{pred}})^{-1} r_{ij}\big], \qquad k_i \mid J_i = j \sim \mathcal{N}\big(\Sigma_i^+(\Sigma_i^{-1}\mu_i + S_i^{Q\top}\widetilde{W}_i S_e^K k_j), \Sigma_i^+\big), \tag{6.43}$$

$$\mathbb{E}[k_i \mid y_{-i}, o] = \Sigma_i^+\Sigma_i^{-1}\mu_i + \sum_j w_{ij} \mathsf{L}_e^{\mathrm{val}} k_j, \quad \mathsf{L}_e^{\mathrm{val}} = \Sigma_i^+ S_i^{Q\top}\widetilde{W}_i S_e^K, \quad \Sigma_i^+\Sigma_i^{-1} = I - \Sigma_i^+ S_i^{Q\top}\widetilde{W}_i S_i^Q. \tag{6.44}$$

The value map is thus fixed by the query map, the key map, the record precision, and the receiver covariance, and the skip connection is replaced by the precision-weighted term $\Sigma_i^+\Sigma_i^{-1}\mu_i$; equivalently the update is $\mu_i + \sum_j w_{ij}\Sigma_i^+ S_i^{Q\top}\widetilde{W}_i(S_e^K k_j - S_i^Q \mu_i)$. (b) The weights use $\Sigma_i^{\mathrm{pred}}$ in place of $\widetilde{W}_i^{-1}$ and need not equal the posterior row (6.4) at $k_i = \mu_i$: for a scalar receiver with $S_i^Q = S_e^K = \widetilde{W}_i = \Sigma_i = 1$, $\mu_i = 0$, a uniform prior, and sources $k_1 = 0$ and $k_2 = 2$, one has $w_{i2} = e^{-1}/(1 + e^{-1})$ and $\beta_{i2}^P = e^{-2}/(1 + e^{-2})$ at $k_i = 0$, while $\mathsf{L}_e^{\mathrm{val}} = \Sigma_i^+\Sigma_i^{-1} = 1/2$. (c) The Woodbury identity $\Sigma_i^+ S_i^{Q\top}\widetilde{W}_i = \Sigma_i S_i^{Q\top}(\Sigma_i^{\mathrm{pred}})^{-1}$ holds, and for full-row-rank $S_i^Q$ and $\widetilde{W}_i^{-1} \to 0$ the value map tends to $R_e S_e^K$ with $R_e = \Sigma_i S_i^{Q\top}(S_i^Q \Sigma_i S_i^{Q\top})^{-1}$, the $\Sigma_i^{-1}$-minimal right inverse of $S_i^Q$. (d) With copy-dependent query maps, precisions, or comparison spaces, the same computation gives $w_{ij} \propto \pi_{ij}\mathcal{N}(0; r_{ij}, \Sigma_{ij}^{\mathrm{pred}})$ and component covariances $\Sigma_{ij}^+$, so the normalizer of the predictive law enters the weights. [ESTABLISHED]

The proof is in appendix B.1. Table 6.1 compares the pointwise attention network with the attention layer of Vaswani et al. [80] row by row; the metric lift $R_e S_e^K$ of (c) is developed for Gaussian beliefs in Supplement G.

**Proposition 6.20 (A shared receiver state couples head labels).** Give two heads $h = 1, 2$ of one receiver their own exclusive labels $J_{i,h}$ and comparison records of the form (6.41), conditionally independent given the latent state. At fixed $k_i$ the label posterior factors by head. For a scalar receiver with $k_i \sim \mathcal{N}(0, 1)$ given the rest, $S^Q = S^K = \widetilde{W} = 1$ for both heads, a uniform prior, and sources with states $k_1 = 0$ and $k_2 = a \neq 0$, the posterior mass of $(J_{i,1}, J_{i,2}) = (j, l)$ after $k_i$ is integrated is proportional to

$$\exp\big[-\tfrac{1}{3}(k_j^2 + k_l^2) + \tfrac{1}{3} k_j k_l\big], \tag{6.45}$$

**Table 6.1:** Correspondence between transformer attention and the pointwise attention network at $c_*$.

| Transformer | Pointwise attention network | Standing |
|---|---|---|
| Token position | Agent $i \in V$ at $c_*$, its embedding read as a coordinate of its latent state. | declared |
| Query and key projections | $S_e^Q = S_{\mathrm{ref}}^Q U_i^{-1}$, $S_e^K = S_{\mathrm{ref}}^K U_j^{-1}$ and the role divergence (6.35); with a common projected key covariance, a dot-product logit plus a key-norm bias, (6.39). | derived, corollary 6.17 |
| One key map shared by all edges | Independent links enter $S_e^K = S_{\mathrm{ref}}^K U_i^{-1}\Theta_e$ only through the represented edge twist $U_i^{-1}\Theta_e U_j$; the shared map is the identity-twist case $\Theta_e = U_i U_j^{-1}$, and a nonidentity twist, including nontrivial link holonomy, has no counterpart in that architecture. | derived, proposition 2.12 |
| Softmax row | Generative posterior of a categorical source label, (6.4); over law-point divergences, the posterior row at a latent configuration, (6.38), not the recognition optimum. | derived, proposition 6.2 and corollary 6.17 |
| Attention mask | Zero prior weight: the source lies outside the support of $\pi_i$, hence outside the interaction multigraph, where the posterior row vanishes. | derived, proposition 6.2 |
| Temperature, inverse square root of key width | The temperature $\tau_i > 0$ of (6.2); for $\Sigma^K = \varsigma^2 I$ the dot-product part of the logit $-d_{ij}/\tau_i$ is $(\mu_i^Q)^\top \mu_j^K/(\varsigma^2\tau_i)$, and $\varsigma^2\tau_i = \sqrt{\dim H_i^{\mathrm{cmp}}}$ reproduces the scaled dot product, a choice not derived here. | declared |
| Value and output maps, residual connection | Declared in (6.36); in the comparison-record model the exact update fixes the value map $\mathsf{L}_e^{\mathrm{val}}$ and weights the residual by $\Sigma_i^+\Sigma_i^{-1}$. | derived, proposition 6.19 |
| Multiple heads | Declared per-head rows $\beta_{ij}^{(h)}$; no label-level construction, and no head index in the event law (6.13). | absent, open problem 6.22 |
| Feed-forward block, layer normalization | No counterpart is modeled. | absent |

and its odds ratio between agreeing and disagreeing label pairs is $e^{a^2/3} \neq 1$, which excludes a product law. [ESTABLISHED]

The proof is in appendix B.1.

**Remark 6.21 (The Gaussian criterion for head-label coupling).** For two heads with independent positive label priors, let their stacked predictive record covariance $\Sigma^{\mathrm{rec}}$ be positive definite and independent of the labels. Write $\Lambda_{12}^{\mathrm{rec}}$ for the off-diagonal block of its inverse, and $\mathsf{a}_j, \mathsf{b}_l$ for the head-specific key means. The positive integrated posterior table $p_{jl}$ satisfies

$$\log \frac{p_{jl}p_{j'l'}}{p_{jl'}p_{j'l}} = -(\mathsf{a}_j - \mathsf{a}_{j'})^\top \Lambda_{12}^{\mathrm{rec}}(\mathsf{b}_l - \mathsf{b}_{l'}). \tag{6.46}$$

Thus the labels are independent exactly when every displayed pairing vanishes. Nonzero covariance between projected queries alone does not imply label dependence. [ESTABLISHED]

*Proof.* Expand the Gaussian log likelihood with precision $(\Sigma^{\mathrm{rec}})^{-1}$. Its mixed term is $-\mathsf{a}_j^\top \Lambda_{12}^{\mathrm{rec}} \mathsf{b}_l$; the prior and one-head terms cancel from the cross ratio. A positive table factors exactly when all its

cross ratios equal one. For example, a receiver $k_i \sim \mathcal{N}(0, I_3)$ with query maps $S_1^Q k_i = (k_{i,1}, k_{i,2})$ and $S_2^Q k_i = (k_{i,1}, k_{i,3})$, independent record noises of covariance $I_2$, and key means $(0,0), (0,a)$ and $(0,0), (0,b)$ has query cross-covariance $\mathrm{diag}(1,0)$ but every mixed key pairing zero. For $a, b \neq 0$ both labels remain informative. □

**Open Problem 6.22 (Label-level multi-head attention).** The per-head rows of (6.36) are declared, and the event law (6.13) carries no head index. Give a receiver $i$ at least two heads $h$, each an exclusive label $J_{i,h}$ with a comparison record of the form (6.41) and its own query map $S_{i,h}^Q$, key maps, precision, and temperature, the records conditionally independent given the latent state, in the setting of proposition 6.19, where $k_i$ has the conditional law $\mathcal{N}(\mu_i, \Sigma_i)$ given the rest. The heads read a common receiver coordinate when the projected queries $S_{i,h}^Q k_i$ and $S_{i,h'}^Q k_i$ are correlated under that law for some pair $h \neq h'$; require also that the integrated positive label table does not factor by head. These conditions define the coupled class of the problem. For two heads with label-independent predictive covariance, the second condition is the nonvanishing of some pairing in eq. (6.46); proposition 6.20 supplies one instance. The factorizing subclass is excluded, including heads that read disjoint channels through untied links under a receiver law that is a product over channels. What holds in the class is the per-row label posterior at a fixed receiver state, proposition 6.2 with rows $(i, h)$; the event law of definition 6.4 carries no head index, so the pushforward of proposition 11.29 gives a parent no row per head. A parent whose receiver state is one coarse coordinate of its block reads every row of its children through that coordinate, so the exact recursion multiplies the number of rows read through one coordinate at each depth unless a head-merging channel, which the recursion does not declare, is imposed; once declared, its conditional relative entropy enters (9.3). The problem is settled by an exact recursion with bounded head count in the class: a family of comparison-record models, closed under the pushforward of proposition 11.29 composed with a declared head-merging channel without increase of the number of heads per receiving row, in which the parent's receiver state is a coarse coordinate of its block with the pushed-forward conditional law, its per-head rows are the disintegrations of the merged event law, and its posterior mean given its records and the rest is the coarse coordinate of its children's posterior means; negatively, by a member of the class and a partition into blocks for which the posterior mean of a block's coarse coordinate, as a function of the source states, is reproduced by no comparison-record model at the block with fewer heads than the block has rows.

## 6.5 Compositions as records

Choosing $s = s(y_j)$ and $v = v(X, y_{\partial e})$ measurably in (6.26) makes the output one owned record of hypothesis 5.5, and replacing record kernels by these compositions leaves the joint (5.11) normalized on cycles by proposition 4.4, with distinct sending and receiving interfaces and under the typing of hypothesis 4.3. Records that share a realized random message retain it once as a latent or declare their joint kernel, since conditional independence cannot be assumed after it is marginalized; independent draws from one latent law point are independent given that point. Fixed-record density formulas keep the family domination and regular-slice hypotheses of hypothesis 5.5, under which

the collective VFE of definition 5.6 and the local–global potential identity of theorem 5.23 apply. Composed records realize rectangular endpoint maps on cycles (a linear Gaussian realization is in Supplement G) but settle neither a canonical compression of such maps nor coarse-graining through moving interfaces (proposition 6.26).

**Definition 6.23 (Learned and random interfaces).** If interfaces or links vary as generative parameters, their score includes every varying normalizer, baseline, and record term. If they are random relation coordinates that evolve during inference, they are latent coordinates with a normalized prior and declared dynamics in the augmented state, and natural-gradient evolution needs a regular configuration metric or a justified quotient; endpoint covariance specifies neither the metric nor the learning rule. In both cases the dynamics move link values and interface maps on the edge copies of the support of the source prior; creating or deleting edge copies is open problem 6.9.

## 6.6 Information loss and descent

**Proposition 6.24 (Conditional data processing through an interface).** For a source law and a side-information law, let $S_j$ be the source sample in $\mathsf{S}_j^x$, $\Upsilon$ the random receiver side information with values $v$, $\xi^-$ and $\xi^+$ the random messages, and $O$ the decoded record. Construction (6.26) makes $S_j \to \xi^- \to \xi^+ \to O$ a Markov chain conditional on $\Upsilon$, and

$$I(S_j; O \mid \Upsilon) \le I(S_j; \xi^+ \mid \Upsilon) \le I(S_j; \xi^- \mid \Upsilon) \qquad \text{in } [0, +\infty]. \tag{6.47}$$

[ESTABLISHED]

The proof is in appendix B.1.

**Proposition 6.25 (Communicated agreement leaves conditional information).** Let $q, p$ be laws on one standard Borel space with $q \ll p$ and finite relative entropy, and let $M$ be one measurable measurement into a standard Borel space with output laws $q_M, p_M$. Then

$$D_{\mathrm{KL}}(q\|p) = D_{\mathrm{KL}}(q_M\|p_M) + \int D_{\mathrm{KL}}\big(q(\cdot \mid M = z)\|p(\cdot \mid M = z)\big) q_M(dz). \tag{6.48}$$

For $M = (1, 0)$, the laws $\mathcal{N}((0, a), I_2)$ and $\mathcal{N}((0, b), I_2)$ have identical communicated laws and full relative entropy $(a - b)^2/2$. [ESTABLISHED]

*Proof.* Append the deterministic measurement to both laws, which leaves the relative entropy unchanged, and apply the disintegration form of lemma 3.12 over $M$. The example has equal first coordinates and Gaussian relative entropy $(a - b)^2/2$. □

The identity needs one measurement: distinct roles $S_e^Q \neq S_e^K$ need not ask the same question, and separate head marginals need not fix their joint law (proposition 3.17).

**Proposition 6.26 (Descent through a selective interface).** Let $E_j, E_i, E$ and $\overline{E}_j, \overline{E}_i, \overline{E}$ be finite-dimensional real vector spaces, and let $M_j : E_j \to \overline{E}_j$, $M_i : E_i \to \overline{E}_i$, and $M : E \to \overline{E}$ be surjective linear measurements. (a) Static intertwiner: for a linear $\mathsf{L}_e : E_j \to E_i$, a linear $\overline{\mathsf{L}}_e : \overline{E}_j \to \overline{E}_i$ with $M_i\mathsf{L}_e = \overline{\mathsf{L}}_e M_j$ exists exactly when

$$\mathsf{L}_e \ker M_j \subseteq \ker M_i, \tag{6.49}$$

and it is then unique. (b) Autonomous linear flow: for $\mathcal{A} \in \mathrm{End}(E)$, a linear $\overline{\mathcal{A}} \in \mathrm{End}(\overline{E})$ with $Me^{t\mathcal{A}} = e^{t\overline{\mathcal{A}}}M$ for every $t \in \mathbb{R}$, equivalently with $M\mathsf{u}$ solving $\dot{\bar{\mathsf{u}}} = \overline{\mathcal{A}}\bar{\mathsf{u}}$ for every solution of $\dot{\mathsf{u}} = \mathcal{A}\mathsf{u}$, exists exactly when

$$M\mathcal{A} = \overline{\mathcal{A}}M \tag{6.50}$$

is solvable, that is, exactly when $\mathcal{A} \ker M \subseteq \ker M$, and $\overline{\mathcal{A}}$ is then unique. (c) Moving measurement: on an open interval $(t_0, t_1)$, let $t \mapsto M_t : E \to \overline{E}$ be a $C^1$ family of surjections and $t \mapsto \mathcal{A}_t$ continuous. A family $\overline{\mathcal{A}}_t \in \mathrm{End}(\overline{E})$ with $\frac{d}{dt}(M_t\mathsf{u}(t)) = \overline{\mathcal{A}}_t M_t \mathsf{u}(t)$ for every solution of $\dot{\mathsf{u}} = \mathcal{A}_t\mathsf{u}$ on $(t_0, t_1)$ exists exactly when

$$\dot{M}_t + M_t\mathcal{A}_t = \overline{\mathcal{A}}_t M_t \qquad (t_0 < t < t_1) \tag{6.51}$$

is solvable, that is, exactly when $\dot{M}_t + M_t\mathcal{A}_t$ vanishes on $\ker M_t$ for every $t$; the family is then unique and continuous, with $\overline{\mathcal{A}}_t = (\dot{M}_t + M_t\mathcal{A}_t)M_t^+$ for any continuous right inverse $M_t^+$ of $M_t$. (d) If $\mathsf{L}_e \ker M_j \not\subseteq \ker M_i$, no linear map of the retained coordinates reproduces $M_i\mathsf{L}_e$; a coordinate swap with $M_i = M_j = (1, 0)$ is an instance, in which a hidden coordinate becomes visible. If $\mathcal{A} \ker M \not\subseteq \ker M$, two solutions with the same retained initial value have different retained velocities, so the retained coordinates do not evolve autonomously, and the same holds at each time for a moving measurement that violates (6.51). (e) A low-rank interface does not by itself give an instantaneous Markov blanket: for a joint law of a private state $\mathsf{z}^{\mathrm{in}}$, an interface $\mathsf{z}^{\mathrm{bd}}$, and an exterior $\mathsf{z}^{\mathrm{ex}}$, the blanket condition is $P(d\mathsf{z}^{\mathrm{in}}, d\mathsf{z}^{\mathrm{ex}} \mid \mathsf{z}^{\mathrm{bd}}) = P(d\mathsf{z}^{\mathrm{in}} \mid \mathsf{z}^{\mathrm{bd}})P(d\mathsf{z}^{\mathrm{ex}} \mid \mathsf{z}^{\mathrm{bd}})$, and if $\mathsf{z}^{\mathrm{in}}$ is nondegenerate and independent of $\mathsf{z}^{\mathrm{bd}}$ while $\mathsf{z}^{\mathrm{ex}} = \mathsf{z}^{\mathrm{in}}$, the conditional law of $(\mathsf{z}^{\mathrm{in}}, \mathsf{z}^{\mathrm{ex}})$ given $\mathsf{z}^{\mathrm{bd}}$ is concentrated on the diagonal and is not a product. [ESTABLISHED]

The proof is in appendix B.1.

**Open Problem 6.27 (Formation with evolving selective interfaces).** Take a connected block of agents on the vector tier with Gaussian belief laws whose means are the fast coordinates, linear factored communication (definition 6.12) through surjective linear interfaces $M_i$ of one fixed rank below the fiber dimension, and comparison records (6.41) whose query and key maps read the receiver and the transported source through the receiver's interface, $S_i^Q = M_i$ and $S_e^K = M_i\Theta_e$, with the labels of definition 6.1. The means follow the gradient flow of the sample-level energy of the current label assignment with identity mobility and diffusion covariance $2\mathrm{T}I dt$ at a temperature $\mathrm{T} > 0$, and the assignment is redrawn at a rate-one clock from the receiver posterior rows, as in definition 13.21; the interfaces are generative parameters in the sense of definition 6.23 that move by the natural-gradient flow of the collective VFE under a declared metric on interfaces of that rank, at a mobility $\epsilon_{\mathrm{M}} > 0$. This is the class of the problem. Along such a path the $\dot{M}_t$ term of (6.51) is not controlled by a fixed-interface construction, because the path reads the state, and neither is the label switching it induces. A parent for the block is a retained mechanism in the sense of

section 8.3: a Markov process on a state space of fixed dimension, below the dimension of the block's means and assignment and independent of $\mathrm{T}$, $\epsilon_{\mathrm{M}}$, and the horizon, whose state determines the communicated coordinates $M_{i,t}\mu_i$ of every agent of the block. The admitted state space is a smooth manifold, the stochastic coefficients between jumps are locally Lipschitz, initialization and jump maps are continuous in their native inputs, and the communicated-coordinate readout is $C^1$ on each assignment stratum. Any additional conditional-law filter obeys definition 11.37. These are explicit regularity restrictions on this problem; dimension alone would allow arbitrary Borel encodings of the full state or history. What holds is the exact alternative that retains the private coordinates with their conditional laws at the dimension of the block (theorem 11.42 in finite state), and the pointwise descent criterion of proposition 6.26(c) for a supplied $C^1$ interface path. The problem is settled by a parent of that kind for which the relative entropy of the fine path law of the communicated coordinates with respect to the parent's path law, on horizons of order $1/\epsilon_{\mathrm{M}}$, on which the interfaces move by an amount of order one, tends to zero as $\epsilon_{\mathrm{M}} \to 0$ uniformly over $\mathrm{T}$ in an interval $(0, \mathrm{T}_0]$; negatively, by a member of the class in which the learned interface path has a rest point at which no parent of that kind is projectable, or in which every parent of that kind has that relative entropy bounded below on those horizons as $\epsilon_{\mathrm{M}} \to 0$ and $\mathrm{T} \to 0$ along a stated relation.

**Physicist's summary.** Attention is inference over a hidden pointer: a categorical label names whom a receiver listens to, its prior's support is the interaction graph, and its record's likelihood falls with a comparison energy. Bayes' formula gives the softmax row, masks are zero prior weights, and the temperature is a free parameter of the record, of which the inverse-square-root key-width scaling is one choice. A recognition row held constant over the latent state is a different softmax, over averaged energies, at an exact expected relative-entropy cost that penalizes uncertain sources. Integrating the label out gives a log-sum-exp potential whose curvature is the frozen-weight curvature minus an attention covariance, which can turn a convex landscape concave, the seed of symmetry breaking in the two-timescale and formation chapters. Communication is a normalized encoder–channel–decoder kernel, covariant under separate recoordinatizations of each message space; it may be singular, and a singular round trip is processing, not holonomy. For a Gaussian comparison record the exact receiver update has the transformer shape, but the value map is fixed by the query and key maps, the residual is precision weighted, and the weights use the predictive covariance. Heads sharing one receiver state can have coupled labels; an exact recursion with bounded head count for the coupled class remains open. Coarse-graining carries the receiver–source event law, not the conditional row, and a coarse source set contains the block itself.

# Chapter 7

# Fiber Information Geometry and Inference Histories

This chapter asks how the collective free energy of chapter 5 selects an inference history of the agents at $c_*$ and what duration the Fisher geometry of their law fibers assigns to it. A chosen metric makes the positive ray of the natural gradient select an oriented orbit, whose Fisher length is an information clock invariant under reparameterization (theorem 7.36). A parameter-independent Markov channel projects scores and loses Fisher information by an exact conditional-covariance defect (theorem 7.8), so a same-path Markov image of a history, such as the law of a meta-agent, accumulates no more duration (corollary 7.42), while a coarse orbit recomputed from its own free energy compares with the fine one only through an oriented semiconjugacy, whose rate cancels from every duration (theorems 7.47 and 7.49). A three-agent parity witness shows that a collective free energy and its Fisher tensor are not functions of the agents' marginal laws. Chapter 8 uses the mobility and clock results, chapter 9 the Fisher-loss theorem, and chapters 11 and 12 the coarse configuration map and the record clocks. Longer proofs are deferred to appendix B.

**Scope.** Law fibers and kernel fibers are distinct, and no manifold, dominating measure, affine chart, cone, or Fisher metric is inferred from the words belief or model. The Fisher metric is chosen: no invariance axiom that characterizes it (proposition 7.4) is imposed, and no uniqueness is asserted. The parity hyperedge (definition 7.27) is supplied generative data, not derived from pairwise locality or a canonical agent inventory, and the promoted parity coordinate (corollary 7.26) is not thereby an agent or an intervention variable. No fine-to-coarse semiconjugacy is asserted for coarse maps that discard directions on which the fine free energy depends. Open are correlated lifts beyond the parity witness (open problem 7.29), global Fisher quotients and canonical block data (open problem 7.15), the quantitative meta-agent clock ratio (open problem 7.43), and the slow-manifold limit of the information-discarding case (open problem 8.7). [NOT-CLAIMED]

**Notation.** Tangent vectors are $Z, Z_1, Z_2, Z_3$, and $\mathrm{s}_Z$ is the score in the direction $Z$; $\xi_{\mathcal{M}}$ is the fundamental vector field of $\xi$ in the Lie algebra $\mathfrak{g}$ of $G$ on a $G$-manifold $\mathcal{M}$. A parameter-independent Markov kernel is $K$, with law map $K_\star p = pK$ and joint law $p \otimes K(dy, d\bar{y}) = p(dy)K(y, d\bar{y})$. A history is $Q$ with curve parameter $t$, $\mathrm{a}$ is a positive rate factor, $\mathcal{I}_6$ indexes the bits of the parity witness and $\mathrm{par}$ is its parity character, $\mathcal{N}(\mu, \varsigma^2)$ is the normal family, and upright letters $\mathrm{c}, \mathrm{c}_1, \mathrm{c}_2, \mathrm{q}$ are constants local to one statement.

## 7.1 Fiber statistical tensors

At $c_*$ a frame $u$ in the $G$-torsor $(\mathscr{P}_G)_{c_*}$ identifies each law fiber $\mathcal{B}_x$ with $(\mathcal{E}_x)_{c_*}$ through $\iota_u : \beta \mapsto [u, \beta]$, and $\iota_{u\cdot g} = \iota_u \circ \widehat{\rho}_x(g)$ by (2.7), so a tensor on $\mathcal{B}_x$ defines one on $(\mathcal{E}_x)_{c_*}$ exactly when it is invariant under the represented action.

**Hypothesis 7.1 (Regular channel models).** Each law fiber $\mathcal{B}_x$, $x \in \{b, m\}$, and each coarse fiber $\bar{\mathcal{B}}_x$ used below is a finite-dimensional smooth parametrized-measure model with smooth represented action, as in hypothesis 2.3, differentiable in quadratic mean with a positive-definite Fisher form; every score direction is integrable to the third power, the second- and third-order score moments depend smoothly on the point, and the model is dominated enough to differentiate the divergences of section 7.2 under the integral through third order. The hypothesis excludes singular statistical strata, families whose support moves without this differentiability, and law fibers given only by a chart without a sample-space action.

**Definition 7.2 (Fiber statistical tensors and the model-parameter Fisher tensor).** With $\mathrm{s}_Z$ the score of $\mathcal{B}_x$ at $p$ in the direction $Z \in T_p\mathcal{B}_x$, the Fisher and Amari–Chentsov tensors of the channel and, for the generative parameter $\theta$, structural data $X$, tangent vectors $Z_1, Z_2$ of the parameter space of $\theta$, and a probability law $\mathsf{r}(dX)$ on structural data, the model-parameter Fisher tensor of the evidence $p_{\theta,V}(o \mid X)$ and its average are

$$g^F_{x,p}(Z_1, Z_2) = \mathbb{E}_p[\mathbf{s}_{Z_1}\mathbf{s}_{Z_2}], \tag{7.1}$$

$$T^F_{x,p}(Z_1, Z_2, Z_3) = \mathbb{E}_p[\mathbf{s}_{Z_1}\mathbf{s}_{Z_2}\mathbf{s}_{Z_3}], \tag{7.2}$$

$$\mathsf{I}_X(\theta)(Z_1, Z_2) = \mathbb{E}_\theta\left[\partial_{Z_1} \log p_{\theta,V}(O \mid X)\partial_{Z_2} \log p_{\theta,V}(O \mid X) \,\middle|\, X\right], \tag{7.3}$$

$$\mathsf{I}_\mathsf{r}(\theta) = \int \mathsf{I}_X(\theta)\mathsf{r}(dX). \tag{7.4}$$

Each of $\mathsf{I}_X$ and $\mathsf{I}_\mathsf{r}$ is called a metric only where it is nondegenerate.

On the latent sample space Y an interaction energy is a function, not a metric; the law fibers and the configuration manifolds of section 7.4 carry $g^F$; the parameter space of $\theta$ carries $\mathsf{I}_X$ or $\mathsf{I}_\mathsf{r}$, of rank depending on $X$ or $\mathsf{r}$, one available metric for $\theta$ in the fast and slow coordinates of definition 8.2, whose model-channel recognition coordinates carry $g^F_m$. Equal components of a quadratic energy on Y and a Fisher metric, as in the Gaussian mean sector of Supplement G, are a coincidence of coordinates, not an identification.

**Proposition 7.3 (Descent of the statistical tensors).** Under hypothesis 7.1, every $\widehat{\rho}_x(g)$ preserves $g^F_x$ and $T^F_x$, so $(\iota_u^{-1})^* g^F_x$ and $(\iota_u^{-1})^* T^F_x$ are tensors on $(\mathcal{E}_x)_{c_*}$ that do not depend on the frame $u$; they are denoted by the same symbols. The sample-space origin of the action cannot be replaced by invariance of the metric alone: for the normal family, whose Fisher metric is $(d\mu^2 + 2d\varsigma^2)/\varsigma^2$, the map induced by $z \mapsto (1+z)/(1-z)$ on $z = \mu/\sqrt{2} + \mathrm{i}\varsigma$ preserves $g^F$, fixes $(\mu, \varsigma) = (0, 1)$, and does not preserve $T^F$. [ESTABLISHED]

The proof is in appendix B.2.

**Proposition 7.4 (Invariance characterizations of the Fisher metric).** On the open simplex of positive laws on a finite set, the Fisher metric is, up to a positive factor, the unique Riemannian metric invariant under congruent Markov embeddings [81, Lemma 11.3, p. 161]. On the cone of positive finite measures on a finite set, with coordinates $\mathsf{x}_i$ and total mass $|\mathsf{x}| = \sum_i \mathsf{x}_i$, the positive-definite metrics invariant under the same maps are exactly those with components $g_{\mathsf{x}}(\partial_i, \partial_j) = \mathrm{f}_2(|\mathsf{x}|) + \delta_{ij}|\mathsf{x}|\mathrm{f}_1(|\mathsf{x}|)/\mathsf{x}_i$ for smooth functions with $\mathrm{f}_1 > 0$ and $\mathrm{f}_1 + \mathrm{f}_2 > 0$ [15]; on tangent vectors of the simplex, where $\sum_i Z_i = 0$, the $\mathrm{f}_2$ term vanishes and the metric is $\mathrm{f}_1(1)$ times the Fisher metric. For parametrized measure models on general sample spaces, every local continuous quadratic form field invariant under sufficient statistics is a combination of the Fisher form and the square of the differential of total mass with coefficients that are functions of the total mass, and on statistical models, where the total mass is one, the Fisher form is the unique weakly continuous such field up to a constant [3, Main Theorem 2.10]; the finite case of this characterization is Campbell's [56, Section 1]. Uniqueness up to a constant also holds for strongly continuous forms that are monotone under statistics [56, Theorem 1.4], for diffeomorphism-invariant smooth weak Riemannian metrics on smooth positive densities on a closed manifold of dimension at least two [7, Theorem in the introduction], and, on exponential families and their derived families, for Riemannian metrics invariant under independent and identically distributed extensions and canonical sufficient statistics [28, Theorem 1]. [STANDARD]

## 7.2 Divergence jets and the orbit Fisher form

**Definition 7.5 (Regular invariant contrast and its jets).** A regular invariant contrast on a law fiber $\mathcal{B}$ is a function $\mathscr{D} : \mathcal{B} \times \mathcal{B} \to \mathbb{R}$, $C^3$ near the diagonal, with $\mathscr{D} \geq 0$ and, near the diagonal, $\mathscr{D}(p, q) = 0$ exactly when $p = q$, invariant under the simultaneous represented action, $\mathscr{D}(\widehat{\rho}(g)p, \widehat{\rho}(g)q) = \mathscr{D}(p, q)$, and with positive-definite mixed second jet. With $(Z, 0)$ and $(0, Z)$ the lifts of $Z \in T_p\mathcal{B}$ to the two factors, extended as constant coordinate fields in one chart, the jets of $\mathscr{D}$ are

$$g_{\mathscr{D}}(Z_1, Z_2) = -[(Z_1, 0)(0, Z_2)\mathscr{D}]_{(p,p)}, \tag{7.5}$$

$$\Gamma_{\mathscr{D}}(Z_1, Z_2, Z_3) = -[(Z_1, 0)(Z_2, 0)(0, Z_3)\mathscr{D}]_{(p,p)}, \tag{7.6}$$

$$\Gamma^*_{\mathscr{D}}(Z_1, Z_2, Z_3) = -[(0, Z_1)(0, Z_2)(Z_3, 0)\mathscr{D}]_{(p,p)}, \tag{7.7}$$

$$T^F_{\mathscr{D}} = \Gamma_{\mathscr{D}} - \Gamma^*_{\mathscr{D}}. \tag{7.8}$$

The two third-order arrays are connection coefficients in the chart, their difference is a symmetric cubic tensor. Invariance makes the metric and cubic jets define tensors, and the connection jets define affine connections, on the associated fiber independently of the frame. The contrast $(p, q) \mapsto (p - q)^4$ on a one-dimensional fiber is positive off the diagonal with zero second jet, so it is excluded.

**Proposition 7.6 (KL divergence jets).** On a fiber satisfying hypothesis 7.1, the relative entropy $D_{\mathrm{KL}}(p\|q)$ is a regular invariant contrast, its metric jet is $g^F$, and its cubic jet $T^F_{D_{\mathrm{KL}}}$ is the Amari–Chentsov tensor $T^F$. [ESTABLISHED]

**Corollary 7.7 (Orbit Fisher form).** Assume hypothesis 7.1 in one channel. For $q \in \mathcal{B}$ define

$$\mathfrak{k}_q(\xi,\eta) = g^F_q\big(\xi_{\mathcal{B}}(q), \eta_{\mathcal{B}}(q)\big), \qquad \mathfrak{g}_q = \{\xi \in \mathfrak{g} : \xi_{\mathcal{B}}(q) = 0\}. \tag{7.9}$$

Then $\mathfrak{k}_q$ is positive semidefinite, its radical is exactly $\mathfrak{g}_q$, and $\mathfrak{g}_q$ is the Lie algebra of the stabilizer $G_q$ of $q$. For every $\xi \in \mathfrak{g}$, as $t \to 0$,

$$D_{\mathrm{KL}}\left(q\|\widehat{\rho}(\exp t\xi)q\right) = \frac{t^2}{2}\mathfrak{k}_q(\xi,\xi) + O(t^3), \tag{7.10}$$

and for $\xi \in \mathfrak{g}_q$ the left side vanishes for every $t$. The same expansion holds for every regular invariant contrast, with $g^F$ replaced by its metric jet $g_{\mathscr{D}}$. [ESTABLISHED]

The proofs of proposition 7.6 and corollary 7.7 are in appendix B.2.

In the running example of section 2.5, a loop holonomy $\exp t\xi$ acting on an agent's law gives the loop a transported divergence $\frac{t^2}{2}\mathfrak{k}_q(\xi,\xi) + O(t^3)$ that vanishes to all orders exactly when $\xi \in \mathfrak{g}_q$, the second-order form of law coherence up to holonomy stabilizers (definition 2.17).

## 7.3 Markov coarse morphisms and the vertical Fisher defect

**Theorem 7.8 (Score projection and Fisher loss).** Let $\{P_\lambda\}$ be a family of laws on $\mathsf{Y}$, indexed by an open subset of $\mathbb{R}^d$, that is differentiable in quadratic mean at $\lambda_0$ with centered square-integrable score $\mathbf{s}_{\lambda_0} : \mathsf{Y} \to \mathbb{R}^d$, and let $K : \mathsf{Y} \rightsquigarrow \bar{\mathsf{Y}}$ be a Markov kernel that does not depend on $\lambda$ and satisfies $K(y,\bar{\mathsf{Y}}) = 1$ for every $y$. Under the joint law $P_{\lambda_0} \otimes K$, the coarse family $\{P_\lambda K\}$ is differentiable in quadratic mean at $\lambda_0$ with score

$$\bar{\mathbf{s}}_{\lambda_0}(\bar{Y}) = \mathbb{E}\left[\mathbf{s}_{\lambda_0}(Y) \mid \bar{Y}\right] \quad P_{\lambda_0}K\text{-almost surely}, \tag{7.11}$$

and its Fisher information obeys

$$\mathsf{I}_{\bar{\mathsf{Y}}}(\lambda_0) \preceq \mathsf{I}_{\mathsf{Y}}(\lambda_0), \qquad \mathsf{I}_{\mathsf{Y}}(\lambda_0) - \mathsf{I}_{\bar{\mathsf{Y}}}(\lambda_0) = \mathbb{E}\,\mathrm{Cov}\left(\mathbf{s}_{\lambda_0}(Y) \mid \bar{Y}\right) \succeq 0. \tag{7.12}$$

Equality in a direction $z \in \mathbb{R}^d$ holds exactly when $z^\top \mathbf{s}_{\lambda_0}(Y)$ is $\bar{Y}$-measurable. No common dominating measure is assumed. [ESTABLISHED]

The proof is in appendix B.2.

**Remark 7.9 (Fisher equality is not recovery).** Equality in (7.12) at one parameter value is local score sufficiency, not recovery of the experiment. Let $Y = (Y_1, Y_2)$ have independent Bernoulli coordinates with success probabilities $\frac{1}{2} + \frac{\lambda}{4}$ and $\frac{1}{2} + \frac{\lambda^2}{4}$, and let $K$ discard $Y_2$. At $\lambda = 0$ the score of

$Y_2$ vanishes, so the fine and coarse Fisher informations both equal $\frac{1}{4}$, yet no parameter-independent kernel $K'$ satisfies $(P_\lambda K)K' = P_\lambda$ for every $\lambda$ in a neighborhood of zero: applying any fixed $K'$ to the Bernoulli law $P_\lambda K$ makes every output event probability affine in $\lambda$, whereas the required output event $Y_2 = 1$ has probability $\frac{1}{2} + \frac{\lambda^2}{4}$. Kagan and Shepp [48] give an insufficient statistic that preserves Fisher information at every parameter value, and Pollard [68, Theorem 7] explains such examples by the criterion of theorem 7.8: preservation at $\lambda$ holds exactly when the score is almost surely a function of the statistic. [ESTABLISHED]

At $c_*$ the theorem applies to maps between fibers, for instance from the joint fiber of a block of agents to the fiber of their meta-agent.

**Definition 7.10 (Markov coarse morphism and its vertical defects).** Let $\bar{\mathcal{B}}_x$ be a second law fiber in channel $x$, with represented action $\widehat{\bar{\rho}}_x$ of $G$, associated bundle $\bar{\mathcal{E}}_x$, and frame maps $\bar{\iota}_u$. A Markov coarse morphism in channel $x$ at $c_*$ is a map $\mathscr{K} : (\mathcal{E}_x)_{c_*} \to (\bar{\mathcal{E}}_x)_{c_*}$ with $\bar{\iota}_u^{-1} \circ \mathscr{K} \circ \iota_u = K_\star|_{\mathcal{B}_x}$ for one, and hence every, frame $u$, where $K : \mathsf{Y}_x \rightsquigarrow \bar{\mathsf{Y}}_x$ is a Markov kernel that does not depend on the law parameter, satisfies $K(y, \bar{\mathsf{Y}}_x) = 1$ for every $y$, maps $\mathcal{B}_x$ smoothly into $\bar{\mathcal{B}}_x$, and is equivariant, $K_\star \circ \widehat{\rho}_x(g) = \widehat{\bar{\rho}}_x(g) \circ K_\star$ for every $g$; no common dominating measure is assumed. The vertical Fisher defect and the Amari defect of $\mathscr{K}$ are

$$\Delta_F^{\mathscr{K}} = g_x^F - (T\mathscr{K})^* \bar{g}_x^F, \tag{7.13}$$

$$\Delta_T^{\mathscr{K}} = T_x^F - (T\mathscr{K})^* \bar{T}_x^F. \tag{7.14}$$

Equivariance makes the frame representative independent of $u$, since $\bar{\iota}_{u\cdot g}^{-1} \mathscr{K} \iota_{u\cdot g} = \widehat{\bar{\rho}}_x(g)^{-1} K_\star \widehat{\rho}_x(g)$, and by proposition 7.3 both defects are tensors on $(\mathcal{E}_x)_{c_*}$.

Restriction to a submodel, precomposition of an energy with a map, a kernel that depends on the law parameter, a fitted coarse approximation, and a comparison across a mismatch of generative and recognition parameters are Markov coarse morphisms, and fall under the results below, only when one kernel that realizes them is exhibited, along the whole history when histories are compared. One kernel serves every law on one fixed-precision fiber of a Gaussian family; Supplement G shows that Gaussian Galerkin aggregation has this form on each fixed-precision fiber and no such form across precisions.

**Corollary 7.11 (Vertical Fisher defect).** For a Markov coarse morphism $\mathscr{K}$ and every tangent vector $Z$ at $p$, under the joint law $p \otimes K$,

$$\Delta_F^{\mathscr{K}} \succeq 0, \tag{7.15}$$

$$\Delta_F^{\mathscr{K}}(Z, Z) = \mathbb{E}\,\mathrm{Var}\left(\mathbf{s}_Z(Y) \mid \bar{Y}\right), \tag{7.16}$$

and $\Delta_F^{\mathscr{K}}(Z, Z) = 0$ exactly when $\mathbf{s}_Z(Y)$ is $\bar{Y}$-measurable. [ESTABLISHED]

*Proof.* In a frame $\mathscr{K}$ is $K_\star$. For a smooth curve $p_t$ in $\mathcal{B}_x$ with velocity $Z$ at $p$, theorem 7.8 gives the image curve $K_\star p_t$ the score $\mathbb{E}[\mathbf{s}_Z(Y) \mid \bar{Y}]$, which by uniqueness of the quadratic-mean score is

the score of $\bar{\mathcal{B}}_x$ in the direction $TK_\star Z$, with second moment $\bar{g}^F_x(TK_\star Z, TK_\star Z)$; the defect identity (7.12) for this one-parameter family is (7.16), with the stated equality clause. □

**Theorem 7.12 (Vertical Fisher defect cocycle).** Let $\mathscr{K}_{01} : (\mathcal{E}_0)_{c_*} \to (\mathcal{E}_1)_{c_*}$ and $\mathscr{K}_{12} : (\mathcal{E}_1)_{c_*} \to (\mathcal{E}_2)_{c_*}$ be smooth maps between fibers carrying Fisher tensors $g^F_0, g^F_1, g^F_2$. Then, unconditionally,

$$\Delta_F^{\mathscr{K}_{12}\circ\mathscr{K}_{01}} = \Delta_F^{\mathscr{K}_{01}} + (T\mathscr{K}_{01})^* \Delta_F^{\mathscr{K}_{12}}, \tag{7.17}$$

and the Amari defects obey the same identity. If both maps are Markov coarse morphisms, both terms of (7.17) are positive semidefinite. No nonzero Amari defect is sign definite, since $\Delta_T^{\mathscr{K}}(-Z,-Z,-Z) = -\Delta_T^{\mathscr{K}}(Z,Z,Z)$; Markov coarse-graining therefore gives no monotonicity for $T^F$, and $\Delta_T^{\mathscr{K}} = 0$ is an additional property of $\mathscr{K}$. [ESTABLISHED]

*Proof.* Contravariance of pullback gives $\Delta_F^{\mathscr{K}_{12}\circ\mathscr{K}_{01}} = g^F_0 - (T\mathscr{K}_{01})^* g^F_1 + (T\mathscr{K}_{01})^*[g^F_1 - (T\mathscr{K}_{12})^* g^F_2]$, which is (7.17), and the same computation applies to $T^F$. Positivity is corollary 7.11 for each Markov arrow, and oddness is homogeneity of degree three. □

## 7.4 Configuration manifolds and the exact free-energy lift

**Definition 7.13 (Configuration manifold at $c_*$).** The configuration manifold $\mathcal{Q}_i$ of agent $i$ at $c_*$ is the product $(\mathcal{E}_b)_{c_*} \times (\mathcal{E}_m)_{c_*}$ of its belief and model law fibers, identified with $\mathcal{B}_b \times \mathcal{B}_m$ through its frame $u_i$, or an embedded submanifold of that product (a regular submodel), with points $Q_i = (q^b_i, q^m_i)$. The collective configuration manifold is $\mathcal{Q}_V = \prod_{i\in V} \mathcal{Q}_i$, extended by further coordinate manifolds, such as open simplices of attention rows, when those coordinates move. A configuration manifold carries a metric $\mathsf{G}$ that must be specified. The channel-weighted metric of agent $i$ is $\mathsf{G}_i = w_b g^F_b \oplus w_m g^F_m$ with positive weights, under which a $C^1$ curve $Q_i(t)$ has squared speed

$$\|\dot{Q}_i\|^2_{\mathsf{G}_i} = w_b g^F_b(\dot{q}^b_i, \dot{q}^b_i) + w_m g^F_m(\dot{q}^m_i, \dot{q}^m_i); \tag{7.18}$$

on $\mathcal{Q}_V$ the metric is either block diagonal, $\bigoplus_i \mathsf{G}_i$, or the pullback of a joint Fisher metric along an exact lift (definition 7.17).

These metrics are frame independent by proposition 7.3; their base-direction version is in Supplement B.

**Theorem 7.14 (Fisher radical and quotient of a pulled-back tensor).** Let $\mathrm{ret} : \mathcal{Q} \to \mathcal{Q}'$ be a smooth map of finite-dimensional manifolds, $g$ a positive-semidefinite symmetric tensor on $\mathcal{Q}'$, and $h = \mathrm{ret}^* g$. Then

$$\operatorname{rad} h_Q = (d\,\mathrm{ret}_Q)^{-1}\left(\operatorname{rad} g_{\mathrm{ret}(Q)}\right), \tag{7.19}$$

so $\operatorname{rad} h_Q = \ker d\,\mathrm{ret}_Q$ exactly when $\operatorname{im} d\,\mathrm{ret}_Q \cap \operatorname{rad} g_{\mathrm{ret}(Q)} = \{0\}$. Where $\operatorname{rad} h$ has constant rank and $h$ is $C^1$, $\operatorname{rad} h$ is a $C^1$ subbundle and $h$ induces a positive-definite form on the quotient bundle $T\mathcal{Q}/\operatorname{rad} h$. A quotient manifold additionally requires the null leaves to form a Hausdorff manifold leaf space; when $\operatorname{rad} h = \ker d\,\mathrm{ret}$ and the connected leaves form such a space $\mathcal{Q}/\!\sim$ with quotient

projection a smooth submersion, $h$ is the pullback of a positive-definite tensor on $\mathcal{Q}/\!\sim$. For a positive $C^3$ family on a finite set, with $\mathrm{ret}$ the law map into the open simplex and $g$ its Fisher metric, $h$ is the Fisher tensor of the family, of class $C^2$, positive definite exactly where $\mathrm{ret}$ is an immersion, and there $\mathrm{ret}$ is locally injective and $h$ has a $C^1$ Levi-Civita connection. For softmax logits $z \in \mathbb{R}^k$ with $p = \mathrm{softmax}(z)$, $h(z) = \mathrm{diag}(p) - pp^\top$ has radical $\mathbb{R}\mathbf{1}$ at every $z$, and for an objective that depends on $z$ only through $p$ the solutions $Z$ of $hZ = \nabla\mathcal{F}$ form one class $Z_0 + \mathbb{R}\mathbf{1}$, whose Moore–Penrose representative $h^+\nabla\mathcal{F}$ is Euclidean-orthogonal to $\mathbf{1}$. [ESTABLISHED]

The proof is in appendix B.2.

**Open Problem 7.15 (Global Fisher quotients and canonical block data).** Theorem 7.14 gives a positive-definite quotient bundle on every constant-rank stratum and a quotient manifold once a Hausdorff leaf space is supplied, and proposition 11.46 shows that the law and its Fisher tensor alone select no splitting into agent blocks. Open are a criterion, checkable on the configuration manifolds of this chapter, under which the null foliation of a pulled-back Fisher tensor has a Hausdorff manifold leaf space, and block data on that quotient determined by a stated typing of agents and channels. Such a criterion with a proof, for instance exhibiting the leaves as the orbits of a free and proper group action, together with a construction of natural block data from the typing or a counterexample showing that the typing does not determine it, settles the problem.

**Hypothesis 7.16 (Standing configuration tier).** Wherever this chapter asserts a natural-gradient vector field or a Fisher duration on a configuration manifold $\mathcal{Q}$, either $\mathcal{Q}$ is a finite-dimensional configuration manifold as in definition 7.13, a finite product of regular statistical manifolds or an embedded submodel of one, with a $C^1$ positive-definite metric $\mathsf{G}$, or $\mathcal{Q}$ is the quotient of such a manifold by the null leaves of a $C^1$ positive-semidefinite Fisher tensor of constant rank with a Hausdorff manifold leaf space, as in theorem 7.14, and the objective is constant on those leaves. The objective $\mathcal{F}$ is $C^2$, and every minimal regular exponential family with its Fisher metric is an instance. The hypothesis excludes spaces of sections over a base, which belong to Supplement B, semidefinite metrics without a constructed quotient, and configuration metrics called exact joint-law Fisher metrics without the lift of definition 7.17.

**Definition 7.17 (Exact multi-agent free-energy lift).** Fix a block $B$ of latent coordinates, an outside recognition marginal $\mathbb{Q}_{B^c}$, and a manifold $\mathscr{R}_B$ of admissible conditional recognition kernels $r_B(dy_B \mid y_{B^c})$. An exact lift of the configuration manifold $\mathcal{Q}_i$ consists of a smooth extraction map $\mathrm{ext}_i : \mathscr{R}_B \to \mathcal{Q}_i$ and a smooth right inverse $\iota_i : \mathcal{Q}_i \to \mathscr{R}_B$, $Q_i \mapsto r_B^{Q_i}$, $\mathrm{ext}_i \circ \iota_i = \mathrm{id}_{\mathcal{Q}_i}$, such that, for every $Q_i$, the support, finiteness, and integrability hypotheses of definition 5.21 and theorem 5.23 hold for $\mathbb{Q}_{B^c} r_B^{Q_i}$, the lifted family is $C^2$, and for every chart $\mathrm{ch}$ of $\mathcal{Q}_i$ with coordinates $\lambda$ and every compact set $\mathsf{C}$ in its image there is $\mathrm{e}_{\mathsf{C}} \in L^1(\mathbb{Q}_{B^c})$ with, for jointly measurable versions of the derivatives,

$$\sup_{\lambda \in \mathsf{C}} \left| \partial_\lambda^\alpha \mathcal{F}^{\mathrm{ext}}_{B,o}(r_B^{\mathrm{ch}^{-1}(\lambda)}; y_{B^c}) \right| \le \mathrm{e}_{\mathsf{C}}(y_{B^c}) \quad \text{for } \mathbb{Q}_{B^c}\text{-almost every } y_{B^c} \text{ and } |\alpha| \le 2. \tag{7.20}$$

The lifted objective is the outside-averaged local free energy

$$\overline{\mathcal{F}}_{B,o}(Q_i) = \mathbb{E}_{\mathbb{Q}_{B^c}}\left[\mathcal{F}^{\mathrm{ext}}_{B,o}\left(r_B^{Q_i}; Y_{B^c}\right)\right], \tag{7.21}$$

a history of it is an exact free-energy history, and the exact recognition Fisher metric of $\mathcal{Q}_i$ is $\mathsf{G}_i = \iota_i^* \mathsf{G}_{\mathscr{R}_B}$ on a nondegenerate stratum, or its quotient as in theorem 7.14.

The lift carries every correlation that the configuration does not determine, and $\mathrm{ext}_i \circ \iota_i = \mathrm{id}$ makes it a lift rather than a family of laws indexed by configurations. With the outside marginal held fixed, (7.21) differs from the collective free energy along $Q_i \mapsto \mathbb{Q}_{B^c} r_B^{Q_i}$ by a term independent of $Q_i$ (theorem 5.23), so the two have the same differential and, for one metric, the same natural-gradient ray.

**Proposition 7.18 (Joint-versus-weighted-marginal score residual and lift dependence).** Let $P_\lambda$ be a regular dominated family of joint laws of blocks $(Y_1, \dots, Y_n)$ with square-integrable joint score $\mathbf{s} = \nabla_\lambda \log p_\lambda(Y)$, embedded marginal scores $\mathbf{s}_i = \mathbb{E}[\mathbf{s} \mid Y_i]$, and scalar weights $w_i$, and define

$$\mathsf{G}_{\mathrm{J}} = \mathbb{E}[\mathbf{s}\mathbf{s}^\top], \qquad \mathsf{G}_i = \mathbb{E}[\mathbf{s}_i\mathbf{s}_i^\top], \qquad \mathsf{G}_w = \sum_i w_i \mathsf{G}_i, \tag{7.22}$$

$$\mathbf{r} = \mathbf{s} - \sum_i \mathbf{s}_i, \qquad \mathfrak{K} = \mathbb{E}[\mathbf{r}\mathbf{r}^\top], \qquad \mathfrak{C} = \sum_{i \neq j} \mathbb{E}[\mathbf{s}_i\mathbf{s}_j^\top]. \tag{7.23}$$

The cross tensor $\mathfrak{C}$ is symmetric and can be indefinite,

$$\mathsf{G}_{\mathrm{J}} - \mathsf{G}_w = \mathfrak{K} - \mathfrak{C} + \sum_i (1 - w_i)\mathsf{G}_i, \tag{7.24}$$

and on a tangent subspace $\mathcal{T}$, $\mathsf{G}_{\mathrm{J}} = \mathsf{G}_w$ holds exactly when the right side of (7.24) vanishes on $\mathcal{T} \times \mathcal{T}$. For two blocks and unit weights, in a direction $Z$ with scores $\mathbf{s}_Z = Z^\top \mathbf{s}$ and $\mathbf{s}_{i,Z} = Z^\top \mathbf{s}_i$,

$$\|\mathbf{s}_Z\|^2 - \|\mathbf{s}_{1,Z}\|^2 - \|\mathbf{s}_{2,Z}\|^2 = \|\mathbf{s}_Z - \mathbf{s}_{1,Z} - \mathbf{s}_{2,Z}\|^2 - 2\langle \mathbf{s}_{1,Z}, \mathbf{s}_{2,Z}\rangle, \tag{7.25}$$

whose first term is nonnegative and whose second is signed. If the two blocks are independent under every $P_\lambda$, then $\mathsf{G}_{\mathrm{J}} = \mathsf{G}_1 + \mathsf{G}_2$, and $\mathsf{G}_{\mathrm{J}} = w_1\mathsf{G}_1 + w_2\mathsf{G}_2$ forces $w_1 = 1$ as soon as some direction has $\mathsf{G}_1(Z,Z) > 0 = \mathsf{G}_2(Z,Z)$, and symmetrically for $w_2$. No Loewner order between $\mathsf{G}_{\mathrm{J}}$ and $\mathsf{G}_1 + \mathsf{G}_2$ holds in general: for the bivariate normal location family with precision $\left(\begin{smallmatrix} 1 & \mathrm{c} \\ \mathrm{c} & 1 \end{smallmatrix}\right)$ and $0 < \mathrm{c} < 1$, the difference is $\left(\begin{smallmatrix} \mathrm{c}^2 & \mathrm{c} \\ \mathrm{c} & \mathrm{c}^2 \end{smallmatrix}\right)$, with value $2\mathrm{c}(1+\mathrm{c}) > 0$ along $(1,1)$ and $2\mathrm{c}(\mathrm{c}-1) < 0$ along $(1,-1)$. Distinct right inverses of one extraction map can induce distinct pullback metrics: for $\mathcal{N}((\lambda_1,\lambda_2), \mathrm{I}_2)$, with Fisher metric $d\lambda_1^2 + d\lambda_2^2$, and the extraction $(\lambda_1,\lambda_2) \mapsto \lambda_1$, the lifts $\lambda_1 \mapsto (\lambda_1, 0)$ and $\lambda_1 \mapsto (\lambda_1, 1)$ induce $d\lambda_1^2$, whereas $\lambda_1 \mapsto (\lambda_1, \lambda_1)$ induces $2d\lambda_1^2$. [ESTABLISHED]

The proof is in appendix B.2.

### 7.4.1 Canonicity boundary of the lift

A lift supplies the dependence that the marginal configurations do not determine. The following finite results show that naturality under local kernels forces the product law and that admitting correlated refinements leaves no natural selector.

**Definition 7.19 (Finite marginal selection problem).** For a finite nonempty typed list $\mathcal{X} = (\mathcal{X}_i)_{i=1}^n$ of finite sets, let $\mathcal{J}_\mathcal{X} = \Delta(\prod_i \mathcal{X}_i)$, $\mathcal{J}^1_\mathcal{X} = \prod_i \Delta(\mathcal{X}_i)$, and $\mathrm{mar}_\mathcal{X} : \mathcal{J}_\mathcal{X} \to \mathcal{J}^1_\mathcal{X}$ the map that extracts the one-coordinate marginals. A selector is a map $\mathrm{sel}_\mathcal{X} : \mathcal{J}^1_\mathcal{X} \to \mathcal{J}_\mathcal{X}$ with $\mathrm{mar}_\mathcal{X} \circ \mathrm{sel}_\mathcal{X} = \mathrm{id}$. A local morphism is a tensor product $K = \bigotimes_i K_i$ of finite Markov kernels of the same arity, and a kernel $K : \prod_i \mathcal{X}_i \rightsquigarrow \prod_j \mathcal{Y}_j$ that changes arity is marginal compatible when $\mathrm{mar}_\mathcal{Y} K_\star = \bar{K} \mathrm{mar}_\mathcal{X}$ for some map $\bar{K} : \mathcal{J}^1_\mathcal{X} \to \mathcal{J}^1_\mathcal{Y}$. A family of selectors is natural on a class of morphisms when $K_\star \mathrm{sel}_\mathcal{X} = \mathrm{sel}_\mathcal{Y} \bar{K}$ for every morphism in the class; the class is part of the data, and not every finite kernel is admitted.

**Theorem 7.20 (Tensor-local naturality selects the product law).** If a family of selectors is natural under every local morphism, then $\mathrm{sel}_\mathcal{X}(\mu_1, \dots, \mu_n) = \bigotimes_{i=1}^n \mu_i$. Among classes of marginal-compatible kernels that contain every local morphism, a natural selector exists exactly when every admitted kernel sends product laws to the product of their image marginals, and then it is the product selector; the product-preserving marginal-compatible kernels therefore form the largest such class. [ESTABLISHED]

*Proof.* For each $i$ let the kernel from a one-point set send its point to $\mu_i$. Their tensor product sends the unique law on the one-point product to $\bigotimes_i \mu_i$, while the induced marginal map sends the unique marginal tuple to $(\mu_i)_i$, so naturality forces the product selector. With it forced, naturality for an added kernel says exactly that it sends each product law to the product of its image marginals. □

**Theorem 7.21 (Admitted correlated refinements refute an absolute selector).** For a source bit $y$ and target bits $(y', y'')$, admit the split kernels $K_\mathrm{c}^{\mathrm{split}}(y', y'' \mid y) = \mathbf{1}_{\{y'=y\}}(1 + \mathrm{c})/2$ if $y'' = y$ and $\mathbf{1}_{\{y'=y\}}(1 - \mathrm{c})/2$ if $y'' \neq y$, for $\mathrm{c} \in \{1/3, 1/2\}$. Both are marginal compatible, and no single-valued selector is natural under both; the argument uses no local naturality. In particular marginalization is not injective on positive laws, so no selector returns every joint law from its marginals. [ESTABLISHED]

*Proof.* The two kernels push the fair source law to laws with fair marginals and atom multisets $\{1/3, 1/3, 1/6, 1/6\}$ and $\{3/8, 3/8, 1/8, 1/8\}$. Naturality would require the value of the target selector at the fair marginal pair to equal both laws, which differ even after relabeling, and the two positive laws share their marginals. □

Selection relative to a positive reference law is a different problem. A finite-relative-entropy minimizer over a convex family of moment-constrained laws, when it exists, is unique by strict convexity of relative entropy. Both that minimizer and, for a deterministic coarse map, the reference conditional completion of a prescribed coarse marginal (lemma 3.12) depend on the reference law as additional data.

## 7.5 A finite correlated lift: the parity witness

**Definition 7.22 (Three-agent parity lift).** Let $\mathcal{I}_6 = \{1,2,3\} \times \{b,m\}$, $y = (y_{i,x}) \in \{0,1\}^{\mathcal{I}_6}$, and $\lambda = (\lambda_{i,x}) \in (0,1)^{\mathcal{I}_6}$, where typed agent permutations preserve the channel label, and set

$$p_0(y;\lambda) = \prod_{(i,x)} \lambda_{i,x}^{y_{i,x}} (1-\lambda_{i,x})^{1-y_{i,x}}, \qquad \mathrm{D}(\lambda) = \prod_{(i,x)} \lambda_{i,x}(1-\lambda_{i,x}), \tag{7.26}$$

$$p_\varkappa(y;\lambda) = p_0(y;\lambda) + \varkappa\,\mathrm{par}(y)\mathrm{D}(\lambda), \qquad \mathrm{par}(y) = (-1)^{\sum_{(i,x)} y_{i,x}}, \qquad |\varkappa| < 1. \tag{7.27}$$

The map $\iota_\varkappa : (0,1)^6 \to \Delta_{63}^\circ$ sends $\lambda$ to $p_\varkappa(\cdot;\lambda)$. The parity amplitude $\varkappa$ is fixed data of the lift, not an inferred coordinate, and its hyperedge correlation is $\bar{\varkappa} = \varkappa/64$.

**Theorem 7.23 (Positive smooth lift with paired relabeling symmetry).** For every $|\varkappa| < 1$, the law (7.27) is strictly positive and normalized, and every proper marginal is the corresponding marginal of $p_0$. Consequently $\iota_\varkappa$ is a smooth right inverse of the map that extracts the six one-bit marginals, and it is an immersion. For each marginal tuple $\lambda$, the positive laws with those marginals form a $57$-dimensional set, of which varying $\varkappa$ traces one line. Let $H_{\mathrm{pair}} = (\mathbb{Z}_2)^3 \rtimes S_3$ act by complementing both bits of a selected agent and by typed agent permutations, and act on $\lambda$ by the same complements and permutations. Then

$$\iota_\varkappa(g\lambda) = g_\# \iota_\varkappa(\lambda) \qquad \text{for every } g \in H_{\mathrm{pair}}, \tag{7.28}$$

whereas the complement $F_{i,x}$ of one bit satisfies

$$(F_{i,x})_\# p_\varkappa(\cdot;\lambda) = p_{-\varkappa}(\cdot; F_{i,x}\lambda). \tag{7.29}$$

Odd complements are therefore equivariant only if $\varkappa$ changes sign with them, or if $\varkappa = 0$. [ESTABLISHED]

The proof is in appendix B.2.

**Corollary 7.24 (Agent-block local–global free-energy identity).** Fix an agent block $B_i = \{(i,b),(i,m)\}$ and the four outside parameters $\lambda_{B_i^c}$, and let $r_{i,\lambda_{B_i}} = p_\varkappa(y;\lambda)/p_{0,B_i^c}(y_{B_i^c};\lambda_{B_i^c})$ be the conditional law of $y_{B_i}$ given $y_{B_i^c}$. For any interior block parameters $\lambda_{B_i}$ and $\lambda'_{B_i}$ and any strictly positive posterior $\mathbf{\Pi}_o$ on the Boolean cube, all relative entropies are finite, and for every block tangent $Z$

$$\begin{aligned} &\mathcal{F}_o^{\mathrm{ext}}\big(p_\varkappa(\cdot;\lambda'_{B_i},\lambda_{B_i^c})\big) - \mathcal{F}_o^{\mathrm{ext}}\big(p_\varkappa(\cdot;\lambda_{B_i},\lambda_{B_i^c})\big) \\ &\qquad = \mathbb{E}_{p_{0,B_i^c}}\big[\mathcal{F}_{B_i,o}^{\mathrm{ext}}(r_{i,\lambda'_{B_i}};Y_{B_i^c}) - \mathcal{F}_{B_i,o}^{\mathrm{ext}}(r_{i,\lambda_{B_i}};Y_{B_i^c})\big], \end{aligned} \tag{7.30}$$

$$d\mathcal{F}_o^{\mathrm{ext}}[Z] = \mathbb{E}_{p_{0,B_i^c}}\big[d\mathcal{F}_{B_i,o}^{\mathrm{ext}}[Z]\big], \tag{7.31}$$

with both differentials evaluated at the same lifted conditional law. The outside-averaged block free energy is thus the restriction of one collective free energy up to a constant. [ESTABLISHED]

*Proof.* By theorem 7.23 the outside marginal is exactly $p_{0,B_i^c}$, independent of $\lambda_{B_i}$, and strict positivity makes $r_{i,\lambda_{B_i}}$ a positive conditional law. The finite state space and the positive posterior make every relative entropy finite, so (7.30) is theorem 5.23 with this outside marginal, and differentiating the finite smooth identity gives (7.31). □

**Corollary 7.25 (Parity-family Fisher residual).** For the family of definition 7.22, every pair of bits is independent, so $\mathfrak{C} = 0$, and with unit weights $\mathsf{G}_{\mathrm{J}} - \sum_{(i,x)\in\mathcal{I}_6} \mathsf{G}_{i,x} = \mathfrak{K} \succ 0$ when $\varkappa \neq 0$. At $\lambda = \frac{1}{2}\mathbf{1}$, and there with $\varkappa = \frac{1}{2}$,

$$\mathsf{G}_{\mathrm{J}} = \frac{4}{1-\bar{\varkappa}^2}\mathrm{I}_6, \qquad \mathsf{G}_{\mathrm{J}} - \sum_{(i,x)\in\mathcal{I}_6} \mathsf{G}_{i,x} = \frac{4\bar{\varkappa}^2}{1-\bar{\varkappa}^2}\mathrm{I}_6, \tag{7.32}$$

$$\mathsf{G}_{\mathrm{J}} = \frac{65536}{16383}\mathrm{I}_6, \qquad \mathsf{G}_{\mathrm{J}} - \sum_{(i,x)\in\mathcal{I}_6} \mathsf{G}_{i,x} = \frac{4}{16383}\mathrm{I}_6, \tag{7.33}$$

and at the center a weighted marginal metric equals the joint one on the full tangent space exactly when

$$w_{i,x} = \frac{1}{1-\bar{\varkappa}^2} \qquad \text{for every } (i,x) \in \mathcal{I}_6, \tag{7.34}$$

so unit weights are exact only at $\varkappa = 0$. Since the one-bit marginal families do not depend on $\varkappa$, no rule that reads only them returns every joint Fisher tensor. [ESTABLISHED]

The proof is in appendix B.2.

**Corollary 7.26 (Promoted parity has one hidden direction).** For the seven-parameter family $(\lambda, \varkappa) \mapsto p_\varkappa(\cdot;\lambda)$ on $(0,1)^6 \times (-1,1)$, the derivative of the joint law has rank seven everywhere, while the one-bit marginal map has rank six and

$$\ker d\,(\mathrm{mar} \circ p) = \mathrm{span}\{\partial_\varkappa\}. \tag{7.35}$$

Thus $\varkappa$ is identifiable from the joint law and invisible to the one-bit marginals. [ESTABLISHED]

*Proof.* The one-bit success probabilities are exactly $\lambda$, so a vanishing joint derivative has zero $\lambda$-component, and the remaining derivative is a multiple of $\partial_\varkappa p_\varkappa = \mathrm{par}\,\mathrm{D}$, which is nowhere zero on the open domain; hence the joint derivative is injective. The same identity gives rank six for the marginal map and (7.35). □

At the center $\lambda = \frac{1}{2}\mathbf{1}$, where $\mathrm{D} = 4^{-6}$, the lift (7.27) is

$$p_\varkappa\big(y; \tfrac{1}{2}\mathbf{1}\big) = \tfrac{1}{64}\left(1 + \bar{\varkappa}\,\mathrm{par}(y)\right). \tag{7.36}$$

**Definition 7.27 (Three-agent parity hyperedge).** At the center (7.36), take the uniform latent prior $\mathrm{Unif}(y) = 1/64$ and the binary record kernel

$$K_{\mathrm{hyp}}(1 \mid y) = \tfrac{1}{2}\left(1 + \bar{\varkappa}\,\mathrm{par}(y)\right), \qquad K_{\mathrm{hyp}}(0 \mid y) = \tfrac{1}{2}\left(1 - \bar{\varkappa}\,\mathrm{par}(y)\right). \tag{7.37}$$

**Proposition 7.28 (Exact posterior and tight free energy for the hyperedge).** The record $o = 1$ has evidence $\frac{1}{2}$, and its posterior is exactly $p_{\varkappa}(\cdot; \frac{1}{2}\mathbf{1})$. For the free energy $\mathcal{F}_1(Q) = D_{\mathrm{KL}}(Q \| \mathrm{Unif}) - \mathbb{E}_Q \log K_{\mathrm{hyp}}(1 \mid Y)$ the posterior attains the tight value $\mathcal{F}_1 = \log 2$, the product law of its one-bit marginals is $\mathrm{Unif}$, and

$$\mathcal{F}_1(\mathrm{Unif}) - \log 2 = -\tfrac{1}{2}\log(1 - \bar{\varkappa}^2). \tag{7.38}$$

Consequently, when $\varkappa \neq 0$, no function of the one-bit marginals of the recognition law, the one-bit marginals of the posterior, and the evidence equals the free energy on every recognition law. [ESTABLISHED]

The proof is in appendix B.2.

**Open Problem 7.29 (Scope beyond the finite parity witness).** Theorem 7.23 gives one smooth correlated right inverse of the one-bit marginal map, equivariant under paired complements and typed permutations, and theorem 7.21 refutes an absolute selector on the admitted split kernels, so that question is closed. Open are the existence, for a stated regular family of agent laws and a stated class of morphisms that contains non-product-preserving kernels, of a smooth right inverse natural under that class, and the existence of correlated lifts equivariant under a connected structure group acting on continuous law fibers. A construction with proof for a stated family and class, or a no-go theorem for it, settles each part.

## 7.6 Natural-gradient orbits

**Definition 7.30 (Oriented unparameterized history and natural-gradient history).** A representative of a history in a configuration manifold $\mathcal{Q}$ is a $C^1$ curve $Q : \mathcal{T} \to \mathcal{Q}$ on an open interval with nowhere-vanishing velocity, and two representatives define the same oriented history when

$$\widetilde{Q} = Q \circ \phi, \qquad \phi : \widetilde{\mathcal{T}} \longrightarrow \mathcal{T} \text{ an orientation-preserving } C^1 \text{ diffeomorphism.} \tag{7.39}$$

The resulting abstract oriented curve $\mathscr{H}$ comes with its traversal map $\jmath : \mathscr{H} \to \mathcal{Q}$, which retains order and multiplicity; for an embedded representative $\mathscr{H}$ is its oriented image. Under hypothesis 7.16 the natural gradient of $\mathcal{F}$ is defined by the first line below, the descent ray at a noncritical $Q$ is $\mathcal{R}^-_{\mathcal{F}}(Q) = \{-\mathsf{a}\,\mathrm{grad}^{\mathsf{G}}\,\mathcal{F}(Q) : \mathsf{a} > 0\}$, and a natural-gradient history is an oriented history whose representatives satisfy the second:

$$\mathsf{G}\big(\mathrm{grad}^{\mathsf{G}}\,\mathcal{F}, Z\big) = d\mathcal{F}[Z] \quad \text{for every } Z \in T\mathcal{Q}, \tag{7.40}$$

$$\dot{Q}(t) = -\mathsf{a}(t)\,\mathrm{grad}^{\mathsf{G}}\,\mathcal{F}(Q(t)), \qquad \mathsf{a}(t) > 0. \tag{7.41}$$

When $\mathcal{F}$ is the lifted objective of definition 7.17, such a history is an exact free-energy history.

**Proposition 7.31 (Scalar mobility preserves the oriented orbit).** On a noncritical domain, multiplying the natural-gradient field by a smooth positive function changes only the parameterization of its orbits. A positive-definite anisotropic mobility can change the orbit: on the Euclidean plane,

for $\mathcal{F} = \frac{1}{2}(z_1^2 + z_2^2)$ and $\mathsf{M} = \mathrm{diag}(1, 2)$, gradient descent from $(z_1^0, z_2^0)$ with $z_1^0 z_2^0 \neq 0$ follows the line $z_2 = (z_2^0/z_1^0)z_1$, whereas $\dot{z} = -\mathsf{M}\nabla\mathcal{F}$ follows

$$z_2 = z_2^0 \left(z_1/z_1^0\right)^2 . \tag{7.42}$$

A natural-gradient history is not in general a geodesic: on the same Euclidean plane the natural-gradient flow of $\mathcal{F} = \frac{1}{2}(z_1^2 + 2z_2^2)$ is $\dot{z}_1 = -z_1$, $\dot{z}_2 = -2z_2$, whose orbits are the curves (7.42), while the geodesics are straight lines. [ESTABLISHED]

*Proof.* If $\mathscr{V} = -\operatorname{grad}^{\mathsf{G}} \mathcal{F}$ and $\mathscr{V}' = \mathsf{a}\mathscr{V}$ with $\mathsf{a} > 0$, the integral curves have the same tangent lines and orientation, and the change of parameter defined by a positive integral of $\mathsf{a}^{-1}$ carries one equation into the other. In the example, $\dot{z}_1 = -z_1$ and $\dot{z}_2 = -2z_2$ keep $z_2/z_1^2$ constant, which is (7.42), while ordinary gradient descent keeps $z_2/z_1$ constant, and for $z_1^0 z_2^0 \neq 0$ the images differ; the same equations are the Euclidean natural-gradient flow of $\frac{1}{2}(z_1^2 + 2z_2^2)$. □

A rate ratio between coordinate blocks, such as fast beliefs and slow models, is an anisotropic mobility (the example is the ratio $1 : 2$), so the selected orbit can depend on it; the rate-separated flow is in section 8.2.

**Proposition 7.32 (Noncritical free-energy orbits are embedded).** Let $\mathcal{Q}$ be finite-dimensional, $\mathcal{F}$ of class $C^2$, and $\mathsf{M}$ a $C^1$ field of positive-definite symmetric forms on the cotangent spaces, for instance $\mathsf{M} = \mathsf{a}\mathsf{G}^{-1}$ for a $C^1$ metric and a $C^1$ function $\mathsf{a} > 0$. Let $\gamma$ be a maximal integral curve of $\mathscr{V} = -\mathsf{M}(d\mathcal{F})$ through a noncritical point. Then $\gamma$ meets no critical point, $\mathcal{F} \circ \gamma$ is strictly decreasing, and $\gamma$ is an injective immersion that is a homeomorphism onto its image, an embedded $C^1$ curve on which $\mathcal{F}$ is a strictly monotone continuous coordinate, independent of parameterization and gauge invariant whenever $\mathcal{F}$ is. In particular the multiplicity clause of definition 7.30 is vacuous for noncritical histories, and for the natural gradient the Fisher clock of definition 7.35 satisfies $d\mathfrak{t} = -d\mathcal{F}/\|\operatorname{grad}^{\mathsf{G}} \mathcal{F}\|_{\mathsf{G}}$ along the orbit. [ESTABLISHED]

*Proof.* The field $\mathscr{V}$ is $C^1$, so its integral curves are unique, and its zeros are the critical points of $\mathcal{F}$; since the constant curve at a zero is an integral curve, a curve through a nonzero of $\mathscr{V}$ never reaches one. Along $\gamma$, $\frac{d}{dt}\mathcal{F}(\gamma) = -\mathsf{M}(d\mathcal{F}, d\mathcal{F}) < 0$, which makes $\mathcal{F} \circ \gamma$ strictly decreasing, $\gamma$ injective, and $\dot{\gamma} \neq 0$. On the image the inverse of $\gamma$ is $Q \mapsto (\mathcal{F} \circ \gamma)^{-1}(\mathcal{F}(Q))$, a composite of continuous maps, so $\gamma$ is an embedding. The clock relation is (7.52). □

**Proposition 7.33 (A semidefinite tensor cannot be silently inverted).** If the proposed metric on a configuration space is only semidefinite, the natural-gradient equation can have no solution or several. [ESTABLISHED]

*Proof.* On $\mathbb{R}^2$ let $\mathsf{G} = dz_1^2$. For $\mathcal{F} = z_2$ the equation $\mathsf{G}(\operatorname{grad}\mathcal{F}, \cdot) = dz_2$ has no solution because the left side annihilates $\partial_{z_2}$, and for $\mathcal{F} = z_1$ every $\partial_{z_1} + \mathrm{c}\partial_{z_2}$ solves it. □

A quotient constructed as in theorem 7.14, or a separate regularization, must precede any use of an inverse or a pseudoinverse.

**Remark 7.34 (Discrete histories).** In a chart where the metric has a constant coefficient matrix $\mathsf{G}$, the explicit Euler step of (7.41) with $\mathsf{a} \equiv 1$ and step $\Delta t > 0$ is $Q_{n+1} = Q_n - \Delta t \mathsf{G}^{-1} \nabla \mathcal{F}(Q_n)$. If $\nabla^2 \mathcal{F} \preceq \mathrm{c}_{\mathsf{G}} \mathsf{G}$ on the segment $[Q_n, Q_{n+1}]$, with $\mathrm{c}_{\mathsf{G}} > 0$, then

$$\mathcal{F}(Q_{n+1}) \leq \mathcal{F}(Q_n) - \Delta t \left(1 - \frac{\mathrm{c}_{\mathsf{G}} \Delta t}{2}\right) \left\| \mathrm{grad}^{\mathsf{G}} \mathcal{F}(Q_n) \right\|_{\mathsf{G}}^2, \tag{7.43}$$

so every step with $0 < \Delta t < 2/\mathrm{c}_{\mathsf{G}}$ decreases the objective off the critical set, and the bound is sharp: for a quadratic objective whose Hessian has largest generalized eigenvalue $\mathrm{c}_{\mathsf{G}}$ relative to $\mathsf{G}$, an iterate on that eigendirection is multiplied by $1 - \Delta t \mathrm{c}_{\mathsf{G}}$, so $\Delta t > 2/\mathrm{c}_{\mathsf{G}}$ increases the objective there. If, in the same chart, $\mathscr{V} = -\mathsf{G}^{-1} \nabla \mathcal{F}$ is Lipschitz with constant $\mathrm{c}_{\mathrm{L}} > 0$ on the closed $\delta$-neighborhood of an exact arc $Q([0, t_1])$ and $\|\ddot{Q}\| \leq \mathrm{c}_2$ on the arc, then

$$\|Q_n - Q(n\Delta t)\| \leq \frac{\mathrm{c}_2 \Delta t}{2\mathrm{c}_{\mathrm{L}}} \left(e^{\mathrm{c}_{\mathrm{L}} n \Delta t} - 1\right), \qquad n\Delta t \leq t_1, \tag{7.44}$$

for every $\Delta t$ small enough that the right side stays below $\delta$ at $n\Delta t = t_1$. Euler iterates follow the selected orbit to first order in $\Delta t$ on compact parameter intervals without lying on it. [ESTABLISHED]

*Proof.* Taylor's formula with integral remainder along the segment gives $\mathcal{F}(Q_{n+1}) \leq \mathcal{F}(Q_n) + \nabla \mathcal{F}^\top \Delta Q + \frac{1}{2} \mathrm{c}_{\mathsf{G}} \Delta Q^\top \mathsf{G} \Delta Q$ with $\Delta Q = -\Delta t \mathsf{G}^{-1} \nabla \mathcal{F}$, which is (7.43); sharpness is the stated eigendirection computation. The local truncation error is at most $\mathrm{c}_2 \Delta t^2/2$, so $e_n = \|Q_n - Q(n\Delta t)\|$ obeys $e_{n+1} \leq (1 + \Delta t \mathrm{c}_{\mathrm{L}}) e_n + \mathrm{c}_2 \Delta t^2/2$ while the iterates stay in the neighborhood, the discrete Grönwall inequality gives (7.44), and induction on $n$ keeps the iterates in the neighborhood. □

## 7.7 Fisher duration

**Definition 7.35 (Fisher length and information clock).** Let $Q : \mathcal{T} \to \mathcal{Q}$ be a $C^1$ curve in a configuration manifold with metric $\mathsf{G}$. At $c_*$ its velocity lies in the fibers, $\mathrm{ver}(\dot{Q}) = \dot{Q} \in T_Q \mathcal{Q}$, so no connection enters. Its Fisher speed, its length on $[t_0, t_1]$, its cumulative clock from a chosen origin $t_0$, and, for a curve $\Gamma$ in one law fiber with its Fisher metric, the vertical Fisher length are

$$\dot{\mathfrak{t}}(t) = \sqrt{\mathsf{G}\big(\dot{Q}(t), \dot{Q}(t)\big)}, \tag{7.45}$$

$$\mathrm{len}_{\mathsf{G}}(Q; [t_0, t_1]) = \int_{t_0}^{t_1} \dot{\mathfrak{t}}(t) dt, \tag{7.46}$$

$$\mathfrak{t}_{Q,t_0}(t) = \int_{t_0}^{t} \dot{\mathfrak{t}}(t') dt', \tag{7.47}$$

$$\mathrm{len}_{g^F}(\Gamma) = \int \sqrt{g^F\big(\mathrm{ver}(\dot{\Gamma}), \mathrm{ver}(\dot{\Gamma})\big)} dt. \tag{7.48}$$

The belief clock and the model clock of agent $i$ are the Fisher lengths of $t \mapsto q_i^b(t)$ and $t \mapsto q_i^m(t)$, combined by the channel-weighted speed (7.18). Along a history of $\mathcal{Q}_V$ with block-diagonal metric $\bigoplus_i \mathsf{G}_i$, the clock $\mathfrak{t}_i$ of agent $i$ is the $\mathsf{G}_i$-length of $t \mapsto Q_i(t)$. The origin is part of the clock, and the parameter $t$ does not enter the length.

**Theorem 7.36 (Reparameterization invariance and unit-speed form).** Let $\widetilde{Q} = Q \circ \phi$ with $\phi : \widetilde{\mathcal{T}} \to \mathcal{T}$ an orientation-preserving $C^1$ diffeomorphism. Corresponding subarcs have the same length, and

$$\mathfrak{t}_{\widetilde{Q},\widetilde{t}_0}(\widetilde{t}) = \mathfrak{t}_{Q,t_0}(\phi(\widetilde{t})), \qquad t_0 = \phi(\widetilde{t}_0). \tag{7.49}$$

If $\dot{\mathfrak{t}} > 0$, the clock is strictly increasing and gives the unit-speed representative of the oriented history; on a subinterval where $\dot{\mathfrak{t}}$ vanishes, as on a critical or null segment, the clock is constant and is not a coordinate. If $Q$ satisfies (7.41), then on its noncritical part

$$\frac{d\mathfrak{t}}{dt} = \mathsf{a} \left\| \operatorname{grad}^{\mathsf{G}} \mathcal{F} \right\|_{\mathsf{G}}, \tag{7.50}$$

$$\frac{dQ}{d\mathfrak{t}} = -\frac{\operatorname{grad}^{\mathsf{G}} \mathcal{F}}{\left\| \operatorname{grad}^{\mathsf{G}} \mathcal{F} \right\|_{\mathsf{G}}}, \tag{7.51}$$

$$\frac{d\mathcal{F}}{d\mathfrak{t}} = -\left\| \operatorname{grad}^{\mathsf{G}} \mathcal{F} \right\|_{\mathsf{G}}. \tag{7.52}$$

If the clock is read with a second metric $\mathsf{G}'$, for instance the unweighted Fisher metric along a flow whose blocks carry rate weights, then along any $C^1$ history with velocity $Z \neq 0$

$$\left| \frac{d\mathcal{F}}{d\mathfrak{t}'} \right| = \frac{|d\mathcal{F}(Z)|}{\|Z\|_{\mathsf{G}'}} \leq \left\| \operatorname{grad}^{\mathsf{G}'} \mathcal{F} \right\|_{\mathsf{G}'}, \tag{7.53}$$

with equality exactly when $d\mathcal{F} = 0$ or $Z$ is parallel to $\operatorname{grad}^{\mathsf{G}'} \mathcal{F}$. Along a history of $\mathcal{Q}_V$ with block-diagonal metric the squared speeds add, $\dot{\mathfrak{t}}^2 = \sum_i \dot{\mathfrak{t}}_i^2$, and over any common parameter interval $\max_i \mathfrak{t}_i \leq (\sum_i \mathfrak{t}_i^2)^{1/2} \leq \mathfrak{t} \leq \sum_i \mathfrak{t}_i$. [ESTABLISHED]

The proof is in appendix B.2.

**Remark 7.37 (Endpoints do not determine an update duration).** For an update $Q$ from $Q(t_0) = Q'$ to $Q(t_1) = Q''$ in one path component of $\mathcal{Q}$, the distance $\operatorname{dist}_{\mathsf{G}}(Q', Q'')$, the infimum of $\operatorname{len}_{\mathsf{G}}(\Gamma; [t_0, t_1])$ over curves $\Gamma$ with these endpoints, is at most $\operatorname{len}_{\mathsf{G}}(Q; [t_0, t_1])$, because the realized update is one such curve, and the inequality is strict for a curved path in the Euclidean Fisher geometry of the two-parameter normal location family. Duration is also distinct from the amount of curve parameter: in the unit-variance normal location family the path $\mu(t) = e^{-t}$ has length $\int_0^\infty |\dot{\mu}(t)| dt = \int_0^\infty e^{-t} dt = 1$ over an unbounded parameter interval. [ESTABLISHED]

**Proposition 7.38 (Finite duration for analytic objectives).** Let $\mathcal{Q}$ be finite-dimensional with a continuous positive-definite metric $\mathsf{G}$, and let $\gamma : [0, \infty) \to \mathcal{Q}$ satisfy $\dot{\gamma} = -\operatorname{grad}^{\mathsf{G}} \mathcal{F}$. If $\gamma(t)$ has an accumulation point $Q_*$ as $t \to \infty$ and $\mathcal{F}$ is real analytic in a chart around $Q_*$, then $\gamma(t) \to Q_*$, $Q_*$ is critical, and $\operatorname{len}_{\mathsf{G}}(\gamma; [0, \infty))$ is finite. This holds for every forward history with relatively compact image when $\mathcal{Q}$ has a real-analytic atlas in whose charts $\mathcal{F}$ is real analytic. Analyticity cannot be weakened to smoothness: for the $C^\infty$ function of Absil et al. [1, Section 2, Eq. (8)] on the Euclidean plane, the gradient history from the initial data given there is bounded, has the unit circle as its $\omega$-limit set, and has infinite duration. [ESTABLISHED]

**Proposition 7.39 (Invariant free-energy histories are horizontal).** Let a Lie group $\mathbb{G}$ act smoothly on a finite-dimensional configuration manifold $\mathcal{Q}$ with a positive-definite metric $\mathsf{G}$, write $\xi_{\mathcal{Q}}(Q) = \frac{d}{dt}\big|_{t=0} \exp(t\xi)\cdot Q$ and $V_Q = \{\xi_{\mathcal{Q}}(Q) : \xi \in \operatorname{Lie}\mathbb{G}\}$, and let $\mathcal{F}$ be $C^1$ and $\mathbb{G}$-invariant. Then $\operatorname{grad}^{\mathsf{G}}\mathcal{F}(Q)$ is $\mathsf{G}$-orthogonal to $V_Q$ at every $Q$, every natural-gradient history has velocity orthogonal to the orbits, and its gauge-quotient speed $\min_\xi \|\dot{Q} - \xi_{\mathcal{Q}}(Q)\|_{\mathsf{G}}$ equals its Fisher speed. If $\mathsf{G}$ is $C^\infty$ and the action is free, proper, and isometric, the quotient map $\pi : \mathcal{Q} \to \mathcal{Q}/\mathbb{G}$ is a Riemannian submersion for a unique metric $\bar{\mathsf{G}}$ on the quotient, $\mathcal{F} = \bar{\mathcal{F}} \circ \pi$, and $T\pi \operatorname{grad}^{\mathsf{G}}\mathcal{F} = \operatorname{grad}^{\bar{\mathsf{G}}}\bar{\mathcal{F}} \circ \pi$; each natural-gradient history descends to a natural-gradient history of $\bar{\mathcal{F}}$ with the same Fisher duration. [ESTABLISHED]

The proofs of propositions 7.38 and 7.39 are in appendix B.2.

## 7.8 Orthogonal clocks

**Theorem 7.40 (Exact criterion for an orthogonal clock).** Let $\mathcal{U} \subseteq \mathcal{Q}$ be a connected open set without critical points, with $\mathcal{F}$ of class $C^2$ and $\mathsf{G}$ of class $C^1$. Define $\mathsf{n}$, $\mathbf{e}_{\mathcal{F}}$, and $\boldsymbol{\alpha}$ by the first line below, and call a $C^1$ function $\mathrm{t}^\perp : \mathcal{U} \to \mathbb{R}$ that satisfies the second an orthogonal clock:

$$\mathsf{n} = \left\|\operatorname{grad}^{\mathsf{G}}\mathcal{F}\right\|_{\mathsf{G}}, \qquad \mathbf{e}_{\mathcal{F}} = -\frac{\operatorname{grad}^{\mathsf{G}}\mathcal{F}}{\mathsf{n}}, \qquad \boldsymbol{\alpha} = \mathbf{e}_{\mathcal{F}}^{\flat} = -\frac{d\mathcal{F}}{\mathsf{n}}, \tag{7.54}$$

$$d\mathrm{t}^\perp = \boldsymbol{\alpha}. \tag{7.55}$$

An orthogonal clock exists if and only if

$$d\boldsymbol{\alpha} = 0, \qquad \oint_\gamma \boldsymbol{\alpha} = 0 \quad \text{for every piecewise } C^1 \text{ closed curve } \gamma \subset \mathcal{U}, \tag{7.56}$$

and on a simply connected $\mathcal{U}$ closedness suffices. The local obstruction is

$$d\boldsymbol{\alpha} = \frac{1}{\mathsf{n}^2} d\mathsf{n} \wedge d\mathcal{F}, \tag{7.57}$$

so $\boldsymbol{\alpha}$ is closed exactly when $\mathsf{n}$ is locally a function of $\mathcal{F}$, and an orthogonal clock is then constant on every connected piece of every level set of $\mathcal{F}$ and locally a strictly decreasing function of $\mathcal{F}$. Whenever $\mathrm{t}^\perp$ exists, every unit-speed natural-gradient history in $\mathcal{U}$ obeys $d\mathrm{t}^\perp/d\mathrm{t} = 1$, and the level sets of $\mathrm{t}^\perp$ are $\mathsf{G}$-orthogonal to $\mathbf{e}_{\mathcal{F}}$. For $\mathcal{F} = z_1 z_2$ on the open first quadrant with the Euclidean metric,

$$d\boldsymbol{\alpha} = \frac{z_1^2 - z_2^2}{(z_1^2 + z_2^2)^{3/2}} dz_1 \wedge dz_2, \tag{7.58}$$

which vanishes identically on no open set, so no orthogonal clock exists on any open subset of the quadrant. [ESTABLISHED]

**Proposition 7.41 (Unit-speed clocks from a global cross-section).** Let $\mathcal{U}$ be as in theorem 7.40, and let $\mathcal{S} \subset \mathcal{U}$ be a $C^1$ hypersurface that every maximal integral curve of $\mathbf{e}_{\mathcal{F}}$ in $\mathcal{U}$ meets exactly once and transversally. Then the signed Fisher arc length along each orbit, measured from its crossing

with $\mathcal{S}$, is a $C^1$ function $\mathfrak{t}^{\mathcal{S}} : \mathcal{U} \to \mathbb{R}$ with $d\mathfrak{t}^{\mathcal{S}}(\mathbf{e}_{\mathcal{F}}) = 1$, so every unit-speed natural-gradient history obeys $d\mathfrak{t}^{\mathcal{S}}/d\mathfrak{t} = 1$. A level set of $\mathcal{F}$ that every orbit meets is such a cross-section, and the level sets of $\mathfrak{t}^{\mathcal{S}}$ are G-orthogonal to $\mathbf{e}_{\mathcal{F}}$ only if $\boldsymbol{\alpha}$ is closed. For $\mathcal{F} = z_1 z_2$ on the open first quadrant with the Euclidean metric, the hyperbola $\{z_1 z_2 = 1\}$ is a global cross-section, so a unit-speed clock exists on the whole quadrant although no orthogonal one exists on any open subset. [ESTABLISHED]

The proofs of theorem 7.40 and proposition 7.41 are in appendix B.2.

## 7.9 Record clocks

**Corollary 7.42 (Same-path record and meta-agent clock contraction).** Let $t \mapsto P_t$ be a $C^1$ history in a family of laws on Y that is differentiable in quadratic mean along it, with directional score $\mathbf{s}_t$, and let $K : \mathrm{Y} \rightsquigarrow \bar{\mathrm{Y}}$ be a Markov kernel that does not depend on $t$ and satisfies $K(y, \bar{\mathrm{Y}}) = 1$ for every $y$. Then the record history $\bar{P}_t = P_t K$, the image of the same fine path, satisfies

$$\bar{\mathbf{s}}_t(\bar{Y}) = \mathbb{E}_t \left[ \mathbf{s}_t(Y) \mid \bar{Y} \right], \tag{7.59}$$

$$\dot{\mathfrak{t}}_P^2 - \dot{\mathfrak{t}}_{\bar{P}}^2 = \mathbb{E}_t \operatorname{Var}_t \left( \mathbf{s}_t(Y) \mid \bar{Y} \right) \geq 0, \tag{7.60}$$

and on every common parameter interval

$$\operatorname{len}(\bar{P}) \leq \operatorname{len}(P), \qquad \mathfrak{t}_{\bar{P}}(t_1) - \mathfrak{t}_{\bar{P}}(t_0) \leq \mathfrak{t}_P(t_1) - \mathfrak{t}_P(t_0), \tag{7.61}$$

with equality, when the fine length on the interval is finite, exactly when the fine score is $\bar{Y}$-measurable for almost every $t$ in the interval. Equality of two infinite lengths gives no such conclusion. The inequality applies to a meta-agent exactly when its law is the image of the same fine history under one such kernel. Parameter independence of the kernel cannot be dropped: for a dominated channel $K_t$ with density $\mathsf{k}_t(\bar{y} \mid y)$, assume the joint family $P_t(dy)K_t(y, d\bar{y})$ is differentiable in quadratic mean with square-integrable score $\mathbf{s}_t(y) + \partial_t \log \mathsf{k}_t(\bar{y} \mid y)$. Its record score is

$$\bar{\mathbf{s}}_t(\bar{Y}) = \mathbb{E}_t \left[ \mathbf{s}_t(Y) + \partial_t \log \mathsf{k}_t(\bar{Y} \mid Y) \,\middle|\, \bar{Y} \right], \tag{7.62}$$

and a constant input followed by a channel that emits a Bernoulli bit of success probability $(1+e^{-t})^{-1}$ has zero fine speed and record speed $(2\cosh(t/2))^{-1} > 0$. [ESTABLISHED]

The proof is in appendix B.2.

**Open Problem 7.43 (Quantitative meta-agent clock ratio).** Corollary 7.42 shows that a meta-agent whose law is a Markov image of its constituents' joint history accumulates no more Fisher duration than that history, with the exact defect (7.60). Open is the ratio between the meta-agent's duration and the constituents' collective duration for an $N$-constituent meta-agent of chapter 11, as a function of $N$ and of the correlation between the constituents' velocities, for coherent histories (aligned velocities) and incoherent ones (independent random velocities). The expected ratio of arc lengths is not the square root of the expected ratio of relative entropies, so a divergence

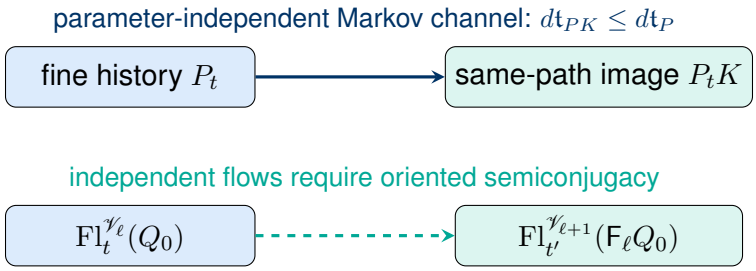



**Figure 7.1:** The upper arrow is the same-path comparison of corollary 7.42, with its conditional-variance defect. The lower arrow compares independently generated flows and exists only under the semiconjugacy condition of definition 7.45.

computation does not settle it; a theorem that computes or bounds the ratio for a stated coarse map and a stated law of constituent velocities, including its dependence on $N$ in both regimes, does.

## 7.10 The coarse configuration map and oriented semiconjugacy

**Definition 7.44 (Coarse configuration map).** A coarse configuration map at depth $\ell$ is a smooth map

$$\mathsf{F}_\ell : \mathcal{Q}_\ell \longrightarrow \mathcal{Q}_{\ell+1} \tag{7.63}$$

from the configuration manifold of the fine agents at $c_*$ to that of their meta-agents, specified separately from the sample kernels that may realize it. The fine and coarse natural-gradient fields are $\mathscr{V}_\ell = -\operatorname{grad}^{\mathsf{G}_\ell} \mathcal{F}_\ell$ and $\mathscr{V}_{\ell+1} = -\operatorname{grad}^{\mathsf{G}_{\ell+1}} \mathcal{F}_{\ell+1}$, with local flows $\mathrm{Fl}_t^{\mathscr{V}_\ell}$ and $\mathrm{Fl}_{t'}^{\mathscr{V}_{\ell+1}}$.

**Definition 7.45 (Oriented semiconjugacy).** For an open $\mathcal{U} \subseteq \mathcal{Q}_\ell$, the pair $(\mathsf{F}_\ell, \mathsf{a}_\ell)$ with $\mathsf{a}_\ell : \mathcal{U} \to (0, \infty)$ continuous is an oriented semiconjugacy of $\mathscr{V}_\ell$ onto $\mathscr{V}_{\ell+1}$ over $\mathcal{U}$ when

$$T_Q \mathsf{F}_\ell \mathscr{V}_\ell(Q) = \mathsf{a}_\ell(Q) \mathscr{V}_{\ell+1}(\mathsf{F}_\ell Q) \quad \text{for every } Q \in \mathcal{U}. \tag{7.64}$$

**Lemma 7.46 (The factor is determined and regular off the coarse critical set).** If $\mathscr{V}_{\ell+1}(\mathsf{F}_\ell Q) \neq 0$, then $\mathsf{a}_\ell(Q)$ is unique and

$$\mathsf{a}_\ell(Q) = \frac{\mathsf{G}_{\ell+1}\left(T_Q \mathsf{F}_\ell \mathscr{V}_\ell(Q), \mathscr{V}_{\ell+1}(\mathsf{F}_\ell Q)\right)}{\|\mathscr{V}_{\ell+1}(\mathsf{F}_\ell Q)\|^2_{\mathsf{G}_{\ell+1}}}, \tag{7.65}$$

so $\mathsf{a}_\ell$ is $C^k$ wherever $Q \mapsto T_Q \mathsf{F}_\ell \mathscr{V}_\ell(Q)$, $Q \mapsto \mathscr{V}_{\ell+1}(\mathsf{F}_\ell Q)$, and the pulled-back coefficients of $\mathsf{G}_{\ell+1}$ are $C^k$, for instance when $\mathsf{F}_\ell \in C^{k+1}$ and $\mathscr{V}_\ell, \mathscr{V}_{\ell+1}, \mathsf{G}_{\ell+1}$ are $C^k$; continuity of $\mathsf{a}_\ell$ is therefore automatic where $\mathscr{V}_{\ell+1} \neq 0$ on $\mathsf{F}_\ell(\mathcal{U})$. If $\mathscr{V}_{\ell+1}(\mathsf{F}_\ell Q) = 0$, then (7.64) forces $T_Q \mathsf{F}_\ell \mathscr{V}_\ell(Q) = 0$ and fixes no value of $\mathsf{a}_\ell(Q)$; limits from the noncritical set may determine it or may obstruct any continuous positive extension. [ESTABLISHED]

*Proof.* Pair (7.64) with $\mathscr{V}_{\ell+1}(\mathsf{F}_\ell Q)$ and divide by its positive squared norm. □

**Theorem 7.47 (Oriented semiconjugacy and maximal intervals).** Assume (7.64) on an open $\mathcal{U}$ with $\mathscr{V}_{\ell+1}$ locally Lipschitz. Fix $Q \in \mathcal{U}$, let $\mathcal{T}_Q$ be the connected component of $0$ in the set of $t$ for which $\mathrm{Fl}^{\mathscr{V}_\ell}_{t''}(Q) \in \mathcal{U}$ for every $t''$ between $0$ and $t$, and set $\psi_Q(t) = \int_0^t \mathsf{a}_\ell(\mathrm{Fl}^{\mathscr{V}_\ell}_{t''} Q)dt''$. Then $\psi_Q$ is a strictly increasing $C^1$ diffeomorphism of $\mathcal{T}_Q$ onto an open interval $\mathcal{T}'_Q \ni 0$, and

$$\mathcal{T}'_Q \subseteq \bar{\mathcal{T}}^{\max}_{\mathsf{F}_\ell Q}, \qquad \mathsf{F}_\ell\left(\mathrm{Fl}^{\mathscr{V}_\ell}_t(Q)\right) = \mathrm{Fl}^{\mathscr{V}_{\ell+1}}_{\psi_Q(t)}(\mathsf{F}_\ell Q) \quad \text{for every } t \in \mathcal{T}_Q, \tag{7.66}$$

where $\bar{\mathcal{T}}^{\max}_{\mathsf{F}_\ell Q} = (\bar{t}^-, \bar{t}^+)$ is the maximal interval of the coarse flow through $\mathsf{F}_\ell Q$. With $\mathcal{T}_Q = (t^-, t^+)$ and endpoints in the extended real line,

$$\mathcal{T}'_Q = \left(\lim_{t \downarrow t^-} \psi_Q(t), \lim_{t \uparrow t^+} \psi_Q(t)\right), \tag{7.67}$$

and the whole maximal coarse orbit is traversed if and only if the two limits equal $\bar{t}^-$ and $\bar{t}^+$; when the fine orbit is complete in $\mathcal{U}$ and the coarse maximal interval is $\mathbb{R}$, this means $\int_0^{+\infty} \mathsf{a}_\ell(\mathrm{Fl}^{\mathscr{V}_\ell}_t Q)dt = \int_{-\infty}^0 \mathsf{a}_\ell(\mathrm{Fl}^{\mathscr{V}_\ell}_t Q)dt = +\infty$. A positive lower bound on $\mathsf{a}_\ell$ suffices and is not necessary: on $\mathbb{R}$ with $\mathscr{V}_\ell = \mathscr{V}_{\ell+1} = \partial_z$ and $\mathsf{F}_\ell = \operatorname{arsinh}$, $\mathsf{a}_\ell = (1+z^2)^{-1/2}$ has infimum zero while $\psi_0 = \operatorname{arsinh}$ maps $\mathbb{R}$ onto $\mathbb{R}$; with $\mathsf{F}_\ell = \arctan$, $\mathsf{a}_\ell = (1+z^2)^{-1}$ is positive while $\mathcal{T}'_0 = (-\pi/2, \pi/2)$ although both flows are complete. Conversely, differentiating a flow relation of the form (7.66) at $t = 0$ gives (7.64). [ESTABLISHED]

The proof is in appendix B.2.

**Proposition 7.48 (Noncollapse is a separate hypothesis).** Condition (7.64) alone permits total collapse: with $\mathsf{F}_\ell$ constant and $\mathscr{V}_{\ell+1} \equiv 0$ it holds for every positive $\mathsf{a}_\ell$, while every fine orbit maps to a point and every coarse duration is zero. A nonconstant shared oriented history additionally requires $T_Q \mathsf{F}_\ell \mathscr{V}_\ell(Q) \neq 0$ on every nontrivial subarc, equivalently $\mathscr{V}_{\ell+1}(\mathsf{F}_\ell Q) \neq 0$ there, and full traversal requires the endpoint equalities of (7.67); neither follows from (7.64). [ESTABLISHED]

*Proof.* Both sides of (7.64) vanish in the example, and the arctan example of theorem 7.47 satisfies it without full traversal. □

**Theorem 7.49 (The semiconjugacy rate cancels from geometric length).** Under (7.64), let $Q(t)$ be a fine orbit. The image speed is

$$\left\| \tfrac{d}{dt} \mathsf{F}_\ell Q(t) \right\|_{\mathsf{G}_{\ell+1}} = \mathsf{a}_\ell\left(Q(t)\right) \left\| \mathscr{V}_{\ell+1}\left(\mathsf{F}_\ell Q(t)\right) \right\|_{\mathsf{G}_{\ell+1}}, \tag{7.68}$$

and for $t_0 < t_1$ in $\mathcal{T}_Q$ the accumulated coarse duration is

$$\mathfrak{t}^{(\ell+1)}(t_1) - \mathfrak{t}^{(\ell+1)}(t_0) = \int_{\psi_Q(t_0)}^{\psi_Q(t_1)} \left\| \mathscr{V}_{\ell+1}\left(\mathrm{Fl}^{\mathscr{V}_{\ell+1}}_{t'}(\mathsf{F}_\ell Q)\right) \right\|_{\mathsf{G}_{\ell+1}} dt', \tag{7.69}$$

the intrinsic $\mathsf{G}_{\ell+1}$-length of the coarse orbit arc, which does not depend on $\mathsf{a}_\ell$. [ESTABLISHED]

*Proof.* Equation (7.68) is (7.64) with positive homogeneity of the norm and $\mathsf{a}_\ell > 0$. Substituting it and changing variables $t' = \psi_Q(t)$, whose Jacobian is $\mathsf{a}_\ell(\mathrm{Fl}_t^{\mathscr{V}_\ell} Q)$, cancels the factor. □

**Remark 7.50 (Exact duration comparison).** Under (7.64), along one orbit and from a common origin, the coarse and fine durations agree identically if and only if $\|T\mathsf{F}_\ell \mathscr{V}_\ell\|_{\mathsf{G}_{\ell+1}} = \|\mathscr{V}_\ell\|_{\mathsf{G}_\ell}$ almost everywhere along the orbit, they agree up to a factor $\mathrm{c} > 0$ if and only if the same holds with that factor, and the coarse duration has no larger increments on every subinterval if and only if

$$\|T\mathsf{F}_\ell \mathscr{V}_\ell\|_{\mathsf{G}_{\ell+1}} \le \|\mathscr{V}_\ell\|_{\mathsf{G}_\ell} \quad \text{along the orbit,} \tag{7.70}$$

because both durations are integrals of continuous nonnegative densities from a common origin, and two such integrals agree for every upper limit exactly when the densities agree almost everywhere. The condition is scalar and concerns the single direction $\mathscr{V}_\ell$; the Loewner condition $\mathsf{F}_\ell^* \mathsf{G}_{\ell+1} \preceq \mathsf{G}_\ell$ is sufficient and stronger, and corollary 7.51 supplies it for channelwise Markov maps. Fiberwise Fisher contraction alone gives neither, for two independent reasons. The configuration metrics are separate data: on $\mathbb{R}$ with $\mathscr{V}_\ell = \mathscr{V}_{\ell+1} = \partial_z$ and $\mathsf{F}_\ell$ the identity, the semiconjugacy holds with $\mathsf{a}_\ell \equiv 1$, and the metrics $dz^2$ and $4dz^2$ double the duration, whether realized with the identity fiber map and unequal channel weights $1$ and $4$ (zero Fisher defect) or with fibers $\mathcal{N}(\mu, 1)$ and $\mathcal{N}(\mu, \frac{1}{4})$ and the identity parameter map (Fisher defect $1 - 4 = -3$, so no Markov morphism). And contraction compares the image of the fine orbit, not a recomputed coarse orbit: on $\mathbb{R}^2$ with $\mathsf{F}_\ell$ the identity and the Euclidean metric at both depths, $\mathcal{F}_\ell = \frac{1}{2}(z_1^2 + z_2^2)$ and $\mathcal{F}_{\ell+1} = \frac{1}{2}(z_1^2 + 4z_2^2)$ give the fine history $(e^{-t}, e^{-t})$ of duration $\sqrt{2}$ and the recomputed coarse history $(e^{-t}, e^{-4t})$ of duration $\int_0^1 \sqrt{1 + 16z^6} dz > \sqrt{2}$, and (7.64) fails because $(z_1, z_2) = \mathsf{a}(z_1, 4z_2)$ forces $\mathsf{a} = 1$ and $z_2 = 0$. [ESTABLISHED]

**Corollary 7.51 (When fiberwise contraction does lift).** Let $\mathcal{Q}_\ell$ and $\mathcal{Q}_{\ell+1}$ carry the channel-weighted metrics $\mathsf{G}_\ell = \sum_x w_x g_x^F$ and $\mathsf{G}_{\ell+1} = \sum_x w_x \bar{g}_x^F$ with the same positive weights, and let $\mathsf{F}_\ell$ act channelwise by Markov coarse morphisms, $(\mathsf{F}_\ell Q)_x = \mathscr{K}_x(Q_x)$. Then for every $Z \in T_Q \mathcal{Q}_\ell$,

$$\|Z\|_{\mathsf{G}_\ell}^2 - \|T\mathsf{F}_\ell Z\|_{\mathsf{G}_{\ell+1}}^2 = \sum_x w_x \Delta_F^{\mathscr{K}_x}(Z_x, Z_x) \ge 0, \tag{7.71}$$

so the image of every $C^1$ fine curve has no greater speed and no greater duration on every subinterval, and (7.70) holds along every orbit of an oriented semiconjugacy. On a subinterval of finite fine duration, equality holds exactly when, for every channel and almost every parameter value, the fine score in the direction $Z_x$ is measurable with respect to the coarse output. [ESTABLISHED]

*Proof.* Differentiate the channelwise action, insert it into $\mathsf{G}_{\ell+1}$, and subtract from $\mathsf{G}_\ell$; the difference is the weighted sum of the vertical defects, which are positive semidefinite with the stated equality case by corollary 7.11. Monotonicity of the square root and of the integral gives the speed and duration statements. □

## 7.11 Semiconjugacy on symmetry quotients

Throughout this section $\mathscr{V}_\ell$ and $\mathscr{V}_{\ell+1}$ are natural-gradient fields of $C^2$ objectives for positive-definite metrics.

**Definition 7.52 (Functional compatibility).** A coarse configuration map $\mathsf{F}_\ell$ is functionally compatible on an open set $\mathcal{U} \subseteq \mathcal{Q}_\ell$ when there is a $C^1$ function $\chi_\ell$ with $\chi_\ell' > 0$ such that $\mathcal{F}_\ell = \chi_\ell \circ \mathcal{F}_{\ell+1} \circ \mathsf{F}_\ell$ on $\mathcal{U}$.

**Proposition 7.53 (Semiconjugacy on symmetry quotients).** Equal objectives do not intertwine natural gradients: with $\mathsf{F}_\ell$ the identity of $\mathbb{R}^2$, $\mathcal{F}_\ell = \mathcal{F}_{\ell+1} = \frac{1}{2}(z_1^2 + 2z_2^2)$, and metrics $\mathrm{diag}(1,1)$ and $\mathrm{diag}(1,\mathrm{c})$ with $\mathrm{c} \neq 1$, (7.64) fails wherever $z_1 z_2 \neq 0$. Functional compatibility on $\mathcal{U}$ forces $d\mathcal{F}_\ell$ to vanish on $\ker T\mathsf{F}_\ell$ at every point of $\mathcal{U}$, so when $\mathsf{F}_\ell$ is a submersion, $\mathcal{F}_\ell$ is constant on each connected piece of each fiber of $\mathsf{F}_\ell$ in $\mathcal{U}$. If $\mathsf{F}_\ell$ is a surjective submersion, functionally compatible on $\mathcal{U}$, and horizontally conformal with dilation $\mathsf{d}_\ell > 0$, meaning that $\mathsf{G}_{\ell+1}(T\mathsf{F}_\ell Z_1, T\mathsf{F}_\ell Z_2) = \mathsf{d}_\ell^2 \mathsf{G}_\ell(Z_1, Z_2)$ for $Z_1, Z_2 \in \mathrm{Hor}_Q = (\ker T_Q \mathsf{F}_\ell)^\perp$, then (7.64) holds on $\mathcal{U}$ with

$$\mathsf{a}_\ell(Q) = \chi_\ell'\left(\mathcal{F}_{\ell+1}(\mathsf{F}_\ell Q)\right) \mathsf{d}_\ell(Q)^2 > 0. \tag{7.72}$$

In particular the quotient map of a free, proper, isometric action that preserves $\mathcal{F}_\ell$, with the quotient metric and descended objective of proposition 7.39, is an oriented semiconjugacy with $\mathsf{a}_\ell \equiv 1$; and at $c_*$ the identity map from the mean coordinate of the fiber $\{\mathcal{N}(\mu,1)\}$ to that of $\{\mathcal{N}(\mu,2)\}$, realized by the kernel $y \mapsto \mathcal{N}(y,1)$, is horizontally conformal with $\mathsf{d}_\ell^2 = \frac{1}{2}$ and gives $\mathsf{a}_\ell = \frac{1}{2}$ for $\mathcal{F}_\ell = \mathcal{F}_{\ell+1} = \mu^2/2$, an injective case that does not exercise the fiber condition. The route does not reach maps that discard directions on which the fine objective depends: if $\mathcal{F}_{\ell+1}(Q') = \inf\{\mathcal{F}_\ell(Q) : \mathsf{F}_\ell Q = Q'\}$, $\mathsf{F}_\ell$ is a submersion with positive-dimensional fibers, and each fiber contains at most one minimizer of $\mathcal{F}_\ell$, then compatibility with $\chi_\ell = \mathrm{id}$ holds on no open set; for $\mathsf{F}_\ell(f,s) = s$ and $\mathcal{F}_\ell(f,s) = \frac{1}{2}(f-s)^2 + \frac{1}{2}s^2$ it holds for no $\chi_\ell$ on any open set. [ESTABLISHED]

The proof is in appendix B.2.

For maps that discard directions on which the fine objective depends, the exact route runs through the reduced slow flow, which on the set of conditional minimizers is conjugate to the coarse natural-gradient flow (proposition 8.4, (8.8)); whether the fine or rate-separated flow follows it is open problem 8.7.

## 7.12 Coarse-graining depth, orbit parameter, and duration

A tower of coarse configurations $Q^{(\ell+1)} = \mathsf{F}_\ell(Q^{(\ell)})$, $\ell = 0, 1, 2, \ldots$, is indexed by the coarse-graining depth $\ell$, which is not a parameter on a history. A history can be present at every depth, as $Q^{(\ell)}(t)$ with accumulated duration $\mathfrak{t}^{(\ell)}(t)$; when $Q^{(\ell+1)}(t) = \mathsf{F}_\ell(Q^{(\ell)}(t))$ and $\mathsf{F}_\ell$ acts by Markov coarse morphisms as in corollary 7.51, the record-clock comparison applies between adjacent depths, while for a general oriented semiconjugacy the exact comparison is remark 7.50, and independently

recomputed histories also need the noncollapse conditions of proposition 7.48. The Markov property of a configuration arrow can fail even for natural meta-agent maps: no parameter-independent Markov kernel sends every product of two normal laws to their joint mean-and-variance barycenter, as Supplement G proves.

**Proposition 7.54 (Typed non-identification of depth, orbit parameter, and duration).** Coarse-graining depth $\ell$, an orbit parameter $t$, and Fisher duration $\mathfrak{t}^{(\ell)}$ are distinct: $\ell$ is a discrete index, $t$ a chosen oriented parameter on a selected orbit, and $\mathfrak{t}^{(\ell)}$ the arc length for a chosen metric from a chosen origin. Once the metric, origin, orientation, and parameterized path are chosen, $\mathfrak{t}^{(\ell)}$ is a function of $t$, but it is not determined by the unparameterized orbit or by $\ell$: changing the metric rescales it and changing the origin shifts it. Strict monotonicity of $\mathfrak{t}^{(\ell)}$ does not make it a regular coordinate: for $\mathsf{G} = dz^2$ on $\mathbb{R}$ and the path $z(t) = t^2$ on $[0, 1]$, $\mathfrak{t}(t) = t^2$, whose inverse is not differentiable at the origin, and a zero-speed subinterval destroys strict monotonicity. [ESTABLISHED]

*Proof.* Only the metric, origin, orientation, and parameterized path produce a duration function. Metric dependence is homogeneity of the norm in (7.45), and the example is $\int_0^t 2t''dt'' = t^2$ with the square root not differentiable at zero. □

## Physicist's summary

The Fisher metric on an agent's belief and model law fibers is a gauge-invariant ruler and the second-order term of the relative entropy, so a loop holonomy near the identity appears as a quadratic form on the Lie algebra that vanishes on the isotropy algebra of the law. A Markov channel loses exactly the conditional variance of the score, so a meta-agent that is a channel image of its constituents' history has an information clock no faster than theirs, by a ratio that is open. A chosen metric turns the free energy into an orbit, Fisher arc length removes its parameterization, and a belief-model rate ratio, unlike a scalar mobility, can bend the orbit; an orthogonal clock needs the gradient norm to be locally a function of the free energy, a unit-speed clock only a global cross-section. By the parity witness, joint laws, Fisher metrics, and free energies are not functions of marginals. A recomputed coarse flow is the image of the fine flow only under an oriented semiconjugacy, which holds with rate one on symmetry quotients; information-discarding maps are compared through the reduced slow flow of the next chapter.

# Chapter 8

# Two-Timescale Dynamics

This chapter asks what the collective free energy of chapter 5 supplies for separating fast inference of beliefs from slower evolution of models, priors, and parents: descent for every positive-definite block mobility, local in a product family (proposition 8.1); a reduced slow flow conjugate to its tangent lift on the graph of the conditional optimum (lemma 8.3 and proposition 8.4), the exact route of chapter 7 for information-discarding maps; and a Markov parent history under instantaneous within-parent refresh (theorem 8.11), whose failure under finite refresh (remark 8.12) is measured in chapter 11. The slow block (definition 8.2) moves by accepted M-coordinates (definition 5.28) or by the rate-separated flow; the formation family of chapter 13 realizes one scalar slow parameter (13.25), not slow per-agent model laws, while on the four-cycle of that chapter, whose link translations this chapter writes $\tilde{\omega}$, slow model means are a supplied rate, and the model channel lags the belief channel exactly when the links are slow as well (proposition 8.6, the held-sheet instance of open problem 8.7). The finite model of the last section uses fine states $\xi \in \Xi$, parent values $z \in \mathsf{Z}$, a parent map $\mathrm{par}$, a prior $p^0$, record laws $\psi_\xi$, a refresh fraction $\omega$, and sans-serif kernels acting on row vectors from the right.

**Scope.** The collective free energy does not choose the fast–slow split, the rate ratio, or the block mobilities, and a rate ratio is an anisotropic mobility that can change the selected orbit (proposition 7.31). The chapter does not claim that the rate-separated flow tracks the reduced flow for a recognition family of the program beyond the held matching sheets of proposition 8.6, that the sampled dynamics with source jumps track it, that a sampled episode with source jumps at the temperature of definition 13.21 keeps a matched pair belief-coherent and model-incoherent on an interval of order $1/\varepsilon_{\mathrm{m}}$ with probability bounded below, that the discrete alternating scheme converges, that the parent sweep converges at a rate uniform in the number of fine states or in the smallest posterior weight, or that a static parent datum has closed dynamics (open problems 5.32, 8.7 and 11.38). [NOT-CLAIMED]

## 8.1 Block descent at fixed parameter

**Proposition 8.1 (Block descent of the collective free energy).** Let $\lambda \mapsto Q_\lambda$ be a family of recognition laws on an open domain $\Lambda = \prod_{i \in V} \Lambda_i$ with block coordinates $\lambda = (\lambda_i)_i$, such that $F(\lambda) := \mathcal{F}_o^{\mathrm{ext}}(Q_\lambda)$ is real and $C^1$ on $\Lambda$, and let $\mathsf{G}_i(\lambda)$ be continuous fields of symmetric positive-definite block metrics, with positive block rates absorbed into the metrics and blocks of zero rate

held fixed. Let $\lambda(t)$ solve

$$\dot{\lambda}_i = -\mathsf{G}_i^{-1}\nabla_{\lambda_i}F \tag{8.1}$$

for the moving blocks $i \in V' \subseteq V$ and remain in $\Lambda$.

(i) Along the solution

$$\frac{d}{dt}F = -\sum_{i\in V'}(\nabla_{\lambda_i}F)^{\top}\mathsf{G}_i^{-1}(\nabla_{\lambda_i}F) \le 0. \tag{8.2}$$

(ii) If the family is a regular product family $Q_\lambda = \bigotimes_i Q_{i,\lambda_i}$ with $\mathbb{E}_{Q_{-i}}|\log Z_{\{i\}}(Y_{-i})| < \infty$ for every $i$ and $\lambda$, then $\nabla_{\lambda_i}F = \nabla_{\lambda_i}\overline{\mathcal{F}}_i$ for the outside-averaged local free energy (5.53), so an agent that descends its own local free energy moves along a block of a descent flow of the collective one. The Fisher metric of such a family is block diagonal, so with $\mathsf{G}_i$ equal to the block Fisher metrics the flow is the Fisher natural gradient of $F$.

(iii) Outside product families moving $\lambda_i$ can change $Q_{-i}$, and the identification in (ii) needs the conditional lift of definition 7.17, whose outside-averaged objective (7.21) has the differential of the collective free energy restricted to fixed outside marginal. The Fisher metric then need not be block diagonal: for the bivariate normal location family with unit variances and correlation $\mathrm{c}_2$, $0 < |\mathrm{c}_2| < 1$, its off-diagonal entry is $-\mathrm{c}_2/(1-\mathrm{c}_2^2)$.

[ESTABLISHED]

*Proof.* Part (i) is the chain rule for $F$ along the solution. For (ii), a change of $\lambda_i$ leaves $Q_{-i}$ fixed, and theorem 5.23 with $B = \{i\}$, whose finiteness hypotheses hold because $F$ is real, shows that $F$ and $\overline{\mathcal{F}}_i$ differ by a quantity independent of $\lambda_i$. Scores of distinct factors of a product law are independent with mean zero, so their cross covariances vanish. For (iii), the Fisher metric of the location family is the inverse covariance matrix. □

## 8.2 Fast recognition and slow model coordinates

**Definition 8.2 (Fast and slow coordinate blocks).** Let $\mathcal{U}_{\mathrm{f}} \subseteq \mathbb{R}^{n_{\mathrm{f}}}$ and $\mathcal{U}_{\mathrm{s}} \subseteq \mathbb{R}^{n_{\mathrm{s}}}$ be open, and write $s = (s_{\mathrm{m}}, \theta) \in \mathcal{U}_{\mathrm{s}}$. A fast–slow coordinatization is a map $(f, s) \mapsto (Q_{f,s_{\mathrm{m}}}, \theta)$ into population recognition laws and generative parameters, together with the real-valued objective

$$F(f,s) := \mathcal{F}_o^{\mathrm{ext}}(Q_{f,s_{\mathrm{m}}};\theta) = -\log p_{\theta,V}(o \mid X) + D_{\mathrm{KL}}\big(Q_{f,s_{\mathrm{m}}} \| \mathbf{\Pi}_{\theta,V,o,X}\big), \tag{8.3}$$

in which the generative argument of (5.17) is explicit. The fast coordinate $f$ carries the belief-channel recognition coordinates and, when present, the source-label rows of section 6.1; the slow coordinate carries the model-channel recognition coordinates $s_{\mathrm{m}}$, which move a recognition law, and the generative parameter $\theta$, which moves the fixed joint between evaluations and never inside one, so a composite $\theta(Q)$ inside one evaluation of (8.3) stays excluded by hypothesis 4.3. The assignment of coordinates to the two blocks is a declaration.

**Lemma 8.3 (Envelope identity and reduced slow descent).** Let $F \in C^1(\mathcal{U}_\mathrm{f} \times \mathcal{U}_\mathrm{s})$, let $\mathcal{U}_0 \subseteq \mathcal{U}_\mathrm{s}$ be open, and let $f^\star : \mathcal{U}_0 \to \mathcal{U}_\mathrm{f}$ be a $C^1$ map with $F(f^\star(s), s) = \min_{f \in \mathcal{U}_\mathrm{f}} F(f, s)$ for every $s \in \mathcal{U}_0$. The reduced functional

$$F_\mathrm{red}(s) := \min_{f \in \mathcal{U}_\mathrm{f}} F(f, s) = F(f^\star(s), s) \tag{8.4}$$

satisfies the envelope identity

$$dF_\mathrm{red}(s) = \partial_s F(f^\star(s), s), \qquad s \in \mathcal{U}_0. \tag{8.5}$$

For every continuous field $\mathsf{G}_\mathrm{s}$ of symmetric positive-definite metrics on $\mathcal{U}_\mathrm{s}$, the reduced slow flow

$$\dot{s} = -\mathsf{G}_\mathrm{s}(s)^{-1} \partial_s F(f^\star(s), s) \tag{8.6}$$

is the $\mathsf{G}_\mathrm{s}$ natural-gradient flow of $F_\mathrm{red}$, and along the lifted curve $(f^\star(s(t)), s(t))$ the free energy decreases at the rate $dF_\mathrm{red} \cdot \mathsf{G}_\mathrm{s}^{-1} dF_\mathrm{red}$. A $C^1$ branch exists near $s_0 \in \mathcal{U}_0$ whenever $F \in C^2$, $\partial_{ff}^2 F(f^\star(s_0), s_0)$ is positive definite, and $f^\star$ is continuous at $s_0$. If $\mathcal{U}_\mathrm{f}$ is replaced by a compact set on which $F$ and $\partial_s F$ are continuous and the conditional minimizer is unique, (8.5) holds without differentiability of $f^\star$. [ESTABLISHED]

*Proof.* At the minimizer $\partial_f F(f^\star(s), s) = 0$, so the chain rule gives $dF_\mathrm{red} = \partial_f F \cdot Df^\star + \partial_s F = \partial_s F$, and substitution in (8.6) gives the flow and its rate. For the branch, the implicit function theorem applied to $\partial_f F = 0$ at $(f^\star(s_0), s_0)$ gives a $C^1$ map whose graph is the whole critical set near that point, and continuity places $f^\star(s)$ on it for $s$ near $s_0$. In the compact case, uniqueness and compactness make $f^\star$ continuous, since a limit point $f'$ of $f^\star(s_n)$ with $s_n \to s$ satisfies $F(f', s) \le F(f^\star(s), s)$. With $s' = s + \Delta s$, the difference $F_\mathrm{red}(s') - F_\mathrm{red}(s)$ lies between $F(f^\star(s'), s') - F(f^\star(s'), s)$ and $F(f^\star(s), s') - F(f^\star(s), s)$, and by the mean value theorem and uniform continuity of $\partial_s F$ on a compact neighborhood both equal $\partial_s F(f^\star(s), s)\Delta s + o(|\Delta s|)$; this is the theorem of Danskin [25, Ch. III] in the unique-minimizer case. □

**Proposition 8.4 (Rate-separated natural gradient and its reduced flow).** Let $F \in C^1(\mathcal{U}_\mathrm{f} \times \mathcal{U}_\mathrm{s})$, let $\mathsf{G}_\mathrm{f}(f, s)$ and $\mathsf{G}_\mathrm{s}(s)$ be continuous symmetric positive-definite block metrics with $\mathsf{G}_\mathrm{s}$ independent of $f$, and for $\varepsilon > 0$ set $\mathsf{G}_\varepsilon = \varepsilon \mathsf{G}_\mathrm{f} \oplus \mathsf{G}_\mathrm{s}$. The $\mathsf{G}_\varepsilon$ natural gradient of $F$ is the rate-separated flow

$$\varepsilon \dot{f} = -\mathsf{G}_\mathrm{f}^{-1} \partial_f F, \qquad \dot{s} = -\mathsf{G}_\mathrm{s}^{-1} \partial_s F, \tag{8.7}$$

which decreases $F$ at the rate $\varepsilon^{-1} \partial_f F \cdot \mathsf{G}_\mathrm{f}^{-1} \partial_f F + \partial_s F \cdot \mathsf{G}_\mathrm{s}^{-1} \partial_s F$ for every $\varepsilon > 0$; its orbit can depend on $\varepsilon$. Under the hypotheses of lemma 8.3, let $\mathcal{S}^\star = \{(f^\star(s), s) : s \in \mathcal{U}_0\}$. At every point of $\mathcal{S}^\star$ and for every $\varepsilon > 0$, the fast component of (8.7) vanishes and the slow component equals $-\mathsf{G}_\mathrm{s}^{-1} dF_\mathrm{red}$. The reduced field $\mathscr{V}^\star = (Df^\star v, v)$ with $v = -\mathsf{G}_\mathrm{s}^{-1} dF_\mathrm{red}$ is tangent to $\mathcal{S}^\star$, and

$$d \operatorname{pr}_\mathrm{s} \mathscr{V}^\star = -\operatorname{grad}^{\mathsf{G}_\mathrm{s}} F_\mathrm{red} \circ \operatorname{pr}_\mathrm{s} \quad \text{on } \mathcal{S}^\star. \tag{8.8}$$

Thus $(\operatorname{pr}_\mathrm{s}|_{\mathcal{S}^\star}, 1)$ is an oriented semiconjugacy in the sense of definition 7.45, with the manifold $\mathcal{S}^\star$ in place of an open fine domain, and a conjugacy, because $\operatorname{pr}_\mathrm{s}|_{\mathcal{S}^\star}$ is a diffeomorphism onto $\mathcal{U}_0$.

The speeds satisfy $\|\mathscr{V}^\star\|^2_{\mathsf{G}_\varepsilon} = \varepsilon\|Df^\star v\|^2_{\mathsf{G}_\mathrm{f}} + \|v\|^2_{\mathsf{G}_\mathrm{s}}$, so the $\mathsf{G}_\varepsilon$-length of a reduced orbit is at least the $\mathsf{G}_\mathrm{s}$-length of its image. These lengths decrease to the latter as $\varepsilon \downarrow 0$ on every compact time interval of the orbit, and on any interval where the lifted length is finite at some positive $\varepsilon_0$. [ESTABLISHED]

*Proof.* Because $\mathsf{G}_\varepsilon^{-1} = \varepsilon^{-1}\mathsf{G}_\mathrm{f}^{-1} \oplus \mathsf{G}_\mathrm{s}^{-1}$, the flow is (8.7) with the stated dissipation, and the orbit statement is proposition 7.31, since $\mathsf{G}_\varepsilon$ differs from $\mathsf{G}_1$ by an anisotropic rescaling. On $\mathcal{S}^\star$, $\partial_f F = 0$ and $\partial_s F = dF_{\mathrm{red}}$ by (8.5), and $\mathsf{G}_\mathrm{s}$ does not depend on $f$. The field $\mathscr{V}^\star$ is the differential of $s \mapsto (f^\star(s), s)$ applied to $v$, hence tangent to $\mathcal{S}^\star$ with slow component $v$. The length identity is the block form of $\mathsf{G}_\varepsilon$. For $0 < \varepsilon \le \varepsilon_0$ the lifted speed is bounded by its $\varepsilon_0$ value, which is integrable under the stated hypothesis and on compact time intervals by continuity. Dominated convergence gives the length limit. □

**Remark 8.5 (The slow graph off the optimum and under correlated metrics).** Off $\mathcal{S}^\star$ the slow component $-\mathsf{G}_\mathrm{s}^{-1}\partial_s F(f,s)$ need not be a positive multiple of $-\operatorname{grad}^{\mathsf{G}_\mathrm{s}} F_{\mathrm{red}}(s)$: for $F = [(f - s_1)^2 + s_1^2 + s_2^2]/2$ with $\mathsf{G}_\mathrm{s} = I$, $\partial_s F = (2s_1 - f, s_2)$ and $\operatorname{grad} F_{\mathrm{red}} = (s_1, s_2)$ are not parallel when $f \neq s_1$ and $s_2 \neq 0$, so the projection semiconjugates the finite-$\varepsilon$ flow to the reduced flow only where that proportionality holds. At points of $\mathcal{S}^\star$ the finite-$\varepsilon$ velocity $(0, v)$ is tangent to $\mathcal{S}^\star$ only where $Df^\star v = 0$, so $\mathcal{S}^\star$ is not invariant in general. A correlated belief–model metric requires the rate declaration to name its slow manifold. For $F(f,s) = [(f-s)^2 + s^2]/2$, whose conditional minimizer is $f = s$, and the metric $\mathsf{G}$ with unit diagonal and off-diagonal entry $\mathrm{c}_2 \in (0,1)$, the row-rated field $\varepsilon\dot{f} = -(\mathsf{G}^{-1}dF)_f$, $\dot{s} = -(\mathsf{G}^{-1}dF)_s$ has, in the fast time $t/\varepsilon$, the critical set

$$\{(\mathsf{G}^{-1}dF)_f = 0\} = \Big\{f = \frac{1+2\mathrm{c}_2}{1+\mathrm{c}_2}s\Big\}, \tag{8.9}$$

which attracts at rate $1/(1-\mathrm{c}_2)$ and differs from the conditional minimizer. This field need not descend $F$: at $\mathrm{c}_2 = 1/2$, every $\varepsilon \in (0, 7-4\sqrt{3})$, and $(f,s) = (1+2\delta, 1+\delta)$ with $\delta = (1+\varepsilon)/4$, it gives $dF/dt = (1+\varepsilon)^2/(12\varepsilon) - 4/3 > 0$, so for those rates it is the natural gradient of no metric. The congruence rescaling $\mathsf{G}_\varepsilon = \Delta_\varepsilon \mathsf{G} \Delta_\varepsilon$ with $\Delta_\varepsilon = \operatorname{diag}(\sqrt{\varepsilon}, 1)$ gives a natural gradient that descends $F$ for every $\varepsilon$, and the critical set of its $\varepsilon = 0$ fast subsystem is the conditional-minimizer set $\{\partial_f F = 0\}$. [ESTABLISHED]

*Proof.* The first two statements are direct differentiation. For the correlated metric, $\mathsf{G}^{-1}$ has diagonal entries $1/(1-\mathrm{c}_2^2)$ and off-diagonal entries $-\mathrm{c}_2/(1-\mathrm{c}_2^2)$, and $dF = (f-s, 2s-f)$, so $(\mathsf{G}^{-1}dF)_f$ vanishes on the displayed line, and its derivative in $f$ is $(1+\mathrm{c}_2)/(1-\mathrm{c}_2^2) = 1/(1-\mathrm{c}_2)$. Substituting the stated point into $dF/dt = -\varepsilon^{-1}\partial_f F(\mathsf{G}^{-1}dF)_f - \partial_s F(\mathsf{G}^{-1}dF)_s$ at $\mathrm{c}_2 = 1/2$ gives the displayed value, which is positive exactly when $\varepsilon^2 - 14\varepsilon + 1 > 0$. Under congruence, $\mathsf{G}_\varepsilon$ is positive definite for every $\varepsilon > 0$, and in the fast time the fast component of its natural gradient is $-(1-\mathrm{c}_2^2)^{-1}(\partial_f F - \mathrm{c}_2\sqrt{\varepsilon}\partial_s F)$, whose zero set at $\varepsilon = 0$ is $\{\partial_f F = 0\}$. □

**Proposition 8.6 (Slow model means on the four-cycle and the model-link dichotomy).** In definition 13.21 let the belief means $\bar{k}$ move at rate one, the model means $\bar{m}$ at rate $\varepsilon_\mathrm{m} \in [0,1]$, and

the link translations $\check{\omega}$ at rate $\varepsilon_{\check{\omega}} \in [0,1]$: between clock times

$$d\mathbf{x} = -\big(\partial_{\bar{k}} U_S, \varepsilon_{\mathrm{m}} \partial_{\bar{m}} U_S, \varepsilon_{\check{\omega}} \partial_{\check{\omega}} U_S\big) dt + \sqrt{2\mathrm{T}}\big(dW^{\bar{k}}, \varepsilon_{\mathrm{m}}^{1/2} dW^{\bar{m}}, \varepsilon_{\check{\omega}}^{1/2} dW^{\check{\omega}}\big) \tag{8.10}$$

with one standard Brownian motion $W$ in $\mathbb{R}^{12}$, and the clock redraws the assignment as before, leaving $\mathbf{x}$ unchanged. The rates are supplied, since the Fisher metric of $\mathcal{N}(\bar{m}_i, \sigma_{\mathrm{ev}}^2)$ in $\bar{m}_i$ is $1/\sigma_{\mathrm{ev}}^2$ and a Fisher natural gradient would move the model means at rate $1/2$. For $i \in \{0,2\}$ write $\bar{k}_i^{\pm} = (\bar{k}_{i+1} \pm \bar{k}_i)/\sqrt{2}$ and $\bar{m}_i^{\pm} = (\bar{m}_{i+1} \pm \bar{m}_i)/\sqrt{2}$, so that $r_i^b = \sqrt{2}\bar{k}_i^- - \check{\omega}_i$ and $r_i^m = \sqrt{2}\bar{m}_i^- - \check{\omega}_i$.

(i) *Held matching.* On the held assignment $S^{(0)}$ a Gaussian law of $\mathbf{x}$ stays Gaussian and its mean follows (8.10) without the noise; for $\varepsilon_{\mathrm{m}}, \varepsilon_{\check{\omega}} > 0$, in the slow time $\tau = \varepsilon_{\mathrm{m}} t$, this mean flow is (8.7) for $F = U_{S^{(0)}}$, which is the collective free energy of proposition 13.22 at the Gaussian law up to terms independent of the mean, with $\varepsilon = \varepsilon_{\mathrm{m}}$, $f = \bar{k}$, $s = (\bar{m}, \check{\omega})$, $\mathsf{G}_{\mathrm{f}} = I_4$ and $\mathsf{G}_{\mathrm{s}} = I_4 \oplus (\varepsilon_{\mathrm{m}}/\varepsilon_{\check{\omega}}) I_4$. In the orthonormal coordinates $(\bar{k}_i^-, \bar{m}_i^-, \check{\omega}_i)$ and $(\bar{k}_i^+, \bar{m}_i^+)$ for $i \in \{0,2\}$, $\check{\omega}_1$, and $\check{\omega}_3$, the Hessian of $U_{S^{(0)}}$ and the mobility are block diagonal, with the blocks

$$\begin{pmatrix} 4+3\mathrm{T} & -2\mathrm{T} & -2\sqrt{2} \\ -2\mathrm{T} & 4+3\mathrm{T} & -2\sqrt{2} \\ -2\sqrt{2} & -2\sqrt{2} & 4+\mathrm{T} \end{pmatrix}, \qquad \mathrm{T}\begin{pmatrix} 3 & -2 \\ -2 & 3 \end{pmatrix}, \qquad \mathrm{T}, \tag{8.11}$$

so the mean flow is a direct sum of six flows and $r_i^b, r_i^m$ read only the block of $e_i$. With $\varepsilon_{\check{\omega}} = \varepsilon_{\mathrm{m}}$ the split by mobility, $f = \bar{k}$ and $s = (\bar{m}, \check{\omega})$, is the only split for which the metric has the form $\mathsf{G}_\varepsilon$ with blocks independent of $\varepsilon_{\mathrm{m}}$; restricted to a selected-edge block it satisfies the three hypotheses of open problem 8.7 with $\mathrm{c}_1 = 4$ for every $\mathrm{T} > 0$, on every closed ball whose interior contains the reduced orbit, while on the full sheet its fast-block curvature along the pair means $\bar{k}_i^+$ is $3\mathrm{T}$.

(ii) *Dichotomy at zero temperature.* Hold $S^{(0)}$, drop the temperature terms, and start the mean flow $\dot{\mathbf{x}} = -(\partial_{\bar{k}} V, \varepsilon_{\mathrm{m}} \partial_{\bar{m}} V, \varepsilon_{\check{\omega}} \partial_{\check{\omega}} V)$, $V = V_{S^{(0)}}$, from $\mathbf{x}^\circ$, where every residual equals $-1$. Unless $(\varepsilon_{\mathrm{m}}, \varepsilon_{\check{\omega}}) \in \{(0,0),(1,0)\}$, the flow matrix of a selected-edge block has the simple eigenvalue $0$ on $(1,1,\sqrt{2})$, which both residual rows annihilate, and the two rates

$$\varsigma_\pm = 2(1 + \varepsilon_{\mathrm{m}} + \varepsilon_{\check{\omega}}) \pm 2\sqrt{(1-\varepsilon_{\mathrm{m}})^2 + \varepsilon_{\check{\omega}}^2}, \tag{8.12}$$

and each internal residual is the combination of $e^{-\varsigma_- t}$ and $e^{-\varsigma_+ t}$ with value $-1$ and derivative $4(\varepsilon_{\mathrm{m}} + \varepsilon_{\check{\omega}})$ (model) or $4(1 + \varepsilon_{\check{\omega}})$ (belief) at $t = 0$. With a fast link, $\varepsilon_{\check{\omega}} = 1$, one has $\varsigma_- \geq 4 - 2\sqrt{2}$ for every $\varepsilon_{\mathrm{m}} \in [0,1]$ and $|r_i^m(t)| \leq e^{-(4-2\sqrt{2})t}$: the link removes the model disagreement on the belief clock even when the model means are frozen. With fast model means, $\varepsilon_{\mathrm{m}} = 1$ and $\varepsilon_{\check{\omega}} \in (0,1]$, $r_i^m(t) = -e^{-(4+4\varepsilon_{\check{\omega}})t}$. With slow model means and links, $\varepsilon_{\check{\omega}} = \varepsilon_{\mathrm{m}} \leq 1/2$, one has $6\varepsilon_{\mathrm{m}} - 2\varepsilon_{\mathrm{m}}^2 \leq \varsigma_- \leq 6\varepsilon_{\mathrm{m}}$, $\varsigma_+ \geq 4$, and

$$|r_i^b(t)| \leq (1+\varepsilon_{\mathrm{m}})e^{-4t} + \varepsilon_{\mathrm{m}}, \qquad \big(1 - \tfrac{3}{2}\varepsilon_{\mathrm{m}}\big)e^{-6\varepsilon_{\mathrm{m}} t} - \tfrac{3}{2}\varepsilon_{\mathrm{m}} \leq |r_i^m(t)| \leq e^{-(6\varepsilon_{\mathrm{m}} - 2\varepsilon_{\mathrm{m}}^2)t} + \tfrac{3}{2}\varepsilon_{\mathrm{m}}, \tag{8.13}$$

so for $\varepsilon_{\mathrm{m}} \leq 1/25$ each matched pair has $|r_i^b| \leq 2\varepsilon_{\mathrm{m}}$ and $|r_i^m| \geq 1/2$ at every time from $\frac{1}{4}\log(4/\varepsilon_{\mathrm{m}})$ to $1/(12\varepsilon_{\mathrm{m}})$.

(iii) *Persistence along every history.* Let $\mathrm{T} > 0$, let $\mathbf{x}_t$ be the sampled process from any initial configuration and assignment, and let $\bar{\varepsilon} = \max(\varepsilon_{\mathrm{m}}, \varepsilon_{\breve{\omega}}) > 0$. For an edge $e_i$ with $|r_i^m(\mathbf{x}_0)| \geq 1$, numbers $\mathrm{R} \geq 1$ and $\mathrm{R}' > 0$, and $0 < t \leq [4\bar{\varepsilon}(12\mathrm{R} + 9\mathrm{T}\mathrm{R}')]^{-1}$, the event that $|r_i^m(\mathbf{x}_u)| < 1/2$ for some $u \leq t$ while every residual stays bounded by $\mathrm{R}$ and $|\mathbf{x}_u| \leq \mathrm{R}'$ on $[0, t]$ has probability at most $2\exp[-1/(192\mathrm{T}\bar{\varepsilon}t)]$, whatever the assignment path, and at most $2e^{-1/(4\mathrm{T})}$ at the largest such $t$.

In these arms a belief-agreeing, model-disagreeing interval of order $1/\varepsilon_{\mathrm{m}}$ exists exactly when the model means and the links are both slow: with both slow it is the interval of (ii), whose model disagreement (iii) protects along every assignment path up to a time of order $1/\varepsilon_{\mathrm{m}}$ while residuals and configuration stay bounded, and a fast link or fast model means remove the disagreement at a rate at least $4 - 2\sqrt{2}$ whatever the other rate. [ESTABLISHED]

*Proof.* For (i), both endpoints of $e_0$ and of $e_2$ attend them under $S^{(0)}$ and no receiver attends $e_1$ or $e_3$, so $V_{S^{(0)}} = 2\mathcal{D}_{e_0} + 2\mathcal{D}_{e_2}$ with $2\mathcal{D}_{e_i} = (\sqrt{2}\bar{k}_i^- - \breve{\omega}_i)^2 + (\sqrt{2}\bar{m}_i^- - \breve{\omega}_i)^2$, the reference term is invariant under the orthonormal change of coordinates, and the evaluator term is $\mathrm{T}\sum_{i\in\{0,2\}}[(\bar{k}_i^- - \bar{m}_i^-)^2 + (\bar{k}_i^+ - \bar{m}_i^+)^2]$; collecting second derivatives gives (8.11), and the mobility stays diagonal because the new coordinates mix one channel at a time. On a held assignment (8.10) is linear with constant coefficients, so a Gaussian law stays Gaussian and its mean solves the noiseless equation. Since $U_S(\mathbf{x}) = \frac{1}{2}\mathbf{x}^\top\nabla^2 U_S\mathbf{x}$, the identity $\mathcal{F}_{\mathrm{T}}(q) = \mathrm{T}D_{\mathrm{KL}}(q\|q^\star) + \mathrm{B}_{\mathrm{T}} - \mathrm{T}\log Z_{\mathrm{opt}}$ of proposition 13.22 at $q = \mathcal{N}(\mu, \Sigma)\otimes\delta_S$ gives $\mathcal{F}_{\mathrm{T}} = U_S(\mu) + \frac{1}{2}\operatorname{Tr}(\nabla^2 U_S\Sigma) - \frac{\mathrm{T}}{2}\log\det\Sigma$ up to a constant. Dividing the mean equation by $\varepsilon_{\mathrm{m}}$ gives the slow-time form with the metric $\varepsilon_{\mathrm{m}}I_4 \oplus I_4 \oplus (\varepsilon_{\mathrm{m}}/\varepsilon_{\breve{\omega}})I_4$, whose blocks are independent of $\varepsilon_{\mathrm{m}}$ exactly when $f$ consists of the unit-rate coordinates. On a selected-edge block with $\varepsilon_{\breve{\omega}} = \varepsilon_{\mathrm{m}}$, $F$ is the quadratic form of the first block of (8.11), positive definite for $\mathrm{T} > 0$ since its leading minors are $4 + 3\mathrm{T}$, $(4 + \mathrm{T})(4 + 5\mathrm{T})$ and $\mathrm{T}(8 + \mathrm{T})(4 + 5\mathrm{T})$; $\partial_{ff}^2 F = 4 + 3\mathrm{T} \geq 4\mathsf{G}_{\mathrm{f}}$; the reduced matrix is a Schur complement, hence positive definite, so $|s|$ is nonincreasing along the reduced flow and every closed ball whose interior contains the initial slow state serves as $\mathcal{K}$; and the frozen fast flow is linear with coefficient $4 + 3\mathrm{T}$, so every fast state lies in its basin. The pair-mean block gives the curvature $3\mathrm{T}$ along $\bar{k}_i^+$.

For (ii), the block flow matrix is $\operatorname{diag}(1, \varepsilon_{\mathrm{m}}, \varepsilon_{\breve{\omega}})$ times the first block of (8.11) at $\mathrm{T} = 0$, whose kernel is spanned by $(1, 1, \sqrt{2})$, on which $r_i^b = r_i^m = 0$. Expanding the determinant, its characteristic polynomial is $\varsigma[\varsigma^2 - 4(1 + \varepsilon_{\mathrm{m}} + \varepsilon_{\breve{\omega}})\varsigma + 8(2\varepsilon_{\mathrm{m}} + \varepsilon_{\breve{\omega}} + \varepsilon_{\mathrm{m}}\varepsilon_{\breve{\omega}})]$, whose quadratic factor has discriminant $16[(1 - \varepsilon_{\mathrm{m}})^2 + \varepsilon_{\breve{\omega}}^2]$ and constant term vanishing only at $(0, 0)$, so its roots are (8.12) and the eigenvalues are distinct outside the two excluded pairs. The matrix is then diagonalizable and the residual rows annihilate the kernel, so each internal residual is $\mathrm{c}_2 e^{-\varsigma_- t} + \mathrm{c}_3 e^{-\varsigma_+ t}$ with $\mathrm{c}_2 + \mathrm{c}_3 = -1$ and $\varsigma_-\mathrm{c}_2 + \varsigma_+\mathrm{c}_3$ equal to minus the derivative at $t = 0$, which is $a \cdot \operatorname{diag}(1, \varepsilon_{\mathrm{m}}, \varepsilon_{\breve{\omega}})(-2\sqrt{2}, -2\sqrt{2}, 4)^\top$ for the residual row $a$; hence $\mathrm{c}_3 = (\varsigma_- - \mathrm{c}_5)/(\varsigma_+ - \varsigma_-)$ with $\mathrm{c}_5 = 4(\varepsilon_{\mathrm{m}} + \varepsilon_{\breve{\omega}})$ for $r_i^m$ and $\mathrm{c}_5 = 4(1 + \varepsilon_{\breve{\omega}})$ for $r_i^b$. For $\varepsilon_{\breve{\omega}} = 1$ put $\mathrm{c}_4 = \sqrt{(1 - \varepsilon_{\mathrm{m}})^2 + 1}$, so $\varsigma_\pm = 4 + 2\varepsilon_{\mathrm{m}} \pm 2\mathrm{c}_4$; $\varsigma_-$ increases in $\varepsilon_{\mathrm{m}}$ from $4 - 2\sqrt{2}$, and $\varsigma_- \leq 4 + 4\varepsilon_{\mathrm{m}} \leq \varsigma_+$ because $\mathrm{c}_4 \geq \varepsilon_{\mathrm{m}}$, so both coefficients of $r_i^m$ are nonpositive with sum $-1$. For $\varepsilon_{\mathrm{m}} = 1$, $\varsigma_- = 4 = \mathrm{c}_5 - 4\varepsilon_{\breve{\omega}}$ gives $\mathrm{c}_3 = -1$. For $\varepsilon_{\breve{\omega}} = \varepsilon_{\mathrm{m}} \leq 1/2$ put $\mathrm{c}_4 = \sqrt{(1 - \varepsilon_{\mathrm{m}})^2 + \varepsilon_{\mathrm{m}}^2}$, so that $1 - \varepsilon_{\mathrm{m}} \leq \mathrm{c}_4 \leq 1 - \varepsilon_{\mathrm{m}} + \varepsilon_{\mathrm{m}}^2$ by squaring, $\varsigma_\pm = 2 + 4\varepsilon_{\mathrm{m}} \pm 2\mathrm{c}_4$ and $\varsigma_+ - \varsigma_- = 4\mathrm{c}_4 \geq 4(1 - \varepsilon_{\mathrm{m}})$, which gives the stated brackets; the coefficient of $e^{-\varsigma_+ t}$ in $r_i^m$ is $(\varsigma_- - 8\varepsilon_{\mathrm{m}})/(4\mathrm{c}_4)$, of absolute value at

most $\varepsilon_{\mathrm{m}}(1+\varepsilon_{\mathrm{m}})/(2(1-\varepsilon_{\mathrm{m}})) \le \frac{3}{2}\varepsilon_{\mathrm{m}}$, and the coefficient of $e^{-\varsigma_- t}$ in $r_i^b$ is $(\varsigma_- - 4\varepsilon_{\mathrm{m}})/(4\mathrm{c}_4) \in [0, \varepsilon_{\mathrm{m}}]$. Since each pair of coefficients sums to $-1$, (8.13) follows, and the last claim is (8.13) at the two times, using $(1+\varepsilon_{\mathrm{m}})/4 \le 1$ and $e^{-1/2}(1-\frac{3}{50}) - \frac{3}{50} > \frac{1}{2}$.

For (iii), write $a$ for the row of $r_i^m$, with entries $-1$ and $1$ at $\bar{m}_i$ and $\bar{m}_{i+1}$ and $-1$ at $\check{\omega}_i$. Since the clock leaves $\mathbf{x}$ unchanged, (8.10) gives, along every assignment path, $r_i^m(\mathbf{x}_t) - r_i^m(\mathbf{x}_0) = -\int_0^t a \cdot (\partial_{\bar{k}} U_{S_u}, \varepsilon_{\mathrm{m}} \partial_{\bar{m}} U_{S_u}, \varepsilon_{\check{\omega}} \partial_{\check{\omega}} U_{S_u})(\mathbf{x}_u) du + \sqrt{2\mathrm{T}(2\varepsilon_{\mathrm{m}} + \varepsilon_{\check{\omega}})} B_t$ with $B$ a standard Brownian motion. With $\mathrm{n}_j(S) \in \{0,1,2\}$ the number of receivers attending $e_j$ under $S$, $\partial_{\bar{m}_j} U_S = \mathrm{n}_{j-1}(S) r_{j-1}^m - \mathrm{n}_j(S) r_j^m + 3\mathrm{T}\bar{m}_j - 2\mathrm{T}\bar{k}_j$ and $\partial_{\check{\omega}_j} U_S = -\mathrm{n}_j(S)(r_j^b + r_j^m) + \mathrm{T}\check{\omega}_j$, so on the stated event the integrand is bounded by $\varepsilon_{\mathrm{m}}(8\mathrm{R} + 8\mathrm{TR}') + \varepsilon_{\check{\omega}}(4\mathrm{R} + \mathrm{TR}') \le \bar{\varepsilon}(12\mathrm{R} + 9\mathrm{TR}')$, using $3|\bar{m}_j| + 2|\bar{k}_j| \le \sqrt{13}|\mathbf{x}|$. For $t$ as stated the drift moves the residual by at most $1/4$, so a crossing below $1/2$ requires $\sup_{u \le t} |B_u| \ge [4\sqrt{2\mathrm{T}(2\varepsilon_{\mathrm{m}} + \varepsilon_{\check{\omega}})}]^{-1} \ge [4\sqrt{6\mathrm{T}\bar{\varepsilon}}]^{-1}$. Doob's maximal inequality applied to the martingale $\exp(\theta B_u - \theta^2 u/2)$ with $\theta = \mathrm{y}/t$ gives $\Pr(\sup_{u \le t} B_u \ge \mathrm{y}) \le e^{-\mathrm{y}^2/(2t)}$, and with $\mathrm{y}^2 = 1/(96\mathrm{T}\bar{\varepsilon})$ the two-sided bound is $2e^{-1/(192\mathrm{T}\bar{\varepsilon}t)}$; at the largest $t$ the exponent is $-(12\mathrm{R} + 9\mathrm{TR}')/(48\mathrm{T}) \le -1/(4\mathrm{T})$. The closing sentence collects (ii) and (iii), the fast-means case having rate $4 + 4\varepsilon_{\check{\omega}} \ge 4$. □

**Open Problem 8.7 (Slow-manifold reduction of the collective free energy).** Consider the rate-separated flow (8.7) under three hypotheses. First, $F \in C^2$ and the block metric $\mathsf{G}_{\mathrm{f}}(f,s) \oplus \mathsf{G}_{\mathrm{s}}(s)$ has $C^1$ symmetric positive-definite blocks with $\mathsf{G}_{\mathrm{s}}$ independent of $f$. Second, there are a compact set $\mathcal{K} \subset \mathcal{U}_0$, a $C^1$ branch $f^\star$ of conditional minimizers, and $\mathrm{c}_1 > 0$ with $\partial_{ff}^2 F(f^\star(s), s) \succeq \mathrm{c}_1 \mathsf{G}_{\mathrm{f}}(f^\star(s), s)$ for all $s \in \mathcal{K}$. Third, the reduced flow (8.6) from $s_0$ stays in the interior of $\mathcal{K}$ on $[0, t_1]$, and the initial fast state lies in the basin of $f^\star(s_0)$ for the frozen fast flow $\dot{f} = -\mathsf{G}_{\mathrm{f}}^{-1} \partial_f F(f, s_0)$. Under the second hypothesis the linearized frozen fast flow has spectrum in $(-\infty, -\mathrm{c}_1]$, so $f^\star(s)$ is isolated and uniformly asymptotically stable on $\mathcal{K}$, and Tikhonov's theorem gives $s_\varepsilon(t) \to s(t)$ for $0 \le t \le t_1$ and $f_\varepsilon(t) \to f^\star(s(t))$ for $0 < t \le t_1$ as $\varepsilon \to 0$ [76]; for smooth data Fenichel's theory gives a locally invariant slow manifold $O(\varepsilon)$-close to a compact normally hyperbolic piece of $\mathcal{S}^\star$ [30]. Their application at $c_*$ is open beyond the held matching sheets of proposition 8.6(i), whose selected-edge blocks meet the three hypotheses with a curvature constant uniform in the temperature, so that Tikhonov's theorem applies to them, while the split by mobility on the full sheet has a constant of order $\mathrm{T}$: proposition 6.8 shows that the second hypothesis can fail after the source rows are optimized even when every source energy and every frozen-weight curvature is positive, the sampled dynamics with source jumps of chapter 13 lie outside this deterministic statement, and no explicit bound on the slow orbit or its Fisher duration in terms of $\varepsilon$ is available. It is settled in the class of the sampled dynamics of definition 13.21 under the mobility of proposition 8.6 with $\varepsilon_{\check{\omega}} = \varepsilon_{\mathrm{m}}$ and no parameter learning, where the split by mobility, $f = \bar{k}$ and $s = (\bar{m}, \check{\omega})$, and (8.6) for $F = U_S$ under the current assignment $S$ define the reduced slow flow of an assignment path: positively, by a bound on the distance between the slow coordinates of the sampled process and that reduced flow over a bounded slow-time interval, explicit in $\varepsilon_{\mathrm{m}}$, $\mathrm{T}$, and the number of assignment changes, tending to zero with $\varepsilon_{\mathrm{m}}$ and $\mathrm{T}$, and holding, for every positive integer $\mathrm{n}$, with probability tending to one on the event that the assignment path has at most $\mathrm{n}$ changes per unit slow time; negatively, by a positive integer $\mathrm{n}$, a sequence of mobility ratios and temperatures tending to zero, and initial states from which, with probability bounded below, the assignment path has at most $\mathrm{n}$ changes

per unit slow time and that distance stays bounded below over a bounded slow-time interval. The learning case, in which the generative parameter is the slow coordinate and the reduced object is the recognition law at the frozen parameter, lies outside this class.

## 8.3 Parent-sensitive inference and the prepared parent

**Definition 8.8 (Parent-sensitive sweep).** Let $\Xi$ be a nonempty finite set of fine states with a parent map $\mathrm{par} : \Xi \to \mathsf{Z}$ that retains the complete decorated parent, keeping only parent values with nonempty fibers. Take a prior $p^0$ on $\Xi$ with $p^0(\xi) > 0$, normalized conditional laws $\psi_\xi$ of an unobserved record $r$ in a finite set $\mathsf{O}'$, and record likelihoods $L_o(\xi) > 0$ at the observation, none of which reads recognition. The generative law of $(\xi, r, o)$ has density $p^0(\xi)\psi_\xi(r)L_o(\xi)$, and its evidence and posterior are

$$\mathcal{Z}_o = \sum_\xi p^0(\xi)L_o(\xi), \qquad p(\xi) = \frac{p^0(\xi)L_o(\xi)}{\mathcal{Z}_o}. \tag{8.14}$$

On an undirected graph of allowed moves take symmetric proposal weights $\mathsf{a}_{\xi\xi'} = \mathsf{a}_{\xi'\xi} \geq 0$ with off-diagonal row sums at most $1/2$. The Metropolis kernel, the parent projection, the posterior lift, the refresh, and the sweep are

$$\begin{aligned} &\mathsf{M}_{\xi\xi'} = \mathsf{a}_{\xi\xi'} \min\bigl(1, p(\xi')/p(\xi)\bigr) \quad (\xi \neq \xi'), \quad \mathsf{M}_{\xi\xi} = 1 - \sum_{\xi' \neq \xi} \mathsf{M}_{\xi\xi'}, \\ &\mathsf{C}_{\xi z} = \mathbf{1}_{\mathrm{par}(\xi)=z}, \qquad \mathsf{J}_{z\xi} = \frac{p(\xi)\mathbf{1}_{\mathrm{par}(\xi)=z}}{p_{\mathsf{Z}}(z)}, \qquad p_{\mathsf{Z}} = p\mathsf{C}, \\ &\mathsf{D} = \mathsf{CJ}, \qquad \mathsf{T} = \mathsf{MD}. \end{aligned} \tag{8.15}$$

The sweep depends on the parent through the conditional laws $p(\cdot \mid z)$. The marked version of a kernel $\mathsf{K}$ on $\Xi$ is $\widetilde{\mathsf{K}}((\xi, r), (\xi', r')) = \mathsf{K}_{\xi\xi'}\psi_{\xi'}(r')$. The proposal and the update order are declared choices.

**Hypothesis 8.9 (Instantaneous within-parent refresh).** In this section the refresh is the exact kernel $\mathsf{D} = \mathsf{CJ}$: each application resamples the fine state from its exact conditional law given the current complete parent. This is the adiabatic elimination of the within-parent fine state, whose timescale is infinitely shorter than the parent's. It excludes finite within-fiber mixing, for example the partial refresh $\mathsf{D}_\omega = (1-\omega)I + \omega\mathsf{CJ}$ with $0 \leq \omega < 1$, which is equally normalized, $p$-reversible, and parent preserving.

**Theorem 8.10 (Stationarity, exact descent, and finite-component convergence).**

(i) For a recognition law $q(\xi)\psi'_\xi(r)$ the free energy is

$$\mathcal{F}(q \otimes \psi'; o) = D_{\mathrm{KL}}(q\|p) + \mathbb{E}_q D_{\mathrm{KL}}(\psi'_\xi\|\psi_\xi) - \log \mathcal{Z}_o, \tag{8.16}$$

so matching the record conditionals, $\psi' = \psi$, is a specialization of the joint objective and not the removal of its dependence terms.

(ii) The kernels M, D, T are normalized and preserve $p$, and M and D are $p$-reversible; their marked versions are normalized and preserve $p(\xi)\psi_\xi(r)$.

(iii) $q$D minimizes $D_{\mathrm{KL}}(\cdot\|p)$ among laws with parent marginal $q$C, with $D_{\mathrm{KL}}(q\mathsf{D}\|p) = D_{\mathrm{KL}}(q\mathsf{C}\|p_Z)$.

(iv) Call a component a connected component of the union of the positive off-diagonal edges of M and D. For every initial $q$, T preserves component masses, and for every $p$-stationary kernel K on $\Xi$, with $q' = q\mathsf{K}$ and reversal $\mathsf{K}^*_{\xi'\xi} = p(\xi)\mathsf{K}_{\xi\xi'}/p(\xi')$,

$$D_{\mathrm{KL}}(q\|p) - D_{\mathrm{KL}}(q\mathsf{K}\|p) = \sum_{\xi':q'(\xi')>0} q'(\xi') D_{\mathrm{KL}}\big(q(\cdot \mid \xi')\|\mathsf{K}^*_{\xi'\cdot}\big), \tag{8.17}$$

where $q(\xi \mid \xi') = q(\xi)\mathsf{K}_{\xi\xi'}/q'(\xi')$. Hence T decreases (8.16), strictly unless $q(\xi) = q(\Xi')p(\xi \mid \Xi')$ on each component $\Xi'$.

(v) On each component, $q\mathsf{T}^n$ converges geometrically to that conditional target.

[ESTABLISHED]

*Proof.* For (i), the posterior of $(\xi, r)$ given $o$ is $p(\xi)\psi_\xi(r)$ because $r$ is unobserved and conditionally independent of $o$, and the evidence is $\mathcal{Z}_o$; definition 5.6 and the chain rule lemma 3.12 give the display. For (ii), $p(\xi)\mathsf{M}_{\xi\xi'} = \mathsf{a}_{\xi\xi'}\min(p(\xi), p(\xi')) = p(\xi')\mathsf{M}_{\xi'\xi}$; $\mathsf{JC} = I$ gives $p\mathsf{D} = p$, and $p(\xi)\mathsf{D}_{\xi\xi'} = p(\xi)p(\xi')/p_Z(z)$ on a common parent fiber $z$. Stationarity survives composition; summing the terminal record of a marked kernel proves normalization and summing the initial record proves invariance. Part (iii) is corollary 5.22 for the block of fine coordinates given the parent: $q$D has parent marginal $q\mathsf{CJC} = q\mathsf{C}$ and conditionals $p(\cdot \mid z)$, and the chain rule over the parent gives the value. For (iv), the input–output laws $q(\xi)\mathsf{K}_{\xi\xi'}$ and $p(\xi)\mathsf{K}_{\xi\xi'}$ have divergence $D_{\mathrm{KL}}(q\|p)$, and disintegrating them over the output with lemma 3.12 gives (8.17). With positive holding, equality forces $q/p$ to be constant across every positive edge; equality for M makes $q$ stationary under M, and equality for the subsequent D makes the ratio constant on parent fibers, which is constancy on components. For (v), positive holding makes every M or D edge a positive T edge, so on a component $\Xi'$ of $n$ states $\mathsf{T}^k > 0$ for $k = \max(1, n-1)$. With target $p_{\Xi'} = p(\cdot \mid \Xi')$ and $\mathrm{c}_1 = \min_{\xi\xi'}(\mathsf{T}^k)_{\xi\xi'}/p_{\Xi'}(\xi') \in (0, 1]$, one has $\mathsf{T}^k = \mathrm{c}_1\mathbf{1}p_{\Xi'} + (1-\mathrm{c}_1)\mathsf{K}'$ on $\Xi'$ with a $p_{\Xi'}$-stationary stochastic K′ (or $\mathsf{T}^k = \mathbf{1}p_{\Xi'}$ when $\mathrm{c}_1 = 1$), and convexity with data processing give $D_{\mathrm{KL}}(q\mathsf{T}^k\|p_{\Xi'}) \le (1-\mathrm{c}_1)D_{\mathrm{KL}}(q\|p_{\Xi'})$ for laws $q$ on $\Xi'$. □

**Theorem 8.11 (The prepared parent is Markov).** Under hypothesis 8.9, the reduced sweep $\overline{\mathsf{T}} := \mathsf{JMC}$ is reversible under $p_Z$ and

$$\mathsf{JT} = \overline{\mathsf{T}}\mathsf{J}, \qquad (\nu\mathsf{J})\mathsf{T}^n = \nu\overline{\mathsf{T}}^n\mathsf{J}. \tag{8.18}$$

For a start $q_0 = \nu\mathsf{J}$ the parent history $Z_t = \mathrm{par}(\xi_t)$ of the chain with kernel T is Markov with transition $\overline{\mathsf{T}}$, and for an arbitrary start it is so after the first completed refresh. [ESTABLISHED]

*Proof.* With $\mathsf{T} = \mathsf{MCJ}$ and $\mathsf{JC} = I$, $\mathsf{JT} = \mathsf{JMCJ} = \overline{\mathsf{T}}\mathsf{J}$, and induction gives the second identity. Moreover $p_Z(z)\overline{\mathsf{T}}_{zz'}$ is the sum of $p(\xi)\mathsf{M}_{\xi\xi'}$ over $\mathrm{par}(\xi) = z$ and $\mathrm{par}(\xi') = z'$, which is symmetric by reversibility of M. The final refresh makes the fine law given the whole parent history $\mathsf{J}_{z\cdot}$, with $z$ the last parent, so the next parent row is $\mathsf{J}_{z\cdot}\mathsf{MC} = \overline{\mathsf{T}}_{z\cdot}$, and induction over histories gives the Markov property. □

The refresh holds the fast fine state at its conditional optimum given the slow parent (theorem 8.10(iii)), so $\mathsf{JT} = \overline{\mathsf{T}}\mathsf{J}$ is the discrete instantaneous-equilibration case of (8.8).

**Remark 8.12 (Finite refresh leaves parent memory).** Hypothesis 8.9 cannot be dropped. Take $\Xi = \{1, 2, 3\}$ with $\mathrm{par}(1) = \mathrm{par}(2) = z_1$ and $\mathrm{par}(3) = z_2$, uniform $p$, the single proposal $\mathsf{a}_{23} = 1/2$, the partial sweep $\mathsf{T}_\omega = \mathsf{MD}_\omega$, and a start $\nu\mathsf{J}$ with $\nu(z_1), \nu(z_2) > 0$. Then

$$\mathbb{P}(Z_2 = z_2 \mid Z_1 = z_1, Z_0 = z_1) = 1/6 + \omega/12, \qquad \mathbb{P}(Z_2 = z_2 \mid Z_1 = z_1, Z_0 = z_2) = 1/2 - \omega/4. \tag{8.19}$$

The dependence on $Z_0$ is $(1-\omega)/3$, equal to $7/30$ at $\omega = 3/10$, and it vanishes only at $\omega = 1$, where both values equal $\overline{\mathsf{T}}_{z_1 z_2} = 1/4$. For $\omega < 1$ the parent history is not Markov. It remains a deterministic function of the fine Markov chain with kernel $\mathsf{T}_\omega$, so the hidden-history law of section 11.8.1 retains it exactly, and the memory cost of a first-order parent closure is the defect $D_N^{(1)}$ of theorem 11.44, whose value $I(Z_2; Z_0 \mid Z_1)$ at $N = 3$ is positive. [ESTABLISHED]

*Proof.* Only state $2$ moves under M, to $3$ with probability $1/2$, and $\mathsf{D}_\omega$ keeps a state of $\{1, 2\}$ with probability $1 - \omega/2$. Given $Z_1 = z_1$, the state before the second sweep is $2$ with probability $1 - \omega/2$ if $Z_0 = z_2$, since it moved from $3$, and with probability $1/3 + \omega/6$ if $Z_0 = z_1$, since after M it is $2$ with probability $1/3$ and the refresh redraws it uniformly with probability $\omega$; multiplying by $1/2$ gives (8.19). Both conditioning events have positive probability, so $I(Z_2; Z_0 \mid Z_1) > 0$. □

## Physicist's summary

Fast and slow are declared, not derived. The collective free energy is then a potential for every positive-definite mobility. The reduced gradient flow is conjugate to its tangent lift on the graph of the conditional optimum; the finite-rate flow shares its slow component there but can leave that graph. Instantaneous within-parent refresh gives the Markov parent proved here. On the four-cycle, slow model means lag the beliefs only if the links are slow too; a fast link carries the model agreement on the belief clock.

# Part II

# Exact Coarse-Graining at One Base Point

# Chapter 9

# Coarse Channels

Part II asks what a coarse description of the agent network at $c_*$ keeps of the joint of Chapter 4, of its posterior, and of a recognition law. For one coarse channel the answer is Theorem 9.6: a channel that pushes all three laws preserves the evidence and the posterior version, and splits the free energy into the coarse free energy plus the expected relative entropy of the discarded conditional laws. Corollary 9.10 specializes it to a candidate parent of a block of agents, and the exact results of Chapters 10 to 13 rest on both. The chapter also separates Markov pushforward, precomposition, and Galerkin restriction (Section 9.4) and closes with finite labeled families, in which a retained label carries information of order $N^2$ at a tie among $N$ agents (Section 9.8). Fisher contraction under a channel is Theorem 7.8 and is not restated. The longer proofs of the chapter are deferred to appendix B.

**Scope.** The common-channel identity concerns one channel applied to all three laws with the observation unchanged. It does not compare different generative and recognition channels, a coarse joint fitted rather than pushed forward, or a channel that also coarsens the observation, even an invertible or evidence-sufficient one, because each changes the comparison target or the evidence event; nor is a Fisher tensor of a labeled family used as the metric of a dynamics. [NOT-CLAIMED] A natural selector of partitions or membership laws is Open Problem 9.19, and equivariant Bayes versions beyond the finite and compact cases are Open Problem 9.18.

## 9.1 Markov coarse maps

**Definition 9.1 (Coarse channel).** A coarse channel between standard Borel spaces X and Y is a Markov kernel $K : \mathsf{X} \rightsquigarrow \mathsf{Y}$ in the sense of Section 1.4. It sends a law $P$ on X to $(PK)(B) = \int_{\mathsf{X}} K(x, B)P(dx)$ and a statistical experiment $(\mathsf{X}, \{P_\theta\}_{\theta\in\Theta})$ to $(\mathsf{Y}, \{P_\theta K\}_{\theta\in\Theta})$. A measurable statistic $c : \mathsf{X} \to \mathsf{Y}$ is the deterministic kernel $K_c(x, B) = \mathbf{1}_B(c(x))$. A family $K_\theta$ that varies with the parameter is a parameter-dependent model transformation, not one channel between experiments.

**Proposition 9.2 (Markov kernels form a category).** If $K : \mathsf{X} \rightsquigarrow \mathsf{Y}$ and $L : \mathsf{Y} \rightsquigarrow \mathsf{Z}$ are Markov kernels, then $(KL)(x, D) = \int_{\mathsf{Y}} L(y, D)K(x, dy)$ is a Markov kernel, composition is associative with the deterministic identity as unit, and

$$P(KL) = (PK)L \tag{9.1}$$

for every law $P$ on X. [ESTABLISHED]

*Proof.* Kernel integration gives measurability in $x$, and monotone convergence gives countable additivity in $D$. By Tonelli's theorem $[P(KL)](D) = \int_{\mathsf{X}} \int_{\mathsf{Y}} L(y, D)K(x, dy)P(dx)$, which equals $\int_{\mathsf{Y}} L(y, D)(PK)(dy)$; this is (9.1), and applied to point masses and a third kernel it proves associativity. These are the morphisms and the composition of the Kleisli category of the Giry probability monad [35]. □

## 9.2 KL contraction, equality, and recovery

Every Markov kernel contracts relative entropy in $[0, +\infty]$ (Lemma 3.12); the equality case needs finiteness.

**Theorem 9.3 (Data processing with its exact equality condition).** Let $P, Q$ be laws on X with $P \ll Q$ and $D_{\mathrm{KL}}(P\|Q) < \infty$, let $K : \mathsf{X} \rightsquigarrow \mathsf{Y}$ be a Markov kernel, and put $\mathbb{J}(dx, dy) = Q(dx)K(x, dy)$, $r = dP/dQ$, and $\bar{r}(y) = \mathbb{E}_{\mathbb{J}}[r(X) \mid Y = y]$. Then $d(PK)/d(QK) = \bar{r}$ and $D_{\mathrm{KL}}(PK\|QK) \leq D_{\mathrm{KL}}(P\|Q)$, with equality exactly when $r(X) = \bar{r}(Y)$ $\mathbb{J}$-almost surely. [ESTABLISHED]

*Proof.* For bounded measurable $f$ on Y, $\int f d(PK) = \mathbb{E}_{\mathbb{J}}[f(Y)r(X)] = \mathbb{E}_{\mathbb{J}}[f(Y)\bar{r}(Y)]$, which is the Radon-Nikodym formula. Conditional Jensen for the strictly convex function $t \mapsto t \log t$ gives $\mathbb{E}_{\mathbb{J}}[\bar{r}(Y) \log \bar{r}(Y)] \leq \mathbb{E}_{\mathbb{J}}[r(X) \log r(X)]$, and finite divergence and strict convexity make equality equivalent to $r(X) = \bar{r}(Y)$ almost surely. This is the information-loss theorem of Csiszár [24], Kullback and Leibler [53] in kernel form. □

**Corollary 9.4 (Pairwise Bayes recovery).** In the setting of Theorem 9.3, let $R_Q(y, dx)$ be a regular conditional law of $X$ given $Y = y$ under $\mathbb{J}$. Then $QKR_Q = Q$, and at equality in the data-processing inequality also $PKR_Q = P$. Conversely, any one Markov kernel $R$ that recovers both $P$ and $Q$ forces that equality. For a dominated experiment $\{P_\theta\}$ with one common dominating law $P_0$, simultaneous finite equalities $D_{\mathrm{KL}}(P_\theta K\|P_0K) = D_{\mathrm{KL}}(P_\theta\|P_0)$ for every $\theta$ make the single kernel $R_{P_0}$ recover every $P_\theta$. [ESTABLISHED]

*Proof.* Disintegration gives $QKR_Q = Q$. At equality, $(PKR_Q)(A) = \int \bar{r}(y)R_Q(y, A)(QK)(dy) = \iint \mathbf{1}_A(x)r(x)\mathbb{J}(dx, dy) = P(A)$. The converse follows by data processing through $K$ and then through $R$, and the family statement is the pairwise statement with $Q = P_0$ for each $\theta$. This is the commutative case of the channel-sufficiency theorem of Petz [67], whose reference-dependent reverse map is here the Bayes kernel $R_Q$. □

**Corollary 9.5 (Infinite equality carries no recovery conclusion).** The finiteness hypothesis of Theorem 9.3 cannot be removed from its equality characterization: $D_{\mathrm{KL}}(PK\|QK) = D_{\mathrm{KL}}(P\|Q) = +\infty$ does not imply that one reverse kernel recovers $P$ and $Q$. [ESTABLISHED]

*Proof.* Let $\mathsf{X} = \{a, b, c\}$ and $\mathsf{Y} = \{u, v\}$, let the deterministic channel send $a$ to $u$ and $b, c$ to $v$, and set $P = \frac{1}{2}\delta_a + \frac{1}{2}\delta_b$ and $Q = \frac{1}{2}\delta_b + \frac{1}{2}\delta_c$. Both divergences are $+\infty$, because $P(a) > 0 = Q(a)$ and

$(PK)(u) > 0 = (QK)(u)$. A reverse kernel $R$ recovering $Q$ has $R(v,\cdot) = Q$, since $QK = \delta_v$; then $(PK)R = \frac{1}{2}R(u,\cdot) + \frac{1}{2}Q$ assigns at least $1/4$ to $c$, while $P(c) = 0$. □

## 9.3 The common-channel VFE identity and the pointwise parent datum

Fix the joint $\mathbb{P}(do, dy)$ on O × Y of Chapter 4 with $(\theta, V, X)$ suppressed, an admitted observation $o$ with evidence density representative $p(o)$, the selected posterior-version kernel $o \mapsto \mathbf{\Pi}_o$, a recognition law $\mathbb{Q}_o$ on Y, and a coarse channel $C : \mathsf{Y} \rightsquigarrow \mathsf{Z}$ on the latent coordinates, which reads neither $\mathbb{Q}_o$ nor the observation. The common pushforward of the three laws is

$$\mathbb{P}^c(do, dz) = \int_{\mathsf{Y}} C(y, dz)\mathbb{P}(do, dy), \qquad \mathbf{\Pi}^c_o = \mathbf{\Pi}_o C, \qquad \mathbb{Q}^c_o = \mathbb{Q}_o C, \tag{9.2}$$

and free energies are the extended values $\mathcal{F}_{\mathbb{P}}(\mathbb{Q}_o) = -\log p(o) + D_{\mathrm{KL}}(\mathbb{Q}_o \| \mathbf{\Pi}_o)$ of Definition 5.6, defined for every law on the latent space.

**Theorem 9.6 (Exact coarse VFE and discarded conditional information).** In the setting of (9.2), the following hold. (i) The law $\mathbb{P}^c$ is normalized and has the observation marginal of $\mathbb{P}$, so the same evidence representative $p(o)$ serves both scales. The kernel $o \mapsto \mathbf{\Pi}_o C$ is a selected posterior version of $\mathbb{P}^c$, including at the exceptional observations of the fine version. If $\mathbb{Q}_o \ll \mathbf{\Pi}_o$, then $\mathbb{Q}^c_o \ll \mathbf{\Pi}^c_o$. (ii) With the bridge lifts $\widehat{\mathbb{Q}}_o(dy, dz) = \mathbb{Q}_o(dy)C(y, dz)$ and $\widehat{\mathbf{\Pi}}_o(dy, dz) = \mathbf{\Pi}_o(dy)C(y, dz)$ and their disintegrations over $z$, the free energies satisfy, in $(-\infty, +\infty]$ and with no absolute-continuity hypothesis,

$$\mathcal{F}_{\mathbb{P}}(\mathbb{Q}_o) = \mathcal{F}_{\mathbb{P}^c}(\mathbb{Q}^c_o) + \Delta(o), \qquad \Delta(o) = \int_{\mathsf{Z}} D_{\mathrm{KL}}\big(\widehat{\mathbb{Q}}_o(dy \mid z) \| \widehat{\mathbf{\Pi}}_o(dy \mid z)\big)\mathbb{Q}^c_o(dz) \in [0, +\infty], \tag{9.3}$$

so the fine free energy is at least the coarse one. (iii) $\Delta(o) = 0$ exactly when the two discarded conditional laws in (9.3) agree $\mathbb{Q}^c_o$-almost surely. (iv) If $D_{\mathrm{KL}}(\mathbb{Q}_o \| \mathbf{\Pi}_o) < \infty$, then $\mathcal{F}_{\mathbb{P}}(\mathbb{Q}_o) - \mathcal{F}_{\mathbb{P}^c}(\mathbb{Q}^c_o) = \Delta(o)$ is an ordinary difference of real numbers, and $\Delta(o) = 0$ exactly when one Markov kernel $R : \mathsf{Z} \rightsquigarrow \mathsf{Y}$ recovers both laws, $\mathbb{Q}^c_o R = \mathbb{Q}_o$ and $\mathbf{\Pi}^c_o R = \mathbf{\Pi}_o$. This recovery is pairwise; one kernel recovering every member of a family of laws requires the hypothesis for each member simultaneously. [ESTABLISHED]

The proof is in appendix B.3.

**Corollary 9.7 (Evidence-preserving coarse channel).** In the setting of Theorem 9.6, $\mathbb{P}^c(do \times \mathsf{Z}) = \mathbb{P}(do \times \mathsf{Y})$, $p^c(o) = p(o)$, and $\mathbf{\Pi}^c_o = \mathbf{\Pi}_o C$ for every declared $o$, and the extended ELBO of Definition 5.6 satisfies, in $[-\infty, \log p(o)]$,

$$\mathcal{L}^{\mathrm{ext}}_{\mathbb{P}^c}(\mathbb{Q}_o C) = \log p(o) - D_{\mathrm{KL}}(\mathbb{Q}_o C \| \mathbf{\Pi}_o C) \geq \log p(o) - D_{\mathrm{KL}}(\mathbb{Q}_o \| \mathbf{\Pi}_o) = \mathcal{L}^{\mathrm{ext}}_{\mathbb{P}}(\mathbb{Q}_o). \tag{9.4}$$

When the fine ELBO is finite, the difference of the two ELBOs is $\Delta(o)$: the rise is the discarded conditional information at unchanged evidence, neither model improvement nor evidence ascent. For two separately declared models with evidences $p_1(o), p_2(o)$ and finite gaps, the difference of their ELBOs is $\log p_1(o) - \log p_2(o)$ minus the difference of their gaps, so it compares variational tightness only when the evidences agree, as they do under one evidence-preserving common pushforward. [ESTABLISHED]

*Proof.* The first three identities are part (i) of the theorem. The inequality is data processing in $[0, +\infty]$ (Lemma 3.12), and the difference is part (ii) written for $\mathcal{L}^{\mathrm{ext}} = -\mathcal{F}$. The statement on separately declared models is the gap identity of Corollary 5.7 for each model. □

**Remark 9.8 (Support boundaries of the discarded term).** If $\mathbb{Q}_o \not\ll \mathbf{\Pi}_o$ while $D_{\mathrm{KL}}(\mathbb{Q}_o^c \| \mathbf{\Pi}_o^c) < \infty$, then the fine free energy is $+\infty$, the coarse one is finite, and $\Delta(o) = +\infty$. If $\mathbb{Q}_o^c \not\ll \mathbf{\Pi}_o^c$, both free energies are $+\infty$; the conditional law $\widehat{\mathbf{\Pi}}_o(dy \mid z)$ is then undetermined on a $\mathbb{Q}_o^c$-positive set of targets, so $\Delta(o)$ depends there on an arbitrary version and the difference of the two infinite free energies is not a numerical loss. [ESTABLISHED]

*Proof.* In the first case (9.3) reads $+\infty = \text{finite} + \Delta(o)$. In the second, a $\mathbf{\Pi}_o^c$-null set $B$ with $\mathbb{Q}_o^c(B) > 0$ gives $\widehat{\mathbb{Q}}_o(\mathsf{Y} \times B) > 0 = \widehat{\mathbf{\Pi}}_o(\mathsf{Y} \times B)$, so both divergences are infinite, and a disintegration of $\widehat{\mathbf{\Pi}}_o$ over $z$ is determined only $\mathbf{\Pi}_o^c$-almost surely. □

**Proposition 9.9 (Stochastic loss and deterministic minimization).** Let Y and Z be finite and write $\mathcal{F}^c(q') = -\log p(o) + D_{\mathrm{KL}}(q' \| \mathbf{\Pi}_o C)$ for a law $q'$ on Z. (i) For a deterministic statistic $c : \mathsf{Y} \to \mathsf{Z}$ and $q' \ll c_{\#}\mathbf{\Pi}_o$, the fine law $\mathbb{Q}_*(y) = q'(c(y))\mathbf{\Pi}_o(y \mid c(y))$, defined to be zero where $q'(c(y)) = 0$ without evaluating a posterior fiber at a null target, has image $q'$ and satisfies $\mathcal{F}_{\mathbb{P}}(\mathbb{Q}_*) = \mathcal{F}^c(q') = \min\{\mathcal{F}_{\mathbb{P}}(\mathbb{Q}) : \mathbb{Q}K_c = q'\}$. (ii) For a stochastic channel the statement $\mathcal{F}^c(q') = \inf_{\mathbb{Q}:\mathbb{Q}C=q'} \mathcal{F}_{\mathbb{P}}(\mathbb{Q})$ can fail: with a one-point observation alphabet, $\mathbf{\Pi}_o = (1/2, 1/2)$, $C = \left(\begin{smallmatrix} 3/4 & 1/4 \\ 1/4 & 3/4 \end{smallmatrix}\right)$, $\mathbb{Q} = (1, 0)$, and $q' = \mathbb{Q}C = (3/4, 1/4)$, the law $\mathbb{Q}$ is the only fine law with image $q'$, its free energy is $\log 2$, the coarse free energy is $\frac{3}{4}\log\frac{3}{2} + \frac{1}{4}\log\frac{1}{2}$, and the gap is $2\log 2 - \frac{3}{4}\log 3 > 0$. [ESTABLISHED]

*Proof.* (i) Summing $\mathbb{Q}_*$ over a fiber $c^{-1}(z)$ with $q'(z) > 0$ gives $q'(z)$, because the posterior fiber $\mathbf{\Pi}_o(\cdot \mid z)$ is a law on $c^{-1}(z)$; so $\mathbb{Q}_*$ is a probability law with image $q'$. By (9.3), every $\mathbb{Q}$ with image $q'$ has $\mathcal{F}_{\mathbb{P}}(\mathbb{Q}) = \mathcal{F}^c(q') + \Delta \geq \mathcal{F}^c(q')$, and $\Delta = 0$ for $\mathbb{Q}_*$ because its occupied reverse fibers are the posterior fibers. (ii) The matrix $C$ has determinant $1/2$, so $\mathbb{Q}$ is the unique probability vector with image $q'$. The evidence is one, $\mathbf{\Pi}_o C = (1/2, 1/2)$, and direct evaluation gives the two free energies, whose difference $2\log 2 - \frac{3}{4}\log 3$ is positive because $2^8 > 3^3$. □

**Corollary 9.10 (Full pointwise probabilistic datum for a candidate parent).** Fix a nonempty finite block $I \subseteq V$ of agents with $c_* \in \bigcap_{i \in I} \mathcal{C}_i$, a parent label $A$, and structural data $X$ with $X_A = \chi_A(X)$ kept outside the random channel. Let the parent state space $\mathsf{Z}_A = \mathsf{B}_A \times \mathsf{M}_A \times \mathbf{\Xi}_A \times \mathsf{H}_A$ have nonempty standard Borel factors for parent belief samples, parent model presentations, the inputs read by the parent evaluator, and retained marks such as the rooted-holonomy and dressed-boundary data of Section 10.2. Supply the joint $\mathbb{P}_I(do, dy \mid X)$ of the block with its selected

posterior-version kernel $o \mapsto \mathbf{\Pi}_{I,o,X}$, a recognition law $\mathbb{Q}_{I,o,X} \ll \mathbf{\Pi}_{I,o,X}$ at an admitted $o$, and one parent coarse kernel $C_A : \mathsf{Y}_I \rightsquigarrow \mathsf{Z}_A$, a declared input like $I$, not a canonical selection. Define

$$\mathbb{P}_A(do, dz \mid X) = \int_{\mathsf{Y}_I} C_A(y, dz)\mathbb{P}_I(do, dy \mid X), \tag{9.5}$$

$$\mathbf{\Pi}_{A,o,X} = \mathbf{\Pi}_{I,o,X}C_A, \qquad \mathbb{Q}_{A,o,X} = \mathbb{Q}_{I,o,X}C_A. \tag{9.6}$$

Then the following hold. (i) All three parent laws are normalized, the parent observation marginal and evidence equal the fine ones, $o \mapsto \mathbf{\Pi}_{A,o,X}$ is the selected parent posterior version, $\mathbb{Q}_{A,o,X} \ll \mathbf{\Pi}_{A,o,X}$, and in $[0, +\infty]$ $D_{\mathrm{KL}}(\mathbb{Q}_{I,o,X} \| \mathbf{\Pi}_{I,o,X}) = D_{\mathrm{KL}}(\mathbb{Q}_{A,o,X} \| \mathbf{\Pi}_{A,o,X}) + \Delta_A(o, X)$, where $\Delta_A(o, X) \geq 0$ is the discarded term $\Delta(o)$ of (9.3) for $\mathbb{P} = \mathbb{P}_I$ and $C = C_A$. Adding $-\log p_X(o)$ gives $\mathcal{F}_I = \mathcal{F}_A + \Delta_A$ in the extended reals, where a finite free energy may be negative. The zero-defect and recovery statements of Theorem 9.6 hold with $C = C_A$; a bare equality $+\infty = +\infty$ supplies neither. (ii) The parent model coordinate is generative only through a jointly measurable normalized evaluator $\mathrm{ev}_A : \mathsf{M}_A \to \mathrm{Kern}(\mathbf{\Xi}_A, \mathsf{B}_A \times \mathsf{O} \times \mathsf{H}_A)$, $m_A \mapsto K^{X_A}_{A,m_A}$. A selected standard-Borel disintegration of $\mathbb{P}_A$ over $(m_A, \xi_A)$ supplies one such evaluator; an evaluator declared independently of $\mathbb{P}_A$ is compatible with it only under the separate hypothesis that $\mathbb{P}_A(db_A, do, dh_A \mid \xi_A, m_A, X) = K^{X_A}_{A,m_A}(\xi_A; db_A, do, dh_A)$ almost surely for the parent $(m_A, \xi_A)$ marginal. Disintegration neither validates an arbitrary declared evaluator nor selects its null extension or makes evaluation injective. (iii) With $\mathbb{P}^Z_A(D \mid X) = \mathbb{P}_A(\mathsf{O} \times D \mid X)$, the parent recognition, prior, and posterior marginals in channel $x \in \{b, m\}$ are the coordinate pushforwards $q^x_A = (\mathrm{pr}_x)_\# \mathbb{Q}_{A,o,X}$, $p^x_A = (\mathrm{pr}_x)_\# \mathbb{P}^Z_A(\cdot \mid X)$, and $\mathbf{\Pi}^x_{A,o,X} = (\mathrm{pr}_x)_\# \mathbf{\Pi}_{A,o,X}$. If both laws of a marginal pair are non-Dirac, or if $\mathbf{\Xi}_A \times \mathsf{H}_A$ has at least two points, the pair does not determine the full law on $\mathsf{Z}_A$. If $\mathbf{\Xi}_A \times \mathsf{H}_A$ is a point and one marginal is a point mass, the coupling class of the pair is a single law. [ESTABLISHED]

The proof is in appendix B.3.

## 9.4 Pushforward, precomposition, and Galerkin restriction

Three operations share the name coarse-graining. A Markov pushforward (Definition 9.1) moves laws and obeys data processing; precomposition identifies variables inside an energy and defines a new model; Galerkin restriction compresses an operator and needs further data before it propagates signals. A proof about one of them does not transfer to another.

**Definition 9.11 (Energy precomposition).** For a measurable injection $\iota : \mathsf{Z} \to \mathsf{X}$ and an energy $E$ on $\mathsf{X}$, the precomposed model is

$$\bar{E}(z) = E(\iota z), \qquad \bar{P}(dz) = \frac{e^{-\bar{E}(z)}}{\bar{Z}}\bar{\nu}(dz), \tag{9.7}$$

defined once a reference $\bar{\nu}$ on $\mathsf{Z}$ and a finite normalizer $\bar{Z}$ are declared; the reference, its Radon-Nikodym or Jacobian rule, and the normalizer are part of the construction.

**Proposition 9.12 (Hard identification is conditioning).** (i) On a finite or countable X, with counting references on Z and X and $E = -\log p$ for the mass function of a law $P$, the precomposed law is $\bar{P}(z) = p(\iota z)/P(\iota(\mathrm{Z}))$: the law of $X$ conditioned on $X \in \iota(\mathrm{Z})$, transported by $\iota^{-1}$. (ii) Conditioning can increase relative entropy. On $\mathrm{X} = \{1,2,3\}$ with $\iota(\mathrm{Z}) = \{1,2\}$, the laws $P = (1,1,98)/100$ and $Q = (1,2,97)/100$ have $D_{\mathrm{KL}}(P\|Q) \approx 0.0031$, while their conditioned laws $(1/2,1/2)$ and $(1/3,2/3)$ have relative entropy $\approx 0.0589$. Hence no Markov kernel sends $P$ and $Q$ to their precomposed laws, and data processing, evidence preservation, and (9.4) do not transfer to precomposition. (iii) When $\iota(\mathrm{Z})$ is null for the fine reference, as for block-constant configurations of a continuous fiber, the fine law defines no conditional law on it, and different references give different precomposed laws: for $\bar{E}(z) = z^2$ on $\mathrm{Z} = \mathbb{R}$, Lebesgue reference gives $\mathcal{N}(0,1/2)$ and the reference $e^{-z^2}dz$ gives $\mathcal{N}(0,1/4)$. Exact marginalization is a Markov pushforward, and restricting a recognition family keeps the joint and optimizes the same ELBO over a subset, which is neither pushforward nor precomposition. [ESTABLISHED]

*Proof.* (i) Substitution in (9.7) gives $\bar{P}(z) \propto p(\iota z)$, with normalizer $P(\iota(\mathrm{Z}))$. (ii) The values are direct evaluation; every Markov kernel contracts relative entropy (Lemma 3.12), and the conditioned pair has larger relative entropy. (iii) Conditioning on a null event is not defined by the law, and the densities $e^{-z^2}$ and $e^{-2z^2}$ are proportional to the stated Gaussians. The last sentence restates the definitions. □

### 9.4.1 Galerkin operators and message passing

**Definition 9.13 (Galerkin data).** Let $S : V_c \to V_f$ be an injective linear map of finite-dimensional real spaces, let $M \succ 0$ be a fine reference form, and let $L$ be a fine bilinear form. Put $M_c = S^\top M S$, $S^+ = M_c^{-1} S^\top M$, $H = M^{-1}L$, and $H_c = S^+ H S = M_c^{-1} S^\top L S$. The form $S^\top L S$ is the Galerkin coarse form of aggregation-based algebraic multigrid [65, Eqs. (2.1) and (2.2)], [58, Sec. 2], here for an arbitrary injective prolongation as in smoothed aggregation [79]; the map $H_c$ is the coarse propagator.

**Proposition 9.14 (Exact message-passing residual).** The data of Definition 9.13 satisfy $S^+ S = I$ and $S^+ H - H_c S^+ = S^+ H(I - SS^+)$. The idempotent $SS^+$ is the $M$-orthogonal projector onto $\operatorname{range} S$, with kernel $\ker S^+ = (\operatorname{range} S)^{\perp_M}$. Hence exact coarse message passing, $S^+ H = H_c S^+$ on all fine signals, holds exactly when

$$H\big((\operatorname{range} S)^{\perp_M}\big) \subseteq (\operatorname{range} S)^{\perp_M}. \tag{9.8}$$

The operator $H$ is self-adjoint in the $M$-inner product exactly when $L$ is symmetric, and then (9.8) is equivalent to $H$-invariance of $\operatorname{range} S$. Without symmetry the two invariance conditions are independent: with $M = I$, $S = e_1$, and $L = \left(\begin{smallmatrix} 2 & 1 & 0 \\ 0 & 3 & 0 \\ 0 & 0 & 4 \end{smallmatrix}\right)$, the subspace $\operatorname{range} S$ is invariant but the residual $S^+ H(I - SS^+) = (0,1,0)$ is nonzero, while $L^\top$ leaves the complement invariant and gives zero residual although $\operatorname{range} S$ is not invariant. [ESTABLISHED]

*Proof.* $S^+ S = M_c^{-1} S^\top M S = I$, and substituting $H_c = S^+ H S$ gives the residual identity. The map $SS^+$ is idempotent with range $\operatorname{range} S$ and kernel $\ker S^+$; since $M_c$ is invertible, $S^+ y = 0$ exactly

when $S^\top My = 0$, that is, when $y$ is $M$-orthogonal to $\operatorname{range} S$. The residual vanishes exactly when $S^+Hy = 0$ for every $y \in \operatorname{range}(I - SS^+) = \ker S^+$, which is (9.8). Next $\langle Hx, y\rangle_M = x^\top L^\top y$ and $\langle x, Hy\rangle_M = x^\top Ly$, and for an $M$-self-adjoint operator a subspace is invariant exactly when its $M$-orthogonal complement is. The two witnesses are direct computations. □

## 9.5 Gauge and permutation equivariance

**Proposition 9.15 (Equivariant Markov channels).** Let a group $\mathbb{G}$ act bimeasurably on X and Y, and let $K : \mathsf{X} \rightsquigarrow \mathsf{Y}$ be equivariant, $K(gx, B) = K(x, g^{-1}B)$. Then $(g_\# P)K = g_\#(PK)$ for every law $P$ on X, and composites of equivariant kernels are equivariant. [ESTABLISHED]

*Proof.* For measurable $B$ the law $(g_\# P)K$ gives $B$ the mass $\int K(gx, B)P(dx)$, which by equivariance equals $\int K(x, g^{-1}B)P(dx) = [g_\#(PK)](B)$. Substitution in the composition of Proposition 9.2 proves closure under composition. □

At $c_*$ a principal gauge transformation acts on belief and model coordinates through $\rho_b(g)$ and $\rho_m(g)$ of one $g \in G$ (Proposition 2.6), and each passive re-expression $g_i$ acts on both channels; a coarse channel on joint belief and model coordinates is gauge covariant when it is equivariant for these actions, and for agent relabelings the acting group is a permutation group.

**Proposition 9.16 (Equivariant transport of Bayes recovery kernels).** Let $\mathbb{G}$ act on X and Y by bimeasurable bijections, let $K$ be equivariant as in Proposition 9.15, let $Q$ be a law on X, let $R_Q$ be a regular conditional law of $X$ given $Y$ under $\mathbb{J}_Q(dx, dy) = Q(dx)K(x, dy)$, and put $R_Q^g(y, A) = R_Q(g^{-1}y, g^{-1}A)$ for $g \in \mathbb{G}$. (i) $R_Q^g$ is a regular conditional law of $X$ given $Y$ under $\mathbb{J}_{g_\# Q}$ for every $g$. (ii) If $PKR_Q = P$, then $(g_\# P)KR_Q^g = g_\# P$, and $D_{\mathrm{KL}}((g_\# P)K\|(g_\# Q)K) = D_{\mathrm{KL}}(PK\|QK)$ and $D_{\mathrm{KL}}(g_\# P\|g_\# Q) = D_{\mathrm{KL}}(P\|Q)$. (iii) If X and Y are finite and $\mathbb{G}$ acts by permutations, the rule $R_Q(y, \{x\}) = Q(\{x\})K(x, \{y\})/(QK)(\{y\})$ where $(QK)(\{y\}) > 0$ and $R_Q(y, \{x\}) = 1/|\mathsf{X}|$ where $(QK)(\{y\}) = 0$ selects a version for every $Q$, depends Borel measurably on $(Q, y)$, and satisfies $R_{g_\# Q}(gy, \{gx\}) = R_Q(y, \{x\})$ for all $g, Q, y, x$. (iv) If $\mathbb{G}$ is a topological group acting jointly measurably and the stabilizer $\mathbb{G}_Q = \{h : h_\# Q = Q\}$ is a compact metrizable group with Haar probability $dh$, then $\bar{R}_Q(y, A) = \int_{\mathbb{G}_Q} R_Q^h(y, A)dh$ is a version of $R_Q$ with $\bar{R}_Q(hy, hA) = \bar{R}_Q(y, A)$ for $h \in \mathbb{G}_Q$, and $R_{g_\# Q} := \bar{R}_Q^g$ is well defined on the orbit of $Q$ and satisfies $R_{k_\# Q'}(ky, kA) = R_{Q'}(y, A)$ for every $Q'$ in that orbit and every $k \in \mathbb{G}$. [ESTABLISHED]

The proof is in appendix B.3.

**Corollary 9.17 (Continuous posterior versions are equivariant).** In the setting of Proposition 9.16, let X and Y be Polish, let every $h \in \mathbb{G}_Q$ act on both by homeomorphisms, and let $S$ be the support of $QK$. If some version of $R_Q$ is weakly continuous on $S$, then $R(y, \cdot) = R_Q(y, \cdot)$ for $y \in S$ and $R(y, \cdot) = Q$ for $y \notin S$ is a version of $R_Q$ with $R(hy, hA) = R(y, A)$ for every $h \in \mathbb{G}_Q$, every $y \in \mathsf{Y}$, and every measurable $A$, whether or not the stabilizer is compact. [ESTABLISHED]

*Proof.* By Proposition 9.15, $h_\#(QK) = (h_\# Q)K = QK$ for $h \in \mathbb{G}_Q$, so $S$ is a closed $\mathbb{G}_Q$-invariant set, and $R$ is a Markov kernel that agrees with $R_Q$ on the full-measure set $S$. For $h \in \mathbb{G}_Q$ the

kernel $R_Q^h$ is a version of $R_Q$ by part (i) of the proposition, and $y \mapsto R_Q^h(y,\cdot) = h_{\#}R_Q(h^{-1}y,\cdot)$ is weakly continuous on $S$, because $h^{-1}$ is a homeomorphism of Y preserving $S$ and pushforward by the homeomorphism $h$ of X is weakly continuous. Two versions agree $QK$-almost everywhere, and the set where two weakly continuous maps from $S$ into the Hausdorff space of laws agree is closed in $S$; its complement in $S$ is an open $QK$-null subset of the support, hence empty. So $R_Q^h = R_Q$ on $S$, which is the equivariance identity at $y \in S$; for $y \notin S$ also $h^{-1}y \notin S$, and $R(h^{-1}y, h^{-1}A) = Q(h^{-1}A) = Q(A) = R(y,A)$. □

**Open Problem 9.18 (Equivariant Bayes versions beyond the finite and compact cases).** Proposition 9.16 settles equivariant Bayes versions for finite spaces and along orbits with compact stabilizer, and Corollary 9.17 settles them for every stabilizer when the posterior has a weakly continuous version on the support of $QK$. Two cases remain open in the class of a locally compact second countable group acting continuously on Polish spaces, an equivariant channel $K$, and laws $Q$ whose stabilizer is noncompact, so that no Haar probability averages versions, and whose posterior has no weakly continuous version on the support of $QK$: an everywhere-defined version equivariant under the stabilizer; and a version that is simultaneously equivariant and jointly measurable across a measurable family of such laws meeting several orbits, which also needs a measurable choice of orbit representatives. The problem is settled by an equivariant measurable selection theorem giving, for every member of that class, an everywhere-defined version equivariant under the stabilizer, jointly measurable across every measurable family of such laws meeting several orbits, or by a pair $(Q,K)$ of the class with no everywhere-defined version equivariant under the stabilizer.

**Open Problem 9.19 (A natural selector of partitions and membership laws).** A data-dependent hard partition selector $S(X)$ must be natural under agent relabeling: for every fine permutation $P$ there must exist a coarse relabeling $Q = Q(P,X)$ with $S(P\cdot X) = PS(X)Q^\top$, that is, the unlabeled partition is equivariant. The quantifier on $Q$ is existential because coarse labels are arbitrary; a universal one would force $S(X) = S(X)Q^\top$ for every $Q$ at $P = \mathrm{id}$, which no partition into two or more blocks satisfies. Gauge-carrying inputs must enter through invariant assignment scores, while represented messages transform equivariantly. These conditions type an admissible selector of a partition, a membership law, or a structural channel, and select none. The costs analyzed so far do not select one: an unpenalized sum of nonnegative within-block divergences is minimized by the singleton partition (Remark 10.20), the Gaussian recognition costs of the Gaussian supplement select trivial endpoints or need an external coefficient, and one common recovery kernel for a whole statistical experiment, (10.20), is available only under simultaneous family hypotheses. The problem is settled by a relabeling-natural, gauge-compatible functional of the laws, not monotone under refinement, with a nondegenerate optimizer, or by a theorem that no such selector exists in a stated class.

## 9.6 Scale diagrams

**Definition 9.20 (Scale diagram).** Let $\mathsf{S}$ be a category whose objects are resolutions $\ell$, and choose one target category $\mathscr{K}$: sets and maps for a deterministic state, standard Borel spaces and Markov kernels for a stochastic state, or a declared category of Banach spaces and bounded linear maps, or a declared product of such categories, for an operator or mixed state. A scale diagram is a functor $\mathfrak{X} : \mathsf{S} \to \mathscr{K}$, with objects $\mathfrak{X}_\ell$ and arrows $C_{\ell k} \in \mathrm{Hom}_{\mathscr{K}}(\mathfrak{X}_\ell, \mathfrak{X}_k)$ satisfying $C_{kr} \circ C_{\ell k} = C_{\ell r}$, where $\circ$ is composition in $\mathscr{K}$, componentwise in a product. For a stochastic component, kernels act on the right,

$$K_{\ell r} := K_{\ell k} K_{kr}, \qquad (\mu K_{\ell k}) K_{kr} = \mu K_{\ell r}, \tag{9.9}$$

so that $K_{\ell k} K_{kr}$ represents the abstract composite $K_{kr} \circ K_{\ell k}$. One symbol is never at once a map, a kernel, and an operator. An expression on a reference space needs declared isomorphisms, and transformations of reference measures stay explicit.

**Proposition 9.21 (When functoriality has content).** (i) On a chain of resolutions $0 < 1 < \cdots < L$, any family of adjacent arrows $C_{\ell,\ell+1}$ extends uniquely to a scale diagram, by $C_{\ell k} = C_{k-1,k} \circ \cdots \circ C_{\ell,\ell+1}$ and $C_{\ell\ell} = \mathrm{id}$; the composition law imposes no condition. (ii) On a poset of resolutions that is not a chain, functoriality requires any two routes between the same resolutions to agree. For a family of kernels generated by one blocking rule indexed by block factors, functoriality is the semigroup law $K_{b_1 b_2} = K_{b_1} K_{b_2}$. (iii) Deterministic arrows that keep the complete block configuration, or an additive block statistic such as block sums, satisfy both conditions. Majority with ties resolved by the first block member does not: on six agents it sends $(0, 0, 1, 0, 1, 0)$ to $1$ by pairs and then triples, and to $0$ by triples and then pairs. (iv) Mutually inverse Markov kernels between finite spaces are permutation matrices, so no stochastic isomorphism exists between finite levels with different numbers of configurations. [ESTABLISHED]

The proof is in appendix B.3.

## 9.7 Labeled families and retained label information

A finite mixture that keeps its label in the sample space is the elementary model of a random partition or random link pattern among agents at $c_*$; its Fisher tensor splits into a label part and a weighted branch part.

**Definition 9.22 (Finite labeled family).** Let $\Theta \subseteq \mathbb{R}^d$ be open, $1 \le d < \infty$, and let $\mathfrak{G}$ be a finite nonempty set of labels $\gamma$, for instance partitions of a block of agents or link patterns. Each label has a finite nonempty carrier $\mathsf{X}_\gamma$, and the sample space is the disjoint union of the $\{\gamma\} \times \mathsf{X}_\gamma$, so that labels stay distinct when branch laws coincide. A finite labeled family on $\Theta$ consists of effective log masses $\lambda_\gamma \in C^3(\Theta)$ and branch laws $p_{\gamma,\theta}$ on $\mathsf{X}_\gamma$, with

$$P_\theta(\gamma, x) = W_\gamma(\theta) p_{\gamma,\theta}(x), \qquad W_\gamma = \frac{e^{\lambda_\gamma}}{\sum_{\delta \in \mathfrak{G}} e^{\lambda_\delta}}, \tag{9.10}$$

on a fixed support stratum: every branch mass is either identically zero on $\Theta$, a structural zero that carries no score, or positive and $C^3$ in $\theta$. A common additive function of $\theta$ in all $\lambda_\gamma$ is a change of representative. When the family comes from a normalization model, the effective log masses absorb its label weights and branch partitions (Proposition 12.13). For a fixed-support finite family $\theta \mapsto R_\theta$ the score is $\mathbf{s}_R = \nabla_\theta \log R_\theta$ on positive atoms and the Fisher tensor is $g^F(R) = \sum_{z:R_\theta(z)>0} R_\theta(z)\mathbf{s}_R(z)\mathbf{s}_R(z)^\top$; write $\xi_\gamma = \nabla\lambda_\gamma$ and $\bar\xi = \sum_\gamma W_\gamma \xi_\gamma$. An observation model is a record kernel $K^{\mathrm{obs}}(o \mid \gamma, x)$ on a finite alphabet O, of the same fixed-support $C^3$ type and independent of recognition; it defines the complete observation law $P^{\mathrm{obs}}_\theta(\gamma, x, o) = P_\theta(\gamma, x)K^{\mathrm{obs}}(o \mid \gamma, x)$ and the evidence $p_\theta(o) = \sum_{\gamma,x} P^{\mathrm{obs}}_\theta(\gamma, x, o)$. The complete latent, complete observation, and observed-data experiments are $\theta \mapsto P_\theta$, $\theta \mapsto P^{\mathrm{obs}}_\theta$, and $\theta \mapsto p_\theta$. An extensive labeled family is a sequence $P_N = (W_{\gamma,N}p_{\gamma,N})$ of finite labeled families indexed by an extensive size $N \geq 1$, such as the number of agents, with carriers that may depend on $N$. On a compact $\mathsf{K} \subset \Theta$ it has extensive label growth with bulk log-mass densities $\mathsf{s}_\gamma$ if, on an open neighborhood of K,

$$\|\lambda_{\gamma,N}\|_{C^3(\mathsf{K})} \leq \mathrm{C}_{\mathsf{K}} N, \qquad \lambda_{\gamma,N}/N \longrightarrow \mathsf{s}_\gamma \quad \text{in } C^3(\mathsf{K}), \tag{9.11}$$

where $\|f\|_{C^k(\mathsf{K})}$ is the largest supremum over K of the ordinary coordinate derivatives of $f$ through order $k$, including order zero, with fixed norms on vector and matrix entries.

**Proposition 9.23 (Complete labeled Fisher decomposition).** For a finite labeled family, the branch score $\mathbf{s}_{p_\gamma} = \nabla \log p_\gamma$ on the positive atoms of $\mathsf{X}_\gamma$ satisfies $\sum_x p_\gamma(x)\mathbf{s}_{p_\gamma}(x) = 0$, and the latent score, the Fisher tensor, and, for every multi-index $\alpha$ with $|\alpha| \leq 2$, the derivatives of its weighted branch term satisfy

$$\mathbf{s}_P(\gamma, x) = \xi_\gamma - \bar\xi + \mathbf{s}_{p_\gamma}(x), \tag{9.12}$$

$$g^F(P) = \mathrm{Cov}_W(\xi_\gamma) + \sum_\gamma W_\gamma g^F(p_\gamma), \tag{9.13}$$

$$\partial^\alpha\Big(\sum_\gamma W_\gamma g^F(p_\gamma)\Big) = \sum_\gamma \sum_{\beta\leq\alpha} \binom{\alpha}{\beta} (\partial^\beta W_\gamma)(\partial^{\alpha-\beta} g^F(p_\gamma)). \tag{9.14}$$

The first term of $g^F(P)$ is the Fisher tensor of the label marginal $W$. A common additive function in every $\lambda_\gamma$ leaves these identities and their derivatives unchanged, and $g^F(P) \succeq 0$ with no rank assumption. [ESTABLISHED]

The proof is in appendix B.3.

**Proposition 9.24 (Selected posterior and its complete score).** For a finite labeled family with an observation model and a record $o$, let $e_\gamma(o) = \sum_x p_\gamma(x)K^{\mathrm{obs}}(o \mid \gamma, x)$ be the branch evidence, so that $p_\theta(o) = \sum_\gamma W_\gamma e_\gamma(o)$, and where $p_\theta(o) > 0$ define the selected posterior $\mathbf{\Pi}_o(\gamma, x) = W_\gamma p_\gamma(x)K^{\mathrm{obs}}(o \mid \gamma, x)/p_\theta(o)$ and its label weights $W^o_\gamma = W_\gamma e_\gamma(o)/p_\theta(o)$. A label with $e_\gamma(o) = 0$ stays on the carrier with zero posterior mass and has no conditional branch; zero evidence leaves the posterior undefined. On occupied labels $W^o$ is the softmax of $\lambda^o_\gamma = \lambda_\gamma + \log e_\gamma(o)$. On the positive posterior support the complete score and, within an occupied branch $p^o_\gamma(x) = p_\gamma(x)K^{\mathrm{obs}}(o \mid \gamma, x)/e_\gamma(o)$, the

branch score are

$$\mathbf{s}_{\mathbf{\Pi}_o}(\gamma, x) = \mathbf{s}_P(\gamma, x) + \nabla \log K^{\mathrm{obs}}(o \mid \gamma, x) - \nabla \log p_\theta(o), \tag{9.15}$$

$$\nabla \log p^o_\gamma(x) = \mathbf{s}_{p_\gamma}(x) + \nabla \log K^{\mathrm{obs}}(o \mid \gamma, x) - \nabla \log e_\gamma(o), \tag{9.16}$$

and the decomposition (9.13) holds for the posterior family $\theta \mapsto \mathbf{\Pi}_{o,\theta}$, with the weights $W^o$ and the branches $p^o_\gamma$. [ESTABLISHED]

*Proof.* Summing the posterior over all labeled atoms gives one, and the label weights are proportional to $e^{\lambda_\gamma} e_\gamma(o)$. Differentiating the logarithm of the positive posterior mass gives (9.15), and the same computation within one branch gives (9.16). The fixed zero patterns of branch laws and record kernel keep the occupied labels and atoms fixed on the stratum, so no division by zero occurs, and the proof of Proposition 9.23 applies to the posterior family. □

**Proposition 9.25 (Observed-data information and missing information).** For a finite labeled family with an observation model, let $\mathsf{O}^+ = \{o : p_\theta(o) > 0\}$, fixed on the stratum, let $\mathbf{s}^c = \nabla \log P^{\mathrm{obs}}$ be the complete score on the positive support of the complete observation law, and let $g^F(p) = \sum_{o \in \mathsf{O}^+} p_\theta(o) \nabla \log p_\theta(o) \nabla \log p_\theta(o)^\top$ be the observed-data Fisher tensor. Then

$$g^F(P^{\mathrm{obs}}) = g^F(P) + \sum_\gamma W_\gamma \sum_x p_\gamma(x) g^F\big(K^{\mathrm{obs}}(\cdot \mid \gamma, x)\big), \tag{9.17}$$

$$\nabla \log p_\theta(o) = \mathbb{E}_{\mathbf{\Pi}_o}\big[\mathbf{s}^c(\gamma, X, o)\big], \qquad g^F(p) = g^F(P^{\mathrm{obs}}) - \sum_{o \in \mathsf{O}^+} p_\theta(o) \operatorname{Cov}_{\mathbf{\Pi}_o}(\mathbf{s}^c). \tag{9.18}$$

In particular $g^F(p) \preceq g^F(P^{\mathrm{obs}})$, and positive definiteness of $g^F(P)$ or of $g^F(P^{\mathrm{obs}})$ does not imply that of $g^F(p)$: for a one-point alphabet $p_\theta(o) = 1$ identically and $g^F(p) = 0$ for all branch data. [ESTABLISHED]

The proof is in appendix B.3.

Fisher contraction requires a channel that does not depend on the parameter; the constant-input Bernoulli witness of Corollary 7.42 shows that a parameter-dependent channel can create information that its input does not carry.

**Corollary 9.26 (Stochastic score projection and retained-label loss).** A label-preserving channel is given by row-stochastic matrices $C_\gamma(x, y)$ from $\mathsf{X}_\gamma$ to finite nonempty targets $\mathsf{Y}_\gamma$, independent of $\theta$ and of recognition, and lifted as zero between different labels, so that $P^C(\gamma, y) = \sum_x P(\gamma, x) C_\gamma(x, y)$. The source law may be the latent family or a positive selected posterior; on the complete observation law the channel acts on $x$ and keeps $o$. For such a channel, on positive targets, the score projection $\mathbf{s}_{P^C}(\gamma, y) = \mathbb{E}_{P \otimes C}[\mathbf{s}_P(\gamma, X) \mid \gamma, y]$ and the Fisher loss $g^F(P) - g^F(P^C) = \mathbb{E}_{P \otimes C}[(\mathbf{s}_P - \mathbf{s}_{P^C})(\mathbf{s}_P - \mathbf{s}_{P^C})^\top] \succeq 0$ hold, and for $P = (W_\gamma p_\gamma)$

$$g^F(P) - g^F(P^C) = \sum_{\gamma : W_\gamma > 0} W_\gamma \big[g^F(p_\gamma) - g^F(p_\gamma C_\gamma)\big], \qquad 0 \preceq g^F(P) - g^F(P^C) \preceq \sum_\gamma W_\gamma g^F(p_\gamma), \tag{9.19}$$

while the label marginal and all its derivatives are unchanged. [ESTABLISHED]

*Proof.* The score projection and the Fisher loss are Theorem 7.8 for the finite fixed-support family, or directly: differentiating $P^C$ at a positive target and dividing by its mass gives the conditional expectation, and the residual $\mathbf{s}_P - \mathbf{s}_{P^C}$ has zero conditional mean given the target, so the cross terms vanish. Row sums give $\sum_y P^C(\gamma, y) = W_\gamma$ as functions of $\theta$. Subtracting the two decompositions (9.13) gives the first part of (9.19), and $g^F(p_\gamma C_\gamma) \succeq 0$ gives the upper bound. A zero-mass target receives no incidence from a positive atom and is a structural zero. □

**Proposition 9.27 (Finite-word score and VFE telescopes).** Let $C_1, \dots, C_m$ be parameter-independent Markov kernels between finite spaces, let $P_0$ be a fixed-support finite family, and put $P_j = P_0 C_1 \cdots C_j$ on its fixed positive support, with score $\mathbf{s}_j$ and Fisher tensor $g_j = g^F(P_j)$. Then $g_0 - g_m = \sum_{j=1}^m D_j$ with $D_j = g_{j-1} - g_j \succeq 0$, and on the path law $P_0(x_0)C_1(x_0, x_1) \cdots C_m(x_{m-1}, x_m)$ the increments $\varepsilon_j^{\mathrm{sc}} = \mathbf{s}_{j-1}(X_{j-1}) - \mathbf{s}_j(X_j)$ are orthogonal with $\mathbb{E}[\varepsilon_j^{\mathrm{sc}}(\varepsilon_j^{\mathrm{sc}})^\top] = D_j$, so $\mathbb{E}[(\mathbf{s}_0(X_0) - \mathbf{s}_m(X_m))(\mathbf{s}_0(X_0) - \mathbf{s}_m(X_m))^\top] = \sum_j D_j$. For a record with positive evidence, a selected posterior $\mathbf{\Pi}_0$, and a recognition law $\mathbb{Q}_0 \ll \mathbf{\Pi}_0$, push both laws and the joint by the same word and set $\mathcal{F}_j = -\log p(o) + D_{\mathrm{KL}}(\mathbb{Q}_j \| \mathbf{\Pi}_j)$. Then $\mathcal{F}_0 - \mathcal{F}_m = \sum_{j=1}^m K_j^{\mathrm{KL}}$, where $K_j^{\mathrm{KL}} \geq 0$ is the discarded conditional relative entropy of step $j$. If the fine and coarse families have extensive sizes $N_0$ and $N_m$, the per-agent bookkeeping reads $\frac{N_0}{N_m} \cdot \frac{\mathcal{F}_0}{N_0} = \frac{\mathcal{F}_m}{N_m} + \frac{1}{N_m} \sum_{j=1}^m K_j^{\mathrm{KL}}$. [ESTABLISHED]

The proof is in appendix B.3.

## 9.8 Retained-label Fisher asymptotics

When log masses grow like the number of agents $N$, the softmax calculus of Proposition 9.23 has three regimes with different hypotheses and coordinates: a strict winner, a tie with converging weights, and a window of width $1/N$ at a tie. Fix an extensive labeled family with weights $W_{\gamma,N}$, log masses $\lambda_{\gamma,N}$, gradients $\xi_{\gamma,N} = \nabla \lambda_{\gamma,N}$, Fisher density $g_N = N^{-1} g^F(P_N)$, and branch densities $\mathcal{I}_{\gamma,N} = N^{-1} g^F(p_{\gamma,N})$ with limits $\mathcal{I}_\gamma$ when they converge; $\mathrm{C}_\mathsf{K}$ depends on K and the family and may change from line to line, and limits are taken at a fixed label set.

**Theorem 9.28 (Strict-phase convergence with both metric derivatives).** Assume extensive label growth (9.11) on a neighborhood of a compact K, and $\mathcal{I}_{\gamma,N} \to \mathcal{I}_\gamma$ in $C^2(\mathsf{K})$ for every label. Suppose a label $\gamma_*$ satisfies $\mathsf{s}_{\gamma_*} - \mathsf{s}_\gamma \geq \delta_\mathsf{K} > 0$ on K for every $\gamma \neq \gamma_*$. Then $\|g_N - \mathcal{I}_{\gamma_*}\|_{C^2(\mathsf{K})} \to 0$. More precisely, for $j = 0, 1, 2$ and all large $N$,

$$\left\| D^j \left[ N^{-1} \operatorname{Cov}_{W_N}(\xi_{\gamma,N}) \right] \right\|_\mathsf{K} \leq \mathrm{C}_\mathsf{K} N^{j+1} e^{-\delta_\mathsf{K} N/2}, \tag{9.20}$$

and the weighted branch error $\sum_\gamma W_{\gamma,N} \mathcal{I}_{\gamma,N} - \mathcal{I}_{\gamma_*,N}$ has $C^2(\mathsf{K})$ norm at most $\mathrm{C}_\mathsf{K} N^2 e^{-\delta_\mathsf{K} N/2}$. Only three derivatives of the effective log masses are used. If moreover $\mathcal{I}_{\gamma_*} \succeq \mathrm{c}_{\min} \mathrm{Id}_d$ on K for a constant $\mathrm{c}_{\min} > 0$, then $g_N(\theta) \succeq \frac{1}{2} \mathrm{c}_{\min} \mathrm{Id}_d$ for all $\theta \in \mathsf{K}$ and all large $N$. [ESTABLISHED]

The proof is in appendix B.3.

**Proposition 9.29 (Squared-size Fisher tensor at a finite-weight tie).** Fix $\theta_0$ and assume (9.11) on a neighborhood of $\theta_0$. Let $S$ be the nonempty set of labels maximizing $\mathsf{s}_\gamma(\theta_0)$, suppose the weights satisfy $W_{\gamma,N}(\theta_0) \to \bar{\mathsf{p}}_\gamma$, and assume $g^F(p_{\gamma,N})(\theta_0) = O(N)$ for every label. Write $\mathsf{v}_\gamma = \nabla \mathsf{s}_\gamma(\theta_0)$ and $\bar{\mathsf{v}} = \sum_\gamma \bar{\mathsf{p}}_\gamma \mathsf{v}_\gamma$. Then $\bar{\mathsf{p}}$ is supported on $S$ and $N^{-2} g^F(P_N)(\theta_0) \to \mathrm{Cov}_{\bar{\mathsf{p}}}(\mathsf{v}_\gamma) = \sum_\gamma \bar{\mathsf{p}}_\gamma (\mathsf{v}_\gamma - \bar{\mathsf{v}})(\mathsf{v}_\gamma - \bar{\mathsf{v}})^\top$. Its rank is the dimension of the span of the differences $\mathsf{v}_\gamma - \mathsf{v}_\eta$ over pairs with positive limiting weights, at most $\min(d, |\operatorname{supp} \bar{\mathsf{p}}| - 1)$. A direction with positive variance has Fisher information asymptotic to a positive multiple of $N^2$; zero variance gives only $o(N^2)$ in that direction. [ESTABLISHED]

The proof is in appendix B.3.

**Theorem 9.30 (Center ratios, smooth weights, and the full Fisher limit).** Assume (9.11) and $\mathcal{I}_{\gamma,N} \to \mathcal{I}_\gamma$ in $C^2$ on a neighborhood of $\theta_0$. Let $S$ be the maximizing set of $\mathsf{s}_\gamma(\theta_0)$, choose $\gamma_0 \in S$, and assume the center-ratio condition

$$\exp\bigl[\lambda_{\gamma,N}(\theta_0) - \lambda_{\gamma_0,N}(\theta_0)\bigr] \longrightarrow \mathrm{c}_\gamma \in (0,\infty) \quad (\gamma \in S), \qquad \mathrm{c}_{\gamma_0} = 1. \tag{9.21}$$

Set $\mathsf{v}_\gamma = \nabla \mathsf{s}_\gamma(\theta_0)$ and define the window experiment on the original labeled carrier by $\theta_N(\mathsf{w}) = \theta_0 + \mathsf{w}/N$ and $\mathcal{E}_N(\mathsf{w}) = P_{N,\theta_N(\mathsf{w})}$. On every compact set of $\mathsf{w}$ its label law converges in $C^3$ to

$$\bar{\mathsf{p}}_\gamma(\mathsf{w}) = \frac{\mathrm{c}_\gamma e^{\mathsf{w}\cdot\mathsf{v}_\gamma}}{\sum_{\eta\in S} \mathrm{c}_\eta e^{\mathsf{w}\cdot\mathsf{v}_\eta}} \quad (\gamma \in S), \qquad \bar{\mathsf{p}}_\gamma(\mathsf{w}) = 0 \quad (\gamma \notin S), \tag{9.22}$$

and its Fisher tensor in the coordinate $\mathsf{w}$ converges in $C^2$:

$$g^F_{\mathsf{w}}(\mathcal{E}_N)(\mathsf{w}) = N^{-2} g^F(P_N)(\theta_N(\mathsf{w})) \longrightarrow \mathrm{Cov}_{\bar{\mathsf{p}}(\mathsf{w})}(\mathsf{v}_\gamma) = D^2_{\mathsf{w}} \log \sum_{\gamma\in S} \mathrm{c}_\gamma e^{\mathsf{w}\cdot\mathsf{v}_\gamma}. \tag{9.23}$$

The center-ratio condition is additional data at finite $N$; equality of the bulk densities supplies no such condition. [ESTABLISHED]

The proof is in appendix B.3.

**Proposition 9.31 (Ties that diverge or do not converge).** Extensive label growth, even with identical parameter-independent branches, implies neither a finite extensive Fisher density at a tie, nor convergence of the label law, nor $C^2$ convergence of a uniformly convergent Fisher density. For two labels $0, 1$ with one common branch law independent of $\theta$ and log-odds $\omega_N = \lambda_{1,N} - \lambda_{0,N}$,

$$W_{1,N} = \frac{e^{\omega_N}}{1 + e^{\omega_N}}, \qquad g^F(P_N) = \mathrm{bv}(\omega_N)(\omega_N')^2, \qquad \mathrm{bv}(x) = \frac{e^x}{(1+e^x)^2} = \frac{1}{4\cosh^2(x/2)}, \tag{9.24}$$

and each of the following families has extensive label growth on every compact interval. (a) For $\omega_N = N\theta$ the tie at $\theta = 0$ has $g^F(P_N)(\theta) = N^2/[4\cosh^2(N\theta/2)]$ and $g^F(P_N)(0)/N = N/4$, so the extensive density diverges, and $N^{-2} g^F(P_N)$ tends to $1/4$ at zero and to zero elsewhere. (b) For $\omega_N$ equal to $0$ at even $N$ and $\log 2$ at odd $N$, the Fisher tensor vanishes identically while $W_{1,N}$

alternates between $1/2$ and $2/3$. (c) For $\omega_N = N^{\mathrm{c}}\theta$ with $0 < \mathrm{c} < 1$, both bulk densities vanish and

$$\frac{g^F(P_N)(0)}{N} = \frac{1}{4}N^{2\mathrm{c}-1}, \tag{9.25}$$

which is below, at, or above the extensive scale according as $\mathrm{c} < 1/2$, $\mathrm{c} = 1/2$, or $\mathrm{c} > 1/2$; for $\omega_N = \mathrm{c}_N + \sqrt{N}\theta$ with $\mathrm{c}_N$ alternating between $0$ and $\log 2$, the density at zero alternates between $1/4$ and $2/9$. (d) For $\omega_N = N\theta - \mathrm{c}\log N$ with $\mathrm{c} > 0$, the tied label is rare, $W_{1,N}(0) = (1 + N^{\mathrm{c}})^{-1} \to 0$, yet

$$\frac{g^F(P_N)(0)}{N} = \frac{N^{1+\mathrm{c}}}{(1+N^{\mathrm{c}})^2} \sim N^{1-\mathrm{c}}, \tag{9.26}$$

so the limiting $N^2$ coefficient is zero while the density diverges for $\mathrm{c} < 1$, tends to one for $\mathrm{c} = 1$, and vanishes for $\mathrm{c} > 1$. (e) For $\omega_N(\theta) = N^{-1/6}\sin(N^{1/3}\theta)$ both weights converge uniformly to $1/2$ and $F_N = N^{-1}g^F(P_N)$ converges uniformly to zero, while

$$F_N'' = \frac{1}{N}\Big[\mathrm{bv}''(\omega_N)(\omega_N')^4 + 5\,\mathrm{bv}'(\omega_N)(\omega_N')^2\omega_N'' + 2\,\mathrm{bv}(\omega_N)(\omega_N'')^2 + 2\,\mathrm{bv}(\omega_N)\omega_N'\omega_N'''\Big] \tag{9.27}$$

satisfies $F_N''(0) = -\frac{1}{2} - \frac{1}{8}N^{-1/3} \to -\frac{1}{2}$. [ESTABLISHED]

The proof is in appendix B.3.

**Remark 9.32 (The retained label is observed information).** The covariance term of (9.13) is the Fisher tensor of the label marginal, so the $N^2$ scales of Proposition 9.29 and theorem 9.30 are information about the label. When all carriers are one common set, the experiment that observes only the configuration, through $(\gamma, x) \mapsto x$, loses exactly the expected conditional covariance of $\mathrm{s}_P$ given $X$, which vanishes when $X$ determines the label; for a record the loss is given by Proposition 9.25. In the binary families of Proposition 9.31 the configuration alone has a law independent of $\theta$ and zero Fisher information, while for $\omega_N = N\theta$ the labeled window experiment has Fisher tensor $\mathrm{bv}(\mathrm{w})$ in w for every $N$. [ESTABLISHED]

*Proof.* The label marginal has score $\xi_\gamma - \bar{\xi}$. The map $(\gamma, x) \mapsto x$ is a parameter-independent statistic, so Theorem 7.8 gives the loss as the expected conditional covariance, which vanishes when the score is a function of $X$, in particular when $X$ determines $(\gamma, x)$. With one common $\theta$-independent branch law $p$, the configuration has law $\sum_\gamma W_\gamma p = p$. For $\omega_N = N\theta$ at $\theta = \mathrm{w}/N$, (9.24) gives $N^{-2}\,\mathrm{bv}(\mathrm{w})N^2$. □

**Physicist's summary.** A coarse channel is a noisy detector on latent coordinates, under which relative entropy and Fisher information can only drop. Applied to the generative joint, the posterior, and the recognition law at once, it leaves the evidence untouched, and the free energy drops by exactly what it discarded about the disagreement of recognition and posterior inside each coarse cell, which is bookkeeping, not learning. A retained discrete label, such as which partition is present, carries its own Fisher information: with $N$ agents a strict winner hides it, a tie exposes information of order $N^2$, and near a tie the natural coordinate is $N(\theta - \theta_0)$, as in the finite-volume rounding of a first-order transition.

# Chapter 10

# Gauge-Retained Network State and Fixed Sectors

Grouping the agents at $c_*$ into blocks does not average away their gauge data: links are independent group elements and closed walks carry holonomy. This chapter asks what a coarse description of a block must retain of that data and which part of the block's state survives exact aggregation. A rooted spanning tree retains tree transports, the based holonomy, and block-dressed boundary edges as one root-framed datum (Proposition 10.5). The fixed sector of a block is the kernel of its connection Laplacian, with positive-definite weights the holonomy-fixed root values (Proposition 10.17), and transported marginal laws admit a parent law exactly when they occupy one holonomy-stabilized mode (Theorem 10.9). Cut links compress exactly into rectangular endpoint maps (Theorem 10.25), and into one group-valued link with new self terms exactly when an explicit matrix is invertible (Lemma 10.22), with one balanced choice of link and weight (Proposition 10.23) that no presentation at the level of laws reproduces once the cut has excess (Proposition 10.24) and that admits no continuous extension to all singular cuts (Proposition 10.26); the coarse frame is invisible to the constituents and cannot be computed from their frames by a symmetric equivariant rule (Proposition 10.27); and a record at $c_*$ distinguishes loop holonomies at the level of laws (Theorem 10.28). These results supply the retained gauge data of Chapters 11 and 12, the formation witnesses of Chapter 13, and the testability discussion of Chapter 14; Supplement G realizes them for Gaussian families.

**Scope.** The fixed-sector rank is an admissibility criterion: it assigns no partition cost, keeps singletons admissible, and does not require constituent beliefs to coincide. The marginal-law score takes its block and weights as inputs, selects no partition, supplies no Markov kernel, and does not recover the dependence structure of a joint law. A cross map parallel for the links neither identifies the belief and model bundles nor orders their two scores, and the operational record, which sees a loop only through its represented holonomy at the probe, recovers no link values. [NOT-CLAIMED] The obstruction to a continuous extension of balanced cut links is Proposition 10.26.

## 10.1 Levelwise fiber state and scale intertwiners at $c_*$

**Definition 10.1 (Levelwise fiber state).** At depth $\ell$ fix a Lie group $G_\ell$ and a right $G_\ell$-torsor $\mathscr{P}_\ell$, the fiber over $c_*$ of the principal bundle of that depth, containing one frame of each agent or meta-agent

of depth $\ell$, used for both channels. Law fibers $\mathcal{B}_{\ell,x}$, $x \in \{b,m\}$, carry left actions $\widehat{\rho}_{\ell,x}$ of $G_\ell$, which may be inequivalent, with associated fibers $\mathcal{E}_{\ell,x} = \mathscr{P}_\ell \times_{\widehat{\rho}_{\ell,x}} \mathcal{B}_{\ell,x}$. The depth-$\ell$ interaction multigraph $\mathfrak{I}_\ell$ carries links $\Theta_e^x \in G_\ell$ with $\Theta_{\bar{e}}^x = (\Theta_e^x)^{-1}$ (Definition 2.11), every block $I$ of depth $\ell$ carries its based holonomy groups $\mathrm{Hol}_I^x$ (Definition 10.4), and declared cross maps $\Phi_\ell : \mathcal{E}_{\ell,b} \to \mathcal{E}_{\ell,m}$ and $\widetilde{\Phi}_\ell : \mathcal{E}_{\ell,m} \to \mathcal{E}_{\ell,b}$ need not be invertible or mutually inverse. The fiber state of depth $\ell$ is

$$\mathfrak{X}_\ell = \Big(G_\ell, \mathscr{P}_\ell;\ \mathcal{E}_{\ell,b}, \mathcal{E}_{\ell,m};\ \mathfrak{I}_\ell, (\Theta_e^b, \Theta_e^m)_e, (\mathrm{Hol}_I^b, \mathrm{Hol}_I^m)_I;\ \Phi_\ell, \widetilde{\Phi}_\ell;\ \ldots\Big), \tag{10.1}$$

where the ellipsis stands for law, kernel, or operator components declared by a realization.

**Proposition 10.2 (Scale maps and intertwiners at $c_*$).** (i) A map $\Phi_\ell : \mathcal{E}_{\ell,b} \to \mathcal{E}_{\ell,m}$, expressed in a frame $u$ by the fiber map $f_u$ with $\Phi_\ell[u,w] = [u, f_u(w)]$, has one expression $f_\ell^{bm}$ in every frame exactly when $f_\ell^{bm} \circ \widehat{\rho}_{\ell,b}(g) = \widehat{\rho}_{\ell,m}(g) \circ f_\ell^{bm}$ for every $g \in G_\ell$, which is also the condition that $\Phi_\ell$ commutes with every automorphism of $\mathscr{P}_\ell$ and constrains the covariance of a cross map, not its existence; the reverse direction is analogous. For a homomorphism $\kappa_\ell : G_\ell \to G_{\ell+1}$, points $p_0 \in \mathscr{P}_\ell$ and $p_1 \in \mathscr{P}_{\ell+1}$, and fiber maps $q_{\ell,x} : \mathcal{B}_{\ell,x} \to \mathcal{B}_{\ell+1,x}$, $x \in \{b,m\}$, consider

$$\mathsf{P}_\ell(p \cdot g) = \mathsf{P}_\ell(p) \cdot \kappa_\ell(g), \tag{10.2}$$

$$\widehat{q}_{\ell,x}[p,w] = [\mathsf{P}_\ell(p), q_{\ell,x}(w)], \tag{10.3}$$

$$q_{\ell,x} \circ \widehat{\rho}_{\ell,x}(g) = \widehat{\rho}_{\ell+1,x}(\kappa_\ell(g)) \circ q_{\ell,x} \qquad \text{for every } g \in G_\ell. \tag{10.4}$$

(ii) Exactly one map $\mathsf{P}_\ell : \mathscr{P}_\ell \to \mathscr{P}_{\ell+1}$ satisfies $\mathsf{P}_\ell(p_0) = p_1$ and (10.2). (iii) The formula (10.3) defines a map $\widehat{q}_{\ell,x} : \mathcal{E}_{\ell,x} \to \mathcal{E}_{\ell+1,x}$ exactly when (10.4) holds. (iv) For composable steps the homomorphisms, torsor maps, fiber maps, and induced maps compose, the rightmost factor acting first, and on law fibers given by pushforward the composite is the pushforward of the composed kernel, by (9.1). [ESTABLISHED]

*Proof.* (i) Each point of $\mathcal{E}_{\ell,b}$ has one representative with first entry $u$, and $[u \cdot g, w] = [u, \widehat{\rho}_{\ell,b}(g)w]$ gives $f_{u\cdot g} = \widehat{\rho}_{\ell,m}(g)^{-1} \circ f_u \circ \widehat{\rho}_{\ell,b}(g)$, as in (2.22). An automorphism of the torsor, $u \mapsto u \cdot \kappa$, acts by $[u,w] \mapsto [u, \widehat{\rho}(\kappa)w]$, so commuting with all of them is the same condition. (ii) The action is free and transitive, so $\mathsf{P}_\ell(p_0 \cdot g) := p_1 \cdot \kappa_\ell(g)$ is well defined, satisfies (10.2), and is forced by it. (iii) The representatives $(p,w)$ and $(p \cdot g, \widehat{\rho}_{\ell,x}(g)^{-1}w)$ of one point have images $[\mathsf{P}_\ell(p), q_{\ell,x}(w)]$ and $[\mathsf{P}_\ell(p), \widehat{\rho}_{\ell+1,x}(\kappa_\ell(g))q_{\ell,x}(\widehat{\rho}_{\ell,x}(g)^{-1}w)]$, which agree for every $(p,w,g)$ exactly under (10.4). (iv) The composites satisfy the defining identities by substitution. □

On a chain of depths these composites form a scale diagram (Definition 9.20 and proposition 9.21); at the information-geometric tier an induced map also needs the pushed laws to stay in the coarse model and $q_{\ell,x}$ to be smooth. For an edge copy $e : j \to i$ of $\mathfrak{I}_\ell$, the cross-scale morphism defects and the link-intertwining defects are the ordered pairs

$$\begin{aligned} \mathbf{D}_\ell^\Phi &= \big(\widehat{q}_{\ell,m} \circ \Phi_\ell,\ \Phi_{\ell+1} \circ \widehat{q}_{\ell,b}\big), & \mathbf{D}_\ell^{\widetilde{\Phi}} &= \big(\widehat{q}_{\ell,b} \circ \widetilde{\Phi}_\ell,\ \widetilde{\Phi}_{\ell+1} \circ \widehat{q}_{\ell,m}\big), \\ \mathbf{D}_e^\Phi &= \big(\widehat{\mathsf{t}}_e^m \circ \Phi_\ell,\ \Phi_\ell \circ \widehat{\mathsf{t}}_e^b\big), & \mathbf{D}_e^{\widetilde{\Phi}} &= \big(\widehat{\mathsf{t}}_e^b \circ \widetilde{\Phi}_\ell,\ \widetilde{\Phi}_\ell \circ \widehat{\mathsf{t}}_e^m\big), \end{aligned} \tag{10.5}$$

where $\widehat{\mathsf{t}}^x_e[p,w] = [\mathsf{t}^x_e(p), w]$ is induced by the link map $\mathsf{t}^x_e$ of Proposition 2.12. A defect vanishes when its two components are equal; subtracting them needs a vector or affine realization.

**Proposition 10.3 (Covariance of the defect diagrams).** The maps $\widehat{\mathsf{t}}^x_e$ are well defined and act from the frame of $j$ to the frame of $i$ as $\widehat{\rho}_x(\Theta^x_e)$. Under passive re-expressions at depths $\ell$ and $\ell+1$, both components of every pair in (10.5) acquire the same coordinate changes at source and target, so the equality locus of each pair is gauge invariant; in a fiberwise-linear realization the difference of the components transforms by the same left and right laws. [ESTABLISHED]

*Proof.* Right equivariance of $\mathsf{t}^x_e$ gives $[\mathsf{t}^x_e(p\cdot g), \widehat{\rho}_x(g)^{-1}w] = [\mathsf{t}^x_e(p), w]$, and $\mathsf{t}^x_e(u_j\cdot g) = u_i\cdot\Theta^x_e g$ gives the coordinate form. Both components of each pair are frame-free maps between the same two associated fibers, so their frame expressions are conjugated by the same coordinate changes, and equality of two maps does not depend on the frames. □

## 10.2 Rooted trees, based holonomy, and dressed edges

Fix a depth and suppress its index. Let $\mathcal{P}$ be a partition of the agents into blocks and $\mathfrak{I}_I$ the internal multigraph of a block $I$, with its loops and parallel copies. A passive re-expression $g=(g_i)$ acts by $\Theta^x_e \mapsto g_i\Theta^x_e g_j^{-1}$ for $e: j\to i$ (2.28) and by $z^x_i \mapsto \rho_x(g_i)z^x_i$ on sample coordinates (2.9).

**Definition 10.4 (Rooted spanning-tree data).** The component-indexed meta-agent set of $\mathcal{P}$ is

$$\widehat{\mathcal{P}} = \{(I,\mathsf{c}) : I\in\mathcal{P},\ \mathsf{c}\in\pi_0(\mathfrak{I}_I)\}. \tag{10.6}$$

From here on $I, J$ denote elements of $\widehat{\mathcal{P}}$, each connected with one root $r_I$; the state of a disconnected block is the product over its components, and collapsing that product to one root, row, or feature is an additional coarse channel. Choose a rooted spanning tree in each $I$, nested across depths. With the path products $\Theta^x_p$ and holonomies $\mathrm{Hol}^x(p)$ of (2.30), the tree transport of $i\in I$ along the tree path $p_{I,i}$ from $i$ to $r_I$, the holonomy of the fundamental cycle of a non-tree copy $e: j\to i$ of $\mathfrak{I}_I$, the based holonomy group of $I$ at a vertex $r$, and the block-dressed edge of a boundary copy $e: j\to i$ from $J$ to $I$ are

$$\Theta^x_{I\leftarrow i} = \Theta^x_{p_{I,i}}, \qquad \mathrm{Hol}^x(p_e) = \Theta^x_{I\leftarrow i}\Theta^x_e(\Theta^x_{I\leftarrow j})^{-1}, \tag{10.7}$$

$$\mathrm{Hol}^x_I(r) = \{\Theta^x_p : p \text{ a closed walk of } \mathfrak{I}_I \text{ at } r\}\subseteq G, \qquad \mathrm{Hol}^x_I = \mathrm{Hol}^x_I(r_I), \tag{10.8}$$

$$\widehat{\Theta}^x_e = \Theta^x_{I\leftarrow i}\Theta^x_e(\Theta^x_{J\leftarrow j})^{-1}. \tag{10.9}$$

For an action $\rho_x$ of $G$ on the channel-$x$ sample fiber, the root-framed child tuple $\bar{z}^x_I = (\rho_x(\Theta^x_{I\leftarrow i})z^x_i)_{i\in I}$ collects the member coordinates in the root frame, and a law $Q$ on the root sample fiber of $I$ is holonomy stabilized when

$$\widehat{\rho}_x(\mathsf{k})Q = Q \qquad \text{for every } \mathsf{k}\in\mathrm{Hol}^x_I, \tag{10.10}$$

which is law coherence up to holonomy stabilizers on the block (Definition 2.17).

**Proposition 10.5 (Gauge laws of the rooted datum).** The map $p \mapsto \Theta^x_p$ on closed walks at $r_I$ is multiplicative under concatenation and invariant under insertion or removal of backtracks, so it is a homomorphism on the fundamental group of $\mathfrak{I}_I$ at $r_I$, with image $\mathrm{Hol}^x_I$ generated by the fundamental-cycle holonomies $\mathrm{Hol}^x(p_e)$. Under a passive re-expression $g$ the tree transports, this homomorphism, and, for a boundary copy $e : j \to i$ from $J$ to $I$, the link and the dressed edge transform as

$$\Theta^x_{I\leftarrow i} \longmapsto g_{r_I}\Theta^x_{I\leftarrow i}g_i^{-1}, \tag{10.11}$$

$$\pi_1(\mathfrak{I}_I, r_I) \longrightarrow G, \quad p \longmapsto \Theta^x_p, \qquad \Theta^x_p \longmapsto g_{r_I}\Theta^x_p g_{r_I}^{-1}, \tag{10.12}$$

$$\Theta^x_e \longmapsto g_i\Theta^x_e g_j^{-1}, \qquad \widehat{\Theta}^x_e \longmapsto g_{r_I}\widehat{\Theta}^x_e g_{r_J}^{-1}, \tag{10.13}$$

and $\bar{z}^x_I \mapsto \rho_x(g_{r_I})\bar{z}^x_I$, one root element acting on every entry. The datum consisting of $\bar{z}^x_I$, the homomorphism (10.12), and the dressed edges is therefore root framed and equivariant under one root re-expression per block. The homomorphism does not depend on the trees, and a change of tree multiplies each tree transport on the left by a based holonomy. The tree links can be set to the identity by a passive re-expression unique up to one element per block, after which each non-tree internal link equals the conjugated fundamental-cycle holonomy it closes. [ESTABLISHED]

*Proof.* Substitution cancels interior factors of path products and gives the law of $\bar{z}^x_I$, and a backtrack contributes $\Theta^x_e\Theta^x_{\bar{e}} = e$. Every internal copy satisfies $\Theta^x_e = (\Theta^x_{I\leftarrow i})^{-1}\mathrm{Hol}^x(p_e)\Theta^x_{I\leftarrow j}$, with $\mathrm{Hol}^x(p_e) = e$ on tree edges, so the product along a closed walk at $r_I$ telescopes to a product of fundamental-cycle holonomies and their inverses; walks and holonomy do not involve the trees, and for a second tree $\widetilde{\Theta}^x_{I\leftarrow i}(\Theta^x_{I\leftarrow i})^{-1}$ is the product along a closed walk at $r_I$. Finally $g_i = g_{r_I}\Theta^x_{I\leftarrow i}$ makes every tree transport, hence every tree link, the identity, is forced up to the choice of $g_{r_I}$, and turns a non-tree link into $g_{r_I}\mathrm{Hol}^x(p_e)g_{r_I}^{-1}$. □

The last statement is the maximal-tree gauge fixing of lattice gauge theory [22]. The measure tier retains the raw root-framed datum, since quotienting each holonomy separately would lose its orientation relative to the root features and boundary legs, and an orbit space would need a quotient theorem not supplied here.

**Proposition 10.6 (Two full-law treatments of holonomy for a candidate parent).** In the setting of Corollary 9.10 two alternatives are available, which need not exclude each other at different retained-coordinate or quotient levels. (i) Groupoid-covariant parent: supply bimeasurable actions of a typed groupoid on the fine random space, on $\mathsf{Z}_A$, on the observation space, and on the structural data; let the fine generative and recognition laws be covariant and the selected posterior versions compatible with the actions; and require $C_A(g\cdot y, g\cdot D) = C_A(y, D)$ for every admitted arrow $g$ and measurable $D \subseteq \mathsf{Z}_A$, together with covariance of the evaluator for the actions on $(m_A, \xi_A)$ and on $(b_A, o, h_A)$. Then the parent generative, posterior, and recognition laws are covariant. Invariance at one slice $(o, X)$ holds for the arrows that fix $(o, X)$ and the selected posterior versions; invariance of $q^b_A$ and $q^m_A$ separately is only a projection consequence and does not replace these hypotheses. (ii) Retention: the component roots, the based holonomy homomorphisms, and the dressed boundary

edges of Proposition 10.5 are placed in $\mathsf{H}_A$, and their joint law is pushed forward with no invariance claim. Neither alternative selects $I$, $A$, a partition, or a membership kernel. [ESTABLISHED]

*Proof.* (i) Substituting the invariance of $C_A$ and the covariant fine laws into (9.5) and (9.6) gives covariance of the three parent laws, and evaluator covariance carries it to the evaluator tier; at a fixed slice only arrows fixing $(o, X)$ and the versions map the slice to itself. (ii) defines the retained coordinates, and the pushforward is that of Corollary 9.10. □

## 10.3 Linear feature tier and tree dependence

**Proposition 10.7 (Linear cross-scale maps and their tree dependence).** Let $\rho_x^{\mathbb{W}} : G \to \mathrm{GL}(\mathbb{W}_x)$ be a linear representation on a feature fiber, choose weights $w_{Ii} \geq 0$ with $\sum_{i\in I} w_{Ii} = 1$, and set

$$(\mathsf{C}_x z)_I = \sum_{i\in I} w_{Ii}\rho_x^{\mathbb{W}}(\Theta^x_{I\leftarrow i})z_i, \qquad (\mathsf{S}_x \hat{z})_i = \rho_x^{\mathbb{W}}(\Theta^x_{I\leftarrow i})^{-1}\hat{z}_I \tag{10.14}$$

for a fine feature $z$ and a coarse feature $\hat{z} = (\hat{z}_I)_I$, with $\rho^{\mathbb{W}}_{x,f}(g) = \bigoplus_i \rho_x^{\mathbb{W}}(g_i)$ and $\rho^{\mathbb{W}}_{x,c}(g) = \bigoplus_I \rho_x^{\mathbb{W}}(g_{r_I})$. (i) $\mathsf{C}_x\mathsf{S}_x = I$, the averaged feature $(\mathsf{C}_x z)_I$ is a weighted mean of the entries of the root-framed tuple, which it compresses, and $\mathsf{C}'_x = \rho^{\mathbb{W}}_{x,c}(g)\mathsf{C}_x\rho^{\mathbb{W}}_{x,f}(g)^{-1}$, $\mathsf{S}'_x = \rho^{\mathbb{W}}_{x,f}(g)\mathsf{S}_x\rho^{\mathbb{W}}_{x,c}(g)^{-1}$. (ii) At nested depths $0 \to 1 \to 2$, with $\Theta^{x,12}_{A\leftarrow I}$ the tree transport of the coarse vertex $I$ to the root of its depth-two block $A$, the conditions $\Theta^{x,02}_{A\leftarrow i} = \Theta^{x,12}_{A\leftarrow I}\Theta^{x,01}_{I\leftarrow i}$ and $w^{02}_{Ai} = w^{12}_{AI}w^{01}_{Ii}$ give $\mathsf{C}^{02}_x = \mathsf{C}^{12}_x\mathsf{C}^{01}_x$ and $\mathsf{S}^{02}_x = \mathsf{S}^{01}_x\mathsf{S}^{12}_x$; nested forests alone give neither equality. (iii) For a fine message operator $\mathcal{A}_x$ with $\mathcal{A}'_x = \rho^{\mathbb{W}}_{x,f}(g)\mathcal{A}_x\rho^{\mathbb{W}}_{x,f}(g)^{-1}$, the fine-to-meta, meta-to-fine, and meta-to-meta operators $\mathcal{A}^x_{I\leftarrow j} = (\mathsf{C}_x\mathcal{A}_x)_{Ij}$, $\mathcal{A}^x_{j\leftarrow I} = (\mathcal{A}_x\mathsf{S}_x)_{jI}$, and $\mathcal{A}^x_{I\leftarrow J} = (\mathsf{C}_x\mathcal{A}_x\mathsf{S}_x)_{IJ}$ transform between the corresponding endpoint frames, and for a linear cross map $\Phi_f : \mathbb{W}_{b,f} \to \mathbb{W}_{m,f}$ with $\Phi'_f = \rho^{\mathbb{W}}_{m,f}(g)\Phi_f\rho^{\mathbb{W}}_{b,f}(g)^{-1}$ the compressed map $\Phi^c = \mathsf{C}_m\Phi_f\mathsf{S}_b$ satisfies $\Phi^{c\prime} = \rho^{\mathbb{W}}_{m,c}(g)\Phi^c\rho^{\mathbb{W}}_{b,c}(g)^{-1}$. (iv) A second family of trees gives transports $\widetilde{\Theta}^x_{I\leftarrow i} = \mathrm{Hol}^x(\gamma_i)\Theta^x_{I\leftarrow i}$, with $\gamma_i$ the closed walk from $r_I$ to $i$ along the first tree and back along the second, and

$$\begin{aligned}(\widetilde{\mathsf{C}}_x z)_I - (\mathsf{C}_x z)_I &= \sum_{i\in I} w_{Ii}\big(\rho_x^{\mathbb{W}}(\mathrm{Hol}^x(\gamma_i)) - I\big)\rho_x^{\mathbb{W}}(\Theta^x_{I\leftarrow i})z_i,\\ (\widetilde{\mathsf{S}}_x\hat{z})_i &= \rho_x^{\mathbb{W}}(\Theta^x_{I\leftarrow i})^{-1}\rho_x^{\mathbb{W}}(\mathrm{Hol}^x(\gamma_i))^{-1}\hat{z}_I.\end{aligned} \tag{10.15}$$

The difference is not a root re-expression, so $\mathsf{C}_x$, $\mathsf{S}_x$, and the message operators depend on the trees through gauge-invariant quantities: for two vertices joined by two edge copies with $\mathrm{SO}(2)$ links the identity and the rotation by $\vartheta$, the defining representation, equal weights, equal features $e_1$, and root at the first vertex, the two trees give averaged features of norms $1$ and $|\cos(\vartheta/2)|$, although these norms are invariant under every root re-expression. Tree independence holds for inputs whose root-framed features lie in the fixed space $\mathrm{Fix}(\rho_x^{\mathbb{W}} \circ \mathrm{Hol}^x_I) = \{v \in \mathbb{W}_x : \rho_x^{\mathbb{W}}(\mathsf{k})v = v \text{ for every } \mathsf{k} \in \mathrm{Hol}^x_I\}$, for $\mathsf{S}_x$ on coarse features in that subspace, and for all inputs when $\rho_x^{\mathbb{W}} \circ \mathrm{Hol}^x_I$ is trivial. [ESTABLISHED]

*Proof.* (i) $(\mathsf{C}_x\mathsf{S}_x\hat{z})_I = \sum_i w_{Ii}\hat{z}_I = \hat{z}_I$, and substituting (10.11) gives the covariance laws. (ii) Substitute the two conditions into the composites. (iii) follows from (i). (iv) The closed walk $\gamma_i$ has path product $\widetilde{\Theta}^x_{I\leftarrow i}(\Theta^x_{I\leftarrow i})^{-1}$, which gives (10.15). In the example the two trees give $e_1$ and

$\frac{1}{2}(e_1 + R_\vartheta e_1)$, of norm $|\cos(\vartheta/2)|$. If every $\rho_x^{\mathbb{W}}(\Theta^x_{I\leftarrow i})z_i$ is fixed by the represented holonomy, the difference vanishes, and likewise for $\mathsf{S}_x$. □

This fixed space is the feature-level counterpart of (10.10); a tree choice is a gauge fixing for the retained datum but part of the model for the averaged feature $\mathsf{C}_x z$.

## 10.4 Holonomy-conditioned marginal-law modes

**Definition 10.8 (Holonomy-conditioned marginal-law score).** Fix a channel $x$, a connected block $I$ with root $r$, weights $w_i > 0$ with $\sum_{i\in I} w_i = 1$, and marginal laws $q_i^x$ of the agents of $I$ on the sample fiber $\mathsf{Z}^x$, on which $G$ acts by bimeasurable bijections $\rho_x(g)$ with pushforward $\widehat{\rho}_x(g)$. Let parent-law families $\mathscr{M}_i^x \subseteq \mathrm{Prob}(\mathsf{Z}^x)$ be equivariant: $\widehat{\rho}_x(\Theta_p^x)\mathscr{M}_j^x = \mathscr{M}_i^x$ for every walk $p$ of $\mathfrak{I}_I$ from $j$ to $i$, and $\mathscr{M}_i^{x\prime} = \widehat{\rho}_x(g_i)\mathscr{M}_i^x$ under a passive re-expression. For walks $p_i$ from $i$ to $r$ put $\bar{q}_i^x = \widehat{\rho}_x(\Theta^x_{p_i})q_i^x$, and with the based holonomy group $\mathrm{Hol}_I^x(r)$ of (10.8) define the fixed parent family and the score

$$\mathscr{Q}^x_{I,\mathrm{fix}}(r) = \big\{Q \in \mathscr{M}_r^x : \widehat{\rho}_x(\mathsf{k})Q = Q \text{ for every } \mathsf{k} \in \mathrm{Hol}_I^x(r)\big\}, \tag{10.16}$$

$$\mathfrak{D}_I^x = \inf_{Q\in\mathscr{Q}^x_{I,\mathrm{fix}}(r)} \sum_{i\in I} w_i D_{\mathrm{KL}}(\bar{q}_i^x \| Q), \qquad \inf\varnothing := +\infty. \tag{10.17}$$

A minimizer, when one exists, is a holonomy-conditioned marginal-law mode.

**Theorem 10.9 (Invariant holonomy-conditioned marginal-law mode).** The score $\mathfrak{D}_I^x$ does not depend on the walks, the root, or the passive frames. If the infimum is attained, then $\mathfrak{D}_I^x = 0$ exactly when some $Q_\star \in \mathscr{Q}^x_{I,\mathrm{fix}}(r)$ satisfies $\bar{q}_i^x = Q_\star$ for every $i \in I$, and the laws $\widehat{\rho}_x(\Theta^x_{p_i})^{-1}Q_\star$ then form a holonomy-stabilized parallel marginal-law section. If $\mathscr{Q}^x_{I,\mathrm{fix}}(r) = \varnothing$, then $\mathfrak{D}_I^x = +\infty$. If $\mathfrak{D}_I^x = 0$ without attainment, the conclusion is only forward-KL approximability by fixed parent laws. [ESTABLISHED]

*Proof.* Relative entropy is invariant under bimeasurable bijections (3.24). Another walk $p_i'$ from $i$ to $r$ replaces $\bar{q}_i^x$ by $\widehat{\rho}_x(\mathsf{k})\bar{q}_i^x$ with $\mathsf{k} = \Theta^x_{p_i'}(\Theta^x_{p_i})^{-1} \in \mathrm{Hol}_I^x(r)$, and for every fixed parent $Q$, $D_{\mathrm{KL}}(\widehat{\rho}_x(\mathsf{k})\bar{q}_i^x\|Q) = D_{\mathrm{KL}}(\widehat{\rho}_x(\mathsf{k})\bar{q}_i^x\|\widehat{\rho}_x(\mathsf{k})Q) = D_{\mathrm{KL}}(\bar{q}_i^x\|Q)$, so every objective value is path independent. A walk $\delta$ from $r$ to $r'$ conjugates $\mathrm{Hol}_I^x(r)$ onto $\mathrm{Hol}_I^x(r')$, and pushforward by $\Theta_\delta^x$ maps the two fixed families bijectively, by equivariance, and the transported laws of $r$ to those of $r'$, preserving every term. A passive re-expression sends $\Theta_p^x$ to $g_i\Theta_p^x g_j^{-1}$ for $p$ from $j$ to $i$ and $q_i^x$ to $\widehat{\rho}_x(g_i)q_i^x$, so it pushes every transported law by $\widehat{\rho}_x(g_r)$, conjugates the holonomy group by $g_r$, and carries the fixed family bijectively, and invariance of relative entropy gives gauge invariance. Relative entropy vanishes exactly between equal laws and every $w_i$ is positive, which gives the zero case; since $Q_\star$ is holonomy fixed, the reconstructed laws do not depend on the walks and are law coherent by Proposition 2.18. The empty case is the convention, and without attainment the infimum supplies $Q_n$ with score below $1/n$ and no minimizer. □

**Proposition 10.10 (Compact holonomy: closed form and attainment).** In the setting of Theorem 10.9, let the represented holonomy group $\mathfrak{H}_r = \rho_x(\mathrm{Hol}_I^x(r))$ be a compact metrizable group acting jointly measurably on $\mathsf{Z}^x$, with Haar probability $d\mathsf{k}$, and put

$$\bar{P} = \sum_{i\in I} w_i \bar{q}_i^x, \qquad \bar{M}(B) = \int_{\mathfrak{H}_r} \bar{P}(\mathsf{k}^{-1}B) d\mathsf{k}. \tag{10.18}$$

Then $\bar{M}$ is $\mathfrak{H}_r$-invariant, and every $\mathfrak{H}_r$-invariant law $Q$ satisfies, in $[0,+\infty]$,

$$\sum_{i\in I} w_i D_{\mathrm{KL}}(\bar{q}_i^x \| Q) = \sum_{i\in I} w_i D_{\mathrm{KL}}(\bar{q}_i^x \| \bar{M}) + D_{\mathrm{KL}}(\bar{M}\|Q). \tag{10.19}$$

Consequently $\mathfrak{D}_I^x = \sum_i w_i D_{\mathrm{KL}}(\bar{q}_i^x\|\bar{M}) + \inf_{Q\in\mathscr{Q}_{I,\mathrm{fix}}^x(r)} D_{\mathrm{KL}}(\bar{M}\|Q)$. If $\bar{M}\in\mathscr{M}_r^x$, for instance when $\mathscr{M}_r^x = \mathrm{Prob}(\mathsf{Z}^x)$, the infimum is attained at $\bar{M}$, the score is the holonomy-symmetrized information radius $\sum_i w_i D_{\mathrm{KL}}(\bar{q}_i^x\|\bar{M})$, and $\bar{M}$ is the unique minimizer when that value is finite. If instead $\mathsf{Z}^x = \mathbb{R}^K$, the holonomy acts linearly, $\mathscr{M}_r^x$ is the nondegenerate Gaussian family, and $\bar{M}$ has finite second moments, mean $\bar{m}$, nonsingular covariance $\bar{S}$, and $D_{\mathrm{KL}}(\bar{M}\|\mathcal{N}(\bar{m},\bar{S})) < \infty$, then $\mathcal{N}(\bar{m},\bar{S})$ is holonomy invariant and attains the infimum, uniquely when $\mathfrak{D}_I^x < \infty$. For a finite holonomy group $d\mathsf{k}$ is uniform; a holonomy group that is only relatively compact in a group acting jointly continuously on a Polish fiber has the same fixed laws as its closure, to which the statement applies; and when every $\bar{q}_i^x$ is a nondegenerate Gaussian and the holonomy acts linearly with compact closure, the Gaussian hypotheses hold and the score is finite. [ESTABLISHED]

*Proof.* Left invariance of $d\mathsf{k}$ gives invariance of $\bar{M}$. Since $d\bar{q}_i^x/d\bar{P} \le w_i^{-1}$, splitting $\log(d\bar{q}_i^x/dQ)$ through $\bar{P}$ gives, for every $Q$, the compensation identity $\sum_i w_i D_{\mathrm{KL}}(\bar{q}_i^x\|Q) = \sum_i w_i D_{\mathrm{KL}}(\bar{q}_i^x\|\bar{P}) + D_{\mathrm{KL}}(\bar{P}\|Q)$ in $[0,+\infty]$: the negative part of the second term is integrable because $t(\log t)^- \le e^{-1}$, and both sides are infinite unless $\bar{P}\ll Q$. Applied with $Q$ and with $\bar{M}$, it reduces (10.19) to $D_{\mathrm{KL}}(\bar{P}\|Q) = D_{\mathrm{KL}}(\bar{P}\|\bar{M}) + D_{\mathrm{KL}}(\bar{M}\|Q)$ for invariant $Q$, where both sides are infinite unless $\bar{P}\ll Q$. If $\bar{P}\ll Q$ with density $f$, invariance of $Q$ gives $\bar{M}$ the invariant density $\bar{f}(z) = \int f(\mathsf{k}^{-1}z)d\mathsf{k}$; the invariant set $\{\bar{f}=0\}$ is null for $\bar{M}$ and hence for $\bar{P}$, so $\bar{P}\ll\bar{M}\ll Q$. Averaging the invariant function $\log\bar{f}$ over the group by Tonelli's theorem gives $\int\log\bar{f}d\bar{P} = \int\log\bar{f}d\bar{M} = D_{\mathrm{KL}}(\bar{M}\|Q)$ with finite negative parts, and integrating $\log f = \log(f/\bar{f}) + \log\bar{f}$ against $\bar{P}$ proves the identity. Attainment and uniqueness follow because relative entropy vanishes only between equal laws. In the Gaussian case, invariance of $\bar{M}$ and linearity give $\mathsf{k}\bar{m} = \bar{m}$ and $\mathsf{k}\bar{S}\mathsf{k}^\top = \bar{S}$, and $\log[d\mathcal{N}(\bar{m},\bar{S})/d\mathcal{N}(\mu,\Sigma)]$ is a quadratic polynomial whose $\bar{M}$- and $\mathcal{N}(\bar{m},\bar{S})$-expectations agree, which gives $D_{\mathrm{KL}}(\bar{M}\|\mathcal{N}(\mu,\Sigma)) = D_{\mathrm{KL}}(\bar{M}\|\mathcal{N}(\bar{m},\bar{S})) + D_{\mathrm{KL}}(\mathcal{N}(\bar{m},\bar{S})\|\mathcal{N}(\mu,\Sigma))$. For a relatively compact group, $\mathsf{k}\mapsto\int\phi(\mathsf{k}z)Q(dz)$ is continuous for bounded continuous $\phi$ by dominated convergence, so fixed laws of the group and of its closure agree. For Gaussian constituents with compact linear holonomy, $\bar{M}$ is a mixture of Gaussians with covariances in a compact set of positive-definite matrices, so its density is bounded and its second moments are finite, and convexity of relative entropy in its second argument bounds each $D_{\mathrm{KL}}(\bar{q}_i^x\|\bar{M})$ by an average of Gaussian relative entropies continuous on the compact group. □

**Corollary 10.11 (Noncompact Gaussian holonomy leaves no admissible parent).** Let $\mathsf{Z}^x = \mathbb{R}^K$, let the holonomy act by invertible linear maps, and let $\mathscr{M}_r^x$ be the nondegenerate Gaussian family. Then $\mathscr{Q}^x_{I,\mathrm{fix}}(r) \neq \varnothing$ if and only if the represented holonomy group is conjugate into $\mathrm{O}(K)$. If one loop holonomy has an eigenvalue off the unit circle, for instance $|\det| \neq 1$, or is not diagonalizable, then $\mathfrak{D}_I^x = +\infty$ for every configuration of marginals. For a block containing a loop with independent $\mathrm{GL}(K)$ links, the link configurations with $|\det| = 1$ on that loop form a Lebesgue-null set, so the score is $+\infty$ off a null set of configurations; it distinguishes configurations only after the links are restricted to a group conjugate into $\mathrm{O}(K)$ or a parent family with a nonempty fixed set is declared. [ESTABLISHED]

*Proof.* $\mathcal{N}(\mu, \Sigma)$ with $\Sigma \succ 0$ is fixed by $\mathsf{k}$ exactly when $\mathsf{k}\mu = \mu$ and $\mathsf{k}\Sigma\mathsf{k}^\top = \Sigma$. If one such $\Sigma$ is fixed by the whole group, then $\Sigma^{-1/2}\mathsf{k}\Sigma^{1/2}$ is orthogonal for every element; conversely, if $S^{-1}\mathsf{k}S \in \mathrm{O}(K)$ for every element, then $\mathcal{N}(0, SS^\top)$ is fixed. An element conjugate into $\mathrm{O}(K)$ is diagonalizable with eigenvalues of modulus one, and $\det(\mathsf{k}\Sigma\mathsf{k}^\top) = (\det \mathsf{k})^2 \det \Sigma$; an empty fixed family gives $+\infty$ by convention. The holonomy of a simple loop is a product of link matrices in which one chosen link $\Theta$ enters once, as $A_1\Theta^{\pm 1}A_2$ with $A_1, A_2$ fixed by the other links; this is a diffeomorphism of $\mathrm{GL}(K)$, so the preimage of the null set $\{|\det| = 1\}$ is null, and Fubini's theorem over the remaining links completes the argument. □

**Remark 10.12 (Holonomy score on a three-agent loop).** On the triangle of the running example (Section 2.5), now with belief sample fiber $\mathbb{R}^2$ and links in $\mathrm{SO}(2)$ acting by rotation, root the block at the first agent, take the tree edges joining it to the other two, and let the loop holonomy be a rotation by $\phi$ with $\sin\phi \neq 0$. For transported marginals $\bar{q}_i^b = \mathcal{N}(m_i, \sigma^2 I_2)$ with weights $w_i$ and the nondegenerate Gaussian parent family, the fixed family is $\{\mathcal{N}(0, sI_2) : s > 0\}$ and $\mathfrak{D}_I^b = \log(1 + \frac{1}{2\sigma^2}\sum_i w_i|m_i|^2)$, attained at $s = \sigma^2 + \frac{1}{2}\sum_i w_i|m_i|^2$, the moment-matched Gaussian of the Haar-symmetrized mixture; rotations preserve norms, so the score does not depend on the walks. With frame-aligned links on the same triangle the holonomy is trivial, every nondegenerate Gaussian is fixed, and the score is $\frac{1}{2}\log\det(I_2 + \mathrm{Cov}_w(m)/\sigma^2)$, which never exceeds the value above. Replacing one link by a matrix of determinant $2$ empties the fixed family and makes the score $+\infty$ (Corollary 10.11). [ESTABLISHED]

*Proof.* A fixed Gaussian has mean fixed by $R_\phi$, hence zero because $R_\phi - I$ is invertible, and a symmetric covariance commuting with $R_\phi$; the linear conditions on its off-diagonal entry and the difference of its diagonal entries have determinant $2\sin^2\phi \neq 0$, so the covariance is scalar. In dimension two, $D_{\mathrm{KL}}(\mathcal{N}(m, \sigma^2 I_2)\|\mathcal{N}(0, sI_2)) = \sigma^2/s + |m|^2/(2s) - 1 + \log(s/\sigma^2)$, whose weighted sum is minimized at the stated $s$ with the stated value. In the flat case moment matching gives $\mathcal{N}(\bar{m}, \sigma^2 I_2 + \mathrm{Cov}_w(m))$ and the stated value. With $c_1, c_2$ the eigenvalues of $\mathrm{Cov}_w(m)/\sigma^2$ and $\sum_i w_i|m_i|^2 = |\bar{m}|^2 + \mathrm{tr}\,\mathrm{Cov}_w(m)$, the inequality of arithmetic and geometric means gives $\frac{1}{2}\log((1 + c_1)(1 + c_2)) \leq \log(1 + \frac{1}{2}(c_1 + c_2))$, which is at most the holonomy value. A link of determinant $2$ gives the loop holonomy a determinant of modulus $2$. □

**Definition 10.13 (Exact common recovery).** A statistical experiment $\{P_\theta\}_{\theta\in\Theta}$ admits exact common recovery through a parameter-independent Markov kernel $K_\downarrow$ when a parameter-independent

Markov kernel $K_\uparrow$ satisfies

$$P_\theta K_\downarrow K_\uparrow = P_\theta \qquad \text{for every } \theta \in \Theta. \tag{10.20}$$

**Remark 10.14 (A zero score is not common recovery).** The score (10.17) depends on the laws of the block only through their transported marginals. When both marginals of a pair of agents are non-Dirac, there are joint laws $P_1 \neq P_2$ with the same marginals (Proposition 3.17), and every kernel $K_\downarrow$ for which $PK_\downarrow$ depends on $P$ only through its marginals sends them to one law, so the experiment $\{P_1, P_2\}$ admits no common recovery through it. [ESTABLISHED]

*Proof.* If $P_1K_\downarrow = P_2K_\downarrow$, a kernel $K_\uparrow$ recovering both would give $P_1 = P_1K_\downarrow K_\uparrow = P_2K_\downarrow K_\uparrow = P_2$. □

**Corollary 10.15 (Parallel cross morphisms preserve fixed marginal-law sections).** Let $\Phi : \mathcal{E}_b \to \mathcal{E}_m$ be a cross map, expressed in the frame of agent $i$ by $\Phi_i : \mathcal{B}_b \to \mathcal{B}_m$, whose link-intertwining defects (10.5) vanish on the internal copies of the block, so that $\widehat{\rho}_m(\Theta_p^m) \circ \Phi_j = \Phi_i \circ \widehat{\rho}_b(\Theta_p^b)$ for every walk $p$ from $j$ to $i$, and assume $\Phi_r(\mathscr{M}_r^b) \subseteq \mathscr{M}_r^m$. Then $\Phi_r(\mathscr{Q}_{I,\mathrm{fix}}^b(r)) \subseteq \mathscr{Q}_{I,\mathrm{fix}}^m(r)$, and applying $\Phi_i$ agentwise to a holonomy-stabilized parallel belief-law section gives such a model-law section. If $\widetilde{\Phi} : \mathcal{E}_m \to \mathcal{E}_b$ likewise has vanishing link-intertwining defects and $\widetilde{\Phi}_r(\mathscr{M}_r^m) \subseteq \mathscr{M}_r^b$, then $\widetilde{\Phi}_r(\mathscr{Q}_{I,\mathrm{fix}}^m(r)) \subseteq \mathscr{Q}_{I,\mathrm{fix}}^b(r)$, with the analogous statement for sections. [ESTABLISHED]

*Proof.* Vanishing defects on the edges give the intertwining for walks by composition, using the coordinate form of Proposition 10.3. If $Q \in \mathscr{Q}_{I,\mathrm{fix}}^b(r)$ and $p$ is a closed walk at $r$, then $\widehat{\rho}_m(\Theta_p^m)\Phi_r(Q) = \Phi_r(\widehat{\rho}_b(\Theta_p^b)Q) = \Phi_r(Q)$, and family compatibility gives the first inclusion. If the belief marginals transport to $Q$, the same square shows that their images transport to $\Phi_r(Q)$. Interchanging the channels and replacing $\Phi$ by $\widetilde{\Phi}$ gives the rest. □

## 10.5 Structural fixed sectors

**Definition 10.16 (Connection Laplacian and fixed fiber).** Fix a channel and a block $I$ with internal multigraph $\mathfrak{I}_I$, let the channel's sample fiber be a real vector space $\mathbb{V}$ of dimension $K$ with a linear representation $\rho$, and write $\Theta_e$ for the represented link $\rho(\Theta_e^x) \in \mathrm{GL}(\mathbb{V})$. For each unoriented internal copy choose an orientation $e : j \to i$ and a positive-semidefinite form $W_e$ on $\mathbb{V}$. The connection Laplacian of the block is the positive-semidefinite form $L_I$ on $\mathbb{V}^I$ with

$$z^\top L_I z = \sum_{e:j\to i} (z_i - \Theta_e z_j)^\top W_e (z_i - \Theta_e z_j), \tag{10.21}$$

in which loops and parallel copies each contribute one term. Its fixed fiber and rank are $\mathbb{V}_{I,\mathrm{fix}} = \ker L_I$ and $f_I = \dim \mathbb{V}_{I,\mathrm{fix}}$. For a root $r$ the represented based holonomy group $\mathfrak{H}_r = \rho(\mathrm{Hol}_I^x(r))$ has the fixed space $\mathrm{Fix}(\mathfrak{H}_r) = \{v \in \mathbb{V} : \mathsf{k}v = v \text{ for every } \mathsf{k} \in \mathfrak{H}_r\}$.

**Proposition 10.17 (Kernel-holonomy fixed-space isomorphism).** (i) The energy (10.21) does not depend on the chosen orientations when the reversed copy carries $W_{\bar{e}} = \Theta_e^\top W_e \Theta_e$. For

positive-semidefinite weights, $z \in \ker L_I$ exactly when

$$W_e^{1/2}(z_i - \Theta_e z_j) = 0 \qquad \text{for every internal copy } e : j \to i, \tag{10.22}$$

and $\ker L_I$ is the space $H^0$ of global sections of the cellular sheaf with vertex stalks $\mathbb{V}$ and restriction maps $W_e^{1/2}$ at $i$ and $W_e^{1/2}\Theta_e$ at $j$, of which $L_I$ is the degree-zero Laplacian [37, Sec. 3.2 and Thm. 3.1]; for orthogonal transports and scalar weights $L_I$ is the vector-diffusion operator of Singer and Wu [73, Sec. 3]. A passive re-expression, with the weights re-expressed as $W_e \mapsto \rho(g_i)^{-\top} W_e \rho(g_i)^{-1}$, carries $\ker L_I$ to $\rho_f(g) \ker L_I$, with $\rho_f(g) = \bigoplus_{i \in I} \rho(g_i)$. (ii) If $W_e \succ 0$ on every internal copy, then on every connected component $I_\alpha$ with root $r_\alpha$ root evaluation $z \mapsto z_{r_\alpha}$ is a linear isomorphism $\ker L_{I_\alpha} \to \mathrm{Fix}(\mathfrak{H}_{r_\alpha})$, so $\ker L_I \cong \bigoplus_\alpha \mathrm{Fix}(\mathfrak{H}_{r_\alpha})$ and $f_I = \sum_\alpha \dim \mathrm{Fix}(\mathfrak{H}_{r_\alpha})$. (iii) With positive-definite weights a connected block has $f_I = K$ exactly when every represented holonomy element is the identity, equivalently when the represented links are graph flat on $I$ (Proposition 2.13); this is the full fixed sector, $0 < f_I < K$ is a partial fixed sector, and $f_I = 0$ leaves no retained degree. A nonfaithful representation can make the represented holonomy trivial when the principal holonomy is not. (iv) With positive-semidefinite weights holonomy may survive in weight-null directions: on a triangle with $W_{12} = \mathrm{diag}(1, 0)$, $\Theta_{12} = \mathrm{diag}(1, 2)$, $W_{23} = \mathrm{diag}(0, 1)$, $\Theta_{23} = \mathrm{diag}(2, 1)$, and $W_{31} = \Theta_{31} = I$, the loop product is $2I$ and fixes no nonzero vector, while $\ker L_I$ is the two-dimensional space of constant configurations. [ESTABLISHED]

*Proof.* (i) Substituting $\Theta_{\bar{e}} = \Theta_e^{-1}$ and $W_{\bar{e}}$ reproduces the term of $e$. Since $L_I \succeq 0$, $z \in \ker L_I$ exactly when $z^\top L_I z = \sum_e |W_e^{1/2}(z_i - \Theta_e z_j)|^2 = 0$. The diagonal blocks of $L_I$ are sums of $W_e$ and $\Theta_e^\top W_e \Theta_e$ and its off-diagonal blocks are $-W_e\Theta_e$, the blocks of the sheaf Laplacian with the stated restriction maps, whose kernel is $H^0$ by the discrete Hodge theorem. The energy is unchanged under $z_i \mapsto \rho(g_i) z_i$, $\Theta_e \mapsto \rho(g_i)\Theta_e \rho(g_j)^{-1}$, and the weight law, so $L_I' = \rho_f(g)^{-\top} L_I \rho_f(g)^{-1}$. (ii) With positive-definite weights, $z \in \ker L_{I_\alpha}$ exactly when $z_i = \Theta_e z_j$ on every copy. Such a section is determined by its root value, and transporting that value around a closed walk returns it, so root evaluation is injective with image in $\mathrm{Fix}(\mathfrak{H}_{r_\alpha})$. Conversely, for $u \in \mathrm{Fix}(\mathfrak{H}_{r_\alpha})$ put $z_v = \rho(\Theta^x_{I_\alpha \leftarrow v})^{-1} u$. Tree constraints hold by construction, and for a non-tree copy $e : j \to i$ the represented fundamental-cycle holonomy $\rho(\mathrm{Hol}^x(p_e)) = \rho(\Theta^x_{I_\alpha \leftarrow i})\Theta_e \rho(\Theta^x_{I_\alpha \leftarrow j})^{-1}$ of (10.7) fixes $u$, which gives $\Theta_e z_j = z_i$. The Laplacian is the direct sum over components. (iii) $\mathrm{Fix}(\mathfrak{H}_r) = \mathbb{V}$ exactly when every element is the identity, the fundamental-cycle holonomies generate the group (Proposition 10.5), and elements of the kernel of $\rho$ are represented by the identity. (iv) The loop product is $\mathrm{diag}(1, 2)\,\mathrm{diag}(2, 1) = 2I$. The copy $31$ forces $z_3 = z_1$, the copy $12$ equal first coordinates of $z_1$ and $z_2$, and the copy $23$ equal second coordinates of $z_2$ and $z_3$, so the kernel consists of the constant configurations. □

**Proposition 10.18 (Fixed-section compression).** Let $\Lambda$ be a positive-semidefinite form on the fine configuration space, for instance the connection Laplacian of the network plus positive-semidefinite self terms. For a partition $\mathcal{P}$ into blocks with positive-semidefinite internal weights, let $\iota_I$ be a full-column matrix whose range is $\ker L_I$, and $\iota_\mathcal{P} = \mathrm{blkdiag}_{I \in \mathcal{P}} \iota_I$. Then

$$L_I \iota_I = 0, \qquad \Lambda_c = \iota_\mathcal{P}^\top \Lambda \iota_\mathcal{P} \succeq 0, \qquad \Lambda_c \succ 0 \iff \mathrm{range}(\iota_\mathcal{P}) \cap \ker \Lambda = \{0\}. \tag{10.23}$$

Internal edges contribute nothing to $\Lambda_c$, and cut edges become the rectangular endpoint data of Theorem 10.25. With positive-definite weights, the tree construction in the proof of Proposition 10.17 extends a basis of root values to one such matrix. [ESTABLISHED]

*Proof.* The first identity holds by construction and the second by congruence; $y^\top \Lambda_c y = 0$ exactly when $\iota_{\mathcal{P}} y \in \ker \Lambda$, and $\iota_{\mathcal{P}}$ is injective. An internal term vanishes on $\operatorname{range} \iota_I = \ker L_I$, and a cut term restricted to the ranges is the energy of Theorem 10.25 with the blocks of $\iota$ as injections. □

**Proposition 10.19 (Nested fixed sections compose exactly).** Let a larger block $A$ be a union of first-level blocks with positive-semidefinite internal weights, and let $\iota_{A\leftarrow\mathcal{P}} = \operatorname{blkdiag}_{I\subset A} \iota_I$ and $L_A^{(1)} = \iota_{A\leftarrow\mathcal{P}}^\top L_A \iota_{A\leftarrow\mathcal{P}}$. Then $\ker L_A = \iota_{A\leftarrow\mathcal{P}} \ker L_A^{(1)}$, and for every injection $\iota_A^{(2)} : \mathbb{V}_{A,\mathrm{fix}} \hookrightarrow \bigoplus_{I\subset A} \mathbb{V}_{I,\mathrm{fix}}$ with range $\ker L_A^{(1)}$ the fine fixed-section injection is $\iota_A^{\mathrm{fine}} = \iota_{A\leftarrow\mathcal{P}} \iota_A^{(2)}$. [ESTABLISHED]

*Proof.* A zero mode of $L_A$ has zero energy in every positive-semidefinite edge summand, so its restriction to each first-level block lies in $\ker L_I$ and it equals $\iota_{A\leftarrow\mathcal{P}} y$ for some $y$. Since $L_A \succeq 0$, $y \in \ker L_A^{(1)}$ exactly when $y^\top L_A^{(1)} y = 0$, that is, exactly when $\iota_{A\leftarrow\mathcal{P}} y \in \ker L_A$, which proves the first identity; substituting the range of $\iota_A^{(2)}$ gives the second, and associativity of matrix composition iterates the statement. □

**Remark 10.20 (Structural fixed sector versus marginal law occupation).** Relative entropy between the marginals of two agents needs one of them transported into the frame of the other and depends on the walk: a second walk changes the transported law by a loop-holonomy action, and the two values agree for every pair of laws only under stabilization; with $\mathsf{k} = \operatorname{diag}(1,-1,-1) \in \mathrm{SO}(3)$ and $q_i = q_j = \mathcal{N}(e_2, I_3)$, the identity walk gives zero and a walk differing by $\mathsf{k}$ gives $2$. An unpenalized sum of nonnegative within-block pairwise divergences vanishes at, and is minimized by, the singleton partition; it contains only marginal information and does not test common recovery (Remark 10.14). The structural fixed sector and the marginal-law score are independent: on a two-vertex tree with identity transport and $W = I_K$ the represented holonomy is trivial, yet $D_{\mathrm{KL}}(\mathcal{N}(-ae_1, I_K) \| \mathcal{N}(ae_1, I_K)) = 2a^2$ is arbitrarily large; on a triangle with $K = 3$, unit weights, identity transports on two edges, and closing transport $\mathsf{k}$, the represented holonomy is nontrivial and $\operatorname{Fix}(\mathsf{k}) = \operatorname{span}(e_1)$, so $f_I = 1 < K$, while $\mathsf{k}$ preserves $\mathcal{N}(0, \sigma^2 I_3)$ and assigning that law to every vertex makes every transported marginal divergence zero. Transported divergence thus measures the occupation of the available law modes, not the dimension of the structural fixed sector. [ESTABLISHED]

*Proof.* For equal covariances $D_{\mathrm{KL}}(\mathcal{N}(m_1, I) \| \mathcal{N}(m_2, I)) = \frac{1}{2}|m_1 - m_2|^2$, which gives $\frac{1}{2}|e_2 - \mathsf{k}e_2|^2 = 2$ and $2a^2$. A sum of nonnegative terms over pairs inside blocks is zero for singleton blocks. The fixed space of $\mathsf{k}$ is $\operatorname{span}(e_1)$, and $\mathsf{k}\sigma^2 I_3 \mathsf{k}^\top = \sigma^2 I_3$. □

## 10.6 Cut links and rectangular sheaf closure

Let $I$ and $J$ be blocks reduced to root coordinates $z_I, z_J \in \mathbb{V}$, with represented dressed boundary edges $\widehat{\Theta}_e \in \mathrm{GL}(\mathbb{V})$ (Definition 10.4) and positive-definite weights $W_e$ on the cut copies $e$ from $J$ to $I$.

The cut energy $\mathcal{E}_{\mathrm{cut}}(z_I, z_J) = \sum_e (z_I - \widehat{\Theta}_e z_J)^\top W_e (z_I - \widehat{\Theta}_e z_J)$ equals $z_I^\top \Lambda_{II} z_I - 2 z_I^\top \Lambda_{IJ} z_J + z_J^\top \Lambda_{JJ} z_J$ with $\Lambda_{II} = \sum_e W_e$, $\Lambda_{IJ} = \sum_e W_e \widehat{\Theta}_e$, and $\Lambda_{JJ} = \sum_e \widehat{\Theta}_e^\top W_e \widehat{\Theta}_e$. Coarse families admit newly generated positive-semidefinite self terms.

**Definition 10.21 (Anchored and no-new-anchor presentations of a cut).** An anchored one-link presentation of the cut is a tuple $(\Theta, W, \Lambda_I^{\mathrm{s}}, \Lambda_J^{\mathrm{s}})$ with $W \succ 0$, $\Theta \in \mathrm{GL}(\mathbb{V})$, and $\Lambda_I^{\mathrm{s}}, \Lambda_J^{\mathrm{s}} \succeq 0$ such that, for all $z_I, z_J$,

$$\mathcal{E}_{\mathrm{cut}}(z_I, z_J) = (z_I - \Theta z_J)^\top W (z_I - \Theta z_J) + z_I^\top \Lambda_I^{\mathrm{s}} z_I + z_J^\top \Lambda_J^{\mathrm{s}} z_J. \tag{10.24}$$

The no-new-anchor pair subfamily consists of the presentations with $\Lambda_I^{\mathrm{s}} = \Lambda_J^{\mathrm{s}} = 0$: one link and one weight, with no new self terms.

**Lemma 10.22 (Completed square for a cut).** One has $\Lambda_{JJ} \succ 0$ and, with $\bar{\Theta} = \Lambda_{II}^{-1} \Lambda_{IJ}$,

$$\mathcal{E}_{\mathrm{cut}}(z_I, z_J) = (z_I - \bar{\Theta} z_J)^\top \Lambda_{II} (z_I - \bar{\Theta} z_J) + z_J^\top \Delta z_J, \tag{10.25}$$

$$\Delta = \Lambda_{JJ} - \Lambda_{IJ}^\top \Lambda_{II}^{-1} \Lambda_{IJ} = \sum_e (\widehat{\Theta}_e - \bar{\Theta})^\top W_e (\widehat{\Theta}_e - \bar{\Theta}) \succeq 0. \tag{10.26}$$

(i) An anchored presentation exists if and only if $\det \Lambda_{IJ} \neq 0$; with links restricted to $\mathrm{GL}^+(\mathbb{V})$ the condition is $\det \Lambda_{IJ} > 0$. (ii) When $\det \Lambda_{IJ} \neq 0$, the anchored presentations are exactly $\Theta = W^{-1} \Lambda_{IJ}$, $\Lambda_I^{\mathrm{s}} = \Lambda_{II} - W$, and $\Lambda_J^{\mathrm{s}} = \Lambda_{JJ} - \Lambda_{IJ}^\top W^{-1} \Lambda_{IJ}$, with $W$ in the Loewner interval

$$\mathfrak{W}_{\mathrm{cut}} = \{W : \Lambda_{IJ} \Lambda_{JJ}^{-1} \Lambda_{IJ}^\top \preceq W \preceq \Lambda_{II}\} = \{W \succ 0 : W \preceq \Lambda_{II},\ \Lambda_{IJ}^\top W^{-1} \Lambda_{IJ} \preceq \Lambda_{JJ}\}. \tag{10.27}$$

The choice $W = \Lambda_{II}$ gives $\Lambda_I^{\mathrm{s}} = 0$ and $\Lambda_J^{\mathrm{s}} = \Delta$, and $W = \Lambda_{IJ} \Lambda_{JJ}^{-1} \Lambda_{IJ}^\top$ gives $\Lambda_J^{\mathrm{s}} = 0$, the reversed completion with link $\Lambda_{JJ}^{-1} \Lambda_{IJ}^\top$ from $I$ to $J$ and weight $\Lambda_{JJ}$. (iii) When $\det \Lambda_{IJ} \neq 0$, the anchored presentation is unique if and only if $\Delta = 0$, and it is then the bare pair $(\bar{\Theta}, \Lambda_{II}, 0, 0)$; otherwise there are infinitely many. The condition $\Delta = 0$ itself forces $\det \Lambda_{IJ} \neq 0$. (iv) The no-new-anchor pair subfamily contains a presentation of the cut if and only if $\Delta = 0$, that is, if and only if all cut twists $\widehat{\Theta}_e$ coincide. (v) The weighted mean $\bar{\Theta}$ has the gauge law of a dressed edge, $\bar{\Theta} \mapsto \rho(g_{r_I}) \bar{\Theta} \rho(g_{r_J})^{-1}$, but it can be singular: equal weights with twists $I$ and $-I$ give $\bar{\Theta} = 0$. [ESTABLISHED]

*Proof.* Each $\widehat{\Theta}_e$ is invertible and $W_e \succ 0$, so $\Lambda_{JJ} \succ 0$. Completing the square in $z_I$ gives (10.25), and expanding the sum in (10.26) gives $\Lambda_{JJ} - \Lambda_{IJ}^\top \Lambda_{II}^{-1} \Lambda_{IJ}$. Matching the three blocks of (10.24) gives $W\Theta = \Lambda_{IJ}$, $W + \Lambda_I^{\mathrm{s}} = \Lambda_{II}$, and $\Theta^\top W \Theta + \Lambda_J^{\mathrm{s}} = \Lambda_{JJ}$, hence the formulas of (ii); both self terms are positive semidefinite exactly on the second description in (10.27). The link is invertible exactly when $\Lambda_{IJ}$ is, and $\det \Theta = \det \Lambda_{IJ} / \det W$ has its sign. If $\det \Lambda_{IJ} \neq 0$, then $\Lambda_{IJ}^\top W^{-1} \Lambda_{IJ} \preceq \Lambda_{JJ}$ is equivalent to $W^{-1} \preceq \Lambda_{IJ}^{-\top} \Lambda_{JJ} \Lambda_{IJ}^{-1}$ and, since inversion reverses the Loewner order on positive-definite matrices, to $W \succeq \Lambda_{IJ} \Lambda_{JJ}^{-1} \Lambda_{IJ}^\top$; this is the first description, which contains $W = \Lambda_{II}$ because $\Delta \succeq 0$, and proves (i). The reversed completion is the same computation with $z_I$ and $z_J$ exchanged. Expanding the right side shows $\Lambda_{II} - \Lambda_{IJ} \Lambda_{JJ}^{-1} \Lambda_{IJ}^\top = \Lambda_{II} \Lambda_{IJ}^{-\top} \Delta \Lambda_{JJ}^{-1} \Lambda_{IJ}^\top$, so the interval is a single point exactly when $\Delta = 0$; otherwise every point of the segment between its endpoints is admissible and distinct weights give distinct presentations. If $\Delta = 0$ then $\Lambda_{IJ}^\top \Lambda_{II}^{-1} \Lambda_{IJ} = \Lambda_{JJ} \succ 0$, so $\Lambda_{IJ}$ is invertible. For (iv), a presentation without self terms needs $W = \Lambda_{II}$, $W\Theta = \Lambda_{IJ}$, and

$\Theta^{\top} W \Theta = \Lambda_{JJ}$, that is, $\Delta = 0$, and the sum in (10.26) with $W_e \succ 0$ vanishes exactly when every $\widehat{\Theta}_e = \bar{\Theta}$. For (v), invariance of the energy under $z_I \mapsto \rho(g_{r_I}) z_I$ and $z_J \mapsto \rho(g_{r_J}) z_J$ transforms $W_e$ at the receiving end and $\widehat{\Theta}_e$ as in (10.13), and substitution gives the law of $\bar{\Theta}$; the example is direct. □

**Proposition 10.23 (Balanced presentation of a cut).** In the setting of Lemma 10.22 with $\det \Lambda_{IJ} \neq 0$, write $W_{\mathrm{lo}} = \Lambda_{IJ} \Lambda_{JJ}^{-1} \Lambda_{IJ}^{\top}$ and $W_{\mathrm{hi}} = \Lambda_{II}$ for the endpoints of $\mathfrak{W}_{\mathrm{cut}}$, and call an anchored presentation balanced when the self term at $J$, carried to $I$ along the presentation's own link, equals the self term at $I$,

$$\Theta^{-\top} \Lambda_J^{\mathrm{s}} \Theta^{-1} = \Lambda_I^{\mathrm{s}}, \qquad \text{equivalently} \qquad \Theta^{\top} \Lambda_{II} \Theta = \Lambda_{JJ}, \tag{10.28}$$

so that the balanced link carries the form $\Lambda_{JJ}$ to the form $\Lambda_{II}$. (i) Exactly one anchored presentation is balanced. Its weight is the matrix geometric mean of the endpoints,

$$W_{\mathrm{bal}} = W_{\mathrm{lo}}^{1/2} \big( W_{\mathrm{lo}}^{-1/2} W_{\mathrm{hi}} W_{\mathrm{lo}}^{-1/2} \big)^{1/2} W_{\mathrm{lo}}^{1/2}, \tag{10.29}$$

the unique positive-definite solution of $W W_{\mathrm{lo}}^{-1} W = W_{\mathrm{hi}}$, and its link is $\Theta_{\mathrm{bal}} = W_{\mathrm{bal}}^{-1} \Lambda_{IJ}$. (ii) Under root re-expressions the weight transforms as a weight at $I$, $W_{\mathrm{bal}} \mapsto \rho(g_{r_I})^{-\top} W_{\mathrm{bal}} \rho(g_{r_I})^{-1}$, the link as a dressed edge, $\Theta_{\mathrm{bal}} \mapsto \rho(g_{r_I}) \Theta_{\mathrm{bal}} \rho(g_{r_J})^{-1}$, and each self term as a form at its own block. The presentation is end symmetric: read from $J$, with link $\Theta_{\mathrm{bal}}^{-1}$ and weight $\Theta_{\mathrm{bal}}^{\top} W_{\mathrm{bal}} \Theta_{\mathrm{bal}}$, it is the balanced presentation of the reversed cut, so that weight is the geometric mean of the endpoints $\Lambda_{IJ}^{\top} \Lambda_{II}^{-1} \Lambda_{IJ}$ and $\Lambda_{JJ}$ of $J$. (iii) For scalar weights $W_e = w_e I$ and orthogonal twists, an anchored presentation is balanced exactly when its link is orthogonal; then $W_{\mathrm{bal}} = (\Lambda_{IJ} \Lambda_{IJ}^{\top})^{1/2}$ and $\Theta_{\mathrm{bal}}$ is the orthogonal polar factor of $\bar{\Theta}$. [ESTABLISHED]

*Proof.* By part (ii) of the lemma an anchored presentation has $\Theta = W^{-1} \Lambda_{IJ}$, $\Lambda_I^{\mathrm{s}} = \Lambda_{II} - W$, and $\Lambda_J^{\mathrm{s}} = \Lambda_{JJ} - \Lambda_{IJ}^{\top} W^{-1} \Lambda_{IJ}$, so $\Theta^{-\top} \Lambda_J^{\mathrm{s}} \Theta^{-1} = W \Lambda_{IJ}^{-\top} \Lambda_{JJ} \Lambda_{IJ}^{-1} W - W = W W_{\mathrm{lo}}^{-1} W - W$, and the first form of (10.28) reads $W W_{\mathrm{lo}}^{-1} W = W_{\mathrm{hi}}$. The second form is $\Lambda_J^{\mathrm{s}} = \Theta^{\top} \Lambda_I^{\mathrm{s}} \Theta$ with the common term $\Theta^{\top} W \Theta$ removed from both sides, and after the substitution $\Theta = W^{-1} \Lambda_{IJ}$ it reads $W^{-1} W_{\mathrm{hi}} W^{-1} = W_{\mathrm{lo}}^{-1}$, the same equation. (i) The substitution $W = W_{\mathrm{lo}}^{1/2} Y W_{\mathrm{lo}}^{1/2}$ is a bijection of positive-definite matrices and turns the equation into $Y^2 = M$ with $M = W_{\mathrm{lo}}^{-1/2} W_{\mathrm{hi}} W_{\mathrm{lo}}^{-1/2} \succ 0$. A positive-definite $Y$ with $Y^2 = M$ commutes with $M$, so it is diagonal in an eigenbasis of $M$ with positive eigenvalues whose squares are those of $M$; hence $Y = M^{1/2}$, which is (10.29). Since $\Delta \succeq 0$ gives $W_{\mathrm{lo}} \preceq W_{\mathrm{hi}}$, one has $M \succeq I$, hence $I \preceq M^{1/2} \preceq M$, and congruence by $W_{\mathrm{lo}}^{1/2}$ gives $W_{\mathrm{lo}} \preceq W_{\mathrm{bal}} \preceq W_{\mathrm{hi}}$; the solution lies in $\mathfrak{W}_{\mathrm{cut}}$ and is the weight of exactly one anchored presentation. (ii) As in the proof of part (v) of the lemma, a root re-expression acts by $W_e \mapsto g_I^{-\top} W_e g_I^{-1}$ and $\widehat{\Theta}_e \mapsto g_I \widehat{\Theta}_e g_J^{-1}$ with $g_I = \rho(g_{r_I})$ and $g_J = \rho(g_{r_J})$, so $\Lambda_{II}$, $\Lambda_{IJ}$, and $\Lambda_{JJ}$ transform by $g_I^{-\top}(\cdot) g_I^{-1}$, $g_I^{-\top}(\cdot) g_J^{-1}$, and $g_J^{-\top}(\cdot) g_J^{-1}$, and both endpoints by $g_I^{-\top}(\cdot) g_I^{-1}$. If $W$ solves $W W_{\mathrm{lo}}^{-1} W = W_{\mathrm{hi}}$, then $g_I^{-\top} W g_I^{-1}$ solves the transformed equation, so uniqueness gives the law of $W_{\mathrm{bal}}$, and substitution gives the laws of the link and of the self terms. For end symmetry, $(z_I - \Theta z_J)^{\top} W (z_I - \Theta z_J) = (z_J - \Theta^{-1} z_I)^{\top} \Theta^{\top} W \Theta (z_J - \Theta^{-1} z_I)$, so an anchored presentation of the cut, read from $J$, is the anchored presentation $(\Theta^{-1}, \Theta^{\top} W \Theta, \Lambda_J^{\mathrm{s}}, \Lambda_I^{\mathrm{s}})$ of the reversed cut, whose blocks are $\Lambda_{JJ}$, $\Lambda_{IJ}^{\top}$, and $\Lambda_{II}$. The second form of (10.28) is equivalent to $\Theta^{-\top} \Lambda_{JJ} \Theta^{-1} = \Lambda_{II}$, the

balance condition of the reversed cut, so the balanced presentation read from $J$ is balanced for the reversed cut, and (i) applied to that cut identifies its weight. (iii) Here $\Lambda_{II} = \Lambda_{JJ} = (\sum_e w_e)I$, so the second form of (10.28) says that $\Theta$ is orthogonal. For orthogonal $\Theta$ the weight $W = \Lambda_{IJ}\Theta^\top$ is positive definite exactly when $\Lambda_{IJ} = W\Theta$ is a polar decomposition of $\Lambda_{IJ}$, and then $W^2 = \Lambda_{IJ}\Lambda_{IJ}^\top$ forces $W = (\Lambda_{IJ}\Lambda_{IJ}^\top)^{1/2}$ and $\Theta = W^{-1}\Lambda_{IJ}$, the polar factor of $\Lambda_{IJ}$ and of $\bar\Theta = \Lambda_{IJ}/\sum_e w_e$. □

**Proposition 10.24 (No balanced presentation of a cut at the level of laws).** Let $\mathrm{Gau}(\mathbb{V})$ be the fiber of nondegenerate Gaussian laws on $\mathbb{V}$, and let $k$ copies $e$ from $J$ to $I$ carry dressed edges $\widehat\Theta_e \in \mathrm{GL}(\mathbb{V})$ and weights $w_e > 0$ with $\sum_e w_e \widehat\Theta_e^{-1}$ invertible. For $q_J = \mathcal{N}(m_J, \Sigma_J) \in \mathrm{Gau}(\mathbb{V})$ the transported laws $(\widehat\Theta_e)_\# q_J = \mathcal{N}(\widehat\Theta_e m_J, \Sigma_e)$, $\Sigma_e = \widehat\Theta_e \Sigma_J \widehat\Theta_e^\top$, form the cut of Lemma 10.22 with weights $W_e = \frac{1}{2} w_e \Sigma_e^{-1}$ and twists $\widehat\Theta_e$; its blocks $\Lambda_{II}(\Sigma_J)$, $\Lambda_{IJ}(\Sigma_J)$, $\Lambda_{JJ}(\Sigma_J)$, its excess $\Delta(\Sigma_J)$ (10.26), and its balanced weight and link $W_{\mathrm{bal}}(\Sigma_J)$ and $\Theta_{\mathrm{bal}}(\Sigma_J)$ (10.29) are functions of $\Sigma_J$, with $\Lambda_{IJ}(\Sigma_J)$ invertible for every $\Sigma_J$. The cut has excess, $\Delta(\Sigma_J) \neq 0$ for one and hence for every $\Sigma_J$, exactly when the dressed edges are not all equal. Let $\mathscr{C} \subseteq \mathrm{Gau}(\mathbb{V})$ be nonempty and closed under rescaling of the covariance, so that $\mathcal{N}(m, \Sigma) \in \mathscr{C}$ implies $\mathcal{N}(m, t\Sigma) \in \mathscr{C}$ for every $t > 0$. A presentation of the copies at the level of laws on $\mathscr{C}$ is a weight $W' > 0$, an image map $\mathfrak{T} : \mathscr{C} \to \mathrm{Gau}(\mathbb{V})$, and self terms $\mathfrak{S}_I : \mathrm{Gau}(\mathbb{V}) \to \mathbb{R}$ and $\mathfrak{S}_J : \mathscr{C} \to \mathbb{R}$ with

$$\sum_e w_e D_{\mathrm{KL}}\big(q_I \| (\widehat\Theta_e)_\# q_J\big) = W' D_{\mathrm{KL}}\big(q_I \| \mathfrak{T}[q_J]\big) + \mathfrak{S}_I(q_I) + \mathfrak{S}_J(q_J) \quad \text{for all } q_I \in \mathrm{Gau}(\mathbb{V}) \text{ and } q_J \in \mathscr{C}, \tag{10.30}$$

and it is balanced when its Gaussian mean part is the link term of the balanced presentation of the cut of $\Sigma_J$, that is, writing $\mathfrak{T}[q_J] = \mathcal{N}(\mu_{\mathfrak{T}}, \Sigma_{\mathfrak{T}})$,

$$\tfrac{1}{2} W' \Sigma_{\mathfrak{T}}^{-1} = W_{\mathrm{bal}}(\Sigma_J) \quad \text{and} \quad \mu_{\mathfrak{T}} = \Theta_{\mathrm{bal}}(\Sigma_J) m_J \qquad \text{for every } q_J = \mathcal{N}(m_J, \Sigma_J) \in \mathscr{C}. \tag{10.31}$$

If the cut has excess, no presentation at the level of laws on $\mathscr{C}$ is balanced. Two copies with dressed edges $I$ and $2I$ and weights $w_1, w_2$ have $\Delta(\Sigma_J) = \frac{w_1 w_2}{2(4w_1 + w_2)} \Sigma_J^{-1}$, so they admit no balanced presentation on any such $\mathscr{C}$. [ESTABLISHED]

*Proof.* Since $\Sigma_e^{-1} = \widehat\Theta_e^{-\top} \Sigma_J^{-1} \widehat\Theta_e^{-1}$, the blocks are $\Lambda_{II}(\Sigma_J) = \frac{1}{2}\sum_e w_e \widehat\Theta_e^{-\top} \Sigma_J^{-1} \widehat\Theta_e^{-1}$, $\Lambda_{IJ}(\Sigma_J) = \frac{1}{2}\big(\sum_e w_e \widehat\Theta_e^{-\top}\big)\Sigma_J^{-1}$, and $\Lambda_{JJ}(\Sigma_J) = \frac{1}{2}\big(\sum_e w_e\big)\Sigma_J^{-1}$, so $\Lambda_{IJ}(\Sigma_J)$ is invertible exactly when $\sum_e w_e \widehat\Theta_e^{-1}$ is, and each block is homogeneous of degree $-1$ in $\Sigma_J$. Hence so are the endpoints $W_{\mathrm{lo}} = \Lambda_{IJ}\Lambda_{JJ}^{-1}\Lambda_{IJ}^\top$ and $W_{\mathrm{hi}} = \Lambda_{II}$, the excess $\Delta$, and, since (10.29) is jointly homogeneous of degree one in its two arguments, the balanced weight, $W_{\mathrm{bal}}(t\Sigma_J) = t^{-1} W_{\mathrm{bal}}(\Sigma_J)$. By part (iv) of the lemma, $\Delta(\Sigma_J) = 0$ exactly when all $\widehat\Theta_e$ coincide, a condition that does not involve $\Sigma_J$.

Fix a presentation, a law $q_J = \mathcal{N}(m_J, \Sigma_J) \in \mathscr{C}$, and $q_I = \mathcal{N}(m_I, \Sigma_I) \in \mathrm{Gau}(\mathbb{V})$; write $\mathfrak{T}[q_J] = \mathcal{N}(\mu_{\mathfrak{T}}, \Sigma_{\mathfrak{T}})$, $B = \frac{1}{2} W' \Sigma_{\mathfrak{T}}^{-1}$, and abbreviate $\Lambda_{II} = \Lambda_{II}(\Sigma_J)$, $\Lambda_{IJ} = \Lambda_{IJ}(\Sigma_J)$. For Gaussian laws $D_{\mathrm{KL}}(\mathcal{N}(m_1, \Sigma_1) \| \mathcal{N}(m_2, \Sigma_2)) = \frac{1}{2}(m_1 - m_2)^\top \Sigma_2^{-1}(m_1 - m_2) + \frac{1}{2}[\mathrm{Tr}(\Sigma_2^{-1}\Sigma_1) - \dim \mathbb{V} + \log\det \Sigma_2 - \log\det \Sigma_1]$, so both sides of (10.30) are quadratic polynomials in $m_I$, and their difference is

$$\begin{aligned} 0 = {} & m_I^\top (\Lambda_{II} - B) m_I - 2 m_I^\top (\Lambda_{IJ} m_J - B\mu_{\mathfrak{T}}) + \mathrm{Tr}[(\Lambda_{II} - B)\Sigma_I] \\ & + \tfrac{1}{2}\big(W' - \textstyle\sum_e w_e\big) \log\det \Sigma_I - \mathfrak{S}_I(q_I) + \mathrm{c}(q_J), \end{aligned}$$

where $\mathrm{c}(q_J)$ collects the terms that read $q_J$ alone. Subtracting this identity for a second law $q'_J \in \mathscr{C}$ removes $\mathfrak{S}_I(q_I)$ and the $\log\det\Sigma_I$ term and leaves a quadratic polynomial in $m_I$ that vanishes for every $m_I \in \mathbb{V}$ and every $\Sigma_I \succ 0$, so its quadratic and linear coefficients vanish: $M = \Lambda_{II}(\Sigma_J) - B$ and $v = \Lambda_{IJ}(\Sigma_J)m_J - B\mu_{\mathfrak{T}}$ take the same value at every $q_J \in \mathscr{C}$. If the presentation is balanced, (10.31) gives $B = W_{\mathrm{bal}}(\Sigma_J)$ and $B\mu_{\mathfrak{T}} = W_{\mathrm{bal}}(\Sigma_J)\Theta_{\mathrm{bal}}(\Sigma_J)m_J = \Lambda_{IJ}(\Sigma_J)m_J$, so $v = 0$ and $M = \Lambda_{II}(\Sigma_J) - W_{\mathrm{bal}}(\Sigma_J)$ for every $q_J \in \mathscr{C}$. Comparing $q_J$ with its rescaling $\mathcal{N}(m_J, t\Sigma_J) \in \mathscr{C}$, the homogeneity gives $M = t^{-1}M$ for every $t > 0$, hence $M = 0$ and $W_{\mathrm{bal}}(\Sigma_J) = W_{\mathrm{hi}}(\Sigma_J)$. Substituting into $WW_{\mathrm{lo}}^{-1}W = W_{\mathrm{hi}}$ gives $W_{\mathrm{hi}}W_{\mathrm{lo}}^{-1}W_{\mathrm{hi}} = W_{\mathrm{hi}}$, so $W_{\mathrm{lo}} = W_{\mathrm{hi}}$, and the identity $\Lambda_{II} - \Lambda_{IJ}\Lambda_{JJ}^{-1}\Lambda_{IJ}^{\top} = \Lambda_{II}\Lambda_{IJ}^{-\top}\Delta\Lambda_{JJ}^{-1}\Lambda_{IJ}^{\top}$ of the proof of the lemma, whose outer factors are invertible, gives $\Delta(\Sigma_J) = 0$: the cut has no excess. For the two copies $I$ and $2I$, $\Sigma_1 = \Sigma_J$ and $\Sigma_2 = 4\Sigma_J$, so $\Lambda_{II} = \frac{1}{8}(4w_1 + w_2)\Sigma_J^{-1}$, $\Lambda_{IJ} = \frac{1}{4}(2w_1 + w_2)\Sigma_J^{-1}$, and $\Lambda_{JJ} = \frac{1}{2}(w_1 + w_2)\Sigma_J^{-1}$, whence $\Delta = \big[\frac{1}{2}(w_1 + w_2) - \frac{(2w_1+w_2)^2}{2(4w_1+w_2)}\big]\Sigma_J^{-1} = \frac{w_1 w_2}{2(4w_1+w_2)}\Sigma_J^{-1}$. □

The balance condition is imposed only on the mean part of the link term, so the conclusion does not depend on whether the image map is gauge covariant or the self terms are end symmetric: a self term that reads one law cannot carry $\Lambda_{II} - W_{\mathrm{bal}}$, which is homogeneous of degree $-1$ in the covariance of the other law, and the excess of the cut is exactly what keeps that difference from vanishing.

The exact closed category is larger than one invertible link; its endpoint maps and stalks use the notation of the comparison maps and spaces of Chapter 6, which are an instance.

**Theorem 10.25 (Rectangular endpoint-map closure).** Choose a minimal factor $W_e = (S_e^Q)^{\top}S_e^Q$ with $S_e^Q : \mathbb{V} \to H_e$ of full row rank, and put $S_e^K = S_e^Q\Theta_e$. Then $(z_i - \Theta_e z_j)^{\top}W_e(z_i - \Theta_e z_j) = \|S_e^Q z_i - S_e^K z_j\|_{H_e}^2$. For linear maps $\iota_{iI} : \mathbb{V}_I \to \mathbb{V}$ from block coordinates, the coarse endpoint maps $S_{I,e}^{Q,c} = S_e^Q\iota_{iI}$ and $S_{J,e}^{K,c} = S_e^K\iota_{jJ}$ reproduce the Galerkin compression of the cut energy exactly, parallel cut copies compress losslessly into the direct sum of their stalks $H_e$, and nested blocking is associative. Parallel copies between the same two coarse vertices also admit a minimal exact compression: with coboundary rows $\delta_e = (S_{I,e}^{Q,c}, -S_{J,e}^{K,c})$, one edge with stalk $H$ and coboundary $\delta$ reproduces their summed energy exactly when $\delta^{\top}\delta = \sum_e \delta_e^{\top}\delta_e$; the minimal stalk dimension is $\operatorname{rank}\sum_e \delta_e^{\top}\delta_e \le \dim\mathbb{V}_I + \dim\mathbb{V}_J$, attained by any full-row-rank factor, and two full-row-rank factors differ by an orthogonal change of edge basis. [ESTABLISHED]

*Proof.* Expanding the square gives the pair form. Precomposition gives the coarse maps, whose cross block $-(S_{I,e}^{Q,c})^{\top}S_{J,e}^{K,c} = -\iota_{iI}^{\top}W_e\Theta_e\iota_{jJ}$ and diagonal blocks agree with the compressed energy. Stacking in $\bigoplus_e H_e$ sums the energies, and a second injection $\iota'$ gives $(S\iota)\iota' = S(\iota\iota')$. Two quadratic forms agree exactly when their symmetric matrices agree, which gives the compression criterion and the minimal stalk dimension; for full-row-rank factors $\delta_1, \delta_2$ of one matrix, $O = \delta_2\delta_1^{+}$ satisfies $O^{\top}O = I$ and $\delta_2 = O\delta_1$. □

This is the degree-zero cellular-sheaf Laplacian category [37], in which vertex spaces may have unequal dimensions and automorphism groups $\mathrm{GL}(\mathbb{V}_I)$ and edge basis changes act on both endpoint maps and the edge metric. Twice the quadratic disagreement $V_e$ of Definition 6.15 is the edge term of such an energy once its comparison precision is absorbed into the endpoint maps, as

in the discourse sheaves of Hansen and Ghrist [38]; the closure selects no decoder and no law, and the normalized realization for fixed maps is (6.41).

**Proposition 10.26 (No continuous extension of balanced cut links).** There is no rule assigning a matrix-valued link and a positive-semidefinite weight and self terms to every cut that is continuous in the constituent weights and twists and agrees with Proposition 10.23 whenever $\Lambda_{IJ}$ is invertible. This already fails for two parallel copies with identity weights and twists in $\mathrm{SO}(2) \subset \mathrm{GL}^+(2)$. Consequently adding gauge covariance, invariance under reversal, or a rank requirement on singular links cannot produce such a rule on all singular and rectangular cuts. This obstruction concerns extension of the balanced presentation; the exact rectangular endpoint-map closure of Theorem 10.25 remains available. [ESTABLISHED]

*Proof.* For $-1 < t < 1$, take $W_1 = W_2 = I_2$ and

$$\widehat{\Theta}_1(t) = \begin{pmatrix} t & -\sqrt{1-t^2} \\ \sqrt{1-t^2} & t \end{pmatrix}, \qquad \widehat{\Theta}_2(t) = \widehat{\Theta}_1(t)^\top. \tag{10.32}$$

The twists lie in $\mathrm{SO}(2)$ and depend continuously on $t$, while $\Lambda_{II} = \Lambda_{JJ} = 2I_2$ and $\Lambda_{IJ} = 2tI_2$. For $t \neq 0$, the interval endpoints are $W_{\mathrm{lo}} = 2t^2 I_2$ and $W_{\mathrm{hi}} = 2I_2$, so the balanced weight and link are $W_{\mathrm{bal}} = 2|t|I_2$ and $\Theta_{\mathrm{bal}} = \mathrm{sign}(t)I_2$. The same limiting cut at $t = 0$ is approached from both sides, but the links tend to $I_2$ and $-I_2$. A matrix-valued continuous extension would have to equal both limits. □

## 10.7 The frame sector

**Proposition 10.27 (The frame sector).** Fix a block $I$ whose constituent frames are $u_i = \sigma \cdot U_i^{-1}$ relative to a reference point $\sigma$ (Proposition 2.5), and a coarse datum $(u_I, q_I)$ with $q_I \in \mathcal{B}_x$ and $u_I = \sigma \cdot U_I^{-1}$ a point of the same torsor, so that $u_I = u_i \cdot T_{iI}$ with $T_{iI} = U_i U_I^{-1}$. (i) The constituent expressions $\widehat{\rho}_x(T_{iI})q_I = \widehat{\rho}_x(U_i)\bar{q}_I$ of the coarse law, with $\bar{q}_I = \widehat{\rho}_x(U_I)^{-1}q_I$ the coordinate in $\sigma$ of the point $[u_I, q_I]$ of the associated fiber, depend on $(u_I, q_I)$ only through that point; for the Gaussian law fiber of the defining representation of $\mathrm{GL}^+(K)$ the constituent expression of $q_I = \mathcal{N}(\mu_I, \Sigma_I)$ is $\mathcal{N}(U_i\bar{m}, U_i\bar{S}U_i^\top)$ with $\bar{m} = U_I^{-1}\mu_I$ and $\bar{S} = U_I^{-1}\Sigma_I U_I^{-\top}$. The data with the same constituent expressions form exactly one orbit $\{(u_I \cdot g, \widehat{\rho}_x(g)^{-1}q_I) : g \in G\}$, on which $G$ acts freely, so the coarse frame occupies a copy of $G$, of dimension $K^2$ for $G = \mathrm{GL}^+(K)$, that the constituents do not see. If only the represented frame $\rho_x(U_I)$ is recorded, a nonfaithful representation determines the principal coarse frame only modulo its kernel. (ii) If $G$ contains an element $v \neq e$ with $v^n = e$, no map $F : G^n \to G$ is both left equivariant, $F(gg_1, \ldots, gg_n) = gF(g_1, \ldots, g_n)$, and invariant under permutations of its arguments. For $G = \mathrm{GL}^+(K)$ with $K \geq 2$ and every $n \geq 2$ such an element exists, a rotation of order $n$ in one coordinate plane. The torsion hypothesis cannot be dropped: on the additive group $\mathbb{R}^d$ the arithmetic mean is left equivariant and symmetric. [ESTABLISHED]

*Proof.* (i) By the quotient convention (2.7), $[u_I, q_I] = [u_i \cdot T_{iI}, q_I] = [u_i, \widehat{\rho}_x(T_{iI})q_I]$, so the constituent expression is the coordinate of $[u_I, q_I]$ in the frame $u_i$ and determines that point. Likewise $[u_I, q_I] =$

$[\sigma, \widehat{\rho}_x(U_I)^{-1} q_I] = [u_i \cdot U_i, \bar{q}_I] = [u_i, \widehat{\rho}_x(U_i)\bar{q}_I]$, and the Gaussian pushforward by the linear map $U_i$ gives the Gaussian form. Two data define the same point exactly when they differ by $(u_I, q_I) \mapsto (u_I \cdot g, \widehat{\rho}_x(g)^{-1} q_I)$, and $u_I \cdot g = u_I$ forces $g = e$. Elements of the kernel of $\rho_x$ give equal represented frames. (ii) Left multiplication by $v$ permutes the tuple $(e, v, \dots, v^{n-1})$ cyclically, because $v \cdot v^{n-1} = e$. Symmetry gives $F$ the same value at the permuted tuple, while equivariance multiplies it by $v$; hence $vF = F$ and $v = e$, a contradiction. A rotation by $2\pi/n$ in one coordinate plane lies in $\mathrm{GL}^+(K)$ and has order $n$. The mean on $\mathbb{R}^d$ satisfies $\frac{1}{n}\sum_i (g + g_i) = g + \frac{1}{n}\sum_i g_i$. □

A rule $u_I = \sigma \cdot F(a_1, \dots, a_n)$ with $u_i = \sigma \cdot a_i$ is independent of $\sigma$ exactly when $F$ is left equivariant, because $\sigma' = \sigma \cdot g_0$ replaces every $a_i$ by $g_0^{-1} a_i$; (ii) therefore excludes every symmetric rule of this kind for a meta-agent's frame.

## 10.8 An operational record at $c_*$ sensitive to loop holonomy

**Theorem 10.28 (An operational record sensitive to network-loop holonomy).** Let $p$ be a closed walk of the interaction multigraph based at a probe agent $i_0$ whose belief and model sample fibers are real vector spaces $\mathbb{V}_b, \mathbb{V}_m$ with linear representations $\rho_b, \rho_m$, and write $z_{i_0}(y) \in \mathbb{V}_{\mathrm{pr}} = \mathbb{V}_b \oplus \mathbb{V}_m$ for the probe coordinates of a latent configuration $y$ and $\rho_{\mathrm{pr}} = \rho_b \oplus \rho_m$. Let $\Omega$ be a finite set of labels $\omega$ with prior masses $\psi_\omega > 0$, each prescribing link values on the edges of $p$, with represented loop holonomy $\mathrm{Hol}^{\mathrm{pr}}_\omega = \rho_b(\mathrm{Hol}^b_\omega(p)) \oplus \rho_m(\mathrm{Hol}^m_\omega(p))$. Keep the joint $\mathbb{P}(dy, do)$ independent of $\omega$, with selected posterior version $o \mapsto \mathbf{\Pi}_o$, the slice $\mathbf{\Pi}_{\theta,V,o,X}$ at fixed parameters and structural data; fix $\Sigma_{\mathrm{r}} \succ 0$ on $\mathbb{V}_{\mathrm{pr}}$, and declare the record joint

$$P^{\mathrm{rec}}(\omega, dy, do, dr) = \psi_\omega \mathbb{P}(dy, do) \mathcal{N}\big(r; (\mathrm{Hol}^{\mathrm{pr}}_\omega - I) z_{i_0}(y), \Sigma_{\mathrm{r}}\big) dr.$$

(i) $P^{\mathrm{rec}}$ is a normalized joint. At an admitted $o$ and given $\omega$, the record is distributed as $(\mathrm{Hol}^{\mathrm{pr}}_\omega - I)Z + \epsilon$, with $Z \sim (z_{i_0})_{\#}\mathbf{\Pi}_o$ and $\epsilon \sim \mathcal{N}(0, \Sigma_{\mathrm{r}})$ independent of $Z$. Two labels $\omega, \omega'$ have the same record law exactly when $(\mathrm{Hol}^{\mathrm{pr}}_\omega - I)Z$ and $(\mathrm{Hol}^{\mathrm{pr}}_{\omega'} - I)Z$ have the same law. (ii) If $\mathrm{Hol}^{\mathrm{pr}}_{\omega_0} = I$ and $\mathrm{Hol}^{\mathrm{pr}}_{\omega_1} - I$ is invertible, as for $\mathrm{SO}(2)$ links in both channels whose loop rotation angles are not multiples of $2\pi$, the record laws of $\omega_0$ and $\omega_1$ differ exactly when $(z_{i_0})_{\#}\mathbf{\Pi}_o$ is not the point mass at $0$. Under the same condition the statistic $\mathsf{q} = r^\top \Sigma_{\mathrm{r}}^{-1} r$ has different laws: its mean is $\dim \mathbb{V}_{\mathrm{pr}}$ under $\omega_0$ and $\dim \mathbb{V}_{\mathrm{pr}} + \mathbb{E}\big[Z^\top (\mathrm{Hol}^{\mathrm{pr}}_{\omega_1} - I)^\top \Sigma_{\mathrm{r}}^{-1} (\mathrm{Hol}^{\mathrm{pr}}_{\omega_1} - I) Z\big] \in (\dim \mathbb{V}_{\mathrm{pr}}, +\infty]$ under $\omega_1$. (iii) Under a passive re-expression $g$, with the record expressed in the probe frame, $r \mapsto \rho_{\mathrm{pr}}(g_{i_0}) r$ and $\Sigma_{\mathrm{r}} \mapsto \rho_{\mathrm{pr}}(g_{i_0}) \Sigma_{\mathrm{r}} \rho_{\mathrm{pr}}(g_{i_0})^\top$, the record joint is the image of the original under a bimeasurable bijection; the posterior over $\omega$ given $(o, r)$ and the law of $\mathsf{q}$ are unchanged. [ESTABLISHED]

*Proof.* (i) The Gaussian factor integrates to one in $r$, and $\psi$ and $\mathbb{P}$ are normalized. The label is independent of $(y, o)$, so given $\omega$ and $o$ the configuration has law $\mathbf{\Pi}_o$ and the record kernel is $\mathcal{N}((\mathrm{Hol}^{\mathrm{pr}}_\omega - I) z_{i_0}(y), \Sigma_{\mathrm{r}})$. The characteristic function of the record is $\mathbb{E}[e^{is^\top (\mathrm{Hol}^{\mathrm{pr}}_\omega - I)Z}] e^{-s^\top \Sigma_{\mathrm{r}} s/2}$, whose Gaussian factor never vanishes, so two record laws agree exactly when the laws of the two images of $Z$ agree. (ii) Under $\omega_0$ the image is $0$. Under $\omega_1$ invertibility gives $(\mathrm{Hol}^{\mathrm{pr}}_{\omega_1} - I)Z = 0$ almost surely exactly when $Z = 0$ almost surely; a rotation by $\alpha$ in the plane has $\det(R_\alpha - I) = 2 - 2\cos\alpha$. Given $Z$, the conditional mean of $\mathsf{q}$ is $|\Sigma_{\mathrm{r}}^{-1/2}(\mathrm{Hol}^{\mathrm{pr}}_\omega - I)Z|^2 + \dim \mathbb{V}_{\mathrm{pr}}$, and Tonelli's theorem gives

the two means; the second exceeds the first exactly when the image of $Z$ is nonzero with positive probability. (iii) The loop product telescopes under $\Theta_e^x \mapsto g_i \Theta_e^x g_j^{-1}$, so with $U = \rho_{\mathrm{pr}}(g_{i_0})$ one has $\mathrm{Hol}_\omega^{\mathrm{pr}} \mapsto U \mathrm{Hol}_\omega^{\mathrm{pr}} U^{-1}$ and $z_{i_0} \mapsto U z_{i_0}$, hence the new mean is $U(\mathrm{Hol}_\omega^{\mathrm{pr}} - I) z_{i_0}$ and the new record kernel is the image of the old one under $r \mapsto Ur$. The joint is thus the image under the bijection that changes the latent coordinates and multiplies $r$ by $U$, which leaves the posterior over $\omega$ unchanged, and $r^\top U^\top (U \Sigma_{\mathrm{r}} U^\top)^{-1} U r = r^\top \Sigma_{\mathrm{r}}^{-1} r$. □

**Physicist's summary.** A block of agents at one base point is a small lattice gauge system whose invariants are loop products and edge twists modulo one global conjugation. A coarse description keeps a root frame, the tree transports to it, the holonomy of every independent loop, and the tree-dressed boundary links; the tree is a gauge fixing, as in lattice gauge theory, but averaged member features depend on it whenever the holonomy acts on them. The configurations that survive exact aggregation form the kernel of the connection Laplacian, for positive-definite couplings the holonomy-fixed vectors, and the members' marginal laws have one parent exactly when their transported laws coincide in a holonomy-invariant law, so a planar rotation loop forces a zero parent mean and a volume-stretching loop allows no Gaussian parent. Links between blocks compress exactly into rectangular endpoint maps, and into one link with new self-energies when an explicit matrix is invertible, uniquely once the self-energies are required to balance across the link; at the level of the laws themselves no such balanced form exists once the copies disagree. The parent frame is pure gauge for the constituents, and a noisy record of how the loop rotates the probe's state distinguishes holonomies unless that state is pinned at zero; Supplement B computes that holonomy from a connection along a base loop.

# Chapter 11

# Exact Agent-Network Coarse-Graining

This chapter asks what one exact coarse-graining step does to the interacting agent network of Chapter 5 at $c_*$, and which state must be kept so that the step closes. The effective action of a coarse agent network is a conditional partition function (Theorem 11.3), real analytic in bounded perturbations (Theorem 11.4); in Hoeffding interaction coordinates the step is an exact nonlinear map between spaces of different dimension whose finite-step multipliers are bounded by $3^{|V_n|}-1$ at every depth (Proposition 11.19); and exact closure keeps induced hyperedges, bridge kernels, joint attention-event laws, gauge-retained marks, and path laws with their memory (Table 11.1). The multiscale free energy of Chapter 12 is built from these objects, and the formation witnesses of Chapter 13 are exact marked parents in the sense fixed here.

At depth $\ell$ a coarse channel $C_\ell : \mathsf{Y}_\ell \rightsquigarrow \mathsf{Y}_{\ell+1}$ (Definition 9.1) maps the product state space of the finite agent set $V_\ell$ to the next depth; for one arrow the spaces are Y and Z with samples $y$ and $z$. Local symbols are $\varrho_\ell, \mathsf{m}_\ell, M_\ell, \bar{\mathsf{m}}_\ell, H_\ell$ for the measure pair; $C^\dagger, \mathsf{U}, \mathcal{A}$ for the reverse kernel, its conditional expectation, and the action increment; $\mathcal{G}_\ell, \mathsf{g}, \mathsf{E}_{\ell,B}, \mathsf{P}_{\ell,B}, \mathrm{asm}_\ell, \mathrm{hd}_\ell$ for interaction coordinates; $\mathcal{T}_\ell, D_\ell, D_{n\leftarrow\ell}$ for the exact step and its multipliers; $\chi$ and T for a deterministic coarse map and a transition kernel; and a superscript $c$ for coarse objects. Longer proofs are deferred to appendix B.

Scope. The retained-state table of Section 11.9 catalogs exact operations under their admission hypotheses and is not a completeness theorem; the coarse channel and the child blocks are inputs, not canonical selections; and the auxiliary path experiment of Section 11.8.1 has its own evidence and free energy, not identified with those of the original sensor model. [NOT-CLAIMED] Open items are the invariant retained subspaces of the interaction step (Open Problem 11.25), the dynamic closure of the pointwise parent (Open Problem 11.38), and jointly measurable equivariant conditional versions beyond the compact and finite cases (Open Problem 9.18).

## 11.1 Exact effective action

**Definition 11.1 (Scale-indexed measure pair).** For each depth $\ell$ let $\varrho_\ell$ be a probability measure on $\mathsf{Y}_\ell$, let $\mathsf{m}_\ell \ll \varrho_\ell$ be a finite positive measure, and let a depth arrow be a parameter-independent

Markov kernel $C_\ell : \mathsf{Y}_\ell \rightsquigarrow \mathsf{Y}_{\ell+1}$ acting on both components:

$$\mathsf{m}_\ell(dy) = \exp[-H_\ell(y)]\varrho_\ell(dy), \qquad 0 < M_\ell := \mathsf{m}_\ell(\mathsf{Y}_\ell) < \infty, \qquad \bar{\mathsf{m}}_\ell = \mathsf{m}_\ell/M_\ell, \tag{11.1}$$

$$\varrho_{\ell+1} = \varrho_\ell C_\ell, \qquad \mathsf{m}_{\ell+1} = \mathsf{m}_\ell C_\ell, \qquad H_{\ell+1} = -\log \frac{d\mathsf{m}_{\ell+1}}{d\varrho_{\ell+1}} \quad \varrho_{\ell+1}\text{-almost everywhere}, \tag{11.2}$$

with $\exp(-\infty) = 0$, so that each action is an extended real function defined almost everywhere. At a record $o$ and for one arrow $C : \mathsf{Y} \rightsquigarrow \mathsf{Z}$ with its index suppressed, the fine pair is $(\varrho, \mathsf{m}_o)$ with $\mathsf{m}_o(dy) = L_o(y)\varrho(dy)$, $L_o = \exp(-H_o)$, $0 < M_o := \mathsf{m}_o(\mathsf{Y}) < \infty$, and $\bar{\mathsf{m}}_o = M_o^{-1}\mathsf{m}_o$, and the pushed pair, effective likelihood, and effective action are

$$\varrho^c = \varrho C, \quad \mathsf{m}_o^c = \mathsf{m}_o C, \quad \bar{\mathsf{m}}_o^c = \bar{\mathsf{m}}_o C = M_o^{-1}\mathsf{m}_o^c, \quad L_o^c = \frac{d\mathsf{m}_o^c}{d\varrho^c}, \quad H_o^c = -\log L_o^c. \tag{11.3}$$

When $\varrho$ is the latent prior and $L_o$ the record likelihood, $\bar{\mathsf{m}}_o$ is the posterior slice and $M_o$ the evidence.

**Proposition 11.2 (Exact composition and normalized law).** Composable arrows act on the pair componentwise: $(\varrho_\ell, \mathsf{m}_\ell)(C_\ell C_{\ell+1}) = ((\varrho_\ell C_\ell)C_{\ell+1}, (\mathsf{m}_\ell C_\ell)C_{\ell+1})$. Moreover $\mathsf{m}_{\ell+1} \ll \varrho_{\ell+1}$, so that (11.2) is defined, $\mathsf{m}_{\ell+1}(\mathsf{Y}_{\ell+1}) = M_\ell$, and $\bar{\mathsf{m}}_{\ell+1} = \bar{\mathsf{m}}_\ell C_\ell$; the same statements hold for the single-arrow pair (11.3). [ESTABLISHED]

*Proof.* Composition on each component is Proposition 9.2. If $\varrho_{\ell+1}(A) = 0$, then $C_\ell(y, A) = 0$ for $\varrho_\ell$-almost every $y$, hence for $\mathsf{m}_\ell$-almost every $y$, and $\mathsf{m}_{\ell+1}(A) = 0$. A Markov kernel preserves total mass, and dividing by that mass gives the normalized law. □

**Theorem 11.3 (Conditional-partition formula and composition).** Let $C_\varrho^\dagger(z, dy)$ be a reverse conditional of $Y$ given $Z = z$ under the joint law $\varrho(dy)C(y, dz)$. Then for $\varrho^c$-almost every $z$

$$\exp[-H_o^c(z)] = \int \exp[-H_o(y)]C_\varrho^\dagger(z, dy), \qquad M_o = \int \exp[-H_o^c(z)]\varrho^c(dz). \tag{11.4}$$

For a second channel $C'$ on Z, the direct effective likelihood of $CC'$ and the staged one agree $\varrho CC'$-almost everywhere when their reverse conditionals are chosen from one joint tower. If $0 \le L_o \le 1$, as for a binary success record, then $0 \le L_o^c \le 1$, so one effective binary record realizes the whole coarse action; for a general record density $L_o^c$ need not be bounded by one. [ESTABLISHED]

*Proof.* For bounded measurable $f$, $\int f(z)\mathsf{m}_o^c(dz) = \mathbb{E}_\varrho[f(Z)L_o(Y)] = \int f(z)\mathbb{E}_\varrho[L_o(Y) \mid Z = z]\varrho^c(dz)$ under the joint law, which identifies the Radon–Nikodym derivative, as in the density formula of Theorem 9.3 with $L_o = d\mathsf{m}_o/d\varrho$; $f = 1$ gives the evidence. The tower property proves composition, and a conditional expectation of a function with values in $[0, 1]$ takes values in $[0, 1]$. □

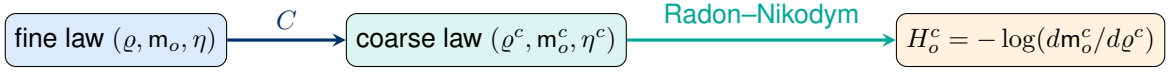



**Figure 11.1:** One exact step acts on the reference probability, the evidence measure, and the joint attention-event law; the effective action is the Radon–Nikodym derivative of the pushed evidence measure with respect to the pushed reference.

## 11.2 Bounded action calculus and contraction

**Theorem 11.4 (Exact bounded action calculus).** Fix a jointly measurable reverse kernel $C^\dagger(z, dy)$ of the normalized joint law and its reverse conditional expectation,

$$\bar{\mathsf{m}}_o(dy)C(y, dz) = \bar{\mathsf{m}}_o^c(dz)C^\dagger(z, dy), \qquad (\mathsf{U}h)(z) = \int h(y)C^\dagger(z, dy), \tag{11.5}$$

whose values on $\bar{\mathsf{m}}_o^c$-null sets are not used. A bounded $\psi$ perturbs the evidence measure to $e^{-\psi}\mathsf{m}_o$; on $\mathfrak{B} = L^\infty(\bar{\mathsf{m}}_o; \mathbb{R})$ the exact coarse action increment $\mathcal{A} : \mathfrak{B} \to L^\infty(\bar{\mathsf{m}}_o^c; \mathbb{R})$ is

$$\mathcal{A}(\psi)(z) = -\log \frac{d\big((e^{-\psi}\mathsf{m}_o)C\big)}{d(\mathsf{m}_o C)}(z) = -\log \int e^{-\psi(y)}C^\dagger(z, dy), \tag{11.6}$$

the difference of the coarse actions of $\mathsf{m}_o$ and $e^{-\psi}\mathsf{m}_o$ where both are finite; elsewhere the calculus supplies no action difference and the Radon–Nikodym ratio is the definition.

(a) $e^{-\|\psi\|_\infty} \le \mathsf{U}(e^{-\psi}) \le e^{\|\psi\|_\infty}$, so $\|\mathcal{A}(\psi)\|_\infty \le \|\psi\|_\infty$, and $\mathcal{A}$ is real analytic on the whole of $\mathfrak{B}$.

(b) For $\psi \in \mathfrak{B}$ the tilted reverse kernel and its operator

$$C_\psi^\dagger(z, dy) = \frac{e^{-\psi(y)}C^\dagger(z, dy)}{\int e^{-\psi(u)}C^\dagger(z, du)}, \qquad \mathsf{U}^\psi h = \frac{\mathsf{U}(e^{-\psi}h)}{\mathsf{U}(e^{-\psi})}, \tag{11.7}$$

are positive, unital, and bounded from $L^\infty(\bar{\mathsf{m}}_o)$ to $L^\infty(\bar{\mathsf{m}}_o^c)$, and for bounded $h$ and $k$

$$\mathcal{A}(\psi + h) - \mathcal{A}(\psi) = -\log \mathsf{U}^\psi(e^{-h}), \tag{11.8}$$

$$D\mathcal{A}(\psi)[h](z) = \int h(y)C_\psi^\dagger(z, dy) = (\mathsf{U}^\psi h)(z), \tag{11.9}$$

$$D^2\mathcal{A}(\psi)[h, k](z) = -\operatorname{Cov}_{C_\psi^\dagger(z,\cdot)}(h, k), \tag{11.10}$$

which at the origin read $D\mathcal{A}(0)[h] = \mathsf{U}h$ and $D^2\mathcal{A}(0)[h, k] = -\operatorname{Cov}_{C^\dagger(z,\cdot)}(h, k)$.

(c) The map is order preserving and additively homogeneous: $\psi \le \psi' + c$ implies $\mathcal{A}(\psi) \le \mathcal{A}(\psi') + c$, and $\mathcal{A}(\psi + c) = \mathcal{A}(\psi) + c$ for constants $c$. Consequently

$$\|\mathcal{A}(\psi) - \mathcal{A}(\psi')\|_\infty \le \|\psi - \psi'\|_\infty. \tag{11.11}$$

(d) Both $\mathcal{A}$ and $\mathsf{U}$ descend to the projective action space $\overline{\mathfrak{B}} = \mathfrak{B}/\mathbb{R}1$, which forgets the perturbed mass $(e^{-\psi}\mathsf{m}_o)(\mathsf{Y})$.

(e) Normalizing the perturbed fine law instead gives

$$\widehat{\mathsf{m}}^\psi = \frac{e^{-\psi}\bar{\mathsf{m}}_o}{\bar{\mathsf{m}}_o(e^{-\psi})}, \qquad \widetilde{\mathcal{A}}(\psi) := -\log \frac{d(\widehat{\mathsf{m}}^\psi C)}{d\bar{\mathsf{m}}_o^c} = \mathcal{A}(\psi) + \log \bar{\mathsf{m}}_o(e^{-\psi}), \tag{11.12}$$

which differs from $\mathcal{A}$ by a constant, with derivatives $D\widetilde{\mathcal{A}}(0)[h] = \mathsf{U}h - \bar{\mathsf{m}}_o(h)$ and $D^2\widetilde{\mathcal{A}}(0)[h,k] = -\operatorname{Cov}_{C^\dagger(z,\cdot)}(h,k) + \operatorname{Cov}_{\bar{\mathsf{m}}_o}(h,k)$. [ESTABLISHED]

**Theorem 11.5 (Conditional-expectation contraction and exact defect).** For every $1 \le p \le \infty$ the operator $\mathsf{U}$ of (11.5) extends uniquely to a mean-preserving contraction $\mathsf{U} : L^p(\bar{\mathsf{m}}_o;\mathbb{R}) \to L^p(\bar{\mathsf{m}}_o^c;\mathbb{R})$, and at $p = 2$ the loss is a conditional variance:

$$\|\mathsf{U}h\|_{L^p(\bar{\mathsf{m}}_o^c)} \le \|h\|_{L^p(\bar{\mathsf{m}}_o)}, \tag{11.13}$$

$$\|h\|^2_{L^2(\bar{\mathsf{m}}_o)} - \|\mathsf{U}h\|^2_{L^2(\bar{\mathsf{m}}_o^c)} = \int \operatorname{Var}_{C^\dagger(z,\cdot)}(h)\bar{\mathsf{m}}_o^c(dz) \ge 0, \tag{11.14}$$

with equality exactly when $h(y) = f(z)$ for $\bar{\mathsf{m}}_o(dy)C(y,dz)$-almost every $(y,z)$ and some measurable $f$ on Z; for a deterministic coarse map $z = \chi(y)$ this means $h = f \circ \chi$ $\bar{\mathsf{m}}_o$-almost everywhere. With the oscillation norm on $L^\infty(\bar{\mathsf{m}}_o)/\mathbb{R}1$ and the contraction coefficient of Dobrushin [27] of the reverse kernel,

$$\|[h]\|_{\mathrm{osc}} = \inf_{c\in\mathbb{R}} \|h - c\|_\infty = \tfrac{1}{2}\big(\operatorname{ess\,sup} h - \operatorname{ess\,inf} h\big), \tag{11.15}$$

$$\delta = \operatorname*{ess\,sup}_{(z,z')\sim(\bar{\mathsf{m}}_o^c)^{\otimes 2}} \sup_A \big|C^\dagger(z,A) - C^\dagger(z',A)\big|, \tag{11.16}$$

the inner supremum running over Borel sets, one has $\|[\mathsf{U}h]\|_{\mathrm{osc}} \le \delta\|[h]\|_{\mathrm{osc}}$, and along a finite composition of arrows the oscillation norm of the composed reverse operators is at most the product of the one-step coefficients. [ESTABLISHED]

The proofs of Theorems 11.4 and 11.5 are in appendix B.4.

## 11.3 Finite-network interaction coordinates

**Definition 11.6 (Product-referenced network scale).** At depth $\ell$ agent $i \in V_\ell$ has state space $\mathsf{Y}_{\ell i}$ and a probability law $\nu_{\ell i}$, and the normalized law $\bar{\mathsf{m}}_\ell$ of Definition 11.1 on $\mathsf{Y}_\ell = \prod_{i\in V_\ell} \mathsf{Y}_{\ell i}$ is mutually absolutely continuous, $\bar{\mathsf{m}}_\ell \sim \nu_\ell$, with the product reference $\nu_\ell = \bigotimes_{i\in V_\ell} \nu_{\ell i}$. Equivalence identifies $L^\infty(\bar{\mathsf{m}}_\ell)$ with $L^\infty(\nu_\ell)$ isometrically, with no bound on the Radon–Nikodym ratio, and the product structure of $\nu_\ell$ makes the coordinate expectations below well defined on equivalence classes by Fubini's theorem. For an arrow $C_\ell$, equivalence of $\bar{\mathsf{m}}_{\ell+1} = \bar{\mathsf{m}}_\ell C_\ell$ to a product reference $\nu_{\ell+1}$ is a separate premise (Proposition 11.8). If $V_\ell = \varnothing$, then $\mathsf{Y}_\ell$ is a point, $\mathcal{G}_\ell = \{0\}$, and every map below is zero; estimates involving $3^{|V_\ell|} - 1$ assume $V_\ell \neq \varnothing$.

**Remark 11.7 (The bounded tier).** Every interaction component below is essentially bounded. The bounded tier contains every interaction on finite agent spaces and every continuous interaction on compact ones, hence the natural couplings of categorical agents and of compact group-valued states. On $\mathsf{Y}_{\ell i} = \mathbb{R}^K$ it excludes the polynomial couplings: for a matrix $\mathsf{N} \neq 0$ and every $M > 0$ the set where $|y_i^\top \mathsf{N} y_j| > M$ is a nonempty open subset of $\mathbb{R}^K \times \mathbb{R}^K$, so it has positive mass under every reference with full support, and $y_i^\top \mathsf{N} y_j \notin L^\infty(\nu_\ell)$. Such couplings need a separate tier,

exact Gaussian marginalization or a weighted action space, which this chapter does not construct. [ESTABLISHED]

**Proposition 11.8 (Product equivalence is an admitted, not an automatic, scale premise).** Let $\mathsf{Y}_0 = \{0,1\}$ with its uniform law $\bar{\mathsf{m}}_0 = \nu_0$, and let $C$ be the deterministic diagonal-cloning channel $y \mapsto (y,y)$ from $\{0,1\}$ to $\{0,1\}^2$. Then $\bar{\mathsf{m}}_1 = \bar{\mathsf{m}}_0 C$ is supported on $\{(0,0),(1,1)\}$, no product probability $\nu_{11} \otimes \nu_{12}$ on $\{0,1\}^2$ is equivalent to it, and no target Hoeffding decomposition relative to a product reference is admitted at this target. [ESTABLISHED]

*Proof.* Equivalence would make both diagonal atoms positive under the product, so each marginal would charge both $0$ and $1$ and both off-diagonal atoms would have positive product mass, while $\bar{\mathsf{m}}_1$ assigns them mass zero. □

**Definition 11.9 (Hoeffding coordinates).** Under Definition 11.6, for $B \subseteq V_\ell$ let $(\mathbb{E}_{\ell,B} f)(y_B) = \int f(y_B, y'_{B^c}) \nu_{\ell,B^c}(dy'_{B^c})$ integrate out the complementary coordinates against the product reference, with $\mathbb{E}_{\ell,\varnothing} f = \nu_\ell(f)$. The Hoeffding projectors are the Möbius inversions

$$\mathsf{P}_{\ell,B} = \sum_{B' \subseteq B} (-1)^{|B|-|B'|} \mathbb{E}_{\ell,B'} = \prod_{i \in B} \big(I - \mathbb{E}_{\ell,V_\ell \setminus \{i\}}\big) \mathbb{E}_{\ell,B}, \qquad B \subseteq V_\ell, \tag{11.17}$$

where the alternating sum is the definition and the product is the identity of Theorem 11.10. The component spaces $\mathcal{H}_{\ell,B} = \mathsf{P}_{\ell,B} L^\infty(\nu_\ell)$ form the interaction space $\mathcal{G}_\ell = \bigoplus_{\varnothing \neq B \subseteq V_\ell}^{\ell^1} \mathcal{H}_{\ell,B}$ with norm $\|\mathsf{g}\|_{\mathcal{G}_\ell} = \sum_{\varnothing \neq B \subseteq V_\ell} \|\mathsf{g}_B\|_\infty$, so every nonempty hyperedge has its component and no bound on $|V_\ell|$ is imposed. With the action quotient $\overline{\mathfrak{B}}_\ell = L^\infty(\nu_\ell)/\mathbb{R}1$, normed by (11.15), the assembly and extraction maps are

$$\psi_{\mathsf{g}} = \sum_{\varnothing \neq B \subseteq V_\ell} \mathsf{g}_B, \qquad \mathrm{asm}_\ell\, \mathsf{g} = [\psi_{\mathsf{g}}], \qquad \mathrm{hd}_\ell[f] = (\mathsf{P}_{\ell,B} f)_{\varnothing \neq B \subseteq V_\ell}. \tag{11.18}$$

**Theorem 11.10 (Hoeffding assembly is an exact finite-network action isomorphism).** For every finite $V_\ell$ the coordinate expectations commute and satisfy $\mathbb{E}_{\ell,B}\mathbb{E}_{\ell,B'} = \mathbb{E}_{\ell,B\cap B'}$, the product formula in (11.17) holds, and $\mathsf{P}_{\ell,B}\mathsf{P}_{\ell,B'} = \mathbf{1}_{B=B'}\mathsf{P}_{\ell,B}$ with $\sum_{B \subseteq V_\ell} \mathsf{P}_{\ell,B} = I$. Each component $\mathsf{g}_B$ depends only on $y_B$ and has zero $\nu_\ell$-mean in every coordinate of $B$. The maps of (11.18) are mutually inverse bounded linear isomorphisms,

$$\mathrm{hd}_\ell \,\mathrm{asm}_\ell = I_{\mathcal{G}_\ell}, \qquad \mathrm{asm}_\ell \,\mathrm{hd}_\ell = I_{\overline{\mathfrak{B}}_\ell}, \qquad \|\mathrm{asm}_\ell\| \le 1, \qquad \|\mathrm{hd}_\ell\| \le 3^{|V_\ell|} - 1. \tag{11.19}$$

The last bound is not dimension free, and its exponential dependence is sharp: for independent $\{-1,1\}$-valued coordinates with $\Pr(Y_i = 1) = p$ with $1/2 \le p < 1$ and $f = \prod_i Y_i$ on $n = |V_\ell|$ agents, the sum of the component norms is $(4p-1)^n - (2p-1)^n$, which tends to $3^n - 1$ as $p \uparrow 1$. [ESTABLISHED]

The proof is in appendix B.4.

**Corollary 11.11 (Gauge covariance of interaction components).** Let a passive change $g = (g_i)$ of Chapter 2, one element per agent acting on both the belief and the model components of that agent's state, be realized at depth $\ell$ by coordinate Borel isomorphisms and a Borel permutation $\varsigma$ of the agents, whose product map $\kappa_\ell : \mathsf{Y}_\ell \to \widehat{\mathsf{Y}}_\ell$ sends $\nu_\ell$ to a product reference $\widehat{\nu}_\ell$. With $\mathsf{K}_\ell f = f \circ \kappa_\ell^{-1}$, its action $\overline{\mathsf{K}}_\ell$ on the quotient by constants, and its action $\mathsf{K}_\ell^{\mathcal{G}}$ on components,

$$\mathsf{K}_\ell \mathbb{E}^{\nu}_{\ell,B} = \mathbb{E}^{\widehat{\nu}}_{\ell,\varsigma B} \mathsf{K}_\ell, \qquad \mathsf{K}_\ell \mathsf{P}^{\nu}_{\ell,B} = \mathsf{P}^{\widehat{\nu}}_{\ell,\varsigma B} \mathsf{K}_\ell, \qquad \overline{\mathsf{K}}_\ell \operatorname{asm}_\ell = \widehat{\operatorname{asm}}_\ell \mathsf{K}_\ell^{\mathcal{G}}, \qquad \mathsf{K}_\ell^{\mathcal{G}} \operatorname{hd}_\ell = \widehat{\operatorname{hd}}_\ell \overline{\mathsf{K}}_\ell. \tag{11.20}$$

Preservation of the product reference alone is insufficient: the product-Haar-preserving shear $(y_1, y_2) \mapsto (y_1, y_1 + y_2)$ of the torus $\mathbb{T}^2$ sends a function of $y_2$ to one depending on both coordinates and moves a singleton component into a pair component. If only almost-everywhere reverse-kernel equivariance is supplied, the covariance holds for $L^\infty$ equivalence classes and not for a selected pointwise version. [ESTABLISHED]

*Proof.* Fubini's theorem proves the first identity, and the others follow by the finite Möbius sums. The shear maps $f(y_2)$ to $f(y_2 - y_1)$, which is not a function of one coordinate when $f$ is nonconstant. □

## 11.4 The exact interaction step, its residual, and finite-step multipliers

**Definition 11.12 (Exact interaction step).** Assume Definition 11.6 at depths $\ell$ and $\ell+1$, let $\mathcal{A}_\ell$ be the action increment of Theorem 11.4 for the arrow $C_\ell$ at the base law $\bar{\mathsf{m}}_\ell$, and let $\overline{\mathcal{A}}_\ell : \overline{\mathfrak{B}}_\ell \to \overline{\mathfrak{B}}_{\ell+1}$ be its quotient map, defined on the whole bounded space by additive homogeneity. The exact interaction step is

$$\mathcal{T}_\ell = \operatorname{hd}_{\ell+1} \overline{\mathcal{A}}_\ell \operatorname{asm}_\ell : \mathcal{G}_\ell \longrightarrow \mathcal{G}_{\ell+1}, \qquad \mathsf{g}^{\mathrm{ex}}_{\ell+1} = \mathcal{T}_\ell(\mathsf{g}_\ell). \tag{11.21}$$

**Lemma 11.13 (Mass, origin, and derivative of the exact step).** The sum $\psi_{\mathsf{g}}$ is the representative of $\operatorname{asm}_\ell \mathsf{g}$ with zero $\nu_\ell$-mean, and the perturbed evidence mass is carried separately and preserved by the arrow,

$$M_\ell(\mathsf{g}) := M_\ell \bar{\mathsf{m}}_\ell\big(e^{-\psi_{\mathsf{g}}}\big) = \big((e^{-\psi_{\mathsf{g}}} \mathsf{m}_\ell) C_\ell\big)(\mathsf{Y}_{\ell+1}) =: M_{\ell+1}(\mathsf{g}). \tag{11.22}$$

The step fixes the origin, $\mathcal{T}_\ell(0) = 0$, for every model and every channel. The tilted operator $\mathsf{U}_\ell^{\psi_{\mathsf{g}}}$ of (11.7) is unchanged when a constant is added to $\psi_{\mathsf{g}}$, descends to the action quotient, and gives the derivative

$$D\mathcal{T}_\ell(\mathsf{g}) = \operatorname{hd}_{\ell+1} \overline{\mathsf{U}_\ell^{\psi_{\mathsf{g}}}} \operatorname{asm}_\ell : \mathcal{G}_\ell \longrightarrow \mathcal{G}_{\ell+1}; \tag{11.23}$$

only at $\mathsf{g} = 0$ may the tilted operator be replaced by $\mathsf{U}_\ell$. [ESTABLISHED]

*Proof.* A Markov kernel preserves total mass, which gives (11.22), and $\mathcal{A}_\ell(0) = -\log \mathsf{U}_\ell(1) = 0$. A constant added to $\psi_{\mathsf{g}}$ multiplies numerator and denominator of $\mathsf{U}_\ell^{\psi_{\mathsf{g}}} h$ by one positive number, and the chain rule with part (b) of Theorem 11.4 gives (11.23). □

**Remark 11.14 (Two agents and one exact step).** Let $V_0 = \{1, 2\}$ with $\mathsf{Y}_{0i} = \{-1, 1\}$ and uniform product base $\bar{\mathsf{m}}_0 = \nu_0$, and perturb by $\psi(y) = -\xi_1 y_1 - \xi_2 y_2 - \xi_{12} y_1 y_2$, whose Hoeffding components are $-\xi_1 y_1$, $-\xi_2 y_2$, and $-\xi_{12} y_1 y_2$, so $\dim \mathcal{G}_0 = 3$. The channel $C_0(y, dz) = \delta_{y_1}(dz)$ keeps the first agent, so $V_1 = \{I\}$, $\mathsf{Y}_1 = \{-1, 1\}$, the pushed base is uniform, and $\dim \mathcal{G}_1 = 1$. The reverse kernel sets $y_1 = z$ and leaves $y_2$ uniform, so $\mathcal{A}_0(\psi)(z) = -\xi_1 z - \log\cosh(\xi_2 + \xi_{12} z)$, and the exact step is $\mathcal{T}_0(\mathsf{g})_{\{I\}}(z) = -\xi_1' z$ with

$$\xi_1' = \xi_1 + \tfrac{1}{2} \log \frac{\cosh(\xi_2 + \xi_{12})}{\cosh(\xi_2 - \xi_{12})} = \xi_1 + \operatorname{artanh}\big(\tanh \xi_2 \tanh \xi_{12}\big).$$

The pair coupling enters the coarse field only through the polarization of the eliminated agent, $\partial \xi_1' / \partial \xi_{12} = \tanh \xi_2$ at $\xi_{12} = 0$. At the origin $D\mathcal{T}_0(0)$ keeps the field of agent $1$ and sends the field of agent $2$ and the pair direction to zero, since $\mathsf{U}_0(y_2) = \mathsf{U}_0(y_1 y_2) = 0$; at a point with $\xi_2 \neq 0$ the pair direction has the nonzero image $-\tanh(\xi_2) z$. The drop from dimension $3$ to $1$ is the obstruction of Proposition 11.20, and the multipliers obey the bound $3^2 - 1 = 8$ of Proposition 11.19. [ESTABLISHED]

*Proof.* Each monomial has zero mean in each of its coordinates, which gives the components by (11.17). Averaging $e^{-\psi}$ over $y_2$ at $y_1 = z$ gives $e^{\xi_1 z} \cosh(\xi_2 + \xi_{12} z)$, and subtracting the mean of $\mathcal{A}_0(\psi)$ over $z = \pm 1$ gives the coefficient $-\xi_1'$; the artanh form follows from the addition formula for $\cosh$. Differentiating at $\xi_{12} = 0$ gives $\tanh \xi_2$, and (11.23) at $\psi = -\xi_2 y_2$ gives $\mathsf{U}_0^{\psi}(y_1 y_2)(z) = z \tanh \xi_2$. □

**Proposition 11.15 (A product base and a block-product channel preserve the intra-block sector).** Suppose that $\bar{\mathsf{m}}_\ell = \nu_\ell$ is a product law, that blocks $I \in V_{\ell+1}$ partition $V_\ell$, each identified with its coarse agent, and that the channel factors as

$$C_\ell(y, dz) = \prod_{I \in V_{\ell+1}} C_I(y_I, dz_I) \tag{11.24}$$

with Markov kernels $C_I$. Then $\bar{\mathsf{m}}_{\ell+1} = \bar{\mathsf{m}}_\ell C_\ell$ is a product law, so the convention $\bar{\mathsf{m}}_{\ell+1} = \nu_{\ell+1}$ is admissible at the next depth. Let $\mathcal{G}_\ell^{\mathrm{intra}} = \{\mathsf{g} \in \mathcal{G}_\ell : \mathsf{g}_B = 0 \text{ unless } B \subseteq I \text{ for some } I\}$. For every $\mathsf{g} \in \mathcal{G}_\ell^{\mathrm{intra}}$ and every product reference $\nu_{\ell+1} \sim \bar{\mathsf{m}}_{\ell+1}$, the components of $\mathcal{T}_\ell(\mathsf{g})$ vanish on every $B \subseteq V_{\ell+1}$ with $|B| \geq 2$; in particular the singleton sector $\{\mathsf{g} : \mathsf{g}_B = 0 \text{ for } |B| \geq 2\}$, which lies in $\mathcal{G}_\ell^{\mathrm{intra}}$, is mapped into the singleton sector of the next depth. [ESTABLISHED]

The proof is in appendix B.4.

**Definition 11.16 (Scale-indexed evolution and its derivative cocycle).** Let $(\mathscr{X}_\ell, \|\cdot\|_\ell)_{\ell \geq 0}$ be Banach spaces and let $\mathcal{T}_\ell : \mathscr{O}_\ell \to \mathscr{X}_{\ell+1}$ be continuously Fréchet differentiable on an open set $\mathscr{O}_\ell \subseteq \mathscr{X}_\ell$. The ordered evolution $\mathcal{T}_{n \leftarrow \ell} = \mathcal{T}_{n-1} \circ \cdots \circ \mathcal{T}_\ell$, with $\mathcal{T}_{\ell \leftarrow \ell} = \mathrm{id}_{\mathscr{X}_\ell}$, is defined at the points whose successive images stay in the successive domains. Along an exact orbit $\xi_{\ell+1} = \mathcal{T}_\ell(\xi_\ell)$ in those domains, the finite-step multipliers are

$$D_\ell = D\mathcal{T}_\ell(\xi_\ell) \in \mathcal{L}(\mathscr{X}_\ell, \mathscr{X}_{\ell+1}), \qquad D_{n \leftarrow \ell} = D_{n-1} \cdots D_\ell, \qquad D_{\ell \leftarrow \ell} = I_{\mathscr{X}_\ell}, \tag{11.25}$$

with the rightmost factor acting first.

**Proposition 11.17 (Ordered composition of the evolution and its derivative).** For $\ell \le n \le r$ and an exact orbit as in Definition 11.16,

$$\mathcal{T}_{r\leftarrow n} \circ \mathcal{T}_{n\leftarrow \ell} = \mathcal{T}_{r\leftarrow \ell}, \qquad D_{n\leftarrow \ell} = D\mathcal{T}_{n\leftarrow \ell}(\xi_\ell), \qquad D_{r\leftarrow n} D_{n\leftarrow \ell} = D_{r\leftarrow \ell}. \tag{11.26}$$

[ESTABLISHED]

*Proof.* The first and third identities are associativity. The second follows by induction on $n-\ell$ from the Fréchet chain rule applied to $\mathcal{T}_{n+1\leftarrow \ell} = \mathcal{T}_n \circ \mathcal{T}_{n\leftarrow \ell}$ at $\xi_\ell$. □

**Definition 11.18 (Finite-step multipliers of the interaction step).** Instantiate Definition 11.16 with $\mathscr{X}_\ell = \mathcal{G}_\ell$, $\mathcal{T}_\ell$ the exact interaction step, and an exact interaction orbit $\mathrm{g}_{\ell+1} = \mathcal{T}_\ell(\mathrm{g}_\ell)$. By Lemma 11.13 the one-step multipliers are

$$D_\ell = D\mathcal{T}_\ell(\mathrm{g}_\ell) = \mathrm{hd}_{\ell+1}\, \overline{\mathsf{U}_\ell^{\psi_{\mathrm{g}_\ell}}}\, \mathrm{asm}_\ell : \mathcal{G}_\ell \longrightarrow \mathcal{G}_{\ell+1}, \tag{11.27}$$

and $D_{n\leftarrow \ell}$ are their ordered products over the finite depth.

**Proposition 11.19 (Native growth bound for the exact interaction cocycle).** Equip every action quotient $\overline{\mathfrak{B}}_k$ with the oscillation norm (11.15). Along an exact interaction orbit with multipliers (11.27) and for $n > \ell$, the ordered product telescopes, and its norm is bounded by the target size:

$$D_{n\leftarrow \ell} = \mathrm{hd}_n\, \overline{\mathsf{U}_{n-1}^{\psi_{\mathrm{g}_{n-1}}}} \cdots \overline{\mathsf{U}_\ell^{\psi_{\mathrm{g}_\ell}}}\, \mathrm{asm}_\ell, \tag{11.28}$$

$$\|D_{n\leftarrow \ell}\|_{\mathcal{G}_\ell \to \mathcal{G}_n} \le 3^{|V_n|} - 1. \tag{11.29}$$

The bound does not depend on the number of steps, and it is zero when $V_n = \varnothing$. [ESTABLISHED]

*Proof.* By (11.19), $\mathrm{asm}_{k+1}\,\mathrm{hd}_{k+1} = I$ on $\overline{\mathfrak{B}}_{k+1}$, so every interior pair cancels in $D_{n-1}\cdots D_\ell$, which gives (11.28). Each tilted operator is positive and unital by Theorem 11.4, so $\mathsf{U}_k^\psi(f-c) = \mathsf{U}_k^\psi f - c$ and $\|\mathsf{U}_k^\psi(f-c)\|_\infty \le \|f-c\|_\infty$ for every constant $c$, and the infimum over $c$ shows that its quotient has norm at most one. Multiplying by the bounds on $\mathrm{asm}_\ell$ and $\mathrm{hd}_n$ in (11.19) gives (11.29). □

**Proposition 11.20 (Finite networks admit no typed interaction comparison across a reducing step).** Let every agent space be finite. Then $\mathcal{G}_\ell \cong L^\infty(\nu_\ell)/\mathbb{R}1$ has dimension $|\mathrm{supp}\,\nu_\ell| - 1 = \prod_{i\in V_\ell} |\mathrm{supp}\,\nu_{\ell i}| - 1$. If a step reduces this number, no space admits linear isomorphisms onto both $\mathcal{G}_\ell$ and $\mathcal{G}_{\ell+1}$, so no comparison identification of the interaction spaces exists across that step; six binary agents blocked into three binary blocks give $\dim \mathcal{G}_\ell = 63$ and $\dim \mathcal{G}_{\ell+1} = 7$. If the blocking ratios are the vertex ratios $b_k = |V_k|/|V_{k+1}| \ge b > 1$, then $|V_n| \le |V_\ell| b^{-(n-\ell)}$, and a nonempty $V_n$ forces $n - \ell \le \log_b |V_\ell|$, so a finite network admits only finitely many such steps. [ESTABLISHED]

*Proof.* By Theorem 11.10, $\mathcal{G}_\ell$ is linearly isomorphic to $L^\infty(\nu_\ell)/\mathbb{R}1$. On a finite space $L^\infty(\nu_\ell)$ is the space of functions on the support of $\nu_\ell$, the product of the coordinate supports, and the quotient

lowers the dimension by one; linear isomorphisms preserve dimension. The vertex ratios telescope to $|V_\ell|/|V_n| = \prod_{k=\ell}^{n-1} b_k \geq b^{n-\ell}$, and $|V_n| \geq 1$ bounds $n - \ell$. □

On finite networks the depth objects are therefore the exact step, its residual (Proposition 11.22), and the multipliers, none of which needs a comparison isomorphism; blocking sequences with independently declared ratios, and infinite networks, are treated in Supplement B.

**Definition 11.21 (Retained projection and closure residual).** Let $\mathsf{P}^{\mathrm{ret}}_\ell : \mathcal{G}_\ell \to \mathcal{G}_\ell$ be a bounded idempotent that intertwines the gauge action of Corollary 11.11. The retained output, the coordinate residual, and the action residual are

$$\mathsf{g}^{\mathrm{ret}}_{\ell+1} = \mathsf{P}^{\mathrm{ret}}_{\ell+1}\mathcal{T}_\ell(\mathsf{g}_\ell), \qquad r^{\mathcal{G}}_{\ell+1}(\mathsf{g}_\ell) = \big(I - \mathsf{P}^{\mathrm{ret}}_{\ell+1}\big)\mathcal{T}_\ell(\mathsf{g}_\ell), \tag{11.30}$$

$$\overline{r}^{\mathcal{A}}_{\ell+1}(\mathsf{g}_\ell) = \overline{\mathcal{A}}_\ell\,\mathrm{asm}_\ell\,\mathsf{g}_\ell - \mathrm{asm}_{\ell+1}\,\mathsf{P}^{\mathrm{ret}}_{\ell+1}\,\mathrm{hd}_{\ell+1}\,\overline{\mathcal{A}}_\ell\,\mathrm{asm}_\ell\,\mathsf{g}_\ell. \tag{11.31}$$

**Proposition 11.22 (Residual identity and exact-image invariance).** The residuals satisfy $\overline{r}^{\mathcal{A}}_{\ell+1} = \mathrm{asm}_{\ell+1}\, r^{\mathcal{G}}_{\ell+1}$, so one vanishes exactly when the other does, and

$$\mathcal{T}_\ell\big(\mathrm{Ran}\,\mathsf{P}^{\mathrm{ret}}_\ell\big) \subseteq \mathrm{Ran}\,\mathsf{P}^{\mathrm{ret}}_{\ell+1} \quad\Longleftrightarrow\quad r^{\mathcal{G}}_{\ell+1}(\mathsf{g}) = 0 \quad \text{for every } \mathsf{g} \in \mathrm{Ran}\,\mathsf{P}^{\mathrm{ret}}_\ell. \tag{11.32}$$

A retained update is exact precisely under this exact-image invariance; otherwise it is a projection scheme with the displayed residual. For $V_{\ell+1} \neq \varnothing$,

$$\mathrm{hd}_{\ell+1}\,\overline{r}^{\mathcal{A}}_{\ell+1} = r^{\mathcal{G}}_{\ell+1}, \qquad \frac{\|r^{\mathcal{G}}_{\ell+1}\|_{\mathcal{G}_{\ell+1}}}{3^{|V_{\ell+1}|}-1} \leq \|\overline{r}^{\mathcal{A}}_{\ell+1}\|_{\overline{\mathfrak{B}}_{\ell+1}} \leq \|r^{\mathcal{G}}_{\ell+1}\|_{\mathcal{G}_{\ell+1}}. \tag{11.33}$$

[ESTABLISHED]

*Proof.* Since $\mathrm{asm}_{\ell+1}\,\mathrm{hd}_{\ell+1} = I$, the first term of (11.31) is $\mathrm{asm}_{\ell+1}\,\mathcal{T}_\ell(\mathsf{g}_\ell)$, so the action residual is $\mathrm{asm}_{\ell+1}$ applied to $(I - \mathsf{P}^{\mathrm{ret}}_{\ell+1})\mathcal{T}_\ell(\mathsf{g}_\ell)$, and $\mathrm{asm}_{\ell+1}$ is injective. An element lies in the range of an idempotent exactly when the complementary idempotent annihilates it, which is (11.32). Applying $\mathrm{hd}_{\ell+1}$ gives the coordinate identity, and the norm bounds are (11.19). □

**Proposition 11.23 (Radon–Nikodym covariance of the nonlinear interaction step).** At adjacent depths let $\kappa_\ell$ and $\kappa_{\ell+1}$ be realizations as in Corollary 11.11, transform the measure pair and the normalized law by pushforward, $\widehat{\varrho}_\ell = (\kappa_\ell)_{\#}\varrho_\ell$, $\widehat{\mathsf{m}}_\ell = (\kappa_\ell)_{\#}\mathsf{m}_\ell$, and $\widehat{\bar{\mathsf{m}}}_\ell = (\kappa_\ell)_{\#}\bar{\mathsf{m}}_\ell$, and suppose that the transformed channel satisfies $\widehat{C}_\ell(\kappa_\ell y, \kappa_{\ell+1}A) = C_\ell(y, A)$ for every Borel set $A$. Then, as Radon–Nikodym equivalence classes,

$$\widehat{\mathcal{A}}_\ell(\mathsf{K}_\ell\psi) = \mathsf{K}_{\ell+1}\mathcal{A}_\ell(\psi), \tag{11.34}$$

$$\widehat{\mathsf{U}}^{\mathsf{K}_\ell\psi}_\ell\mathsf{K}_\ell h = \mathsf{K}_{\ell+1}\mathsf{U}^{\psi}_\ell h, \tag{11.35}$$

$$\widehat{\mathcal{T}}_\ell\mathsf{K}^{\mathcal{G}}_\ell = \mathsf{K}^{\mathcal{G}}_{\ell+1}\mathcal{T}_\ell. \tag{11.36}$$

If the retained projections also intertwine, $\widehat{\mathsf{P}}^{\mathrm{ret}}_k \mathsf{K}^{\mathcal{G}}_k = \mathsf{K}^{\mathcal{G}}_k \mathsf{P}^{\mathrm{ret}}_k$ for $k = \ell, \ell+1$, then both residuals are covariant,

$$\widehat{r}^{\mathcal{G}}_{\ell+1}(\mathsf{K}^{\mathcal{G}}_\ell \mathsf{g}) = \mathsf{K}^{\mathcal{G}}_{\ell+1} r^{\mathcal{G}}_{\ell+1}(\mathsf{g}), \tag{11.37}$$

$$\widehat{\overline{r}}^{\mathcal{A}}_{\ell+1}(\mathsf{K}^{\mathcal{G}}_\ell \mathsf{g}) = \overline{\mathsf{K}}_{\ell+1} \overline{r}^{\mathcal{A}}_{\ell+1}(\mathsf{g}). \tag{11.38}$$

[ESTABLISHED]

The proof is in appendix B.4.

**Remark 11.24 (Grouping, symmetry, and pullback pairs).** Under Definition 11.6 at depths $\ell$ and $\ell+1$, three families of retained projections satisfy (11.32). (a) Let blocks $I \in V_{\ell+1}$ partition $V_\ell$, let the blocks be grouped into at least two groups, let the channel be block product (11.24), and let $\bar{\mathsf{m}}_\ell$ be a product over the groups. The grouping sector at either depth consists of the elements whose components vanish on every set that meets two groups; the two grouping sectors form an invariant pair, and when a group contains two blocks the depth-$\ell$ sector contains components coupling agents of different blocks. (b) Let a finite group act at both depths by product maps as in Corollary 11.11 that preserve the references, the measure pair, and the channel. The fixed sectors of its two actions on components form an invariant pair onto which the averaging operators are bounded idempotents; the base law may couple agents of different blocks, as for a law on $\{-1,1\}^{V_\ell}$ whose density is even under $y \mapsto -y$ with a channel that keeps one agent of each block, where the fixed sector at depth $\ell$ is spanned by the components of even order. (c) Let the channel be a deterministic Borel surjection $\chi$ and put $\chi^*[h] = [h \circ \chi]$ on $\overline{\mathfrak{B}}_{\ell+1}$. Then $\overline{\mathcal{A}}_\ell \chi^* = I$ and $\overline{\mathsf{U}}_\ell \chi^* = I$ for every base law, so for every retained projection $\mathsf{P}^{\mathrm{ret}}_{\ell+1}$ the operator

$$\mathsf{P}^{\mathrm{ret}}_\ell = \mathrm{hd}_\ell \, \chi^* \, \mathrm{asm}_{\ell+1} \, \mathsf{P}^{\mathrm{ret}}_{\ell+1} \, \mathrm{hd}_{\ell+1} \, \overline{\mathsf{U}_\ell} \, \mathrm{asm}_\ell \tag{11.39}$$

is a bounded idempotent whose range, the pullback sector $\mathrm{hd}_\ell \, \chi^* \, \mathrm{asm}_{\ell+1} \, \mathrm{Ran} \, \mathsf{P}^{\mathrm{ret}}_{\ell+1}$, is mapped by $\mathcal{T}_\ell$ onto $\mathrm{Ran} \, \mathsf{P}^{\mathrm{ret}}_{\ell+1}$. Transported with the gauge together with the grouping, the group action, and the channel, each projection intertwines the gauge action. [ESTABLISHED]

*Proof.* (a) The joint law $\bar{\mathsf{m}}_\ell(dy) C_\ell(y, dz)$ is a product over the groups, so a product of group reverse kernels is a version of the reverse kernel, and for $\mathsf{g}$ in the depth-$\ell$ sector $\psi_{\mathsf{g}}$ is a sum over groups of functions of the agents of one group, so that $\mathcal{A}_\ell(\psi_{\mathsf{g}})$ is by Fubini's theorem a sum over groups of functions of the coarse agents of one group. If $B$ meets two groups, choose $j \in B$ outside the group of one summand; $\mathbb{E}_{\ell+1,B}$ of that summand does not depend on $z_j$, so the factor $I - \mathbb{E}_{\ell+1, V_{\ell+1} \setminus \{j\}}$ of (11.17) annihilates it, as in the proof of Proposition 11.15. (b) Proposition 11.23 with every hatted object equal to the original gives $\mathcal{T}_\ell \mathsf{K}^{\mathcal{G}}_\ell = \mathsf{K}^{\mathcal{G}}_{\ell+1} \mathcal{T}_\ell$ for each group element, so a fixed $\mathsf{g}$ has a fixed image, and the average of a representation of a finite group by bounded operators is a bounded idempotent onto its fixed space. In the example the components of a set $B$ of binary agents are the multiples of $\prod_{i \in B} y_i$, on which the flip acts by $(-1)^{|B|}$, and the channel commutes with the flips at the two depths. (c) The reverse kernel of a deterministic channel is carried by the fibers of $\chi$, so $\int e^{-h(\chi(y))} C^\dagger(z, dy) = e^{-h(z)}$ and $\mathsf{U}_\ell(h \circ \chi) = h$ for $\bar{\mathsf{m}}_{\ell+1}$-almost every $z$. Idempotence follows

from $\overline{\mathrm{U}_\ell}\chi^* = I$ and $\mathrm{asm}_{\ell+1}\,\mathrm{hd}_{\ell+1} = I$, the same identities show that $\mathsf{P}_\ell^{\mathrm{ret}}$ fixes the pullback sector, and $\mathcal{T}_\ell\,\mathrm{hd}_\ell\,\chi^*\,\mathrm{asm}_{\ell+1} = \mathrm{hd}_{\ell+1}\,\overline{\mathcal{A}}_\ell\chi^*\,\mathrm{asm}_{\ell+1} = I$ on $\mathcal{G}_{\ell+1}$. A gauge change permutes the sets that index components, conjugates the group action, and turns $\chi$ into $\kappa_{\ell+1}\chi\kappa_\ell^{-1}$, so the transported projections intertwine by (11.20) and, for (c), (11.35) at $\psi = 0$. □

**Open Problem 11.25 (Invariant retained subspaces).** Under a product base law and a block-product channel, Proposition 11.15 gives (11.32) for the retained projections onto the intra-block sector and the singleton sector of the next depth, so the residual vanishes on them, and Remark 11.24 adds the grouping sectors of a base law that is a product over groups of blocks, the fixed sectors of a symmetry of the measure pair and the channel, and the pullback sectors of a deterministic channel under every base law. Unproved is an invariant pair beyond these. The problem is posed for finite agent spaces with fixed product references of full support, at least two blocks, and deterministic block-product channels each of whose block maps depends on every coordinate of its block, in each of two classes of base laws of full support, the product laws and all laws. It is settled by a proper invariant pair, with fixed retained projections, that is invariant under the step of every base law in a nonempty open subset of the class and whose depth-$\ell$ range is not contained in the sum of the pullback sector of $\mathcal{G}_{\ell+1}$ and the grouping sector of any partition of the blocks into at least two groups; or by a proof that, for base laws outside a closed nowhere dense subset of the class, every proper invariant pair has its depth-$\ell$ range in the sum of the pullback sector of $\mathcal{G}_{\ell+1}$ and the grouping sector of some such partition.

## 11.5 Exact closure generically requires hyperedges

**Proposition 11.26 (Elimination cliques and hyperedge generation).** (a) Suppose the complete nonnegative density of the law, baseline included, is factorized against a product or block-product reference. Marginalizing an internal variable multiplies every incident factor and integrates their product, giving one factor on the union of the remaining incident scopes. Iterating represents the normalized law by generally unnormalized factors on the elimination cliques, whose product is normalized by the inherited total normalizer; the joint law does not depend on the elimination order, although the displayed factorization does. A correlated baseline without such a factorization is kept as one global factor and pushed at the level of laws.

(b) For a finite coarse agent set $V^c$, a finite-valued coarse action $H_o^c$, and an anchor $z^\circ$, the anchored components satisfy, for $B \subseteq V^c$,

$$H_{o,B}^{c,\circ}(z_B) = \sum_{B' \subseteq B} (-1)^{|B|-|B'|} H_o^c(z_{B'}, z^\circ_{V^c \setminus B'}), \qquad H_o^c(z) = \sum_{B \subseteq V^c} H_{o,B}^{c,\circ}(z_B). \tag{11.40}$$

The expansion exists for every finite-valued function on a finite product, so closure of the full hypergraph family is automatic. The anchored components are indexed by the same subsets as the Hoeffding components (11.17) and also depend only on $z_B$, but their values differ, and only the Hoeffding components enter the exact step. An extended-valued coarse action is kept as one top-order factor, which closes without forming undefined alternating sums.

(c) The anchor is a covariant background: if the action is gauge invariant, each component satisfies $H^{c,\circ}_{o,B}(gz_B; gz^\circ) = H^{c,\circ}_{o,B}(z_B; z^\circ)$, while an anchor held fixed in one coordinate chart gives covariance of the sum only.

(d) Pairwise closure fails in general. Eliminating the center $y_0$ of a star of $\{-1, 1\}$-valued agents with field $\xi_0$ and couplings $\xi_1, \xi_2, \xi_3$ gives the factor $\sum_{y_0=\pm1} \exp[y_0(\xi_0 + \sum_{r=1}^3 \xi_r y_r)] = 2\cosh(\xi_0 + \sum_{r=1}^3 \xi_r y_r)$, whose negative logarithm has the coefficient $2\,\mathrm{sech}^2(\xi_0)\tanh(\xi_0)\xi_1\xi_2\xi_3 + O(\xi^5)$ on $y_1y_2y_3$, nonzero for all sufficiently small nonzero couplings when $\xi_0 \neq 0$. At $\xi_0 = 0$ this coefficient vanishes identically, because $\log 2\cosh(\sum_r \xi_r y_r)$ is even under $y \mapsto -y$, so hyperedge generation is generic, not universal. [ESTABLISHED]

The proof is in appendix B.4.

## 11.6 Fine–meta bridge kernels

**Definition 11.27 (Posterior bridge and bridge kernels).** At a record $o$ with posterior slice $\mathbf{\Pi}_o$ on the fine latent space and a coarse channel $C$, the posterior bridge is the joint law $\mathbb{B}_o(dy, dz) = \mathbf{\Pi}_o(dy)C(y, dz)$, with coarse marginal $\mathbf{\Pi}^c_o = \mathbf{\Pi}_o C$ and full reverse kernel $\mathbb{B}_o(dy \mid z)$. For a fine agent $j$ and a meta-agent $I$, the $(Y_j, Z_I)$ marginal of $\mathbb{B}_o$ disintegrated in both directions,

$$\mathbf{\Pi}_{o,j}(dy_j)K_{I\leftarrow j}(y_j, dz_I) = \mathbf{\Pi}^c_{o,I}(dz_I)K_{j\leftarrow I}(z_I, dy_j), \tag{11.41}$$

defines the agent-to-meta and meta-to-agent bridge kernels. For a deterministic statistic $z = \chi(y)$, the lift and the posterior conditional expectation are

$$(\mathsf{J}f)(y) = f(\chi(y)), \qquad (\mathsf{C}g)(z) = \mathbb{E}_{\mathbf{\Pi}_o}[g(Y) \mid \chi(Y) = z], \tag{11.42}$$

and for a randomized channel $(\mathsf{J}f)(y) = \int f(z)C(y, dz)$ and $(\mathsf{C}g)(z) = \int g(y)\mathbb{B}_o(dy \mid z)$.

**Proposition 11.28 (Bridge-kernel identities).** For bounded observables $f$ and $g$ the bridge kernels are adjoint, $\int f(y_j)(K_{I\leftarrow j}g)(y_j)\mathbf{\Pi}_{o,j}(dy_j) = \int (K_{j\leftarrow I}f)(z_I)g(z_I)\mathbf{\Pi}^c_{o,I}(dz_I)$. The kernel $K_{j\leftarrow I}$ does not reconstruct the fine posterior, whose full reverse kernel generally reads all meta variables: a product baseline and a block-product channel make the baseline reverse kernel factor over blocks, but a likelihood that couples blocks can destroy this factorization in the posterior, and posterior-local refinement requires $Y_I \perp Z_{-I} \mid Z_I$ under $\mathbb{B}_o$. The operators of (11.42) are adjoint in $L^2$; for a deterministic statistic $\mathsf{CJ} = I$ in $L^p(\mathbf{\Pi}^c_o)$, and the operators compose by the tower property along a nested filtration, while for a randomized channel $\mathsf{CJ} = I$ fails in general. [ESTABLISHED]

The proof is in appendix B.4.

## 11.7 Exact attention between meta-agents

A row-stochastic attention matrix is a conditional law of a source given a receiver and cannot be coarse-grained without the receiver law. The recursive attention state is therefore the joint edge-event law $\eta$ of (6.13) over the oriented edge copies $e : j \to i$, with receiver occupancy $\alpha$ supplied by a structural receiver-selection law or by the single-edge model of Definition 6.5.

**Proposition 11.29 (Exact meta-attention pushforward).** For a posterior bridge $\mathbb{B}_o$ and a partition of the agents into blocks, put

$$\eta^c_{IJ}(z) = \mathbb{E}_{\mathbb{B}_o}\Big[\sum_{\substack{e:j\to i\\ i\in I,\ j\in J}} \eta_e(Y)\ \Big|\ Z = z\Big], \qquad \alpha^c_I(z) = \sum_J \eta^c_{IJ}(z), \qquad \beta^c_{IJ}(z) = \frac{\eta^c_{IJ}(z)}{\alpha^c_I(z)}, \tag{11.43}$$

the last ratio on $\{\alpha^c_I > 0\}$. The coarse event law is normalized, pushing then disintegrating is associative under nested partitions, and the next depth pushes $\eta^c$, not $\beta^c$. If the channel and the selected posterior versions are gauge equivariant, $\eta^c$, $\alpha^c$, and $\beta^c$ are invariant as values at corresponding coarse states and covariant as functions of $z$. Events with both endpoints in $I$ become meta-self events in $\eta^c_{II}$, so the coarse source set of $I$ contains $I$ whenever $\eta^c_{II} > 0$, even when the fine rows exclude self-edges; for rows that do not depend on the latent state, $\eta^c$ is the block sum, and corollary 6.6 gives the uniform case. [ESTABLISHED]

*Proof.* Conditional expectation preserves $\sum_e \eta_e = 1$, and the tower property gives associativity. Equivariance transports $\mathbb{B}_o$ and leaves the scalar event probabilities unchanged. When the rows do not depend on the latent state the conditional expectation is the block sum. □

**Proposition 11.30 (Evidence-weighted meta-attention in the single-edge model).** Assume Definition 6.5, fix the latent state $y$, let $\mathcal{Z}_i(y) = \sum_j \pi_{ij} e^{-D_{ij}(y)/\tau_i}$ be the normalizer of the row $\beta^P_{ij}(y)$ of (6.4), and put $\overline{\mathcal{Z}}_I(y) = \sum_{k\in I} a_{k|I}\widetilde{c}_k(o,y)\mathcal{Z}_k(y)$ for a receiver block $I$ with $a_I = \sum_{k\in I} a_k > 0$ and $a_{i|I} = a_i/a_I$. On $\{\overline{\mathcal{Z}}_I(y) > 0\}$ the posterior receiver law conditional on the receiver lying in $I$ and the posterior block row are

$$\alpha'_i(y) = \frac{a_{i|I}\widetilde{c}_i(o,y)\mathcal{Z}_i(y)}{\overline{\mathcal{Z}}_I(y)}, \qquad \beta^c_{IJ}(y) = \sum_{i\in I} \alpha'_i(y) \sum_{j\in J} \beta^P_{ij}(y). \tag{11.44}$$

At that $y$ they coincide with the block sums and disintegration of (11.43) applied to the posterior edge-label law $\eta_{ij}(y) = \alpha^P_i(y)\beta^P_{ij}(y)$, whose receiver occupancy is $\alpha^P_i(y) \propto a_i\widetilde{c}_i(o,y)\mathcal{Z}_i(y)$. The formula holds posterior-almost everywhere; on a zero-evidence block the row is an arbitrary conditional version and is removed from the effective support. The equality $\alpha'_i = a_{i|I}$ holds exactly when $\widetilde{c}_i(o,y)\mathcal{Z}_i(y)$ is constant on the conditional support in $I$, which guarantees equality with the pre-observation coarse mixture, although the block rows can also agree accidentally after source coarsening. [ESTABLISHED]

*Proof.* Bayes' formula gives $P(i,j \mid y,o) \propto a_i \pi_{ij} \widetilde{c}_i(o,y) \exp[-D_{ij}(y)/\tau_i]$, because the label-free likelihood factor cancels. Summing over $j$ gives $P(i \mid y,o) \propto a_i \widetilde{c}_i(o,y) \mathcal{Z}_i(y)$, and conditioning on $i \in I$ gives $\alpha'_i$. Given $i$ the source posterior is $\beta^P_{ij}(y)$, and the block sums of $P(i,j \mid y,o)$ over $I \times J$ and over $I \times V$ have the stated ratio. □

**Proposition 11.31 (Prior-weighted collapse in the per-receiver model).** In the per-receiver label model of Section 6.1, under its exclusivity hypothesis (6.3), let $I$ be a set of receivers and let a structural coarse channel draw a receiver $i^\star \in I$ with the law $a_{\cdot|I}$, independently of the latent state, the labels, and the record, and output the source label of $i^\star$. Then the posterior law of $i^\star$ given $(y,o)$ is $a_{\cdot|I}$, and the posterior block row is

$$\Pr(\text{source of } i^\star \in J \mid y,o) = \sum_{i \in I} a_{i|I} \sum_{j \in J} \beta^P_{ij}(y), \tag{11.45}$$

with $\beta^P_{ij}$ of (6.4); its receiver occupancy is the selection law, and (11.44) does not apply to this model. The two block rows agree when $\widetilde{c}_i \mathcal{Z}_i$ is constant on the conditional support in $I$ and otherwise can differ: for two receivers with $a_{\cdot|I} = (1/2, 1/2)$, block-source masses $\sum_{j \in J} \beta^P_{ij}$ equal to $1$ and $0$, and $\widetilde{c}_i \mathcal{Z}_i$ equal to $3$ and $1$, the evidence-weighted mass of $J$ is $3/4$ and the prior-weighted mass is $1/2$. If all receiver labels are kept, the exact meta-label is their vector; the collapse to one categorical row is a further coarse channel, and its conditional relative entropy enters (9.3). [ESTABLISHED]

*Proof.* The selector enters neither the baseline nor the likelihood, so its posterior equals its prior and it is conditionally independent of the labels given $(y,o)$. By Proposition 6.2 the label posterior factors by receiver with row $i$ equal to $\beta^P_i(y)$, which gives (11.45). In the example the evidence weights are $3/4$ and $1/4$. □

**Remark 11.32 (Coarse prior and log-sum-exp energy of one row).** For one categorical row with prior $\pi_j$, row energies $D_j$, temperature $\tau$, and a source partition $j \mapsto J$, the block posterior is the Gibbs row of the coarse prior and coarse energy

$$\pi^c_J = \sum_{j \in J} \pi_j, \qquad D^c_J = -\tau \log\Big[\frac{1}{\pi^c_J} \sum_{j \in J} \pi_j \exp(-D_j/\tau)\Big], \tag{11.46}$$

for every block with $\pi^c_J > 0$; a block of zero prior has zero posterior mass and leaves the effective support. The rule associates for nested partitions of one row. Merging independently normalized receiver rows needs a receiver law: the evidence-weighted occupancy of Proposition 11.30 in the single-edge model and the selection law of Proposition 11.31 in the per-receiver model. [ESTABLISHED]

*Proof.* By construction $\pi^c_J \exp(-D^c_J/\tau) = \sum_{j \in J} \pi_j \exp(-D_j/\tau)$, so the block sums of the Gibbs row equal the Gibbs row of $(\pi^c, D^c)$ with the same normalizer. Applying the rule to a coarser partition of the blocks gives the same sums, which is associativity. □

**Proposition 11.33 (Transported meta-messages as marked moments).** Let each edge copy $e : j \to i$ with $i \in I$ and $j \in J$ carry its dressed link $\widehat{\Theta}^x_e$ of (10.13), let $\rho^{\mathbb{W}}_x : G \to \mathrm{GL}(\mathbb{W}_x)$ be the representation of the linear feature tier of (10.14), and assume the conditional Bochner-integrability condition $\mathbb{E}_{\mathbb{B}_o}[\sum_e \eta_e(Y)\|\rho^{\mathbb{W}}_x(\widehat{\Theta}^x_e)\| \mid Z] < \infty$ almost everywhere, the sum running over these edge copies. Define

$$\begin{aligned} \mathcal{M}^x_{IJ}(z) &= \mathbb{E}_{\mathbb{B}_o}\Big[ \sum_{\substack{e:j\to i \\ i\in I,\ j\in J}} \eta_e(Y)\rho^{\mathbb{W}}_x(\widehat{\Theta}^x_e) \;\Big|\; Z = z\Big], \\ \overline{\mathcal{M}}^x_{IJ} &= \frac{\mathcal{M}^x_{IJ}}{\eta^c_{IJ}} \quad \text{on } \{\eta^c_{IJ} > 0\}, \qquad \mathcal{M}^{x,\mathrm{att}}_{IJ} = \frac{\mathcal{M}^x_{IJ}}{\alpha^c_I} \quad \text{on } \{\alpha^c_I > 0\}. \end{aligned} \tag{11.47}$$

Then $\mathcal{M}^{x,\mathrm{att}}_{IJ} = \beta^c_{IJ}\overline{\mathcal{M}}^x_{IJ}$ on $\{\eta^c_{IJ} > 0\}$, and the attention-weighted block vanishes on $\{\alpha^c_I > 0, \eta^c_{IJ} = 0\}$, where $\overline{\mathcal{M}}^x_{IJ}$ is a null-event conditional left undefined. Under a passive root-frame change with no relabeling of agents, corresponding conditional versions, and the transformed coarse state $z'$, $\mathcal{M}^{x'}_{IJ}(z') = \rho^{\mathbb{W}}_x(g_{r_I})\mathcal{M}^x_{IJ}(z)\rho^{\mathbb{W}}_x(g_{r_J})^{-1}$ for $\mathbf{\Pi}^c_o$-almost every $z$, so these operators map the $J$ root frame to the $I$ root frame. The conditional mean need not lie in $\rho^{\mathbb{W}}_x(G)$: the average of the rotations by $\pi/2$ and $-\pi/2$ is zero. The moment is exact only as the first operator moment on a common source input: for equally likely marks and features $(I, v)$ and $(-I, -v)$ the separate means vanish while the mean product is $v$, so exact closure keeps the marked operator–feature kernel, with (11.47) as a derived moment. Since neither the definition nor the covariance uses invertibility, both hold for any measurable edge mark with values in $\mathrm{End}(\mathbb{W}_x)$ that obeys the two-endpoint law, under the corresponding integrability condition. In particular, when $\mathbb{W}_x$ is the vector sample fiber of Section 6.3 and $\rho^{\mathbb{W}}_x$ is its representation $\rho_x$, the selective links of (6.28) obey that law by (6.30), and their dressed event-weighted first moment is a covariant operator between root frames that need not be invertible. [ESTABLISHED]

*Proof.* The first two statements are the definitions of $\beta^c$ and of the ratios. The dressed link transforms as $\widehat{\Theta}^{x\prime}_e = g_{r_I}\widehat{\Theta}^x_e g^{-1}_{r_J}$ by (10.13), the event probabilities are invariant, and conditional expectation commutes with the fixed linear maps $\rho^{\mathbb{W}}_x(g_{r_I})$ and $\rho^{\mathbb{W}}_x(g_{r_J})^{-1}$. For a general mark, substituting its two-endpoint law and the tree transport law (10.11) into the dressed mark gives the same root-frame law. The two witnesses are direct computations. □

**Proposition 11.34 (Meta-attention and meta-agent kernels split one conditional divergence).** At a fixed outside context $z_{-I}$ let $\beta^{\mathbb{Q}}_I(J \mid z_{-I}) = \mathbb{Q}(J \mid z_{-I})$ and $K^{\mathbb{Q}}_{I\leftarrow J}(dz_I \mid z_{-I}) = \mathbb{Q}(dz_I \mid J, z_{-I})$, and define the generative counterparts under $\mathbb{P}(\cdot \mid z_{-I}, o)$. If the recognition row and kernels are absolutely continuous with respect to their generative counterparts, then, in $[0, +\infty]$,

$$\begin{aligned} D_{\mathrm{KL}}\big(\mathbb{Q}(J, Z_I \mid z_{-I})\|\mathbb{P}(J, Z_I \mid z_{-I}, o)\big) &= D_{\mathrm{KL}}\big(\beta^{\mathbb{Q}}_I(\cdot \mid z_{-I})\|\beta^{\mathbb{P}}_I(\cdot \mid z_{-I}, o)\big) \\ &+ \sum_J \beta^{\mathbb{Q}}_I(J \mid z_{-I}) D_{\mathrm{KL}}\big(K^{\mathbb{Q}}_{I\leftarrow J}(\cdot \mid z_{-I})\|K^{\mathbb{P}}_{I\leftarrow J}(\cdot \mid z_{-I}, o)\big). \end{aligned} \tag{11.48}$$

[ESTABLISHED]

*Proof.* This is the chain rule Lemma 3.12 for the categorical coordinate $J$ followed by $Z_I$. □

## 11.8 Path laws and memory

On a finite ordered path interval the pointwise coarse channel applied to the joint path law is again a coarse channel, so Theorem 9.6 applies on path space; the order parameter is an inference index, not the Fisher duration of Chapter 7. The coarse path law need not be first-order Markov (Remark 11.36); this section says when it is and gives the exact non-Markov dynamics and the cost of finite memory.

**Theorem 11.35 (Standard-Borel strong lumpability).** Let $\chi : \mathsf{Y} \to \mathsf{Z}$ be Borel and surjective, and let $\mathsf{T}$ be a Markov kernel on $\mathsf{Y}$. Consider

$$\mathsf{T}(y, \chi^{-1}A) = \mathsf{T}(y', \chi^{-1}A) \quad \text{whenever } \chi(y) = \chi(y'), \qquad \text{for every Borel } A \subseteq \mathsf{Z}. \tag{11.49}$$

If a Markov kernel $\mathsf{T}^\chi$ on $\mathsf{Z}$ satisfies

$$\chi_\#(\mu\mathsf{T}) = (\chi_\#\mu)\mathsf{T}^\chi \qquad \text{for every probability } \mu \text{ on } \mathsf{Y}, \tag{11.50}$$

then (11.49) holds and $\mathsf{T}^\chi$ is unique. Conversely, if (11.49) holds and $\chi$ has a Borel right inverse $\iota$, then $\mathsf{T}^\chi(z, A) = \mathsf{T}(\iota(z), \chi^{-1}A)$ is a Markov kernel satisfying (11.50). A Borel right inverse exists for a countable coarse space and for a product projection $\mathsf{Y} = \mathsf{Z} \times \mathsf{Y}'$ with $\mathsf{Y}'$ nonempty. For finite spaces the biconditional is the lumpability criterion of Kemeny and Snell [51, Thm. 6.3.2]. [ESTABLISHED]

The proof is in appendix B.4.

**Remark 11.36 (Two memory witnesses).** Strong lumpability is the every-initial-law condition and is not necessary for a single initial law. On $\mathsf{Y} = \{1, 2, 3\}$ with $\chi(1) = \chi(2) = 0$ and $\chi(3) = 1$, let $\mathsf{T}$ send $1$ to $3$, send $2$ to $1$ or $2$ with equal probability, and send $3$ to $1$ or $3$ with equal probability. Then $\mathsf{T}(1, \chi^{-1}\{0\}) = 0$ while $\mathsf{T}(2, \chi^{-1}\{0\}) = 1$, so (11.49) fails, but started at $\delta_3$ the chain stays in $\{1, 3\}$, on which $\chi$ is injective, and the coarse process is Markov. A stationary coarse process can also have memory. On the deterministic cycle $1 \to 2 \to 3 \to 1$ with uniform stationary law and labels $\chi = (0, 0, 1)$, the observed histories $(0, 0)$ and $(1, 0)$ both have positive probability and the same last symbol, but predict the next symbols $1$ and $0$ respectively, so first-order Markov closure fails under this stationary preparation. [ESTABLISHED]

*Proof.* In the first witness the reachable set from $3$ is $\{1, 3\}$, and the coarse transitions are $0 \mapsto 1$ surely and $1 \mapsto 0$ or $1$ with equal probability. In the second, the history $(0, 0)$ is the state pair $(1, 2)$, followed by $3$, and the history $(1, 0)$ is $(3, 1)$, followed by $2$. □

**Definition 11.37 (Regular finite-dimensional filter class).** For the dynamic-closure problem below, an admitted filter is a deterministic functional of the retained path taking values in a finite-dimensional smooth manifold, with dimension bounded independently of $\epsilon$ and the horizon. Between assignment jumps it has a time-homogeneous finite-dimensional stochastic realization with locally Lipschitz coefficients; its initialization and jump-update maps are continuous in their declared native

inputs. A decoded conditional law of the hidden coordinates has a density with respect to one declared sigma-finite measure on the hidden-coordinate space, and the decoder is continuously Fréchet differentiable from the filter manifold into that $L^1$ space. Discrete assignment values are treated as separate strata. This is an additional regularity restriction for this problem: dimension alone does not rule out Borel encodings of an entire history or of a probability-valued filter into one real coordinate. The restriction excludes such arbitrary encodings and singular decoded families without a common dominating measure, while allowing exact finite-dimensional filters that satisfy these conditions. Changing the regularity class changes the closure question.

**Open Problem 11.38 (Dynamic closure of the pointwise parent).** The static construction of Corollary 9.10 supplies no fine evolution and no parent evolution. In a differentiable special case with fine vector field $\mathscr{V}_t$, moving deterministic coarse map $\mathsf{F}_t$, and candidate parent field $\overline{\mathscr{V}}_t$, the exact semiconjugacy defect is

$$\delta_t = \partial_t \mathsf{F}_t + D\mathsf{F}_t \mathscr{V}_t - \overline{\mathscr{V}}_t \circ \mathsf{F}_t. \tag{11.51}$$

For a Markov evolution the corresponding exact condition is semigroup intertwining; a generator identity is an alternative criterion only on an invariant domain or core, with generation and uniqueness hypotheses sufficient to lift it to the semigroups. Neither condition is proved for the static datum together with any autonomy, moving-membership, path-law, stochastic, delayed, memory, or nonequilibrium hypothesis. Without a class the question is decided in both directions by statements already present: strongly lumpable kernels lump (Theorem 11.35), the prepared parent intertwines under exact refresh (Theorem 8.11), a parent carrying the terminal posterior of Theorem 11.42 reproduces every path law, and the stationary cycle of Remark 11.36 and the finite refresh of Remark 8.12 admit no first-order parent on their coarse labels. The problem is therefore posed for switching diffusions at $c_*$ of the type of Definition 13.21, indexed by the regime parameter $\epsilon$: finitely many source assignments, a positive-definite quadratic sheet energy for each assignment with Hessian bounded uniformly in $\epsilon$, diffusion of the configuration with identity mobility on the sheet of the current assignment, and one rate-one clock at which the assignment is redrawn from its conditional law under the restricted optimizer. The channels keep the assignment and an orthogonal projection of the configuration, fixed or depending on the assignment, as in part (b) of Proposition 13.22, and on some assignment whose weight under the restricted optimizer is bounded below in $\epsilon$ they hide coordinates on which the receiver rows of that assignment depend. The parents are time-homogeneous Markov evolutions of the retained assignment and coordinates, or of these together with a filter in the explicitly restricted class of Definition 11.37, each activated from the pushforward of the recognition law at the activation time. Two levels are asked: the retained and assignment path law on windows of fixed length, uniformly over windows placed up to the switching time, the time after activation by which the expected number of assignment changes reaches one; and the complete retained and assignment history over horizons on which the expected number of assignment changes diverges as $\epsilon \to 0$. At either level the problem is settled by a proof that every member of the class admits such a parent whose retained and assignment path law is within a total variation tending to zero with $\epsilon$ of the native one, or by a member for which every such parent keeps that distance bounded below.

**Theorem 11.39 (Exact discrete projection-memory recurrence).** Let X and W be vector spaces, let $\mathsf{T} : \mathsf{X} \to \mathsf{X}$, $\mathsf{C} : \mathsf{X} \to \mathsf{W}$, and $\mathsf{J} : \mathsf{W} \to \mathsf{X}$ be linear with $\mathsf{CJ} = I_{\mathsf{W}}$, and set $\mathsf{P}^{\mathrm{res}} = \mathsf{JC}$ and $\mathsf{Q} = I - \mathsf{P}^{\mathrm{res}}$, so that $\mathsf{P}^{\mathrm{res}}$ is idempotent. For the fine recursion $\xi_{n+1} = \mathsf{T}\xi_n$ write $w_n = \mathsf{C}\xi_n$, so that $\mathsf{P}^{\mathrm{res}}\xi_n = \mathsf{J}w_n$. Then for every $n \geq 0$ the unresolved part and the resolved coordinate obey

$$\mathsf{Q}\xi_n = (\mathsf{QTQ})^n \mathsf{Q}\xi_0 + \sum_{k=0}^{n-1} (\mathsf{QTQ})^{n-1-k} \mathsf{QTJ} w_k, \tag{11.52}$$

$$w_{n+1} = \mathsf{CTJ} w_n + \sum_{k=0}^{n-1} \mathsf{CTQ}(\mathsf{QTQ})^{n-1-k} \mathsf{QTJ} w_k + \mathsf{CTQ}(\mathsf{QTQ})^n \mathsf{Q}\xi_0, \tag{11.53}$$

an exact closed recurrence for $w_n$. The operators $\mathsf{CTQ}(\mathsf{QTQ})^m\mathsf{QTJ}$, $m \geq 0$, are the memory kernels and the last term of (11.53) is the unresolved initial-state term; no closure or Markov assumption is used. [ESTABLISHED]

*Proof.* Applying $\mathsf{P}^{\mathrm{res}}$ and Q to $\xi_{n+1} = \mathsf{T}\xi_n$ with $\xi_n = \mathsf{J}w_n + \mathsf{Q}\xi_n$ and $\mathsf{Q}^2 = \mathsf{Q}$ gives $w_{n+1} = \mathsf{CTJ}w_n + \mathsf{CTQQ}\xi_n$ and $\mathsf{Q}\xi_{n+1} = \mathsf{QTQQ}\xi_n + \mathsf{QTJ}w_n$. The second relation gives (11.52) by induction on $n$, and substituting it into the first gives (11.53). □

For the slow block of a fast–slow coordinatization (Definition 8.2) of a linear evolution, (11.53) is the exact discrete slow dynamics; the projection-operator method goes back to Nakajima [63] and Zwanzig [89], in continuous time on ensemble densities.

**Corollary 11.40 (Autonomy of the resolved recursion).** Let $\mathscr{S}_0 \subseteq \mathsf{X}$ be a class of initial states. The resolved recursion is autonomous on $\mathscr{S}_0$, meaning $w_{n+1} = \mathsf{CTJ}w_n$ for every $n \geq 0$ and every $\xi_0 \in \mathscr{S}_0$, exactly when the total correction in (11.53) vanishes for every such $n$ and $\xi_0$. It suffices that all memory kernels vanish and that $\mathsf{CTQ}(\mathsf{QTQ})^n\mathsf{Q}\xi_0 = 0$ for $n \geq 0$ throughout $\mathscr{S}_0$. On the resolved class $\mathsf{Q}\xi_0 = 0$ the second family is vacuous, and there $\mathsf{QTJ} = 0$ is sufficient but not necessary: on $\mathsf{X} = \mathbb{R}^2$ with $\mathsf{W} = \mathbb{R}$, $\mathsf{C}(u, v) = u$, $\mathsf{J}(w) = (w, 0)$, and $\mathsf{T}(u, v) = (u, u)$, one has $\mathsf{QTJ}w = (0, w) \neq 0$ for $w \neq 0$ while $\mathsf{CTQ} = 0$, so every memory kernel vanishes and $w_{n+1} = w_n$ is exactly autonomous. [ESTABLISHED]

*Proof.* The biconditional is (11.53) read as an identity, and the two families are its two correction terms. In the witness $\mathsf{Q}(u, v) = (0, v)$ and $\mathsf{T}(0, v) = (0, 0)$ give $\mathsf{CTQ} = 0$, while $\mathsf{T}(w, 0) = (w, w)$ gives $\mathsf{QTJ}w = (0, w)$. □

### 11.8.1 The hidden history and its filter

**Definition 11.41 (Hidden path experiment).** Let $Y_t$ take values in a finite set Y with initial law $p_0$ and transition matrix T, and let $Z_t = \chi(Y_t)$ take values in a finite set Z, with $\mathsf{P}_z$ the diagonal projector onto $\chi^{-1}(z)$. The joint path law over $N$ steps and its observed marginal are

$$p_0(y_0) \prod_{t=0}^{N-2} \mathsf{T}_{y_t y_{t+1}} \prod_{t=0}^{N-1} \mathbf{1}_{\chi(y_t)=z_t}, \qquad \mathbb{P}_N^{\mathrm{path}}(z_{0:N-1}) = p_0 \mathsf{P}_{z_0} \mathsf{T} \mathsf{P}_{z_1} \cdots \mathsf{T} \mathsf{P}_{z_{N-1}} \mathbf{1}. \tag{11.54}$$

This auxiliary model, the law-level counterpart of Theorem 11.39, is fixed before a recognition law on its paths is varied.

**Theorem 11.42 (Exact filtering without a projected Markov assumption).** For an observed word $z_{0:n}$ of positive probability put

$$\mathsf{f}_{z_{0:n}} = p_0 \mathsf{P}_{z_0} \mathsf{T}\mathsf{P}_{z_1} \cdots \mathsf{T}\mathsf{P}_{z_n}, \qquad \widehat{p}_{z_{0:n}} = \frac{\mathsf{f}_{z_{0:n}}}{\mathsf{f}_{z_{0:n}}\mathbf{1}}, \qquad \mathbb{P}(Z_{n+1} = z \mid z_{0:n}) = \widehat{p}_{z_{0:n}} \mathsf{T}\mathsf{P}_z \mathbf{1}. \tag{11.55}$$

These are the exact terminal posterior and next-symbol prediction. A word of zero probability has zero messages for all its extensions and undefined conditional readouts. The terminal posterior determines the law of every future suffix, but its attainable values need not form a finite set. [ESTABLISHED]

*Proof.* Summing the joint path law with the final hidden state fixed gives $\mathsf{f}_{z_{0:n}}$. Replacing $\mathsf{P}_z$ by $I$ and using $\mathsf{T}\mathbf{1} = \mathbf{1}$ proves normalization and prefix consistency, division by the positive word mass gives the posterior, and a further matrix product gives any suffix. Nonnegative zero messages stay zero under multiplication. A posterior is a probability vector on Y, and sufficiency does not assert that its attainable subset is finite. □

### 11.8.2 Finite-memory projection

**Definition 11.43 (Finite-memory projection of a path law).** Let $\mathbb{P}_N^{\mathrm{path}}$ be any law on $\mathsf{Z}^N$, stationary or not. For an integer $k \geq 0$ and a position $t$ let $\mathbf{z}_t = z_{\max(0,t-k):t-1}$ be the retained context. On contexts of positive mass define the rows and the projected law

$$\mathsf{T}_t^{(k)}(z \mid \mathbf{z}) = \mathbb{P}_N^{\mathrm{path}}(Z_t = z \mid \mathbf{z}_t = \mathbf{z}), \qquad \widehat{\mathbb{P}}_N^{(k)}(z_{0:N-1}) = \prod_{t=0}^{N-1} \mathsf{T}_t^{(k)}(z_t \mid \mathbf{z}_t), \tag{11.56}$$

recursively, so that a prefix of zero mass has zero extensions and no row is chosen on a null context.

**Theorem 11.44 (Forward-KL optimality and its exact information defect).** The projected law (11.56) is normalized, has the same context-plus-symbol marginals as $\mathbb{P}_N^{\mathrm{path}}$, and is positive on every word of positive $\mathbb{P}_N^{\mathrm{path}}$ probability. Among all laws with time-dependent $k$-context rows it uniquely minimizes the forward path divergence, and the minimum is

$$D_N^{(k)} := D_{\mathrm{KL}}\big(\mathbb{P}_N^{\mathrm{path}} \| \widehat{\mathbb{P}}_N^{(k)}\big) = \sum_{t=0}^{N-1} I\big(Z_t; Z_{0:t-k-1} \mid Z_{\max(0,t-k):t-1}\big), \tag{11.57}$$

an empty discarded history contributing zero. For a competing law $\widetilde{\mathbb{P}}$ with rows $\widetilde{\mathsf{T}}_t$,

$$D_{\mathrm{KL}}\big(\mathbb{P}_N^{\mathrm{path}} \| \widetilde{\mathbb{P}}\big) - D_N^{(k)} = \sum_{t,\mathbf{z}:\ \mathbb{P}_N^{\mathrm{path}}(\mathbf{z}_t = \mathbf{z}) > 0} \mathbb{P}_N^{\mathrm{path}}(\mathbf{z}_t = \mathbf{z}) D_{\mathrm{KL}}\big(\mathsf{T}_t^{(k)}(\cdot \mid \mathbf{z}) \| \widetilde{\mathsf{T}}_t(\cdot \mid \mathbf{z})\big). \tag{11.58}$$

The defect decreases in $k$ and vanishes for $k \geq N-1$; the projection is idempotent and commutes with restriction to a shorter prefix; and $2\,\mathrm{TV}(\mathbb{P}_N^{\mathrm{path}}, \widehat{\mathbb{P}}_N^{(k)})^2 \leq D_N^{(k)}$. The reverse divergence is not controlled: the stationary cycle of Remark 11.36 has $\mathbb{P}_3^{\mathrm{path}}(000) = 0$ but $\widehat{\mathbb{P}}_3^{(1)}(000) = 1/6$. For a path recognition law $\mathbb{Q}^{\mathrm{path}}$ supported on the support of $\mathbb{P}_N^{\mathrm{path}}$,

$$D_{\mathrm{KL}}\big(\mathbb{Q}^{\mathrm{path}} \| \widehat{\mathbb{P}}_N^{(k)}\big) - D_{\mathrm{KL}}\big(\mathbb{Q}^{\mathrm{path}} \| \mathbb{P}_N^{\mathrm{path}}\big) = \mathbb{E}_{\mathbb{Q}^{\mathrm{path}}} \log \frac{\mathbb{P}_N^{\mathrm{path}}}{\widehat{\mathbb{P}}_N^{(k)}}, \tag{11.59}$$

whose worst case over such laws is set by the extrema of the log ratio, not by its $\mathbb{P}_N^{\mathrm{path}}$-average. [ESTABLISHED]

**Theorem 11.45 (Stationary projection and the finite hidden information budget).** Let the hidden chain of Definition 11.41 start from a stationary law $p_0 = \bar{p}$, $\bar{p}\mathsf{T} = \bar{p}$, restricted to its positive support, which is closed under $\mathsf{T}$; arbitrary initial preparations are excluded. Write $\mathrm{h}_j = H(Z_j \mid Z_{0:j-1})$ in natural logarithms, with $\mathrm{h}_0 = H(Z_0)$, and $\mathrm{c}_Y = H(Y_0)$. Then the stationary projection is normalized, stationary, and prefix consistent, and agrees with the source on every window of length at most $k+1$. The sequence $\mathrm{h}_j$ decreases to a limit $\mathrm{h}_\infty$, and with $\mathrm{h}_k^{\mathrm{ex}} = \mathrm{h}_k - \mathrm{h}_\infty$,

$$\begin{aligned} D_N^{(k)} &= \sum_{t=k+1}^{N-1} (\mathrm{h}_k - \mathrm{h}_t), \quad \lim_{N\to\infty} \frac{D_N^{(k)}}{N} = \mathrm{h}_k^{\mathrm{ex}}, \\ \sum_{j=0}^{\infty} \mathrm{h}_j^{\mathrm{ex}} &\leq \mathrm{c}_Y, \qquad\qquad \mathrm{h}_k^{\mathrm{ex}} \leq \frac{\mathrm{c}_Y}{k+1}, \\ 0 \leq \mathrm{h}_k^{\mathrm{ex}} - \frac{D_N^{(k)}}{N} &\leq \frac{\mathrm{c}_Y}{N}, \qquad\qquad D_N^{(k)} \leq (N-k-1)_+ \mathrm{h}_k^{\mathrm{ex}}, \end{aligned} \tag{11.60}$$

an empty sum being zero. The series $\sum_j \mathrm{h}_j^{\mathrm{ex}}$ is the excess entropy of the observed process in the sense of Crutchfield and Feldman [23], so the second line bounds it by the hidden-state entropy. No irreducibility or forgetting assumption is used. [ESTABLISHED]

The proofs of Theorems 11.44 and 11.45 are in appendix B.4.

The memory cost of a prepared parent whose within-parent refresh is only partial (Section 8.3) is the defect $D_N^{(1)}$ of Theorem 11.44 for the parent process.

## 11.9 The exact retained state

Exact closure keeps every object of Table 11.1, listed with its exact operation, source statement, and admission hypotheses, uniformly in the number and arrangement of the agents.

The blocks and the channel enter the table as inputs, and the smallest case in which three two-dimensional blocks are asked for shows that the law cannot supply them.

**Table 11.1:** The retained state of exact agent-network coarse-graining at $c_*$.

| Retained object | Exact operation and standing | Admission hypotheses |
|---|---|---|
| Normalized generative law, posterior versions, recognition law | One channel pushes all three; the variational free energy decomposes by the chain rule. Derived: Theorem 9.6, Proposition 11.2. | Recognition-independent channel; finite positive evidence; recognition law absolutely continuous with respect to the posterior. |
| Reference probability, evidence measure, effective action | Radon–Nikodym action by the conditional partition function, composing along towers. Derived: Theorem 11.3. | Probability reference pushed with the evidence measure; reverse conditionals from one joint tower. |
| Bounded action calculus | Analytic action increment, tilted derivatives, contraction and oscillation bounds. Derived: Theorem 11.4, Theorem 11.5. | Bounded perturbations; jointly measurable reverse kernel. |
| Interaction coordinates and the exact step | Hoeffding isomorphism and exact nonlinear step with its closure residual. Derived: Theorem 11.10, Lemma 11.13, Proposition 11.22. Retained projection: declared. | Product-equivalent normalized laws at both depths (Definition 11.6). |
| Finite-step multipliers | Ordered products of tilted reverse operators, bounded by $3^{\lvert V_n\rvert} - 1$. Derived: Definition 11.18, Proposition 11.19. | An exact orbit in the bounded tier. |
| Induced hyperedges | Elimination cliques and Möbius components; generically nonpairwise. Derived: Proposition 11.26. | Finite-valued coarse action, or one top-order factor. |
| Bridge kernels | Pair-marginal disintegrations of the posterior bridge. Derived: Proposition 11.28. Gauge-covariant versions: declared beyond the compact and finite cases (Open Problem 9.18). | Regular disintegration on standard Borel spaces. |
| Roots, based holonomy, dressed boundary marks | Raw root-framed marks and their joint law, pushed without a holonomy-blind claim. Derived: Section 10.2. Trees: declared. | Rooted spanning forests of the blocks. |
| Joint attention-event law | Event pushforward and disintegration on positive occupancy. Derived: Proposition 11.29, Proposition 11.30, Proposition 11.31. | Receiver law of the label model. |
| Marked operator–feature kernels | Event-weighted first operator moments. Derived: Proposition 11.33. | Conditional Bochner integrability. |
| Path laws and memory | Lumped kernel when strongly lumpable; otherwise memory kernels, filters, and finite-memory projections with their defects. Derived: Theorem 11.35, Theorem 11.39, Theorem 11.44, Theorem 11.45. | The full path law, or the memory kernels and initial-state term. |

**Proposition 11.46 (Law and Fisher geometry do not canonically agentize blocks).** At the uniform law on seven outcomes, no rule that depends only on the unlabeled law and its Fisher metric and is natural under every relabeling of the outcomes can select an unordered decomposition of the six-dimensional tangent space into three two-dimensional blocks. Declared typing or other symmetry-breaking data can supply blocks; the law and its Fisher tensor alone do not. [ESTABLISHED]

The proof is in appendix B.4.

## Summary for physicists

One exact coarse-graining step pushes normalized laws through a Markov channel; the effective action is the Radon–Nikodym derivative of the pushed evidence measure with respect to a pushed reference probability, the conditional partition function of the fine action over each coarse state, analytic in bounded changes of the fine action and never expanding the supremum norm. In Hoeffding coordinates, one component per set of coupled agents, the step is an exact nonlinear map whose finite-step multipliers telescope to one extraction, a product of tilted conditional expectations, and one assembly, so they are bounded by $3^{|V_n|} - 1$ at every depth. Exactness costs state: eliminated agents generate hyperedges unless a symmetry forbids them, attention is carried as the joint event law, transported messages are moments of marked kernels, and the coarse path law generally has memory, with an exact memory-kernel recurrence and, for memory length $k$, a finite-memory cost of at most the hidden-state entropy divided by $k + 1$ per step.

# Chapter 12

# The Multiscale Graph VFE and Closing the Tower

This chapter asks what probabilistic structure places adaptive event laws and links on the support of the source prior, random partitions, and the influence of parents on their children inside one variational free energy, and how a tower of such depths closes. The answer is the free energy of one normalized generative law and one recognition law over the whole tower (Definition 12.1). For an adaptive tower both laws factor top-down, each depth conditioned on the whole next-depth state and the structural mark (Lemma 12.5), and the free energy splits into observation, top-prior, partition, graph, and parent–child terms (Theorem 12.7). A finite acyclic tower needs a declared latent root (Proposition 12.20); participation across depths enters through the inference loop, parameter learning (Proposition 12.22), and delayed feedback (Propositions 12.24 and 12.25). The formation witnesses of Chapter 13 are adaptive towers, and Chapter 14 reads the three loci of participation.

The depth $\ell$ runs from $0$ to a finite $L$. At depth $\ell$, $Z_\ell$ are the coordinates, $\mathfrak{N}_\ell$ the graph datum, $Y_\ell = (Z_\ell, \mathfrak{N}_\ell)$ the state on a standard Borel space $\mathsf{Y}_\ell$ that here includes the graph datum, $R_\ell \in \mathsf{R}_\ell$ the structural mark with realized value $\gamma$, and $\mathfrak{h}_\ell$ the retained root, holonomy, and boundary marks. The whole tower is $\mathbf{Y}$, row normalizers are $\mathcal{Z}$, and $\mathbb{M}$ is an input law: generative, posterior, or recognition. In the episode-indexed models of Section 12.7.2, $\widetilde{o}_n$ is a latent copy of the record $o_n$, $\epsilon_{\mathrm{e}}$ is the constant step of the aggregate parameter update, and $\mathrm{T}$ is the temperature of the descent kinetics, the symbol of Chapter 13.

Scope. In an adaptive tower the downward kernel is a reversal kernel, and the coupling it carries is not read as causal downward influence; declining that reading does not refute a causal reading of a separately declared top-down hierarchy. The composite potential of Definition 12.10 is not claimed to be an evidence bound, and delayed feedback is not claimed to support a sustained nonequilibrium reading. [NOT-CLAIMED] Open items are the persistent hierarchy (Open Problem 12.27), the accounting across episodes (Open Problem 12.26), the dynamic closure of the pointwise parent (Open Problem 11.38), and a natural partition selector (Open Problem 9.19).

## 12.1 The full answer before decomposition

**Definition 12.1 (Tower variational free energy).** Write $X_* = X(c_*)$, fix an admitted observation $o \in \mathsf{O}^{\mathrm{reg}}_{\theta, X_*}$ and a finite depth $L$, and let $\mathbf{Y}$ contain every random object of the multiscale model: belief

and model coordinates at depths $0, \dots, L$, links, event laws, retained marks, structural marks, and auxiliary latents. Declare one normalized generative kernel $\mathbb{P}_\theta(do, d\mathbf{Y} \mid X_*)$ and one recognition kernel $\mathbb{Q}_\lambda(d\mathbf{Y} \mid o, X_*)$. With posterior slice $\mathbf{\Pi}_{\theta,o,X_*}$ and evidence density $p_\theta(o \mid X_*) \in (0, \infty)$, the tower variational free energy is

$$\boxed{\mathcal{F}^{\mathrm{ext}}_{\mathrm{tower}}[\mathbb{Q}_\lambda; X_*, o] = -\log p_\theta(o \mid X_*) + D_{\mathrm{KL}}\big(\mathbb{Q}_\lambda(\cdot \mid o, X_*) \| \mathbf{\Pi}_{\theta,o,X_*}\big).} \tag{12.1}$$

By the evidence identity of Chapter 5, on the finite-density domain it equals

$$\mathcal{F}_{\mathrm{tower}} = \mathbb{E}_{\mathbb{Q}_\lambda}\Big[\log \frac{q_\lambda(\mathbf{Y} \mid o, X_*)}{p_\theta(o, \mathbf{Y} \mid X_*)}\Big]. \tag{12.2}$$

Every later term is an accounting decomposition of this object or is named a composite potential.

## 12.2 A normalized hierarchical graph model

**Definition 12.2 (Closed scale datum).** At depth $\ell$ let $V_\ell$ be the vertex set and let $Z_\ell = (Z_{i,\ell})_{i \in V_\ell}$ contain the full probabilistic coordinates of that depth, belief and model presentations rather than only their marginals. The graph datum lives on a finite interaction multigraph $\mathfrak{I}_\ell = (V_\ell, E_\ell, {}^-)$ (Definition 2.11), at depth $0$ the support of the source prior and at deeper depths induced by the link lift of Definition 12.17. Put

$$\eta_\ell = (\eta^{b,\ell}, \eta^{m,\ell}), \qquad \Theta^{x,\ell} = (\Theta^{x,\ell}_e)_{e \in E_\ell}, \qquad \mathfrak{N}_\ell = (\eta_\ell, \Theta^{b,\ell}, \Theta^{m,\ell}, \mathfrak{h}_\ell), \tag{12.3}$$

where $\eta^{x,\ell} = (\eta^{x,\ell}_e)_{e \in E_\ell}$ is a normalized event law on the copies, $\Theta^{x,\ell}_e \in G$ is the channel-$x$ link of the copy $e$ with $\Theta^{x,\ell}_{\bar{e}} = (\Theta^{x,\ell}_e)^{-1}$, and $\mathfrak{h}_\ell$ contains the root, based-holonomy, and dressed boundary marks of Section 10.2. Each copy carries one link, one event mass, and one divergence per channel; endpoint notation abbreviates a copy only when it is unique, and $\eta^{x,\ell}_{ij}$ is the total mass of the copies $j \to i$. Occupancies and rows are derived from $\eta_\ell$ as in (6.13) and do not replace it in the recursive state. The observation and $X_*$ stay outside every coarse channel.

**Definition 12.3 (Common structural channel).** A structural step is one recognition-independent Markov channel

$$C^{\mathrm{struct}}_{\ell+1 \leftarrow \ell} : \mathsf{Y}_\ell \rightsquigarrow \mathsf{R}_\ell \times \mathsf{Y}_{\ell+1}, \tag{12.4}$$

whose output includes $(R_\ell, \eta_{\ell+1}, \mathfrak{h}_{\ell+1})$, with augmented transition state

$$\widehat{\mathsf{Y}}_{\ell+1 \leftarrow \ell} := (\mathsf{Y}_\ell, R_\ell, \mathsf{Y}_{\ell+1}), \tag{12.5}$$

which does not duplicate $R_\ell$. The same channel is applied to the fine generative marginal, the selected posterior, and the recognition law,

$$\widehat{\mathbb{M}}^{\mathrm{struct}}_{\ell+1 \leftarrow \ell}(dy_\ell, d\gamma, dy_{\ell+1}) := \mathbb{M}_\ell(dy_\ell) C^{\mathrm{struct}}_{\ell+1 \leftarrow \ell}(y_\ell; d\gamma, dy_{\ell+1}), \qquad \mathbb{M} \in \{\mathbb{P}, \mathbf{\Pi}, \mathbb{Q}\}, \tag{12.6}$$

so the three laws of $R_\ell$ may differ while the channel does not. If the channel is realized by an auxiliary seed, the same measurable map and seed law define all three pushforwards; a blocking rule that depends on the input law is not admitted.

**Hypothesis 12.4 (Depth-Markov hierarchy).** The generative law factors top-down as

$$\mathbb{P}_\theta(do, d\mathbf{Y} \mid X_*) = \mathbb{P}_\theta^{\mathrm{obs}}(do \mid Z_0, X_*)\mathbb{P}_L(dY_L \mid X_*) \times \prod_{\ell=0}^{L-1} \mathbb{P}_R^\ell(dR_\ell \mid Y_{\ell+1}, X_*)\mathbb{P}_{\mathfrak{N}}^\ell(d\mathfrak{N}_\ell \mid Y_{\ell+1}, R_\ell, X_*)K_\downarrow^\ell(dZ_\ell \mid Y_{\ell+1}, R_\ell, \mathfrak{N}_\ell, X_*), \tag{12.7}$$

with the kernels composed from depth $L$ downward, and recognition factors along the same order,

$$\mathbb{Q}_\lambda(d\mathbf{Y} \mid o, X_*) = \mathbb{Q}_L(dY_L \mid o, X_*) \prod_{\ell=0}^{L-1} \mathbb{Q}_R^\ell(dR_\ell \mid Y_{\ell+1}, o, X_*) \times \mathbb{Q}_{\mathfrak{N}}^\ell(d\mathfrak{N}_\ell \mid Y_{\ell+1}, R_\ell, o, X_*)\mathbb{Q}_Z^\ell(dZ_\ell \mid Y_{\ell+1}, R_\ell, \mathfrak{N}_\ell, o, X_*). \tag{12.8}$$

Each depth-$\ell$ factor reads the complete next-depth state $Y_{\ell+1} = (Z_{\ell+1}, \mathfrak{N}_{\ell+1})$ and no deeper variable; this excludes direct dependence on depths $\ell+2, \dots, L$ given $(Y_{\ell+1}, R_\ell)$, leaves arbitrary dependence within a depth and between adjacent depths, and is not a mean-field assumption. In a declared top-down hierarchy $K_\downarrow^\ell$ is where downward influence enters, and the bottom-up coarse map belongs to recognition or to an induced pushforward, never to a second generative arrow at the same depth.

**Lemma 12.5 (Depth-Markov reversal of an adaptive tower).** Let $\mathbb{M}_0$ be a probability law on $\mathsf{Y}_0$, and generate the adaptive tower bottom-up through the common structural channels,

$$\mathbb{M}_0(dy_0) \prod_{\ell=0}^{L-1} C_{\ell+1\leftarrow\ell}^{\mathrm{struct}}(y_\ell; d\gamma_\ell, dy_{\ell+1}). \tag{12.9}$$

For every $\ell$, the conditional law of $R_\ell$ given all variables $(Y_{\ell+1}, R_{\ell+1}, \dots, Y_L)$ of greater depth has a version that depends only on $Y_{\ell+1}$, and the conditional law of $Y_\ell$ given $R_\ell$ and those variables has a version that depends only on $(R_\ell, Y_{\ell+1})$. Hence (12.9) factors as in (12.7) without its observation factor, with $\mathbb{P}_R^\ell$, $\mathbb{P}_{\mathfrak{N}}^\ell$, and $K_\downarrow^\ell$ selected versions of these conditional laws. The same holds for the recognition input $\mathbb{Q}_0(dy_0 \mid o, X_*)$, for the selected posterior slice, and for the generative input $\mathbb{P}_0(do, dy_0 \mid X_*)$ whenever its observation kernel reads only $Z_0$, so Hypothesis 12.4 holds automatically for an adaptive tower. These factors are reversal, or bridge, kernels in the sense of (11.41), not independently chosen structural channels, and when the parent coordinates are a function of the child coordinates and the structural mark, $K_\downarrow^\ell$ is supported on the child configurations compatible with the realized parent. [ESTABLISHED]

*Proof.* For each $\ell$, the variables $(R_{\ell+1}, Y_{\ell+2}, \dots, Y_L)$ are produced from $Y_{\ell+1}$ by kernels that read neither the observation nor any of $(Y_0, R_0, \dots, Y_\ell, R_\ell)$, so they are conditionally independent of the latter given $Y_{\ell+1}$, which gives both version statements. Regular conditional laws exist on standard Borel spaces [49], and successive disintegration from depth $L$ downward gives the factorization. A conditional law of $Z_\ell$ given a value of a function of $(Z_\ell, R_\ell)$ charges only the preimage of that value. □

**Proposition 12.6 (The next-depth graph datum cannot be dropped).** If the depth-$\ell$ factors of an adaptive tower read only $(Z_{\ell+1}, R_\ell)$, the product $\mathbb{P}^{\mathrm{red}}$ of the corresponding selected conditional laws of the tower law $\mathbb{P}^{\mathrm{tow}}$ of (12.9) differs from it by

$$D_{\mathrm{KL}}\big(\mathbb{P}^{\mathrm{tow}} \| \mathbb{P}^{\mathrm{red}}\big) = \sum_{\ell=0}^{L-1} I_{\mathbb{P}^{\mathrm{tow}}}\big(Y_\ell; \mathfrak{N}_{\ell+1} \mid Z_{\ell+1}, R_\ell\big). \tag{12.10}$$

The cost vanishes exactly when every $\mathfrak{N}_{\ell+1}$ is conditionally independent of $Y_\ell$ given $(Z_{\ell+1}, R_\ell)$. For a deterministic event pushforward it is positive as soon as two event laws in the support, at the same $(Z_{\ell+1}, R_\ell)$, are pushed to different coarse event laws: for four fine vertices with the partition $\{0,1\} \mapsto I$, $\{2,3\} \mapsto J$, constant coordinates, trivial links and marks, and $\eta_0$ equal with probability $1/2$ each to the unit mass on a copy $2 \to 0$ or on a copy $1 \to 0$, the pushed $\eta_1$ is the unit mass on $(I, J)$ or on $(I, I)$, and the divergence is $\log 2$. [ESTABLISHED]

*Proof.* By Lemma 12.5 the reduced product replaces each conditional law of $Y_\ell$ given $(R_\ell, Y_{\ell+1})$ by its conditional law given $(R_\ell, Z_{\ell+1})$, and the chain rule Lemma 3.12 turns each replacement into the displayed conditional mutual information. In the witness the joint law of $(\mathfrak{N}_0, \mathfrak{N}_1)$ has two atoms of mass $1/2$ while the reduced product is the product of their marginals, which puts mass $1/2$ on pairs that violate the pushforward. □

**Theorem 12.7 (Exact hierarchical decomposition).** Assume Hypothesis 12.4, which holds for an adaptive tower by Lemma 12.5, together with common domination and the integrability needed for the finite log-density form. Then

$$\begin{aligned}\mathcal{F}_{\mathrm{tower}} = \; & -\mathbb{E}_{\mathbb{Q}_\lambda} \log p_\theta^{\mathrm{obs}}(o \mid Z_0, X_*) + D_{\mathrm{KL}}(\mathbb{Q}_L \| \mathbb{P}_L) \\ & + \sum_{\ell=0}^{L-1} \mathbb{E}_{\mathbb{Q}_\lambda}\big[D_{\mathrm{KL}}(\mathbb{Q}_R^\ell \| \mathbb{P}_R^\ell) + D_{\mathrm{KL}}(\mathbb{Q}_{\mathfrak{N}}^\ell \| \mathbb{P}_{\mathfrak{N}}^\ell) + D_{\mathrm{KL}}(\mathbb{Q}_Z^\ell \| K_\downarrow^\ell)\big],\end{aligned} \tag{12.11}$$

where each divergence compares the recognition and generative factors of (12.8) and (12.7) at their common conditioning variables, the recognition factor also reading $o$, and is averaged over those variables under $\mathbb{Q}_\lambda$. These are the observation, top-prior, partition, graph, and parent–child contributions. [ESTABLISHED]

*Proof.* Substitute the two factorizations into the Radon–Nikodym ratio of (12.2) and apply the chain rule for conditional relative entropy, Lemma 3.12 and Cover and Thomas [21, Sec. 2.5], in its measure-level form, which subtracts no infinities. □

## 12.3 The scale-local two-channel graph energy

**Definition 12.8 (Transported divergences and row free energies).** For every copy $e : j \to i$ in $E_\ell$ and $x \in \{b, m\}$ put

$$D_e^{x,\ell} := D_{\mathrm{KL}}\big(q_{i,\ell}^x \| (\Theta_e^{x,\ell})_\# q_{j,\ell}^x\big), \tag{12.12}$$

where $(\Theta)_\#$ is the pushforward action $\widehat{\rho}_x(\Theta)$ on channel-$x$ laws; parallel copies carry separate divergences, and these are law-point energies in the sense of Section 3.8. A selectively communicating variant uses the role divergences (6.35) or a message channel instead, which changes the energy, and its recursive state keeps the interface and noise data and any private conditional law or memory that later kernels read (Section 6.5, Proposition 6.26). With $E_\ell(i)$ the copies received by $i$, the occupancies and rows of an event law are

$$\alpha_i^{x,\ell} := \sum_{e \in E_\ell(i)} \eta_e^{x,\ell}, \qquad \beta_e^{x,\ell} := \frac{\eta_e^{x,\ell}}{\alpha_i^{x,\ell}} \quad (e \in E_\ell(i),\ \alpha_i^{x,\ell} > 0), \tag{12.13}$$

with pair rows $\beta_{ij}^{x,\ell} = \sum_{e:j\to i} \beta_e^{x,\ell}$. The rows are the conditional source laws of whichever event law is displayed, generative, posterior, or recognition; a row of zero occupancy is not determined by the joint law. With normalized copy priors $\pi_i^{x,\ell}$ on $E_\ell(i)$ and temperatures $\tau_{i,\ell}^x > 0$, the row free energies in energy units are

$$\mathcal{F}_{i,\ell}^{\mathrm{att},x} = \sum_{e \in E_\ell(i)} \beta_e^{x,\ell} D_e^{x,\ell} + \tau_{i,\ell}^x D_{\mathrm{KL}}\big(\beta_i^{x,\ell} \| \pi_i^{x,\ell}\big), \tag{12.14}$$

and their dimensionless form is $\mathcal{F}_{i,\ell}^{\mathrm{att},x}/\tau_{i,\ell}^x$.

**Proposition 12.9 (Gibbs rows).** With $\mathcal{Z}_i^{x,\ell} = \sum_{e\in E_\ell(i)} \pi_e^{x,\ell} e^{-D_e^{x,\ell}/\tau_{i,\ell}^x}$, assume $\mathcal{Z}_i^{x,\ell} > 0$, equivalently that at least one edge of positive prior mass has finite energy. On the closed simplex of rows supported on $\{e : \pi_e^{x,\ell} > 0\}$, with $0 \cdot (+\infty) = 0$, the extended-real identity is $\mathcal{F}_{i,\ell}^{\mathrm{att},x} = \tau_{i,\ell}^x D_{\mathrm{KL}}(\beta_i^{x,\ell} \| \beta_i^{x,\ell,*}) - \tau_{i,\ell}^x \log \mathcal{Z}_i^{x,\ell}$ with

$$\beta_e^{x,\ell,*} = \frac{\pi_e^{x,\ell} e^{-D_e^{x,\ell}/\tau_{i,\ell}^x}}{\mathcal{Z}_i^{x,\ell}}, \tag{12.15}$$

so the minimizer is unique and the minimum is $-\tau_{i,\ell}^x \log \mathcal{Z}_i^{x,\ell}$, with zero optimal mass on every infinite-energy edge. If $\mathcal{Z}_i^{x,\ell} = 0$, every prior-supported row has infinite free energy, and the Gibbs formula defines no probability row or finite minimizer. When all prior-supported energies are finite, the optimum has positive mass on every prior-supported edge. The optimal rows move with the transported divergences, and omitting the row relative entropy changes the objective. The dimensionless form is the row free energy (6.7) with prior $\pi_i^{x,\ell}$ and the fixed energies $D_e^{x,\ell}$ in place of the expected divergences, with the same minimizer (6.8) and minimum $-\log \mathcal{Z}_i^{x,\ell}$. The energy-unit form with $\tau_{i,\ell}^x \neq 1$ is not an independently weighted sector of one evidence bound (Proposition 6.2), and the composite potential below inherits this scaling. [ESTABLISHED]

*Proof.* Expanding $D_{\mathrm{KL}}(\beta\|\beta^*)$ with (12.15) gives $\sum_e \beta_e \log(\beta_e/\pi_e) + \sum_e \beta_e D_e/\tau + \log \mathcal{Z}$, which is the stated identity after multiplication by $\tau$. Relative entropy is nonnegative and vanishes only at $\beta = \beta^*$. A row charging an infinite-energy edge has both sides equal to $+\infty$; otherwise the expansion uses only finite-energy edges. If $\mathcal{Z} = 0$, every edge of positive prior mass has infinite energy, giving the stated exceptional case. □

**Definition 12.10 (Composite multiscale potential).** If a scale-local prior is a normalized Gibbs law containing the self terms and $\sum_{e\in E_\ell}(\eta_e^{b,\ell} D_e^{b,\ell} + \eta_e^{m,\ell} D_e^{m,\ell})$, those terms and the global normalizer belong to the tower variational free energy. If they are appended to a fixed population free energy $\mathcal{F}_{\mathrm{pop}}$ of Chapter 5 instead, the resulting objective is the composite potential

$$\mathcal{J}_{\mathrm{comp}} = \mathcal{F}_{\mathrm{pop}} + \sum_\ell \sum_i (\alpha_i^{b,\ell} \mathcal{F}_{i,\ell}^{\mathrm{att},b} + \alpha_i^{m,\ell} \mathcal{F}_{i,\ell}^{\mathrm{att},m}) + \sum_\ell \mathcal{F}_{\mathrm{part}}^\ell + \sum_\ell \mathcal{F}_{\mathrm{cross}}^\ell + \sum_\ell \mathcal{J}_{\mathrm{link}}^\ell, \tag{12.16}$$

with the partition and cross-scale terms of (12.17) and (12.33) and a declared link regularizer $\mathcal{J}_{\mathrm{link}}^\ell$. It is a descent potential and carries no evidence interpretation of its own.

## 12.4 Partition variables and event pushforwards

Attention rows do not create a parent random variable. A hierarchy meant to emerge under the same variational principle puts the structural mark $R_\ell$ inside both the generative and the recognition law.

**Definition 12.11 (Partition factor).** A finite-temperature partition term has the form

$$\mathcal{F}_{\mathrm{part}}^\ell = \mathbb{E}_{\mathbb{Q}_R^\ell}\big[H_{\mathrm{block}}^\ell(R_\ell; Z_\ell, \mathfrak{N}_\ell)\big] + \tau_R^\ell D_{\mathrm{KL}}\big(\mathbb{Q}_R^\ell \| \mathbb{P}_R^{\ell,\circ}\big), \tag{12.17}$$

and enters the variational free energy only when the normalized Gibbs factor that it represents is part of the generative law. The block action $H_{\mathrm{block}}^\ell$ may trade internal transported mismatch against boundary retention, evaluator closure, and residual cost. No canonical partition selector follows from the global free energy alone (Open Problem 9.19).

### 12.4.1 A random partition catalog

The structural mark may be random with a finite catalog of partitions, or membership configurations, of $V_\ell$; each label fixes the branch carrier of the next depth, $X$ stays outside the channel, no edge is created, and the two normalizations of the catalog are distinct generative models.

**Definition 12.12 (Finite partition catalog).** Let $\mathfrak{G}$ be a finite nonempty catalog of values of $R_\ell$. For each $\gamma \in \mathfrak{G}$ declare a finite nonempty carrier $\mathsf{Y}_\gamma$, a strictly positive probability reference $\nu_\gamma$, a positive raw factor $\widetilde{p}_\gamma$, a positive weight $\mathsf{a}_\gamma$, and a normalized, possibly vanishing, likelihood $L_\gamma(o \mid y)$ on a common finite observation alphabet, all fixed generative data independent of recognition. On

the disjoint union of the carriers put

$$\mathcal{Z}_\gamma = \sum_y \nu_\gamma(y)\widetilde{p}_\gamma(y), \qquad p_\gamma(y) = \frac{\nu_\gamma(y)\widetilde{p}_\gamma(y)}{\mathcal{Z}_\gamma}, \qquad \bar{\mathsf{a}}_\gamma = \frac{\mathsf{a}_\gamma}{\sum_\delta \mathsf{a}_\delta}. \tag{12.18}$$

The hierarchically normalized and the globally normalized joints are

$$\mathbb{P}_{\mathrm{h}}(\gamma, y, o) = \bar{\mathsf{a}}_\gamma p_\gamma(y) L_\gamma(o \mid y), \tag{12.19}$$

$$\mathbb{P}_{\mathrm{f}}(\gamma, y, o) = \frac{\mathsf{a}_\gamma \nu_\gamma(y)\widetilde{p}_\gamma(y) L_\gamma(o \mid y)}{\mathcal{Z}_{\mathrm{all}}}, \qquad \mathcal{Z}_{\mathrm{all}} = \sum_\delta \mathsf{a}_\delta \mathcal{Z}_\delta. \tag{12.20}$$

**Proposition 12.13 (Branch normalization changes label weights).** Both joints are normalized and have the same conditional branch laws, while the label weight of the globally normalized joint is $\mathsf{a}_\gamma \mathcal{Z}_\gamma / \mathcal{Z}_{\mathrm{all}}$. The label weights of the two joints are the softmax weights of Definition 9.22 with effective log masses $\log \mathsf{a}_\gamma$ and $\log \mathsf{a}_\gamma + \log \mathcal{Z}_\gamma$ respectively. For the same positive weights the joints coincide exactly when all $\mathcal{Z}_\gamma$ coincide. Replacing $\widetilde{p}_\gamma$ by $\mathrm{c}_\gamma \widetilde{p}_\gamma$ with $\mathrm{c}_\gamma > 0$ leaves the hierarchical joint unchanged and reweights the global one by $\mathrm{c}_\gamma$, and a common factor cancels in both. For a strictly positive probability reference $\nu_{\mathfrak{G}}$ on the catalog write $\mathsf{a}_\gamma = \nu_{\mathfrak{G}}(\gamma) e^{-H_\gamma^{\mathrm{lab}}}$ and $\widetilde{p}_\gamma = e^{-H_\gamma}$; every positive family of weights has this form. Fix positive evidence $M_o = \sum_{\gamma,y} \mathbb{P}(\gamma, y, o)$ and a recognition law $\mathbb{Q}(\gamma, y) = \mathbb{Q}^{\mathrm{lab}}(\gamma)\mathbb{Q}_\gamma(y)$, with $\mathbb{Q}_\gamma$ defined where $\mathbb{Q}^{\mathrm{lab}}(\gamma) > 0$. Then, with $\mathcal{F}_o = D_{\mathrm{KL}}(\mathbb{Q} \| \mathbf{\Pi}_o) - \log M_o$,

$$\begin{aligned} \mathcal{F}_{\mathrm{f},o}(\mathbb{Q}) &= D_{\mathrm{KL}}(\mathbb{Q}^{\mathrm{lab}} \| \nu_{\mathfrak{G}}) + \mathbb{E}_{\mathbb{Q}^{\mathrm{lab}}} H_\gamma^{\mathrm{lab}} + \log \mathcal{Z}_{\mathrm{all}} \\ &\quad + \sum_\gamma \mathbb{Q}^{\mathrm{lab}}(\gamma) \big[ D_{\mathrm{KL}}(\mathbb{Q}_\gamma \| \nu_\gamma) + \mathbb{E}_{\mathbb{Q}_\gamma} H_\gamma - \mathbb{E}_{\mathbb{Q}_\gamma} \log L_\gamma(o \mid Y) \big], \end{aligned} \tag{12.21}$$

$$\mathcal{F}_{\mathrm{h},o}(\mathbb{Q}) = \mathcal{F}_{\mathrm{f},o}(\mathbb{Q}) - \log \mathcal{Z}_{\mathrm{all}} + \log \sum_\delta \mathsf{a}_\delta + \sum_\gamma \mathbb{Q}^{\mathrm{lab}}(\gamma) \log \mathcal{Z}_\gamma, \tag{12.22}$$

with the convention $0 \log 0 = 0$; a recognition-positive atom with zero likelihood gives $+\infty$. Each expansion contains one full recognition entropy, decomposed by its conditional chain, and no further correlation-entropy term. [ESTABLISHED]

*Proof.* Sum either joint over $o$, then $y$, then $\gamma$. Equality of the joints implies equality of the label marginals, hence $\mathcal{Z}_\gamma = \mathcal{Z}_{\mathrm{all}} / \sum_\delta \mathsf{a}_\delta$ for every $\gamma$; the converse is substitution. A branch rescaling changes $\mathcal{Z}_\gamma$ to $\mathrm{c}_\gamma \mathcal{Z}_\gamma$ and cancels in $p_\gamma$. The softmax statement is the label marginal of each joint written as the exponential of its log mass divided by the sum of these exponentials. The expansions follow by writing $\log \mathbb{P}$ with the displayed factors and taking the expectation under $\mathbb{Q}$; the two differ only in the normalizers, $\log \sum_\delta \mathsf{a}_\delta + \log \mathcal{Z}_\gamma$ against $\log \mathcal{Z}_{\mathrm{all}}$. □

Reference, raw factor, label weight, evidence, and posterior are distinct objects, and the score of either joint in a parameter is the complete labeled score of Proposition 9.23, which contains every varying generative factor, the $\mathcal{Z}_\gamma$ of the global model included.

### 12.4.2 Membership kernels and event pushforwards

**Definition 12.14 (Membership kernels).** Overlapping assignment is a membership Markov kernel

$$\mathfrak{m}_\ell(I \mid i, R_\ell) \geq 0, \qquad \sum_{I \in V_{\ell+1}} \mathfrak{m}_\ell(I \mid i, R_\ell) = 1, \tag{12.23}$$

of which a hard partition is the deterministic case $\mathfrak{m}_\ell(I \mid i, R_\ell) = \mathbf{1}\{R_\ell(i) = I\}$. Correlated endpoints use one joint endpoint kernel

$$\mathfrak{M}_\ell(I, J \mid i, j, R_\ell) \geq 0, \qquad \sum_{I,J} \mathfrak{M}_\ell(I, J \mid i, j, R_\ell) = 1, \tag{12.24}$$

whose product form requires a declared conditional independence of the two draws.

**Proposition 12.15 (Event pushforwards).** Fix $y \in \mathsf{Y}_\ell$ and a realized mark $\gamma \in \mathsf{R}_\ell$. If every vertex receives one membership draw and the draws of distinct vertices are conditionally independent at $(y, \gamma)$, the conditional coarse event law is

$$\eta^{x,\ell+1}_{IJ}(y, \gamma) = \sum_{\substack{i,j \in V_\ell \\ i \neq j}} \eta^{x,\ell}_{ij}(y)\mathfrak{m}_\ell(I \mid i, \gamma)\mathfrak{m}_\ell(J \mid j, \gamma) + \delta_{IJ} \sum_{i \in V_\ell} \eta^{x,\ell}_{ii}(y)\mathfrak{m}_\ell(I \mid i, \gamma), \tag{12.25}$$

where the loop term uses the diagonal endpoint kernel, since both endpoints of a loop follow one draw. With correlated endpoints,

$$\eta^{x,\ell+1}_{IJ}(y, \gamma) = \sum_{i,j \in V_\ell} \eta^{x,\ell}_{ij}(y)\mathfrak{M}_\ell(I, J \mid i, j, \gamma). \tag{12.26}$$

With the assignment marginal $C^R_\ell(y, d\gamma) = C^{\mathrm{struct}}_{\ell+1 \leftarrow \ell}(y; d\gamma \times \mathsf{Y}_{\ell+1})$ of the structural channel and an input law $\mathbb{M}$, the unconditional event mass averages the whole conditional expression,

$$\overline{\eta}^{x,\ell+1}_{IJ}[\mathbb{M}] = \sum_{i,j \in V_\ell} \int \eta^{x,\ell}_{ij}(y)\mathfrak{M}_\ell(I, J \mid i, j, \gamma) C^R_\ell(y, d\gamma)\mathbb{M}(dy). \tag{12.27}$$

All three coarse event laws sum to one. Conditional endpoint independence permits the product in (12.25) but not the factoring of the expectation of a random event mass and a random assignment, so (12.27) applies to all three input laws with no independence assumed between event and membership marks. The random $\eta_{\ell+1}(y, \gamma)$ is part of $\mathfrak{N}_{\ell+1}$ with rows recovered by (12.13), and its unconditional mean does not replace it. Loop copies arise at every depth after the first link lift (12.31); they are induced by blocking and are not a new support declaration, as in Proposition 11.29. For a hard partition the loop and product terms agree, and under the posterior bridge the pushforward is (11.43). Literal unit-weight membership in several parents has total mass greater than one; it is finite-measure replication, not a probability pushforward. [ESTABLISHED]

*Proof.* Summing the first term of (12.25) over $I, J$ gives $\sum_{i \neq j} \eta_{ij}$ and the loop term gives $\sum_i \eta_{ii}$, which add to one; the correlated and unconditional cases use normalization of $\mathfrak{M}_\ell$, $\eta_\ell$, $C^R_\ell$, and $\mathbb{M}$.

The expectation of a product of dependent random quantities is not the product of expectations. A vertex with unit weight in $k$ parents contributes its mass $k$ times. □

**Definition 12.16 (Chapman–Kolmogorov composition).** Nested assignments compose by

$$\mathfrak{m}_{\ell+2\leftarrow\ell}(A \mid i, R_\ell, R_{\ell+1}) = \sum_I \mathfrak{m}_{\ell+1}(A \mid I, R_{\ell+1})\mathfrak{m}_\ell(I \mid i, R_\ell), \tag{12.28}$$

$$\mathfrak{M}_{\ell+2\leftarrow\ell}(A, A' \mid i, j, R_\ell, R_{\ell+1}) = \sum_{I,J} \mathfrak{M}_{\ell+1}(A, A' \mid I, J, R_{\ell+1})\mathfrak{M}_\ell(I, J \mid i, j, R_\ell), \tag{12.29}$$

and the augmented structural channels by

$$C^{\mathrm{struct}}_{\ell+2\leftarrow\ell} = C^{\mathrm{struct}}_{\ell+2\leftarrow\ell+1} \circ C^{\mathrm{struct}}_{\ell+1\leftarrow\ell}, \tag{12.30}$$

where composition integrates every intermediate assignment, event law, and mark, and the root, holonomy, and boundary marks stay in the target of each composite. An independently declared multi-depth channel is admitted only when it coincides with the composite, and the composite is the exact multi-depth step only when the augmented state carries every mark that the next structural kernel reads.

## 12.5 Coarse links and what the parent retains

The links are not assumed flat: $\mathfrak{h}_\ell$ keeps the root-framed datum of Proposition 10.5 in every connected block with its joint law, and on a non-flat graph a reduced mean of the dressed marks must carry a declared residual, whose penalization to zero is a modeling preference, not a consequence of gauge covariance.

**Definition 12.17 (Coarse links under a hard partition).** Let the realized mark $R_\ell = \gamma$ be a hard partition whose blocks are replaced by their component-indexed meta-agents (10.6), so that each coarse vertex $I$ has one root $r_I$ and one rooted spanning tree. This excludes overlapping membership kernels, for which a copy would need a root transport into every block that contains one of its endpoints with positive membership; the event pushforwards of Proposition 12.15 hold for general membership kernels. The depth-$(\ell+1)$ multigraph $\mathfrak{I}_{\ell+1} = (V_{\ell+1}, E_{\ell+1}, {}^{-})$ has the coarse vertices, the same copies $E_{\ell+1} := E_\ell$ with the same reversal, and endpoints $e : J \to I$ whenever $e : j \to i$ with $j \in J$ and $i \in I$; its event law is $\eta_e^{x,\ell+1} := \eta_e^{x,\ell}$, and its links are the dressed marks of all copies,

$$\Theta_e^{x,\ell+1} := \widehat{\Theta}_e^{x,\ell} = \Theta_{I\leftarrow i}^{x,\ell}\Theta_e^{x,\ell}\big(\Theta_{J\leftarrow j}^{x,\ell}\big)^{-1}, \qquad e : j \to i,\ j \in J,\ i \in I, \tag{12.31}$$

internal as well as boundary, so that internal copies become loops at their block. The parent laws $q^x_{I,\ell+1}$ are expressed in the frame of the root $r_I$. The coarse multigraph is induced structure; no edge is created.

**Proposition 12.18 (Properties of the coarse links).** The coarse links are well typed and keep the whole link content of depth $\ell$. An internal tree copy has the identity link $1_G$. An internal non-tree copy carries the based holonomy of its fundamental loop, and for every closed walk $p$ in $I$ based at $r_I$ that traverses the copies $e_0, \dots, e_{n-1}$ in this order, $\mathrm{Hol}_I^{x,\ell}(p) = \Theta_{e_{n-1}}^{x,\ell+1} \cdots \Theta_{e_0}^{x,\ell+1}$. Inverse reversal persists, and the dressed law is the vertex law (2.28) with vertex coordinates $g_{r_I}$, so the divergences, rows, and row free energies of Definition 12.8 are defined and gauge invariant at depth $\ell+1$. On a loop the divergence $D_{\mathrm{KL}}(q_{I,\ell+1}^x \| (\Theta_e^{x,\ell+1})_\# q_{I,\ell+1}^x)$ vanishes on tree copies and, on a non-tree copy, exactly when the parent law is fixed by that holonomy, as in (10.10). The pair marginal $\sum_{e:J\to I} \eta_e^{x,\ell+1}$ is (12.25) for the hard partition, and no averaged group element replaces the parallel copies. The links depend on the trees, not only on the roots: changing the root of a fixed tree is a vertex gauge transformation at depth $\ell+1$, while replacing the tree changes $\Theta_{I\leftarrow i}^{x,\ell}$ to $\mathrm{Hol}_I^{x,\ell}(p_i)\Theta_{I\leftarrow i}^{x,\ell}$ for a based loop $p_i$ that depends on $i$, and a copy that leaves the tree changes its loop link from $1_G$ to the holonomy $\mathrm{Hol}_I^{x,\ell}(p)$ of its new fundamental loop $p$ and its loop divergence from $0$ to $D_{\mathrm{KL}}(q_{I,\ell+1}^x \| \mathrm{Hol}_I^{x,\ell}(p)_\# q_{I,\ell+1}^x)$. Whenever such a holonomy is not the identity, the trees are part of the coarse channel rather than a gauge choice. When $q_{I,\ell+1}^x$ and $q_{J,\ell+1}^x$ lie in the stabilized sector, the divergences $D_e^{x,\ell+1}$ do not depend on the trees. This link-level tree dependence complements the feature-level one of (10.15). [ESTABLISHED]

*Proof.* The dressed link of an internal copy is $1_G$ on a tree copy, since there $\Theta_e = \Theta_{I\leftarrow i}^{-1}\Theta_{I\leftarrow j}$, and is the fundamental-cycle holonomy (10.7) otherwise. Inserting $(\Theta_{I\leftarrow k})^{-1}\Theta_{I\leftarrow k}$ at each intermediate vertex $k$ of a closed walk turns the ordered product of the fine links into that of the coarse links. Reversal and the vertex law follow from (10.11), and relative entropy is invariant under a common bijective pushforward, which gives gauge invariance; a loop divergence vanishes exactly when $q_I = \mathrm{Hol}_I(p)_\# q_I$. A root change multiplies every transport of the block by one element, and a tree change multiplies each tree transport on the left by a based holonomy (Proposition 10.5), so the link of $e : j \to i$ becomes $\mathrm{Hol}_I(p_i)\Theta_e^{x,\ell+1}\mathrm{Hol}_J(p_j)^{-1}$, and when $\mathrm{Hol}_I(p_i)$ fixes $q_I$ and $\mathrm{Hol}_J(p_j)$ fixes $q_J$, invariance under the pushforward by $\mathrm{Hol}_I(p_i)^{-1}$ gives $D_{\mathrm{KL}}(q_I \| (\mathrm{Hol}_I(p_i)\Theta\mathrm{Hol}_J(p_j)^{-1})_\# q_J) = D_{\mathrm{KL}}(q_I \| \Theta_\# q_J)$. □

## 12.6 How a parent influences its children

**Proposition 12.19 (Downward coupling in a declared hierarchy).** Under Hypothesis 12.4 for a declared top-down hierarchy, a factorization such as

$$K_\downarrow^\ell(dZ_\ell \mid Y_{\ell+1}, R_\ell, \mathfrak{N}_\ell, X_*) = \prod_{I\in V_{\ell+1}} K_{\downarrow,I}^\ell\big(dZ_{I,\ell} \mid Z_{I,\ell+1}, \mathfrak{N}_{\ell+1}, R_\ell, \mathfrak{N}_\ell, X_*\big) \tag{12.32}$$

makes the parent belief and model presentations arguments of the children's conditional generative prior, and the corresponding term of (12.11) is

$$\mathcal{F}_{\mathrm{cross}}^\ell = \mathbb{E}_{\mathbb{Q}_\lambda} D_{\mathrm{KL}}\big(\mathbb{Q}_Z^\ell(\cdot \mid Y_{\ell+1}, R_\ell, \mathfrak{N}_\ell, o, X_*) \| K_\downarrow^\ell(\cdot \mid Y_{\ell+1}, R_\ell, \mathfrak{N}_\ell, X_*)\big). \tag{12.33}$$

The generative denominator in this cross-scale divergence reads the parent coordinates through $K_{\downarrow}^{\ell}$. The full tower objective can also couple parent coordinates to child structural marks through $\mathbb{P}_R^{\ell}$ and $\mathbb{P}_N^{\ell}$ in Hypothesis 12.4. A downward kernel that does not depend on the parent coordinates contributes no parent–child cross-derivative through this generative factor, and normalization of the downward kernel alone implies no coupling. When $K_{\downarrow}^{\ell}$ depends on the parent coordinates and the recognition parameterization transmits that dependence, the mixed derivative in a parent and a child parameter can be nonzero, so that optimization is coupled in both directions although the generative arrow is top-down; a universal claim of such coupling for a family requires a proved nonzero cross-derivative in that family. For a binary parent $b$ and child $y$ with product recognition factors $q_b(1) = u$ and $q_y(1) = v$, the mixed derivative of $\mathbb{E}_{q_b q_y}[-\log K_{\downarrow}(y \mid b)]$ in $(u, v)$ is $-\log[K_{\downarrow}(1 \mid 1)K_{\downarrow}(0 \mid 0)/(K_{\downarrow}(0 \mid 1)K_{\downarrow}(1 \mid 0))]$, which vanishes exactly when the odds ratio of the kernel is one. [ESTABLISHED]

*Proof.* The generative denominator in (12.33) is $K_{\downarrow}^{\ell}$. If that kernel does not read the parent coordinates, its negative log density has no parent argument; the expectation and recognition parameterization may still depend on the parent. For the binary case the expectation is bilinear in $(u, v)$, and its $uv$ coefficient is the displayed alternating sum of $-\log K_{\downarrow}(y \mid b)$ over the four configurations. □

This is the static sense in which a meta-agent influences its sub-agents without a generative factor reading a posterior; in an adaptive tower the same kernel is the reversal kernel of Lemma 12.5.

## 12.7 Closing the tower

A tower in which each depth's prior is supplied by the depth above must terminate, and within one finite acyclic joint its top cannot be constituted by the agents it organizes (Proposition 12.20). A same-time reciprocal structure among latents is admissible only as a globally normalized joint (Section 4.2, Proposition 4.6): a flat unanchored energy is improper, an anchored or non-flat energy with its global normalizer is a valid undirected model, and an energy appended without its normalizer is not an evidence bound. A same-time deterministic parent $Z_{\ell+1} = \chi_{\ell}(Z_{\ell})$ and an independently sampled parent that generates $Z_{\ell}$ are therefore not two free descriptions of one joint: one uses the hierarchy above, a deterministic pushforward, or a normalized undirected compatibility factor, and reciprocal same-time influence requires the last with its partition function. Supplement G works out the anchored Gaussian star and the convergence of its coordinate ascent.

### 12.7.1 A declared root

A truncated tower with a declared top prior $\mathbb{P}_L$ satisfies the hypotheses of Chapter 5 without amendment; its top is given, not constituted.

**Proposition 12.20 (A declared root is unavoidable).** Let a finite directed acyclic generative model be typed as in Chapter 4, so that every observed record is a child of at least one latent vertex. Then at least one latent has no parents, and its factor in the joint depends only on the parameters

and the structural data. No finite acyclic construction of this type has every prior constituted by other latents. [ESTABLISHED]

*Proof.* A finite directed acyclic graph admits a topological order, whose first element has no incoming edge. That element is not a record, because every record has a latent parent, so it is a latent, and its factor is conditioned on no other vertex. The typing hypothesis is needed: in the graph $o \to y$ with observed root $o$, the only latent has a parent. □

Anchored, globally normalized cyclic models remain available, but their anchors and partition functions are declared; parameter learning fits only their parameters.

### 12.7.2 Three loci of participation

What a cycle in the model was reaching for is partly available without one, at three loci, each a sequence of evaluations of fixed normalized joints.

**Remark 12.21 (The inference loop is block coordinate ascent).** In a mean-field recognition family, the coordinate update of a factor reads the current factors of every variable that shares a factor of the joint with it [12, Sec. 10.1.1]. In every connected hierarchical model in which the children's conditionals depend on their parents, parent and child updates therefore read each other; at a stationary point of the schedule the factor that acts as a child's prior in its own update is $\exp \mathbb{E}[\log K_{\downarrow}(z_{\text{child}} \mid \cdot)]$ under the parent's current factor, which is itself formed from the factors of all its children. The effective prior is thus formed from the population's current beliefs through a latent inside the joint. This is a schedule of updates against one fixed, acyclic, normalized joint, not participation in time, and no kernel reads a posterior, so Hypothesis 4.3 holds. [STANDARD]

**Proposition 12.22 (The aggregate recognition law is the optimal top prior).** Let episodes $n = 1, \dots, N_{\text{e}}$ be conditionally independent given $\theta$, each with its own top latent in a common space carrying the prior density $p_\theta$ relative to a fixed reference measure, let $\mathbb{Q}_n$ be the recognition law of the top latent of episode $n$, and let every other factor be free of $\theta$. The $\theta$-dependent part of the summed bound is

$$\sum_{n=1}^{N_{\text{e}}} \mathbb{E}_{\mathbb{Q}_n}[\log p_\theta] = N_{\text{e}} \mathbb{E}_{\bar{\mathbb{Q}}}[\log p_\theta], \qquad \bar{\mathbb{Q}} = \frac{1}{N_{\text{e}}} \sum_{n=1}^{N_{\text{e}}} \mathbb{Q}_n. \tag{12.34}$$

If $\bar{\mathbb{Q}}$ has a density $\bar{q}$ with finite entropy relative to the same reference, then $\mathbb{E}_{\bar{\mathbb{Q}}}[\log p] = \mathbb{E}_{\bar{\mathbb{Q}}}[\log \bar{q}] - D_{\text{KL}}(\bar{\mathbb{Q}} \| p)$ for every prior density $p$. Over an unrestricted family of prior densities the bound is maximized exactly at the aggregate recognition law $p = \bar{q}$, and within a restricted family its maximizers are the minimizers of $D_{\text{KL}}(\bar{\mathbb{Q}} \| p)$ over the family, when they exist. Each evaluation uses one fixed normalized joint, and updating $\theta$ at fixed recognition is a coordinate step of the two-argument objective [64], so Hypothesis 4.3 holds. [ESTABLISHED]

*Proof.* The first identity is linearity of expectation in the recognition law. The second is $\log p = \log \bar{q} - \log(\bar{q}/p)$ integrated against $\bar{\mathbb{Q}}$, and Gibbs' inequality gives the maximizers. □

This is the variational counterpart of empirical Bayes [70], whose optimum is the aggregate-posterior prior of Hoffman and Johnson [44] and Tomczak and Welling [77]; only the parameter of the declared prior family is fitted, by a slow-block update in the typing of Section 8.2.

**Definition 12.23 (Episode-indexed joint).** Episodes are indexed by $n = 0, 1, \dots$ with records $o_n$. In the parameter route each episode has its own normalized joint $\mathbb{P}_{\theta_n, V}$, and the parameter is updated by

$$\theta_{n+1} = \mathcal{U}_n\big(\theta_n, \mathbb{Q}^{(n)}_{V, o_n, X}\big), \tag{12.35}$$

where $\mathbb{Q}^{(n)}_{V,o_n,X}$ is the recognition coupling selected at episode $n$ and $\mathcal{U}_n$ is a declared map; within each episode the joint is fixed, and its generative kernels read $\theta_n$ and never a recognition law of the same episode. In the latent route one joint spans episodes $0, \dots, n_1$: episode $n$ has a tower state $\mathbf{Y}^{(n)}$ with parent coordinates $Z^{(n)}_{\mathrm{par}} = \chi^{\mathrm{ep}}(\mathbf{Y}^{(n)})$ for a measurable read-out $\chi^{\mathrm{ep}}$, and

$$\begin{aligned}\mathbb{P}^{\mathrm{ep}}\big(d\mathbf{Y}^{(0:n_1)}, do_{0:n_1} \mid X\big) = \mathbb{P}^{(0)}(d\mathbf{Y}^{(0)} \mid X) \prod_{n=1}^{n_1} \mathbb{P}^{(n \leftarrow n-1)}\big(d\mathbf{Y}^{(n)} \mid Z^{(n-1)}_{\mathrm{par}}, X\big) \\ \times \prod_{n=0}^{n_1} \mathbb{P}^{\mathrm{obs}}\big(do_n \mid \mathbf{Y}^{(n)}, X\big),\end{aligned} \tag{12.36}$$

with a declared root in $\mathbb{P}^{(0)}$; the parent latent of episode $n-1$ parameterizes the conditional law of the tower of episode $n$. The observation kernel is any Markov kernel and need not have a density: the tower state may carry a latent copy $\widetilde{o}_n$ of the record, observed through $\mathbb{P}^{\mathrm{obs}}(do_n \mid \mathbf{Y}^{(n)}, X) = \delta_{\widetilde{o}_n}(do_n)$, in which case the free energy of the joint is its extended form (12.1) and the decomposition (12.38) applies as displayed only for discrete records.

**Proposition 12.24 (Delayed feedback in the episode-indexed joint).** In the latent route, (12.36) is a normalized joint whose kernels read structural data and earlier latents only, so Hypothesis 4.3 holds; its graph is acyclic with every cross-episode edge pointing forward in $n$, so Proposition 12.20 still requires a declared root in episode $0$. The prior of the tower at episode $n+1$ given the records so far is the filtered parent law pushed through the transition,

$$\mathbb{P}^{\mathrm{ep}}\big(\mathbf{Y}^{(n+1)} \in \cdot \mid o_{0:n}\big) = \int \mathbb{P}^{(n+1 \leftarrow n)}(\cdot \mid z, X)\, \mathbb{P}^{\mathrm{ep}}\big(Z^{(n)}_{\mathrm{par}} \in dz \mid o_{0:n}\big), \tag{12.37}$$

so the children's prior at episode $n+1$ is set by the parent inferred from the population at episode $n$. For every recognition law $\mathbb{Q}$ on $\mathbf{Y}^{(0:n_1)}$, on the finite-density domain with observation densities $p^{\mathrm{obs}}$,

$$\begin{aligned}\mathcal{F}^{\mathrm{ep}}[\mathbb{Q}] = \sum_{n=0}^{n_1} \mathbb{E}_{\mathbb{Q}}\big[-\log p^{\mathrm{obs}}(o_n \mid \mathbf{Y}^{(n)}, X)\big] + D_{\mathrm{KL}}\big(\mathbb{Q}^{(0)} \| \mathbb{P}^{(0)}\big) \\ + \sum_{n=1}^{n_1} \mathbb{E}_{\mathbb{Q}} D_{\mathrm{KL}}\big(\mathbb{Q}(\cdot \mid \mathbf{Y}^{(0:n-1)}) \| \mathbb{P}^{(n \leftarrow n-1)}(\cdot \mid Z^{(n-1)}_{\mathrm{par}}, X)\big),\end{aligned} \tag{12.38}$$

where $\mathbb{Q}^{(0)}$ is the law of $\mathbf{Y}^{(0)}$ under $\mathbb{Q}$ and $\mathbb{Q}(\cdot \mid \mathbf{Y}^{(0:n-1)})$ the conditional law of $\mathbf{Y}^{(n)}$. In the parameter route, if $\mathcal{U}_n$ maximizes the bound of episode $n$ over the parameter at fixed $\mathbb{Q}^{(n)}_{V,o_n,X}$, the bound of episode $n$ at $\theta_{n+1}$ is at least its value at $\theta_n$. [ESTABLISHED]

*Proof.* Normalization is composition of Markov kernels, as in Proposition 4.4, and acyclicity follows because every kernel reads only earlier episodes. Given $Z^{(n)}_{\mathrm{par}}$, the tower of episode $n+1$ is independent of $(\mathbf{Y}^{(0:n)}, o_{0:n})$, because its kernel reads only $Z^{(n)}_{\mathrm{par}}$ and the records up to $n$ are generated before it; conditioning on $o_{0:n}$ and integrating out $Z^{(n)}_{\mathrm{par}}$ gives (12.37). The decomposition is the chain rule Lemma 3.12 applied along $n$ to the Radon–Nikodym ratio of $\mathbb{Q}$ against the joint. The last statement is the definition of a maximizer. □

**Proposition 12.25 (One joint for a declared parameter route).** Let the episode joints of the parameter route form a Markov kernel $\theta \mapsto \mathbb{P}_{\theta,V}(dy, do \mid X)$ on the parameter space, with evidence densities $p_{\theta,V}(o \mid X)$ relative to one reference $\nu$ on O and posterior slices $\mathbf{\Pi}_{\theta,V,o,X}$ jointly measurable in $(\theta, o)$; let the coupling $\mathbb{Q}^{\theta}_{V,o,X}$ selected by Definition 3.15 for the joint at $\theta$, so that $\mathbb{Q}^{(n)}_{V,o_n,X} = \mathbb{Q}^{\theta_n}_{V,o_n,X}$, be jointly measurable in $(\theta, o)$; and let each update $\mathcal{U}_n$ of (12.35) be measurable, so that $\mathcal{U}^{\mathrm{sel}}_n(\theta, o) := \mathcal{U}_n(\theta, \mathbb{Q}^{\theta}_{V,o,X})$ is measurable and the route reads $\theta_{n+1} = \mathcal{U}^{\mathrm{sel}}_n(\theta_n, o_n)$ from a declared $\theta_0$. Write $\theta_n(o_{0:n-1})$ for the parameter the route reaches after the records $o_{0:n-1}$, and let each realized record be admitted at that parameter. The lifted joint on the tower states $\mathbf{Y}^{(n)} = (\theta_n, y_n, \widetilde{o}_n)$, $n = 0, \dots, n_1$, and the records,

$$\mathbb{P}^{\mathrm{lift}}\big(d\mathbf{Y}^{(0:n_1)}, do_{0:n_1} \mid X\big) = \delta_{\theta_0}(d\theta_0) \prod_{n=0}^{n_1} \mathbb{P}_{\theta_n,V}(dy_n, d\widetilde{o}_n \mid X)\delta_{\widetilde{o}_n}(do_n) \prod_{n=1}^{n_1} \delta_{\mathcal{U}^{\mathrm{sel}}_{n-1}(\theta_{n-1},\widetilde{o}_{n-1})}(d\theta_n), \tag{12.39}$$

has four properties. (i) It is a latent route of Definition 12.23 with parent coordinates $Z^{(n)}_{\mathrm{par}} = (\theta_n, \widetilde{o}_n)$, transition $\mathbb{P}^{(n\leftarrow n-1)}(d\mathbf{Y}^{(n)} \mid z, X) = \delta_{\mathcal{U}^{\mathrm{sel}}_{n-1}(z)}(d\theta_n)\mathbb{P}_{\theta_n,V}(dy_n, d\widetilde{o}_n \mid X)$, observation kernel $\delta_{\widetilde{o}_n}(do_n)$, and root $\delta_{\theta_0} \otimes \mathbb{P}_{\theta_0,V}$, so Proposition 12.24 applies to it. (ii) Given $o_{0:n-1}$, the conditional law of $(y_n, \widetilde{o}_n)$ is the episode joint $\mathbb{P}_{\theta_n(o_{0:n-1}),V}$ of the parameter route, and the evidence density of the lifted joint relative to $\nu^{\otimes(n_1+1)}$ is the product of the episode evidences,

$$p^{\mathrm{lift}}(o_{0:n_1} \mid X) = \prod_{n=0}^{n_1} p_{\theta_n(o_{0:n-1}),V}(o_n \mid X). \tag{12.40}$$

(iii) Its posterior slice fixes $\widetilde{o}_n = o_n$ and $\theta_n = \theta_n(o_{0:n-1})$ and is the product $\bigotimes_n \mathbf{\Pi}_{\theta_n(o_{0:n-1}),V,o_n,X}$ on $y_{0:n_1}$, so the posterior law of $y_n$ reads $o_{0:n}$ only. (iv) For every recognition law $\mathbb{Q}$ on the tower states, the free energy (12.1) of the lifted joint is

$$\mathcal{F}^{\mathrm{lift}}[\mathbb{Q}] = \sum_{n=0}^{n_1} \Big[ -\log p_{\theta_n,V}(o_n \mid X) + D_{\mathrm{KL}}\big(\mathbb{Q}^y_n \| \mathbf{\Pi}_{\theta_n,V,o_n,X}\big) \Big] + \mathrm{TC}(\mathbb{Q}^y), \qquad \theta_n = \theta_n(o_{0:n-1}), \tag{12.41}$$

when $\mathbb{Q}$ fixes $\widetilde{o}_n = o_n$ and $\theta_n = \theta_n(o_{0:n-1})$ almost surely, and $+\infty$ otherwise, where $\mathbb{Q}^y$ is the law of $y_{0:n_1}$ under $\mathbb{Q}$, $\mathbb{Q}^y_n$ its marginal on $y_n$, each bracket is the extended free energy (5.16) of episode $n$ at the parameter the route reaches, and $\mathrm{TC}$ is the total correlation (5.1) across episodes. At

product laws, and in particular at the selected couplings $\bigotimes_n \mathbb{Q}^{\theta_n}_{V,o_n,X}$, the free energy is the sum of the episode bounds of the parameter route. [ESTABLISHED]

*Proof.* Each factor of (12.39) is a Markov kernel in the coordinates to its left: $\theta \mapsto \mathbb{P}_{\theta,V}$ by hypothesis, and the Dirac factors as kernels of measurable maps, $\mathcal{U}^{\mathrm{sel}}_{n-1}$ being the composition of $\mathcal{U}_{n-1}$ with the jointly measurable selection. Finite composition normalizes the product (Proposition 4.4), and reading the factors in the order of (12.36) gives (i); every kernel reads structural data and latent coordinates only, because the selection is evaluated at the latent copy $\widetilde{o}_{n-1}$ and not at a recognition law of the joint being scored. For (ii), integrate (12.39) against a bounded measurable $f(o_{0:n-1}, y_n, \widetilde{o}_n)$: the factors of episodes $n+1, \dots, n_1$ have unit mass, the Dirac factors identify $o_k = \widetilde{o}_k$ for $k < n$ and, by induction along the recursion, $\theta_n = \theta_n(o_{0:n-1})$, and Tonelli's theorem leaves the integral of $\int f(o_{0:n-1}, y, \widetilde{o}) \mathbb{P}_{\theta_n(o_{0:n-1}),V}(dy, d\widetilde{o} \mid X)$ against the law of $o_{0:n-1}$, which is the stated conditional law. Its record marginal is $p_{\theta_n(o_{0:n-1}),V}(o_n \mid X)\nu(do_n)$, and multiplying the conditional densities along $n$ gives (12.40); the product is measurable because the evidence density is jointly measurable and $o_{0:n-1} \mapsto \theta_n(o_{0:n-1})$ is a composition of measurable maps, and it is finite and positive because each factor is. For (iii), substitute the disintegration $\mathbb{P}_{\theta,V}(dy, d\widetilde{o} \mid X) = \mathbf{\Pi}_{\theta,V,\widetilde{o},X}(dy) p_{\theta,V}(\widetilde{o} \mid X)\nu(d\widetilde{o})$ into each factor and integrate out the Dirac coordinates: for bounded measurable $f$ of the tower states and $g$ of the records,

$$\begin{aligned} \mathbb{E}_{\mathbb{P}^{\mathrm{lift}}}[fg] = \int g(o) \Big[ \int f\big((\theta_n(o_{0:n-1}), y_n, o_n)_{n \le n_1}\big) \bigotimes_{n=0}^{n_1} \mathbf{\Pi}_{\theta_n(o_{0:n-1}),V,o_n,X}(dy_n) \Big] \\ \times\, p^{\mathrm{lift}}(o \mid X) \nu^{\otimes(n_1+1)}(do), \end{aligned} \tag{12.42}$$

and the bracket, measurable in $o$ by the joint measurability of the slices, is a version of the conditional expectation of $f$ given the records. For (iv), the posterior is carried by the set on which $\widetilde{o}_n = o_n$ and $\theta_n = \theta_n(o_{0:n-1})$ for every $n$; a $\mathbb{Q}$ that charges its complement is not absolutely continuous with respect to the posterior, and its relative entropy is $+\infty$. Otherwise disintegrate both laws over $y_{0:n_1}$ with Lemma 3.12(a): the conditional laws of the remaining coordinates given $y_{0:n_1}$ are the same Dirac mass under both, so $D_{\mathrm{KL}}(\mathbb{Q} \| \mathbf{\Pi}^{\mathrm{lift}}) = D_{\mathrm{KL}}(\mathbb{Q}^y \| \bigotimes_n \mathbf{\Pi}_{\theta_n,V,o_n,X})$, and the product-reference identity (5.2) splits the right side into $\mathrm{TC}(\mathbb{Q}^y) + \sum_n D_{\mathrm{KL}}(\mathbb{Q}^y_n \| \mathbf{\Pi}_{\theta_n,V,o_n,X})$. Adding $-\log p^{\mathrm{lift}}(o \mid X)$ from (12.40) gives (12.41), and the total correlation of a product law is zero. □

The route's bookkeeping with the parameter read from the law being scored is the composite that Hypothesis 4.3 excludes; the lifted joint evaluates the selection at the latent copy of the record instead, and its evidence is the sequential empirical-Bayes evidence of the route.

**Open Problem 12.26 (Accounting across episodes).** In the parameter route the bounds of different episodes refer to different joints and records. Within an episode a maximizing update is monotone (Proposition 12.24), and for an update that reads a coupling selected by a declared rule the lifted joint of Proposition 12.25 is one joint whose evidence is the product of the episode evidences and whose free energy is the sum of the episode bounds plus the cross-episode total correlation; an update that reads the recognition law being scored is excluded by Hypothesis 4.3. The lifted joint informs the parent of episode $n+1$ by the record of episode $n$ through its latent copy.

What is unproved is whether the label coordinates alone can carry the feedback. The problem is posed for label-sufficient latent routes of Definition 12.23: the episode model, with a finite label coordinate $y^{\mathrm{lab}}$ whose prior is the parameter, its private latents, and its record kernel, is that of the parameter route with the posterior slice as its recognition coupling and the aggregate update of constant step, $\theta_{n+1} = (1-\epsilon_{\mathrm{e}})\theta_n + \epsilon_{\mathrm{e}}\mathbf{\Pi}^{\mathrm{lab}}_{\theta_n,V,o_n,X}$ with $\mathbf{\Pi}^{\mathrm{lab}}$ the label marginal of the posterior slice; the tower state of episode $n$ consists of the episode model's latents and an auxiliary preference coordinate; each record is emitted from its own episode's latents by the record kernel; and the read-out $\chi^{\mathrm{ep}}$ reads the label and the preference coordinate only. The records either have one value whose density $p_V(o \mid y^{\mathrm{lab}}, X)$ given the label does not depend on the label, while the expected squared total-variation displacement of the label posterior from the prior, under the evidence at that prior, is bounded below uniformly over priors whose smallest weight is at least half that of $\theta_0$; or their likelihood vectors $(p_V(o \mid y^{\mathrm{lab}}, X))_{y^{\mathrm{lab}}}$ span a cone on which no nonzero polynomial vanishes. For such routes, with the per-episode discrepancy the total variation between the route's filtered label law $\mathbb{P}^{\mathrm{ep}}(y^{\mathrm{lab}}_{n+1} \in \cdot \mid o_{0:n})$ and $(1-\epsilon_{\mathrm{e}})\mathbb{P}^{\mathrm{ep}}(y^{\mathrm{lab}}_{n} \in \cdot \mid o_{0:n-1}) + \epsilon_{\mathrm{e}}\mathbb{P}^{\mathrm{ep}}(y^{\mathrm{lab}}_{n} \in \cdot \mid o_{0:n})$, the aggregate update applied to the route's own filtered law and label posterior, the problem is settled by a route, defined for every step, whose discrepancy is $o(\epsilon_{\mathrm{e}})$ uniformly over record histories and over episodes $n$ up to any fixed multiple of $1/\epsilon_{\mathrm{e}}$, or by a lower bound of order $\epsilon_{\mathrm{e}}$ on the supremum of the discrepancy over such histories and episodes for every such route. Whether a declared update intertwines the fine and parent evolutions is part of Open Problem 11.38; $\theta_n$ belongs to the slow block of Definition 8.2.

### 12.7.3 The remaining hierarchy question

The prepared parent of Section 8.3, the memory cost of partial refresh in Section 11.8.2, and the successive formation of Theorem 13.17 are constructions at stated hypotheses, not consequences of arbitrary descent.

**Open Problem 12.27 (Persistent multiscale hierarchy).** Starting from $c_*$ with a finite multigraph datum $\mathfrak{N}_0$, construct a normalized partition law and parent–child kernels such that descent of (12.1) produces persistent nested blocks, the pushforward rules (12.25) to (12.30) and the link lift (12.31) compose across depths, the parent family has a controlled closure residual in the sense of Proposition 11.22, and non-flat holonomy is kept or stabilized without imposing flatness. What holds: every candidate hierarchy has the exact accounting of Theorem 12.7; at one depth the fast–slow typing of Definition 8.2 and its conditional slow-manifold limit, Open Problem 8.7, are available; Theorem 13.17 forms nested blocks under a supplied energy hierarchy whose strengths are tied to the temperature; and Proposition 13.22 forms one level. What is unproved is that descent produces a second level that depends on the first. The problem is posed in the class of models at $c_*$ whose structural marks are drawn from a finite partition catalog with a partition factor inside the generative and the recognition law (Definitions 12.11 and 12.12), whose recognition coordinates descend (12.1) by a reversible diffusion at temperature $\mathrm{T}$ with rate-one redraws of the discrete marks from their conditional laws, as in Definition 13.21, and whose energies, precisions, prior ratios, hyperparameters, and mobilities do not depend on $\mathrm{T}$, as $\mathrm{T} \to 0$. It is settled by three results

together, or by a proof that the first of them cannot hold in this class. The first is a lineage of at least two levels, formed with probability bounded below from every initial law with a density and a finite second moment, whose second level meets two conditions. It depends on the first level: for a declared handoff time, a declared coalescence event whose block is a union of at least two first-level blocks and whose windows are measured from the handoff, one declared statistic of the handoff state (the probability of the event, of its occurrence by declared times, or of a declared block), and a handoff class fixed before the comparison that contains every first-level handoff and, with probability bounded below, the handoffs of histories without first-level formation, the event has probability bounded below after first-level formation, does not have probability tending to one uniformly over the handoff class at the declared windows, and the statistic exceeds its value after a counterfactual handoff in the class by a margin bounded below. It is a coherence episode rather than restored equilibrium or coincidental coherence: a restoring interaction, evaluated on the law of the coordinates given the realized marks or on their marginal over the marks and not on a conditional law held fixed by hand, holds the second-level block's coherence at a declared resolution, with the restoring rate times the squared resolution divided by $\mathrm{T}$ tending to infinity, on horizons short compared with the stationary time of entrance into the event from a set of stationary mass bounded below, or with a probability exceeding its stationary value by a margin bounded below. The second result is a nondegenerate intrinsic partition selector (Open Problem 9.19), and the third is a separated slow manifold on which depth-$(\ell+1)$ variables feed back consistently into depth-$\ell$ kernels. The negative resolution is a proof that in this class no second level produced by descent, apart from a supplied-scale sequence whose event occurs with probability tending to one uniformly over the handoff class at the declared windows, both meets the restoring and horizon conditions and exceeds its counterfactual statistic by such a margin.

## Summary for physicists

The multiscale free energy is that of one generative and one recognition law over the whole tower, not a sum of energies chosen scale by scale. Bayes reversal of a bottom-up structural channel conditions every depth on the whole next-depth state, graph datum included. Partitions enter as a normalized factor, normalized branch by branch or globally, and coarse links are the dressed links of all copies, so internal loops keep their holonomy. A parent can couple to child coordinates through a downward kernel that reads it and to child structural marks through the partition and graph factors. A finite acyclic tower has a declared root, and the agents participate in their own top through coordinate ascent, a learned aggregate prior, and delayed feedback. A sequence of episodes whose parameter update reads a coupling selected by a declared rule is itself one joint, with the product of the episode evidences as its evidence (Proposition 12.25).

# Part III

# Formation and Interpretation

# Chapter 13

# Formation Witnesses

Exact parents are normalized pushforwards of the fine law (Theorem 9.6) but in general not autonomous reduced dynamics (Theorem 11.39). This chapter asks whether, in one normalized experiment at $c_*$, coherent components can form in succession while every parent is activated from the complete preceding recognition law, retains every loop coordinate and source assignment, and learns its own copy of a model parameter from its own law. Two witnesses answer yes under explicit restrictions. Theorem 13.17 treats a triangle chain with a supplied hierarchy of interaction energies: at each finite depth $L$ the coherent components pass through $3 \cdot 2^{L-1} \to 2^{L-1} \to \cdots \to 1$ with probability tending to one, while the discarded conditional losses, parameter covectors, assignment histories, and continuous paths obey separate vanishing bounds. Proposition 13.22 treats a homogeneous four-cycle without supplied hierarchy, in which the source assignment decides which pairs form, with probability bounded below independently of the small parameter. Proposition 13.24 bounds what that mechanism can do on any finite graph: in the gated family of Definition 13.23, the stationary coherent-block partition is, up to the error of eq. (13.59), the partition into the components of a minimum edge cover, and each merge beyond such a cover costs a factor of the temperature. The proposition and the conjecture of the interpretation chapter (Proposition 14.4 and theorem 14.5) are posed on the four-cycle family of the second witness. Both witnesses are approximation statements under a stated preparation and regime, weaker than the exact intertwining that an autonomous parent dynamics would require (Open Problem 11.38). Both realize the general construction for the translation group $(\mathbb{R}, +)$ on scalar sample fibers: agents are local sections at $c_*$ (Definition 2.8), one common frame per agent lets one independent link translate both channels (Definition 2.11), and the free energy is that of Definition 5.6.

Letters of the notation table that would collide with chapter-local objects are replaced by the symbols of table 13.1; upright $\mathrm{J}, \mathrm{Q}, \mathrm{A}, \mathrm{T}, \mathrm{c}_L, \mathrm{C}_L$ are local scalars, the configuration is bold, and $\ell' < \ell$ whenever two levels occur.

**Scope.** For the finite-depth family this chapter declines the following claims: that attention or learning drives formation, since coherence is linear relaxation of the prepared mean, the source record saturates at the coherence scale (Remark 13.18), and learning is calibration dominated up to $\epsilon^{2\mathrm{Q}_1+20}$; that per-agent models are slow, since means and links have unit mobility and $\mathfrak{d}_i$ is a level-one row; that loop holonomy is of order one or nonabelian, since the group is $(\mathbb{R}, +)$ and the retained loops are thermally negligible; that there is an inter-agent model-channel record, since only the belief channel is observed; that parents are coarse agents with their own belief and

**Table 13.1:** Chapter-local symbols of the formation families, shared with the proof appendix.

| Symbol | Meaning |
|---|---|
| $L$, $N_{\triangle} = 2^{L-1}$, $N = 3N_{\triangle}$, $\ell, \ell'$ | depth, number of triangles, number of agents, level indices |
| $\mathbf{x} = (\bar{k}, \bar{m}, \omega)$, $\mathbf{X}_t$ | configuration of law presentations (belief means, model means, link translations) and its process |
| $\mathfrak{d}_i = \bar{k}_i - \bar{m}_i$, $\mathfrak{D}$, $r_e$ | belief–model discrepancy, its row matrix, edge residual |
| $\mathfrak{h}_q$, $h_q$, $O_{\triangle}$ | oriented triangle holonomy, normalized loop coordinate, contrast rows |
| $\mathcal{U}, \mathcal{H}, \mathcal{K}$; $W_\ell, \mathcal{U}_\ell, M_\ell, P_\ell$ | normal, loop, and gauge spaces; retained spaces and projections |
| $H$, $H_\ell$, $\mathrm{Q}_\ell$ | stiffness, its Schur complement on $M_\ell$, orthonormal flag columns |
| $\mathrm{J}_\ell$, $\mathrm{Q}_\ell$, A, T, $t_*$ | energy strengths, their exponents, prepared amplitude, temperature, horizon |
| $\sigma_{\mathrm{ev}}^2$, $R_{\mathrm{ev}}$, $\nu_i^*$ | evaluator variance, precision of the evaluator restriction cost, optimal evaluator coupling |
| $\vartheta$, $\Omega_0(\vartheta)$, M, $\mathsf{m} = \mathrm{TM}$ | learned scalar, prior precision, complete Fisher metric and its scaled form |
| $\kappa$, $z_\star$, $o_{\mathcal{K}}$, $o_{\mathrm{cal}}$, o | root observation precision, root record, root and calibration observations, observed records |
| $S \in \mathcal{J}$, $\mathcal{N}_i^{\triangle}$, $i_0$, $\pi_{ij}(\vartheta)$, $\pi_S(\vartheta)$ | complete source assignment, internal source palette, distinguished receiver, source priors |
| $\chi$, $\phi_S$, $\delta$, $B_0$, $B_*$ | source function, source potential, record coefficient, and its two bounds |
| $\beta_S$, $\beta_{SS'}^{\ell}$, $\bar{\beta}_{jk}$ | complete posterior source row, inherited proposal row, averaged fast row |
| $q_\vartheta^\star$, $q_\vartheta^{\star,\ell}$, $\mathbf{\Pi}_{\vartheta,\mathsf{o}}$, $q_t^\ell$ | restricted optimizer and its level-$\ell$ push, physical posterior, recognition density |
| $C_{\ell \leftarrow \ell'}$, $\mathcal{D}_e^{\ell,\mu}$, $\Delta_{\ell':\ell}$ | marked channel, inherited edge statistic, discarded conditional loss |
| $\mathcal{F}_{\mathrm{T}}$, $\mathrm{C_T}$, $\mathrm{E}_\vartheta$, $U_{\vartheta,S}$, $V_{\ell,S}$, $\psi_{\ell,S}$, $g_\ell$ | scaled restricted free energy, its offset, quadratic and sheet energies, parent potential and tilt, parameter covector |
| $\mathscr{L}_\vartheta^\ell$ | inference generator at level $\ell$ |
| $t_\ell$, $I_\ell$, $t_\ell^{\mathrm{act}}$, $k_L$, $\eta_L$, $\varsigma_\epsilon$, $r_\epsilon$ | window start, window, activation time, entrance constant, tolerance, $1 + \log(1/\epsilon)$, parameter remainder |
| $\mathrm{c}_L$, $\mathrm{C}_L$, $\mathsf{j}_m$ | positive constants depending on $L$ and the fixed data; proposal marks |
| $\mathrm{B}_{sr}$, $\mathsf{j}_\ell$, $\mathfrak{Y}_\ell$, $\mathfrak{y}_\ell$, $\mathfrak{x}_\ell$, $\varrho_s$, $\mathrm{Q_T}$, $\mathrm{Q}_\Delta$ | flag blocks of the residual rows, least eigenvalue of a diagonal Gram block, forcing map and core coefficient of an interface, current-layer transient, strength ratio, exponents of the thermal ratio and of a loss target |
| $(V, E)$, $\mathfrak{n}$, $F$, $\mathrm{cc}(F)$, $\mathrm{att}(S)$ | graph of the gated family, its matching number, an edge set, the number of components of $(V, F)$, the edge set attended under $S$ |
| $\mathbf{a}_e$, $\mathbf{a}_F$, $W_F$, $\mathrm{s}_F$ | residual pair map of an edge, its stack over $F$, the kernel of the stack, the least positive singular value of the stack |
| $H_S$, $\mathrm{k}_S$, $\mathrm{h}_S$, $\Omega$, $Z_S$, $\mathrm{z}_S$, $V_\pi$, $V_{\min}$ | sheet stiffness, its rank and least positive eigenvalue, reference precision, sheet weight and its constant, source-marginal and gated energies |
| $\mathrm{Coh_a}(\mathbf{x})$, $\mathfrak{P}_{\mathrm{a}}(\mathbf{x})$, a | coherent edges and coherent blocks of a configuration; resolution |

model fibers, since a parent removes only normal coordinates; that source learning creates the supplied contrasts; that the doubling recursion $\mathrm{Q}_\ell = 2\mathrm{Q}_{\ell+1} + 20$ is required, as distinct from the doubly exponential separation of inner base $4/3$ that Proposition 13.15 shows the Gaussian core to require; that the continuous path laws converge in total variation or are exactly lumpable; and that the construction saves computation or uses depth-uniform resources. For the four-cycle it declines generic formation, formation of a matching from every initial assignment, a unique block chosen by the invariant full law, all-slow parents, exact closure under arbitrary source switching, permanent membership, repeated coalescence, and a physical posterior conditioned on the episode. The corresponding constructions are Open Problems 13.16, 13.20 and 13.25. [NOT-CLAIMED]

## 13.1 The pointwise carrier and its restoring flag

**Definition 13.1 (Triangle-chain carrier).** Fix an integer $L \geq 2$, set $N_\triangle = 2^{L-1}$, and arrange $N_\triangle$ disjoint triangles in a chain, with one bridge between each consecutive pair. Each triangle is cyclically oriented and every bridge has one fixed orientation, so the graph has $N = 3N_\triangle$ vertices (the agents) and $4N_\triangle - 1$ oriented edge copies. Bridge number $j$, between triangles $j$ and $j+1$, has level $\ell(j) = 2 + v_2(j)$, where $v_2(j)$ is the exponent of two dividing $j$; triangle edges have level one. Energy strengths $\mathrm{J}_1 \geq \cdots \geq \mathrm{J}_L > 0$ are attached to the levels. The group $(\mathbb{R}, +)$ acts by translation on two scalar sample fibers at $c_*$; fix $0 < \sigma_{\mathrm{ev}}^2 < 1$. A configuration and its law presentations are

$$\begin{aligned} \mathbf{x} = (\bar{k}, \bar{m}, \omega) \in \mathbb{R}^{2N+4N_\triangle-1} = \mathbb{R}^{10N_\triangle-1}, \qquad & \Theta_e(k) = k + \omega_e, \\ q_i^b(\cdot \mid \mathbf{x}) = \mathcal{N}(\bar{k}_i, 1), \qquad & q_i^m(\cdot \mid \mathbf{x}) = \mathcal{N}(\bar{m}_i, \sigma_{\mathrm{ev}}^2), \\ \mathfrak{d}_i(\mathbf{x}) = \bar{k}_i - \bar{m}_i, \qquad & r_e(\mathbf{x}) = \bar{k}_i - \bar{k}_j - \omega_e \quad (e : j \to i), \end{aligned} \tag{13.1}$$

with $\omega_{\bar{e}} = -\omega_e$. A passive frame change $g_i$ sends $\bar{k}_i, \bar{m}_i, \omega_e$ to $\bar{k}_i + g_i$, $\bar{m}_i + g_i$, $\omega_e + g_i - g_j$, and re-expresses reference centers and observation records in the same way. For triangle $q$ put $\mathfrak{h}_q = \sum_{e \in q} \omega_e$ and $h_q = \mathfrak{h}_q/\sqrt{3}$, and let $O_\triangle = \begin{pmatrix} 1/\sqrt{2} & -1/\sqrt{2} & 0 \\ 1/\sqrt{6} & 1/\sqrt{6} & -2/\sqrt{6} \end{pmatrix}$. Let $B_1$ stack the $N$ rows $\mathfrak{d}_i$ and the $2N_\triangle$ internal contrast rows $O_\triangle r_q$, where $r_q$ is the vector of the three residuals of triangle $q$; for $\ell \geq 2$ let $B_\ell$ stack the bridge residuals of level $\ell$, and let $B$ stack all of them. With the oriented incidence operator $(\partial t)_e = t_i - t_j$,

$$\begin{aligned} \mathcal{K} &= \{(t, t, \partial t) : t \in \mathbb{R}^N\}, \quad \mathcal{H} = \operatorname{span}\{\text{normalized loop covectors}\}, \\ \mathcal{U} &= \operatorname{range} B^\top, \qquad\qquad H = B^\top \operatorname{diag}(\mathrm{J}_\ell I) B + \mathrm{J}_1 P_{\mathcal{H}}. \end{aligned} \tag{13.2}$$

The retained spaces and their Euclidean orthogonal projections are

$$W_\ell = \ker(B_1, \ldots, B_\ell), \qquad \mathcal{U}_\ell = W_\ell \cap \mathcal{U}, \qquad M_\ell = \mathcal{U}_\ell \oplus \mathcal{H}, \qquad P_\ell : \mathbb{R}^{10N_\triangle-1} \to W_\ell, \qquad P_0 = I. \tag{13.3}$$

**Lemma 13.2 (Carrier algebra).** In Definition 13.1, $\mathfrak{d}_i$ and $r_e$ are invariant under passive frame changes, and every triangle satisfies

$$|r_q|^2 = |O_\triangle r_q|^2 + h_q^2, \tag{13.4}$$

so the loop energy belongs to the original sum of edge mismatches. The rows of $B$ are independent, $\operatorname{rank} B = 6N_\triangle - 1$, the normalized loop covectors are orthonormal and orthogonal to every row of $B$, and $\mathbb{R}^{10N_\triangle - 1} = \mathcal{U} \oplus \mathcal{H} \oplus \mathcal{K}$ is an orthogonal direct sum. The stiffness has kernel $\mathcal{K}$ and $\mathbf{x}^\top H \mathbf{x} = \mathrm{J}_1 \sum_i \mathfrak{d}_i(\mathbf{x})^2 + \sum_e \mathrm{J}_{\ell(e)} r_e(\mathbf{x})^2$. Finally $P_\ell P_{\ell'} = P_\ell$ for $\ell' < \ell$, every $W_\ell$ contains $\mathcal{H} \oplus \mathcal{K}$, and

$$\dim W_\ell = 4N_\triangle - 1 + \frac{N_\triangle}{2^{\ell-1}} \qquad (1 \le \ell \le L). \tag{13.5}$$

[ESTABLISHED]

*Proof.* A passive change adds $g_i - g_i$ to $\mathfrak{d}_i$ and $g_i - g_j - (g_i - g_j)$ to $r_e$. Around an oriented triangle the vertex means telescope, so $\sum_{e \in q} r_e = -\mathfrak{h}_q$; the rows of $O_\triangle$ are orthonormal and orthogonal to $(1,1,1)/\sqrt{3}$, which gives eq. (13.4). In a linear relation among the rows of $B$, the link columns occur only in contrast and bridge rows, with the independent patterns $-O_\triangle$ on the triangle edges and $-1$ on each bridge, so every contrast and bridge coefficient vanishes; the model-mean columns then force the discrepancy coefficients to vanish. A normalized loop covector has link entries $1/\sqrt{3}$ on one triangle and nothing else, so it is orthogonal to the discrepancy rows, to the contrast rows because $O_\triangle(1,1,1)^\top = 0$, and to the bridges. Every row of $B$ and every loop covector annihilates $(t, t, \partial t)$, and $(6N_\triangle - 1) + N_\triangle + N = 10N_\triangle - 1$. The quadratic identity is the sum of $\mathrm{J}_\ell |B_\ell \mathbf{x}|^2$ and $\mathrm{J}_1 |P_{\mathcal{H}} \mathbf{x}|^2$ combined with eq. (13.4); it vanishes exactly when $B\mathbf{x} = 0$ and $P_{\mathcal{H}} \mathbf{x} = 0$, that is on $\mathcal{K}$. The spaces $W_\ell$ decrease, which gives $P_\ell P_{\ell'} = P_\ell$, and $\mathcal{H} \oplus \mathcal{K}$ lies in each of them. There are $N_\triangle / 2^{\ell''-1}$ bridges of level $\ell'' \ge 2$, so $\operatorname{rank}(B_1, \dots, B_\ell) = 3N_\triangle + 2N_\triangle + N_\triangle(1 - 2^{1-\ell})$, which gives eq. (13.5). □

The restoring estimates of the proof appendix rest on the bounds eq. (C.2) of $H$ by weighted energies $\sum_\ell \mathrm{J}_\ell |w_\ell|^2$ along orthonormal flag columns $\mathrm{Q}_\ell$ spanning $\mathcal{U}_{\ell-1} \cap \mathcal{U}_\ell^\perp$.

## 13.2 One normalized experiment with a nonlinear source record

**Definition 13.3 (Normalized formation experiment).** At each vertex the physical model sample is $m_i \sim \mathcal{N}(\bar{m}_i, \sigma_{\mathrm{ev}}^2)$ and the evaluator sends $m_i$ to $\mathcal{N}(m_i, 1 - \sigma_{\mathrm{ev}}^2)$ on the belief sample $k_i$. Among couplings with the prescribed marginals $\mathcal{N}(\bar{m}_i, \sigma_{\mathrm{ev}}^2)$ and $\mathcal{N}(\bar{k}_i, 1)$, the recognition coupling at a fixed configuration is

$$m_i = \bar{m}_i + \xi_i, \qquad k_i = \bar{k}_i + \xi_i + \xi_i', \qquad \xi_i \sim \mathcal{N}(0, \sigma_{\mathrm{ev}}^2), \quad \xi_i' \sim \mathcal{N}(0, 1 - \sigma_{\mathrm{ev}}^2), \tag{13.6}$$

with independent noises; write $\nu_i^*$ for its law and $R_{\mathrm{ev}} = \mathfrak{D}^\top \mathfrak{D} / (1 - \sigma_{\mathrm{ev}}^2)$, where $\mathfrak{D}\mathbf{x} = \bar{k} - \bar{m}$. Let $\vartheta \in [1, 21/10]$ be the learned scalar and $0 < \mathrm{T} \le \mathrm{J}_1(1 - \sigma_{\mathrm{ev}}^2)/2$. The configuration prior and its

Lebesgue normalizer are

$$\begin{aligned}\Omega_0(\vartheta) &= \vartheta H + P_{\mathcal{K}} - \mathrm{T}R_{\mathrm{ev}},\\ \mathbb{P}^0_\vartheta(d\mathbf{x}) &= Z_0(\vartheta)^{-1}\exp\left[-\frac{\mathbf{x}^\top\Omega_0(\vartheta)\mathbf{x}}{2\mathrm{T}}\right]d\mathbf{x},\\ Z_0(\vartheta) &= (2\pi\mathrm{T})^{(10N_\triangle-1)/2}\det\Omega_0(\vartheta)^{-1/2}.\end{aligned} \tag{13.7}$$

Fix $0<\kappa\le 1$ and a nonzero $z_\star\in\mathcal{K}$. The observation kernels are

$$o_{\mathcal{K}}\mid\mathbf{x}\sim\mathcal{N}_{\mathcal{K}}\big(P_{\mathcal{K}}\mathbf{x},(\mathrm{T}/\kappa)I_{\mathcal{K}}\big),\qquad o_{\mathrm{cal}}\mid\vartheta\sim\mathcal{N}(\vartheta,\mathrm{T}), \tag{13.8}$$

with the records $o_{\mathcal{K}}=z_\star$ and $o_{\mathrm{cal}}=2$; the first is a density on the $N$-dimensional carrier $\mathcal{K}$. Each receiver $i$ has the two candidate sources $\mathcal{N}_i^{\triangle}$ inside its own triangle, and the complete assignment is $S=(S_i)_i\in\mathcal{J}=\prod_i\mathcal{N}_i^{\triangle}$, with $|\mathcal{J}|=2^N$; every bridge keeps its interaction energy but lies outside the source palettes. A distinguished receiver $i_0$ has sources labeled $+$ and $-$, and the independent source priors are

$$\begin{aligned}\pi_{i_0,+}(\vartheta)&=\mathrm{sig}\big(4(\vartheta-3/2)\big),\quad \pi_{i_0,-}(\vartheta)=1-\pi_{i_0,+}(\vartheta),\\ \pi_{ij}(\vartheta)&=\tfrac12\quad(i\neq i_0),\qquad\qquad \pi_S(\vartheta)=\prod_i\pi_{iS_i}(\vartheta),\end{aligned} \tag{13.9}$$

where $\mathrm{sig}(u)=(1+e^{-u})^{-1}$. Set $B_0=10^{-4}$, $\delta=B_0/N$, and $\chi(r)=1-e^{-r^2/2}$. At each receiver a binary record $E_i$, conditionally independent given $(\mathbf{x},S)$, has

$$\begin{aligned}\Pr(E_i=1\mid\mathbf{x},S_i=j)&=q_0\exp\left[-\delta\chi\left(\sqrt{\mathrm{J}_1/\mathrm{T}}\,r_{ij}(\mathbf{x})\right)\right],\qquad q_0=\tfrac12,\\ \phi_S(\mathbf{x})&=\delta\sum_i\chi\left(\sqrt{\mathrm{J}_1/\mathrm{T}}\,r_{iS_i}(\mathbf{x})\right),\qquad 0\le\phi_S\le B_0,\end{aligned} \tag{13.10}$$

where $r_{ij}$ is the residual of the edge from source $j$ to receiver $i$, and the observed record is $E_i=1$ for every $i$. Since $\mathrm{J}_1r_{ij}^2/(2\mathrm{T})=(\mathrm{J}_1/\mathrm{T})D_{\mathrm{KL}}(\mathcal{N}(\bar k_i,1)\|\mathcal{N}(\bar k_j+\omega_{ij},1))$, the record is a bounded observation of the standardized edge mismatch. The complete joint is

$$\begin{aligned}\mathbb{P}_\vartheta=\mathbb{P}^0_\vartheta(d\mathbf{x})L_{\mathcal{K}}(do_{\mathcal{K}}\mid\mathbf{x})L_{\mathrm{cal}}(do_{\mathrm{cal}}\mid\vartheta)\\ \times\prod_i\left[\pi_i(dS_i\mid\vartheta)L_i(dE_i\mid\mathbf{x},S_i)\mathcal{N}(dm_i;\bar m_i,\sigma_{\mathrm{ev}}^2)\mathcal{N}(dk_i;m_i,1-\sigma_{\mathrm{ev}}^2)\right].\end{aligned} \tag{13.11}$$

**Proposition 13.4 (Normalization and the evaluator optimum).** In Definition 13.3, for every admissible coupling $\nu_i$ of the two prescribed marginals,

$$D_{\mathrm{KL}}\big(\nu_i\|\mathcal{N}(\bar m_i,\sigma_{\mathrm{ev}}^2)\,\mathrm{ev}_i\big)=D_{\mathrm{KL}}(\nu_i\|\nu_i^*)+\frac{\mathfrak{d}_i^2}{2(1-\sigma_{\mathrm{ev}}^2)}, \tag{13.12}$$

so $\nu_i^*$ is the optimal coupling and the total restriction cost is $\mathbf{x}^\top R_{\mathrm{ev}}\mathbf{x}/2$. Moreover $\mathrm{T}R_{\mathrm{ev}}\preceq H/2$, hence $\Omega_0(\vartheta)\succeq(\vartheta/2)H+P_{\mathcal{K}}\succ 0$ and the prior is a normalized correlated Gaussian, and the joint

eq. (13.11) has total mass one, contains no kernel that takes a recognition law as input, and gives the selected observation a strictly positive density. [ESTABLISHED]

The term $-\mathrm{T}R_{\mathrm{ev}}$ of the prior cancels the restriction cost $\frac{\mathrm{T}}{2}\mathbf{x}^{\top}R_{\mathrm{ev}}\mathbf{x}$ that eq. (13.12) adds, so the restricted optimizer below has a temperature-independent quadratic energy.

## 13.3 Physical posterior, restricted optimizer, and recognition

**Definition 13.5 (Three laws of the formation experiment).** Write $\mathsf{o} = (z_\star, 2, \mathbf{1})$ for the observed records. The recognition family consists of the laws with a smooth positive density $q(\mathbf{x}, S)$ relative to Lebesgue measure times counting measure and conditional evaluator coupling $\prod_i \nu_i^*$; assignments and configurations may be correlated, and the conditional assignment law need not be the source posterior. With

$$\begin{aligned} \widehat{K}_\vartheta &= \vartheta H + (1+\kappa)P_{\mathcal{K}}, & \mathrm{E}_\vartheta(\mathbf{x}) &= \tfrac{1}{2}\mathbf{x}^{\top}\widehat{K}_\vartheta\mathbf{x} - \kappa z_\star^{\top}\mathbf{x}, \\ \mathrm{E}_\vartheta^{\mathrm{phys}}(\mathbf{x}) &= \mathrm{E}_\vartheta(\mathbf{x}) - \tfrac{\mathrm{T}}{2}\mathbf{x}^{\top}R_{\mathrm{ev}}\mathbf{x}, & U_{\vartheta,S}(\mathbf{x}) &= \mathrm{E}_\vartheta(\mathbf{x}) + \mathrm{T}\phi_S(\mathbf{x}), \\ Z_{\mathrm{opt}}(\vartheta) &= \sum_S \pi_S(\vartheta)\int e^{-\mathrm{E}_\vartheta(\mathbf{x})/\mathrm{T}-\phi_S(\mathbf{x})}d\mathbf{x}, & Z_{\mathrm{phys}}(\vartheta) &= \sum_S \pi_S(\vartheta)\int e^{-\mathrm{E}_\vartheta^{\mathrm{phys}}(\mathbf{x})/\mathrm{T}-\phi_S(\mathbf{x})}d\mathbf{x}, \end{aligned} \tag{13.13}$$

the restricted optimizer is $q_\vartheta^\star(\mathbf{x}, S) = \pi_S(\vartheta)e^{-U_{\vartheta,S}(\mathbf{x})/\mathrm{T}}/Z_{\mathrm{opt}}(\vartheta)$, and the physical posterior $\mathbf{\Pi}_{\vartheta,\mathsf{o}}$ on $(\mathbf{x}, S)$ is the same expression with $\mathrm{E}_\vartheta^{\mathrm{phys}}$ and $Z_{\mathrm{phys}}$. The scaled restricted free energy $\mathcal{F}_{\mathrm{T}}(q,\vartheta)$ is $\mathrm{T}$ times the negative ELBO of the evaluator restriction, evaluated at the records $\mathsf{o}$.

**Proposition 13.6 (Laws, restricted free energy, and the physical gap).** The integrals eq. (13.13) are finite and positive, both laws have full support, and their common conditional assignment row at fixed $\mathbf{x}$ is the product of the receiver posterior rows,

$$\beta_S(\mathbf{x},\vartheta) = \frac{\pi_S(\vartheta)e^{-\phi_S(\mathbf{x})}}{\sum_{S'}\pi_{S'}(\vartheta)e^{-\phi_{S'}(\mathbf{x})}}. \tag{13.14}$$

With the scalars $\mathrm{B}_{\mathrm{T}}$, $\mathrm{C}_{\mathrm{T}}$, and $p_\vartheta(\mathsf{o})$ of eq. (C.1),

$$\begin{aligned} \mathcal{F}_{\mathrm{T}}(q,\vartheta) &= \mathrm{T}\sum_S \int q(\mathbf{x},S)\log q(\mathbf{x},S)d\mathbf{x} + \mathbb{E}_q\big[\mathrm{E}_\vartheta(\mathbf{x}) + \mathrm{T}\phi_S(\mathbf{x}) - \mathrm{T}\log\pi_S(\vartheta)\big] + \mathrm{B}_{\mathrm{T}}(\vartheta) \\ &= \mathrm{T}D_{\mathrm{KL}}(q\|q_\vartheta^\star) + \mathrm{C}_{\mathrm{T}}(\vartheta), \end{aligned} \tag{13.15}$$

and the excess over the unconstrained physical optimum is $\mathrm{C}_{\mathrm{T}} + \mathrm{T}\log p_\vartheta(\mathsf{o}) = \mathrm{T}\log(Z_{\mathrm{phys}}/Z_{\mathrm{opt}}) \geq 0$. Whenever $N\mathrm{T} \leq (1-\sigma_{\mathrm{ev}}^2)\mathrm{J}_1$,

$$0 \leq \mathrm{T}\log\frac{Z_{\mathrm{phys}}(\vartheta)}{Z_{\mathrm{opt}}(\vartheta)} \leq \frac{2e^{B_0}N\mathrm{T}^2}{(1-\sigma_{\mathrm{ev}}^2)\mathrm{J}_1}, \qquad D_{\mathrm{KL}}(q_\vartheta^\star\|\mathbf{\Pi}_{\vartheta,\mathsf{o}}) \leq \frac{2e^{B_0}N\mathrm{T}}{(1-\sigma_{\mathrm{ev}}^2)\mathrm{J}_1}. \tag{13.16}$$

[ESTABLISHED]

## 13.4 Exact marked parents and their inherited inference

**Definition 13.7 (Marked parents and inherited edge statistics).** The parent carrier at level $\ell$ is $W_\ell \times \mathcal{J}$. For $\ell' < \ell$ the deterministic Markov kernel

$$C_{\ell\leftarrow\ell'}((y,S),\cdot) = \delta_{(P_{\ell'}y,S)}, \qquad C_\ell = C_{\ell\leftarrow 0}, \tag{13.17}$$

extended by the identity on observations, pushes the generative joint, the physical posterior, the restricted optimizer, and the recognition law, giving $q_\vartheta^{\star,\ell} = (C_\ell)_\# q_\vartheta^\star$, $\mathbf{\Pi}^\ell_{\vartheta,\mathrm{o}}$, and $(C_\ell)_\# q$, and erases no assignment. The evaluator samples $(m_i,k_i)_i$, the records $(E_i)_i$, and the discarded normal coordinates leave the parent carrier through the disintegration of eq. (13.11) given $(P_\ell \mathbf{x}, S)$, so their correlations are retained and no mark is replaced by its value at $P_\ell\mathbf{x}$. For an original edge $e$ and each $\mu_\vartheta \in \{\mathbf{\Pi}_{\vartheta,\mathrm{o}}, q_\vartheta^\star\}$, the inherited edge statistic is

$$\mathcal{D}_e^{\ell,\mu}(y,S) = \mathbb{E}_{\mu_\vartheta}\left[\tfrac{1}{2} r_e(\mathbf{X})^2 \mid P_\ell \mathbf{X} = y, S\right], \tag{13.18}$$

the discrepancy observables for $\mathfrak{d}_i$ are defined in the same way, and at $\ell = 0$ both reduce to the configuration residuals. For a density $q^\ell$ on the marked parent carrier, the parent free energy is

$$\mathcal{F}^\ell_{\mathrm{T}}(q^\ell,\vartheta) = \mathrm{T} D_{\mathrm{KL}}(q^\ell \| q_\vartheta^{\star,\ell}) + \mathrm{C}_{\mathrm{T}}(\vartheta). \tag{13.19}$$

**Proposition 13.8 (Exact parent potential and discarded conditional cost).** The kernels eq. (13.17) compose, $C_{\ell''\leftarrow\ell} \circ C_{\ell\leftarrow\ell'} = C_{\ell''\leftarrow\ell'}$, and positive smooth versions of every pushed law and conditional exist. The root factor of $q_\vartheta^{\star,\ell}$ is independent of the assignment and of $\vartheta$. On $M_\ell$ let $H_\ell$ be the Schur complement of $H|_{\mathcal{K}^\perp}$ and $\gamma^{\mathrm{hid}}_{\ell,\vartheta}(dz \mid y)$ the cumulative hidden Gaussian conditional of the quadratic optimizer, whose mean is independent of $\vartheta$ and whose covariance is $\mathrm{T}/\vartheta$ times a fixed inverse precision. Integration at fixed $S$ gives the sheet potential

$$\begin{aligned} V_{\ell,S}(y,\vartheta) &= \tfrac{1}{2}\vartheta y^\top H_\ell y + \mathrm{C}^{\mathrm{G}}_\ell(\vartheta) + \mathrm{T}\psi_{\ell,S}(y,\vartheta), \\ \psi_{\ell,S}(y,\vartheta) &= -\log \mathbb{E}_{\gamma^{\mathrm{hid}}_{\ell,\vartheta}(\cdot|y)} e^{-\phi_S}, \qquad 0 \le \psi_{\ell,S} \le B_0, \\ \frac{d}{d\vartheta}\mathrm{C}^{\mathrm{G}}_\ell(\vartheta) &= \frac{\mathrm{T}\mathrm{d}_\ell}{2\vartheta}, \end{aligned} \tag{13.20}$$

where $\mathrm{d}_\ell = \dim W_\ell^\perp$ is the cumulative eliminated dimension, $\mathrm{C}^{\mathrm{G}}_0 = 0$, $\psi_{0,S} = \phi_S$, and the unchanged root potential is appended on $W_\ell$. At one common parameter and one common time,

$$\begin{aligned} \Delta_{\ell':\ell}(q^{\ell'},\vartheta) &= \mathbb{E}_{(C_{\ell\leftarrow\ell'})_\# q^{\ell'}} D_{\mathrm{KL}}\left(q^{\ell'}(dz \mid y,S) \middle\| q_\vartheta^{\star,\ell'}(dz \mid y,S)\right), \\ \mathcal{F}^{\ell'}_{\mathrm{T}}(q^{\ell'},\vartheta) - \mathcal{F}^{\ell}_{\mathrm{T}}\big((C_{\ell\leftarrow\ell'})_\# q^{\ell'},\vartheta\big) &= \mathrm{T}\Delta_{\ell':\ell}(q^{\ell'},\vartheta), \end{aligned} \tag{13.21}$$

where $z$ consists only of the discarded normal coordinates, loops and $S$ staying in the retained state. [ESTABLISHED]

*Proof.* Composition follows from $P_{\ell''}P_{\ell} = P_{\ell''}$, and finite Gaussian domination supplies the positive smooth versions. The potential is the Gaussian integral computed in Lemma C.1. The cost identity is Theorem 9.6 for the channel $C_{\ell\leftarrow\ell'}$, since the scalar $\mathrm{C_T}$ is common to both sides and the channel leaves the observations unchanged. □

**Definition 13.9 (Declared inference kinetics).** Every microscopic coordinate has unit continuous mobility, and a rate-one Poisson proposal clock redraws the whole assignment from $\beta(\mathbf{x}, \vartheta)$ without moving any continuous coordinate, a proposal being allowed to return the current assignment. The microscopic generator on global Euclidean space times the finite set $\mathcal{J}$ is

$$\mathscr{L}^0_\vartheta \mathsf{f}(\mathbf{x}, S) = \mathrm{T}\Delta_{\mathbf{x}}\mathsf{f} - \nabla_{\mathbf{x}} U_{\vartheta,S} \cdot \nabla_{\mathbf{x}}\mathsf{f} + \sum_{S'} \beta_{S'}(\mathbf{x}, \vartheta)[\mathsf{f}(\mathbf{x}, S') - \mathsf{f}(\mathbf{x}, S)], \tag{13.22}$$

with diffusion covariance $2\mathrm{T}I\,dt$. The inherited proposal row and parent generator are

$$\begin{aligned} \beta^\ell_{SS'}(y, \vartheta) &= \mathbb{E}_{q^\star_\vartheta}[\beta_{S'}(\mathbf{X}, \vartheta) \mid P_\ell \mathbf{X} = y, S], \\ \mathscr{L}^\ell_\vartheta \mathsf{f}(y, S) &= \mathrm{T}\Delta_y \mathsf{f} - \nabla_y V_{\ell,S} \cdot \nabla_y \mathsf{f} + \sum_{S'} \beta^\ell_{SS'}(y, \vartheta)[\mathsf{f}(y, S') - \mathsf{f}(y, S)], \end{aligned} \tag{13.23}$$

with the unchanged root drift included when $y$ denotes all of $W_\ell$. Self-proposals are kept, so every row sums to one and every exit rate is at most one; these rows declare the parent kinetics rather than project the fine process in time.

**Proposition 13.10 (Reversible inherited rows).** The rows $\beta^\ell$ satisfy detailed balance with $q^{\star,\ell}_\vartheta$, and for $\ell' < \ell$

$$\beta^{\ell'}_{SS'}(y, \vartheta) \geq e^{-B_0}\pi_{S'}(\vartheta) > 0, \qquad \mathbb{E}_{q^{\star,\ell'}_\vartheta}[\beta^{\ell'}_{SS'}(Y, \vartheta) \mid P_\ell Y = z, S] = \beta^\ell_{SS'}(z, \vartheta). \tag{13.24}$$

[ESTABLISHED]

*Proof.* The product $q^\star_\vartheta(\mathbf{x}, S)\beta_{S'}(\mathbf{x}, \vartheta)$ equals the continuous marginal of $q^\star_\vartheta$ times $\beta_S\beta_{S'}$ and is symmetric in $S, S'$; conditional integration over a retained fiber gives detailed balance. The lower bound follows from eq. (13.14) and $0 \leq \phi_S \leq B_0$, and the second identity is the tower property of conditional expectation. □

## 13.5 Complete model scores and dissipation

**Definition 13.11 (Own-law parameter dynamics).** Let $\mathsf{M}(\vartheta)$ be the Fisher information of the complete normalized experiment eq. (13.11) in $\vartheta$, before conditioning on the records, and $\mathsf{m}(\vartheta) = \mathrm{T}\mathsf{M}(\vartheta)$. With the model metric $\mathsf{m}(\vartheta)/\epsilon$, $\epsilon > 0$, each level $\ell$ evolves its own recognition density and its own parameter by

$$\partial_t q^\ell_t = (\mathscr{L}^\ell_{\vartheta_\ell(t)})^* q^\ell_t, \qquad \dot\vartheta_\ell(t) = -\epsilon \mathsf{m}(\vartheta_\ell(t))^{-1} g_\ell(q^\ell_t, \vartheta_\ell(t)), \qquad g_\ell = \partial_\vartheta \mathcal{F}^\ell_{\mathrm{T}}\Big|_{q^\ell}. \tag{13.25}$$

The parameter path is deterministic, since the feedback is an expectation under the law of its level. Parents inherit this metric, with no new rate multiplier at activation.

**Proposition 13.12 (Complete metric, covector, and dissipation).** The complete Fisher metric is

$$\mathsf{M}(\vartheta) = \frac{1}{\mathrm{T}} + \frac{1}{2}\operatorname{Tr}\big[(\Omega_0(\vartheta)^{-1}H)^2\big] + 16\pi_{i_0,+}(\vartheta)\big[1-\pi_{i_0,+}(\vartheta)\big], \tag{13.26}$$

and the full parameter covector of level $\ell$ is

$$\begin{aligned} g_\ell(q^\ell,\vartheta) = \vartheta - 2 + \tfrac{1}{2}\mathbb{E}_{q^\ell} y^\top H_\ell y + \frac{\mathrm{T}\mathrm{d}_\ell}{2\vartheta} + \mathrm{T}\mathbb{E}_{q^\ell}\partial_\vartheta \psi_{\ell,S}(y,\vartheta) \\ - \frac{\mathrm{T}}{2}\operatorname{Tr}\big(\Omega_0(\vartheta)^{-1}H\big) - \mathrm{T}\mathbb{E}_{q^\ell}\partial_\vartheta \log \pi_S(\vartheta), \end{aligned} \tag{13.27}$$

a partial derivative at fixed recognition law, which at $\ell = 0$, where $\mathrm{d}_0 = 0$ and $\partial_\vartheta \phi_S = 0$, is the microscopic score. Along eq. (13.25), with $\mathfrak{r}_S = q^\ell_S / q^{\star,\ell}_{\vartheta_\ell}(\cdot, S)$,

$$\begin{aligned} \frac{d}{dt}\mathcal{F}^\ell_{\mathrm{T}} = & -\mathrm{T}^2 \sum_S \int q^\ell_S \,|\nabla \log \mathfrak{r}_S|^2 \, dy - \frac{\mathrm{T}}{2}\int \sum_{S,S'} q^{\star,\ell}_{\vartheta_\ell}(y,S)\beta^\ell_{SS'}(y,\vartheta_\ell)(\mathfrak{r}_S - \mathfrak{r}_{S'})(\log \mathfrak{r}_S - \log \mathfrak{r}_{S'})dy \\ & - \epsilon \mathsf{m}(\vartheta_\ell)^{-1} g_\ell(q^\ell,\vartheta_\ell)^2 \le 0. \end{aligned} \tag{13.28}$$

At the microscopic level the jump term equals $-\mathrm{T}\int q_{\mathbf{X}}[D_{\mathrm{KL}}(q(\cdot \mid \mathbf{x})\|\beta) + D_{\mathrm{KL}}(\beta\|q(\cdot \mid \mathbf{x}))]d\mathbf{x}$. [ESTABLISHED]

The proofs are in appendix C.1; the identities hold for the solutions admitted by Lemma C.3, whose cutoff passage is Lemma C.6.

## 13.6 Preparation and successive activation

**Definition 13.13 (Regime, preparation, and activation schedule).** For $0 < \epsilon < 1$ the regime is

$$\mathrm{Q}_L = 0, \qquad \mathrm{Q}_\ell = 2\mathrm{Q}_{\ell+1} + 20, \qquad \mathrm{J}_\ell = \epsilon^{-\mathrm{Q}_\ell}, \qquad \mathrm{A}^2 = \epsilon^{3\mathrm{Q}_1+20}, \qquad \mathrm{T} = \mathrm{A}^2\epsilon^6, \qquad t_* = \frac{2}{\epsilon}, \tag{13.29}$$

with unit mobility and proposal rate and the record coefficient $\delta$ of eq. (13.10); the weakest energy is exactly $\mathrm{J}_L = 1$. Put $\varsigma_\epsilon = 1 + \log(1/\epsilon)$ and

$$r_\epsilon = \mathrm{A}^2\mathrm{J}_1 + \mathrm{T}\varsigma^2_\epsilon = \epsilon^{2\mathrm{Q}_1+20} + \epsilon^{3\mathrm{Q}_1+26}\varsigma^2_\epsilon. \tag{13.30}$$

Let $\mathbf{v}_\star \in \mathcal{U}$ be the unit vector of Lemma C.2, with margin $\mathrm{m}_L > 0$, and $\mathbf{v}_0$ any vector of the neighborhood of $\mathbf{v}_\star$ on which the margin is at least $\mathrm{m}_L/2$. The initial parameter is $\vartheta(0) = 1$. The root is independent of the non-root coordinates and the assignments, with law $\mathcal{N}_{\mathcal{K}}(\kappa z_\star/(1+\kappa), \mathrm{T}(1+\kappa)^{-1} I_{\mathcal{K}})$. On every assignment sheet the reference non-root law has common mean $\mathrm{A}\mathbf{v}_0$, zero loop mean, and covariance $\mathrm{T}(H|_{\mathcal{K}^\perp})^{-1}$; the positive assignment masses are near the product prior at $\vartheta = 1$, and the distinguished source mass $\lambda_{i_0,+}(0) = \Pr\{S_{i_0}(0) = +\}$ satisfies

$|\lambda_{i_0,+}(0) - \mathrm{sig}(-2)| < 10^{-2}$. The preparation neighborhood consists of the smooth within-sheet transports of this reference with bounded log-density ratio, retaining the root factor, whose coupling satisfies

$$\left|\widehat{K}_1^{1/2}(\mathbf{X}_0 - \mathbf{X}_0^{\mathrm{G}})\right| < \epsilon\sqrt{\mathrm{T}} \quad \text{almost surely;} \tag{13.31}$$

small compactly supported diffeomorphisms in standardized non-root coordinates supply such transports, the neighborhood may depend on $\epsilon$ and on $L$, and later laws need not stay Gaussian or factor across sheets. Given an entrance constant $k_L$, the windows and activation times are

$$t_\ell = \frac{k_L \log(1/\epsilon)}{\mathrm{J}_\ell}, \qquad I_\ell = [t_\ell, 2t_\ell], \qquad t_0^{\mathrm{act}} = 0, \quad t_\ell^{\mathrm{act}} = 2t_\ell. \tag{13.32}$$

At $t_\ell^{\mathrm{act}}$ the complete level-$(\ell-1)$ law is pushed through $C_{\ell\leftarrow\ell-1}$, its parameter is copied, and the new law then evolves by eq. (13.25); an independently continued source process keeps its own law and parameter.

**Remark 13.14 (The regime in closed form).** Solving the recursion gives $\mathrm{Q}_\ell = 20(2^{L-\ell} - 1)$, so

$$\mathrm{Q}_1 = 20(2^{L-1} - 1), \qquad \frac{\mathrm{J}_1}{\mathrm{J}_L} = \epsilon^{-\mathrm{Q}_1}, \qquad \mathrm{A}^2 = \epsilon^{3\mathrm{Q}_1+20}, \qquad \mathrm{T} = \epsilon^{3\mathrm{Q}_1+26} = \epsilon^{60\cdot 2^{L-1}-34}. \tag{13.33}$$

The regime is a low-temperature limit with doubly exponential energy separation in the depth: $\mathrm{T} = \epsilon^{86}, \epsilon^{206}, \epsilon^{446}$ for $L = 2, 3, 4$, every ratio $\mathrm{J}_\ell/\mathrm{T}$ diverges, $\mathrm{J}_{\ell+1}/\mathrm{J}_\ell = \epsilon^{\mathrm{Q}_{\ell+1}+20}$ and $\mathrm{J}_{\ell+1}^2/\mathrm{J}_\ell = \epsilon^{20}$ for $\ell < L$, and the relative thermal scale is $\sqrt{\mathrm{T}}/\mathrm{A} = \epsilon^3$. There is no common acceleration of every energy, no vanishing source rate, and no decrease of $\delta$. [ESTABLISHED]

The recursion comes from the early estimate of Lemma C.7, which bounds the retained Fisher information by the full budget eq. (C.25) and hence the discarded loss on $I_\ell$ by the term $(\mathrm{A}^2/\mathrm{T})\mathrm{J}_{\ell+1}^2/\mathrm{J}_\ell$ of eq. (C.41). The linear-Gaussian core shows that this term is not sharp.

**Proposition 13.15 (Gaussian core of the discarded conditional loss).** Take $\delta = 0$, $\vartheta \equiv 1$, and an active non-root normal law $\mathcal{N}(\bar{\mathbf{x}}(t), \mathrm{T}H_{\ell'}^{-1})$ at level $\ell' < \ell$, whose covariance is stationary for the Ornstein–Uhlenbeck flow, so $\bar{\mathbf{x}}(t) = e^{-H_{\ell'}t}\bar{\mathbf{x}}(0)$. Split the coordinates into the discarded $z$ and retained $y$, with blocks $H_{zz}$, $H_{zy}$, $H_{yy}$ of $H_{\ell'}$ and Schur complement $H_\ell = H_{yy} - H_{zy}^\top H_{zz}^{-1} H_{zy}$.

(a) *Loss formula and upper bound.* Then

$$\begin{aligned} \Delta_{\ell':\ell}(t) &= \frac{1}{2\mathrm{T}}\mathfrak{v}^\top H_{zz}^{-1}\mathfrak{v}, & \mathfrak{v} &= (H_{\ell'}\bar{\mathbf{x}})_z, \\ \dot{\mathfrak{v}} &= -H_{zz}\mathfrak{v} - H_{zy}(H_{\ell'}\bar{\mathbf{x}})_y, & (H_{\ell'}\bar{\mathbf{x}})_y &= H_\ell \bar{\mathbf{x}}_y + H_{zy}^\top H_{zz}^{-1}\mathfrak{v}, \end{aligned} \tag{13.34}$$

and, for $|\bar{\mathbf{x}}| \le \mathrm{C}_L\mathrm{A}$, $\Delta_{\ell':\ell} \le \mathrm{C}_L(\mathrm{A}^2/\mathrm{T})[\mathrm{J}_{\ell'+1}e^{-\mathrm{c}_L\mathrm{J}_\ell(t-t_{\ell'}^{\mathrm{act}})} + \mathrm{J}_{\ell+1}^4/\mathrm{J}_\ell^3]$.

(b) *Two-sided asymptotic across an early interface.* Let the strengths $\mathrm{J}_1 \ge \cdots \ge \mathrm{J}_L = 1$ depend on $\epsilon$ with ratios $\varrho_s = \mathrm{J}_{s+1}/\mathrm{J}_s \to 0$ for $s < L$, put $\varrho_0 = \varrho_L = 0$, and let $1 \le \ell \le L-1$ and $\ell' = \ell - 1$. Write $\mathsf{B}_{sr} = B_s\mathsf{Q}_r$ for the blocks of the matrix $\mathsf{B}$ of eq. (C.2), which vanish for $r > s$ and are square and invertible for $r = s$ (Lemma C.1), let $\mathrm{j}_\ell = \lambda_{\min}(\mathsf{B}_{\ell\ell}^\top\mathsf{B}_{\ell\ell})$, let $\mathfrak{Y}_\ell = \mathsf{B}_{\ell+1,\ell}^\top\mathsf{B}_{\ell+1,\ell+1}\mathsf{B}_{\ell+1,\ell+1}^\top$ be the

forcing map of the interface, and for $\mathbf{v} \in \mathbb{R}^{10N_\triangle - 1}$ let

$$\mathfrak{y}_\ell(\mathbf{v}) = \tfrac{1}{2}\big|(\mathrm{B}_{\ell\ell}^{\top}\mathrm{B}_{\ell\ell})^{-3/2}\mathfrak{Y}_\ell B_{\ell+1}P_\ell\mathbf{v}\big|^2, \qquad \mathfrak{x}_\ell(t) = \varrho_\ell^{-2}\exp\big[-\tfrac{1}{2}\mathrm{j}_\ell\mathrm{J}_\ell(t - t^{\mathrm{act}}_{\ell-1})\big]. \tag{13.35}$$

Then $\mathrm{B}_{\ell+1,\ell}$ has full row rank for every choice of the bridge endpoints, so $\mathfrak{Y}_\ell$ is injective. There are $\mathrm{c}_0, \mathrm{C} > 0$ depending only on $L$ and the carrier such that, for the fine process started from the reference non-root law of Definition 13.13 with mean $\mathrm{A}\mathbf{v}_0$ and activated at the times eq. (13.32), whenever $\varrho_{\ell-1} + \varrho_\ell + \varrho_{\ell+1} \le \mathrm{c}_0$, the active level $\ell - 1$ satisfies, for $t^{\mathrm{act}}_{\ell-1} \le t \le t^{\mathrm{act}}_\ell$,

$$\begin{aligned} &\Big|\frac{\mathrm{T}\mathrm{J}_\ell^3}{\mathrm{A}^2\mathrm{J}_{\ell+1}^4}\Delta_{\ell-1:\ell}(t) - \mathfrak{y}_\ell\big(\bar{\mathbf{x}}(t)/\mathrm{A}\big)\Big| \le \mathrm{C}\big[\varrho_{\ell-1} + \varrho_\ell + \varrho_{\ell+1} + \mathfrak{x}_\ell(t) + \mathfrak{x}_\ell(t)^2\big], \\ &\qquad |B_{\ell+1}P_\ell\bar{\mathbf{x}}(t) - \mathrm{A}B_{\ell+1}P_\ell\mathbf{v}_0| \le \mathrm{C}\mathrm{J}_{\ell+1}\mathrm{A}t. \end{aligned} \tag{13.36}$$

If moreover $k_L\log(1/\epsilon)\varrho_\ell \to 0$ and $\frac{1}{2}\mathrm{j}_\ell k_L\log(1/\epsilon)(1 - 2\varrho_{\ell-1}) - 2\log(1/\varrho_\ell) \to \infty$, then, uniformly on $I_\ell$,

$$\Delta_{\ell-1:\ell}(t) = \frac{\mathrm{A}^2}{\mathrm{T}}\frac{\mathrm{J}_{\ell+1}^4}{\mathrm{J}_\ell^3}\big[\mathfrak{y}_\ell(\mathbf{v}_0) + o(1)\big], \qquad \mathfrak{y}_\ell(\mathbf{v}_0) \ge \frac{\sigma_{\min}(\mathfrak{Y}_\ell)^2 N_\triangle \mathrm{m}_L^2}{2^{\ell+3}\|\mathrm{B}_{\ell\ell}\|^6} > 0, \tag{13.37}$$

where $\sigma_{\min}$ denotes the least singular value.

(c) *Separation criterion.* With power laws $\mathrm{J}_s = \epsilon^{-\mathrm{Q}_s}$, $\mathrm{Q}_1 > \cdots > \mathrm{Q}_L = 0$, with $\mathrm{A}^2/\mathrm{T} = \epsilon^{-\mathrm{Q}_\mathrm{T}}$ for some $\mathrm{Q}_\mathrm{T} > 0$, which is $\mathrm{Q}_\mathrm{T} = 6$ in eq. (13.29), and with an entrance constant $k_L > 4(\mathrm{Q}_\ell - \mathrm{Q}_{\ell+1})/\mathrm{j}_\ell$ for every $\ell < L$, both $\inf_{I_\ell}\Delta_{\ell-1:\ell}$ and $\sup_{I_\ell}\Delta_{\ell-1:\ell}$ are of exact order $\epsilon^{3\mathrm{Q}_\ell - 4\mathrm{Q}_{\ell+1} - \mathrm{Q}_\mathrm{T}}$. The core loss is $O(\epsilon^{\mathrm{Q}_\triangle})$ on every early window, $\mathrm{Q}_\triangle \ge 0$, exactly when $3\mathrm{Q}_\ell - 4\mathrm{Q}_{\ell+1} \ge \mathrm{Q}_\mathrm{T} + \mathrm{Q}_\triangle$ for every $\ell \le L - 1$, which forces $\mathrm{Q}_1 \ge (\mathrm{Q}_\mathrm{T} + \mathrm{Q}_\triangle)[(4/3)^{L-1} - 1]$, with equality along the recursion $3\mathrm{Q}_\ell = 4\mathrm{Q}_{\ell+1} + \mathrm{Q}_\mathrm{T} + \mathrm{Q}_\triangle$; it vanishes on every early window exactly when $3\mathrm{Q}_\ell - 4\mathrm{Q}_{\ell+1} > \mathrm{Q}_\mathrm{T}$ for every $\ell \le L - 1$. [ESTABLISHED]

The proof of (b) and (c) is in appendix C.6.

*Proof of (a).* Both conditional laws of $z$ given $y$ have precision $H_{zz}/\mathrm{T}$, and their means differ by $H_{zz}^{-1}(H_{\ell'}\bar{\mathbf{x}})_z$ independently of $y$, which gives the first formula. Differentiating $\mathfrak{v} = (H_{\ell'}\bar{\mathbf{x}})_z$ along $\dot{\bar{\mathbf{x}}} = -H_{\ell'}\bar{\mathbf{x}}$ and eliminating $\bar{\mathbf{x}}_z$ give the other two. Put $\mathfrak{u} = H_{zz}^{-1/2}\mathfrak{v}$. Its equation is

$$\dot{\mathfrak{u}} = -\big(H_{zz} + H_{zz}^{-1/2}H_{zy}H_{zy}^{\top}H_{zz}^{-1/2}\big)\mathfrak{u} - H_{zz}^{-1/2}H_{zy}H_\ell\bar{\mathbf{x}}_y.$$

The matrix multiplying $\mathfrak{u}$ is symmetric and bounded below by $\mathrm{c}_L\mathrm{J}_\ell I$. The retained loop block is orthogonal to $H_{zy}$ and contributes no forcing; on the future normal coordinates, eq. (C.71) bounds the forcing by $\mathrm{C}_L\mathrm{A}\mathrm{J}_{\ell+1}^2/\sqrt{\mathrm{J}_\ell}$. At $t_0 = t^{\mathrm{act}}_{\ell'}$, the weighted flag gives $|\mathfrak{u}(t_0)| \le \mathrm{C}_L\mathrm{A}\sqrt{\mathrm{J}_{\ell'+1}}$. Variation of constants therefore gives

$$|\mathfrak{u}(t)| \le \mathrm{C}_L\mathrm{A}\left[\sqrt{\mathrm{J}_{\ell'+1}}e^{-\mathrm{c}_L\mathrm{J}_\ell(t-t_0)} + \frac{\mathrm{J}_{\ell+1}^2}{\mathrm{J}_\ell^{3/2}}\right].$$

Squaring and using $\Delta_{\ell':\ell} = |\mathfrak{u}|^2/(2\mathrm{T})$ proves the bound, with the decay constant adjusted. The forcing vanishes for $\ell = L$. □

**Open Problem 13.16 (Sharpness of the doubly exponential regime).** Decide whether the recursion $\mathrm{Q}_\ell = 2\mathrm{Q}_{\ell+1}+20$ is required for Theorem 13.17, and whether its constants and preparation neighborhoods can be made uniform in the depth. What holds is the theorem at every finite $L$ with constants depending on $L$, and, by Proposition 13.15, that in the Gaussian core the discarded loss on every early window is of exact order $(\mathrm{A}^2/\mathrm{T})\mathrm{J}_{\ell+1}^4/\mathrm{J}_\ell^3$, smaller than the term of the general estimate by $(\mathrm{J}_{\ell+1}/\mathrm{J}_\ell)^2$, so that the early interfaces need exactly $3\mathrm{Q}_\ell \geq 4\mathrm{Q}_{\ell+1}+20$ for the core loss at the scale $\epsilon^{14}$, a separation still doubly exponential in $L$ but of inner base $4/3$, with $\mathrm{Q}_\ell \geq 14$ for the term $\mathrm{C}_{\mathcal{J}}/\mathrm{J}_\ell$ of the jump remainder eq. (C.39). The doubling enters the present proof through the full Fisher budget eq. (C.25) in the early estimate of Lemma C.7 and through the history estimate of Lemma C.11, which uses it at the top interface in eq. (C.66). The depth enters through $\mathrm{Q}_1$ in the regime and every entrance condition, through $\log|\mathcal{J}| = N\log 2$ in the jump remainder, and through the Frobenius constants of eq. (C.2), whose squares grow linearly in $N_\triangle$. The problem is posed in the family of the theorem under (H1), (H3), and (H4), with power-law strengths $\mathrm{J}_\ell = \epsilon^{-\mathrm{Q}_\ell}$, $\mathrm{Q}_L = 0$, satisfying $3\mathrm{Q}_\ell \geq 4\mathrm{Q}_{\ell+1}+20$ and $\mathrm{Q}_\ell \geq 14$ for $\ell < L$, with the amplitude, temperature, horizon, and record coefficient of eq. (13.29) formed from $\mathrm{Q}_1$, and with an entrance constant above the thresholds of the theorem and of Proposition 13.15. It is settled by conclusions (a) to (e) for every schedule of this class, with constants and preparation neighborhoods independent of $L$, through an early estimate of the retained Fisher information by the future-layer energy, a history-environment bound without the top-interface recursion, and depth-uniform flag and jump constants; or by an admitted preparation and a schedule of the class under which one of the conclusions fails, or a lower bound showing that one of the conclusions cannot hold with constants independent of $L$.

## 13.7 The finite-depth formation theorem

Formation in this family is hierarchical relaxation of the prepared mean. The non-root configuration starts near $\mathrm{A}\mathbf{v}_0$, where every residual has order $\mathrm{A}$ and every bridge residual $r_e(P_{\ell'}\mathrm{A}\mathbf{v}_0)$ with $\ell' < \ell(e)$ stays at least $\mathrm{Am}_L/2$ away from zero. The weighted flag makes residuals of level at most $\ell$ relax at rate of order $\mathrm{J}_\ell$, while bridges of higher level feel only couplings of order $\mathrm{J}_{\ell+1}$. On $I_\ell$ the first have decayed by $\epsilon^{\mathrm{c}_L k_L}$ relative to $\mathrm{A}$, whereas the second have moved by at most $\mathrm{C}_L\mathrm{A}\mathrm{J}_{\ell+1}t_\ell \to 0$ and keep their margin, so each window joins pairs of coherent components along the bridges of level $\ell$ and halves their number. The source record reads residuals at the thermal resolution $\sqrt{\mathrm{T}/\mathrm{J}_1}$, far below the coherence scale $\mathrm{A}$, so it is saturated there. Source jumps, the record, and learning are therefore perturbations, and the content of the theorem is that exact marked parents, activation from the complete law, and own-law learning can be carried through this relaxation with separately vanishing errors.

**Theorem 13.17 (Finite-depth formation under a supplied energy hierarchy).** Assume (H1) the experiment of Definitions 13.1, 13.3, 13.5 and 13.7 with the kinetics of Definitions 13.9 and 13.11, at a fixed depth $L \geq 2$; (H2) the regime eq. (13.29); (H3) an initial law in the preparation neighborhood eq. (13.31); (H4) parents activated at the times eq. (13.32). There exist constants $k_{\mathrm{geom},L}, \mathrm{c}_{\mathrm{early},L}, \mathrm{c}_{\mathrm{joint},L}, \mathrm{c}_{\mathrm{fast},L}, \eta_L^{\max} > 0$, depending only on $L$ and the fixed data and constructed

respectively in Lemma C.12 (together with $\eta_L^{\max}$, a fixed fraction of the margin $\mathrm{m}_L$), Lemma C.7, Lemma C.8, and Lemma C.10, with the following property. For every entrance constant $k_L$ with

$$\begin{aligned} k_L > k_{\mathrm{geom},L}, \quad \tfrac{1}{2}\mathrm{c}_{\mathrm{early},L}k_L > \mathrm{Q}_1 + 20, \\ \mathrm{c}_{\mathrm{joint},L}k_L > \mathrm{Q}_1 + 8, \qquad \mathrm{c}_{\mathrm{fast},L}k_L > \mathrm{Q}_1/2 + 12, \end{aligned} \tag{13.38}$$

and every tolerance $0 < \eta_L \le \eta_L^{\max}$, there exist $\epsilon_L, \mathrm{c}_L, \mathrm{C}_L > 0$, depending on $L$, the fixed data, $k_L$, and $\eta_L$, such that for every $0 < \epsilon < \epsilon_L$ conclusions (a) to (e) hold for the fine process and for the lineage of activated parents. The lineage is said to admit all $L$ successive activations when the fine process and every activated parent have a unique nonexplosive solution of eq. (13.25) through $t_*$, with positive smooth densities and parameter in $[1, 21/10]$.

(a) *Well-posedness and parameter.* The lineage admits all $L$ successive activations, and every process satisfies, in absolute time,

$$1 \le \vartheta_\ell(t) \le 21/10, \qquad |\dot{\vartheta}_\ell(t)| \le \mathrm{C}_L\epsilon, \qquad |\vartheta_\ell(t) - (2 - e^{-\epsilon t})| \le \mathrm{C}_L r_\epsilon, \tag{13.39}$$

with $r_\epsilon = O(\epsilon^{2\mathrm{Q}_1+20})$ of eq. (13.30); the comparison curve solves the calibration part of eq. (13.25).

(b) *Coherence.* For a level-$\ell$ description at time $t$, an original edge $e$ is coherent when $\sqrt{2\mathcal{D}_e^{\ell,\mu}} \le \eta_L\mathrm{A}$, with eq. (13.18) evaluated at the current retained state and parameter of that description, and a vertex is internally coherent when the corresponding observable for $\mathfrak{d}_i$ is at most $\eta_L\mathrm{A}$; the coherent components are the connected components of the graph on the $N$ vertices whose edges are the coherent edges. For each of $\mu_\vartheta = \mathbf{\Pi}_{\vartheta,\mathrm{o}}$ and $\mu_\vartheta = q_\vartheta^\star$, except on an event of probability at most

$$\mathrm{C}_L(1 + t_*)\exp(-\mathrm{c}_L/\epsilon^6), \tag{13.40}$$

the microscopic process and its sequential parent descriptions have the successive coherent components

$$3 \cdot 2^{L-1} \longrightarrow 2^{L-1} \longrightarrow 2^{L-2} \longrightarrow \cdots \longrightarrow 1. \tag{13.41}$$

On $I_\ell$ every edge of level at most $\ell$ is coherent, every vertex is internally coherent, and every bridge of level greater than $\ell$ has residual observable at least a fixed positive multiple of $\mathrm{A}$, so the count on $I_\ell$ is $2^{L-\ell}$; the single final component persists through $t_*$. Loop coordinates remain random and retained, with no identity-holonomy constraint.

(c) *Conditional loss and score.* On each window $I_\ell$, at the current law and parameter,

$$\Delta_{\ell-1:\ell} \le \begin{cases} \mathrm{C}_L\epsilon^{14}, & \ell < L, \\ \mathrm{C}_L\epsilon^2, & \ell = L, \end{cases} \qquad \big|g_{\ell-1}(q,\vartheta) - g_\ell\big((C_{\ell\leftarrow\ell-1})_\#q, \vartheta\big)\big| \le \begin{cases} \mathrm{C}_L\mathrm{T}\epsilon^7, & \ell < L, \\ \mathrm{C}_L\mathrm{T}\epsilon, & \ell = L, \end{cases} \tag{13.42}$$

and $\mathrm{T}\Delta_{\ell-1:\ell}$ is the discarded restricted free energy.

(d) *Source and kernel movement.* Let $\lambda^{\ell}_{i_0,+}(t)$ be the distinguished receiver's source mass in a fine, parent, or sequential process. From each activation time at which that process is continued,

$$\lambda^{\ell}_{i_0,+}(t_*) - \lambda^{\ell}_{i_0,+}(t^{\mathrm{act}}_{\ell'}) > \tfrac{1}{2}, \tag{13.43}$$

so the binary recognition source row changes by total variation greater than one half. Between $\vartheta = 1$ and the final parameter, the first hidden physical conditional and the first hidden optimizing conditional each change by relative entropy greater than $1/200$, at every fixed retained input and assignment. Both changes are driven by the learned parameter through eq. (13.9) and $\vartheta H$, the record moving the distinguished log odds by at most $\delta$, and both take place after formation, since $\vartheta \le 1 + \mathrm{C}_L \epsilon \log(1/\epsilon)$ until the last window closes.

(e) *Histories and paths.* For any source process and its inherited parent begun from a common marked activation sample at $t^{\mathrm{act}}_{\ell}$, there is a deterministic path $\bar{\mathbf{x}}_{\ell} : [t^{\mathrm{act}}_{\ell}, t_*] \to W_{\ell}$ with

$$\Pr\Big\{ \sup_{t^{\mathrm{act}}_{\ell} \le t \le t_*} |P_{\ell}\mathbf{X}_t - \bar{\mathbf{x}}_{\ell}(t)| > \mathrm{C}_L \mathrm{A}\epsilon^2 \Big\} + \Pr\Big\{ \sup_{t^{\mathrm{act}}_{\ell} \le t \le t_*} |\overline{\mathbf{X}}_t - \bar{\mathbf{x}}_{\ell}(t)| > \mathrm{C}_L \mathrm{A}\epsilon^2 \Big\} \le \mathrm{C}_L (1 + t_*) e^{-\mathrm{c}_L/\epsilon^2}. \tag{13.44}$$

There is a coupling preserving the activation sample under which the first bound below holds, and the second holds under every coupling of the two processes:

$$\begin{aligned} &\Pr\big\{ S_t \ne \overline{S}_t \text{ for some } t \in [t^{\mathrm{act}}_{\ell}, t_*] \big\} \le \mathrm{C}_L \epsilon^6, \\ &\Pr\Big\{ \sup_{t^{\mathrm{act}}_{\ell} \le t \le t_*} |P_{\ell}\mathbf{X}_t - \overline{\mathbf{X}}_t| > \mathrm{C}_L \mathrm{A}\epsilon^2 \Big\} \le \mathrm{C}_L (1 + t_*) e^{-\mathrm{c}_L/\epsilon^2}. \end{aligned} \tag{13.45}$$

The resolution $\mathrm{C}_L \mathrm{A}\epsilon^2 = \mathrm{C}_L \sqrt{\mathrm{T}}/\epsilon$ is coarser than the thermal scale, and a finite number of successive replacements changes only the constants. [ESTABLISHED]

*Proof.* The exact statements are Propositions 13.4, 13.6, 13.8 and 13.10, and Table C.1 lists the lemmas of appendix C with their inputs and exponents. Conclusion (a) is Lemmas C.3 and C.4, and (b) is Lemma C.12, which uses neither conditional equilibration nor history averaging. For (c), Lemma C.7 gives the early bound from $\mathrm{J}^2_{\ell+1}/\mathrm{J}_{\ell} = \epsilon^{20}$ and the second entrance condition, and Lemma C.8 the terminal bound by joint diffusion-plus-jump relaxation; conditional relative entropy is bounded by the full one through the chain rule, and the complete hidden score bounds the covector difference by $\mathrm{C}_L \mathrm{T}(\sqrt{\Delta} + \Delta)$. Conclusion (d) is Lemma C.13. For (e), Lemma C.11 couples the assignment histories and Lemma C.14 gives eq. (13.44); each of its events concerns the law of one process, so the union bound gives the second line of eq. (13.45) under every coupling. The four entrance conditions are equivalent to the single threshold

$$k_L > \max\left\{ k_{\mathrm{geom},L},\ \frac{2(\mathrm{Q}_1 + 20)}{\mathrm{c}_{\mathrm{early},L}},\ \frac{\mathrm{Q}_1 + 8}{\mathrm{c}_{\mathrm{joint},L}},\ \frac{\mathrm{Q}_1/2 + 12}{\mathrm{c}_{\mathrm{fast},L}} \right\}, \tag{13.46}$$

whose terms supply, in order, the geometric windows, the early entrance, terminal joint relaxation, and the source-history estimate; every lemma uses the same preparation, regime, and entrance choice. The final window has order $k_L \log(1/\epsilon)$ and lies before $t_* - 10 = 2/\epsilon - 10$ once $\epsilon$ is small;

this, the ordering of the windows, and the smallness conditions inside the lemmas fix $\epsilon_L$ after $k_L$ and $\eta_L$. □

**Remark 13.18 (Saturated record, calibrated learning, and thermally negligible loops).** Uniformly in the configuration, the assignment, and the parameter,

$$\begin{aligned} &|\log \beta_S(\mathbf{x},\vartheta) - \log \pi_S(\vartheta)| \le B_0 = 10^{-4}, \\ &\qquad |H^{-1/2}\mathrm{T}\nabla_{\mathbf{x}}\phi_S(\mathbf{x})| \le B_*\sqrt{\mathrm{T}}, \qquad B_* = \delta\sqrt{2N} = 10^{-4}\sqrt{2/N} \le 5.8\times 10^{-5}, \end{aligned} \tag{13.47}$$

with $H^{-1/2}$ acting on $\mathcal{K}^{\perp}$, by eq. (13.14) with $0 \le \phi_S \le B_0$ and by eq. (C.9). At a residual of order $\mathrm{A}$ the argument of $\chi$ has order $\epsilon^{-(3+\mathrm{Q}_1/2)}$, so the record equals its saturated value up to an exponentially small error. The proof of Lemma C.12 uses the record and the assignments only through $0 \le \phi_S \le B_0$ and the pathwise force bound eq. (C.17), valid for every adapted assignment path, and the parameter only through $1 \le \vartheta \le 21/10$, so with $\delta = 0$ and $\vartheta \equiv 1$ it gives the same components. By Lemmas C.3 and C.4 every law-dependent term of the covector is at most $\mathrm{C}_L r_\epsilon$, which gives the parameter bounds of Theorem 13.17. The loop energy $\mathrm{J}_1|h|^2$ belongs to the strongest term, so $\mathrm{J}_1\mathbb{E}|h|^2 \le \mathrm{C}_L\mathrm{T}$ by eq. (C.15), and the loops have root-mean-square size at most $\mathrm{C}_L\sqrt{\mathrm{T}/\mathrm{J}_1} = \mathrm{C}_L\mathrm{A}\epsilon^{3+\mathrm{Q}_1/2}$. [ESTABLISHED]

## 13.8 A two-triangle instance and the remaining scope

**Remark 13.19 (The two-triangle instance).** For $L = 2$, take triangles $(1,2,3)$ and $(4,5,6)$ with the single bridge $(3,4)$. The retained carriers are

$$\mathbb{R}^{19}\times\mathcal{J} \longrightarrow \mathbb{R}^{9}\times\mathcal{J} \longrightarrow \mathbb{R}^{8}\times\mathcal{J}, \qquad |\mathcal{J}| = 64, \tag{13.48}$$

the coherent component counts are $6 \to 2 \to 1$, both loops survive both maps, and the regime becomes

$$\mathrm{J}_1 = \epsilon^{-20}, \qquad \mathrm{J}_2 = 1, \qquad \mathrm{A} = \epsilon^{40}, \qquad \mathrm{T} = \epsilon^{86}. \tag{13.49}$$

The first window has order $\epsilon^{20}\log(1/\epsilon)$, the second $\log(1/\epsilon)$, and the source and model changes occur over the longer horizon $2/\epsilon$. The source-dependent term is not quadratic: along a direction with selected internal residual slope $\lambda \neq 0$, $\chi(\lambda t) = \lambda^2 t^2/2 - \lambda^4 t^4/8 + O(t^6)$, and every assignment selects an internal edge with nonzero first-normal projection, so the hidden log density has a nonzero fourth derivative in some first-normal direction; exact marginalization keeps the full conditional integral and the negative covariance term of its Hessian, and neither first hidden conditional is a Gaussian. [ESTABLISHED]

**Open Problem 13.20 (A family in which learning or holonomy selects the components).** Construct a normalized family at $c_*$ in which parameter learning or loop holonomy changes, at the level of the law, which components form, with exact marked parents and activation from the complete law. What holds is Theorem 13.17, whose components are fixed by the supplied contrasts

and internal source support, and Proposition 13.22, whose pairs are seeded by the random initial assignment without learning, so that in it the source record selects only along sampled paths; Proposition 14.4 shows that its seeded bound is unchanged by learning and that from an initial law invariant under the cycle automorphisms learning leaves the rows uniform, so a law-level selection needs an initial law without that invariance. Such a family needs at least one of a source record standardized at the coherence scale, for example through $\chi(r/(\eta_L \mathrm{A}))$, model means and model links both slow relative to the belief mobility, since a fast link or fast model means remove the model disagreement on the belief clock whatever the other rate (Proposition 8.6), loop coordinates of order $\mathrm{A}$ or nonabelian links, or interacting model-channel records, together with new fast-marginal and history estimates; sources on different bridge levels and source learning that creates the contrasts belong to the same problem. It is posed on the four-cycle of Proposition 14.4, with its temperature, learning mobility, and horizon, for Gaussian initial laws $\mathcal{N}(\mathrm{m}_0, I_{12})$ with $\mathrm{m}_0$ outside a Lebesgue-null set, and for parents of the type of part (b) of Proposition 13.22 activated from the recognition law of the episode prefix; the law of the formed block means the conditional law, given contrast formation, of the first block for which it occurs, the conditioning event having probability bounded below in every arm by part (b) of Proposition 14.4, so that a difference in whether formation occurs is not selection. It is settled by such parents with separately vanishing conditional and dynamical errors together with a total-variation difference, bounded below as $\epsilon \to 0$, between the conditional laws of the formed block in the learned and in the frozen-row arm, or between those of two initial means that differ only in their loop coordinate $\sum_i \omega_i$, the amplified datum being the initial law or its loop coordinate and the other arm its null; or by a proof that, for almost every $\mathrm{m}_0$, each of these two pairs of arms gives conditional laws of the formed block whose difference vanishes as $\epsilon \to 0$.

A separate slower-mobility regime is resolved by proposition C.17: keeping the same contrast event and horizon but taking $\epsilon_{\mathrm{L}} = o(\mathrm{T}^{1/4})$, the learned and frozen conditional first-block laws have vanishing total variation for every fixed Gaussian mean. The original coupled mobility $\epsilon_{\mathrm{L}} = \mathrm{T}^{1/8}$ and the loop-coordinate comparison remain outside that estimate. The late coherent-partition contrast in proposition C.16 concerns restored equilibrium coherence and does not close this first-block problem. For the late coherence event on the joint slow schedule, propositions C.18 and C.20 establish actual conditioned activation of the existing parent with separately vanishing conditional-information and dynamical errors. The activation event is a completed coherence prefix after selection. It is not the first-contrast event in the original regime posed above.

## 13.9 An endogenous four-cycle witness

**Definition 13.21 (Homogeneous four-cycle experiment).** Index four agents by $i \in \mathbb{Z}/4$ and orient edge $e_i$ from $i$ to $i+1$. A configuration is $\mathbf{x} = (\bar{k}, \bar{m}, \omega) \in \mathbb{R}^{12}$, with $q_i^b = \mathcal{N}(\bar{k}_i, 1)$, $q_i^m = \mathcal{N}(\bar{m}_i, \sigma_{\mathrm{ev}}^2)$, $\sigma_{\mathrm{ev}}^2 = 1/2$, four independent link translations $\omega_i$, residuals $r_i^b = \bar{k}_{i+1} - \bar{k}_i - \omega_i$ and $r_i^m = \bar{m}_{i+1} - \bar{m}_i - \omega_i$, and the edge cost

$$\mathcal{D}_{e_i}(\mathbf{x}) = \tfrac{1}{2}\big[(r_i^b)^2 + (r_i^m)^2\big] = D_{\mathrm{KL}}\big(q_{i+1}^b \| (\Theta_{e_i})_{\#} q_i^b\big) + \sigma_{\mathrm{ev}}^2 D_{\mathrm{KL}}\big(q_{i+1}^m \| (\Theta_{e_i})_{\#} q_i^m\big), \tag{13.50}$$

which is the same for both orientations. Every receiver has both neighbors as sources, $S \in \mathcal{J} = \prod_i \{i-1, i+1\}$ with $|\mathcal{J}| = 16$ and uniform prior $\pi_S = 1/16$, and $V_S = \sum_i \mathcal{D}_{e(i,S_i)}$, where $e(i,j)$ joins $i$ and $j$. The configuration prior is $\mathcal{N}(0, I_{12})$; given $(\mathbf{x}, S)$, independent binary records have $\Pr(E_i = 1 \mid \mathbf{x}, S) = \frac{1}{2}\exp[-\mathcal{D}_{e(i,S_i)}(\mathbf{x})/\mathrm{T}]$, and $E = \mathbf{1}$ is observed; the physical evaluator and recognition coupling are those of Definition 13.3. With $\mathfrak{D}\mathbf{x} = \bar{k} - \bar{m}$ and $R_{\mathrm{ev}} = 2\mathfrak{D}^{\top}\mathfrak{D}$, the restricted optimizer is $q^{\star}(\mathbf{x}, S) \propto \pi_S e^{-U_S(\mathbf{x})/\mathrm{T}}$ with $U_S = V_S + \frac{\mathrm{T}}{2}|\mathbf{x}|^2 + \frac{\mathrm{T}}{2}\mathbf{x}^{\top} R_{\mathrm{ev}}\mathbf{x}$, and the physical posterior omits the last term. The kinetics are identity mobility with diffusion covariance $2\mathrm{T}I\,dt$ and one rate-one clock at which the complete assignment is redrawn from the product of the receiver posterior rows $\beta_S(\mathbf{x})$, self-proposals allowed; no generative parameter is learned. The two perfect-matching assignments are $S^{(0)} = (1, 0, 3, 2)$, which selects $e_0, e_2$ from both ends, and $S^{(1)} = (3, 2, 1, 0)$. For $0 < \epsilon \le 1/100$ put

$$\mathrm{T} = \epsilon^8, \quad t_\epsilon = 2\log(1/\epsilon), \quad t_*^{\circ} = t_\epsilon + 1, \quad \varrho_\circ = \tfrac{1}{100}, \quad \sigma_\circ = \varrho_\circ/\sqrt{48}, \quad \mathbf{x}^{\circ} = (0, 0, \mathbf{1}_4). \tag{13.51}$$

The preparation class consists of the laws $\frac{1}{16}\mathfrak{p}_\circ(\mathbf{x})(1 + \mathfrak{z}(\mathbf{x}))d\mathbf{x}$ on $(\mathbf{x}, S)$, where $\mathfrak{p}_\circ$ is the equal mixture of $\mathcal{N}(\mathbf{x}^{\circ}, \sigma_\circ^2 I)$ and $\mathcal{N}(-\mathbf{x}^{\circ}, \sigma_\circ^2 I)$ and $\mathfrak{z}$ is smooth, invariant under the cycle automorphisms (with link signs transformed under reflections), of $\mathfrak{p}_\circ$-mean zero, and with $\|\mathfrak{z}\|_\infty < 1/2$.

The preparation names no block: its assignment is uniform and independent of the configuration, reflections exchange its two centers, which carry opposite loop orientations, and all initial edge mismatch second moments agree. Under the source gating an edge that neither endpoint selects carries no interaction energy, although every edge stays in every palette and every attention weight stays positive.

**Proposition 13.22 (Endogenous pairing on a homogeneous four-cycle).** In Definition 13.21, let $0 < \epsilon \le 1/100$ and let the initial law belong to the preparation class. The joint is normalized, the scaled negative ELBO is $\mathcal{F}_{\mathrm{T}}(q) = \mathrm{T}D_{\mathrm{KL}}(q\|q^{\star}) + \mathrm{B}_{\mathrm{T}} - \mathrm{T}\log Z_{\mathrm{opt}}$ with $\mathrm{B}_{\mathrm{T}} = \mathrm{T}[6\log(2\pi) + 4\log 2]$, and the recognition law of the kinetics has nonincreasing $\mathcal{F}_{\mathrm{T}}$ and is invariant under the cycle automorphisms at every time. Put

$$p_\epsilon = 73728\epsilon^4 t_*^{\circ} + 4t_*^{\circ} e^{-1/\mathrm{T}} < \tfrac{1}{2}, \qquad \eta_\epsilon = \tfrac{9}{2}\varrho_\circ^2\epsilon^8 + \frac{3}{4(1-e^{-8})}\epsilon^{32}, \qquad \bar{\eta}_\epsilon = \frac{\eta_\epsilon + \mathcal{S}_2(p_\epsilon)}{1 - p_\epsilon}, \tag{13.52}$$

with $\mathcal{S}_2(p) = -p\log p - (1-p)\log(1-p)$. Then:

(a) *Episode.* With probability at least $p_0 = 3e^{-1/4}/512$, the complete assignment equals one perfect matching throughout $[0, t_*^{\circ}]$; at time zero the internal belief transported relative entropy of each matched pair exceeds $1/4$, so that the contrast is entered during the episode and is not present at its start; on $[t_\epsilon, t_*^{\circ}]$ both internal belief and model transported relative entropies of each matched pair are at most $\epsilon^4$; and on $[1/4, t_*^{\circ}]$ every exterior belief relative entropy exceeds $1/2$, the model ones being larger. This is a coherence episode held by the two attended edge costs: on the held matching sheet the normal restoring rates of the kinetics are $8 + \mathrm{T}$ and $4 + 5\mathrm{T}$, those of the physical law $8 + \mathrm{T}$ and $4 + \mathrm{T}$, so at the residual resolution $\epsilon^2$ the least restoring rate times the squared resolution, divided by the temperature, is $4\epsilon^{-4}$, and the sheet is selected by the posterior

proposals on the event. The kinetics leave $q^\star$ invariant, and started from $q^\star$ they realize the same event with probability at most $4e^{-1/(4\mathrm{T})}$, so the episode exceeds its stationary probability by at least $p_0 - 4e^{-1/(4\mathrm{T})} > p_0/2$.

(b) *Parent channel.* For each matching $M$, let $Z_M$ be the span of the three orthonormal directions $(n_e^+ - n_f^+)/\sqrt{2}$, $n_e^-$, $n_f^-$, where $e, f$ are its two edges, $n_e^\pm$ the unit normal covectors of eq. (C.94), and $W_M = Z_M^\perp$. The channel $C(\mathbf{x}, S) = (S, P_{W_S}\mathbf{x})$, with $W_S = W_M$ on the two matchings and $W_S = \mathbb{R}^{12}$ otherwise, extended by the identity on observations and evaluator samples, keeps every assignment, the loop $\sum_i \omega_i$, and a fast common-normal coordinate, and pushes the joint, the physical posterior, the restricted optimizer, and the recognition law to normalized parent laws.

(c) *Discarded conditional free energy.* For every $t \in [t_\epsilon, t_*^\circ]$, the recognition law $q_t$ of the episode has mean discarded optimizing conditional relative entropy at most $\bar{\eta}_\epsilon$ and mean discarded physical conditional relative entropy, with the evaluator samples retained, at most $\bar{\eta}_\epsilon + \mathrm{T}$.

(d) *Parent dynamics.* The inherited parent of $C$, activated at $t_\epsilon$ from the complete recognition law of the episode prefix, has retained path law on $[t_\epsilon, t_*^\circ]$, including the complete assignment path, within total variation $2p_\epsilon/(1-p_\epsilon) + 4\sqrt{8/3}e^{-1/\mathrm{T}}$ of the retained native path conditioned on the same prefix.

The estimates hold under the stated source gating, random matching seed, symmetric circulated preparation, temperature $\epsilon^8$, Euclidean kinetic metric, duration-one episode, and fixed generative parameters. [ESTABLISHED]

The proof is in appendix C.7. On the center path eq. (C.95), which holds $S^{(0)}$ from $\mathbf{x}^\circ$, where every residual equals $-1$, the selected residuals decay as $-e^{-(8+\mathrm{T})t}$ while the exterior ones grow toward $-3/2$, so a source-cost gap greater than one develops by time $1/4$, and the posterior proposals then keep the matching except with probability at most $4t_*^\circ e^{-1/\mathrm{T}}$. Since the full recognition law is invariant, the pairing is an event of sampled paths seeded by the initial assignment, not a block chosen by the law.

The witness raises the question whether its mechanism, an edge that carries energy only while an endpoint attends to it, produces a stationary law that favors coherence beyond one level, or a second coalescence of the pairs it forms. The next statement answers the stationary part negatively in the class to which the four-cycle belongs, the family of Definition 13.23 on an arbitrary finite graph; the second coalescence is posed for models outside that family in Open Problem 13.25. Under exclusive gating the source label enters only the weight of a sheet and never its stiffness, so every sheet restores its attended residual directions at rates of order one and leaves every other direction, including every direction that translates one coherent block relative to another, to the thermal reference at rates of order $\mathrm{T}$. This has a spectral and a stationary consequence.

**Definition 13.23 (Gated translation family on a graph).** Let $(V, E)$ be a finite simple graph without isolated vertices, with $N = |V|$ agents and matching number $\mathfrak{n}$, the largest number of pairwise disjoint edges. Orient each edge $e \in E$ once as $e : j \to i$; the interaction multigraph at $c_*$ consists of the copies $e$ and $\bar{e}$ of the edges of $(V, E)$, and the source set $\mathcal{N}_i$ of receiver $i$ is its set of neighbors in $(V, E)$, with $e(i, j)$ the edge joining $i$ and $j$. The agents, links, laws,

evaluator, and recognition coupling are those of Definition 13.21 on this graph: a configuration is $\mathbf{x} = (\bar{k}, \bar{m}, \omega) \in \mathbb{R}^{2N+|E|}$ with one link translation $\omega_e$ per edge, $q_i^b = \mathcal{N}(\bar{k}_i, 1)$, $q_i^m = \mathcal{N}(\bar{m}_i, 1/2)$, and edge $e : j \to i$ carries the residual pair and the cost

$$\mathbf{a}_e\mathbf{x} = (r_e^b, r_e^m) = (\bar{k}_i - \bar{k}_j - \omega_e,\ \bar{m}_i - \bar{m}_j - \omega_e), \qquad \mathcal{D}_e(\mathbf{x}) = \tfrac{1}{2}|\mathbf{a}_e\mathbf{x}|^2, \tag{13.53}$$

so $\mathbf{a}_e : \mathbb{R}^{2N+|E|} \to \mathbb{R}^2$ is linear and invariant under passive frame changes as in Lemma 13.2, and $\mathcal{D}_e$ is the transported relative entropy eq. (13.50). Each receiver carries one source label of Definition 6.1 with source set $\mathcal{N}_i$, so a complete assignment is $S \in \mathcal{J} = \prod_i \mathcal{N}_i$; the labels are exclusive, the prior is the product $\pi_S = \prod_i \pi_{iS_i}$ with every $\pi_{ij} > 0$ independent of $\mathrm{T}$, and given $(\mathbf{x}, S)$ independent binary records have $\Pr(E_i = 1 \mid \mathbf{x}, S) = \frac{1}{2}\exp[-\mathcal{D}_{e(i,S_i)}(\mathbf{x})/\mathrm{T}]$, with $E = \mathbf{1}$ observed. The attended edge set, sheet energy, and sheet stiffness of $S$ are

$$\mathrm{att}(S) = \{e(i, S_i) : i \in V\}, \qquad V_S = \sum_i \mathcal{D}_{e(i,S_i)} = \tfrac{1}{2}\mathbf{x}^\top H_S \mathbf{x}, \qquad H_S = \sum_i \mathbf{a}_{e(i,S_i)}^\top \mathbf{a}_{e(i,S_i)}. \tag{13.54}$$

The configuration prior is $\mathcal{N}(0, I_{2N+|E|})$, and with $R_{\mathrm{ev}} = 2\mathfrak{D}^\top\mathfrak{D}$ the restricted optimizer and the physical posterior are the laws

$$q_\Omega(\mathbf{x}, S) = \frac{\pi_S}{\sum_{S'} \pi_{S'} Z_{S'}} \exp\Big[-\frac{V_S(\mathbf{x})}{\mathrm{T}} - \tfrac{1}{2}\mathbf{x}^\top\Omega\mathbf{x}\Big], \qquad Z_S = \int \exp\Big[-\frac{V_S(\mathbf{x})}{\mathrm{T}} - \tfrac{1}{2}\mathbf{x}^\top\Omega\mathbf{x}\Big] d\mathbf{x}, \tag{13.55}$$

with the reference precision $\Omega = I + R_{\mathrm{ev}}$ for the optimizer and $\Omega = I$ for the physical posterior; $Z_S$ is the sheet weight. For either $\Omega$ the kinetics have identity mobility, diffusion covariance $2\mathrm{T}I\,dt$, drift $-\nabla_{\mathbf{x}} U_S$ with $U_S = V_S + \frac{\mathrm{T}}{2}\mathbf{x}^\top\Omega\mathbf{x}$, and one rate-one clock at which the complete assignment is redrawn from the product of the receiver posterior rows $\beta_i(j \mid \mathbf{x}) = \pi_{ij} e^{-\mathcal{D}_{e(i,j)}(\mathbf{x})/\mathrm{T}} \big/ \sum_{j' \in \mathcal{N}_i} \pi_{ij'} e^{-\mathcal{D}_{e(i,j')}(\mathbf{x})/\mathrm{T}}$, self-proposals allowed; no generative parameter is learned. For an edge set $F \subseteq E$, $\mathbf{a}_F$ stacks the maps $\mathbf{a}_e$, $e \in F$, $W_F = \ker \mathbf{a}_F$, $\mathrm{s}_F$ is the least positive singular value of $\mathbf{a}_F$, and $\mathrm{cc}(F)$ is the number of connected components of the graph $(V, F)$, isolated vertices included. At resolution $\mathrm{a} > 0$ the coherent edges of a configuration are $\mathrm{Coh}_\mathrm{a}(\mathbf{x}) = \{e \in E : |\mathbf{a}_e\mathbf{x}| \le \mathrm{a}\}$, and $\mathfrak{P}_\mathrm{a}(\mathbf{x})$ is the partition of $V$ into the vertex sets of the components of $(V, \mathrm{Coh}_\mathrm{a}(\mathbf{x}))$, the coherent blocks.

The four-cycle of Definition 13.21 is the member with $(V, E)$ the four-cycle and the uniform prior, and its kinetics are those with $\Omega = I + R_{\mathrm{ev}}$. A minimum edge cover of $(V, E)$ is an edge set covering every vertex with the fewest edges; the proposition identifies these covers with the spanning forests of stars with $\mathfrak{n}$ components.

**Proposition 13.24 (Depth one under exclusive source gating).** In Definition 13.23, for either reference precision $\Omega$, let $\mathrm{k}_S = \operatorname{rank} H_S$, let $\mathrm{h}_S$ be the least positive eigenvalue of $H_S$, $\mathrm{h}_* = \min_S \mathrm{h}_S$, and $\mathrm{z}_S = [\mathrm{pdet}\, H_S \cdot \det(\Omega|_{W_{\mathrm{att}(S)}})]^{-1/2}$, where $\mathrm{pdet}$ is the product of the positive eigenvalues and $\Omega|_W = \mathrm{Q}_W^\top \Omega \mathrm{Q}_W$ for a matrix $\mathrm{Q}_W$ whose columns form an orthonormal basis of $W$; let $\mathcal{M}$ be the set of assignments of minimum rank $\mathrm{k}_S$.

(a) *Two restoring bands.* For every $S$ and $\mathrm{T}>0$, the eigenvalues $\lambda_1 \le \cdots \le \lambda_{2N+|E|}$ of $H_S+\mathrm{T}\Omega$ satisfy $\lambda_k \in [\mathrm{T}\|\Omega^{-1}\|^{-1}, \mathrm{T}\|\Omega\|]$ for $k \le 2N+|E|-\mathrm{k}_S$ and $\lambda_k \in [\mathrm{h}_S, \|H_S\|+\mathrm{T}\|\Omega\|]$ for $k > 2N+|E|-\mathrm{k}_S$. The range of $H_S$ is $W_{\mathrm{att}(S)}^{\perp}$, the span of the attended residual covectors, and the prior enters $q_\Omega$ only through the sheet weights $\pi_S \mathrm{Z}_S$, never through $H_S$.

(b) *Gating wells.* With

$$V_\pi(\mathbf{x}) = -\mathrm{T}\sum_i \log \sum_{j\in\mathcal{N}_i} \pi_{ij} e^{-\mathcal{D}_{e(i,j)}(\mathbf{x})/\mathrm{T}}, \qquad V_{\min}(\mathbf{x}) = \sum_i \min_{j\in\mathcal{N}_i} \mathcal{D}_{e(i,j)}(\mathbf{x}),$$

one has

$$0 \le V_\pi(\mathbf{x}) - V_{\min}(\mathbf{x}) \le \mathrm{T}\sum_i \log \frac{1}{\min_{j\in\mathcal{N}_i}\pi_{ij}} \qquad \text{for every } \mathbf{x}, \tag{13.56}$$

the configuration marginal of $q_\Omega$ has density proportional to $\exp[-V_\pi(\mathbf{x})/\mathrm{T} - \tfrac{1}{2}\mathbf{x}^\top\Omega\mathbf{x}]$, and if every edge of an edge set $F$ covering $V$ has zero residual at $\mathbf{x}$, then $V_{\min}$ vanishes at $\mathbf{x}$ and at every configuration obtained from it by translating the belief and model means of the agents of one component of $(V,F)$ by a common amount with the links unchanged.

(c) *Rank, minimum edge covers, and sheet weights.* For every $F\subseteq E$, $\operatorname{rank}\mathbf{a}_F = |F|+N-\mathrm{cc}(F)$. Hence $\mathrm{k}_S = |\operatorname{att}(S)|+N-\mathrm{cc}(\operatorname{att}(S)) \ge 2(N-\mathfrak{n})$, with equality exactly when $\operatorname{att}(S)$ is a minimum edge cover of $(V,E)$; the minimum edge covers are the spanning forests of stars with $\mathfrak{n}$ components, every one of them is attended by some assignment, and an edge outside $\operatorname{att}(S)$ joining two of its components raises the rank by two. The sheet weight satisfies

$$\Big(1+\frac{\mathrm{T}\|\Omega\|}{\mathrm{h}_S}\Big)^{-\mathrm{k}_S/2} \le \frac{Z_S}{(2\pi)^{(2N+|E|)/2}\mathrm{z}_S\mathrm{T}^{\mathrm{k}_S/2}} \le 1, \tag{13.57}$$

and consequently, for $0<\mathrm{T}\le 1$ and every integer $m\ge 1$,

$$q_\Omega\{S : \mathrm{cc}(\operatorname{att}(S)) \le \mathfrak{n}-m\} \le \mathrm{K}_m\mathrm{T}^m, \qquad \mathrm{K}_m = \Big(1+\frac{\|\Omega\|}{\mathrm{h}_*}\Big)^{N-\mathfrak{n}} \frac{\sum_{S:\mathrm{cc}(\operatorname{att}(S))\le\mathfrak{n}-m}\pi_S\mathrm{z}_S}{\sum_{S\in\mathcal{M}}\pi_S\mathrm{z}_S}. \tag{13.58}$$

(d) *Stationary block law.* Let $\mathrm{s}_*$ be the least, over all $S$ and all edges $e\notin\operatorname{att}(S)$ joining two components of $\operatorname{att}(S)$, of the least positive singular value of $\mathbf{a}_e P_{W_{\operatorname{att}(S)}}$, where $P_W$ is the orthogonal projection onto $W$. When no such pair $(S,e)$ exists, put $\mathrm{s}_*=+\infty$ and $1/\mathrm{s}_*^2=0$. For $\mathrm{a}>0$ and $0<\mathrm{T}\le 1$, the probability under $q_\Omega$ that $\mathfrak{P}_{\mathrm{a}}(\mathbf{x})$ is not the partition of $V$ into the components of a minimum edge cover is at most

$$\mathrm{K}_1\mathrm{T} + \frac{|E|\|\Omega\|}{2\mathrm{s}_*^2}\mathrm{a}^2 + 4|E|e^{-\mathrm{a}^2/(4\mathrm{T})}. \tag{13.59}$$

[ESTABLISHED]

The proof is in appendix C.8. In the four-cycle, $\mathfrak{n}=2$ and the minimum edge covers are the two perfect matchings, so (c) gives them all but $O(\mathrm{T})$ of the stationary mass of either law and (d) makes coherence of an unattended bridge at resolution $\mathrm{a}$ cost $O(\mathrm{a}^2)$ per bridge. A second coalescence

that is held by an attended bridge at stationarity with probability bounded below as $\mathrm{T} \to 0$ therefore needs a model outside the family.

**Open Problem 13.25 (Repeated endogenous formation).** Construct, at $c_*$, a second endogenous coalescence with inherited source interactions and retained marks, in a family where a matching or partition forms from initially nonmatching source assignments, generative parameters are learned, and the episode length is not fixed in advance. What holds is Proposition 13.22, one duration-one episode seeded by the initial assignment, Theorem 13.17, whose hierarchy is supplied, and Proposition 13.24, by which in the gated translation family of Definition 13.23 the stationary coherent-block partition is, up to the error of eq. (13.59), the partition into the components of a minimum edge cover, and every further merge at stationarity costs a factor $\mathrm{T}$. The question includes arbitrary families and initial laws, mixed-level source support, permanent membership under later forcing, and control as the depth tends to infinity; the intrinsic selector and the hierarchy question are Open Problems 9.19 and 12.27. It is posed for models at $c_*$ with source-gated attention on a finite connected graph whose energies, precisions, prior ratios, hyperparameters, and mobilities do not depend on the temperature $\mathrm{T}$, as $\mathrm{T} \to 0$, outside the family of Definition 13.23. A second coalescence in such a model is declared by a handoff time, deterministic or a stopping time; a first-stage event of the history up to the handoff, with its block partition; a second-stage event, with windows measured from the handoff, whose block is a union of at least two first-stage blocks; one statistic of the handoff state, namely the probability of the second-stage event, the probability that it has occurred by declared times, or the probability of a declared block; and a handoff class, fixed before the comparison, that contains every first-stage handoff and, with probability bounded below, the handoffs of histories without first-stage formation. The problem is settled by such a coalescence, with learned generative parameters and from initially nonmatching assignments, whose second-stage event does not occur with probability tending to one uniformly over the handoff class, whose statistic after first-stage formation exceeds its value after a counterfactual handoff in the class by a margin bounded below, whose block is held at a declared resolution by a restoring interaction whose rate times the squared resolution, divided by $\mathrm{T}$, tends to infinity on the sheets the kinetics select, on horizons short compared with the stationary recurrence time of the event measured by its rate of entrance from a set of stationary mass bounded below, and with parents whose conditional and dynamical errors vanish separately; or by a proof that in this class every coalescence of first-stage blocks on such horizons fails the margin or the restoring clause or has a second-stage event occurring with probability tending to one uniformly over the handoff class.

## Physicist's summary

In the triangle chain a supplied hierarchy of stiffnesses $\mathrm{J}_1 \gg \cdots \gg \mathrm{J}_L = 1$ makes the prepared displacement relax like overdamped springs of very different stiffness: the stiffest bonds equilibrate first and glue their endpoints, and each softer level glues the clusters already formed (Theorem 13.17). Every parent keeps every loop and source assignment, starts from the complete current law, and learns its own parameter; the price of each reduction vanishes in an extreme regime (temperature

$\epsilon^{86}$ at depth two) in which nothing adaptive decides the outcome. In the homogeneous four-cycle, without a supplied hierarchy, an edge carries energy only when an endpoint attends to it; attended edges contract, the prepared circulation stretches the unattended ones, and the posterior keeps attending to the contracted edges, so a pair forms with probability bounded away from zero (Proposition 13.22). The same gating is why the story stops there in that family (Proposition 13.24): an unattended edge between two formed blocks feels only the thermal reference, so at stationarity the blocks are the components of a minimum edge cover and each further merge costs a factor $\mathrm{T}$.

# Chapter 14

# Interpretation and Its Limits

What do the preceding chapters commit to, and which readings of them could an observation or a proof count against? The chapter's readings are definitions, hypotheses, and open problems, and it proves a proposition and a theorem. Proposition 14.4 shows that on the four-agent cycle of section 13.9 a collective forms from every Gaussian initial law with probability bounded below uniformly in the regime parameter, whether or not the source prior learns, because the random initial assignment seeds the pair; its proof is in appendix C.7. The main statement, theorem 14.5, sharpens the participatory reading: at a sufficiently small fixed temperature, slow learning of the source prior selects one pairing at the level of the law from almost every Gaussian initial law, while a frozen prior leaves the long-time pairing law symmetric. It feeds open problem 14.8.

Beyond section 1.4, the chapter declines the claims that a structural-realist, Neo-Kantian, or noumenal reading follows from the invariance results or is selected by the mathematics; that the lineage citations below prove the probability construction or establish its interpretation; that definition 14.1 follows from constructive empiricism; that participation entails consciousness or that variational free energy is a physical action principle; and that a constructed record sensitive to link holonomy is a canonical observable or an empirical realization. Its open items are open problems 14.7 and 14.8. [NOT-CLAIMED]

## 14.1 Scaffolding, idle posits, and structural readings

The verdict of the preceding chapters is that the architecture is local-first, not background-free. Agent $i$ carries the law points

$$q_i^b \in (\mathcal{E}_b)_{c*}, \qquad q_i^m \in (\mathcal{E}_m)_{c*}, \tag{14.1}$$

which the cross maps $\Phi$ and $\widetilde{\Phi}$ do not identify, and the base, principal bundle, and measurable spaces are coordination scaffolds that make agent-relative interfaces comparable and that local inference laws do not construct. The verdict would change if edges became random structural coordinates updated by inference (open problem 6.9) or links were recovered from population records, with single-point part open problem 14.8. A Neo-Kantian reading takes local frames, model evaluations, and access maps as constitutive but revisable conditions of a record's inferential meaning, in the lineage of Kant, Reichenbach, Cassirer, and Friedman [17, 31, 50, 69]; a noumenal reading takes the base and bundles as structure reached only through agent-side representations. As scaffolding is no inferential target, the base's lack of a law decides neither reading.

**Definition 14.1 (Idle-wheel criterion).** A posit is idle when changing it alters the law of no record and of no statistic of a record in the model. The idle-wheel criterion removes idle posits by parsimony. It is a declared rule of reading.

The criterion is close in spirit to constructive empiricism, on which accepting a theory involves belief only in its empirical adequacy [78, 11–12], but it is declared, not derived. A closed walk $p$ and an edge copy $e$ determine the classes

$$[\mathrm{Hol}^b(p)]_{\mathrm{conj}}, \qquad [\mathrm{Hol}^m(p)]_{\mathrm{conj}}, \qquad [\mathring{\Theta}^x_e]_{\mathrm{conj}} \tag{14.2}$$

of the link holonomies (2.30) and the per-edge twists, invariant under passive and active changes (proposition 2.12). On a graph-flat family the loop classes are trivial (proposition 2.13) and give a noumenal reading no work; a nontrivial one is a referent only within the interaction complex, made operational at $c_*$ by theorem 10.28.

The interface theory of perception permits an arbitrary measurable relation between representation and world [42, 43]. Here the principal fiber at $c_*$ is a $G$-torsor, a frame is the choice of a reference point, and what survives a change of frame is structural data,

$$\mathcal{I}_{\mathrm{struct}} = \Big\{\text{ELBO and relative-entropy values},\ [\mathrm{Hol}^b(p)]_{\mathrm{conj}},\ [\mathrm{Hol}^m(p)]_{\mathrm{conj}},\ \big[(\mathring{\Theta}^b_e, \mathring{\Theta}^m_e)_e\big]\Big\}, \tag{14.3}$$

the last entry being the per-edge twists of both channels modulo one simultaneous conjugation, invariant by proposition 4.9, (3.24), and proposition 2.12. Of the three agreements of proposition 2.18, frame agreement is independent of the other two, law coherence implies mean agreement but not conversely, and none selects a partition (open problem 9.19). These invariants are compatible with epistemic structural realism [86], with moderate structural realism, on which objects share the ontological footing of relations and are characterized only by them [29], and with ontic structural realism, on which structure is all that exists [54, 55]. No argument from theory change is given, since these are invariants of passive frame change within one theory whereas Worrall's structure is retained across theories, and nothing here selects among these positions or the Neo-Kantian reading. Support would need, respectively, an invariant retained across a change of theory, a typed interpretation of the formalism's objects and relations in the moderate ontology, and an argument from formal unidentifiability to ontological elimination with discriminating evidence.

## 14.2 Participation, formation, and testability

**Hypothesis 14.2 (Operational observational closure).** Every datum directly available to an interior agent is a clamped state or a message of another interior agent or of a boundary environment node, an operational environment node in the sense of definition 5.16. The hypothesis excludes data that reach a record kernel other than through such a node. It does not remove the conditioning record, does not require a boundary node to be autonomous, and does not exclude an unobserved environment that affects all agents without becoming the argument of any agent's belief.

By theorem 5.17 the hypothesis is free for kernels but removes no conditioning: a statistic of the record is posterior-sufficient exactly when the conditional mutual information in (5.35) vanishes, and a latent-independent kernel is a zero-information boundary message (5.34), not an agent.

**Hypothesis 14.3 (Every boundary node is called an agent).** Every boundary environment node of hypothesis 14.2 is called an agent. The hypothesis is a naming decision. It adds to no node a persistent state, an action, a Markov blanket, or a local free energy, and it excludes reading the constructed joint $\mathbb{P}_{\theta,V}$ or the selected coupling $\mathbb{Q}_{V,o,X}$ as an agent: these are complete joint laws at the population level. A block conditional is a meta-agent only after its own coarse state, observation interface, evaluated generative mechanism, recognition law, and update rule have been constructed, as in chapters 11 and 12. A strong agent-only ontology would additionally need a criterion of agency for every boundary node, and the formalism supplies no unique one.

Participation is reciprocal inferential dependence through records inside one joint, with no live posterior in a same-step generative factor (hypothesis 4.3). Section 12.7 locates it in three sequences of fixed joints: the same-time inference loop, block coordinate ascent as in any connected hierarchical model under mean-field updates, which resembles Wheeler's participatory proposal [84] without being derived from it or evidence for it; parameter learning, whose optimal prior over an unrestricted family is the aggregate recognition law (12.34); and delayed feedback through (12.35), which for an update reading a declared recognition kernel is the free energy of one joint (proposition 12.25) and whose dynamic closure is part of open problem 11.38. None of these separates the program from a hierarchical model with learned priors; formation could. Theorem 13.17 forms successive collectives under a supplied energy hierarchy and source support, without showing that source learning creates the contrasts, and proposition 13.22 forms one collective on a homogeneous four-cycle with probability bounded below uniformly in its regime parameter, from a symmetric preparation whose random initial source assignment seeds the pair, and without parameter learning. The same seed forms a collective from every Gaussian initial law, whether or not the source prior learns.

**Proposition 14.4 (Seeded contrast formation with idle source learning).** Take the homogeneous four-cycle experiment of definition 13.21, with its regime parameter $0 < \epsilon \leq 1/100$, the constants of eq. (13.51), and its records, evaluator, reference law, and kinetics, and modify it in two ways. First, the uniform prior on assignments becomes a product $\pi_S = \prod_i \pi_{i,S_i}$ of receiver rows, each a generative parameter that starts uniform and learns from the recognition law $q_t$ of the kinetics by

$$\dot{\pi}_{ij} = \epsilon\big(q_t(S_i = j) - \pi_{ij}\big), \tag{14.4}$$

the natural-gradient flow of the free energy in the prior row under the Fisher metric of the simplex at mobility $\epsilon$; its rest point is the recognition label marginal, as for the aggregate prior (12.34), and the assignment is redrawn from posterior rows computed with the current prior. Second, the initial law is $\mathcal{N}(\mathrm{m}_0, I_{12})$ on configurations, with an arbitrary mean $\mathrm{m}_0 \in \mathbb{R}^{12}$, times the uniform law on assignments. Call contrast formation the event that some connected block $I$ with $2 \leq |I| \leq |V| - 1$ has, at time zero, an internal edge whose belief transported relative entropy exceeds $1/4$, and,

on some time interval of length one inside $[0, 2/\epsilon]$, all its internal belief and model transported relative entropies at most $\epsilon^4$ while every belief transported relative entropy on an edge joining it to its complement exceeds $1/2$, the bounds of part (a) of proposition 13.22; the first clause makes the event an entrance into the contrast and excludes a contrast present at time zero that persists. Then:

(a) *Rows.* The kinetics and eq. (14.4) have exactly one joint solution on $[0, \infty)$, with continuous deterministic rows, and $|\pi_{ij}(t) - 1/2| \le \epsilon t/2$ for all $t \ge 0$.

(b) *Formation.* For every $\mathrm{m}_0 \in \mathbb{R}^{12}$ and every $\epsilon \in (0, 1/100]$,

$$\begin{aligned}\Pr(\text{contrast formation}) &\ge \frac{e^{-1/4}}{16} \sum_{\pm} \mathcal{N}(\mathrm{m}_0, I_{12})\big\{|\mathbf{x} \mp \mathbf{x}^\circ| < \varrho_\circ\big\} \\ &\ge \frac{e^{-1/4}\varrho_\circ^{12}}{737280} \exp\Big[-\tfrac{1}{2}\big(|\mathbf{x}^\circ - \mathrm{m}_0| + \varrho_\circ\big)^2\Big] > 0,\end{aligned}$$

and the same bound holds when the rows are frozen at $1/2$. The bound is carried by the sub-event in which $I$ is a matched pair of the initial assignment, held throughout $[0, t_*^\circ]$, and the contrast is that of part (a) of proposition 13.22. This sub-event is a coherence episode with the restoring interaction and rates of that part, and for every fixed row configuration $\pi$ the kinetics at rows $\pi$ leave $q_\pi^\star(\mathbf{x}, S) \propto \pi_S e^{-U_S(\mathbf{x})/\mathrm{T}}$ invariant and, started from it, realize the sub-event with probability at most $4e^{-1/(4\mathrm{T})}$; so for every $\mathrm{m}_0$, in either arm, the sub-event exceeds every such stationary value by a margin bounded below as $\epsilon \to 0$, the first display not depending on $\epsilon$.

(c) *Invariant initial laws.* The initial law is invariant under the cycle automorphisms, with link signs transformed under reflections, exactly when $\mathrm{m}_0 = (\mathrm{c}_1\mathbf{1}_4, \mathrm{c}_2\mathbf{1}_4, 0)$ for some $\mathrm{c}_1, \mathrm{c}_2 \in \mathbb{R}$, and then $\pi_{ij}(t) = 1/2$ for every receiver, source, and $t \ge 0$. [ESTABLISHED]

The proof is in appendix C.7. The probability in part (b) is supplied by the preparation's matching seed, an initial assignment drawn uniformly and independently of a configuration law of full support, and source learning is idle in it. A statement in which learning creates a contrast must therefore compare learned and frozen priors and concern what the law selects, not the probability of a sample-path episode.

**Theorem 14.5 (Pairing selection by slow source learning).** Take the homogeneous four-cycle experiment of definition 13.21, with its records, evaluator, reference law, and kinetics, at a fixed temperature $\mathrm{T} > 0$ in place of $\epsilon^8$; let the product receiver rows of proposition 14.4 start uniform and learn by eq. (14.4) with a mobility $\epsilon_\mathrm{L} > 0$ in place of $\epsilon$; and let the initial law be $\mathcal{N}(\mathrm{m}_0, I_{12})$ on configurations times the uniform law on assignments. Write $q_t^{\epsilon_\mathrm{L}}$ for the recognition law of the kinetics and $q_t^0$ for that of the frozen-prior arm, whose rows stay at $1/2$. First, there is $\mathrm{T}_0 > 0$ such that for every $\mathrm{T} \in (0, \mathrm{T}_0]$ there is $\mathrm{C}_\mathrm{T} < \infty$ for which, for Lebesgue-almost every $\mathrm{m}_0 \in \mathbb{R}^{12}$, one of the two perfect matchings $M \in \{S^{(0)}, S^{(1)}\}$ satisfies

$$\lim_{\epsilon_\mathrm{L} \to 0} \sup_{t_\mathrm{L} \le t \le 2t_\mathrm{L}} q_t^{\epsilon_\mathrm{L}}(S \ne M) = 0, \qquad t_\mathrm{L} = \mathrm{C}_\mathrm{T}\epsilon_\mathrm{L}^{-1}\log(1/\epsilon_\mathrm{L}).$$

Second, for every $\mathrm{T} > 0$ and every $\mathsf{m}_0$, $q_t^0(S = S^{(0)}) - q_t^0(S = S^{(1)}) \to 0$ as $t \to \infty$. [ESTABLISHED]

The proof is in appendix C.9.

**Corollary 14.6 (Joint selection along a temperature schedule).** In the model of theorem 14.5, fix any sequence $\mathrm{T}_n \downarrow 0$ within its temperature interval and any $\eta_n \in (0, 1)$ tending to zero. There are thresholds $d_n \downarrow 0$, independent of the initial mean, such that for Lebesgue-almost every fixed $\mathsf{m}_0$, all sufficiently large $n$, and every $0 < \epsilon_{\mathrm{L}} \le d_n$, a matching $M_n(\mathsf{m}_0)$ satisfies

$$\sup_{t_{\mathrm{L}} \le t \le 2t_{\mathrm{L}}} q_t^{\epsilon_{\mathrm{L}}, \mathrm{T}_n}(S \neq M_n) \le \eta_n, \qquad t_{\mathrm{L}} = 2\epsilon_{\mathrm{L}}^{-1} \log(1/\epsilon_{\mathrm{L}}).$$

The same thresholds can make the learned law's relative entropy to its instantaneous invariant law, and the frozen law's relative entropy to the uniform-row invariant law, at most $\eta_n$ throughout this window. Consequently the frozen matching probabilities tend to $1/2$ each. For any prescribed $a > 7/2$, the thresholds can also satisfy $d_n \le \mathrm{T}_n^a$. They are existence thresholds, rather than a proved power-law selection condition; the selected matching may depend on $n$. [ESTABLISHED]

The proof and an explicit temperature-dependent tracking bound are in appendix C.10. Setting $\mathrm{T}(\epsilon_{\mathrm{L}}) = \mathrm{T}_n$ for $d_{n+1} < \epsilon_{\mathrm{L}} \le d_n$ gives a joint limit with a schedule common to almost every initial mean. Proposition C.16 identifies what persists on this late window. At residual resolution $\mathrm{a} = \mathrm{T}^{1/4}$, the selected pairing remains the coherent partition for a following window of length $\mathrm{T}^{-1/4}$, with probability tending to one when the schedule tolerances satisfy $\eta_n \le \mathrm{T}_n$. The frozen arm approaches the equal mixture of the two persistent pairings. The statement controls the source redraws of the kinetics; its comparison with a held-sheet Gaussian process is an estimate on path laws. This is restored equilibrium coherence.

The first contrast block is a different statistic. Proposition C.17 keeps the contrast event and the early horizon of proposition 14.4, writing $\xi = \mathrm{T}^{1/8}$ and separating the learning mobility $\epsilon_{\mathrm{L}}$ from $\xi$. If $\epsilon_{\mathrm{L}} = o(\mathrm{T}^{1/4})$, the learned and frozen laws of the first block, each conditioned on contrast formation, have total variation tending to zero for every fixed Gaussian initial mean. Thus the slow schedules can select different late pairing laws while leaving the early first-block comparison unchanged. The original regime $\epsilon_{\mathrm{L}} = \xi$ is outside this estimate.

The retained channel in part (b) of proposition 13.22 also applies after this late coherence. Complete a duration-one coherence prefix at $\theta = t_{\mathrm{L}} + 1$, and initialize the parent from the full recognition law conditioned on that prefix, pushed through the channel. Propositions C.18 and C.20 give, along corollary 14.6 with tolerances $\eta, \delta \le \mathrm{T}$ and mobility cap $\epsilon_{\mathrm{L}} \le \mathrm{T}^2$, separately vanishing discarded conditional relative entropy and retained-path total variation over a following interval $\mathrm{T}^{-1/4}$. The parent inherits the rows at activation and then learns from its own actual law. It retains the assignment, the loop, and the common fast normal coordinate; no independent Gaussian preparation is substituted for its activation law. These are averaged guarantees for this declared past event. They do not assert closure conditional on every realized history or activate a parent at the first contrast.

The limit $\epsilon_{\mathrm{L}} \to 0$ at fixed $\mathrm{T}$ places the mobility of the rows below any positive mixing rate of the kinetics at fixed rows, and $t_{\mathrm{L}}$ is, up to its constant, the time in which an imprint of order $\epsilon_{\mathrm{L}}$ on the rows that grows at a rate of order $\epsilon_{\mathrm{L}}$ reaches order one. The asymmetry that learning would amplify is supplied by the initial law, not by the generative model or the kinetics: by part (c) of proposition 14.4, learning from a Gaussian initial law that is invariant under the cycle automorphisms leaves the rows uniform for all time. The theorem amplifies the asymmetry of $\mathcal{N}(\mathrm{m}_0, I_{12})$ and excludes a null set of means. The frozen arm is the null against which selection is measured. This is a result about the one-time assignment law at fixed temperature. The conditional block events, a power-law vanishing-temperature regime, and conditional and dynamical parent errors of open problem 13.20 remain separate obligations. The participatory reading requires collectives constituted by the coupled inference, with a restoring interaction and a dependence on the preceding state. Assignment selection alone does not establish that recursive claim. The noumenal reading buys structure behind the agents' representations with posits that definition 14.1 removes unless an observable such as a loop holonomy carries them.

**Open Problem 14.7 (Physical-law identification).** Identifying an object produced by the coarse-graining of chapter 11, or by its limits, with a physical law requires a named target system, an observable and an estimator, a validation and uncertainty protocol set before the data, a baseline margin, probability-level closure, and held-out evidence that is robust across admissible blocking schemes. None is supplied here, and the mathematics neither proves a physical law nor supports ontological closure. For a named target system whose data the model does not generate, observed through a declared record kernel that, under the model's own law, separates two objects of the coarse-graining that differ in a gauge-invariant datum, with an observable, an estimator, and a validation and uncertainty protocol set before the data, the problem is settled by held-out evidence, robust across admissible blocking schemes, that an object of the coarse-graining meets a baseline margin with probability-level closure, or by held-out evidence under the same protocol that the baseline margin is not met.

Graph flatness (proposition 2.13) and the fixed-sector condition of proposition 10.17 check declared data but detect no bundle topology or base curvature, so an interpretation must keep graph-link holonomy, base-connection holonomy, bundle topology, and gauge redundancy apart.

**Open Problem 14.8 (Single-point testability).** State an observation at $c_*$ that could count against a claim of the program. What holds is the following: theorem 10.28 constructs a population record and a frame-invariant statistic whose law separates two independent-link configurations that differ only in the conjugacy class of a loop holonomy, and theorem 14.5 establishes a learned pairing selection against a symmetric frozen-prior arm in one declared family. What is missing is a target system: sampling units, a noise model, an estimator of gauge-typed parameters with its uncertainty, a held-out statistic, an acceptance tolerance set in advance, and baselines without link holonomy or without formation. A cross-scale test additionally needs two or more admissible partitions, coarse channels, gauge alignment, and held-out coarse statistics predicted from the fine fit without coarse refitting, and its failure would falsify the added cross-scale hypothesis, not the finite-dimensional

theorems. For a named target system at $c_*$ whose data the model does not generate, observed through a declared record kernel, with sampling units and a noise model, that reads the links of a closed walk or the coordinates of a candidate block, the problem is settled by a test fixed before the data, with an estimator of gauge-typed parameters, a held-out statistic, an acceptance tolerance, and baselines without link holonomy or without formation, whose failure would falsify the presence of nontrivial holonomy on that walk or of formation of that block by the inference in the target, together with a proof that under the declared kernel the law of the tested statistic depends on that hypothesis; or, negatively, by a proof that under the declared kernel every statistic of the records has the same law for two link configurations that differ only in the conjugacy class of the walk's holonomy, or the same law whether or not the block is coherent at the declared resolution.

**Physicist's summary.** Beyond scaffolding, frame-change invariants select no realist position and environment messages create no agents; coordinate ascent, learning, and delayed feedback do not distinguish the program from hierarchical modeling, but formation by the inference would.

# Chapter 15

# Related Work

Where does this construction sit among existing work, and what does it owe to it? The central comparison, with Bayesian renormalization in section 15.4, rests on proposition 5.11 and corollary 9.7, and the open bridge problem of Supplement G builds on it.

No construction below is identified with this one or adopted as authority beyond its cited statement, and qualifications of other works are this manuscript's reading. [NOT-CLAIMED]

## 15.1 Gauge theory, free energy, and multi-agent active inference

Sengupta et al. [71] propose, in an essay, a gauge theory of neuronal dynamics in which variational free energy plays the role of a Lagrangian, gauge fields arise from the information geometry of the brain's probability distributions, and attention, the cognitive counterpart of precision weighting, is a force produced by the curvature of that geometry. Here the group acts on the sample fibers of one principal bundle (chapter 2), links are independent group elements on edge copies at one base point (definition 2.11), and attention is a latent source label with a derived posterior (proposition 6.2) whose precision is that of a record (proposition 6.19).

In federated inference [32], agents share a generative model of a shared world and communicate beliefs through a likelihood mapping to overt expressions, and language emerges when agents share an econiche. Heins et al. [41] obtain cohesion, milling, and directed motion from agents minimizing surprise under generative models of their neighbors, recovering social forces as suppressed prediction errors without a group-level model. Thestrup Waade et al. [75] treat a collective whose group-level Markov blanket consists of its sensory agents' observations and its active agents' actions as a group-level agent, fit a model of the same structure to simulated collective behavior, and find the individual and group models nontrivially related. Here agents may carry inequivalent representations and independent links, their encoder, transmission, and decoder kernels enter one joint with one evidence bound (definition 6.10 and chapter 5), parent laws follow from the fine joint by exact coarse channels (theorem 9.6 and corollary 9.10), formation is proved for a stated family (chapter 13), and a low-rank interface does not by itself give a Markov blanket (proposition 6.26).

## 15.2 Connection Laplacians, sheaves, and equivariant networks

Singer and Wu [73, Section 3 and Theorem 5.1] build a graph connection Laplacian from edge weights and orthogonal transformations aligning the tangent spaces of neighboring sample points and prove its convergence to the connection Laplacian of the underlying manifold; Hansen and Ghrist [37, Section 3.6] identify it as the sheaf Laplacian of an orthogonal vector bundle over a graph, in a spectral theory of sheaves whose discrete Hodge theorem identifies harmonic cochains with sheaf cohomology [37, Theorem 3.1]. Our quadratic energies with represented links are of this type, with kernel the holonomy-fixed sector (proposition 10.17), but with links in $G$ acting through representations, semidefinite weights, and normalized laws; the quadratic disagreement of definition 6.15 is the edge term of a weighted sheaf energy with rectangular restriction maps, as in the discourse sheaves of selective expression [38]. Hansen and Ghrist [37, Section 4.2] show that Kron reduction is in general impossible for sheaves, since an eliminated vertex can impose constraints between boundary vertices that no pairwise restriction maps express, and possible for vertex stalks of dimension at most one, where the Laplacians have factor width at most two [37, Theorem 4.3]; this is prior art for the Kron obstructions of matrix-weighted Gaussian energies that Supplement G classifies, whereas Galerkin closure (theorem 10.25) restricts a form to a subspace.

Group-equivariant convolutional networks share weights across the orbits of a global symmetry group [19], gauge-equivariant ones make a network on a manifold independent of the choice of local frames [20], and a gauge-equivariant transformer applies multi-head self-attention to feature fields on meshes with a parallel transport that makes the layer agnostic to the orientation of local coordinate systems [40]. These constrain learned maps over a base manifold with a connection. Here links are independent dynamic data at one base point, a nonidentity per-edge twist has no counterpart in an architecture with one shared key map (table 6.1), and corollary 6.11 and proposition 6.13 describe how declared data transform, not constraints on a learned layer.

## 15.3 Averaging and path-space coarse-graining

Hartmann et al. [39, Section 2.1.4, Theorem 2.9] bound path-space relative entropy for an effective diffusion under regularity, conditional functional inequalities, and Novikov's condition, with the resolved diffusion coefficient independent of unresolved coordinates. Their fixed-time estimates in Theorems 2.4 and 2.7 have a different target. Neither estimate supplies the growing-history, switching-source comparison requested by open problem 11.38. The one-time selection result of theorem 14.5 uses the separate entropy, transient imprint, and saddle-escape estimates in appendix C.9. A change in diffusion covariance can instead make continuous-path laws singular. Backhoff-Veraguas and Unterberger [5, Remark 1 and Theorem 1] study a partition-refinement relative entropy for martingale diffusions; that quantity is distinct from ordinary path-space relative entropy. The filter and memory constructions of section 8.3 retain the dynamical information that a stationary marginal alone omits.

## 15.4 Information-geometric and Bayesian renormalization

Bény and Osborne [8] formulate the renormalization group as a family of quantum channels connecting the microscopic world to the observable world at a given scale, endow the space of Hamiltonians with a metric induced by an information geometry on states, measure the information lost along a flow by the distinguishability of thermal states, and introduce functions of two-point correlations that do not increase along the flow; theorem 7.8 is only the classical Fisher-information statement for one parameter-independent kernel at one parameter value.

Bayesian renormalization, as formulated by Berman, Klinger, and Stapleton, builds on the dynamical Bayesian inference of Berman et al. [10] and uses Fisher distinguishability on model-parameter space to define an information scale and to separate stiff from sloppy directions. In the late-data regime the posterior is a Gaussian whose covariance is the inverse Fisher information times the scale parameter [11, Eq. (38)]. For a differentiable forward map $\mathfrak{O}$ from parameters to data they write the pushforward of that posterior as a Gaussian [11, Eq. (43)] whose covariance is the scale parameter times the pushed inverse Fisher metric

$$D_\theta^{\mathfrak{O}} = \partial\mathfrak{O}(\theta) g^F(\theta)^{-1} \partial\mathfrak{O}(\theta)^\top \tag{15.1}$$

at the posterior center [11, Eq. (44)], the diffusion tensor of a Fokker–Planck equation and a stochastic flow on data space [11, Eqs. (45) and (46)]. In our reading the delta method makes their Gaussian form the leading local approximation of the pushforward, exact when $\mathfrak{O}$ is affine, and its continuation beyond the late-data regime is a renormalization construction, not an exact nonlinear pushforward or a Bayesian asymptotic theorem. It has been inverted into a statistical-inference reading of exact renormalization flows [9] and applied to neural-network field theories [46].

The information scale ranks which directions of a parent model presentation stay distinguishable at a declared precision, but by itself it selects no graph blocks, coarse-grains no nonflat links, and leaves the two-channel interaction law open. Letting a posterior flow does not make a generative kernel read its own posterior, so hypothesis 4.3 does not refute Bayesian renormalization, but inserting a posterior from one scale as the generative top factor at another declares a new model or an explicit transition between models, whose bounds refer to different joints unless a common target and a comparison map are supplied. By proposition 5.11, pointwise identities designate no one member or evidence, and comparing two bounds shows an improvement in one quantity only on an evidence level set; for two independently declared joints with no evidence-preserving channel between them, the difference of optimized bounds mixes evidence and tightness differences and supplies no order, whereas an evidence-preserving recognition-independent Markov kernel applied to the posterior and the recognition law orders the bounds by data processing (corollary 9.7 and lemma 3.12). Relating the widening of the posterior along its scale to the addition of precision under Gaussian aggregation would need a common state space, a map between the scale parameters, and a transported monotone; Supplement G records this bridge as open.

Mehta and Schwab [61] map Kadanoff's variational renormalization group exactly onto deep restricted Boltzmann machines; a deterministic identification of coarse and fine coordinates does not supply its kernel, so the Galerkin coarse form of section 9.4.1 is not thereby a deep network.

## 15.5 Network renormalization, graph coarsening, and multigrid

Network renormalization addresses heterogeneous graphs, which have no lattice, momentum, or geometric length [33]. The Laplacian renormalization group of Villegas et al. [82] uses the diffusion density operator $e^{-tL}/\operatorname{Tr} e^{-tL}$ of a graph Laplacian $L$, its spectral entropy $\mathcal{S}(t)$ normalized by the logarithm of the number of nodes, and the specific heat

$$\chi(t) = -\frac{d\mathcal{S}}{d\log t}, \tag{15.2}$$

whose peaks they use to identify characteristic resolution scales [82, Eqs. (1)–(2)].

**Remark 15.1 (Gauge behavior of heat-kernel diagnostics).** Let $R \succ 0$ and $L$ be block matrices on $\bigoplus_{i\in V} \mathbb{R}^{n_i}$, and let a vertex reframing $P = \operatorname{diag}(P_i)$ with invertible blocks act by congruence, $R \mapsto P^\top RP$ and $L \mapsto P^\top LP$. Then the generator $R^{-1}L$ becomes $P^{-1}R^{-1}LP$, so its spectrum, its heat trace $\operatorname{Tr} e^{-tR^{-1}L}$, and every spectral entropy built from them are invariant, whereas the off-diagonal blocks of $e^{-tR^{-1}L}$ become $P_i^{-1}(e^{-tR^{-1}L})_{ij}P_j$ and are not invariant scalar node affinities. [ESTABLISHED]

*Proof.* $(P^\top RP)^{-1}P^\top LP = P^{-1}R^{-1}LP$, and the exponential of a similar matrix is similar by the same matrix. □

A gauge-invariant block selector built from heat-kernel data therefore needs invariant statistics of the blocks, and a scale diagnostic alone selects no unique partition (open problem 9.19). Garuccio et al. [34] give a renormalization scheme for any hierarchy of heterogeneous coarse-grainings and identify a class of scale-invariant random graphs with a necessary dependence on additive node hidden variables, a characterized family rather than arbitrary hidden-variable networks; for random graphs whose link probabilities are conditioned on other links, Catanzaro et al. [18] derive a closed-form renormalization transformation for maximum coordination number two and show that higher coordination numbers admit none. These candidate partition and closure mechanisms do not by themselves handle directed row-normalized attention, two gauge representations, or nonflat link holonomy, and the aggregation of chapter 10 does not integrate out edge variables. A combined scheme would infer a partition from gauge-invariant multiplex diagnostics of edge-event laws through a structural channel with a named law and push the complete law exactly.

Graph coarsening and pooling in machine learning add hypotheses of their own: restricted spectral approximation controls spectral and cut properties [60], coarsening can carry message-passing guarantees [47], differentiable pooling learns soft cluster assignments end to end [88], and deep graph convolutional networks whose weights satisfy spectral conditions lose all information except connected components and node degrees exponentially in depth [66]. The Galerkin identity pre-

serves the restricted quadratic energy, while spectral approximation, message-passing commutation, and a graph-neural-network realization need the further controls of section 9.4.1 (proposition 9.14). That coarse form is the one of algebraic multigrid: aggregation-based multigrid forms $P^\top AP$ with a piecewise-constant prolongation defined by disjoint aggregates [65, Section 2, Eqs. (2.1)–(2.2)], smoothed aggregation smooths that tentative prolongation [79], and lean algebraic multigrid for graph Laplacians combines piecewise-constant interpolation, node aggregation by an affinity measure, and an energy correction of the coarse systems [58]. What is added here is the placement of that operator inside normalized laws with gauge covariance and exact evidence accounting.

**Physicist's summary.** Here the group acts on sample fibers with independent links at one point, and parent laws are computed from the fine joint rather than shared, omitted, or fitted as in multi-agent active inference. Sheaf Laplacians supply the linear algebra and the Kron obstruction, and renormalization flows meet the exact channels only through a common comparison target.

# Appendix A

# Measure-Theoretic Lemmas and Integrability Witnesses

This appendix proves the measure-theoretic statements of Chapters 3 and 5 and separates the hypotheses of hypothesis 5.3 from factorwise log-integrability (proposition 5.15) by witnesses.

## A.1 Kernel measurability and proofs for the probability chapter

**Theorem A.1 (Jointly measurable Radon–Nikodym version for a dominated kernel).** Let $(\mathsf{E}, \mathscr{E})$ be a measurable space, and let $K$ and $L$ be probability kernels from $\mathsf{E}$ to $(\mathsf{F}, \mathscr{F})$ with $K(x,\cdot) \ll L(x,\cdot)$ for every $x$. Then there is a finite $\mathscr{E} \otimes \mathscr{F}$-measurable $r : \mathsf{E} \times \mathsf{F} \to [0,\infty)$ with $K(x,B) = \int_B r(x,y) L(x,dy)$ for every $x \in \mathsf{E}$ and $B \in \mathscr{F}$. Selecting a derivative separately for each $x$ does not give joint measurability; the theorem supplies one simultaneous version. [ESTABLISHED]

*Proof.* Standard Borelness gives finite refining measurable partitions $(\mathcal{P}_n)_{n\geq 1}$ whose generated sigma-algebras increase to $\mathscr{F}$. With $0/0 := 0$ put

$$r_n(x,y) = \sum_{A \in \mathcal{P}_n} \frac{K(x,A)}{L(x,A)} \mathbf{1}_A(y), \tag{A.1}$$

and let $r = \limsup_n r_n$ where this is finite and $r = 0$ elsewhere; like every $r_n$, $r$ is jointly measurable. For fixed $x$, $r_n(x,\cdot)$ is the conditional expectation of $h_x = dK(x,\cdot)/dL(x,\cdot)$ given $\sigma(\mathcal{P}_n)$ under $L(x,\cdot)$, so martingale convergence gives $r_n(x,\cdot) \to h_x < \infty$ $L(x,\cdot)$-almost everywhere, and $r(x,\cdot)$ is a derivative for every $x$. □

**Corollary A.2 (Joint densities against one reference).** Let $\nu$ be a sigma-finite measure on $(\mathsf{F}, \mathscr{F})$ and $K$ a probability kernel from $(\mathsf{E}, \mathscr{E})$ to $\mathsf{F}$ with $K(x,\cdot) \ll \nu$ for every $x$. Then $K$ has a finite jointly measurable density against $\nu$, and every finite or countable collection of kernels dominated by the same $\nu$ has jointly measurable densities against that common reference; in particular this applies to the generative and recognition kernels (3.6)–(3.7) when they are dominated by the product references. [ESTABLISHED]

*Proof.* Sigma-finiteness supplies measurable $w$ with $0 < w < \infty$ everywhere and $\int w(y)\nu(dy) = 1$, so the constant kernel $L(x, dy) = w(y)\nu(dy)$ is equivalent to $\nu$. Theorem A.1 gives a finite jointly measurable $r = dK/dL$, and $k = rw$ is the density. Apply this to each member of the collection. □

**Proposition A.3 (Kernel integration preserves measurability).** Let $\kappa$ be an s-finite kernel from $(\mathsf{E}, \mathscr{E})$ to $(\mathsf{F}, \mathscr{F})$, that is, a countable sum of finite kernels; probability kernels and fixed sigma-finite measures are instances. For every $\mathscr{E} \otimes \mathscr{F}$-measurable $f : \mathsf{E} \times \mathsf{F} \to [0, \infty]$,

$$x \longmapsto \int_{\mathsf{F}} f(x, y)\kappa(x, dy) \tag{A.2}$$

is $\mathscr{E}$-measurable. For an extended signed integrand the same holds on the measurable set where the positive-part and negative-part integrals are not both infinite, and no expression $+\infty - \infty$ is formed. Consequently the evidence (3.12) and the $\mathbb{Q}_{V,o,X}$-expectations of Chapter 5 are measurable in $(o, X)$ whenever their integrands are jointly measurable. [ESTABLISHED]

*Proof.* For finite $\kappa$, the class of $C \in \mathscr{E} \otimes \mathscr{F}$ for which $x \mapsto \kappa(x, C_x)$ is measurable, with $C_x = \{y : (x, y) \in C\}$, contains the rectangles, since $\kappa(x, (A \times B)_x) = \mathbf{1}_A(x)\kappa(x, B)$, and is closed under proper differences and increasing unions because each $\kappa(x, \cdot)$ is finite. Rectangles form a $\pi$-system generating $\mathscr{E} \otimes \mathscr{F}$, so the $\pi$–$\lambda$ theorem gives every $C$. Linearity and monotone convergence pass to nonnegative measurable $f$, and countable sums to s-finite $\kappa$; a sigma-finite $\nu$ with a measurable partition into sets $F_m$ of finite measure is the constant s-finite kernel $\sum_m \nu(\cdot \cap F_m)$. A signed integrand is split into $f^+$ and $f^-$, whose difference is measurable where they are not both infinite. □

**Corollary A.4 (Measurability of kernel relative entropy).** Let $\kappa$ and $\eta$ be probability kernels from $(\mathsf{E}, \mathscr{E})$ to $(\mathsf{F}, \mathscr{F})$. Then $x \mapsto D_{\mathrm{KL}}(\kappa(x, \cdot)\|\eta(x, \cdot))$ is $\mathscr{E}$-measurable as a $[0, +\infty]$-valued map, and for finite refining partitions $(\mathcal{P}_n)$ generating $\mathscr{F}$,

$$D_{\mathrm{KL}}(\kappa(x, \cdot)\|\eta(x, \cdot)) = \sup_{n \geq 1} \sum_{A \in \mathcal{P}_n} \kappa(x, A) \log \frac{\kappa(x, A)}{\eta(x, A)}, \tag{A.3}$$

with $0 \log(0/b) = 0$ and $a \log(a/0) = +\infty$ for $a > 0$. [ESTABLISHED]

*Proof.* Put $L = (\kappa + \eta)/2$. Theorem A.1 applied to $(\kappa, L)$ and $(\eta, L)$ gives finite jointly measurable densities $a$ and $b$. Fix $x$, write $L_x = L(x, \cdot)$, and let $a_n, b_n$ be the conditional expectations of $a(x, \cdot), b(x, \cdot)$ under $L_x$ given $\sigma(\mathcal{P}_n)$, whose cell values are the cellwise ratios in (A.3). The perspective $\phi_{\mathrm{KL}}(u, v) = u \log(u/v) - u + v$, with $\phi_{\mathrm{KL}}(0, v) = v$, $\phi_{\mathrm{KL}}(u, 0) = +\infty$ for $u > 0$, and $\phi_{\mathrm{KL}}(0, 0) = 0$, is nonnegative, convex, and lower semicontinuous. Since both kernels have unit mass, the $n$-th partition sum is $D_n(x) = \int \phi_{\mathrm{KL}}(a_n, b_n)dL_x$ and the divergence is $\int \phi_{\mathrm{KL}}(a, b)dL_x$. Conditional Jensen makes $D_n(x)$ nondecreasing in $n$ and at most the divergence, and martingale convergence $(a_n, b_n) \to (a, b)$ $L_x$-almost everywhere, with Fatou's lemma and lower semicontinuity, gives $\int \phi_{\mathrm{KL}}(a, b)dL_x \leq \liminf_n D_n(x)$. Each $D_n$ is measurable in $x$, so the countable supremum is measurable. □

*Proof of Theorem 3.5.* If $A_*$ is countable, $\nu_{\mathrm{mix}}$ is sigma-finite, and a $\nu_{\mathrm{mix}}$-null set is Lebesgue null and contains no point of $A_*$, so every term of every $\mu_t$ vanishes on it. Conversely, let a sigma-finite $\xi$ dominate every $\mu_t$. Each $a \in A_*$ has $\mu_t(\{a\}) > 0$ for some $t$, whence $\xi(\{a\}) > 0$, and a sigma-finite measure has at most countably many points of positive mass, as in the proof of Remark 3.6. The moving-atom family has $A_* = \mathbb{R}$. □

**Remark A.5 (A regular-conditional version is not determined at a point).** On $\mathbb{R}^2$ with coordinates $(o, y)$, let $\phi$ be the standard normal density and $\psi \neq \phi$ another probability density. Then $p(o, y) = \phi(o)\phi(y)$ and $\widetilde{p}$, equal to $p$ for $o \neq 0$ and to $\phi(0)\psi(y)$ for $o = 0$, are versions of the same joint density; both evidence representatives equal $\phi(0)$ at $o = 0$, but their density-ratio conditionals there are $\phi(y)dy$ and $\psi(y)dy$. The joint measure therefore determines neither pointwise posterior at $o = 0$. [ESTABLISHED]

*Proof.* The two functions differ only on the Lebesgue-null line $\{0\} \times \mathbb{R}$, and integrating out $y$ at $o = 0$ gives $\phi(0)$ in both cases. □

*Proof of Lemma 3.12.* (a) Measurability of the integrand is Corollary A.4 with $\mathsf{E} = \mathsf{Y}_1$. Suppose first $Q \ll P$. Then $Q_1 \ll P_1$, since $P_1(A) = 0$ gives $P(A \times \mathsf{Y}_2) = 0$. Let $r = dQ/dP$ and $r_1 = dQ_1/dP_1$. For measurable $A \subseteq \mathsf{Y}_1$ and $B \subseteq \mathsf{Y}_2$,

$$\int_A Q_{2|1}(y_1, B) r_1(y_1) P_1(dy_1) = Q(A \times B) = \int_A \int_B r(y_1, y_2) P_{2|1}(y_1, dy_2) P_1(dy_1),$$

so for each $B$ the integrands agree $P_1$-almost everywhere; over a countable generating algebra of $\mathsf{Y}_2$ the exceptional sets form one $P_1$-null set $N$, off which both sides are finite measures in $B$ that agree on the algebra and hence everywhere. Off the $Q_1$-null set $N \cup \{r_1 = 0\}$, $Q_{2|1}(y_1, \cdot)$ therefore has the density $r(y_1, \cdot)/r_1(y_1)$ against $P_{2|1}(y_1, \cdot)$, and $\log r$ splits $Q$-almost surely into $\log r_1$ and the logarithm of that density. Both summands have negative parts of $Q$-integral at most $1/e$, as in Definition 3.10, so expectations add in $(-\infty, +\infty]$; the first gives $D_{\mathrm{KL}}(Q_1 \| P_1)$, and Tonelli's theorem for the positive and negative parts along $Q = Q_1 \otimes Q_{2|1}$ gives the integral of the conditional divergences. Suppose next $Q \not\ll P$. If $Q_1 \not\ll P_1$, the first term on the right is $+\infty$. If $Q_1 \ll P_1$ and the conditional term were finite, then $Q_{2|1}(y_1, \cdot) \ll P_{2|1}(y_1, \cdot)$ for $Q_1$-almost every $y_1$, so every $P$-null $C$, whose sections are $P_{2|1}(y_1, \cdot)$-null for $P_1$-, hence $Q_1$-almost every $y_1$, would be $Q$-null. Conditional versions are unique $Q_1$-almost everywhere when $Q_1 \ll P_1$, and otherwise the right side is $+\infty$ for every version.

(b) Quantize every block by finite refining partitions generating its sigma-algebra. On the finite product alphabet a cell of positive $Q$-mass has marginal cells of positive $Q_b$-mass, so the first term on the right is finite; if every marginal term is finite, every such cell has positive reference mass and the logarithm splits cellwise, and otherwise a marginal cell of positive $Q_b$-mass and zero $\rho_b$-mass makes both sides infinite. By (A.3) for constant kernels, each of the three divergences is the nondecreasing limit of its quantized values, which gives (3.21).

(c) By (a), every $Q$ with $Q_1 = \mu$ has $D_{\mathrm{KL}}(Q\|P) \geq D_{\mathrm{KL}}(\mu\|P_1)$. When $D_{\mathrm{KL}}(\mu\|P_1) < \infty$, equality holds exactly when $Q_{2|1} = P_{2|1}$ $\mu$-almost everywhere, which determines $Q$. When the marginal divergence is infinite, every member of the class has divergence $+\infty$.

(e) If $Q \ll P$ with $r = dQ/dP$, then $A_{\#}Q(B) = \int_{A^{-1}B} r(y)P(dy) = \int_B r(A^{-1}y')(A_{\#}P)(dy')$ for measurable $B$, which is (3.23), and integrating $\log(r \circ A^{-1})$ against $A_{\#}Q$ gives $\int \log r(y)Q(dy)$. If $Q(S) > 0 = P(S)$, then $A(S)$ is measurable because $A^{-1}$ is, and $A_{\#}Q(A(S)) > 0 = A_{\#}P(A(S))$, so both divergences are $+\infty$.

(d) Form the laws $Q(dy)K(y, dz)$ and $P(dy)K(y, dz)$ on $\mathsf{Y} \times \mathsf{Z}$. By (a) with $y$ as the first coordinate their divergence is $D_{\mathrm{KL}}(Q\|P) + \int D_{\mathrm{KL}}(K(y,\cdot)\|K(y,\cdot))Q(dy) = D_{\mathrm{KL}}(Q\|P)$. By (e) for the coordinate swap and (a) with $z$ as the first coordinate, the same divergence is $D_{\mathrm{KL}}(QK\|PK)$ plus a term in $[0, +\infty]$. □

## A.2 The relative-log representation and integrability witnesses

**Proposition A.6 (Relative-log representation).** Under (H1)–(H2) of Hypothesis 5.3, let $M_{\theta,V,o,X}$ be the slice measure (5.14). If $Q \ll M_{\theta,V,o,X}$ and $r = dQ/dM_{\theta,V,o,X}$, then

$$\begin{aligned} \mathcal{L}_V^{\mathrm{ext}}(Q; X, o) &= -\int_{\mathsf{Y}} \log r(y)Q(dy), \\ \int_{\{r<1\}} -\log r(y)Q(dy) &= \int_{\{r<1\}} -r(y)\log r(y)M_{\theta,V,o,X}(dy) \leq \frac{p_{\theta,V}(o \mid X)}{e}, \end{aligned} \tag{A.4}$$

so the integral is well defined in $(-\infty, +\infty]$ and its negative lies in $[-\infty, +\infty)$. If $Q \not\ll M_{\theta,V,o,X}$, then $\mathcal{L}_V^{\mathrm{ext}}(Q; X, o) = -\infty$. [ESTABLISHED]

*Proof.* Since $\mathbf{\Pi}_{\theta,V,o,X} = M_{\theta,V,o,X}/p_{\theta,V}(o \mid X)$, the Radon–Nikodym chain rule gives $dQ/d\mathbf{\Pi}_{\theta,V,o,X} = p_{\theta,V}(o \mid X)r$. The negative part of $\log r$ is $Q$-integrable by the displayed bound, because $-u\log u \leq 1/e$ on $[0, 1]$ and $M_{\theta,V,o,X}(\mathsf{Y}) = p_{\theta,V}(o \mid X)$. Hence $D_{\mathrm{KL}}(Q\|\mathbf{\Pi}_{\theta,V,o,X}) = \log p_{\theta,V}(o \mid X) + \int \log r(y)Q(dy)$ is a legitimate extended identity, and substitution in (5.15) gives the first branch. The slice measure and the posterior have the same null sets, which gives the second. □

**Proposition A.7 (Integrability witnesses).** The integrability hypotheses of Chapter 5 are separated by the following witnesses.

(a) Clauses (H1)–(H3) of Hypothesis 5.3 do not imply (H4). On $\mathbb{R}$ with a standard Gaussian posterior and a standard Cauchy recognition law, (H1)–(H3) hold, $\mathbb{E}_Q \log q = -\log 4\pi$ is finite, $\mathbb{E}_Q|\log p_{\theta,V}(o, Y \mid X)| = +\infty$, the divergence is $+\infty$, and the extended bound is $-\infty$.

(b) Finite relative entropy does not imply (H4). On $\{2, 3, \ldots\}$ with a trivial record, let the posterior and the recognition law both equal $P(\{n\}) = \mathrm{c}/(n(\log n)^2)$, with $\mathrm{c}$ normalizing. The divergence is $0$, while $\sum_n P(\{n\})|\log P(\{n\})| = +\infty$, so both expectations in the classical split equal $-\infty$ and their difference is undefined.

(c) Factorwise log-integrability (proposition 5.15) does not imply (H4). With one latent $k \in \mathbb{R}$, generative factor $\mathcal{N}(0,1)$, and a record whose log density is bounded for $k \in [0,1]$, let the recognition law spread mass $w_n = 2^{-n}$ uniformly over disjoint intervals in $[0,1]$ of widths $w_n e^{-1/w_n}$, whose total is below $e^{-2}$. Every factor logarithm is bounded on the support, while $\log q = 1/w_n$ on the $n$-th interval and $\mathbb{E}_Q \log q = \sum_n w_n/w_n = +\infty$.

(d) Factorwise log-integrability is strictly stronger than the generative half of (H4). Take a root model coordinate $m \in \{2,3,\ldots\}$ with $p(m) = \mathrm{c}/(m(\log m)^2)$, a bridge $k \mid m \sim \mathcal{N}(0,\sigma_m^2)$ with $\sigma_m^2 = p(m)^2/(2\pi)$, a trivial record, and the recognition law equal to this law. Then $\log p(m) + \log p(k \mid m) = -k^2/(2\sigma_m^2)$ has expectation $-\frac{1}{2}$ and is also the recognition log density, so (H4) holds, while $\mathbb{E}|\log p(m)| = +\infty$.

[ESTABLISHED]

*Proof.* (a) Both laws have positive Lebesgue densities, so (H1)–(H3) hold. The log joint differs from $-y^2/2$ by a constant and the Cauchy law has no second moment; the Cauchy differential entropy is $\log 4\pi$; hence the divergence $\mathbb{E}_Q \log q - \mathbb{E}_Q \log p$ is $+\infty$. (b) For large $n$, $P(\{n\}) \le 1/n$, so $|\log P(\{n\})| \ge \log n$ and the sum dominates $\mathrm{c}\sum_n 1/(n \log n) = +\infty$. (c) The widths sum to $\sum_n 2^{-n} e^{-2^n} < e^{-2}$, so the intervals fit in $[0,1]$; the density on the $n$-th interval is $w_n/(w_n e^{-1/w_n}) = e^{1/w_n}$. (d) The bridge density at $k$ is $(2\pi\sigma_m^2)^{-1/2} e^{-k^2/(2\sigma_m^2)} = p(m)^{-1} e^{-k^2/(2\sigma_m^2)}$, which gives the identity and $\mathbb{E}|\log p(m) + \log p(k \mid m)| = \frac{1}{2}$; and $\mathbb{E}|\log p(m)| = +\infty$ as in (b). □

# Appendix B

# Deferred Proofs

This appendix collects the longer proofs of chapters 6, 7, 9 and 11, in the order of the statements they prove.

## B.1 Proofs for Attention and Selective Communication

*Proof of Proposition 6.2.* (a) The product label prior and exclusivity make the posterior of $J$ factor by rows, and Bayes' formula gives (6.4); a weight $\pi_{ij}e^{-D_{ij}/\tau_i}$ is positive exactly when $\pi_{ij} > 0$, since the energies are finite, and with equal prior weights the largest weight has the smallest energy, $D_{ii} = 0$. (b) With $Z_i^{\mathrm{att}}(y) = \sum_k \pi_{ik}e^{-D_{ik}(y)/\tau_i}$, $D_{\mathrm{KL}}(\beta_i^Q \| \beta_i^P(y)) = D_{\mathrm{KL}}(\beta_i^Q \| \pi_i) + \tau_i^{-1}\sum_j \beta_{ij}^Q D_{ij}(y) + \log Z_i^{\mathrm{att}}(y)$. Subtracting $\log Z_i^{\mathrm{att}}(y)$ and the source-independent $\log \mathrm{c}_i(o_i, y)$ and taking the recognition expectation proves (6.6) and (6.7); Lagrange multiplication under $\sum_j \beta_{ij}^Q = 1$ gives (6.8), and strict convexity of $D_{\mathrm{KL}}(\cdot\|\pi_i)$ gives uniqueness, which the factor $\tau_i > 0$ does not change. (c) The label part of the VFE is $\mathbb{E}_{\mathbb{Q}}[\log \mathbb{Q}_{J|Y} - \sum_i \log \pi_{iJ_i} - \sum_i \log L_i(o_i \mid Y, J_i)]$. The product-reference case of lemma 3.12 splits $D_{\mathrm{KL}}(\mathbb{Q}_{J|Y=y} \| \prod_i \pi_i)$ into the total correlation and the row relative entropies, which with (6.2) gives (6.9). At fixed $y$ the expansion of (b), applied to the product reference, equals the label terms other than $-\sum_i \log \mathrm{c}_i$ to $D_{\mathrm{KL}}(\mathbb{Q}_{J|Y=y} \| \prod_i \beta_i^P(y)) - \sum_i \log Z_i^{\mathrm{att}}(y)$, which the product of the posterior rows minimizes. □

*Proof of Proposition 6.3.* Substituting (6.8) into (6.7) gives $R_i^Q$, the minimum by the expansion in the proof of proposition 6.2. At fixed $y$ the same expansion gives $D_{\mathrm{KL}}(\mathbb{Q}_{J_i|Y=y} \| \pi_i) + \tau_i^{-1}\mathbb{E}_{\mathbb{Q}_{J_i|Y=y}} D_{iJ_i}(y) = D_{\mathrm{KL}}(\mathbb{Q}_{J_i|Y=y} \| \beta_i^P(y)) + R_i^P(y)$, minimized by $\beta_i^P(y)$, and the conditional total correlation in (6.9) is nonnegative and vanishes on the product of these rows. Normalization of $\pi_i$ and concavity of the logarithm give $\min_j D_{ij}/\tau_i \le R_i^P \le \sum_j \pi_{ij} D_{ij}/\tau_i$, so $R_i^P(Y)$ is integrable. The softmax formulas give $\log \beta_{ij}^{Q\star} - \log \beta_{ij}^P(y) = (D_{ij}(y) - d_{ij})/\tau_i + R_i^Q - R_i^P(y)$, and averaging over $\beta_i^{Q\star}$ and then over $\mathbb{Q}_Y$ removes the first term, since $\mathbb{E}_{\mathbb{Q}_Y} D_{ij} = d_{ij}$, which proves (6.11). The relative entropy vanishes at $y$ exactly when $D_{ij}(y) - d_{ij}$ does not depend on $j$, and almost-sure constancy of the differences forces those constants to equal $d_{ij} - d_{ik}$. Under the affine hypothesis $d_{ij} - d_{ik} = (\widetilde{D}_{ij} - \widetilde{D}_{ik})(\bar{\psi}_i)$, and a softmax row depends only on logit differences, which gives (6.12); template differences are affine in $y_i$. For the random source, uncorrelatedness gives $\mathbb{E} r_{ij} = \bar{y}_i - \rho_b(\Theta_{ij}^b)\bar{y}_j$ and $\operatorname{Cov} r_{ij} = \operatorname{Cov} Y_i + \rho_b(\Theta_{ij}^b) \operatorname{Cov}(Y_j) \rho_b(\Theta_{ij}^b)^\top$, so $d_{ij}$ exceeds $D_{ij}$ at the recognition mean by $\frac{1}{2}\operatorname{Tr}(W_i \operatorname{Cov} Y_i) + \frac{1}{2}\operatorname{Tr}(W_i \rho_b(\Theta_{ij}^b) \operatorname{Cov}(Y_j) \rho_b(\Theta_{ij}^b)^\top)$, whose first trace is common to the row and cancels from the softmax. □

*Proof of Corollary 6.7.* On the open simplex the Fisher metric is the Shahshahani metric, whose natural gradient of a function with Euclidean gradient $g$ is $-\beta_j(g_j - \sum_k \beta_k g_k)$. The Euclidean gradient of (6.7) is $\mathsf{c}_{ij} + 1$, and the constant cancels on the simplex, which gives (6.17). Pairing the gradient with the flow gives $-\mathsf{r}_i \sum_j \beta_{ij}\mathsf{c}_{ij}(\mathsf{c}_{ij} - \bar{\mathsf{c}}_i) = -\mathsf{r}_i \operatorname{Var}$, since $\sum_j \dot\beta_{ij} = 0$; with moving $\lambda$ the chain rule adds the explicit $\lambda$-derivative, the second term of (6.19). The quotient rule for $\beta = \eta/\alpha$ gives (6.20), and differentiating $\log \beta_j = \log \pi_j - E_j/\tau - \log Z$ gives $\dot\beta_j/\beta_j = \chi_j - \dot Z/Z$ with $\dot Z/Z = \sum_k \beta_k \chi_k$, which is (6.21). □

*Proof of Proposition 6.8.* Differentiating the finite sum gives $\nabla\beta_{ij} = -\tau_i^{-1}\beta_{ij}(\nabla E_{ij} - \bar g)$ and the gradient; once more, $\nabla^2 R_i = \tau_i^{-1}\sum_j(\beta_{ij}\nabla^2 E_{ij} + \nabla E_{ij}\nabla\beta_{ij}^\top)$, and $\sum_j \beta_{ij}\bar g(\nabla E_{ij} - \bar g)^\top = 0$ turns the second sum into $-\tau_i^{-2}\operatorname{Cov}_{\beta_i}(\nabla E_i)$. The joint Hessian of $\mathscr{J}_i$ has blocks $\tau_i^{-1}\sum_j \beta_{ij}\nabla^2 E_{ij}$, the mixed form $(v, w) \mapsto \tau_i^{-1}\sum_j w_j \nabla E_{ij}\cdot v$, and $w \mapsto \sum_j w_j^2/\beta_{ij}$ on the tangent space; minimizing the joint quadratic form over tangent $w$ at fixed $v$ gives $w_j = -\tau_i^{-1}\beta_{ij}(\nabla E_{ij} - \bar g)\cdot v$, which satisfies $\sum_j w_j = 0$, with minimum $v^\top(\tau_i^{-1}\sum_j \beta_{ij}\nabla^2 E_{ij})v - \tau_i^{-2}\operatorname{Var}_{\beta_i}(\nabla E_{ij}\cdot v)$. The constant-row identification is proposition 6.3; for exact label integration, the product label prior and exclusivity give $\sum_j \pi_{ij}\mathrm{c}_i(o_i, y)e^{-D_{ij}(y)/\tau_i} = \mathrm{c}_i(o_i, y)e^{-R_i(y)}$. In the witness, $O$ given $J = j$ is $\mathcal{N}(-j\delta, \tau(1 + \omega^{-1}))$, whose density at zero is the displayed evidence for both labels. At $o = 0$, $\frac{1}{2}\sum_{j=\pm1} e^{-(x-j\delta)^2/(2\tau)} = e^{-(x^2+\delta^2)/(2\tau)}\cosh(\delta x/\tau)$, which with the baseline gives (6.24), and $\frac{d^2}{dx^2}\log\cosh(\delta x/\tau) = \delta^2/\tau^2$ at $x = 0$. The constant-row value follows from (6.23) with the affine energies, whose weights are equal at $\mathsf{m} = 0$, together with the recognition entropy and the baseline term, whose $\mathsf{m}$-curvature is $\omega/\tau$. □

*Proof of Corollary 6.11.* (a) Integrating $o$, then $\xi^+$, then $\xi^-$ gives total mass one, and measurability follows from measurable kernel integration (chapter 3). (b) Substituting $\xi^- = (\phi_e^-)^{-1}\xi'$ and $\xi^+ = (\phi_e^+)^{-1}\xi''$ in the composition cancels the message recoordinatizations. (c) The first identity is the definition of $F_e^v$, and the represented link law $\rho_x(\Theta_e^x)' = \mathsf{g}_i\rho_x(\Theta_e^x)\mathsf{g}_j^{-1}$ combined with (b) gives the last. □

*Proof of Proposition 6.13.* (a) The two internal basis changes cancel in $\mathsf{L}_e^{\mathrm{dec}\prime}\mathsf{L}_e^{\mathrm{tr}\prime}\mathsf{L}_e^{\mathrm{enc}\prime}$. (b) The inclusions and the rank bound follow from composition. For $k \in \ker \mathsf{L}_e^{\mathrm{enc}}$ and any $\mathsf{u}$, $(\mathsf{G}_j^{-1}(\mathsf{L}_e^{\mathrm{enc}})^\top \mathsf{u})^\top \mathsf{G}_j k = \mathsf{u}^\top \mathsf{L}_e^{\mathrm{enc}} k = 0$, and $\dim \operatorname{im}(\mathsf{L}_e^{\mathrm{enc}})^\top = \dim E_j - \dim\ker \mathsf{L}_e^{\mathrm{enc}}$, which identifies the complement; the metric law is the congruence induced by $\mathsf{g}_j$. Coordinates on $\operatorname{im}\mathsf{L}_e$, the induced quotient as encoder, the identity as transmission, and the inclusion of the image as decoder factor $\mathsf{L}_e$ through its rank, which proves existence, not canonicity of an interface. (c) Rank is subadditive and bounded by the target dimension, and the example is direct. □

*Proof of Proposition 6.14.* (a) The factorization holds by definition of $F_e$, and $\mathsf{L}_e' = \mathsf{g}_i\mathsf{L}_e\mathsf{g}_j^{-1}$ with $\Theta_e' = \mathsf{g}_i\Theta_e\mathsf{g}_j^{-1}$ gives the law of $F_e$; for every choice of $\Theta_e$ some $F_e$ exists, so the link is not determined, and a common $F_i$ exists exactly when all $\mathsf{L}_e\Theta_e^{-1}$ coincide. (b) Consecutive factors $\mathsf{g}_{i_{a+1}}^{-1}\mathsf{g}_{i_{a+1}}$ cancel, which uses no inverse of any $\mathsf{L}_e$. (c) and (d) are direct matrix computations; the holonomy of the loop through $b$ and $\bar a$ is $\Theta_a^{-1}\Theta_b = \Theta_b \neq I_2$. □

*Proof of Corollary 6.17.* (a) Expanding $D_{\mathrm{KL}}(\beta_i \| \beta_i^P(y))$ with (6.4) and $D_{ij} = d_{ij}$ gives the identity, and the relative entropy vanishes only at $\beta_i^P(y)$; the point mass is singular with respect to the posterior slice when the baseline gives $\{y\}$ zero mass. (b) The terms of $D_{\mathrm{KL}}(\mathcal{N}(\mu_i^Q, \Sigma^Q)\|\mathcal{N}(\mu_j^K, \Sigma^K))$

depending on $j$ are those of $\frac{1}{2}(\mu_j^K - \mu_i^Q)^\top(\Sigma^K)^{-1}(\mu_j^K - \mu_i^Q)$, and expanding the quadratic form gives (6.39). □

*Proof of Proposition 6.19.* (a) The product label prior and exclusivity make $J_i$ independent of $k_i$ and of every other record given $y_{-i}$, and the records are conditionally independent given the latent state, so the joint posterior density of $(k_i, j)$ is proportional to $\pi_{ij}\mathcal{N}(k_i;\mu_i,\Sigma_i)\mathcal{N}(0;S_i^Q k_i - S_e^K k_j, \widetilde{W}_i^{-1})$. For fixed $j$, adding the two quadratic forms in $k_i$ gives the precision $(\Sigma_i^+)^{-1}$ and the component mean. Integrating $k_i$ instead gives the density at zero of $o_i$ given $J_i = j$, an affine image of independent Gaussian variables with mean $r_{ij}$ and covariance $\Sigma_i^{\mathrm{pred}}$, whose normalizer does not depend on $j$; Bayes' formula gives $w_{ij}$. The mean follows from the tower property, and multiplying $(\Sigma_i^+)^{-1} = \Sigma_i^{-1} + S_i^{Q\top}\widetilde{W}_i S_i^Q$ by $\Sigma_i^+$ gives the last identity. (b) Here $\Sigma_i^{\mathrm{pred}} = 2$, $r_{i2} = -2$, and $\Sigma_i^+ = 1/2$; substitution gives the displayed values. (c) The identity follows from $\Sigma_i^+ S_i^{Q\top}\widetilde{W}_i\Sigma_i^{\mathrm{pred}} = \Sigma_i^+(S_i^{Q\top} + S_i^{Q\top}\widetilde{W}_i S_i^Q\Sigma_i S_i^{Q\top}) = \Sigma_i^+(\Sigma_i^+)^{-1}\Sigma_i S_i^{Q\top}$. Since $S_i^Q\Sigma_i S_i^{Q\top}$ is invertible, $(\Sigma_i^{\mathrm{pred}})^{-1} \to (S_i^Q\Sigma_i S_i^{Q\top})^{-1}$ as $\widetilde{W}_i^{-1} \to 0$, which gives the limit. Directly $S_i^Q R_e = I$, and any $\delta$ with $S_i^Q\delta = \mathsf{y}$ is $R_e\mathsf{y} + w$ with $S_i^Q w = 0$, for which the cross term $(R_e\mathsf{y})^\top\Sigma_i^{-1}w = \mathsf{y}^\top(S_i^Q\Sigma_i S_i^{Q\top})^{-1}S_i^Q w$ vanishes, so $R_e\mathsf{y}$ uniquely minimizes $\delta^\top\Sigma_i^{-1}\delta$. (d) Repeat (a) with the copy-dependent data. □

*Proof of Proposition 6.20.* At fixed $k_i$ the proof of proposition 6.2 applies verbatim to the rows $(i,1)$ and $(i,2)$. Given the labels, the two records at zero are a Gaussian vector with mean $-(k_j, k_l)$ and covariance $I_2 + \mathbf{1}\mathbf{1}^\top$, whose inverse is $I_2 - \frac{1}{3}\mathbf{1}\mathbf{1}^\top$; completing the square gives (6.45). The masses are $1$ for $(1,1)$ and $e^{-a^2/3}$ for each of $(1,2)$, $(2,1)$, and $(2,2)$, so the odds ratio is $e^{-a^2/3}/e^{-2a^2/3} = e^{a^2/3}$. □

*Proof of Proposition 6.24.* The chain rule of lemma 3.12, applied to mutual information as a relative entropy, gives $I(S_j;\xi^+, O \mid \Upsilon) = I(S_j;\xi^+ \mid \Upsilon) = I(S_j;O \mid \Upsilon) + I(S_j;\xi^+ \mid O,\Upsilon)$, because $I(S_j;O \mid \xi^+,\Upsilon) = 0$ by the conditional Markov property; nonnegativity proves the first inequality, including infinite values, and the second is identical. For finite alphabets without side information this is the data-processing inequality [21, Section 2.8, Theorem 2.8.1]. □

*Proof of Proposition 6.26.* (a) Necessity follows by evaluating the commuting identity on $\ker M_j$. For sufficiency define $\overline{\mathsf{L}}_e(M_j\mathsf{u}) = M_i\mathsf{L}_e\mathsf{u}$; the condition makes this independent of the representative, and surjectivity of $M_j$ gives its domain and its uniqueness. (b) If $M\mathcal{A} = \overline{\mathcal{A}}M$, then $M\mathcal{A}^n = \overline{\mathcal{A}}^n M$ for every $n$, and summing the exponential series gives $Me^{t\mathcal{A}} = e^{t\overline{\mathcal{A}}}M$; conversely, differentiating at $t = 0$ returns (6.50), and $M\mathcal{A}$ factors through $M$ exactly when it vanishes on $\ker M$, which is (a) with $\mathsf{L}_e = \mathcal{A}$ and $M_j = M_i = M$. (c) Along any solution $\frac{d}{dt}(M_t\mathsf{u}) = (\dot{M}_t + M_t\mathcal{A}_t)\mathsf{u}$, so (6.51) is sufficient. For necessity fix $t$ and $\mathsf{v} \in E$; the linear equation with continuous coefficients has a solution through $\mathsf{u}(t) = \mathsf{v}$, so $(\dot{M}_t + M_t\mathcal{A}_t)\mathsf{v} = \overline{\mathcal{A}}_t M_t\mathsf{v}$ for every $\mathsf{v}$, and a linear map $E \to \overline{E}$ factors through the surjection $M_t$ exactly when it vanishes on $\ker M_t$, with a unique factor. Multiplying (6.51) on the right by $M_t^+$ and using $M_tM_t^+ = I$ gives the formula; in fixed inner products $M_t^+ = M_t^\top(M_tM_t^\top)^{-1}$ is continuous. (d) For the swap, $\ker M_j$ is spanned by the second coordinate, which the swap sends outside $\ker M_i$. For the flow, choose $\mathsf{u} \in \ker M$ with $M\mathcal{A}\mathsf{u} \neq 0$; the solutions from $\mathsf{u}$ and from $0$ have retained value $0$ at $t = 0$ and retained velocities $M\mathcal{A}\mathsf{u}$ and $0$. (e) The conditional law is the image of the law of $\mathsf{z}^{\mathrm{in}}$ under the diagonal map, which is a product only for a degenerate law. □

## B.2 Proofs for Fiber Information Geometry and Inference Histories

*Proof of Proposition 7.3.* Each $\widehat{\rho}_x(g)$ is pushforward by the bimeasurable sample map $\rho_x(g)$, which does not depend on the law parameter. Pushforward is a unitary map $L^2(p) \to L^2(\rho_x(g)_{\#}p)$, $f \mapsto f \circ \rho_x(g)^{-1}$, differentiability in quadratic mean makes $\mathbf{s}_Z \circ \rho_x(g)^{-1}$ the score of the pushed direction, and the change of variables preserves (7.1) and (7.2); since $\iota_{u\cdot g} = \iota_u \circ \widehat{\rho}_x(g)$, $(\iota_u^{-1})^*\tau$ is independent of $u$ exactly when $\widehat{\rho}_x(g)^*\tau = \tau$ for every $g$. For the counterexample, $\mathrm{X} = \mu/\sqrt{2}$ and $\mathrm{S} = \varsigma$ turn $g^F$ into $2(d\mathrm{X}^2 + d\mathrm{S}^2)/\mathrm{S}^2$, twice the hyperbolic metric of the upper half-plane, of which the real Möbius map $z \mapsto (1+z)/(1-z)$ of determinant $2 > 0$ is an isometry. It fixes $z = \mathrm{i}$ with derivative $2/(1-\mathrm{i})^2 = \mathrm{i}$, a quarter turn, so it sends the $g^F$-unit vector $\partial_\mu$ at $(0,1)$ to $\partial_\varsigma/\sqrt{2}$. There the only nonzero components of $T^F$ are $T^F(\partial_\mu, \partial_\mu, \partial_\varsigma) = 2$ and $T^F(\partial_\varsigma, \partial_\varsigma, \partial_\varsigma) = 8$, so the cubic values of the two unit vectors are $0$ and $8/2^{3/2} = 2\sqrt{2}$. □

*Proof of Proposition 7.6.* Invariance is (3.24) for the sample maps $\rho(g)$. With local parameters $\lambda$ and $\lambda'$ of the two arguments, primed derivatives acting on $\lambda'$, and $\mathbf{s}_i = \partial_i \log p_\lambda$, the identities $\mathbb{E}_\lambda[\mathbf{s}_i] = 0$, $\partial_i p_\lambda = p_\lambda \mathbf{s}_i$, and $\partial_i\partial_j p_\lambda = p_\lambda(\mathbf{s}_i\mathbf{s}_j + \partial_i \mathbf{s}_j)$ and differentiation at $\lambda' = \lambda$ under hypothesis 7.1 give

$$\begin{aligned} -\partial_i\partial'_j D_{\mathrm{KL}}(p_\lambda \| p_{\lambda'}) &= \mathbb{E}_\lambda[\mathbf{s}_i\mathbf{s}_j], \\ -\partial_i\partial_j\partial'_k D_{\mathrm{KL}}(p_\lambda \| p_{\lambda'}) &= \mathbb{E}_\lambda[(\mathbf{s}_i\mathbf{s}_j + \partial_i\mathbf{s}_j)\mathbf{s}_k], \\ -\partial'_i\partial'_j\partial_k D_{\mathrm{KL}}(p_\lambda \| p_{\lambda'}) &= \mathbb{E}_\lambda[\mathbf{s}_k\partial_i\mathbf{s}_j]. \end{aligned} \tag{B.1}$$

The second derivatives of the score are symmetric, so the last two lines differ by $\mathbb{E}_\lambda[\mathbf{s}_i\mathbf{s}_j\mathbf{s}_k]$, and positivity of the metric jet is positivity of $g^F$. The same dualistic structure is derived from contrast functions in Amari and Nagaoka [2], Ay et al. [4]. □

*Proof of Corollary 7.7.* Positive definiteness of $g^F$ identifies the radical of $\mathfrak{k}_q$ with the kernel of $\xi \mapsto \xi_{\mathcal{B}}(q)$, which is $\mathfrak{g}_q$. The curve $\gamma(t) = \widehat{\rho}(\exp t\xi)q$ is the integral curve of $\xi_{\mathcal{B}}$ through $q$, so by uniqueness of integral curves $\xi_{\mathcal{B}}(q) = 0$ exactly when $\gamma \equiv q$, that is, when $\exp t\xi \in G_q$ for every $t$; this gives the vanishing clause, and since $G_q$ is closed, its Lie algebra is the set of $\xi$ with $\exp t\xi \in G_q$ for every $t$, which is $\mathfrak{g}_q$. The function $\phi(t) = \mathscr{D}(q, \gamma(t))$ is $C^3$ with $\phi(0) = \phi'(0) = 0$ because $\mathscr{D}(q,\cdot)$ is minimal at $q$, and differentiating $\partial'_j\mathscr{D}(\lambda,\lambda) = 0$ along the diagonal gives $\partial_i\partial'_j\mathscr{D} = -\partial_i\partial'_j\mathscr{D} = (g_{\mathscr{D}})_{ij}$ there, so $\phi''(0) = g_{\mathscr{D}}(\xi_{\mathcal{B}}(q), \xi_{\mathcal{B}}(q))$ and Taylor's theorem gives the expansion; for relative entropy $g_{\mathscr{D}} = g^F$ by proposition 7.6. □

*Proof of Theorem 7.8.* Pollard's Le Cam decomposition at $\lambda_0$ is $P_{\lambda_0+\delta} = \mathsf{p}_\delta P_{\lambda_0} + P_\delta^\perp$ with $P_\delta^\perp \perp P_{\lambda_0}$, $P_\delta^\perp(\mathrm{Y}) = o(\|\delta\|^2)$, and $\int(\sqrt{\mathsf{p}_\delta} - 1 - \frac{1}{2}\delta^\top \mathbf{s}_{\lambda_0})^2 dP_{\lambda_0} = o(\|\delta\|^2)$. Hence $P_{\lambda_0+\delta} \otimes K = \mathsf{p}_\delta(y)(P_{\lambda_0} \otimes K) + P_\delta^\perp \otimes K$, where the second term has mass $o(\|\delta\|^2)$ and is singular, because a $P_{\lambda_0}$-null set carrying $P_\delta^\perp$, multiplied by $\bar{\mathrm{Y}}$, carries it and is $P_{\lambda_0} \otimes K$-null since each $K(y,\cdot)$ is a probability, and the square-root remainder of the first term is unchanged because the kernel integrates to one. The joint experiment is therefore differentiable in quadratic mean with score $\mathbf{s}_{\lambda_0}(Y)$, without a common dominating measure, and Pollard's preservation theorem for measurable statistics [68, Theorem 7], applied to $(y, \bar{y}) \mapsto \bar{y}$, gives differentiability of $\{P_\lambda K\}$ with the score (7.11) in $L^2(P_{\lambda_0}K; \mathbb{R}^d)$; the projected singular term may acquire an absolutely continuous part, of mass still $o(\|\delta\|^2)$. The law of

total covariance gives (7.12), and since conditional expectation is an orthogonal projection, equality in the direction $z$ holds exactly when $z^\top \mathbf{s}_{\lambda_0}(Y)$ is $\bar{Y}$-measurable. □

*Proof of Theorem 7.14.* Since $h_Q(Z,\cdot) = g(d\,\mathrm{ret}_Q\, Z, d\,\mathrm{ret}_Q\,\cdot)$, the Cauchy–Schwarz inequality for the semidefinite $g$ makes $h_Q(Z,Z) = 0$, $h_Q(Z,\cdot) = 0$, and $d\,\mathrm{ret}_Q\, Z \in \operatorname{rad} g$ equivalent, which is (7.19) and the transversality criterion. The radical is the kernel of the constant-rank $C^1$ bundle map $h^\flat : T\mathcal{Q} \to T^*\mathcal{Q}$, hence a $C^1$ subbundle, and $h$ descends to the quotient because it annihilates its radical. When $\operatorname{rad} h = \ker d\,\mathrm{ret}$, ret has zero derivative along, hence is constant on, each connected leaf and factors as $\bar{r}\circ\pi$ through the leaf space, with $\pi$ a submersion, $\ker d\pi = \ker d\,\mathrm{ret}$, and $\bar{r} = \mathrm{ret}\circ\iota$ smooth for smooth local sections $\iota$ of $\pi$; then $h = \pi^*(\bar{r}^* g)$, $d\bar{r}$ is injective, and $\operatorname{im} d\bar{r} = \operatorname{im} d\,\mathrm{ret}$ meets $\operatorname{rad} g$ only in zero, so $\bar{r}^* g$ is positive definite. For a finite family, $h = \sum_y \partial p_y \partial p_y / p_y$ is the Gram matrix of the Jacobian of ret weighted by $1/p_y$, hence $C^2$ and positive definite exactly where the Jacobian has full column rank; an immersion is locally injective because the inverse function theorem [57, Ch. 4] makes a coordinate projection of ret with invertible square Jacobian block a local diffeomorphism, and the Levi-Civita connection of a $C^2$ metric is $C^1$. For softmax, $Z^\top h(z) Z = \mathrm{Var}_p(Z)$ vanishes exactly on constant vectors, so $\operatorname{rad} h = \mathbb{R}\mathbf{1} = \ker d\,\mathrm{ret}$ with constant rank $k-1$, and an objective of $p$ has $d\mathcal{F}(\mathbf{1}) = 0$, so $\nabla\mathcal{F} \in \mathbf{1}^\perp = \operatorname{im} h$, which gives the solution class and $h^+\nabla\mathcal{F} \in \operatorname{im} h = \mathbf{1}^\perp$. □

*Proof of Proposition 7.18.* The tower property gives $\mathbb{E}[\mathbf{s}\mathbf{s}_i^\top] = \mathsf{G}_i$ and its transpose, so expanding $\mathfrak{K}$ gives $\mathfrak{K} = \mathsf{G}_{\mathrm{J}} - \sum_i \mathsf{G}_i + \mathfrak{C}$, which rearranges to (7.24); the criterion is its restriction, and (7.25) is its quadratic form for $n = 2$ and unit weights. Under independence the joint score is the sum of the block scores, so $\mathbf{r} = 0$, $\mathfrak{C} = 0$, and the weight condition is $(1-w_1)\mathsf{G}_1 + (1-w_2)\mathsf{G}_2 = 0$ in the stated direction. For the normal location family with precision $\Lambda$ the joint Fisher metric is $\Lambda$ and each marginal is a location family of variance $1/(1-\mathrm{c}^2)$, so $\mathsf{G}_1 + \mathsf{G}_2 = (1-\mathrm{c}^2)\mathrm{I}_2$, which gives the displayed difference and its two values. The lift metrics are the pullbacks of $d\lambda_1^2 + d\lambda_2^2$ along the three maps. □

*Proof of Theorem 7.23.* Each atom has $p_0/\mathrm{D} = \prod_{y_{i,x}=1}(1-\lambda_{i,x})^{-1} \prod_{y_{i,x}=0} \lambda_{i,x}^{-1} > 1$, so $p_\varkappa \geq p_0 - |\varkappa|\mathrm{D} > (1-|\varkappa|)\mathrm{D} > 0$. Summing out any one coordinate cancels the parity term, because $\mathrm{D}$ does not depend on $y$ and $\mathrm{par}$ changes sign with that coordinate; this gives normalization and the product proper marginals, hence $\mathrm{mar}\circ\iota_\varkappa = \mathrm{id}$, which with polynomial dependence gives smoothness and injectivity of $d\iota_\varkappa$. The marginal map is the restriction to the open $63$-dimensional simplex of a linear map of rank six, so its fibers are open subsets of $57$-dimensional affine spaces, and $p_\varkappa$ is affine in $\varkappa$. A paired complement changes two bits, so it preserves $\mathrm{par}$ and $\mathrm{D}$ while $p_0$ transforms by Bernoulli relabeling, and a typed permutation preserves all three expressions, which proves (7.28); a single complement changes the sign of $\mathrm{par}$ and sends $\lambda_{i,x}$ to $1-\lambda_{i,x}$, which preserves $\mathrm{D}$ and proves (7.29). □

*Proof of Corollary 7.25.* Two-bit marginals are products by theorem 7.23, so centered scores of distinct bits are uncorrelated. With $\mathbf{s}_{i,x}$ the one-bit Bernoulli score, differentiation gives the residual components

$$\mathbf{r}_{i,x} = \frac{\varkappa\,\mathrm{par}(y)\mathrm{D}(\lambda)}{p_\varkappa(y;\lambda)} \left( \frac{1-y_{i,x}}{\lambda_{i,x}} - \frac{y_{i,x}}{1-\lambda_{i,x}} \right). \tag{B.2}$$

If $Z^{\top}\mathbf{r}$ vanished on every atom, positivity and $\varkappa \mathrm{D} \neq 0$ would make the $Z$-weighted sum of the six bracketed one-bit functions vanish on the cube, and comparing two atoms that differ in one coordinate forces each component of $Z$ to vanish; so $\mathfrak{K} \succ 0$. At the center $\mathrm{D} = 4^{-6}$, its derivative vanishes, and the law is (7.36), so, each one-bit score having square $4$, averaging $(1+\bar{\varkappa}\,\mathrm{par})^{-1} = (1-\bar{\varkappa}\,\mathrm{par})/(1-\bar{\varkappa}^2)$ over the cube gives (7.32), and $\bar{\varkappa} = 1/128$ gives (7.33). The one-bit tensors at the center are $4e_{i,x}e_{i,x}^{\top}$, which gives (7.34), and the last sentence compares $\varkappa = 0$ with $\varkappa \neq 0$. □

*Proof of Proposition 7.28.* Half of the atoms have each parity, so averaging (7.37) under $\mathrm{Unif}$ gives evidence $\frac{1}{2}$, Bayes' rule gives the posterior $64^{-1}(1+\bar{\varkappa}\,\mathrm{par})$, which is (7.36), and the gap identity of the free energy gives the tight value $-\log\frac{1}{2}$. Every posterior one-bit marginal is uniform, so its product law is $\mathrm{Unif}$, and since the two parities have equal mass, $-\mathbb{E}_{\mathrm{Unif}}\log K_{\mathrm{hyp}}(1 \mid Y) = \log 2 - \frac{1}{2}\log(1-\bar{\varkappa}^2)$, which is (7.38). The posterior and $\mathrm{Unif}$ have the same one-bit marginals, the posterior input and the evidence are unchanged, and the two free-energy values differ when $\varkappa \neq 0$. □

*Proof of Theorem 7.36.* The chain rule gives $\dot{\widetilde{Q}} = (\dot{Q}\circ\phi)\phi'$ with $\phi' > 0$, and the substitution $t = \phi(\widetilde{t})$ proves length invariance and (7.49). Positive speed makes the clock invertible, and dividing by $d\mathfrak{t}/dt$ gives unit speed. Substituting (7.41) into (7.45) gives (7.50), and the other two identities follow from (7.40). For a second metric, $d\mathcal{F}(Z) = \mathsf{G}'(\mathrm{grad}^{\mathsf{G}'}\mathcal{F}, Z)$, and the Cauchy–Schwarz inequality gives (7.53). For a block-diagonal metric the squared speeds add, so $\dot{\mathfrak{l}}_i \le \dot{\mathfrak{l}} \le \sum_i \dot{\mathfrak{l}}_i$, and Minkowski's integral inequality for the vector $(\dot{\mathfrak{l}}_i)_i$ gives $\big(\sum_i(\int \dot{\mathfrak{l}}_i)^2\big)^{1/2} \le \int \dot{\mathfrak{l}}$. □

*Proof of Proposition 7.38.* Along the history $\frac{d}{dt}\mathcal{F}(\gamma) = -\|\dot{\gamma}\|_{\mathsf{G}}^2 \le 0$, so $\mathcal{F}(\gamma(t))$ decreases to $\mathcal{F}_* = \mathcal{F}(Q_*)$; if it reaches $\mathcal{F}_*$ at some $t_1$, then $\dot{\gamma} = 0$ on $[t_1,\infty)$ and the duration is that of a compact segment. Otherwise choose a closed Euclidean ball $\bar{B}$ of radius $\delta_0$ around $Q_*$ in the chart on which $\mathrm{c}_1\mathrm{I} \preceq \mathsf{G} \preceq \mathrm{c}_2\mathrm{I}$ with $0 < \mathrm{c}_1 \le \mathrm{c}_2$ and the Łojasiewicz gradient inequality $\|\nabla\mathcal{F}\| \ge \mathrm{c}_0|\mathcal{F}-\mathcal{F}_*|^{\mathrm{q}}$ holds with $\mathrm{c}_0 > 0$ and $\mathrm{q} \in [0,1)$ [1, Lemma 2.1]. While $\gamma$ stays in $\bar{B}$, $\|\dot{\gamma}\|_{\mathsf{G}}^2 = \nabla\mathcal{F}^{\top}\mathsf{G}^{-1}\nabla\mathcal{F} \ge \mathrm{c}_2^{-1}\|\nabla\mathcal{F}\|^2$, hence

$$\frac{d}{dt}\,(\mathcal{F}(\gamma)-\mathcal{F}_*)^{1-\mathrm{q}} = -(1-\mathrm{q})\,(\mathcal{F}(\gamma)-\mathcal{F}_*)^{-\mathrm{q}}\,\|\dot{\gamma}\|_{\mathsf{G}}^2 \le -(1-\mathrm{q})\mathrm{c}_0\mathrm{c}_2^{-1/2}\|\dot{\gamma}\|_{\mathsf{G}}, \tag{B.3}$$

so, with $\mathrm{C} = ((1-\mathrm{q})\mathrm{c}_0\mathrm{c}_2^{-1/2})^{-1}$, a segment $\gamma([t_1,t_2]) \subseteq \bar{B}$ has Fisher length at most $\mathrm{C}(\mathcal{F}(\gamma(t_1)) - \mathcal{F}_*)^{1-\mathrm{q}}$ and Euclidean length at most $\mathrm{c}_1^{-1/2}$ times that. For $t_1$ with $\|\gamma(t_1) - Q_*\| < \delta_0/2$ and $\mathrm{c}_1^{-1/2}\mathrm{C}(\mathcal{F}(\gamma(t_1))-\mathcal{F}_*)^{1-\mathrm{q}} < \delta_0/2$, a first exit from $\bar{B}$ after $t_1$ would give a segment in $\bar{B}$ of Euclidean length both at least and less than $\delta_0/2$; so $\gamma$ stays in $\bar{B}$, has finite length on $[t_1,\infty)$, and is Cauchy, hence converges to $Q_*$. If $Q_*$ were noncritical, $\frac{d}{dt}\mathcal{F}(\gamma)$ would tend to $-\|\mathrm{grad}^{\mathsf{G}}\mathcal{F}(Q_*)\|_{\mathsf{G}}^2 < 0$, contradicting convergence of $\mathcal{F}(\gamma)$. This is the argument of Absil et al. [1, Theorem 2.2], with its angle condition supplied by the metric bounds. In the smooth counterexample the Euclidean plane is the Fisher geometry of the unit-variance normal location family, and a curve of finite length in the plane converges, so a history that does not converge has infinite duration. □

*Proof of Proposition 7.39.* Differentiating $\mathcal{F}(\exp(t\xi)\cdot Q) = \mathcal{F}(Q)$ at $t=0$ gives $\mathsf{G}(\mathrm{grad}^{\mathsf{G}}\mathcal{F}, \xi_{\mathcal{Q}}(Q)) = 0$, and for a velocity $Z \perp V_Q$ the quantity $\|Z - \xi_{\mathcal{Q}}(Q)\|^2 = \|Z\|^2 + \|\xi_{\mathcal{Q}}(Q)\|^2$ is minimal where $\xi_{\mathcal{Q}}(Q) = 0$. In the free and proper case the orbit space is a smooth manifold and $\pi$ a surjective submersion [57, Theorem 21.10]. With $\mathrm{Hor}_Q = (\ker T_Q\pi)^{\perp}$, $T_Q\pi$ maps $\mathrm{Hor}_Q$ isomorphically onto

$T_{\pi(Q)}(\mathcal{Q}/\mathbb{G})$; define $\bar{\mathsf{G}}$ by transporting $\mathsf{G}|_{\mathrm{Hor}_Q}$ through this isomorphism. Each $L_g : Q \mapsto g \cdot Q$ is an isometry with $\pi \circ L_g = \pi$, so $TL_g$ maps $\mathrm{Hor}_Q$ onto $\mathrm{Hor}_{gQ}$ with $T_{gQ}\pi \circ TL_g = T_Q\pi$ and the transported inner products agree along orbits; $\bar{\mathsf{G}}$ is smooth because, through a smooth local section $\iota_\pi$ of $\pi$ given by the rank theorem, its coefficients are those of $\mathsf{G}$ on the horizontal projections of the pushed coordinate fields, and it is the unique metric for which $\pi$ is a Riemannian submersion. The invariant $\mathcal{F}$ is constant on the fibers of $\pi$, so $\bar{\mathcal{F}} = \mathcal{F} \circ \iota_\pi$ is well defined and $C^1$ with $\mathcal{F} = \bar{\mathcal{F}} \circ \pi$; hence $d\mathcal{F}$ vanishes on $\ker T\pi$, the gradient is horizontal, and for horizontal $Z$, $\mathsf{G}(\mathrm{grad}^{\mathsf{G}}\,\mathcal{F}, Z) = d\bar{\mathcal{F}}(T\pi Z) = \bar{\mathsf{G}}(\mathrm{grad}^{\bar{\mathsf{G}}}\,\bar{\mathcal{F}}, T\pi Z)$, which gives the gradient identity. Applying $T\pi$ to (7.41) shows that $\pi \circ Q$ is a natural-gradient history of $\bar{\mathcal{F}}$, and since $T\pi$ is isometric on horizontal vectors and the velocity is horizontal, the two Fisher speeds agree. □

*Proof of Theorem 7.40.* Exact forms are closed and have zero periods; conversely, under (7.56), integrating $\boldsymbol{\alpha}$ from a base point along curves in $\mathcal{U}$ is path independent and defines a function with $d\mathfrak{t}^\perp = \boldsymbol{\alpha}$. Differentiating $\boldsymbol{\alpha} = -\mathsf{n}^{-1} d\mathcal{F}$ gives (7.57). Since $d\mathcal{F} \neq 0$, the form $d\mathsf{n} \wedge d\mathcal{F}$ vanishes exactly when $d\mathsf{n}$ is locally a multiple of $d\mathcal{F}$, that is, when $\mathsf{n}$ does not depend on the other coordinates of a chart whose first coordinate is $\mathcal{F}$. Then $d\mathfrak{t}^\perp \wedge d\mathcal{F} = -\mathsf{n}^{-1} d\mathcal{F} \wedge d\mathcal{F} = 0$, so $\mathfrak{t}^\perp$ is constant on connected pieces of level sets and locally $\mathfrak{t}^\perp = \tau \circ \mathcal{F}$ with $\tau' = -1/\mathsf{n} < 0$. Since $\boldsymbol{\alpha}(\mathbf{e}_\mathcal{F}) = \mathsf{G}(\mathbf{e}_\mathcal{F}, \mathbf{e}_\mathcal{F}) = 1$, the unit rate follows, and the kernel of $\boldsymbol{\alpha} = \mathbf{e}_\mathcal{F}^\flat$ is the orthogonal complement of $\mathbf{e}_\mathcal{F}$. In the example $\mathsf{n} = (z_1^2 + z_2^2)^{1/2}$ and $\boldsymbol{\alpha} = -(z_2 dz_1 + z_1 dz_2)/\mathsf{n}$, whose exterior derivative is (7.58), with coefficient vanishing only on the diagonal. □

*Proof of Proposition 7.41.* The unit field $\mathbf{e}_\mathcal{F}$ is $C^1$ because $\mathsf{n} > 0$. Let $\mathcal{W} \subseteq \mathbb{R} \times \mathcal{S}$ be the open set of pairs $(t, q)$ for which $\mathrm{Fl}_t^{\mathbf{e}_\mathcal{F}}(q)$ is defined in $\mathcal{U}$, and put $\Xi(t, q) = \mathrm{Fl}_t^{\mathbf{e}_\mathcal{F}}(q)$, a $C^1$ map. Each orbit meets $\mathcal{S}$, so $\Xi$ is onto; $\mathcal{F}$ decreases strictly along orbits, so no orbit is closed, and each orbit meets $\mathcal{S}$ once, so $\Xi$ is injective. At $(t, q)$, $T\Xi$ sends $\partial_t$ to $\mathbf{e}_\mathcal{F}$ and $T_q\mathcal{S}$ onto the image of a complement of $\mathbf{e}_\mathcal{F}(q)$ under the invertible map $T\,\mathrm{Fl}_t^{\mathbf{e}_\mathcal{F}}$, which also carries $\mathbf{e}_\mathcal{F}(q)$ to $\mathbf{e}_\mathcal{F}$ at the image point, so transversality makes $T\Xi$ invertible. Hence $\Xi$ is a $C^1$ diffeomorphism, and $\mathfrak{t}^\mathcal{S} = \mathrm{pr}_1 \circ \Xi^{-1}$ satisfies $d\mathfrak{t}^\mathcal{S}(\mathbf{e}_\mathcal{F}) = 1$, with $t$ the Fisher arc length because $\mathbf{e}_\mathcal{F}$ has unit speed. A level set is transversal and met once because $d\mathcal{F}(\mathbf{e}_\mathcal{F}) = -\mathsf{n} < 0$. If the level sets of $\mathfrak{t}^\mathcal{S}$ are orthogonal to $\mathbf{e}_\mathcal{F}$, then $d\mathfrak{t}^\mathcal{S}$ vanishes on the orthogonal complement and equals one on $\mathbf{e}_\mathcal{F}$, so it equals $\boldsymbol{\alpha}$. In the example the Euclidean descent orbits $\dot{z}_1 = -z_2$, $\dot{z}_2 = -z_1$ in the quadrant are $z_1 = \mathrm{c}_1 e^t + \mathrm{c}_2 e^{-t}$, $z_2 = -\mathrm{c}_1 e^t + \mathrm{c}_2 e^{-t}$ with $\mathrm{c}_2 > |\mathrm{c}_1|$, along which $\mathcal{F} = \mathrm{c}_2^2 e^{-2t} - \mathrm{c}_1^2 e^{2t}$ decreases from $+\infty$ toward $0$, so each meets $\{z_1 z_2 = 1\}$ exactly once. □

*Proof of Corollary 7.42.* At each $t$, theorem 7.8 applied to the one-parameter family $\{P_{t'}\}$ gives (7.59) and the defect identity (7.60) with its equality clause; taking nonnegative square roots and integrating gives (7.61). For a dominated channel, differentiating the joint density $p_t(y)\mathsf{k}_t(\bar{y} \mid y)$ gives the joint score $\mathbf{s}_t(y) + \partial_t \log \mathsf{k}_t(\bar{y} \mid y)$, whose projection onto the record is (7.62). In the example the fine law does not move, and a Bernoulli law of success probability $\mathrm{b}(t) = (1 + e^{-t})^{-1}$ has Fisher speed $|\dot{\mathrm{b}}|/\sqrt{\mathrm{b}(1-\mathrm{b})} = \sqrt{\mathrm{b}(1-\mathrm{b})} = (2\cosh(t/2))^{-1}$. □

*Proof of Theorem 7.47.* Put $\gamma(t) = \mathsf{F}_\ell(\mathrm{Fl}_t^{\mathscr{V}_\ell} Q)$, so $\dot{\gamma} = \mathsf{a}_\ell(\mathrm{Fl}_t^{\mathscr{V}_\ell} Q)\mathscr{V}_{\ell+1}(\gamma)$. The integrand of $\psi_Q$ is continuous and positive, so $\psi_Q$ is a $C^1$ diffeomorphism onto an open interval. The curve $\gamma \circ \psi_Q^{-1}$ satisfies $(\gamma \circ \psi_Q^{-1})' = \mathscr{V}_{\ell+1}(\gamma \circ \psi_Q^{-1})$ on $\mathcal{T}_Q'$ and starts at $\mathsf{F}_\ell Q$, so by local Lipschitz continuity it is

contained in the maximal integral curve through $\mathsf{F}_\ell Q$; this gives the inclusion and the flow identity, with the curve constructed before the coarse flow is evaluated at $\psi_Q(t)$. The endpoint formula is monotonicity of $\psi_Q$, and in the complete case the limits are the two integrals. In the examples $\mathsf{a}_\ell = \mathsf{F}'_\ell$ because both fields are $\partial_z$. The converse is the chain rule. □

*Proof of Proposition 7.53.* In the first example, (7.64) reads $(z_1, 2z_2) = \mathsf{a}(z_1, 2z_2/\mathrm{c})$, which forces $\mathsf{a} = 1$ and $\mathrm{c} = 1$ where $z_1 z_2 \neq 0$. Differentiating the compatibility identity gives $d\mathcal{F}_\ell(Z) = \chi'_\ell d\mathcal{F}_{\ell+1}(T\mathsf{F}_\ell Z) = 0$ for $Z \in \ker T\mathsf{F}_\ell$, so $\mathcal{F}_\ell$ is constant along curves tangent to the fibers, $Z_0 = \mathrm{grad}^{\mathsf{G}_\ell} \mathcal{F}_\ell$ lies in $\mathrm{Hor}_Q$, and for $Z \in \mathrm{Hor}_Q$

$$\mathsf{G}_{\ell+1}(T\mathsf{F}_\ell Z_0, T\mathsf{F}_\ell Z) = \mathsf{d}_\ell^2 d\mathcal{F}_\ell[Z] = \mathsf{d}_\ell^2 \chi'_\ell \mathsf{G}_{\ell+1}\left(\mathrm{grad}^{\mathsf{G}_{\ell+1}} \mathcal{F}_{\ell+1} \circ \mathsf{F}_\ell, T\mathsf{F}_\ell Z\right).$$

Since $T\mathsf{F}_\ell(\mathrm{Hor}_Q)$ is the whole tangent space, $T\mathsf{F}_\ell Z_0 = \mathsf{d}_\ell^2 \chi'_\ell \,\mathrm{grad}^{\mathsf{G}_{\ell+1}} \mathcal{F}_{\ell+1} \circ \mathsf{F}_\ell$, and negating gives the claim. The quotient case is $\mathsf{d}_\ell \equiv 1$ and $\chi_\ell = \mathrm{id}$. In the normal example the kernel sends $\mathcal{N}(\mu, 1)$ to $\mathcal{N}(\mu, 2)$, the two Fisher metrics are $d\mu^2$ and $\frac{1}{2}d\mu^2$, and $\mathscr{V}_\ell = -\mu$, $\mathscr{V}_{\ell+1} = -2\mu$. For the exact-contraction objective, $\mathcal{F}_{\ell+1} \circ \mathsf{F}_\ell \leq \mathcal{F}_\ell$ with equality exactly on the set $\mathcal{A}$ of fiberwise minimizers, so compatibility with $\chi_\ell = \mathrm{id}$ holds exactly on $\mathcal{A}$; an open subset of $\mathcal{A}$ would contain, with each of its points, a neighborhood of that point in its fiber, which has more than one point, so $\mathcal{A}$ has empty interior. In the last example $\mathcal{A} = \{f = s\}$ and $d\mathcal{F}_\ell(\partial_f) = f - s$, which vanishes on no open set, so the fiber condition fails for every $\chi_\ell$. □

## B.3 Proofs for Coarse Channels

The proofs for section 9.8 use the notation fixed at the head of that section, including its convention on the constant $\mathrm{C}_\mathsf{K}$.

*Proof of Theorem 9.6.* (i) Since $C(y, \mathsf{Z}) = 1$, a bounded function of $o$ alone has the same integral against $\mathbb{P}^c$ and against $\mathbb{P}$, which gives normalization and the observation marginal. For bounded measurable $f$ on Z and $g$ on O put $h(y) = \int f(z) C(y, dz)$. Then

$$\begin{aligned}\int g\Big(\int f d(\mathbf{\Pi}_o C)\Big)\mathbb{P}^c(do \times \mathsf{Z}) &= \int g\Big(\int h d\mathbf{\Pi}_o\Big)\mathbb{P}(do \times \mathsf{Y}) \\ &= \int g(o)h(y)\mathbb{P}(do, dy) = \int g(o)f(z)\mathbb{P}^c(do, dz),\end{aligned}$$

using the preserved marginal, the defining property of the fine version, and (9.2). Hence $o \mapsto \mathbf{\Pi}_o C$ is a posterior version of $\mathbb{P}^c$ at every observation where it is declared, because $C$ acts on the conditioned variable only. If $\mathbf{\Pi}^c_o(N) = 0$, then $C(\cdot, N) = 0$ $\mathbf{\Pi}_o$-almost surely, hence $\mathbb{Q}_o$-almost surely under $\mathbb{Q}_o \ll \mathbf{\Pi}_o$, and $\mathbb{Q}^c_o(N) = 0$. (ii) Disintegrated over $y$, both lifts have the marginals $\mathbb{Q}_o$ and $\mathbf{\Pi}_o$ and the common conditional $C(y, \cdot)$, so the chain rule of Lemma 3.12, in $[0, +\infty]$, gives $D_{\mathrm{KL}}(\widehat{\mathbb{Q}}_o \| \widehat{\mathbf{\Pi}}_o) = D_{\mathrm{KL}}(\mathbb{Q}_o \| \mathbf{\Pi}_o)$; disintegrated over $z$, it gives $D_{\mathrm{KL}}(\widehat{\mathbb{Q}}_o \| \widehat{\mathbf{\Pi}}_o) = D_{\mathrm{KL}}(\mathbb{Q}^c_o \| \mathbf{\Pi}^c_o) + \Delta(o)$. Adding the finite real $-\log p(o)$, the same at both scales by (i), and using that $\mathbf{\Pi}^c_o$ is the coarse posterior version gives (9.3); every term lies in $[0, +\infty]$ before the finite shift, so no infinite quantity is subtracted. (iii) $\Delta(o)$ is the integral of a nonnegative measurable function, which vanishes exactly

when the integrand vanishes $\mathbb{Q}_o^c$-almost surely, and a relative entropy vanishes exactly when its arguments agree. (iv) Finite fine divergence makes both terms on the right of (9.3) finite. If $\Delta(o) = 0$, the two divergences are equal and finite, and Theorem 9.3 and corollary 9.4 give recovery of both laws by the Bayes kernel of $\mathbf{\Pi}_o$ under $C$. Conversely, a kernel $R$ recovering both laws gives $D_{\mathrm{KL}}(\mathbb{Q}_o \| \mathbf{\Pi}_o) \geq D_{\mathrm{KL}}(\mathbb{Q}_o^c \| \mathbf{\Pi}_o^c) \geq D_{\mathrm{KL}}(\mathbb{Q}_o^c R \| \mathbf{\Pi}_o^c R) = D_{\mathrm{KL}}(\mathbb{Q}_o \| \mathbf{\Pi}_o)$ by data processing through $C$ and then $R$, so $\Delta(o) = 0$. □

*Proof of Corollary 9.10.* (i) is Theorem 9.6 with $\mathbb{P} = \mathbb{P}_I$, $C = C_A$, and the admitted recognition law. (ii) On standard Borel spaces a disintegration of $\mathbb{P}_A(\cdot \mid X)$ over $(m_A, \xi_A)$ exists and is jointly measurable (Chapter 3); it satisfies the compatibility identity by construction, while a family declared independently of $\mathbb{P}_A$ satisfies it only if assumed. (iii) The pushforward formulas hold by testing each parent law with functions of one coordinate. For two non-Dirac marginals, Proposition 3.17 gives distinct couplings with the same pair, and tensoring each with one fixed law on $\mathbf{\Xi}_A \times \mathsf{H}_A$ gives distinct laws on $\mathsf{Z}_A$. For two points $h \neq h'$ of $\mathbf{\Xi}_A \times \mathsf{H}_A$ and the $(b_A, m_A)$ marginal $R$ of the given law, $R \otimes \delta_h$ and $R \otimes \delta_{h'}$ share the pair and differ, so one of them differs from the given law. If that space is a point and one marginal is Dirac, every coupling assigns mass one to the Dirac coordinate. □

*Proof of Proposition 9.16.* (i) By equivariance, $\mathbb{J}_{g_\# Q}(A \times B) = \mathbb{J}_Q(g^{-1}A \times g^{-1}B)$. Bimeasurability makes $R_Q^g$ a Markov kernel, the Y-marginal of $\mathbb{J}_{g_\# Q}$ is $g_\#(QK)$ by Proposition 9.15, and the change of variables $y = gy'$ gives $\int \mathbf{1}_B(y) R_Q^g(y, A)(g_\#(QK))(dy) = \mathbb{J}_{g_\# Q}(A \times B)$. (ii) With $(g_\# P)K = g_\#(PK)$, $((g_\# P)KR_Q^g)(A) = (PKR_Q)(g^{-1}A) = (g_\# P)(A)$; the relative-entropy identities are invariance under bimeasurable bijections, (3.23) and (3.24). (iii) Where $(QK)(\{y\}) > 0$ the first value is the value of every version; the remaining points are $QK$-null, both branches transform as stated because $(g_\# Q)(\{gx\}) = Q(\{x\})$ and $K(gx, \{gy\}) = K(x, \{y\})$, and the first branch is a rational function of $Q$ on a Borel set. (iv) Each $R_Q^h$ with $h \in \mathbb{G}_Q$ is a version of $R_Q$ by (i). Joint measurability makes $(h, y) \mapsto R_Q^h(y, A)$ measurable, so $\bar{R}_Q$ is a Markov kernel by Tonelli's theorem and monotone convergence, and the defining identity, linear in the kernel, passes to the mixture. Left invariance of $dh$ gives $\bar{R}_Q(ky, kA) = \bar{R}_Q(y, A)$ for $k \in \mathbb{G}_Q$. If $g_\# Q = g'_\# Q$ then $g^{-1}g' \in \mathbb{G}_Q$, so $\bar{R}_Q^{g'} = \bar{R}_Q^g$; finally $\bar{R}_Q^{kg}(ky, kA) = \bar{R}_Q^g(y, A)$. □

*Proof of Proposition 9.21.* (i) Functoriality forces the stated composites, and associativity gives $C_{kr} \circ C_{\ell k} = C_{\ell r}$. (ii) is the functor axiom on two routes, and for a rule family it is the axiom on the arrows indexed by $b_1$, $b_2$, and $b_1 b_2$. (iii) Each composite of configuration-keeping or additive arrows is the same statistic of the final partition, and the witness is direct evaluation: by pairs $(0, 0, 1, 0, 1, 0) \mapsto (0, 1, 1) \mapsto 1$, by triples $(0, 0, 1, 0, 1, 0) \mapsto (0, 0) \mapsto 0$. (iv) If stochastic matrices $P, Q \geq 0$ are mutually inverse, rank gives them equal sizes. From $PQ = I$, every positive entry $P_{ik}$ forces row $k$ of $Q$ to vanish off column $i$, so each column of $P$ has at most one positive entry, and unit row sums make $P$ a permutation matrix. □

*Proof of Proposition 9.23.* Differentiating $\sum_x p_\gamma(x) = 1$ on the fixed support gives the branch-score identity, and differentiating the softmax gives $\nabla \log W_\gamma = \xi_\gamma - \bar{\xi}$, hence (9.12). In the outer product of (9.12), summation over $x$ within each branch removes both cross terms by the branch-score identity, which gives (9.13); the label marginal has score $\xi_\gamma - \bar{\xi}$, so its Fisher tensor is $\mathrm{Cov}_W(\xi_\gamma)$. Positive denominators on a finite support make the $C^3$ masses give a $C^2$ Fisher tensor, and the Leibniz

rule gives (9.14). A common shift cancels from the weights and adds one vector to every $\xi_\gamma$, which leaves the covariance unchanged as a function of $\theta$. Finally $v^\top g^F(P)v = \sum_z P(z)(v^\top \mathbf{s}_P(z))^2 \geq 0$. □

*Proof of Proposition 9.25.* On its positive support $\mathbf{s}^c = \mathbf{s}_P + \nabla \log K^{\mathrm{obs}}(o \mid \gamma, x)$, and each record row has fixed support and sums to one, so $\sum_o K^{\mathrm{obs}}(o \mid \gamma, x)\nabla \log K^{\mathrm{obs}}(o \mid \gamma, x) = 0$ and the cross terms of the second moment vanish; this is (9.17). The record is the image of the complete sample under the parameter-independent map $(\gamma, x, o) \mapsto o$, with conditional law $\mathbf{\Pi}_o$ given $o \in \mathsf{O}^+$. A finite fixed-support family with $C^3$ masses is differentiable in quadratic mean, so Theorem 7.8 gives the score projection and the Fisher loss in (9.18). For a one-point alphabet every record row equals one, so $p_\theta(o) = \sum_{\gamma,x} P_\theta(\gamma, x) = 1$. The pointwise observed-information form of this identity is due to Louis [59]; its average over the record is (9.18). □

*Proof of Proposition 9.27.* The score telescope comes first. Differentiating $P_j = P_{j-1}C_j$ at positive atoms gives $\mathbf{s}_j(X_j) = \mathbb{E}[\mathbf{s}_{j-1}(X_{j-1}) \mid X_j]$, so the single-step loss is $D_j \succeq 0$ by Theorem 7.8, and the sum telescopes. The reversed path is Markov, so $\mathbb{E}[\varepsilon_j^{\mathrm{sc}} \mid X_j, \ldots, X_m] = 0$; for $j < k$ the increment $\varepsilon_k^{\mathrm{sc}}$ is measurable with respect to $X_j, \ldots, X_m$, which gives orthogonality, and $\mathbb{E}[\varepsilon_j^{\mathrm{sc}}(\varepsilon_j^{\mathrm{sc}})^\top] = g_{j-1} - g_j$ because $\mathbb{E}[\mathbf{s}_{j-1}\mathbf{s}_j^\top] = g_j$. Domination persists under each kernel on finite spaces, so every divergence is finite, and Theorem 9.6 and corollary 9.7 at each step give $\mathcal{F}_{j-1} - \mathcal{F}_j = K_j^{\mathrm{KL}}$ at unchanged evidence; summing proves the free-energy telescope, and division by $N_m$ gives the per-agent form. □

*Proof of Theorem 9.28.* Uniform convergence of $\lambda/N$ gives $\lambda_{\gamma,N} - \lambda_{\gamma_*,N} \leq -\delta_{\mathsf{K}} N/2$ for each $\gamma \neq \gamma_*$ and large $N$; the softmax denominator contains the winning mass, so $W_{\gamma,N} \leq e^{-\delta_{\mathsf{K}} N/2}$. Suppress $N$ and write $\bar{\xi}_i = \sum_\eta W_\eta \xi_{\eta,i}$. The softmax derivatives $\partial_i W_\gamma = W_\gamma(\xi_{\gamma,i} - \bar{\xi}_i)$ and $\partial_j \bar{\xi}_i = \sum_\eta W_\eta \partial_j \xi_{\eta,i} + \mathrm{Cov}_W(\xi_i, \xi_j)$ show that, for $|\beta| \leq 2$, $\partial^\beta W_\gamma$ is $W_\gamma$ times a polynomial in the weights and in the components of the $\xi_\eta$ and $\partial\xi_\eta$, of degree at most $|\beta|$ in the latter, each of which is $O(N)$ by (9.11). Hence $\|\partial^\beta W_{\gamma,N}\|_{\mathsf{K}} \leq \mathrm{C}_{\mathsf{K}} N^{|\beta|} e^{-\delta_{\mathsf{K}} N/2}$ for $\gamma \neq \gamma_*$ and $|\beta| \leq 2$, and the same bound holds for $W_{\gamma_*,N} - 1$, which is minus the sum of the losing weights. With $\delta\xi_\gamma = \xi_\gamma - \xi_{\gamma_*}$ and $\bar{d} = \sum_{\gamma \neq \gamma_*} W_\gamma \delta\xi_\gamma$, exact centering gives $\mathrm{Cov}_W(\xi) = \sum_{\gamma \neq \gamma_*} W_\gamma \delta\xi_\gamma \delta\xi_\gamma^\top - \bar{d}\bar{d}^\top$. Each derivative of $\delta\xi_\gamma$ through order two is $O(N)$, so a term with $j$ derivatives in the sum is bounded by $\mathrm{C}_{\mathsf{K}} N^{j+2} e^{-\delta_{\mathsf{K}} N/2}$: at most $j$ derivatives fall on a weight and the two score factors contribute at most $N$ each. The product $\bar{d}\bar{d}^\top$ carries two exponential factors and obeys the same bound. Dividing by $N$ proves (9.20). The branch error equals $\sum_{\gamma \neq \gamma_*} W_\gamma(\mathcal{I}_{\gamma,N} - \mathcal{I}_{\gamma_*,N})$, whose branch differences have bounded $C^2$ norms, so the Leibniz sum with the weight bound gives the stated bound. Adding the $C^2$ error of the winning branch and the covariance error, via (9.13), proves the $C^2$ convergence, since exponentials dominate the powers of $N$. For the floor, uniform convergence in operator norm gives $\|g_N - \mathcal{I}_{\gamma_*}\| \leq \frac{1}{2}\mathrm{c}_{\min}$ on K for large $N$, so $z^\top g_N z \geq z^\top \mathcal{I}_{\gamma_*} z - \frac{1}{2}\mathrm{c}_{\min}|z|^2 \geq \frac{1}{2}\mathrm{c}_{\min}|z|^2$. □

*Proof of Proposition 9.29.* The label set is finite, so there is a positive bulk gap from $S$ to every other label at $\theta_0$; comparison with one maximizing label and convergence of $\lambda/N$ make every nonmaximizing weight exponentially small, so its limit is zero. $C^1$ convergence gives $\xi_{\gamma,N}/N \to \mathsf{v}_\gamma$. Dividing (9.13) by $N^2$ makes the branch term $O(1/N)$ and leaves $\mathrm{Cov}_{W_N}(\xi_{\gamma,N}/N)$, which converges by finite summation. For every $z$, $z^\top \mathrm{Cov}_{\bar{\mathsf{p}}}(\mathsf{v}_\gamma) z = \sum_\gamma \bar{\mathsf{p}}_\gamma [z \cdot (\mathsf{v}_\gamma - \bar{\mathsf{v}})]^2$, which vanishes exactly when

$z \cdot \mathsf{v}_\gamma$ is constant on the positive-weight labels, so the kernel is the orthogonal complement of the stated span. A positive limiting quadratic form gives the asymptotic after multiplication by $N^2$, and a zero one gives only the little-o conclusion. □

*Proof of Theorem 9.30.* A compact set of $\mathsf{w}$ is bounded, so its image under $\theta_N$ and the segments to $\theta_0$ lie in the chosen neighborhood for large $N$. Put $\lambda^{\mathrm{inc}}_{\gamma,N}(\mathsf{w}) = \lambda_{\gamma,N}(\theta_0 + \mathsf{w}/N) - \lambda_{\gamma,N}(\theta_0)$. The fundamental theorem of calculus gives

$$\lambda^{\mathrm{inc}}_{\gamma,N}(\mathsf{w}) = \int_0^1 N^{-1} \nabla \lambda_{\gamma,N}(\theta_0 + \tau \mathsf{w}/N) \cdot \mathsf{w} d\tau, \tag{B.4}$$

whose integrand tends uniformly to $\mathsf{v}_\gamma \cdot \mathsf{w}$ by the $C^1$ part of (9.11) and continuity of $\nabla \mathsf{s}_\gamma$. Differentiation in $\mathsf{w}$ gives

$$D_{\mathsf{w}} \lambda^{\mathrm{inc}}_{\gamma,N} \to \mathsf{v}_\gamma, \qquad \|D^2_{\mathsf{w}} \lambda^{\mathrm{inc}}_{\gamma,N}\| = O(N^{-1}), \qquad \|D^3_{\mathsf{w}} \lambda^{\mathrm{inc}}_{\gamma,N}\| = O(N^{-2}), \tag{B.5}$$

so $\lambda^{\mathrm{inc}}_{\gamma,N} \to \mathsf{w} \cdot \mathsf{v}_\gamma$ in $C^3$. For tied labels, the relative log masses $\lambda_{\gamma,N} - \lambda_{\gamma_0,N}$ at $\theta_N(\mathsf{w})$ are the logarithms of the center ratios plus the difference of increments, and converge in $C^3$ to $\log \mathrm{c}_\gamma + \mathsf{w} \cdot (\mathsf{v}_\gamma - \mathsf{v}_{\gamma_0})$; exponentiation and the finite softmax are smooth on the resulting compact ranges, with denominator at least the reference mass one. For a nonmaximizing label the relative center log mass is at most $-\mathrm{c}N$ for some constant $\mathrm{c} > 0$, while its increment and $\mathsf{w}$-derivatives through order three are bounded, so its relative exponential and those derivatives are $O(e^{-\mathrm{c}N})$. Normalizing proves (9.22) in $C^3$, including the zero coordinates outside $S$. With $\hat{\xi}_{\gamma,N}(\mathsf{w}) = N^{-1} \xi_{\gamma,N}(\theta_N(\mathsf{w}))$, (B.5) gives $\hat{\xi}_{\gamma,N} \to \mathsf{v}_\gamma$ in $C^2$, and the affine change $\theta_N$ has derivative $N^{-1}\mathrm{Id}$, so (9.13) becomes

$$g^F_{\mathsf{w}}(\mathcal{E}_N)(\mathsf{w}) = \mathrm{Cov}_{W_N(\theta_N(\mathsf{w}))}\big(\hat{\xi}_{\gamma,N}(\mathsf{w})\big) + \frac{1}{N} \sum_\gamma W_{\gamma,N}(\theta_N(\mathsf{w})) \mathcal{I}_{\gamma,N}(\theta_N(\mathsf{w})). \tag{B.6}$$

The covariance converges in $C^2$ by finite sums and products, without dividing by a weight that tends to zero. The branch densities and their first two $\mathsf{w}$-derivatives are bounded, since a $\mathsf{w}$-derivative is $N^{-1}$ times a $\theta$-derivative, and the weights have bounded $C^2$ norms, so every term of the Leibniz sum for the last term is $O(N^{-1})$. Finally the log-sum-exponential has gradient $\sum_\gamma \bar{\mathsf{p}}_\gamma \mathsf{v}_\gamma$ and Hessian $\mathrm{Cov}_{\bar{\mathsf{p}}}(\mathsf{v}_\gamma)$. The window is the finite-mixture counterpart of the finite-volume rounding of a first-order transition, which Borgs and Kotecký [13] derive from a lattice model as a hyperbolic-tangent crossover of width of order the inverse volume. □

*Proof of Proposition 9.31.* The common branch law cancels from the relative log mass and has zero score, so the label is the only source of information. Then $W_1' = W_1(1 - W_1)\omega_N'$, the centered label scores are $(1 - W_1)\omega_N'$ and $-W_1\omega_N'$, and averaging their squares gives (9.24), with $W_1(1 - W_1) = \mathrm{bv}(\omega_N)$. In (a) the $C^3$ norm of $\omega_N$ on a compact interval is at most $\mathrm{C}N$ and $\omega_N/N = \theta$; the formulas follow by substitution. In (b) the derivatives vanish and $\omega_N/N \to 0$ in $C^3$. In (c) the $C^3$ norms are $O(N^{\mathrm{c}})$ and $O(\sqrt{N})$, both $o(N)$, and substitution gives the two statements, with $\mathrm{bv}(0) = 1/4$ and $\mathrm{bv}(\log 2) = 2/9$. In (d) $\omega_N/N = \theta - \mathrm{c}N^{-1} \log N \to \theta$ in $C^3$ and the $C^3$ norm is $O(N)$; at zero, $\mathrm{bv}(-\mathrm{c} \log N)N^2/N = N^{1+\mathrm{c}}/(1 + N^{\mathrm{c}})^2$. In (e) derivatives through order three are bounded by $N^{-1/6}N = N^{5/6}$, so $\omega_N/N \to 0$ in $C^3$; $|\omega_N| \le N^{-1/6}$ gives uniform convergence of the weights, and $0 \le F_N \le N^{-1}\frac{1}{4}N^{1/3}$. The Leibniz rule gives (9.27). At zero $\omega_N = 0$, $\omega_N' = N^{1/6}$,

$\omega_N'' = 0$, $\omega_N''' = -N^{5/6}$, and $\mathrm{bv}(0) = 1/4$, $\mathrm{bv}'(0) = 0$, $\mathrm{bv}''(0) = -1/8$, which gives the value of $F_N''(0)$; the uniform limit has second derivative zero. □

## B.4 Proofs for Exact Agent-Network Coarse-Graining

*Proof of Theorem 11.4.* Positivity and unitality of $\mathsf{U}$ give the bounds in (a), the bound on $\mathcal{A}(\psi)$, and the boundedness and positivity of $\mathsf{U}^\psi$; substitution in (11.6) gives (11.8). If $\|h\|_\infty < \log 2$, then $\|\mathsf{U}^\psi(e^{-h}) - 1\|_\infty \le e^{\|h\|_\infty} - 1 < 1$, so the exponential series in $L^\infty(\bar{\mathsf{m}}_o)$ and the series of $\log(1+u)$ in the commutative Banach algebra $L^\infty(\bar{\mathsf{m}}_o^c)$ converge, and $h \mapsto \mathcal{A}(\psi + h)$ is real analytic on the ball of radius $\log 2$ about every $\psi$, which proves (a). The Fréchet chain rule for the exponential, $\mathsf{U}^\psi$, and the logarithm gives (11.9); differentiating $\int h(y) C^\dagger_{\psi+tk}(z, dy)$ at $t = 0$ gives (11.10); and $\psi = 0$ gives the origin. For (c), $\psi \le \psi' + c$ implies $e^{-\psi} \ge e^{-c} e^{-\psi'}$, and positivity of $\mathsf{U}$ gives the order statement; applying it in both directions with $c = \|\psi - \psi'\|_\infty$ gives (11.11), and $\mathsf{U}(e^{-\psi - c}) = e^{-c}\mathsf{U}(e^{-\psi})$ gives homogeneity, which with $\mathsf{U}1 = 1$ gives (d). For (e), $\widehat{\mathsf{m}}^\psi C = (e^{-\psi}\mathsf{m}_o)C / \mathsf{m}_o(e^{-\psi})$ and $\bar{\mathsf{m}}_o^c = \mathsf{m}_o C / M_o$ give the density $e^{-\mathcal{A}(\psi)} / \bar{\mathsf{m}}_o(e^{-\psi})$; the first two derivatives of $t \mapsto \log \bar{\mathsf{m}}_o(e^{-th})$ at $t = 0$ are $-\bar{\mathsf{m}}_o(h)$ and $\mathrm{Var}_{\bar{\mathsf{m}}_o}(h)$, and polarization gives the covariance. □

*Proof of Theorem 11.5.* Conditional Jensen gives $|\mathsf{U}h|^p \le \mathsf{U}|h|^p$ for finite $p$, and integrating the disintegration (11.5) proves (11.13) and preservation of the mean; $p = \infty$ follows from the pointwise bound. The conditional variance identity

$$\int h^2 d\bar{\mathsf{m}}_o = \int (\mathsf{U}h)^2 d\bar{\mathsf{m}}_o^c + \int \mathrm{Var}_{C^\dagger(z,\cdot)}(h)\, \bar{\mathsf{m}}_o^c(dz), \tag{B.7}$$

the decomposition behind Theorem 7.8 applied to $h$, proves (11.14), and a nonnegative conditional variance integrates to zero exactly when it vanishes almost everywhere. For the oscillation bound, fix a representative of $h$ and let $N$ be the $\bar{\mathsf{m}}_o$-null set on which it lies outside $[\operatorname{ess\,inf} h, \operatorname{ess\,sup} h]$. Disintegration gives $\int C^\dagger(z, N)\bar{\mathsf{m}}_o^c(dz) = \bar{\mathsf{m}}_o(N) = 0$, so $C^\dagger(z, N) = 0$ for almost every $z$. Thus this fixed $h$ lies between its essential bounds almost surely under those kernels. For such a pair $(z, z')$ the signed measure $\mu = C^\dagger(z, \cdot) - C^\dagger(z', \cdot)$ has total mass zero and positive and negative parts of equal mass $t \le \delta$, so $(\mathsf{U}h)(z) - (\mathsf{U}h)(z') = \int h(y)\mu^+(dy) - \int h(y)\mu^-(dy) \le t(\operatorname{ess\,sup} h - \operatorname{ess\,inf} h)$. The sets where $\mathsf{U}h$ lies within any positive distance of its essential supremum or infimum have positive $\bar{\mathsf{m}}_o^c$ measure, so the bound passes to $\operatorname{ess\,sup} \mathsf{U}h - \operatorname{ess\,inf} \mathsf{U}h$; applying it factor by factor gives the composition statement. □

*Proof of Theorem 11.10.* Fubini's theorem for the product $\nu_\ell$ gives commutation and $\mathbb{E}_{\ell,B}\mathbb{E}_{\ell,B'} = \mathbb{E}_{\ell,B\cap B'}$; expanding the product in (11.17) with this rule gives the alternating sum, and Möbius inversion on the Boolean lattice gives the orthogonality relations and the resolution of the identity. Every term of $\mathsf{P}_{\ell,B} f$ is a function of $y_B$, and the factor $I - \mathbb{E}_{\ell, V_\ell \setminus \{i\}}$ removes its mean in coordinate $i$; hence the nonempty components recover $\mathrm{g}$ and sum to $f - \mathbb{E}_{\ell,\varnothing} f$, the class $[f]$. The triangle inequality and $\|[\psi]\|_{\mathrm{osc}} \le \|\psi\|_\infty$ give $\|\mathrm{asm}_\ell\| \le 1$. For nonempty $B$ the projector $\mathsf{P}_{\ell,B}$ annihilates constants and is an alternating sum of $2^{|B|}$ contractions, so $\|\mathsf{P}_{\ell,B} f\|_\infty \le 2^{|B|}\|f - c\|_\infty$ for every constant $c$, and $\sum_{B \ne \varnothing} 2^{|B|} = 3^{|V_\ell|} - 1$. For the biased coordinates, expanding $\prod_i \big((Y_i - m) + m\big)$

with $m = \mathbb{E}Y_i = 2p-1$ gives the components $m^{n-|B|}\prod_{i\in B}(Y_i - m)$, whose supremum norms $m^{n-|B|}(1+m)^{|B|}$ sum to $(1+2m)^n - m^n = (4p-1)^n - (2p-1)^n$. □

*Proof of Proposition 11.15.* Because $\bar{\mathsf{m}}_\ell$ is a product over agents and hence over blocks, the joint law $\bar{\mathsf{m}}_\ell(dy)C_\ell(y,dz)$ is the product of the block laws $\bar{\mathsf{m}}_\ell^I(dy_I)C_I(y_I,dz_I)$; its $z$-marginal is the product of the block marginals, and the product of the block reverse kernels $C_I^\dagger(z_I,dy_I)$ is a version of the reverse kernel. For $\mathsf{g}\in\mathcal{G}_\ell^{\mathrm{intra}}$, grouping components by block writes $\psi_{\mathsf{g}} = \sum_I \psi_I$ with $\psi_I$ a function of $y_I$, so Fubini's theorem gives $\mathcal{A}_\ell(\psi_{\mathsf{g}}) = \sum_I f_I(z_I)$ with $f_I = -\log \mathsf{U}_I(e^{-\psi_I})$. For $|B|\geq 2$ choose $j\in B\setminus\{I\}$; the function $\mathbb{E}_{\ell+1,B}f_I$ does not depend on $z_j$, so the factor $I - \mathbb{E}_{\ell+1,V_{\ell+1}\setminus\{j\}}$ of the product form of (11.17) annihilates it. Hence $\mathsf{P}_{\ell+1,B}f_I = 0$, and linearity finishes the proof. □

*Proof of Proposition 11.23.* Channel covariance gives $\widehat{\mathsf{m}}_\ell\widehat{C}_\ell = (\kappa_{\ell+1})_\#(\mathsf{m}_\ell C_\ell)$ and $(e^{-\mathsf{K}_\ell\psi}\widehat{\mathsf{m}}_\ell)\widehat{C}_\ell = (\kappa_{\ell+1})_\#((e^{-\psi}\mathsf{m}_\ell)C_\ell)$. The Radon–Nikodym derivative of two measures pushed through one Borel isomorphism is the transported derivative, and applying $-\log$ gives (11.34) without selecting a pointwise reverse kernel. Differentiating at an arbitrary bounded $\psi$ with part (b) of Theorem 11.4 gives (11.35). Composing with the intertwinings (11.20) gives (11.36), and inserting the intertwining of the retained projections into $I - \mathsf{P}^{\mathrm{ret}}_{\ell+1}$ and into Proposition 11.22 gives the residual formulas. □

*Proof of Proposition 11.26.* (a) is Tonelli's theorem applied to the product of nonnegative factors. (b) is Möbius inversion on the Boolean lattice of subsets of $V^c$, and the component indexed by $B$ reads only $z_B$ because every term does. (c) follows by substituting $gz$ and $gz^\circ$ in (11.40). For (d), write $f(u) = -\log 2\cosh u$, so $f''' = 2\,\mathrm{sech}^2\tanh$. With $x = \sum_r \xi_r y_r$ and $y_r^2 = 1$, only odd powers of $x$ contain the monomial $y_1y_2y_3$; the cubic term $f'''(\xi_0)x^3/6$ contributes $f'''(\xi_0)\xi_1\xi_2\xi_3$, and higher odd powers are of order five in the couplings. At $\xi_0 = 0$ the function $f(x)$ is even in $y$, so every odd Walsh coefficient vanishes. □

*Proof of Proposition 11.28.* Adjointness is (11.41) integrated against $f(y_j)g(z_I)$. For the factorization witness take two blocks with one $\{-1,1\}$-valued agent each, a uniform product baseline, a channel that flips each coordinate independently with probability $1/5$, and the likelihood $e^{Jy_1y_2}$. The baseline reverse probability of $y_1 = 1$ given $z_1 = 1$ is $4/5$ whatever $z_2$ is, while with $x = e^{2J}$ the posterior reverse probability is $4(x+4)/(8x+17)$ at $z_2 = -1$ and $4(4x+1)/(17x+8)$ at $z_2 = 1$, which differ exactly when $J\neq 0$. For a deterministic statistic, $\mathsf{CJ}f(z) = \mathbb{E}[f(\chi(Y))\mid \chi(Y) = z] = f(z)$ almost everywhere. A channel whose output does not depend on $y$ gives $\mathsf{CJ}f = \mathbf{\Pi}_o^c(f)$, which differs from $f$ for nonconstant $f$. □

*Proof of Theorem 11.35.* Assume (11.50) and take $\mu = \delta_y$. Then $\mathsf{T}(y,\chi^{-1}A) = \mathsf{T}^\chi(\chi(y),A)$ for every $y$ and Borel $A$, whose right side depends on $y$ only through $\chi(y)$, and surjectivity determines $\mathsf{T}^\chi$ at every $z$. Conversely, for fixed $z$ the map $A\mapsto \mathsf{T}(\iota(z),\chi^{-1}A)$ is a probability measure because $\chi^{-1}$ commutes with countable disjoint unions, and for fixed $A$ it is Borel in $z$ as the composition of $\iota$ with $y\mapsto \mathsf{T}(y,\chi^{-1}A)$. Since $\chi(\iota(\chi(y))) = \chi(y)$, hypothesis (11.49) gives the same pointwise identity, and integrating it against $\mu$ gives (11.50). □

*Proof of Theorem 11.44.* At $t = 0$ the empty context has mass one. If the context at $t$ has its $\mathbb{P}_N^{\mathrm{path}}$ marginal, multiplying by its row gives the correct context-plus-symbol marginal, and marginalizing the oldest symbol gives the next context marginal; by induction no positive prefix enters a null

context, and successive row sums give normalization. Expanding the logarithm of the products and regrouping by contexts gives (11.58); a competitor that excludes a positive word has infinite divergence. Equality forces every occupied row to agree, and recursive reachability forces the same path law. Subtracting the source chain-rule logarithm from the projected product leaves, in expectation, the difference between entropy given the context and entropy given the full past, which is the conditional mutual information (11.57) by the chain rules of Cover and Thomas [21, Sec. 2.5]. Conditioning on more history does not increase entropy, so the defect decreases in $k$, and $k \geq N-1$ keeps the whole past. The projection keeps its own defining marginals, and a shorter prefix only clips the available history. The total-variation bound is Pinsker's inequality [21, Lemma 11.6.1]. For the cycle, $\mathbb{P}(Z_0 = 0) = 2/3$ and both one-step rows from $0$ put mass $1/2$ on $0$, so the projection gives $000$ the mass $1/6$ although the cycle never produces it. Subtracting the two finite log ratios gives (11.59). □

*Proof of Theorem 11.45.* Stationarity and nonnegativity force $\mathsf{T}_{yy'} = 0$ when $\bar{p}_y > 0 = \bar{p}_{y'}$, so the positive support is closed. For $k \geq 1$ the positive $k$-context chain shifts a context to its last $k$ symbols after appending a symbol, with row $\mathbb{P}^{\mathrm{path}}_{k+1}(\mathbf{z}z)/\mathbb{P}^{\mathrm{path}}_{k}(\mathbf{z})$; its incoming mass at a context is the corresponding marginal of $\mathbb{P}^{\mathrm{path}}_{k+1}$, which equals $\mathbb{P}^{\mathrm{path}}_{k}$ of that context by stationarity, and this gives stationarity and exact $(k+1)$-windows; for $k = 0$ the projection is the stationary product of the one-symbol marginal. Conditioning and stationarity give $0 \leq \mathrm{h}_{j+1} \leq \mathrm{h}_j$. The chain rule gives $H(Z_{0:N-1}) = \sum_{j<N} \mathrm{h}_j$, and for $N > k$ the projected cross-entropy is $H(Z_{0:k-1}) + (N-k)\mathrm{h}_k$, which gives the sum and its rate. For two adjacent observed blocks of lengths $n_1$ and $n_2$, past and future are conditionally independent given the intervening hidden state, so their mutual information is at most $\mathrm{c}_Y$; expanding it by stationary entropies gives $H(Z_{0:n_2-1}) - \sum_{j=0}^{n_2-1} \mathrm{h}_{n_1+j} \leq \mathrm{c}_Y$, and letting $n_1$ and then $n_2$ grow gives the series bound, whose monotone terms give $(k+1)\mathrm{h}^{\mathrm{ex}}_k \leq \mathrm{c}_Y$. Finally, for $N > k$, $N\mathrm{h}^{\mathrm{ex}}_k - D^{(k)}_N = k\mathrm{h}^{\mathrm{ex}}_k + \sum_{t=k}^{N-1} \mathrm{h}^{\mathrm{ex}}_t \leq \mathrm{c}_Y$, for $N \leq k$ the bound $N\mathrm{h}^{\mathrm{ex}}_k \leq \mathrm{c}_Y$ holds, and bounding each nonzero summand of $D^{(k)}_N$ by $\mathrm{h}^{\mathrm{ex}}_k$ gives the last inequality. □

*Proof of Proposition 11.46.* Relabelings fix the uniform law and its Fisher metric, so a natural rule would be permuted by them and would induce a homomorphism from the symmetric group $S_7$ to the symmetric group $S_3$ of the three blocks. Its restriction to the alternating group $A_7$, which is perfect, is trivial, because $S_3$ is solvable, so every block would be $A_7$-invariant. The tangent space is the sum-zero representation of $A_7$. Since $A_7$ is doubly transitive on the seven outcomes, the commutant of its permutation representation is spanned by the identity and the all-ones matrix, so every equivariant endomorphism of the sum-zero space is scalar, the representation is irreducible, and it has no invariant two-dimensional subspace. □

# Appendix C

# Proofs for Finite-Depth Formation

This appendix proves the propositions of chapter 13, the lemmas behind Theorem 13.17, Proposition 13.22, and its extension Proposition 14.4 of chapter 14, in the symbols of table 13.1; Table C.1 orders the lemmas. The depth $L \geq 2$ is fixed. Constants $\mathrm{c}_L, \mathrm{C}_L > 0$ may change between displays; they depend on $L$, on the fixed data, and, once chosen, on $k_L$ and $\eta_L$, but not on $\epsilon$, $\mathrm{A}$, $\mathrm{T}$, or the strengths unless the dependence is displayed, and all limits have $\epsilon, \mathrm{T} > 0$. Densities are relative to Lebesgue measure times counting measure, values of the complete assignment are written $j, k \in \mathcal{J}$, and the independent equilibrium root factor of every law is omitted from non-root integrals and appended unchanged. On $M_\ell$, $H_\ell$ is the Gaussian Schur precision and $\mathrm{d}_\ell = \dim M_0 - \dim M_\ell$ the number of eliminated coordinates.

**Table C.1:** Lemma dependencies of the finite-depth formation theorem.

| Lemma | Inputs | Output | Exponent or constant |
|---|---|---|---|
| Lemma C.1 | Lemma 13.2 | parent potentials, rows, covector; force, Hessian, metric bounds | $B_* = \delta\sqrt{2N}$, $\delta_* < 1/2$ |
| Lemma C.2 | Lemma 13.2 | prepared direction, strict margin | $\mathrm{m}_L > 0$ |
| Lemma C.3 | Lemma C.1 | well-posedness, moments, parameter bounds | $\mathrm{A}^2\mathrm{J}_{\ell+1} + \mathrm{T}\varsigma_\epsilon^2$ |
| Lemma C.4 | Lemma C.3 | calibration-dominated parameter | $r_\epsilon = O(\epsilon^{2\mathrm{Q}_1+20})$ |
| Lemma C.5 | cited inequalities | log-Sobolev, Herbst, duality, transport, Pinsker | curvature, $\mathrm{T}$ explicit |
| Lemma C.6 | Lemmas C.1, C.3 and C.5 | activation and Fisher budgets; score moments | $\mathrm{A}^2\mathrm{J}_{\ell+1} + \mathrm{T}\varsigma_\epsilon^2$ |
| Lemma C.7 | Lemmas C.5 and C.6 | early discarded loss; covector error | $\epsilon^{14}$; $\frac{1}{2}\mathrm{c}_{\mathrm{early},L}k_L > \mathrm{Q}_1 + 20$ |
| Lemma C.8 | Lemmas C.1 and C.6 | terminal loss by joint relaxation | $\epsilon^2$; $\mathrm{c}_{\mathrm{joint},L}k_L > \mathrm{Q}_1 + 8$ |
| Lemma C.9 | Lemma C.1 | averaged row; environment errors | $\mathrm{e}_\epsilon(y)$ |
| Lemma C.10 | Lemmas C.3 and C.9 | mixing to the moving fast equilibrium | rate $\mathrm{c}_{\mathrm{fast},L}\mathrm{J}_1$ |
| Lemma C.11 | Lemmas C.4, C.9 and C.10 | assignment histories | $\epsilon^6$; $\mathrm{c}_{\mathrm{fast},L}k_L > \mathrm{Q}_1/2 + 12$ |
| Lemma C.12 | Lemmas C.2 and C.3 | coherent components | $e^{-\mathrm{c}_L/\epsilon^6}$; $k_{\mathrm{geom},L}$, $\eta_L^{\max}$ |
| Lemma C.13 | Lemma C.4 | source movement; kernel change | $57/100$; $1/200$ |
| Lemma C.14 | Lemmas C.4 and C.12 | common-path concentration | $\mathrm{A}\epsilon^2$; $e^{-\mathrm{c}_L/\epsilon^2}$ |

## C.1 Exact marginals, rates, and parameter derivatives

*Proof of Proposition 13.4.* Under $\nu_i^*$, $k_i - m_i = \mathfrak{d}_i + \xi_i'$, whereas the physical evaluator gives $k_i - m_i = \xi_i'$ with the same law of $m_i$. The log-density ratio $[\mathfrak{d}_i(k_i - m_i) - \mathfrak{d}_i^2/2]/(1 - \sigma_{\mathrm{ev}}^2)$ therefore

has expectation $\mathfrak{d}_i^2/(2(1-\sigma_{\mathrm{ev}}^2))$ under every coupling of the two marginals, since $\mathbb{E}[k_i-m_i]=\bar{k}_i-\bar{m}_i$; this proves eq. (13.12), and summation gives $\mathbf{x}^\top R_{\mathrm{ev}}\mathbf{x}/2$. The rows $\mathfrak{d}_i$ belong to $B_1$, so $\mathrm{J}_1\mathfrak{D}^\top\mathfrak{D}\preceq H$, and $\mathrm{T}\le \mathrm{J}_1(1-\sigma_{\mathrm{ev}}^2)/2$ gives $\mathrm{T}R_{\mathrm{ev}}\preceq H/2$; with $\vartheta\ge 1$ and the orthogonal split of Lemma 13.2, $\Omega_0(\vartheta)\succeq(\vartheta/2)H+P_{\mathcal{K}}\succ 0$. Every factor of eq. (13.11) is a positive normalized kernel of inputs none of which is a recognition law, so iterated integration gives total mass one and the records a positive density. □

*Proof of Proposition 13.6.* Since $\widehat{K}_\vartheta-\mathrm{T}R_{\mathrm{ev}}\succeq(\vartheta/2)H+(1+\kappa)P_{\mathcal{K}}\succ 0$ and $0\le\phi_S\le B_0$, both integrals are finite and positive. At fixed $\mathbf{x}$ both laws are proportional to $\pi_S(\vartheta)e^{-\phi_S(\mathbf{x})}$ in $S$, which gives eq. (13.14), a product of receiver rows because $\phi_S$ is a sum and $\pi_S$ a product over receivers. Define

$$\begin{aligned}
\mathrm{n}_{\mathcal{K}}&=\big(\kappa/(2\pi\mathrm{T})\big)^{N/2}, & \mathrm{n}_{\mathrm{cal}}&=(2\pi\mathrm{T})^{-1/2},\\
\mathrm{B}_{\mathrm{T}}(\vartheta)&=\mathrm{T}\log Z_0(\vartheta)-\mathrm{T}\log\big(\mathrm{n}_{\mathcal{K}}\mathrm{n}_{\mathrm{cal}}q_0^N\big)+\tfrac12\kappa|z_\star|^2+\tfrac12(\vartheta-2)^2, & &\\
\mathrm{C}_{\mathrm{T}}(\vartheta)&=\mathrm{B}_{\mathrm{T}}(\vartheta)-\mathrm{T}\log Z_{\mathrm{opt}}(\vartheta), & p_\vartheta(\mathsf{o})&=e^{-\mathrm{B}_{\mathrm{T}}(\vartheta)/\mathrm{T}}Z_{\mathrm{phys}}(\vartheta).
\end{aligned}\tag{C.1}$$

As $z_\star\in\mathcal{K}$, the log density of the joint at the records is $-[\mathrm{E}_\vartheta^{\mathrm{phys}}+\mathrm{T}\phi_S+\tfrac12\kappa|z_\star|^2+\tfrac12(\vartheta-2)^2]/\mathrm{T}-\log Z_0+\log(\mathrm{n}_{\mathcal{K}}\mathrm{n}_{\mathrm{cal}}q_0^N)+\log\pi_S$ plus the physical evaluator density. Hence $\mathrm{T}\mathbb{E}\log(q\prod_i\nu_i^*/p)$, in which eq. (13.12) adds $\frac{\mathrm{T}}{2}\mathbf{x}^\top R_{\mathrm{ev}}\mathbf{x}=\mathrm{E}_\vartheta-\mathrm{E}_\vartheta^{\mathrm{phys}}$, is the first line of eq. (13.15), and completing the relative entropy to $q_\vartheta^\star$ gives the second. Integrating the joint over $(\mathbf{x},S)$ at the records gives $p_\vartheta(\mathsf{o})$, so $\mathrm{C}_{\mathrm{T}}+\mathrm{T}\log p_\vartheta(\mathsf{o})=\mathrm{T}\log(Z_{\mathrm{phys}}/Z_{\mathrm{opt}})\ge 0$, as $\mathrm{E}_\vartheta^{\mathrm{phys}}\le\mathrm{E}_\vartheta$.

For the gap, $Z_{\mathrm{phys}}/Z_{\mathrm{opt}}=\mathbb{E}_{q_\vartheta^\star}e^{\mathbf{x}^\top R_{\mathrm{ev}}\mathbf{x}/2}$. The Gaussian law $\gamma_\vartheta^{\mathrm{E}}\propto e^{-\mathrm{E}_\vartheta/\mathrm{T}}$ has its mean in $\mathcal{K}$, which $R_{\mathrm{ev}}$ annihilates, and since $R_{\mathrm{ev}}$ has rank $N$ the eigenvalues $\lambda_k$ of $\mathrm{T}\widehat{K}_\vartheta^{-1}R_{\mathrm{ev}}$ satisfy $\lambda_k\le\mathrm{T}/(\mathrm{J}_1(1-\sigma_{\mathrm{ev}}^2))\le 1/2$ and $\sum_k\lambda_k\le N\mathrm{T}/(\mathrm{J}_1(1-\sigma_{\mathrm{ev}}^2))\le 1$, so $\mathbb{E}_{\gamma_\vartheta^{\mathrm{E}}}e^{\mathbf{x}^\top R_{\mathrm{ev}}\mathbf{x}/2}=\prod_k(1-\lambda_k)^{-1/2}\le e^{\sum_k\lambda_k}$. The density of $q_\vartheta^\star$ relative to $\gamma_\vartheta^{\mathrm{E}}\otimes\pi(\vartheta)$ is at most $e^{B_0}$ and $e^{\mathbf{x}^\top R_{\mathrm{ev}}\mathbf{x}/2}\ge 1$, so $\mathbb{E}_{q_\vartheta^\star}[e^{\mathbf{x}^\top R_{\mathrm{ev}}\mathbf{x}/2}-1]\le e^{B_0}(e^{\sum_k\lambda_k}-1)\le 2e^{B_0}\sum_k\lambda_k$, and $\log(1+y)\le y$ gives the first bound of eq. (13.16). The second follows from $D_{\mathrm{KL}}(q_\vartheta^\star\|\mathbf{\Pi}_{\vartheta,\mathsf{o}})=\log(Z_{\mathrm{phys}}/Z_{\mathrm{opt}})-\frac12\mathbb{E}_{q_\vartheta^\star}\mathbf{x}^\top R_{\mathrm{ev}}\mathbf{x}$. □

**Lemma C.1 (Weighted geometry and exact parent score).** For the triangle-chain experiment, all marked marginal laws and inherited proposal rows are normalized and compose exactly. With orthonormal flag columns $\mathsf{Q}_\ell$ spanning $\mathcal{U}_{\ell-1}\cap\mathcal{U}_\ell^\perp$, $\mathsf{B}=B[\mathsf{Q}_1\cdots\mathsf{Q}_L]$, and $\mathbf{x}_{\mathcal{U}}=\sum_\ell\mathsf{Q}_\ell w_\ell$,

$$\|\mathsf{B}^{-1}\|_F^{-2}\sum_\ell \mathrm{J}_\ell|w_\ell|^2\le\mathbf{x}_{\mathcal{U}}^\top H\mathbf{x}_{\mathcal{U}}\le\|\mathsf{B}\|_F^2\sum_\ell\mathrm{J}_\ell|w_\ell|^2,\tag{C.2}$$

and the same bounds hold for every Schur complement obtained by minimizing over an initial group of normal layers. On $1\le\vartheta\le 21/10$, the parent potentials are eq. (13.20) with the tilt and eliminated Gaussian scalar of eqs. (C.6) and (C.7) below, the inherited rows are those of eq. (13.23) and satisfy the detailed balance of Proposition 13.10, and the complete covector, including the source-prior score, is eq. (13.27). With $B_*=\delta\sqrt{2N}$ and $\delta_*=2\delta+2N\delta^2<1/2$, the parent forces,

Hessians, and scaled metric obey

$$|H_\ell^{-1/2}\mathrm{T}\nabla\psi_{\ell,j}| \le B_*\sqrt{\mathrm{T}}, \qquad (\vartheta-\delta_*)H_\ell \preceq \nabla^2 V_{\ell,j} \preceq (\vartheta+2\delta)H_\ell, \tag{C.3}$$

$$1 \le \mathsf{m}(\vartheta) \le 1+\mathrm{C}_L\mathrm{T}, \qquad |\mathsf{m}'(\vartheta)| \le \mathrm{C}_L\mathrm{T}, \tag{C.4}$$

for every fixed depth and all sufficiently small $\epsilon$. [ESTABLISHED]

*Proof.* By Lemma 13.2, rows of level at most $\ell$ annihilate $\mathcal{U}_\ell$, so $\mathsf{B}$ and $\mathsf{B}^{-1}$ are block lower triangular with invertible diagonal blocks. With $\mathsf{D}_\mathrm{J} = \mathrm{diag}(\sqrt{\mathrm{J}_\ell} I)$, each nonzero block $(\ell, r)$ of $\mathsf{D}_\mathrm{J}\mathsf{B}^{\pm1}\mathsf{D}_\mathrm{J}^{-1}$ is that of $\mathsf{B}^{\pm1}$ times $\sqrt{\mathrm{J}_\ell/\mathrm{J}_r} \le 1$, which proves eq. (C.2), and $H|_{M_0} = H_\mathcal{U} \oplus \mathrm{J}_1 I_\mathcal{H}$. The Schur quadratic is the infimum over the eliminated coordinates, so minimizing both bounds proves the retained versions; in particular $H_\ell \succeq \mathrm{h}_L I$ for a fixed $\mathrm{h}_L > 0$, since $\mathrm{J}_L = 1$, the normal part of $H_\ell$ is at most $\mathrm{C}_L\mathrm{J}_{\ell+1} I$ with $\mathrm{J}_{L+1} = 0$, and the loop part keeps the weight $\mathrm{J}_1$.

For a cumulative hidden/retained split $(z, y)$ of $M_0$ with blocks $H_{zz}^{(\ell)}$, $H_{zy}^{(\ell)}$ of $H$, completing the square gives the lift in the first line below, with $z$ embedded in the hidden subspace. The centered hidden Gaussian $\gamma_{\ell,\vartheta}^{\mathrm{hid}}$, of covariance $(\mathrm{T}/\vartheta)(H_{zz}^{(\ell)})^{-1}$, has a mean independent of $\vartheta$, and integrating the source factor gives eq. (13.20) with the tilt, the scalar, and the normalized optimizer of the other lines, $\mathrm{C}_0^\mathrm{G} = 0$:

$$\Xi_\ell y = \big(-(H_{zz}^{(\ell)})^{-1}H_{zy}^{(\ell)}y,\ y\big), \qquad \Xi_\ell^\top H \Xi_\ell = H_\ell, \qquad \Xi_\ell^\top H z = 0, \tag{C.5}$$

$$\psi_{\ell,j}(y,\vartheta) = -\log \int e^{-\phi_j(\Xi_\ell y + z)} \gamma_{\ell,\vartheta}^{\mathrm{hid}}(dz), \qquad 0 \le \psi_{\ell,j} \le B_0, \tag{C.6}$$

$$\mathrm{C}_\ell^\mathrm{G}(\vartheta) = -\frac{\mathrm{T}\mathrm{d}_\ell}{2}\log\frac{2\pi\mathrm{T}}{\vartheta} + \frac{\mathrm{T}}{2}\log\det H_{zz}^{(\ell)}, \tag{C.7}$$

$$q_\vartheta^{\star,\ell}(y,j) = \frac{e^{-\psi_{\ell,j}(y,\vartheta)}\gamma_{\ell,\vartheta}(y)\pi_j(\vartheta)}{Z_\ell(\vartheta)}, \quad \gamma_{\ell,\vartheta} = \mathcal{N}\big(0, (\mathrm{T}/\vartheta)H_\ell^{-1}\big), \quad e^{-B_0} \le Z_\ell \le 1. \tag{C.8}$$

Positivity and Gaussian domination justify iterated integration and differentiation on compact parameter intervals. Proposition 13.10 gives detailed balance, composition, and $\beta_{jk}^\ell \ge e^{-B_0}\pi_k(\vartheta)$ for the rows of eq. (13.23), and the physical, optimizing, and recognition laws compose by separate integration through the same channel.

Let $\mathfrak{B}_\mathrm{int}$ have rows $\sqrt{\mathrm{J}_1}\mathbf{r}_e$ on internal edges, $\mathbf{r}_e$ the covector of $r_e$, so $\mathfrak{B}_\mathrm{int}^\top\mathfrak{B}_\mathrm{int} \preceq H$. An edge is selected $\mathrm{n}_e \in \{0,1,2\}$ times, $\sum_e \mathrm{n}_e = N$, so the force of $\phi_j = \delta\sum_i \chi(\sqrt{\mathrm{J}_1/\mathrm{T}}\, r_{ij_i})$ is $\delta\sqrt{\mathrm{T}}\mathfrak{B}_\mathrm{int}^\top\mathfrak{c}$ with $\mathfrak{c}_e = \mathrm{n}_e\chi'(\sqrt{\mathrm{J}_1/\mathrm{T}}\, r_e)$ and, since $|\chi'|, |\chi''| \le 1$, $|\mathfrak{c}|^2 \le 2N$. This gives the first line below on $M_0$, and with the centered Gaussian variable of eq. (C.6) held fixed, the normalized tilted law $\nu_{\ell,y,j}$ gives the other two:

$$|H^{-1/2}\mathrm{T}\nabla\phi_j| \le B_*\sqrt{\mathrm{T}}, \qquad -2\delta H \preceq \mathrm{T}\nabla^2\phi_j \preceq 2\delta H, \qquad B_* = \delta\sqrt{2N}, \tag{C.9}$$

$$\mathrm{T}\nabla\psi_{\ell,j} = \mathbb{E}_{\nu_{\ell,y,j}}\Xi_\ell^\top(\mathrm{T}\nabla\phi_j), \tag{C.10}$$

$$\mathrm{T}\nabla^2\psi_{\ell,j} = \mathbb{E}_{\nu_{\ell,y,j}}\Xi_\ell^\top(\mathrm{T}\nabla^2\phi_j)\Xi_\ell - \mathrm{T}^{-1}\,\mathrm{Cov}_{\nu_{\ell,y,j}}(\Xi_\ell^\top\mathrm{T}\nabla\phi_j). \tag{C.11}$$

Through the isometry $H^{1/2}\Xi_\ell H_\ell^{-1/2}$, the first line bounds the force by $B_*\sqrt{\mathrm{T}}$ in the $H_\ell$ norm, the covariance by $B_*^2 H_\ell = 2N\delta^2 H_\ell$, and the first Hessian term between $\pm 2\delta H_\ell$, which proves eq. (C.3). The source factor is even, so centered Gaussian integration makes $\psi_{\ell,j}$ even in $y$, and every non-root sheet has its unique minimum at zero.

The hidden Gaussian scale score $\mathfrak{s}_\ell = (\mathrm{d}_\ell - |\mathbf{Z}|^2)/(2\vartheta)$, with $\mathbf{Z}$ standard Gaussian, has mean zero and second moment $\mathrm{d}_\ell/(2\vartheta^2)$, and the tilted density ratio lies in $[e^{-B_0}, e^{B_0}]$; as the Gaussian conditional center is independent of $\vartheta$, uniformly in the retained input,

$$\partial_\vartheta \psi_{\ell,j} = -\mathbb{E}_{\nu_{\ell,y,j}}\mathfrak{s}_\ell, \qquad |\partial_\vartheta \psi_{\ell,j}| \le (e^{B_0}-1)\frac{\sqrt{\mathrm{d}_\ell}}{\sqrt{2}\vartheta}. \tag{C.12}$$

By eq. (13.20) the parent free energy is $\mathrm{T}\mathbb{E}_q \log q + \mathbb{E}_q[V_{\ell,S} - \mathrm{T}\log\pi_S] + \mathrm{B}_\mathrm{T}$, with $\mathrm{B}'_\mathrm{T} = \vartheta - 2 - \frac{\mathrm{T}}{2}\operatorname{Tr}(\Omega_0^{-1}H)$, and differentiation at a fixed recognition law, which does not differentiate assignment masses, gives eq. (13.27), whose source-prior term, present although it has no continuous gradient, has absolute unscaled score at most four for the distinguished binary prior. In the normalized joint the prior, calibration, and source-prior scores have zero mean and zero cross expectations, the other kernels have no explicit parameter dependence at fixed inputs, and the three scores have variances $\frac{1}{2}\operatorname{Tr}[(\Omega_0^{-1}H)^2]$, $1/\mathrm{T}$, and, for $4(\mathbf{1}_{\{S_{i_0}=+\}} - \pi_{i_0,+})$, $16\pi_{i_0,+}(1-\pi_{i_0,+})$, which gives eq. (13.26). As $R_{\mathrm{ev}}$ is fixed, annihilates $\mathcal{K}$, and is supported on the first normal layer, $\Omega_0$ is positive for small $\mathrm{T}$, and conjugation by the quadratic core bounds the trace terms and their derivatives uniformly; calibration contributes exactly one to $\mathrm{m}$ and the other two terms are nonnegative and $O(\mathrm{T})$, which proves eq. (C.4). □

*Proof of Proposition 13.12.* The metric and the covector are derived in Lemma C.1. The diffusion part of the generator is reversible with respect to $e^{-V_{\ell,S}/\mathrm{T}}$ on each sheet, which gives the first term of the dissipation after integration by parts. For the jump part, detailed balance (Proposition 13.10) turns $\mathrm{T}\int\sum_S \partial_t q_S \log\mathfrak{r}_S$ into $\mathrm{T}\int\sum_{S,S'} q^{\star,\ell}(y,S)\beta^\ell_{SS'}(\mathfrak{r}_{S'} - \mathfrak{r}_S)\log\mathfrak{r}_S$, and symmetrizing in $S, S'$ gives the second term, whose integrand is nonnegative because the logarithm is increasing. The parameter term is $g_\ell\dot\vartheta_\ell = -\epsilon\mathrm{m}^{-1}g_\ell^2$. At $\ell = 0$ the row $\beta_{S'}$ does not depend on $S$, and with $\mathfrak{r}_S = (q_\mathbf{X}/q^\star_\mathbf{X})q(S\mid\mathbf{x})/\beta_S$ the jump integrand sums to $2[D_{\mathrm{KL}}(q(\cdot\mid\mathbf{x})\|\beta) + D_{\mathrm{KL}}(\beta\|q(\cdot\mid\mathbf{x}))]$. These identities hold for positive smooth densities with finite entropy and moments; the moment bounds of Lemma C.3 and the cutoff passage of Lemma C.6 justify them for the admitted solutions. □

## C.2 Prepared laws and coupled continuation

**Lemma C.2 (Prepared direction with a strict margin).** Consider the list of linear functionals on $\mathcal{U}$ consisting of each internal edge covector and, for each bridge $e$ and each $\ell < \ell(e)$, the functional $r_e \circ P_\ell|_\mathcal{U}$. Every functional in the list is nonzero, and there is a unit vector $\mathbf{v}_\star \in \mathcal{U}$ with

$$\mathrm{m}_L = \min_{e,\ \ell<\ell(e)} |r_e(P_\ell \mathbf{v}_\star)| > 0, \tag{C.13}$$

where internal edges enter only at $\ell = 0$. A fixed neighborhood of $\mathbf{v}_\star$ keeps a margin of at least $\mathrm{m}_L/2$. [ESTABLISHED]

*Proof.* For a bridge $e$ and $\ell < \ell(e)$, $r_e$ is independent of the rows of $B_1, \dots, B_\ell$ (Lemma 13.2), so it does not vanish on $W_\ell$; it annihilates $\mathcal{H} \oplus \mathcal{K}$, hence does not vanish on $\mathcal{U}_\ell$, which $P_\ell$ fixes. An internal residual of triangle $q$ is a nonzero combination of the contrast rows $O_{\triangle} r_q$, which lie in $\mathcal{U}$, plus a multiple of the loop covector, so its restriction to $\mathcal{U}$ is nonzero. Evaluated in an orthonormal basis of $\mathcal{U}$ on $(1, t, \dots, t^{\dim\mathcal{U}-1})$, each of the $\mathrm{n}_0$ functionals is a nonzero polynomial of degree less than $\dim\mathcal{U}$, so some integer $t \in [0, \mathrm{n}_0(\dim\mathcal{U} - 1)]$ is a simultaneous nonroot; normalizing gives $\mathbf{v}_\star$, and continuity gives the neighborhood. □

**Lemma C.3 (Prepared continuation).** For the preparation of Definition 13.13 and every finite sequence of activations, the coupled diffusion, source proposals, and own-law parameter equation have a unique nonexplosive solution through $t_* = 2/\epsilon$, with positive smooth joint densities. On every active level $\ell$ and for all $t \le t_*$,

$$\mathbb{E}_{q_t^\ell} y^\top H_\ell y \le \mathrm{C}_L[\mathrm{A}^2 \mathrm{J}_{\ell+1} + \mathrm{T}\varsigma_\epsilon^2], \tag{C.14}$$

$$\mathbb{E}_{q_t^\ell} |y|^2 \le \mathrm{C}_L[\mathrm{A}^2 + \mathrm{T}\varsigma_\epsilon^2], \qquad \mathrm{J}_1 \mathbb{E}_{q_t^\ell} |h|^2 \le \mathrm{C}_L \mathrm{T}, \tag{C.15}$$

and the parameter satisfies $1 \le \vartheta_\ell(t) \le 21/10$ and $|\dot\vartheta_\ell(t)| \le \mathrm{C}_L\epsilon$; these bounds use no conditional-entropy or history estimate. [ESTABLISHED]

*Proof.* For a fixed positive instance and a prescribed continuous parameter path near $[1, 21/10]$, each sheet drift has a bounded Hessian and linear growth by eq. (C.3), and the proposal rows are smooth with bounded derivatives and a positive lower bound, their Gaussian-integral denominator being at least $e^{-B_0}$; interlacing the sheet diffusions on an independent rate-one Poisson grid gives a unique nonexplosive process with all polynomial moments. For the parameter feedback let $\mathfrak{u}_j^\vartheta(t', \mathbf{x}) = \mathbb{E}_\vartheta[\mathsf{f}(\mathbf{X}_t, S_t) \mid \mathbf{X}_{t'} = \mathbf{x}, S_{t'} = j]$, for smooth $\mathsf{f}$ with polynomially growing derivatives. Given the proposal times and labels, it composes Lipschitz sheet flows with products of positive row weights; flow derivatives are bounded on compact intervals, a differentiated weight is at most a constant times the weight, and the proposal count has all exponential moments, so $\mathfrak{u}$ and $\nabla\mathfrak{u}$ have polynomial bounds. Itô's formula for $\mathfrak{u}^{\widetilde\vartheta}$ along the $\vartheta$ process, stopped in a ball and then using the moments, gives

$$|\mathbb{E}_\vartheta \mathsf{f}(\mathbf{X}_t, S_t) - \mathbb{E}_{\widetilde\vartheta} \mathsf{f}(\mathbf{X}_t, S_t)| \le \mathrm{C} \int_0^t |\vartheta(t') - \widetilde\vartheta(t')| dt', \tag{C.16}$$

since the second-order coefficients agree and the drift and row differences are at most $\mathrm{C}|\vartheta - \widetilde\vartheta|(1 + |\mathbf{x}|)$ and $\mathrm{C}|\vartheta - \widetilde\vartheta|$. The observable terms of eq. (13.27) meet these hypotheses, and their explicit terminal parameter dependence adds a local Lipschitz term, so the integral map for $\dot\vartheta = -\epsilon \mathrm{m}(\vartheta)^{-1} g_\ell(q, \vartheta)$ is a contraction on a short interval, which gives local coupled existence and uniqueness. The step length is uniform: for parameter paths in $[3/4, 5/2]$, $\vartheta - \delta_* \ge 1/4$ and $H_\ell \succeq \mathrm{h}_L I$ make every non-root sheet drift uniformly confining toward zero, the root is a fixed Ornstein–Uhlenbeck factor, proposals fix and projections contract the configuration, and so the generator maps $|\mathbf{x}|^{2k}$ below $\mathrm{C}_k - \mathrm{c}_k|\mathbf{x}|^{2k}$; all moments on $[0, t_*]$, hence the constants of eq. (C.16)

and of the covector, are uniform over such paths, and finitely many steps reach $t_*$ unless $\vartheta$ leaves $[1, 21/10]$, which the stopped argument below excludes.

Stop at the first possible exit of $\vartheta$ from $[1, 21/10]$. The Gaussian comparison $d\mathbf{X}^{\mathrm{G}} = -\vartheta(t)H_\ell \mathbf{X}^{\mathrm{G}}dt + \sqrt{2\mathrm{T}}dW$, with the same parameter and Brownian increments on the active level, starts from the reference of Definition 13.13 coupled by eq. (13.31); at an activation its projected reference continues while the parent starts at the full retained sample. For $\mathbf{e} = \mathbf{X} - \mathbf{X}^{\mathrm{G}}$ the noise cancels and proposals leave $\mathbf{e}$ continuous, so eq. (C.3) gives, for every adapted label history,

$$\frac{d}{dt}|\mathbf{e}|^2 \le -\mathbf{e}^\top H_\ell \mathbf{e} + \mathrm{C}_L \mathrm{T}, \qquad \sup_{t \le t_*} |\mathbf{e}_t| \le |\mathbf{e}_0| + \mathrm{C}_L \sqrt{\mathrm{T} t_*}. \tag{C.17}$$

For the energy, variation of constants with the commuting propagator $\exp[-H_\ell \int_{t'}^{t} \vartheta]$ uses the first bound below, obtained by splitting the integral at the inverse extreme eigenvalues, where the integrand is at most $\lambda_{\max}(H_\ell)$, $1/(e\sigma)$, and $\lambda_{\min}(H_\ell)e^{-\lambda_{\min}(H_\ell)\sigma}$; Euclidean projection and the Schur variational identity contract the two bounds at each activation, and finitely many stages give the second:

$$\int_0^\infty \|H_\ell e^{-H_\ell \sigma}\| d\sigma \le \mathrm{C}_L \left[1 + \log \frac{\lambda_{\max}(H_\ell)}{\lambda_{\min}(H_\ell)}\right], \qquad \sup_t |H_\ell^{1/2} \mathbf{e}_t| \le \mathrm{C}_L \sqrt{\mathrm{T}} \varsigma_\epsilon. \tag{C.18}$$

The Gaussian covariance barrier $\Sigma^{\mathrm{G}} \preceq \mathrm{C}_L \mathrm{T} H_\ell^{-1}$ needs only $\vartheta \ge 1$, by substitution into $\dot\Sigma^{\mathrm{G}} = -\vartheta(H_\ell \Sigma^{\mathrm{G}} + \Sigma^{\mathrm{G}} H_\ell) + 2\mathrm{T} I$, and survives retention because the retained block of $H_\ell^{-1}$ is the inverse Schur complement; the deterministic mean contracts, has norm $O_L(\mathrm{A})$, and has zero loop part. This gives eqs. (C.14) and (C.15), the loop bound following from the core rate $\vartheta \mathrm{J}_1$ and a force of at most $\mathrm{C}_L\sqrt{\mathrm{T}\mathrm{J}_1}$, whose forced error is at most $\mathrm{C}_L\sqrt{\mathrm{T}/\mathrm{J}_1}$.

Substituting these moments in the covector and metric bounds gives $\dot\vartheta = \epsilon(2 - \vartheta) + O_L(\epsilon r_\epsilon)$, with $r_\epsilon$ of eq. (13.30), before any bound on $\dot\vartheta$ is used. The field points inward at both endpoints of $[1, 21/10]$ for small $\epsilon$, which excludes the stopped exit; the moments exclude explosion, so the solution reaches $t_*$ with $|\dot\vartheta| \le \mathrm{C}_L \epsilon$. Positive-temperature elliptic smoothing and positive finite rates keep densities positive and smooth, and the construction restarts after an activation from the full retained law. □

**Lemma C.4 (Comparison of independently learned parameters).** Every process or finite replacement lineage of Lemma C.3 satisfies eq. (13.39) through $t_*$, where every independently evolving descendant uses its own law and copies only the starting parameter. [ESTABLISHED]

*Proof.* The first two bounds are in Lemma C.3. Subtracting $\dot\vartheta_\circ = \epsilon(2 - \vartheta_\circ)$, $\vartheta_\circ(0) = 1$, from the equation of the process and integrating with the factor $e^{\epsilon t}$ bounds the forcing contribution by $\mathrm{C}_L r_\epsilon$ through the horizon, where $r_\epsilon = \epsilon^{2\mathrm{Q}_1+20}(1 + \epsilon^{\mathrm{Q}_1+6}\varsigma_\epsilon^2) \to 0$ by eq. (13.30), and copying the parameter at finitely many activations adds a finite multiple of that error. The comparison trajectory solves the calibration part of the own-law equation, since the calibration record contributes $\vartheta - 2$ to the covector and $1/\mathrm{T}$ to M, and every law-dependent term is of order $\epsilon^{2\mathrm{Q}_1+20}$. As $2 - e^{-2} > 7/4$, the coefficient exceeds $7/4$ on the last ten units of the horizon for small $\epsilon$. □

## C.3 Entropy estimates

**Lemma C.5 (Continuous entropy and score tools).** Let $\mu(d\mathbf{x}) \propto e^{-\mathcal{V}(\mathbf{x})/\mathrm{T}}d\mathbf{x}$ on a finite-dimensional Euclidean space, where $\mathcal{V}$ is smooth with a globally bounded Hessian and $\nabla^2\mathcal{V} \succeq \lambda I$, $\lambda > 0$. For a positive density $\mathsf{f}$ relative to $\mu$, with $\mathrm{Ent}_\mu(\mathsf{f}) = \int \mathsf{f}\log(\mathsf{f}/\int \mathsf{f}d\mu)d\mu$, for $\mathfrak{F}$ with $\mathrm{Lip}(\mathfrak{F}) \le \mathrm{K}$, for every probability law $\nu$, and whenever the right sides are finite,

$$\mathrm{Ent}_\mu(\mathsf{f}) \le \frac{\mathrm{T}}{2\lambda}\int \frac{|\nabla \mathsf{f}|^2}{\mathsf{f}}d\mu, \tag{C.19}$$

$$\log \mathbb{E}_\mu e^{\upsilon(\mathfrak{F}-\mathbb{E}_\mu\mathfrak{F})} \le \frac{\mathrm{TK}^2\upsilon^2}{2\lambda}, \qquad \upsilon \in \mathbb{R}, \tag{C.20}$$

$$\upsilon\mathbb{E}_\nu\mathfrak{F} \le D_{\mathrm{KL}}(\nu\|\mu) + \log \mathbb{E}_\mu e^{\upsilon\mathfrak{F}}, \tag{C.21}$$

so $|\mathbb{E}_\nu\mathfrak{F} - \mathbb{E}_\mu\mathfrak{F}| \le \mathrm{K}\sqrt{2\mathrm{T}D_{\mathrm{KL}}(\nu\|\mu)/\lambda}$. A centered observable $\xi$ that obeys the first line below obeys the second:

$$\log \mathbb{E}_\mu e^{\upsilon\xi} \le \frac{\mathrm{v}\upsilon^2}{2(1-\mathrm{b}|\upsilon|)} \quad \text{for } |\upsilon| < \mathrm{b}^{-1}, \tag{C.22}$$

$$|\mathbb{E}_\nu\xi| \le \sqrt{2\mathrm{v}D_{\mathrm{KL}}(\nu\|\mu)} + \mathrm{b}D_{\mathrm{KL}}(\nu\|\mu). \tag{C.23}$$

Moreover $\|\nu-\mu\|_{\mathrm{TV}} \le \sqrt{D_{\mathrm{KL}}(\nu\|\mu)/2}$. These statements apply to every conditional sheet potential of the construction. [ESTABLISHED]

*Proof.* For the semigroup $P_t$ generated by $\mathrm{T}\Delta - \nabla\mathcal{V}\cdot\nabla$, the derivative of the synchronous stochastic flow has operator norm at most $e^{-\lambda t}$, so weighted Cauchy–Schwarz gives $|\nabla P_t\mathsf{f}|^2/P_t\mathsf{f} \le e^{-2\lambda t}P_t(|\nabla\mathsf{f}|^2/\mathsf{f})$ and invariance of $\mu$ gives $\mathcal{I}_\mu(P_t\mathsf{f}) \le e^{-2\lambda t}\mathcal{I}_\mu(\mathsf{f})$ for the Fisher information $\mathcal{I}_\mu$. Integration by parts gives $d\,\mathrm{Ent}_\mu(P_t\mathsf{f})/dt = -\mathrm{T}\mathcal{I}_\mu(P_t\mathsf{f})$; for bounded positive Lipschitz $\mathsf{f}$, contraction against a stationary initial sample gives $P_t\mathsf{f} \to \int \mathsf{f}d\mu$ and, by bounded convergence, entropy tending to zero, so integrating the Fisher bound proves eq. (C.19). Cutoff, positive truncation, and short semigroup smoothing extend it by lower semicontinuity, and Gaussian tails from strong convexity justify the integrations. This is the curvature argument of Bakry and Émery [6, pp. 187, 194–200] with the diffusion coefficient made explicit; its Gaussian case is Boucheron et al. [14, Thm. 5.4]. Applied to $e^{\upsilon\mathfrak{F}}$, eq. (C.19) gives $\mathrm{Ent}_\mu(e^{\upsilon\mathfrak{F}}) \le (\mathrm{TK}^2\upsilon^2/(2\lambda))\mathbb{E}_\mu e^{\upsilon\mathfrak{F}}$, the hypothesis of Herbst's argument [14, Prop. 6.1] with variance factor $\mathrm{TK}^2/\lambda$; applied to $\pm\mathfrak{F}$ it gives eq. (C.20). The inequality eq. (C.21) is the variational formula for relative entropy [14, Cor. 4.15], and optimizing it over $\upsilon$ with eq. (C.20) gives the Lipschitz consequence. Under eq. (C.22), the transportation lemma [14, Lemma 4.18] applied to $\pm\xi$, with the inverse convex conjugate $\sqrt{2\mathrm{v}t} + \mathrm{b}t$ of the sub-gamma function [14, Sec. 2.4], gives eq. (C.23); the case $D_{\mathrm{KL}} = \infty$ is trivial. The last inequality is Pinsker's inequality [14, Thm. 4.19], with the constant $1/2$ due to Csiszár [24]. □

**Lemma C.6 (Entropy at activations).** For the law of the active level $\ell$, let $\mathfrak{E}_\ell(t) = \mathrm{T} D_{\mathrm{KL}}(q_t^\ell \| q^{\star,\ell}_{\vartheta_\ell(t)})$, with the root factor canceled. At an activation time and on each ensuing active interval,

$$\mathfrak{E}_\ell(t_\ell^{\mathrm{act}}) \le \mathrm{C}_L[\mathrm{A}^2 \mathrm{J}_{\ell+1} + \mathrm{T}\varsigma_\epsilon^2], \tag{C.24}$$

$$\mathrm{T}^2 \int_{t_\ell^{\mathrm{act}}}^{t} \mathcal{I}_\ell(\sigma) d\sigma \le \mathrm{C}_L[\mathfrak{E}_\ell(t_\ell^{\mathrm{act}}) + \mathrm{T}], \tag{C.25}$$

where $\mathcal{I}_\ell$ is the full continuous relative Fisher information of the joint law. The normalized parameter scores of the hidden Gaussian conditionals and of the joint satisfy eq. (C.22) with constants $\mathrm{v}_L, \mathrm{b}_L$ independent of the retained input, of $\mathrm{T}$, of the strengths, and of the source masses. [ESTABLISHED]

*Proof.* The normalized hidden Gaussian conditionals have centers independent of $\vartheta$, and their logarithmic parameter scores, bounded marginal tilt included, are centered versions of $-|\mathbf{Z}|^2/(2\vartheta) - \partial_\vartheta \psi_{\ell,j}$; the joint score adds $\partial_\vartheta \log \pi_j(\vartheta)$, of absolute value at most four. A bounded tilt of a standard Gaussian satisfies

$$\mathbb{E}|\mathbf{Z}|^{2m} \le e^{B_0} 2^{\dim M_0/2} 4^m m!, \qquad m \ge 0, \tag{C.26}$$

because $r^m \le m! 4^m e^{r/4}$ and $\mathbb{E} e^{|\mathbf{Z}|^2/4} \le 2^{\dim M_0/2}$, and the other score terms are bounded by eq. (C.12). Centering at most doubles a geometric moment constant, and summing the exponential series without its zero linear term, with $\log(1+x) \le x$, gives eq. (C.22), which conditional expectations keep by Jensen's inequality.

For the differential and counting entropy $\mathcal{S}(q) = -\sum_j \int q_j \log q_j$ across changes of dimension, write $q_j = \lambda_j q_{|j}$, omitting terms of zero mass. The entropy–Fisher bound $\mathcal{I}(\mathrm{f}) \ge 2\pi e \mathrm{d} e^{-2\mathcal{S}(\mathrm{f})/\mathrm{d}}$ of Stam [74], equivalent to the Gaussian logarithmic Sobolev inequality [16], follows from the Gaussian case of eq. (C.19) by expanding both sides relative to $\mathcal{N}(0, \sigma^2 I)$, integrating $\int \mathbf{x} \cdot \nabla \mathrm{f} = -\mathrm{d}$, canceling the second moments, and minimizing at $\sigma^2 = \mathrm{d}/\mathcal{I}(\mathrm{f})$; Jensen's inequality and $\mathcal{S}(q) = \sum_j \lambda_j \mathcal{S}(q_{|j}) + \mathcal{S}(\lambda)$, with $\mathcal{S}(\lambda) \ge 0$, extend it to $\sum_j \lambda_j \mathcal{I}(q_{|j}) \ge 2\pi e \mathrm{d} e^{-2\mathcal{S}(q)/\mathrm{d}}$. At fixed $\mathbf{x}$, a proposal moves the conditional label row $\mathfrak{a}$ in the direction $\mathfrak{a}\beta^\ell - \mathfrak{a}$ and leaves the continuous marginal unchanged, and concavity of categorical entropy bounds the directional derivative below by $\mathcal{S}(\mathfrak{a}\beta^\ell) - \mathcal{S}(\mathfrak{a}) \ge -\log|\mathcal{J}|$. On a $\mathrm{d}$-dimensional sheet, diffusion therefore gives the first line below, and the two bounds together give the second:

$$\mathcal{S}'(q)_{\mathrm{diff}} = \mathrm{T} \sum_j \lambda_j \mathcal{I}(q_{|j}) - \mathbb{E}_q \Delta V_{\ell,S} \ge \mathrm{T} \sum_j \lambda_j \mathcal{I}(q_{|j}) - \mathrm{d}\Lambda_\ell, \qquad \Lambda_\ell \le \mathrm{C}_L \mathrm{J}_1, \tag{C.27}$$

$$\mathcal{S}'(q) \ge 2\pi e \mathrm{d} \mathrm{T} e^{-2\mathcal{S}(q)/\mathrm{d}} - \mathrm{d}\Lambda_\ell - \log|\mathcal{J}|. \tag{C.28}$$

Scalar comparison bounds $\mathcal{S}(q)$ below by the lesser of its initial value and $\frac{\mathrm{d}}{2}\log[2\pi e \mathrm{d}\mathrm{T}/(\mathrm{d}\Lambda_\ell + \log|\mathcal{J}|)]$. At a projection dropping $k$ coordinates $z$, the chain rule and the maximum-entropy Gaussian bound, applied conditionally on $(Y, S)$ and then with Jensen's inequality, give

$$\mathcal{S}(Y,S) \ge \mathcal{S}(\mathbf{X},S) - \frac{k}{2}\log\frac{2\pi e \mathbb{E}|z-\bar{z}|^2}{k}, \tag{C.29}$$

where $\bar{z}$ can be the projected comparison mean, whose centered covariance the moment proof of Lemma C.3 bounds by $\mathrm{C}_L \mathrm{T} \varsigma_\epsilon^2 H_\ell^{-1}$. Initial entropies and all logarithms in these estimates are $O_L(\varsigma_\epsilon^2)$, because the covariance eigenvalues, $\mathrm{A}$, and $\mathrm{T}$ are fixed powers of $\epsilon$, the admitted initial density ratios are bounded, and there are finitely many projections. Expanding $\mathrm{T} D_{\mathrm{KL}}(q \| \gamma_{\ell,\vartheta} \otimes \pi(\vartheta))$ and using eq. (C.14) proves eq. (C.24); the bounded optimizer tilt adds at most $\mathrm{T} B_0$, and the source prior has a positive minimum on $[1, 21/10]$, so its logarithmic contribution does not depend on the masses $\lambda_j$.

For the time-integrated estimate put $\mathcal{V} = y^\top H_\ell y/2$ and $\mathrm{R}_L = e^{B_0} 2^{\dim M_0/2}$. The reference satisfies the first bound below. Entropy duality gives $\vartheta \mathbb{E}_q \mathcal{V} \le 2\mathfrak{E}_\ell + 2\mathrm{T} \log \mathrm{R}_L$, Jensen's inequality gives $\vartheta \mathbb{E}_{q_\vartheta^{\star,\ell}} \mathcal{V} \le 2\mathrm{T} \log \mathrm{R}_L$, and the other centered terms of the normalized full score are at most $2\mathrm{C}_{\mathrm{sc},L}$, which gives the second. Diffusion dissipates $\mathrm{T}^2 \mathcal{I}_\ell$ and detailed balance makes the source contribution nonpositive, which gives the third for a fixed $\mathrm{C}_{0,L}$:

$$\mathbb{E}_{q_\vartheta^{\star,\ell}} \exp \frac{\vartheta \mathcal{V}}{2\mathrm{T}} \le \mathrm{R}_L, \tag{C.30}$$

$$|\mathrm{T} \mathbb{E}_q \partial_\vartheta \log q_\vartheta^{\star,\ell}| \le \frac{2\mathfrak{E}_\ell + 4\mathrm{T} \log \mathrm{R}_L}{\vartheta} + 2\mathrm{T}\mathrm{C}_{\mathrm{sc},L}, \tag{C.31}$$

$$\mathfrak{E}_\ell' \le -\mathrm{T}^2 \mathcal{I}_\ell + \frac{2|\dot{\vartheta}|}{\vartheta} \big( \mathfrak{E}_\ell + \mathrm{C}_{0,L} \mathrm{T} \big). \tag{C.32}$$

The integrating factor $\exp[-2 \int |\dot{\vartheta}|/\vartheta]$ has exponent bounded over $t_*$ by the speed bound, and integration proves eq. (C.25). At an activation the law and the reference are pushed through the same channel at the same parameter, so the chain rule makes $\mathfrak{E}$ decrease along the finite lineage. All entropy derivatives are first taken for positive truncated densities on bounded domains, and the moment and entropy bounds remove the cutoffs. □

**Lemma C.7 (Conditional entropy and complete score error).** There is a constant $\mathrm{c}_{\mathrm{early},L} > 0$ such that, for every retention $\ell - 1 \to \ell$, the conditional relative entropy $\Delta_\ell$ with respect to the optimizing hidden conditional satisfies

$$\sup_{t \in I_\ell \text{ and later active times}} \Delta_\ell(t) \le \begin{cases} \mathrm{C}_L \epsilon^{14}, & \ell < L, \\ \mathrm{C}_L \epsilon^2, & \ell = L, \end{cases} \tag{C.33}$$

where the early bound also applies, until the next activation, to the corresponding unreplaced comparison process, and the terminal bound is Lemma C.8. At every such time the complete fine/parent covector difference is at most $\mathrm{C}_L \mathrm{T}(\sqrt{\Delta_\ell} + \Delta_\ell)$. [ESTABLISHED]

*Proof of the early bound and the covector estimate.* For an early retention use the auxiliary continuous split of the first line below. The true parent retains the loops, so the chain rule bounds its hidden conditional defect $\Delta_\ell$ by the auxiliary defect $\widehat{\Delta}_\ell$, which also integrates $h$; the independent root cancels, and conditional strong convexity in $Z$ is at least $\lambda_\ell = \mathrm{c}_{\mathrm{early},L} \mathrm{J}_\ell$ after reducing the

constant if necessary. The mixed Hessian obeys the second line:

$$Z = (\mathcal{U}_{\ell-1} \cap \mathcal{U}_\ell^\perp) \oplus \mathcal{H}, \qquad Y = \mathcal{U}_\ell, \qquad S \text{ retained}, \tag{C.34}$$

$$\|\nabla_Z \nabla_Y V_{\ell-1,j}\| \le \Gamma_\ell = \mathrm{C}_L \mathrm{J}_{\ell+1}, \qquad \Gamma_L = 0. \tag{C.35}$$

The nonlinear residuals lie in the first normal layer and the loop space, and at $\ell - 1 = 0$ the nonlinear potential has no future normal dependence. For $\ell - 1 \ge 1$, in coordinates $r_E = B_E w_E$ on the eliminated normal rows, the Gaussian conditional precision $\Omega_E = \mathrm{J}_E + \mathcal{C}^\top \mathrm{J}_{\mathrm{fut}} \mathcal{C}$, with the diagonal strengths $\mathrm{J}_E$ and $\mathrm{J}_{\mathrm{fut}}$ of the eliminated and future rows, has $\|\Omega_E - \mathrm{J}_E\| \le \mathrm{C}_L \mathrm{J}_\ell$ and $\|\Omega_E^{-1}\| \le \mathrm{C}_L/\mathrm{J}_{\ell-1}$, so the identity in the first line below bounds the first block row of $\Omega_E^{-1}$ by $\mathrm{C}_L/\mathrm{J}_1$. A future force column of level $r$ has size $\mathrm{C}_L \mathrm{J}_r$, so the first conditional mean responds by $\mathrm{C}_L \mathrm{J}_r/\mathrm{J}_1$, with conditional covariance comparable to $\mathrm{T}/\mathrm{J}_1$. In the derivatives of the normalized Gaussian integral defining $\psi_{\ell-1,j}$, a first-layer or loop derivative of $\phi_j$ has size $\mathrm{C}_L\sqrt{\mathrm{J}_1/\mathrm{T}}$ and a second derivative size $\mathrm{C}_L \mathrm{J}_1/\mathrm{T}$, each future derivative brings that mean response, and the covariance term is bounded by products of first-derivative bounds, which gives the second line:

$$\Omega_E^{-1} = \mathrm{J}_E^{-1} - \mathrm{J}_E^{-1}(\Omega_E - \mathrm{J}_E)\Omega_E^{-1}, \tag{C.36}$$

$$\mathrm{T}|\partial_h \partial_{w_r} \psi_{\ell-1,j}| \le \mathrm{C}_L \mathrm{J}_r, \qquad \mathrm{T}|\partial_{w_i} \partial_{w_r} \psi_{\ell-1,j}| \le \mathrm{C}_L \frac{\mathrm{J}_i \mathrm{J}_r}{\mathrm{J}_1}. \tag{C.37}$$

The Gaussian Schur cross block is at most $\mathrm{C}_L \mathrm{J}_{\ell+1}$ by the lower triangular residual form and its variational bound, and summing finitely many blocks, generated conditional covariance included, proves eq. (C.35).

Write $q = q_Y(y,j) q_{Z|Y}(z \mid y,j)$ and $q^\star = q_Y^\star(y,j) q_{Z|Y}^\star(z \mid y,j)$, and let $\mathfrak{w}_Y = \nabla_Y \log(q/q^\star)$, $\mathfrak{v}_Y = \nabla_Y \log(q_Y/q_Y^\star)$, and $\mathfrak{e}_Y = \mathbb{E}_{q_{Z|Y}} \nabla_Y \log q_{Z|Y}^\star$. Normalization gives $\mathbb{E}_{q_{Z|Y}^\star} \nabla_Y \log q_{Z|Y}^\star = 0$ and $\mathbb{E}_{q_{Z|Y}} \mathfrak{w}_Y = \mathfrak{v}_Y - \mathfrak{e}_Y$, the marginal flux is $\mathrm{T} q_Y (\mathfrak{v}_Y - \mathfrak{e}_Y)$, and subtracting the marginal entropy derivative from the full one yields

$$\begin{aligned} \widehat{\Delta}'_\ell|_{\mathrm{cont}} = & -\mathrm{T}\mathcal{I}_Z - \mathrm{T}\mathbb{E}_{q_Y} \mathrm{Var}_{q_{Z|Y}}(\mathfrak{w}_Y) + \mathrm{T}\mathbb{E}_{q_Y}\left[\mathfrak{e}_Y \cdot \mathbb{E}_{q_{Z|Y}} \mathfrak{w}_Y\right] \\ & - \dot{\vartheta} \mathbb{E}_q \partial_\vartheta \log q_{Z|Y}^\star. \end{aligned} \tag{C.38}$$

The conditional logarithmic Sobolev inequality gives $\mathrm{T}\mathcal{I}_Z \ge 2\lambda_\ell \widehat{\Delta}_\ell$. Applying eq. (C.20) to each unit-vector component of $\nabla_Y \log q_{Z|Y}^\star$, whose Lipschitz constant in $Z$ is at most $\Gamma_\ell/\mathrm{T}$, gives $|\mathfrak{e}_Y| \le \Gamma_\ell \sqrt{2 D_{\mathrm{KL}}(q_{Z|Y} \| q_{Z|Y}^\star)/(\mathrm{T}\lambda_\ell)}$, so Cauchy–Schwarz and Young's inequality bound the mixed term by $\lambda_\ell \widehat{\Delta}_\ell/2 + \mathrm{T}\Gamma_\ell^2 \mathcal{I}_Y/\lambda_\ell^2$, with $\mathcal{I}_Y$ the $Y$-component of the full relative Fisher information. The normalized conditional parameter score, sub-gamma by Lemma C.6, contributes at most $\lambda_\ell \widehat{\Delta}_\ell/2 + \mathrm{v}_L |\dot{\vartheta}|^2/\lambda_\ell + \mathrm{b}_L |\dot{\vartheta}| \widehat{\Delta}_\ell$.

For the source updates, at fixed $Y = y$ each $\log(\gamma^{\mathrm{hid}}/q_{Z|Y,j}^\star)$, with $\gamma^{\mathrm{hid}}$ the common label-independent hidden Gaussian, lies in $[-B_0, B_0]$, and the conditional entropy relative to $\gamma^{\mathrm{hid}}$ splits as $D_{\mathrm{KL}}(q_Z \| \gamma^{\mathrm{hid}}) + \mathbb{I}(Z;S)$. A proposal leaves $q_Z$ unchanged, and the mutual information is convex in the conditional label channel at fixed $q_Z$ by joint convexity of relative entropy, so its directional

derivative toward the new channel is at most $\log|\mathcal{J}|$; the difference from the optimizing target, a linear expectation of $\log(\gamma^{\mathrm{hid}}/q^{\star}_{Z|Y,j})$, changes at rate at most $2B_0$. This divides by no source mass and extends to vanishing masses by convex approximation. With $\mathrm{b}_L|\dot{\vartheta}| \le \lambda_\ell/2$ for small $\epsilon$ and the negative variance term dropped,

$$\widehat{\Delta}'_\ell|_{\mathrm{jump}} \le \mathrm{C}_{\mathcal{J}} := \log|\mathcal{J}| + 2B_0, \tag{C.39}$$

$$\widehat{\Delta}'_\ell \le -\frac{\lambda_\ell}{2}\widehat{\Delta}_\ell + \frac{\mathrm{T}\Gamma_\ell^2}{\lambda_\ell^2}\mathcal{I}_Y + \frac{\mathrm{v}_L|\dot{\vartheta}|^2}{\lambda_\ell} + \mathrm{C}_{\mathcal{J}}. \tag{C.40}$$

The initial conditional entropy is at most $\mathfrak{E}_{\ell-1}(t^{\mathrm{act}}_{\ell-1})/\mathrm{T}$. Integrating with eqs. (C.24) and (C.25) and $\mathrm{A}^2/\mathrm{T} = \epsilon^{-6}$ gives

$$\widehat{\Delta}_\ell(t) \le \mathrm{C}_L\Big[e^{-\mathrm{c}_{\mathrm{early},L}\mathrm{J}_\ell(t-t^{\mathrm{act}}_{\ell-1})}(\epsilon^{-6}\mathrm{J}_\ell + \varsigma_\epsilon^2) + \epsilon^{-6}\frac{\mathrm{J}^2_{\ell+1}}{\mathrm{J}_\ell} + \varsigma_\epsilon^2\frac{\mathrm{J}^2_{\ell+1}}{\mathrm{J}^2_\ell} + \frac{\epsilon^2}{\mathrm{J}^2_\ell} + \frac{\mathrm{C}_{\mathcal{J}}}{\mathrm{J}_\ell}\Big], \tag{C.41}$$

with $\mathrm{c}_{\mathrm{early},L}$ absorbing the factor $1/2$. For $\ell < L$, $\mathrm{J}^2_{\ell+1}/\mathrm{J}_\ell = \epsilon^{20}$ and $\mathrm{J}_\ell \ge \epsilon^{-20}$; on $I_\ell$ the entrance time and $\mathrm{c}_{\mathrm{early},L}k_L/2 > \mathrm{Q}_1 + 20$ bound the first term by $\mathrm{C}_L\epsilon^{14}$, and the other terms meet the same bound; at $\ell = L$ the last term would be $\mathrm{C}_{\mathcal{J}}$.

For the true retention, eq. (13.21) with $\ell' = \ell - 1$, differentiated at fixed recognition law and channel, gives $g_{\ell-1} - g_\ell = -\mathrm{T}\mathbb{E}_q\partial_\vartheta \log q^{\star}_{Z|Y,\mathrm{true}}$, the common scalar and the source-prior derivative canceling. The normalized conditional score has the same sub-gamma bound, so eq. (C.23), applied conditionally, with Jensen's inequality gives the covector bound $|g_{\ell-1}-g_\ell| \le \mathrm{C}_L\mathrm{T}(\sqrt{\Delta_\ell}+\Delta_\ell)$. □

**Lemma C.8 (Joint relaxation at the weakest strength).** There is a constant $\mathrm{c}_{\mathrm{joint},L} > 0$, independent of $\epsilon$, such that every lineage satisfies

$$D_{\mathrm{KL}}\big(q^\ell_t \| q^{\star,\ell}_{\vartheta_\ell(t)}\big) \le \mathrm{C}_L e^{-\mathrm{c}_{\mathrm{joint},L}t}[\epsilon^{-(\mathrm{Q}_1+6)} + \varsigma_\epsilon^2] + \mathrm{C}_L\epsilon^2. \tag{C.42}$$

Consequently the terminal part of eq. (C.33) holds once $t \ge t_L = k_L\log(1/\epsilon)$. [ESTABLISHED]

*Proof.* Fix a level and a parameter, write $q^\star = q^{\star,\ell}_\vartheta$ and $\mathfrak{r}_j = q_j/q^\star_j$. The bounded tilt gives $e^{-B_0} \le q^\star/(\gamma_{\ell,\vartheta}\otimes\pi) \le e^{B_0}$, and eq. (13.23) gives the conductance comparison in the first line below. For a probability measure $\mu$, $\mathrm{Ent}_\mu(\mathfrak{r}) = \inf_{r>0}\int[\mathfrak{r}\log(\mathfrak{r}/r) - \mathfrak{r} + r]d\mu$ with a nonnegative integrand, so $\mathrm{Ent}_{q^\star}(\mathfrak{r}) \le e^{B_0}\mathrm{Ent}_{\gamma\otimes\pi}(\mathfrak{r})$, and the product chain rule and convexity give the second line, whose first term the Gaussian entropy inequality bounds by $\mathrm{T}/(2\vartheta\mathrm{h}_L)$ times its continuous Fisher information:

$$q^\star_j(y)\beta^\ell_{jk}(y) \ge e^{-2B_0}\gamma_{\ell,\vartheta}(y)\pi_j(\vartheta)\pi_k(\vartheta), \tag{C.43}$$

$$\mathrm{Ent}_{\gamma\otimes\pi}(\mathfrak{r}) \le \sum_j \pi_j\,\mathrm{Ent}_\gamma(\mathfrak{r}_j) + \mathbb{E}_\gamma\,\mathrm{Ent}_\pi(\mathfrak{r}(y)). \tag{C.44}$$

For any positive finite vector $\mathbf{a}$, Jensen's inequality for the logarithm gives

$$\frac{1}{2}\sum_{j,k}\pi_j\pi_k(\mathbf{a}_j-\mathbf{a}_k)(\log\mathbf{a}_j-\log\mathbf{a}_k)=\sum_j\pi_j\mathbf{a}_j\log\mathbf{a}_j-\bar{\mathbf{a}}\sum_j\pi_j\log\mathbf{a}_j\geq\mathrm{Ent}_\pi(\mathbf{a}),\tag{C.45}$$

with $\bar{\mathbf{a}}=\sum_j\pi_j\mathbf{a}_j$. Comparing the Fisher measures and the conductances through their pointwise bounds, with the jump dissipation of the first line below, gives the second:

$$\mathcal{D}_{\mathrm{jump}}=\frac{1}{2}\int\sum_{j,k}q_j^\star\beta_{jk}^\ell(\mathfrak{r}_j-\mathfrak{r}_k)(\log\mathfrak{r}_j-\log\mathfrak{r}_k)dy,\tag{C.46}$$

$$D_{\mathrm{KL}}(q\|q^\star)\leq e^{2B_0}\frac{\mathrm{T}}{2\vartheta\mathrm{h}_L}\mathcal{I}_\ell+e^{3B_0}\mathcal{D}_{\mathrm{jump}}.\tag{C.47}$$

With $\mathrm{c}_{0,L}=\min\{2\mathrm{h}_Le^{-2B_0},e^{-3B_0}\}>0$, $\mathrm{T}\mathcal{I}_\ell+\mathcal{D}_{\mathrm{jump}}\geq\mathrm{c}_{0,L}D_{\mathrm{KL}}$. The time-dependent normalized full score satisfies eq. (C.22), hence, for small $\epsilon$,

$$D_{\mathrm{KL}}'\leq-\mathrm{c}_{0,L}D_{\mathrm{KL}}+\mathrm{C}_L\epsilon(\sqrt{D_{\mathrm{KL}}}+D_{\mathrm{KL}})\leq-\frac{\mathrm{c}_{0,L}}{2}D_{\mathrm{KL}}+\mathrm{C}_L\epsilon^2,\tag{C.48}$$

and $\mathrm{c}_{\mathrm{joint},L}=\mathrm{c}_{0,L}/2$. Activations decrease the relative entropy at the unchanged parameter, and each descendant then obeys the same inequality with its own parameter speed. The initial relative entropy is at most $\mathrm{C}_L[\epsilon^{-(\mathrm{Q}_1+6)}+\varsigma_\epsilon^2]$, and integration through finitely many decreases proves eq. (C.42); the condition $\mathrm{c}_{\mathrm{joint},L}k_L>\mathrm{Q}_1+8$ makes its decaying term $O_L(\epsilon^2)$ at $t_L$. Both the auxiliary and the true hidden conditional entropies are bounded by the full joint relative entropy, which proves the terminal assertion. □

## C.4 Exact fast marginals and complete source histories

**Lemma C.9 (Common fast marginal and environmental error).** In the flag coordinates $w=(w_1,\dots,w_L)$ of eq. (C.2) and the loop coordinate $h$, put $\Omega_*=\mathrm{J}_1\mathrm{B}_{11}^\top\mathrm{B}_{11}$ and

$$\mu_{\vartheta,j}^*(dw_1,dh)=\frac{1}{Z_{\vartheta,j}^*}\exp\left\{-\frac{\vartheta}{2\mathrm{T}}(w_1^\top\Omega_*w_1+\mathrm{J}_1|h|^2)-\phi_j(w_1,h)\right\}dw_1dh,\tag{C.49}$$

with loop marginal $\mu_{\vartheta,j}^{*,h}$, first-normal conditional $\mu_{\vartheta,j}^{*,1}(dw_1\mid h)$, and, with environment scales $\varrho_\epsilon$, $\varrho_\epsilon'$,

$$\beta_{jk}^{*,h}(h,\vartheta)=\int\beta_k(w_1,h,\vartheta)\mu_{\vartheta,j}^{*,1}(dw_1\mid h),\tag{C.50}$$

$$\bar{\beta}_{jk}(\vartheta)=\int\beta_k(w_1,h,\vartheta)\mu_{\vartheta,j}^*(dw_1,dh)=\int\beta_{jk}^{*,h}(h,\vartheta)\mu_{\vartheta,j}^{*,h}(dh),\tag{C.51}$$

$$\varrho_\epsilon=\frac{\mathrm{J}_2}{\mathrm{J}_1},\qquad\varrho_\epsilon'=\frac{\mathrm{J}_2}{\sqrt{\mathrm{J}_1\mathrm{T}}},\qquad\mathrm{e}_\epsilon(y)=\varrho_\epsilon+\varrho_\epsilon'|y|.\tag{C.52}$$

Then eq. (C.49) is the exact $(w_1, h)$-marginal of the optimizer conditional on $S = j$; averaging the inherited row $\beta^{\ell}_{jk}$ against the optimizer conditional on $j$ gives the same $\bar{\beta}_{jk}(\vartheta)$ for every $\ell$; and the exact parent rows and loop forces satisfy eqs. (C.56) to (C.58) below. [ESTABLISHED]

*Proof.* In the flag coordinates

$$w^{\top} H_{\mathcal{U}} w = \mathrm{J}_1 |\mathsf{B}_{11} w_1|^2 + \sum_{\ell=2}^{L} \mathrm{J}_{\ell} \left| \sum_{r \le \ell} \mathsf{B}_{\ell r} w_r \right|^2 , \tag{C.53}$$

and $\phi_j$ depends only on $(w_1, h)$. For fixed $w_1, h$, change the later $w_\ell$ to the residuals $r_\ell = \sum_{r \le \ell} \mathsf{B}_{\ell r} w_r$; the domain stays the whole Euclidean product and the Jacobian is the constant $\prod_{\ell \ge 2} |\det \mathsf{B}_{\ell\ell}|$, so integrating those Gaussian residuals contributes a scalar independent of $w_1, h, j$, which proves the first claim, the prior factor $\pi_j(\vartheta)$ canceling on conditioning. The change of variables and the tower property give the second.

For an intermediate $\ell \ge 1$, keep $y = (w_{\ell+1}, \dots, w_L)$ fixed and integrate $w_2, \dots, w_\ell$. The unperturbed conditional first-normal law $\gamma^{1}_{\ell,\vartheta,y}$ and its divergence from $\gamma^{1,*}_{\vartheta} = \mathcal{N}(0, (\mathrm{T}/\vartheta)\Omega_*^{-1})$ are

$$\gamma^{1}_{\ell,\vartheta,y} = \mathcal{N}\left( \mu_\ell(y), \frac{\mathrm{T}}{\vartheta} (\Omega_\ell^{(1)})^{-1} \right), \qquad \Omega_* \preceq \Omega_\ell^{(1)} \preceq \Omega_* + \mathrm{C}_L \mathrm{J}_2 I, \qquad |\mu_\ell(y)| \le \mathrm{C}_L \frac{\mathrm{J}_2}{\mathrm{J}_1} |y|, \tag{C.54}$$

$$\frac{1}{2} \left[ \frac{\vartheta}{\mathrm{T}} \mu_\ell^{\top} \Omega_* \mu_\ell + \mathrm{Tr}\big(\Omega_* (\Omega_\ell^{(1)})^{-1}\big) - \dim w_1 - \log\det\big(\Omega_* (\Omega_\ell^{(1)})^{-1}\big) \right] \le \mathrm{C}_L [\varrho_\epsilon^2 + \varrho_\epsilon'^2 |y|^2]. \tag{C.55}$$

For the lower precision bound, minimize the nonnegative later residual squares at $y = 0$; for the upper, set the eliminated later coordinates to zero and use $\mathrm{J}_\ell \le \mathrm{J}_2$ for $\ell \ge 2$; at $\ell = L$, $\Omega_L^{(1)} = \Omega_*$ and $\mu_L = 0$. The mean bound is the first-block-row estimate eq. (C.36), each future coordinate of level $\ell$ contributing at most $\mathrm{C}_L \mathrm{J}_\ell / \mathrm{J}_1$. The divergence bound holds because the whitened covariance eigenvalues lie between $(1 + \mathrm{C}_L \varrho_\epsilon)^{-1}$ and one, where $t - 1 - \log t = O((t-1)^2)$, and Pinsker's inequality turns it into a total variation at most $\mathrm{C}_L \mathrm{e}_\epsilon(y)$.

Tilting two probability measures by one weight $w \in [e^{-B_0}, 1]$ multiplies their total-variation distance by at most $2e^{B_0}$, since each term of $w\mu/\mu(w) - w\nu/\nu(w) = w(\mu - \nu)/\mu(w) + w\nu[1/\mu(w) - 1/\nu(w)]$ has total variation at most $e^{B_0} \|\mu - \nu\|_{\mathrm{TV}}$. Applied to $w = e^{-\phi_j(w_1,h)}$ uniformly in $h$, this gives the first bound below; the second follows through the common Gaussian force $\vartheta \mathrm{J}_1 h$ and $\mathrm{T}\mathbb{E}\nabla_h \phi_j$, and the third because $\gamma^{1,*}_{\vartheta}$ does not depend on $h$, so differentiation under its normalized tilt gives

$$\nabla_h \beta^{*,h}_{jk} = \mathbb{E}_{\mu^{*,1}_{\vartheta,j}(\cdot|h)} \nabla_h \beta_k - \mathrm{Cov}_{\mu^{*,1}_{\vartheta,j}(\cdot|h)} (\beta_k, \nabla_h \phi_j),$$

and both terms have the asserted Lipschitz bound: the standardized source potential has bounded first derivatives and $\sum_k \beta_k = 1$. Thus

$$\|\beta^{\ell}_{j\cdot}(y, h, \vartheta) - \beta^{*,h}_{j\cdot}(h, \vartheta)\|_{\mathrm{TV}} \le \mathrm{C}_L \mathrm{e}_\epsilon(y), \tag{C.56}$$

$$|\nabla_h V_{\ell,j}(y, h, \vartheta) - \nabla_h V^{*,h}_{j}(h, \vartheta)| \le \mathrm{C}_L \sqrt{\mathrm{T}\mathrm{J}_1}\, \mathrm{e}_\epsilon(y), \tag{C.57}$$

$$\mathrm{Lip}_{(w_1,h)} \beta_{j\cdot} + \mathrm{Lip}_h \beta^{*,h}_{j\cdot} \le \mathrm{C}_L \sqrt{\mathrm{J}_1/\mathrm{T}}. \tag{C.58}$$

□

**Lemma C.10 (Fast mixing with a learned coefficient).** There is a constant $\mathrm{c}_{\mathrm{fast},L} > 0$ with the following property. Let the auxiliary process have, on a proposal-free interval, fixed assignment $j$ and either the full fast potential of eq. (C.49) or its exact loop marginal, and a deterministic coefficient with $1 \le \vartheta \le 21/10$ and $|\dot\vartheta| \le \mathrm{C}_L\epsilon$. For a point start $\xi$ at time $t'$ and $t - t' \ge (\mathrm{c}_{\mathrm{fast},L}\mathrm{J}_1)^{-1}$,

$$\|P^*_{t',t}(\xi,\cdot) - \mu^*_{\vartheta(t),j}\|_{\mathrm{TV}} \le \mathrm{C}_L\big(1 + \sqrt{\mathrm{J}_1/\mathrm{T}}|\xi|\big)e^{-\mathrm{c}_{\mathrm{fast},L}\mathrm{J}_1(t-t')} + \mathrm{C}_L\epsilon/\mathrm{J}_1, \tag{C.59}$$

where $\mu^*$ means its loop marginal in the loop case. Also $\bar\beta_{jk}(\vartheta)$ is uniformly Lipschitz in $\vartheta$. [ESTABLISHED]

*Proof.* The relative Hessian bounds and the normalized marginal covariance calculation give $\mathrm{c}_L\mathrm{J}_1 I \preceq \nabla^2 V^* \preceq \mathrm{C}_L\mathrm{J}_1 I$ for both potentials, which are even, minimize at zero, and have equilibrium second moments at most $\mathrm{C}_L\mathrm{T}/\mathrm{J}_1$. For two point starts $\xi, \xi'$, reflection coupling before meeting gives their distance a one-dimensional noise of size $2\sqrt{2\mathrm{T}}$ and a drift at most $-\lambda$ times the distance, $\lambda = \mathrm{c}_L\mathrm{J}_1$; multiplying by $e^{\lambda(t-t')}$ gives a time-changed Brownian motion of variance $4\mathrm{T}(e^{2\lambda(t-t')} - 1)/\lambda$, and the reflection principle bounds the probability of not meeting by

$$\mathrm{C}|\xi - \xi'|\sqrt{\lambda/\mathrm{T}}\big(e^{2\lambda(t-t')} - 1\big)^{-1/2}. \tag{C.60}$$

The coupling reflects one Brownian motion orthogonally and predictably, so both marginals have the required law; the drifts are globally Lipschitz for each instance, so the copies are constructed away from the diagonal and joined at meeting, the comparison and hitting bound following by stopping before the diagonal and letting the stopping radius tend to zero. Averaging $\xi'$ against the equilibrium at time $t'$ gives the first term of eq. (C.59).

For the second, compare that evolving equilibrium start with the instantaneous equilibrium at time $t$. After standardizing $(w_1, h)$ by $\sqrt{\mathrm{J}_1/\mathrm{T}}$, the normalized parameter score is a centered quadratic with uniform exponential moments, and the loop-marginal score, its conditional expectation, has the same bound by Jensen's inequality. The moving-reference entropy calculation with the continuous entropy inequality at curvature $\mathrm{c}_L\mathrm{J}_1$ gives, for the relative entropy $\mathrm{D}_{\mathrm{lag}}$ of the evolved law to the current equilibrium,

$$\mathrm{D}'_{\mathrm{lag}} \le -\mathrm{c}_L\mathrm{J}_1\mathrm{D}_{\mathrm{lag}} + \mathrm{C}_L\epsilon^2/\mathrm{J}_1, \qquad \mathrm{D}_{\mathrm{lag}}(t') = 0, \tag{C.61}$$

the parameter term linear in $\mathrm{D}_{\mathrm{lag}}$ being absorbed for small $\epsilon$; hence $\mathrm{D}_{\mathrm{lag}} \le \mathrm{C}_L\epsilon^2/\mathrm{J}_1^2$, and Pinsker's inequality gives the second term. Finally, differentiating $\bar\beta_{jk} = \mathbb{E}_{\mu^*_{\vartheta,j}}\beta_k$ gives the explicit row derivative, whose absolute values sum to at most eight, and a covariance with the centered scale score, whose standardized first moment is bounded, so $\sum_k |\partial_\vartheta \bar\beta_{jk}| \le \mathrm{C}_L$. □

**Lemma C.11 (Coupling of complete source histories).** The source-assignment histories of a process and of an activated exact parent can be coupled, together with their common initial marked sample, so that the probability of any disagreement through $t_*$ is at most $\mathrm{C}_L\epsilon^6$. This holds through

the finite activation lineage, for the full finite-state path with arbitrary label correlations and the rate-one proposal clock. [ESTABLISHED]

*Proof.* First bound a proposal row on one proposal-free interval. For a microscopic process let $\xi = (w_1, h)$, and couple its auxiliary fast diffusion to the same initial fast point, Brownian increments, assignment, and parameter path. The omitted Gaussian environmental force $\vartheta[(H_{11} - \Omega_*)w_1 + H_{1,>1}w_{>1}]$ has norm at most $\mathrm{C}_L\mathrm{J}_2|w|$, and strong contraction from the most recent proposal or activation time $t_{\mathrm{prev}}$ gives the first bound below. For a parent at level $\ell \geq 1$, coupling its retained loop to the auxiliary loop process, eq. (C.57) gives the second:

$$|\xi_t - \xi_t^*| \leq \mathrm{C}_L\mathrm{J}_2 \int_{t_{\mathrm{prev}}}^{t} e^{-\mathrm{c}_L\mathrm{J}_1(t-\sigma)}|w_\sigma|d\sigma, \tag{C.62}$$

$$|h_t - h_t^*| \leq \mathrm{C}_L\sqrt{\mathrm{TJ}_1} \int_{t_{\mathrm{prev}}}^{t} e^{-\mathrm{c}_L\mathrm{J}_1(t-\sigma)}\mathrm{e}_\epsilon(y_\sigma)d\sigma, \tag{C.63}$$

with the adapted environment of the process. The row Lipschitz constant and the direct parent row error eq. (C.56) convert these into proposal-row errors.

At random proposal times, condition on the whole independent future Poisson grid, but not on its marks or on any future continuous path, so that the Brownian law is unchanged. Evenness and the uniform lower curvature on $M_\ell$ give the first bound below on every sheet and level; proposals preserve $y$, projections contract its norm, and the preparation has second moment at most $\mathrm{C}_L(\mathrm{A}^2 + \mathrm{T})$, which gives the second:

$$(\mathrm{T}\Delta - \nabla V_{\ell,j} \cdot \nabla)|y|^2 \leq 2\mathrm{T}\dim M_0 - 2\mathrm{c}_L|y|^2, \tag{C.64}$$

$$\sup_{t \leq t_*} \mathbb{E}\big[|\mathbf{X}_{M_\ell,t}|^2 \;\big|\; \text{Poisson grid}\big] \leq \mathrm{C}_L(\mathrm{A}^2 + \mathrm{T}). \tag{C.65}$$

If the initial sample is fixed as well, the right side first contains its squared norm, and integrating that sample returns the bound; the future grid after a deterministic activation is independent of the past, so the argument applies at each activation and at grid-selected times. Averaging the row errors of eqs. (C.62) and (C.63) with eq. (C.65), where the exponential convolution has integral at most $\mathrm{C}_L/\mathrm{J}_1$, over the distribution of every past prefix, and dividing by the probability of no rare prefix, gives the common expected per-proposal bound

$$\eta_\epsilon = \mathrm{C}_L\left[\frac{\mathrm{J}_2}{\mathrm{J}_1} + \frac{\mathrm{J}_2}{\sqrt{\mathrm{J}_1}}\left(1 + \frac{\mathrm{A}}{\sqrt{\mathrm{T}}}\right)\right] = O_L(\epsilon^7), \tag{C.66}$$

since $\mathrm{Q}_1 = 2\mathrm{Q}_2 + 20$ and $\mathrm{A}/\sqrt{\mathrm{T}} = \epsilon^{-3}$.

Define one comparison source process on the same proposal grid, started at the initial assignment and using the row $\bar{\beta}_{j\cdot}(\vartheta_\circ(t))$, $\vartheta_\circ(t) = 2 - e^{-\epsilon t}$, a single comparison law common to both processes while each process keeps its own law and coefficient. At a proposal time $t$, condition on the common starting sample, the full future grid, and all previous marks, self marks included; the assignment and level are fixed on $(t_{\mathrm{prev}}, t)$, and no future continuous trajectory is conditioned on. Starting the auxiliary diffusion from the fast point at $t_{\mathrm{prev}}$, the environmental coupling, then eq. (C.59), then the

Lipschitz comparison of $\vartheta(t)$ with $\vartheta_\circ(t)$ bound the averaged conditional row error, for gaps at least $t_{\mathrm{gap}}$, by the second line below:

$$t_{\mathrm{gap}} = \frac{k_L}{\mathrm{J}_1}\log(1/\epsilon), \tag{C.67}$$

$$\mathrm{C}_L\left[\left(1+\mathrm{A}\sqrt{\mathrm{J}_1/\mathrm{T}}\right)e^{-\mathrm{c}_{\mathrm{fast},L}\mathrm{J}_1 t_{\mathrm{gap}}} + \epsilon/\mathrm{J}_1 + \eta_\epsilon + r_\epsilon\right]; \tag{C.68}$$

the moment of the point-start mixing factor, which includes $\sqrt{\mathrm{J}_1/\mathrm{T}}(\mathrm{A}+\sqrt{\mathrm{T}})$, is absorbed because $\mathrm{A}/\sqrt{\mathrm{T}} \geq 1$, and averaging over the past marks returns the grid-conditioned moments exactly. Short gaps use the trivial bound one. For a rate-one Poisson grid the expected number of ordered pairs of proposal times less than $t_{\mathrm{gap}}$ apart is at most $t_* t_{\mathrm{gap}}$, by integrating the two-time intensity; proposals within $t_{\mathrm{gap}}$ after the initial time or one of the finitely many activations add at most $(L+1)t_{\mathrm{gap}}$, and the expected total count is $t_*$.

To pass from conditional row errors to the whole history, condition on a grid with $\mathrm{n}$ proposals, write the two joint mark laws as products of their conditional kernels $\mathsf{K}_m$ and $\overline{\mathsf{K}}_m$, and replace kernels in order, keeping the earlier prefix distributed according to the process and integrating the comparison tail, whose mass is one. With $\mathsf{j}_m \in \mathcal{J}$ the $m$th proposal mark,

$$\left\|\mathrm{Law}(\mathsf{j}_1,\dots,\mathsf{j}_{\mathrm{n}}) - \mathrm{Law}(\bar{\mathsf{j}}_1,\dots,\bar{\mathsf{j}}_{\mathrm{n}})\right\|_{\mathrm{TV}} \leq \sum_{m=1}^{\mathrm{n}} \mathbb{E}_{\mathsf{prefix}}\|\mathsf{K}_m(\cdot \mid \mathsf{j}_{<m}) - \overline{\mathsf{K}}_m(\cdot \mid \mathsf{j}_{<m})\|_{\mathrm{TV}}, \tag{C.69}$$

and integrating over the grid with the preceding bounds gives

$$\begin{aligned}\mathrm{Err}_{\mathrm{hist}} \leq \mathrm{C}_L\big[&(t_*+L+1)t_{\mathrm{gap}} + t_*\big(1+\mathrm{A}\sqrt{\mathrm{J}_1/\mathrm{T}}\big)e^{-\mathrm{c}_{\mathrm{fast},L}\mathrm{J}_1 t_{\mathrm{gap}}} \\ &+ t_*\epsilon/\mathrm{J}_1 + t_*\eta_\epsilon + t_* r_\epsilon\big].\end{aligned} \tag{C.70}$$

With $\mathrm{c}_{\mathrm{fast},L}k_L > \mathrm{Q}_1/2 + 12$, the five contributions are bounded by constants times $\epsilon^{19}\varsigma_\epsilon$, $\epsilon^8$, $\epsilon^{\mathrm{Q}_1}$, $\epsilon^6$, and $\epsilon^{2\mathrm{Q}_1+19} + \epsilon^{3\mathrm{Q}_1+25}\varsigma_\epsilon^2$, so $\mathrm{Err}_{\mathrm{hist}} \leq \mathrm{C}_L\epsilon^6$.

Fix the common initial sample and the grid, couple each mark sequence maximally with the comparison sequence, and glue the two couplings along it; the glued sequences disagree with probability at most $\mathrm{C}_L\epsilon^6$. The path spaces, continuous paths in a Euclidean space and right-continuous paths in $\mathcal{J}$, are standard Borel, so each continuous trajectory can be drawn from its regular conditional law given the initial sample, grid, and marks, which preserves both full-process marginals and the parent projection of the initial sample. Once self-proposals are discarded, each source path is a function of grid and marks, which proves the first bound of eq. (13.45); finitely many activations change only $\mathrm{C}_L$. □

## C.5 Coherence, source movement, and continuous paths

**Lemma C.12 (Successive coherence with the original conditional marks).** Under the preparation and regime of eqs. (C.13), (13.29) and (13.31), there are a finite $k_{\mathrm{geom},L}$ and a tolerance

bound $\eta_L^{\max} > 0$, a fixed fraction of $\mathrm{m}_L$, both depending only on $L$ and the fixed data, with the following property. For every $k_L > k_{\mathrm{geom},L}$ and every $0 < \eta_L \le \eta_L^{\max}$ there are $\epsilon_L, \mathrm{c}_L, \mathrm{C}_L > 0$ such that, for $0 < \epsilon < \epsilon_L$, conclusion (b) of Theorem 13.17 holds for the fine process and its successive parent descriptions on the windows $I_\ell = [t_\ell, 2t_\ell]$, $t_\ell = k_L \log(1/\epsilon)/\mathrm{J}_\ell$, with the probability bound eq. (13.40). On $I_\ell$ the discrepancies $\mathfrak{d}_i$ and the residuals of level at most $\ell$ are below $\eta_L \mathrm{A}$ and future bridges stay separated by a fixed positive multiple of $\mathrm{A}$, in the conditional root-mean-square original residuals of eq. (13.18) for eliminated coordinates, separately for the physical posterior and the restricted optimizer. [ESTABLISHED]

*Proof.* On an active level $\ell' < \ell$, write the non-root Gaussian comparison of Lemma C.3 through a current normal component $z \in \mathcal{U}_{\ell'} \cap \mathcal{U}_\ell^\perp$, a future component $y \in \mathcal{U}_\ell$, and the separate loop component. The weighted flag gives, for the current and future blocks of $H_{\ell'}$,

$$H_{zz} \succeq \mathrm{c}_L \mathrm{J}_\ell I, \qquad \|H_{zy}\| \le \mathrm{C}_L \mathrm{J}_{\ell+1}, \qquad \|H_{yy}\| \le \mathrm{C}_L \mathrm{J}_{\ell+1}, \qquad \mathrm{J}_{L+1} = 0. \tag{C.71}$$

Indeed rows of levels at most $\ell$ annihilate $\mathcal{U}_\ell$, so the original cross and future blocks contain only strengths at most $\mathrm{J}_{\ell+1}$; in an earlier elimination through $\ell''$, the eliminated-current block is $O_L(\mathrm{J}_{\ell''+1})$, the inverse hidden block $O_L(\mathrm{J}_{\ell''}^{-1})$, and the eliminated-future block $O_L(\mathrm{J}_{\ell+1})$, so the Schur correction is still $O_L(\mathrm{J}_{\ell+1})$; at level zero there is no correction, and loop/normal cross blocks vanish in the Gaussian core.

The deterministic reference mean $\mu^{\mathrm{G}}$ has norm at most $\mathrm{C}_L \mathrm{A}$ and zero loop component. Integrating $\dot{\mu}_z^{\mathrm{G}} = -\vartheta(t)(H_{zz}\mu_z^{\mathrm{G}} + H_{zy}\mu_y^{\mathrm{G}})$ with $1 \le \vartheta \le 21/10$ gives the first bound below; integrating the full future row, of norm $O_L(\mathrm{J}_{\ell+1})$, from the common center, earlier projections preserving $P_\ell \mu^{\mathrm{G}}$, gives the second for a sequential process begun at $t_{\ell-1}^{\mathrm{act}}$:

$$|\mu_z^{\mathrm{G}}(t)| \le \mathrm{C}_L \mathrm{A} \left[ e^{-\mathrm{c}_L \mathrm{J}_\ell (t - t_0)} + \frac{\mathrm{J}_{\ell+1}}{\mathrm{J}_\ell} \right], \tag{C.72}$$

$$|\mu^{\mathrm{G}}(t) - \mathrm{A} P_\ell \mathbf{v}_0| \le \mathrm{C}_L \mathrm{A} \left[ e^{-\mathrm{c}_L \mathrm{J}_\ell (t - t_{\ell-1}^{\mathrm{act}})} + \frac{\mathrm{J}_{\ell+1}}{\mathrm{J}_\ell} + \mathrm{J}_{\ell+1} t \right]. \tag{C.73}$$

The future terms are absent when $\ell = L$, and a fine process uses $t_0 = 0$.

For the fluctuations, on each active non-root space $\mathbf{Z} = \mathbf{X}^{\mathrm{G}} - \mu^{\mathrm{G}}$ obeys $d\mathbf{Z} = -\vartheta(t) H_{\ell'} \mathbf{Z} dt + \sqrt{2\mathrm{T}} dW$, with $\vartheta H_{\ell'} \succeq \mathrm{h}_L I$ and initial covariance at most $\mathrm{C}_L \mathrm{T} H_{\ell'}^{-1}$. Choose $\mathrm{a}_* > 0$ below $\mathrm{h}_L$ and small enough for the initial Gaussian exponential moment. For $\mathcal{Y}(\mathbf{z}) = e^{\upsilon}$ with $\upsilon = \mathrm{a}_* |\mathbf{z}|^2/(4\mathrm{T})$, direct differentiation gives the first line below; the initial expectation of $\mathcal{Y}$ is bounded, orthogonal retention decreases $\mathcal{Y}$, and Itô's formula up to the first exit from a ball, localization, and the nonnegative stopped value of $\mathcal{Y}$ yield the second for every $\varrho > 0$, higher stiffnesses only strengthening the negative term:

$$\mathscr{L}_t \mathcal{Y} = \left[ \frac{\mathrm{a}_* \mathrm{d}}{2} + \frac{\mathrm{a}_*^2 |\mathbf{z}|^2}{4\mathrm{T}} - \frac{\mathrm{a}_* \vartheta \mathbf{z}^\top H_{\ell'} \mathbf{z}}{2\mathrm{T}} \right] \mathcal{Y} \le \mathrm{a}_* (\mathrm{d}/2 - \upsilon) e^{\upsilon} \le \mathrm{C}_L, \tag{C.74}$$

$$\Pr\Big\{ \sup_{t \le t_*} |\mathbf{X}_t^{\mathrm{G}} - \mu^{\mathrm{G}}(t)| > \varrho \mathrm{A} \Big\} \le \mathrm{C}_L (1 + t_*) \exp\left[ -\mathrm{c}_L \varrho^2 \mathrm{A}^2 / \mathrm{T} \right]. \tag{C.75}$$

Let $\bar{r}_L$ bound the norms of the finitely many residual and discrepancy covectors, let $\eta_L^{\max}$ be a fixed fraction of $\mathrm{m}_L$, fix $\eta_L \in (0, \eta_L^{\max}]$, and choose a fluctuation radius $\varrho\mathrm{A}$ with $\bar{r}_L \varrho\mathrm{A}$ smaller than the reserved tolerance margin, $\mathrm{c}_L$ in the probability depending on $\varrho$. Take for example $k_{\mathrm{geom},L} = 2/\mathrm{c}_L$ for the restoring constant of eq. (C.72); increasing this finite constant is allowed. For small $\epsilon$, $t_{\ell-1}^{\mathrm{act}} < t_\ell/2$, so the exponential term of eq. (C.73) tends to zero uniformly on $I_\ell$, and

$$\frac{\mathrm{J}_{\ell+1}}{\mathrm{J}_\ell} \longrightarrow 0, \qquad \mathrm{J}_{\ell+1} t_\ell \longrightarrow 0 \quad (\ell < L), \qquad \frac{t_L^{\mathrm{act}}}{t_*} \longrightarrow 0. \tag{C.76}$$

The nonlinear correction of eq. (C.17) is at most $\mathrm{C}_L \mathrm{A} \epsilon^{5/2}$, including the admitted initial transport and finitely many inherited errors; it enters through the source force only via the uniform bound $B_* \sqrt{\mathrm{T}}$, so it holds for every adapted assignment path. Since $\mathrm{A}^2/\mathrm{T} = \epsilon^{-6}$, the strict prepared residual margins and eq. (C.75) give the asserted probability and separation on every window. The stable current-block equation keeps bounding formed residuals as their future forcing decreases, and after the last normal block there is no future forcing, so the final regime persists through $t_*$ on the same event, with every loop retained.

For an eliminated edge and a cumulative elimination through $\ell' \geq 1$, write a fine non-root state as $\mathbf{z} + y$ with $\mathbf{z} \in W_{\ell'}^{\perp}$ and $y \in M_{\ell'}$. The optimizing and physical Gaussian precisions are $\vartheta H$ and $\vartheta H - \mathrm{T} R_{\mathrm{ev}}$. Every row of $\mathfrak{D}$ lies in the first normal layer, so $R_{\mathrm{ev}}$ annihilates $\mathcal{U}_1$, $\mathcal{H}$, and $\mathcal{K}$, and with $\mathrm{T} R_{\mathrm{ev}} \preceq H/2$ both laws have hidden precision at least $\mathrm{c}_L \mathrm{J}_{\ell'} I$, hidden/future block $O_L(\mathrm{J}_{\ell'+1})$, and zero hidden/loop block. The Gaussian conditional mean therefore has norm at most $\mathrm{C}_L(\mathrm{J}_{\ell'+1}/\mathrm{J}_{\ell'})|y_{\mathcal{U}_{\ell'}}|$ and covariance trace at most $\mathrm{C}_L \mathrm{T}/\mathrm{J}_{\ell'}$, and, given the full assignment, each exact hidden density is a normalized tilt of its Gaussian conditional with log oscillation at most $B_0$, hence density ratio at most $e^{B_0}$. For $\mu_\vartheta \in \{\mathbf{\Pi}_{\vartheta,\mathrm{o}}, q_\vartheta^\star\}$ and every retained state with $|y_{\mathcal{U}_{\ell'}}| \leq \mathrm{C}_L \mathrm{A}$, this gives the first bound below, and for every original residual covector the triangle inequality and its reverse in conditional $L^2$ give the second:

$$\mathbb{E}_{\mu_\vartheta}[|\mathbf{z}|^2 \mid P_{\ell'}\mathbf{X} = y, S = j] \leq \mathrm{C}_L e^{B_0} \left[ \mathrm{A}^2 \left( \frac{\mathrm{J}_{\ell'+1}}{\mathrm{J}_{\ell'}} \right)^2 + \frac{\mathrm{T}}{\mathrm{J}_{\ell'}} \right], \tag{C.77}$$

$$\left| \sqrt{\mathbb{E}_{\mu_\vartheta}[r_e(\mathbf{X})^2 \mid P_{\ell'}\mathbf{X} = y, S = j]} - |r_e(y)| \right| \leq \|\mathbf{r}_e\| \sqrt{\mathbb{E}_{\mu_\vartheta}[|\mathbf{z}|^2 \mid P_{\ell'}\mathbf{X} = y, S = j]}. \tag{C.78}$$

On the event already constructed the right side of eq. (C.78) is $o(\mathrm{A})$, uniformly in the assignment and separately for the two laws, and the same holds for $\mathfrak{d}_i$; this transfers the small formed and the separated future residuals to eq. (13.18), with no erased mark replaced by its value at $P_{\ell'}\mathbf{x}$. The bridge levels join successive pairs of already formed components, giving the stated counts. □

**Lemma C.13 (Source envelope and changing hidden kernels).** Every process or finite replacement lineage of Lemma C.3 satisfies eq. (13.43) for the distinguished binary source, from any of its activation times. Between $\vartheta = 1$ and its final parameter, the first hidden physical conditional and the first hidden optimizing conditional each change by relative entropy greater than $1/200$, uniformly at the same retained input and assignment. [ESTABLISHED]

*Proof.* The distinguished receiver's posterior row $\beta_{i_0,\pm}$, a factor of the product row eq. (13.14), has the log odds of the first line below. Since $0 \le \chi \le 1$, it lies between $\operatorname{sig}(4(\vartheta - 3/2) \mp \delta)$, and summing the inherited row over the assignments with distinguished label $+$ and taking the expectation preserves these bounds, so the generator applied to the indicator of that label gives the second line, with $\lambda(t) = \Pr\{S_{i_0}(t) = +\}$:

$$\log \frac{\beta_{i_0,+}(\mathbf{x},\vartheta)}{\beta_{i_0,-}(\mathbf{x},\vartheta)} = 4(\vartheta - 3/2) - \delta \left[\chi\big(\sqrt{\mathrm{J}_1/\mathrm{T}}\, r_{i_0,+}(\mathbf{x})\big) - \chi\big(\sqrt{\mathrm{J}_1/\mathrm{T}}\, r_{i_0,-}(\mathbf{x})\big)\right], \tag{C.79}$$

$$\lambda' = \bar{\beta}_\lambda - \lambda, \qquad \operatorname{sig}(4(\vartheta - 3/2) - \delta) \le \bar{\beta}_\lambda(t) \le \operatorname{sig}(4(\vartheta - 3/2) + \delta). \tag{C.80}$$

Neither diffusion nor a marked projection changes this indicator, so variation of constants remains valid across every activation, even when the inherited row depends on the current assignment and continuous state.

By Lemma C.4, uniformly before $t_L^{\mathrm{act}} = 2k_L \log(1/\epsilon)$, $\vartheta \le 1 + \mathrm{C}_L \epsilon \log(1/\epsilon) + \mathrm{C}_L r_\epsilon$. A solution of eq. (C.80) cannot exceed the larger of its starting value and the supremum forcing, so the preparation tolerance and $\| \operatorname{sig}' \|_\infty \le 1/4$ give the first bound below. For small $\epsilon$, every activation precedes $t_* - 10$ and $\vartheta \ge 7/4$ throughout $[t_* - 10, t_*]$, so nonnegativity of the earlier contribution in variation of constants yields the second; for $\delta < 10^{-4}$, the elementary bounds $e^{1-\delta} > 249/100$, $e^2 > 7$, and $e^{10} > 2^{10}$ give the third:

$$\lambda(t_\ell^{\mathrm{act}}) \le \operatorname{sig}(-2) + 10^{-2} + \delta/4 + \mathrm{C}_L[\epsilon \log(1/\epsilon) + r_\epsilon], \qquad 0 \le \ell \le L, \tag{C.81}$$

$$\lambda(t_*) \ge (1 - e^{-10}) \operatorname{sig}(1 - \delta), \tag{C.82}$$

$$(1 - e^{-10}) \operatorname{sig}(1 - \delta) - \operatorname{sig}(-2) - 10^{-2} - \delta/4 > \frac{1023}{1024}\frac{249}{349} - \frac{1}{8} - \frac{1}{100} - \frac{1}{40000} > \frac{57}{100}. \tag{C.83}$$

The remaining error in eq. (C.81) tends to zero, hence $\lambda(t_*) - \lambda(t_\ell^{\mathrm{act}}) > 1/2$ for every activation, which for a binary row is its total-variation change.

For the hidden kernels, fix the retained input and the assignment, and first omit the bounded source factor. On the first hidden normal space let $\Omega_1, \Omega_{\mathrm{f}}$ be the initial and final Gaussian precisions, with covariances $\mathrm{T}\Omega_1^{-1}, \mathrm{T}\Omega_{\mathrm{f}}^{-1}$. For the optimizer the precision scales with $\vartheta$; for the physical law, on the non-root space,

$$\Omega_1 = H - \mathrm{T}R_{\mathrm{ev}} \succ 0, \qquad \Omega_{\mathrm{f}} = \Omega_1 + (\vartheta_{\mathrm{f}} - 1)H \succeq \vartheta_{\mathrm{f}} \Omega_1, \qquad \vartheta_{\mathrm{f}} \ge 7/4. \tag{C.84}$$

Restriction to a conditional block preserves the order, and so does an earlier Gaussian marginal, whose Schur quadratic is the minimum over the discarded coordinates, so in either case $\Omega_{\mathrm{f}} \succeq (5/4)\Omega_1$. For the Gaussian laws $\gamma_1, \gamma_{\mathrm{f}}$ on a hidden space of dimension $\mathrm{d_h} \ge 1$, whitening $\Omega_1$ and diagonalizing the conjugated final precision, with eigenvalues $\varrho_i \ge 5/4$, and completing the square in the Hellinger affinity give the first line below; each factor decreases for $\varrho_i \ge 1$, so the affinity is at most $[2\sqrt{5/4}/(1 + 5/4)]^{\mathrm{d_h}/2}$. Each exact conditional $\nu_1, \nu_{\mathrm{f}}$ is a normalized source tilt of its Gaussian conditional bounded above by $e^{B_0}$, also after exact hidden integration, so $\mathcal{A}(\nu_1, \nu_{\mathrm{f}}) \le e^{B_0} \mathcal{A}(\gamma_1, \gamma_{\mathrm{f}})$, and Jensen's inequality under $\nu_1$ gives $-D_{\mathrm{KL}}(\nu_1 \| \nu_{\mathrm{f}})/2 \le \log \mathcal{A}(\nu_1, \nu_{\mathrm{f}})$. With $\log(1 + 1/80) \ge 1/81$

this gives the second line, separately for the physical and optimizing conditionals and uniformly in the retained input and assignment:

$$\mathcal{A}(\gamma_1,\gamma_{\mathrm{f}}) = \prod_{i=1}^{\mathrm{d_h}} \left(\frac{2\sqrt{\varrho_i}}{1+\varrho_i}\right)^{1/2} e^{-\mathrm{M}}, \qquad \mathrm{M}\geq 0, \tag{C.85}$$

$$D_{\mathrm{KL}}(\nu_1\|\nu_{\mathrm{f}}) \geq \frac{\mathrm{d_h}}{2}\log\frac{81}{80} - 2B_0 \geq \frac{1}{162} - \frac{1}{5000} > \frac{1}{200}. \tag{C.86}$$

□

**Lemma C.14 (Concentration of continuous paths on a common deterministic path).** Let a source process and its exact parent start from a common retained sample and parameter at $t_\ell^{\mathrm{act}}$. There is a deterministic path $\bar{\mathbf{x}}_\ell : [t_\ell^{\mathrm{act}}, t_*] \to W_\ell$, namely the parent's Gaussian reference mean on the non-root coordinates and the constant equilibrium root mean, such that

$$\Pr\Big\{\sup_{t_\ell^{\mathrm{act}}\leq t\leq t_*} |P_\ell \mathbf{X}_t - \bar{\mathbf{x}}_\ell(t)| > \mathrm{C}_L\mathrm{A}\epsilon^2\Big\} + \Pr\Big\{\sup_{t_\ell^{\mathrm{act}}\leq t\leq t_*} |\overline{\mathbf{X}}_t - \bar{\mathbf{x}}_\ell(t)| > \mathrm{C}_L\mathrm{A}\epsilon^2\Big\} \leq \mathrm{C}_L(1+t_*)e^{-\mathrm{c}_L/\epsilon^2}. \tag{C.87}$$

Consequently the displacement bound of eq. (13.45) holds under every coupling of the two processes, in particular under the coupling of Lemma C.11. A finite number of successive replacements changes only the constants. [ESTABLISHED]

*Proof.* First compare Gaussian means at one deterministic parameter path. In the split of eq. (C.71) the inherited normal precision is $H_\ell = H_{yy} - H_{zy}^\top H_{zz}^{-1} H_{zy}$. With $\mathfrak{w} = \mu_z^{\mathrm{G}} + H_{zz}^{-1}H_{zy}\mu_y^{\mathrm{G}}$, the future mean of the source process satisfies the first line below, the bound following from eq. (C.72) and $\|H_{zz}^{-1}H_{zy}\| \leq \mathrm{C}_L\mathrm{J}_{\ell+1}/\mathrm{J}_\ell$, and integrating the contractive future propagator from a matching retained mean gives the second:

$$\dot{\mu}_y^{\mathrm{G}} = -\vartheta H_\ell \mu_y^{\mathrm{G}} - \vartheta H_{zy}^\top \mathfrak{w}, \qquad |\mathfrak{w}(t)| \leq \mathrm{C}_L\mathrm{A}\left[e^{-\mathrm{c}_L\mathrm{J}_\ell(t-t_\ell^{\mathrm{act}})} + \frac{\mathrm{J}_{\ell+1}}{\mathrm{J}_\ell}\right], \tag{C.88}$$

$$\sup_{t_\ell^{\mathrm{act}}\leq t\leq t_*} |\mu_y^{\mathrm{G}}(t) - \bar{\mu}^{\mathrm{G}}(t)| \leq \mathrm{C}_L\mathrm{A}\left[\frac{\mathrm{J}_{\ell+1}}{\mathrm{J}_\ell} + t_*\frac{\mathrm{J}_{\ell+1}^2}{\mathrm{J}_\ell}\right]. \tag{C.89}$$

There is no future normal contribution when $\ell = L$, and the loop Gaussian core is unchanged by the reduction and has zero mean. If the independently learned parameters differ, with $\mathrm{e}_\vartheta = \sup|\vartheta - \bar{\vartheta}| \leq \mathrm{C}_L r_\epsilon$ by Lemma C.4, Young's inequality in the $H_\ell$ quadratic form gives, for homogeneous parent means $\mu, \bar{\mu}$ with difference $\mathfrak{e}$,

$$\frac{d}{dt}|\mathfrak{e}|^2 \leq -\vartheta \mathfrak{e}^\top H_\ell \mathfrak{e} + \mathrm{C}_L \mathrm{e}_\vartheta^2 \bar{\vartheta}\bar{\mu}^\top H_\ell \bar{\mu}, \tag{C.90}$$

and since $\frac{d}{dt}|\bar{\mu}|^2 = -2\bar{\vartheta}\bar{\mu}^\top H_\ell \bar{\mu}$, the time integral of the final energy is at most $\mathrm{C}_L\mathrm{A}^2$, so the extra displacement is at most $\mathrm{C}_L\mathrm{A}\mathrm{e}_\vartheta$.

Attach to each process its own Gaussian comparison from Lemma C.3; the discrepancy between the parent's projected starting reference and its full retained starting sample is included in eq. (C.17).

On the non-root coordinates let $\bar{\mathbf{x}}_\ell = \bar{\mu}^{\mathrm{G}}$ be the parent's reference mean. For the parent, eqs. (C.17) and (C.75) with $\varrho = \epsilon^2$ give $\sup_t |\overline{\mathbf{X}}_t - \bar{\mathbf{x}}_\ell(t)| \le \mathrm{C}_L\mathrm{A}(\epsilon^2 + \epsilon^{5/2})$ outside a set of probability $\mathrm{C}_L(1 + t_*)e^{-\mathrm{c}_L/\epsilon^2}$. For the source process the same two estimates for its own reference, then eq. (C.89) and the parameter comparison, give, outside a set of the same form,

$$\sup_{t_\ell^{\mathrm{act}} \le t \le t_*} |P_\ell \mathbf{X}_t - \bar{\mathbf{x}}_\ell(t)| \le \mathrm{C}_L\mathrm{A}\left[\epsilon^2 + \epsilon^{5/2} + r_\epsilon + \frac{\mathrm{J}_{\ell+1}}{\mathrm{J}_\ell} + t_* \frac{\mathrm{J}_{\ell+1}^2}{\mathrm{J}_\ell}\right]. \tag{C.91}$$

For $\ell < L$ the last term equals $2\epsilon^{19}$, $\mathrm{J}_{\ell+1}/\mathrm{J}_\ell = \epsilon^{\mathrm{Q}_{\ell+1}+20}$, and $r_\epsilon = o(\epsilon^2)$; both future terms vanish at $\ell = L$. The root coordinates of each process form the same autonomous Ornstein–Uhlenbeck process in law, with stiffness $1 + \kappa \ge 1$, started at its equilibrium and independent of the source history, and the argument of eq. (C.75) keeps each within $\epsilon^2\mathrm{A}$ of the constant mean $\kappa z_\star/(1+\kappa)$ outside a set of the same form. Taking that constant as the root component of $\bar{\mathbf{x}}_\ell$ proves eq. (C.87). Each of its events concerns the law of one process alone, so for any coupling the displacement event of eq. (13.45), after doubling $\mathrm{C}_L$, lies in their union; a finite sequence of replacements uses a finite triangle inequality and union bound. □

## C.6 The Gaussian core across an early interface

*Proof of parts (b) and (c) of Proposition 13.15.* Fix $1 \le \ell \le L-1$. In this section $\mathrm{C}$ and $\mathrm{c}$ denote positive constants depending only on $L$ and the carrier, never on the strengths, $\mathrm{A}$, $\mathrm{T}$, $t$, or $\epsilon$; put $\mathrm{J}_{L+1} = 0$ and $\bar{\varrho}_\ell = \varrho_{\ell-1} + \varrho_\ell + \varrho_{\ell+1}$.

*Flag layers and block scales.* In the flag coordinates $w = (w_1, \dots, w_L)$ of eq. (C.2), $\mathbf{x}_{\mathcal{U}} = \sum_s \mathsf{Q}_s w_s$, the quadratic form $\mathbf{x}_{\mathcal{U}}^\top H \mathbf{x}_{\mathcal{U}} = \sum_s \mathrm{J}_s |\sum_{r \le s} \mathsf{B}_{sr} w_r|^2$ has the blocks $(H_{\mathcal{U}})_{rr'} = \sum_{s \ge \max(r,r')} \mathrm{J}_s \mathsf{B}_{sr}^\top \mathsf{B}_{sr'}$, each of norm at most $\mathrm{C}\mathrm{J}_{\max(r,r')}$; the loop block $\mathrm{J}_1 I_{\mathcal{H}}$ is orthogonal to every layer and plays no part. On the active level $\ell - 1$ the normal part of $H_{\ell-1}$ is the Schur complement of $H_{\mathcal{U}}$ onto the layers $r \ge \ell$, obtained by eliminating the layers $\mathcal{E} = \{1, \dots, \ell-1\}$, with the discarded block $z = w_\ell$ and the retained block $y = (w_{\ell+1}, \dots, w_L)$. The lower bound of eq. (C.2) on $\mathcal{E}$ gives $H_{\mathcal{E}\mathcal{E}} \succeq \mathrm{c}\mathrm{J}_{\ell-1} I$, so the Schur correction to a block $(r, r')$ with $r, r' \ge \ell$ has norm at most $\mathrm{C}\mathrm{J}_r\mathrm{J}_{r'}/\mathrm{J}_{\ell-1} \le \mathrm{C}\mathrm{J}_{\max(r,r')}\varrho_{\ell-1}$. Since in each block the summand with the largest admissible strength leads and the remainder carries the next strength,

$$\begin{aligned} H_{zz} &= \mathrm{J}_\ell\big[\mathsf{B}_{\ell\ell}^\top \mathsf{B}_{\ell\ell} + O(\varrho_\ell + \varrho_{\ell-1})\big], \\ H_{z,\ell+1} &= \mathrm{J}_{\ell+1}\big[\mathsf{B}_{\ell+1,\ell}^\top \mathsf{B}_{\ell+1,\ell+1} + O(\varrho_{\ell+1} + \varrho_{\ell-1})\big], \\ H_{\ell+1,\ell+1} &= \mathrm{J}_{\ell+1}\big[\mathsf{B}_{\ell+1,\ell+1}^\top \mathsf{B}_{\ell+1,\ell+1} + O(\varrho_{\ell+1} + \varrho_{\ell-1})\big], \end{aligned} \tag{C.92}$$

with the constants of the remainders depending only on $L$ and the carrier, together with $\|H_{z,r}\| \le \mathrm{C}\mathrm{J}_r$ and $\|H_{\ell+1,r}\| \le \mathrm{C}\mathrm{J}_{\ell+1}(\varrho_{\ell+1} + \varrho_{\ell-1})$ for $r \ge \ell + 2$, and $\|H_{r,\cdot}\| \le \mathrm{C}\mathrm{J}_r$ for every row block $r \ge \ell + 1$, in particular $\|H_{zy}\| \le \mathrm{C}\mathrm{J}_{\ell+1}$. Hence, for $\bar{\varrho}_\ell \le \mathrm{c}_0$ with $\mathrm{c}_0$ small, $\lambda_{\min}(H_{zz}) \ge \frac{3}{4}\mathrm{j}_\ell\mathrm{J}_\ell$, $\|H_{zz}^{-1}\| \le \mathrm{C}/\mathrm{J}_\ell$, and $\|H_{zz}^{-1}H_{zy}\| \le \mathrm{C}\varrho_\ell$. The same computation at every earlier active level $p \le \ell - 1$ bounds the rows of layer $\ell + 1$ of $H_p$ by $\mathrm{C}\mathrm{J}_{\ell+1}$.

*Full row rank.* Write $\mathbf{r}_e$ for the covector of $r_e$. Since $\mathcal{U}_\ell = \mathcal{U}_{\ell-1} \cap \bigcap_e \ker \mathbf{r}_e$ over the rows $e$ of $B_\ell$, the layer spanned by $\mathsf{Q}_\ell$ is the span of the projections $P_{\mathcal{U}_{\ell-1}}\mathbf{r}_e$, and $\mathsf{B}_{\ell+1,\ell}$ has full row rank exactly when the matrix $\mathrm{M}_{fe} = \mathbf{r}_f \cdot P_{\mathcal{U}_{\ell-1}}\mathbf{r}_e$, with $f$ over the rows of $B_{\ell+1}$ and $e$ over those of $B_\ell$, has full row rank, because $\mathrm{M}_{fe} = (P_{\mathcal{U}_{\ell-1}}\mathbf{r}_f)\cdot(P_{\mathcal{U}_{\ell-1}}\mathbf{r}_e)$, the component of $P_{\mathcal{U}_{\ell-1}}\mathbf{r}_f$ in $\mathcal{U}_\ell$ pairs to zero with the layer, and the second factors span the layer. For $\ell = 1$, $\mathcal{U}_0 = \mathcal{U}$ contains every row of $B$, so $\mathrm{M}_{fe} = \mathbf{r}_f \cdot \mathbf{r}_e$; in the column of the discrepancy row of agent $i$ the entry of a bridge $f$ is $[i = \mathrm{head}(f)] - [i = \mathrm{tail}(f)]$, and the level-two bridges, those of odd index, join disjoint pairs of triangles and have pairwise disjoint endpoint sets, so the columns of the discrepancy rows of their heads form an identity submatrix. For $\ell \geq 2$ let $S$ be the set of bridges of level at least $\ell$, let $\boldsymbol{\partial}_f \in \mathbb{R}^V$ be the vector with entry $1$ at the head of the bridge $f$, $-1$ at its tail, and $0$ elsewhere, let $\partial_S$ have the rows $\boldsymbol{\partial}_f^\top$, $f \in S$, and let $\partial$ be the oriented incidence operator of the whole graph, so that $\partial^\top\partial$ is its Laplacian. For a vector $\mathsf{b}$ on the bridges put $\mathbf{x}_{\mathrm{br}}(\mathsf{b}) = (0, 0, -\mathsf{b})$, the configuration with only the bridge links, and $\mathbf{x}(\mathsf{b}) = P_{\mathcal{K}^\perp}\mathbf{x}_{\mathrm{br}}(\mathsf{b})$. A configuration of $\mathcal{U}_1$ has $\bar{k} = \bar{m} =: t$ and, in each triangle, three equal residuals which sum to $-\mathfrak{h}_q = 0$ by orthogonality to $\mathcal{H}$, so its internal links are $\partial t$ and its bridge links are $(\partial t)_f - \mathsf{b}_f$ with $\mathsf{b}_f = r_f$; it is therefore $(t, t, \partial t) + \mathbf{x}_{\mathrm{br}}(\mathsf{b})$ and, being orthogonal to $\mathcal{K}$, equals $\mathbf{x}(\mathsf{b})$. Conversely $\mathbf{x}_{\mathrm{br}}(\mathsf{b}) \in W_1 \cap \mathcal{H}^\perp$ and $\mathcal{K} \subseteq W_1 \cap \mathcal{H}^\perp$, so $\mathbf{x}(\mathsf{b}) \in \mathcal{U}_1$, and $r_f(\mathbf{x}(\mathsf{b})) = r_f(\mathbf{x}_{\mathrm{br}}(\mathsf{b})) = \mathsf{b}_f$ because residuals annihilate $\mathcal{K}$. Hence $\mathcal{U}_{\ell-1} = \{\mathbf{x}(\mathsf{b}) : \operatorname{supp}\mathsf{b} \subseteq S\}$ with $\mathsf{b} \mapsto \mathbf{x}(\mathsf{b})$ injective, and from $\langle (t,t,\partial t), (t',t',\partial t')\rangle = t^\top(2I + \partial^\top\partial)t'$ and $\langle (t,t,\partial t), \mathbf{x}_{\mathrm{br}}(\mathsf{b})\rangle = -t^\top\partial_S^\top\mathsf{b}$ the Gram matrix is $\langle \mathbf{x}(\mathsf{b}), \mathbf{x}(\mathsf{b}')\rangle = \mathsf{b}^\top \mathrm{G}\mathsf{b}'$ with $\mathrm{G} = I - \partial_S(2I + \partial^\top\partial)^{-1}\partial_S^\top$ on $\mathbb{R}^S$. The functional $\mathsf{b} \mapsto \mathsf{b}_f$, $f \in S$, has the Riesz representative $\mathbf{x}(\mathsf{b}^{(f)})$ with $\mathsf{b}^{(f)}$ the $f$th column of $\mathrm{G}^{-1}$, so $P_{\mathcal{U}_{\ell-1}}\mathbf{r}_f = \mathbf{x}(\mathsf{b}^{(f)})$ and $\mathrm{M}_{fe} = (\mathrm{G}^{-1})_{fe}$. Since $2I + \partial^\top\partial - \partial_S^\top\partial_S = 2I + \partial_{\setminus S}^\top\partial_{\setminus S}$, with $\partial_{\setminus S}$ the incidence operator of the graph with the bridges of $S$ removed, is positive definite, the Woodbury identity gives $\mathrm{M}_{fe} = \delta_{fe} + \boldsymbol{\partial}_f^\top(2I + \partial_{\setminus S}^\top\partial_{\setminus S})^{-1}\boldsymbol{\partial}_e$. The graph with the bridges of $S$ removed is the disjoint union of $2^{L-\ell+1}$ connected blocks of $2^{\ell-2}$ consecutive triangles, and the inverse is block diagonal with entrywise strictly positive blocks: on a block $B$ with Laplacian $\partial_B^\top\partial_B$ and maximal degree $\mathrm{D}_B$, the matrix $(\mathrm{D}_B + 2)I - (2I + \partial_B^\top\partial_B)$ is entrywise nonnegative, irreducible, and of spectral radius at most $\mathrm{D}_B < \mathrm{D}_B + 2$, since the Laplacian's eigenvalues lie in $[0, 2\mathrm{D}_B]$, so the Neumann series of $(2I + \partial_B^\top\partial_B)^{-1}$ in it converges and is entrywise positive. The bridges of $S$ are those of index $j = \mathrm{w}2^{\ell-2}$, $1 \leq \mathrm{w} \leq 2^{L-\ell+1} - 1$, of level $\ell + v_2(\mathrm{w})$, the bridge at position $\mathrm{w}$ joining block $\mathrm{w}$ to block $\mathrm{w}+1$; level $\ell$ means $\mathrm{w}$ odd and level $\ell+1$ means $\mathrm{w} \equiv 2 \pmod 4$. For $f \neq e$ in $S$, with positions $\mathrm{w}_f, \mathrm{w}_e$, the entry $\boldsymbol{\partial}_f^\top(2I + \partial_{\setminus S}^\top\partial_{\setminus S})^{-1}\boldsymbol{\partial}_e$ is a sum over the blocks both bridges touch, so it vanishes unless $|\mathrm{w}_f - \mathrm{w}_e| = 1$, and then it is plus or minus one entry of the positive inverse on the shared block, hence nonzero. Pair each level-$(\ell+1)$ bridge $f$ with the level-$\ell$ bridge $e(f)$ at position $\mathrm{w}_f + 1$, which exists because $4$ divides $2^{L-\ell+1}$; then $\mathrm{M}_{f,e(f')} \neq 0$ exactly when $|\mathrm{w}_f - \mathrm{w}_{f'} - 1| = 1$, that is $\mathrm{w}_f - \mathrm{w}_{f'} \in \{0, 2\}$, and as $\mathrm{w}_f - \mathrm{w}_{f'} \in 4\mathbb{Z}$ exactly when $f = f'$. The square submatrix $(\mathrm{M}_{f,e(f')})_{f,f'}$ is diagonal with nonzero entries, so the rows of $\mathrm{M}$ of level $\ell+1$ are independent and $\mathsf{B}_{\ell+1,\ell}$ has full row rank; as $\mathsf{B}_{\ell+1,\ell+1}\mathsf{B}_{\ell+1,\ell+1}^\top$ is invertible and $\mathsf{B}_{\ell+1,\ell}^\top$ is injective, $\mathfrak{Y}_\ell$ is injective. Nothing used the choice of bridge endpoints.

*Slaving on the active interval.* Let level $\ell - 1$ be active on $[t_0, t_\ell^{\mathrm{act}}]$ with $t_0 = t_{\ell-1}^{\mathrm{act}}$, with mean $\bar{\mathbf{x}}(t)$ in the layers $r \geq \ell$, $\dot{\bar{\mathbf{x}}} = -H_{\ell-1}\bar{\mathbf{x}}$, and $|\bar{\mathbf{x}}(t_0)| \leq \mathrm{CA}$, which holds because $|\bar{\mathbf{x}}(0)| = \mathrm{A}|\mathbf{v}_0|$, the flow contracts $|\bar{\mathbf{x}}|$, and the projections do not increase it. Write $\mathbf{g} = H_{\ell-1}\bar{\mathbf{x}}$, so that $\mathfrak{v} = \mathbf{g}_z$

and $\mathrm{T}\Delta_{\ell-1:\ell} = \frac{1}{2}\mathbf{g}_z^\top H_{zz}^{-1}\mathbf{g}_z$ by eq. (13.34), and $\mathrm{b}_{\ell+1}(t) = B_{\ell+1}P_\ell\bar{\mathbf{x}}(t) = \mathrm{B}_{\ell+1,\ell+1}\bar{\mathbf{x}}_{\ell+1}(t)$, since $P_\ell$ removes the layer $\ell$ and $\mathrm{B}_{\ell+1,r} = 0$ for $r > \ell+1$. The row bound gives $|\dot{\bar{\mathbf{x}}}_{\ell+1}| = |\mathbf{g}_{\ell+1}| \le \mathrm{C}\mathrm{J}_{\ell+1}\mathrm{A}$, and the same bound at the earlier levels, whose projections $P_p$, $p \le \ell$, leave the layer $\ell+1$ unchanged, gives the second line of eq. (13.36). Since $\dot{\mathbf{g}} = -H_{\ell-1}\mathbf{g}$, the variable $\mathbf{u} = \mathbf{g}_z + H_{zz}^{-1}H_{zy}\mathbf{g}_y$ satisfies $\dot{\mathbf{u}} = -H_{zz}\mathbf{u} - H_{zz}^{-1}H_{zy}(H_{\ell-1}\mathbf{g})_y$, where $|(H_{\ell-1}\mathbf{g})_y| \le \mathrm{C}\mathrm{J}_{\ell+1}(|\mathbf{g}_z| + |\mathbf{g}_y|)$, $|\mathbf{g}_y| \le \mathrm{C}\mathrm{J}_{\ell+1}\mathrm{A}$, and $|\mathbf{g}_z| \le |\mathbf{u}| + \mathrm{C}\varrho_\ell\mathrm{J}_{\ell+1}\mathrm{A}$. Variation of constants with $\|e^{-H_{zz}s}\| \le e^{-\lambda s}$, $\lambda = \lambda_{\min}(H_{zz})$, gives

$$|\mathbf{u}(t)| \le e^{-\lambda(t-t_0)}|\mathbf{u}(t_0)| + \int_{t_0}^{t} e^{-\lambda(t-s)}\big(\mathrm{C}\varrho_\ell^2\mathrm{J}_\ell|\mathbf{u}(s)| + \mathrm{C}\varrho_\ell\mathrm{J}_{\ell+1}^2\mathrm{A}\big)ds, \tag{C.93}$$

and iterating the positive Volterra operator of the right side compares $|\mathbf{u}|$ with the solution of the scalar equation with decay rate $\lambda - \mathrm{C}\varrho_\ell^2\mathrm{J}_\ell$, constant forcing $\mathrm{C}\varrho_\ell\mathrm{J}_{\ell+1}^2\mathrm{A}$, and initial value $|\mathbf{u}(t_0)|$, so that $|\mathbf{u}(t)| \le e^{-(\lambda-\mathrm{C}\varrho_\ell^2\mathrm{J}_\ell)(t-t_0)}|\mathbf{u}(t_0)| + \mathrm{C}\varrho_\ell\mathrm{J}_{\ell+1}^2\mathrm{A}/(\lambda - \mathrm{C}\varrho_\ell^2\mathrm{J}_\ell)$. For $\bar{\varrho}_\ell \le \mathrm{c}_0$, $\lambda - \mathrm{C}\varrho_\ell^2\mathrm{J}_\ell \ge \frac{1}{2}\mathrm{j}_\ell\mathrm{J}_\ell$, and $|\mathbf{u}(t_0)| \le \mathrm{C}\mathrm{J}_\ell\mathrm{A}$, so $|\mathbf{u}(t)| \le \mathrm{C}(\mathrm{J}_{\ell+1}^2\mathrm{A}/\mathrm{J}_\ell)[\varrho_\ell + \mathfrak{x}_\ell(t)]$. For the forcing, $H_{zy}\mathbf{g}_y = H_{z,\ell+1}\mathbf{g}_{\ell+1} + \sum_{r\ge\ell+2} H_{z,r}\mathbf{g}_r$, the sum being at most $\mathrm{C}\mathrm{J}_{\ell+2}^2\mathrm{A} \le \mathrm{C}\mathrm{J}_{\ell+1}^2\varrho_{\ell+1}^2\mathrm{A}$; from $\mathbf{g}_z = H_{zz}\bar{\mathbf{x}}_z + H_{zy}\bar{\mathbf{x}}_y$, $|\bar{\mathbf{x}}_z| \le \mathrm{C}\mathrm{A}(\varrho_\ell + \mathfrak{x}_\ell)$, and eq. (C.92) then gives $\mathbf{g}_{\ell+1} = \mathrm{J}_{\ell+1}[\mathrm{B}_{\ell+1,\ell+1}^\top\mathrm{b}_{\ell+1} + O(\mathrm{A}(\bar{\varrho}_\ell + \mathfrak{x}_\ell))]$ and $H_{zy}\mathbf{g}_y = \mathrm{J}_{\ell+1}^2[\mathfrak{Y}_\ell\mathrm{b}_{\ell+1} + O(\mathrm{A}(\bar{\varrho}_\ell + \mathfrak{x}_\ell))]$. Writing $\mathbf{g}_z = -H_{zz}^{-1}H_{zy}\mathbf{g}_y + \mathbf{u}$ with $H_{zz}^{-1} = \mathrm{J}_\ell^{-1}[(\mathrm{B}_{\ell\ell}^\top\mathrm{B}_{\ell\ell})^{-1} + O(\bar{\varrho}_\ell)]$ and inserting in $\frac{1}{2}\mathbf{g}_z^\top H_{zz}^{-1}\mathbf{g}_z$, with $|\mathrm{b}_{\ell+1}| \le \mathrm{C}\mathrm{A}$, gives $\mathrm{T}\Delta_{\ell-1:\ell} = (\mathrm{A}^2\mathrm{J}_{\ell+1}^4/\mathrm{J}_\ell^3)[\frac{1}{2}|(\mathrm{B}_{\ell\ell}^\top\mathrm{B}_{\ell\ell})^{-3/2}\mathfrak{Y}_\ell\mathrm{b}_{\ell+1}/\mathrm{A}|^2 + O(\bar{\varrho}_\ell + \mathfrak{x}_\ell + \mathfrak{x}_\ell^2)]$, which is the first line of eq. (13.36) by the definition of $\mathfrak{y}_\ell$.

*Windows.* On $I_\ell$, $\mathrm{J}_\ell(t-t_0) \ge \mathrm{J}_\ell(t_\ell - 2t_{\ell-1}) = k_L\log(1/\epsilon)(1-2\varrho_{\ell-1})$, so $\mathfrak{x}_\ell \to 0$ uniformly on $I_\ell$ under the stated condition, and $\mathrm{J}_{\ell+1}t \le 2k_L\log(1/\epsilon)\varrho_\ell \to 0$ gives $\mathrm{b}_{\ell+1}(t)/\mathrm{A} = B_{\ell+1}P_\ell\mathbf{v}_0 + o(1)$, so the quadratic $\mathfrak{y}_\ell$ satisfies $\mathfrak{y}_\ell(\bar{\mathbf{x}}(t)/\mathrm{A}) = \mathfrak{y}_\ell(\mathbf{v}_0) + o(1)$. For the lower bound, $|(\mathrm{B}_{\ell\ell}^\top\mathrm{B}_{\ell\ell})^{-3/2}\mathfrak{Y}_\ell\mathrm{b}_{\ell+1}| \ge \|\mathrm{B}_{\ell\ell}\|^{-3}\sigma_{\min}(\mathfrak{Y}_\ell)|\mathrm{b}_{\ell+1}|$, and $|B_{\ell+1}P_\ell\mathbf{v}_0|^2 = \sum_e r_e(P_\ell\mathbf{v}_0)^2$ over the $N_\triangle/2^\ell$ bridges $e$ of level $\ell+1$ is at least $N_\triangle\mathrm{m}_L^2/2^{\ell+2}$ by the margin of Lemma C.2, which gives eq. (13.37).

*Power laws.* With $\mathrm{J}_s = \epsilon^{-\mathrm{Q}_s}$, $\varrho_s = \epsilon^{\mathrm{Q}_s - \mathrm{Q}_{s+1}} \to 0$, $k_L\log(1/\epsilon)\varrho_\ell \to 0$, and on $I_\ell$

$$\mathfrak{x}_\ell \le \epsilon^{\mathrm{j}_\ell k_L(1-2\varrho_{\ell-1})/2 - 2(\mathrm{Q}_\ell - \mathrm{Q}_{\ell+1})},$$

which tends to zero exactly under the stated threshold on $k_L$; the window statement then places $\Delta_{\ell-1:\ell}$ between $\frac{1}{2}\mathfrak{y}_\ell(\mathbf{v}_0)$ and $2\mathfrak{y}_\ell(\mathbf{v}_0)$ times $(\mathrm{A}^2/\mathrm{T})\mathrm{J}_{\ell+1}^4/\mathrm{J}_\ell^3 = \epsilon^{3\mathrm{Q}_\ell - 4\mathrm{Q}_{\ell+1} - \mathrm{Q}_\mathrm{T}}$ on $I_\ell$ for small $\epsilon$. It is $O(\epsilon^{\mathrm{Q}_\triangle})$ on every early window exactly when $3\mathrm{Q}_\ell - 4\mathrm{Q}_{\ell+1} - \mathrm{Q}_\mathrm{T} \ge \mathrm{Q}_\triangle$ for every $\ell \le L-1$, that is $\mathrm{Q}_\ell + \mathrm{Q}_\mathrm{T} + \mathrm{Q}_\triangle \ge \frac{4}{3}(\mathrm{Q}_{\ell+1} + \mathrm{Q}_\mathrm{T} + \mathrm{Q}_\triangle)$, and iterating from $\mathrm{Q}_L = 0$ gives $\mathrm{Q}_1 + \mathrm{Q}_\mathrm{T} + \mathrm{Q}_\triangle \ge (4/3)^{L-1}(\mathrm{Q}_\mathrm{T} + \mathrm{Q}_\triangle)$, with equality along the recursion; vanishing is the case $\mathrm{Q}_\triangle = 0$ with strict inequality. □

## C.7 The endogenous four-cycle

*Proof of Proposition 13.22. Normalization, free energy, and descent.* Every factor of the joint is a normalized kernel, so the joint has mass one, $E = \mathbf{1}$ has positive probability, and every candidate source has positive prior and posterior mass at every finite configuration. The reference density is $(2\pi)^{-6}e^{-|\mathbf{x}|^2/2}$, the four likelihood factors contribute $4\mathrm{T}\log 2$, and eq. (13.12) with $\sigma_{\mathrm{ev}}^2 = 1/2$ adds

$\frac{\mathrm{T}}{2}\mathbf{x}^{\top}R_{\mathrm{ev}}\mathbf{x}$, so the scaled negative ELBO is $\mathrm{T}\mathbb{E}_q\log(q/\pi_S)+\mathbb{E}_qU_S+\mathrm{B_T}=\mathrm{T}D_{\mathrm{KL}}(q\|q^{\star})+\mathrm{B_T}-\mathrm{T}\log Z_{\mathrm{opt}}$, and the evidence is $e^{-\mathrm{B_T}/\mathrm{T}}Z_{\mathrm{phys}}$, the physical energy omitting $R_{\mathrm{ev}}$. Every $U_S$ is a positive definite quadratic form with Hessian at least $\mathrm{T}I$, whose gradient contains both the sending and the receiving appearances of each agent. Between proposals the process is a linear diffusion with one of sixteen matrices, and proposals are finitely many on bounded intervals, so it exists, is pathwise unique with common random inputs, and does not explode. Since $q^{\star}(\mathbf{x},S)\beta_{S'}(\mathbf{x})=q^{\star}_{\mathbf{X}}(\mathbf{x})\beta_S\beta_{S'}$ is symmetric in $S,S'$, the dissipation identity eq. (13.28) at level zero, without a parameter term, shows that $\mathcal{F}_{\mathrm{T}}$ does not increase. A cycle automorphism, with link signs transformed under reflections, preserves the reference, the kernels, the rows, the kinetics, and the preparation class, so the law at every time is invariant.

*Matching sector.* Fix $S^{(0)}$, with selected edges $e=e_0$ and $f=e_2$; $S^{(1)}$ follows by rotation. For each edge the covectors $\mathbf{a}^b_g,\mathbf{a}^m_g$ of $r^b_g,r^m_g$ have Gram matrix $\begin{pmatrix}3&1\\1&3\end{pmatrix}$, and the two selected edges have disjoint supports. The interaction Hessian of $V_{S^{(0)}}$ is $H_M=2\sum_{g\in\{e,f\}}(\mathbf{a}^b_g\mathbf{a}^{b\top}_g+\mathbf{a}^m_g\mathbf{a}^{m\top}_g)$, and the unit covectors

$$n^{+}_g=\frac{\mathbf{a}^b_g+\mathbf{a}^m_g}{2\sqrt{2}},\qquad n^{-}_g=\frac{\mathbf{a}^b_g-\mathbf{a}^m_g}{2},\qquad g\in\{e,f\},\tag{C.94}$$

are orthonormal eigenvectors with eigenvalues $8$ and $4$; the kernel of $H_M$ has dimension eight. With $\sigma^2_{\mathrm{ev}}=1/2$, $R_{\mathrm{ev}}n^{+}_g=0$, $R_{\mathrm{ev}}n^{-}_g=4n^{-}_g$, and $R_{\mathrm{ev}}$ leaves the kernel invariant, so the held-matching drift matrix $K_M=H_M+\mathrm{T}(I+R_{\mathrm{ev}})$ has normal rates $8+\mathrm{T}$ (twice) and $4+5\mathrm{T}$ (twice), all rates are at least $\mathrm{T}$, and the physical matrix has normal rates $8+\mathrm{T}$ and $4+\mathrm{T}$. With $\mathbf{p}=\sum_{g\in\{e,f\}}(\mathbf{a}^b_g+\mathbf{a}^m_g)$, the vector $\mathbf{x}^{\circ}+\mathbf{p}/4$ lies in the kernel of $H_M$, $H_M\mathbf{p}=8\mathbf{p}$, and $R_{\mathrm{ev}}$ annihilates $\mathbf{x}^{\circ}$ and $\mathbf{p}$, so the deterministic held path from $\mathbf{x}^{\circ}$ is

$$\bar{\mathbf{x}}(t)=e^{-\mathrm{T}t}\Big[\mathbf{x}^{\circ}+\frac{1-e^{-8t}}{4}\mathbf{p}\Big],\qquad r_{\mathrm{int}}(\bar{\mathbf{x}}(t))=-e^{-(8+\mathrm{T})t},\qquad r_{\mathrm{ext}}(\bar{\mathbf{x}}(t))=e^{-\mathrm{T}t}\Big[-\frac{3}{2}+\frac{e^{-8t}}{2}\Big],\tag{C.95}$$

in both channels, where $r_{\mathrm{int}}$ is a selected and $r_{\mathrm{ext}}$ an exterior residual; along it $\bar{k}=\bar{m}$, so the evaluator force vanishes.

*Episode.* Let $\mathscr{B}$ be the event that $|\mathbf{X}_0-\mathbf{x}^{\circ}|<\varrho_{\circ}$, $S_0=S^{(0)}$, and no proposal occurs on $[0,1/4]$. The centered Gaussian component has expected squared norm $12\sigma^2_{\circ}=\varrho^2_{\circ}/4$, so it puts mass at least $3/4$ on the ball, and with the mixture weight $1/2$, the bound $1+\mathfrak{z}\geq 1/2$, the assignment mass $1/16$, and the independent clock, $\Pr(\mathscr{B})\geq\frac{3}{256}e^{-1/4}$. On $[0,1/4]$ the process is the held Ornstein–Uhlenbeck process $\widetilde{\mathbf{X}}_t=e^{-K_Mt}\mathbf{X}_0+\mathbf{Z}_t$, $\mathbf{Z}_t=\sqrt{2\mathrm{T}}\int_0^t e^{-K_M(t-\sigma)}dW_{\sigma}$, and afterwards the two coincide until the first proposal returning another assignment. In an eigenbasis of $K_M$, integration by parts gives $\sup_{t\leq t^{\circ}_*}|Z^k_t|\leq 2\sqrt{2\mathrm{T}}\sup_{t\leq t^{\circ}_*}|W^k_t|$, and $\mathbb{E}\sup_{t\leq t^{\circ}_*}|W^k_t|^2\leq 4t^{\circ}_*$ by Doob's $L^2$ maximal inequality: with $W^*=\sup|W^k|$, optional stopping at the hitting time of a level $\mathrm{y}>0$ gives $\mathrm{y}\Pr(W^*\geq\mathrm{y})\leq\mathbb{E}[|W^k_{t^{\circ}_*}|\mathbf{1}_{W^*\geq\mathrm{y}}]$, and integrating against $2d\mathrm{y}$ with truncation gives the bound. Summing twelve coordinates, $\mathbb{E}\sup|\mathbf{Z}|^2\leq 384\mathrm{T}t^{\circ}_*$, so with $\delta_{\epsilon}=\epsilon^2/(8\sqrt{3})$, Chebyshev's inequality gives $\Pr(\sup_{t\leq t^{\circ}_*}|\mathbf{Z}_t|>\delta_{\epsilon})\leq 73728\epsilon^4t^{\circ}_*$. On $\mathscr{B}$ and the complementary noise event, contraction gives $|\widetilde{\mathbf{X}}_t-\bar{\mathbf{x}}(t)|\leq\varrho_{\circ}+\delta_{\epsilon}$. Every residual covector has norm $\sqrt{3}$, so on $\mathscr{B}$ each residual lies within $\sqrt{3}\varrho_{\circ}$ of $-1$ at time zero and each internal belief relative entropy exceeds $(1-\sqrt{3}/100)^2/2>1/4$,

which is the entrance clause of (a); the restoring rates of (a) are those of $K_M$ computed above. Since $\mathrm{T}t_*^\circ < 1/100$, for $1/4 \le t \le t_*^\circ$ the selected residuals have magnitude below $e^{-2} + \sqrt{3}(\varrho_\circ + \delta_\epsilon) < 1/5$ and the exterior ones above $e^{-\mathrm{T}t_*^\circ}(3/2 - e^{-2}/2) - \sqrt{3}(\varrho_\circ + \delta_\epsilon) > 13/10$. Each receiver's exterior cost therefore exceeds its partner cost by more than one, its posterior row gives its exterior source probability at most $e^{-1/\mathrm{T}}$, and a proposal leaves $S^{(0)}$ with probability at most $4e^{-1/\mathrm{T}}$; for the rate-one clock the expected number of departing proposals up to $t_*^\circ$ is at most $4t_*^\circ e^{-1/\mathrm{T}}$, with self-proposals unrestricted. For $t \ge t_\epsilon$, every residual annihilates the kernel of $H_M$, so a selected residual is at most $e^{-8t} + \sqrt{3}\varrho_\circ e^{-4t} + \sqrt{3}\delta_\epsilon \le \epsilon^{16} + \sqrt{3}\varrho_\circ \epsilon^8 + \epsilon^2/8 \le \epsilon^2$, the belief relative entropy $(r^b)^2/2$ and the model relative entropy $(r^m)^2/(2\sigma_{\mathrm{ev}}^2)$ are at most $\epsilon^4$, and every exterior belief relative entropy exceeds $(13/10)^2/2 > 1/2$, the model ones being larger. Let $\mathscr{A}_t$ be the prefix event that the noise stays within $\delta_\epsilon$ up to $t$ and every proposal in $[1/4, t]$ returns $S^{(0)}$; then $\Pr(\mathscr{A}_{t_*^\circ}^c \mid \mathscr{B}) \le p_\epsilon$. Since $\epsilon^4 t_*^\circ$ increases on $(0, 1/100]$, the first term of $p_\epsilon$ is at most $73728 \cdot 10^{-8}(2\log 100 + 1) < 0.008$, and $e^{-1/\mathrm{T}} < \mathrm{T}$ makes the second below $10^{-14}$, so $p_\epsilon < 1/2$ and $\Pr(\mathscr{B} \cap \mathscr{A}_{t_*^\circ}) \ge \frac{3}{256} e^{-1/4}(1 - p_\epsilon) \ge p_0$. The stopping rule is prefix-measurable, and the same argument holds for $S^{(1)}$ and at the reflected center, which proves the probability bound of (a).

*Stationary comparison.* On each sheet the diffusion with drift $-\nabla U_S$ and covariance $2\mathrm{T}I$ is reversible for $e^{-U_S/\mathrm{T}}d\mathbf{x}$, and the symmetric measure $q^\star(\mathbf{x}, S)\beta_{S'}(\mathbf{x})$ makes the proposal step reversible for $q^\star$, so $q^\star$ is invariant. Under $q^\star$, given $S_0 = M$, the configuration is $\mathcal{N}(0, \mathrm{T}K_M^{-1})$, and for a matched edge $e$ with belief covector $\mathbf{a}_e^b$, $K_M \succeq H_M \succeq 2\mathbf{a}_e^b \mathbf{a}_e^{b\top}$, so $r_e^b = \mathbf{a}_e^b \cdot \mathbf{x}$ has variance $\mathrm{T}\mathbf{a}_e^{b\top} K_M^{-1} \mathbf{a}_e^b \le \mathrm{T}/2$ and $|r_e^b| > 1/\sqrt{2}$ has probability at most $2e^{-1/(2\mathrm{T})}$. The event of (a) requires $S_0$ to be a matching whose matched pairs carry internal belief relative entropy above $1/4$ at time zero, that is $|r_e^b| > 1/\sqrt{2}$ on a matched edge, so under the process started from $q^\star$ it has probability at most $2 \cdot 2e^{-1/(2\mathrm{T})} \le 4e^{-1/(4\mathrm{T})}$, and $4e^{-1/(4\mathrm{T})} < p_0/2$ because $\mathrm{T} \le 10^{-16}$. This completes (a).

*Parent channel.* The covector of $\sum_i \omega_i$ pairs to zero with $n_g^-$ and to $-1/\sqrt{2}$ with $n_g^+$, so it annihilates $Z_M$ and the loop is a function of $P_{W_M}\mathbf{x}$; $W_M$, of dimension nine, contains the kernel of $H_M$ and the fast common-normal direction $(n_e^+ + n_f^+)/\sqrt{2}$. Both $H_M$ and $R_{\mathrm{ev}}$ preserve $W_M$ and $Z_M$, so on a matching fiber the hidden optimizing conditional is $\nu_{\mathrm{T}} = \mathcal{N}(0, \mathrm{T}K_Z^{-1})$ with $K_Z = \mathrm{diag}(8 + \mathrm{T}, 4 + 5\mathrm{T}, 4 + 5\mathrm{T})$, independent of the retained state, and the physical one has precision $\mathrm{diag}(8{+}\mathrm{T}, 4{+}\mathrm{T}, 4{+}\mathrm{T})$. The channel $C$ is one deterministic measurable kernel onto the disjoint union $\coprod_S \{S\} \times W_S$, with Lebesgue measure on each $W_S$ and counting measure on $S$; its pushforwards of the joint, the physical posterior, the optimizer, and the recognition law are normalized, it keeps every assignment, and the chain rule of Theorem 9.6 splits $D_{\mathrm{KL}}(q\|q^\star)$ into the retained term and the mean discarded conditional term. This proves (b).

*Discarded conditional free energy.* Under the held process the hidden coordinates are $z_t = e^{-K_Z t} z_0 + \xi_t$ with $\xi_t \sim \mathcal{N}(0, \mathrm{T}K_Z^{-1}(I - e^{-2K_Z t}))$ independent of the retained coordinates and of the initial data, and $|z_0| \le \varrho_\circ$ on $\mathscr{B}$ because $\mathbf{x}^\circ$ has zero hidden component. Given the retained state, the hidden law is a mixture over $z_0$ of these Gaussians, and for one coordinate with rate $\lambda_k$ and $a = e^{-2\lambda_k t}$ the relative entropy to the stationary law is $\frac{1}{2}[-a - \log(1-a) + \lambda_k a z_0^2/\mathrm{T}]$. Convexity of relative entropy in its first argument, $4 \le \lambda_k \le 9$, and $-a - \log(1-a) \le a^2/(2(1-a))$ bound the mean conditional relative entropy of the held law conditioned on $\mathscr{B}$ by $\frac{9\varrho_\circ^2}{2\mathrm{T}} e^{-8t} + \frac{3}{4(1-e^{-8})} e^{-16t} \le \eta_\epsilon$ for $t \ge t_\epsilon$.

To pass to the episode, let $\widehat{Q}$ be this held law, $\mathscr{R} = \widehat{Q}_Y \otimes \nu_{\mathrm{T}}$, and $Q_{\mathscr{A}}$, $Q_{\mathscr{A}^c}$ its conditionals on $\mathscr{A}_t$ and its complement, of probabilities $1-p$ and $p \le p_\epsilon$. The relative entropy of the joint law of $(\mathbf{1}_{\mathscr{A}_t}, Y, Z)$ to $\mathrm{Bernoulli}(1-p) \otimes \mathscr{R}$, computed in both orders, gives $(1-p)D_{\mathrm{KL}}(Q_{\mathscr{A}}\|\mathscr{R}) + pD_{\mathrm{KL}}(Q_{\mathscr{A}^c}\|\mathscr{R}) = D_{\mathrm{KL}}(\widehat{Q}\|\mathscr{R}) + \mathbb{I}(\mathbf{1}_{\mathscr{A}_t}; (Y,Z))$, and the mutual information is at most $\mathcal{S}_2(p)$. Dropping the second term on the left and the nonnegative retained-marginal term of $D_{\mathrm{KL}}(Q_{\mathscr{A}}\|\mathscr{R})$, and using that $\mathcal{S}_2$ increases on $[0,1/2]$, bounds the mean discarded optimizing conditional relative entropy of $Q_{\mathscr{A}}$ by $\bar{\eta}_\epsilon$; on $\mathscr{B}\cap\mathscr{A}_t$ the held and native processes coincide, so $Q_{\mathscr{A}}$ is the episode law. The physical parent conditions the discarded coordinates also on the evaluator samples, so compare with the normalized law $\Pi(dz \mid y)P_{\mathrm{ev}}(dm, dk \mid \mathbf{x}) \exp[(\mathfrak{d}_y \cdot (k-m) - |\mathfrak{d}_y|^2/2)/(1-\sigma_{\mathrm{ev}}^2)]$, where $\mathfrak{d}_y = \mathfrak{D}y$ and $\mathfrak{d}_z = \mathfrak{D}z$ are orthogonal; the multiplier has expectation one and depends only on retained variables, so this law has the physical hidden conditional given $(y,m,k)$. The recognition evaluator differs from it by the mean shift $\mathfrak{d}_z$, at relative entropy $|\mathfrak{d}_z|^2/(2(1-\sigma_{\mathrm{ev}}^2))$, and the chain rule in both orders bounds the mean discarded physical conditional relative entropy by $D_{\mathrm{KL}}(q(dz \mid y)\|\Pi(dz \mid y)) + \mathbb{E}_q|\mathfrak{d}_z|^2/(2(1-\sigma_{\mathrm{ev}}^2)) = D_{\mathrm{KL}}(q(dz \mid y)\|\nu_{\mathrm{T}}) + \log\frac{4+5\mathrm{T}}{4+\mathrm{T}}$, where $\log\frac{4+5\mathrm{T}}{4+\mathrm{T}} \le \frac{4\mathrm{T}}{4+\mathrm{T}} \le \mathrm{T}$. This proves (c).

*Parent dynamics.* The inherited proposal kernel

$$\beta_C((S,y),\cdot) = \int q^\star(d\mathbf{x} \mid C = (S,y)) \sum_{S'} \beta_{S'}(\mathbf{x}) \delta_{(S', P_{W_{S'}}\mathbf{x})}$$

is a probability kernel with the rate-one clock, reversible with respect to $C_\# q^\star$ because the symmetric measure $q^\star(d\mathbf{x}, S)\beta_{S'}(\mathbf{x})$ is pushed through $(C,C)$; on a matching stratum the retained diffusion is the exact projection of the held process, and elsewhere $C$ is the identity. On a matching fiber, with $z_+ = ((n_e^+ + n_f^+)/\sqrt{2}) \cdot y$ and the hidden coordinates $z_1, z_2, z_3$ of $Z_M$, the selected costs are $\mathcal{D}_e = (z_+ + z_1)^2 + z_2^2$ and $\mathcal{D}_f = (z_+ - z_1)^2 + z_3^2$, and each exterior common residual $(r_g^b + r_g^m)/2$ annihilates $Z_M$ and satisfies $\mathcal{D}_g \ge ((r_g^b + r_g^m)/2)^2$. If both exterior common residuals have magnitude at least $\mathrm{r}_\circ$, each receiver chooses its exterior source with probability at most $\exp[(\mathcal{D}_{\mathrm{in}} - \mathcal{D}_{\mathrm{out}})/\mathrm{T}]$, and the Gaussian identity $\mathbb{E}e^{(a+Z)^2/\mathrm{T}} = (1-2\sigma^2/\mathrm{T})^{-1/2}\exp[a^2/(\mathrm{T}(1-2\sigma^2/\mathrm{T}))]$ with $\sigma^2 = \mathrm{T}/(8+\mathrm{T})$ and $\mathrm{T}/(4+5\mathrm{T})$, whose factors are at most $\sqrt{4/3}$ and $\sqrt{2}$, gives $\beta_C((S^{(0)},y), \{S' \ne S^{(0)}\}) \le 4\sqrt{8/3}\exp[-(\mathrm{r}_\circ^2 - 4z_+^2/3)/\mathrm{T}]$. On the episode $\mathrm{r}_\circ = 13/10$ and $|z_+| \le \epsilon^2$ on $[t_\epsilon, t_*^\circ]$, so this is at most $4\sqrt{8/3}e^{-1/\mathrm{T}}$. Condition on $\mathscr{B}\cap\mathscr{A}_{t_\epsilon}$ and let $\mu_{\mathrm{act}}$ be the law of $C(\mathbf{X}_{t_\epsilon}, S_{t_\epsilon})$ under this prefix. The Brownian increments after $t_\epsilon$ are independent of the prefix, and the native future equals the held process until a departure or a tube exit, whose conditional probability is at most $\Pr(\mathscr{A}_{t_*^\circ}^c \mid \mathscr{B})/\Pr(\mathscr{A}_{t_\epsilon} \mid \mathscr{B}) \le p_\epsilon/(1-p_\epsilon)$. The inherited parent started from $\mu_{\mathrm{act}}$ and driven by the same increments has the retained drift and diffusion of the projected held process until its first departure; the retained tube fails with conditional probability at most $p_\epsilon/(1-p_\epsilon)$, and on the tube the parent departs during the unit interval with probability at most $4\sqrt{8/3}e^{-1/\mathrm{T}}$. The union bound and the triangle inequality give (d). □

*Proof of Proposition 14.4. Rows.* For a continuous row path $\pi$ whose entries lie in $[\mathrm{c}_0, 1-\mathrm{c}_0]$ for some $\mathrm{c}_0 \in (0,1/2]$, construct the kinetics pathwise as in the proof above: the configuration drift $-\nabla U_S$ does not involve the prior, and at a clock time $\tau$ receiver $i$ takes source $i+1$ exactly when

a uniform variable on $[0,1]$, independent and fresh for each clock time and receiver, is at most its posterior row

$$\beta_{i,i+1}(\mathbf{x};\pi(\tau)) = \frac{\pi_{i,i+1}(\tau)e^{-\mathcal{D}_{e(i,i+1)}(\mathbf{x})/\mathrm{T}}}{\sum_{j=i\pm1}\pi_{ij}(\tau)e^{-\mathcal{D}_{e(i,j)}(\mathbf{x})/\mathrm{T}}}$$

at $\mathbf{x}=\mathbf{X}_{\tau-}$. Write $q_t^{\pi}$ for the law at time $t$. Drive the constructions for two such row paths $\pi,\pi'$ by the same initial data, Brownian motion, clock, and uniform variables. They coincide until the first clock time at which the uniform variable of some receiver falls between its posterior rows under $\pi$ and $\pi'$, and, given agreement before a clock time $\tau$, this happens at $\tau$ with conditional probability at most $\sum_i|\beta_{i,i+1}(\mathbf{x};\pi(\tau))-\beta_{i,i+1}(\mathbf{x};\pi'(\tau))|$ at the common configuration. With $\mathrm{p}=\pi_{i,i+1}$ and $\pi_{i,i-1}=1-\mathrm{p}$, the derivative of a posterior row in $\mathrm{p}$ is $\beta_{i,i+1}(1-\beta_{i,i+1})/(\mathrm{p}(1-\mathrm{p}))\le 1/(4\mathrm{c}_0(1-\mathrm{c}_0))\le 1/(2\mathrm{c}_0)$, and the clock has rate one, so

$$\left|q_t^{\pi}(S_i=j)-q_t^{\pi'}(S_i=j)\right| \le \frac{2}{\mathrm{c}_0}\int_0^t \|\pi(s)-\pi'(s)\|_\infty ds,$$

where $\|\cdot\|_\infty$ is the maximum over receivers and sources. A solution of eq. (14.4) that starts uniform satisfies

$$\pi_{ij}(t)-\tfrac12 = \epsilon\int_0^t e^{-\epsilon(t-s)}\big(q_s(S_i=j)-\tfrac12\big)ds, \qquad \left|\pi_{ij}(t)-\tfrac12\right| \le \tfrac12\big(1-e^{-\epsilon t}\big)\le \tfrac{\epsilon t}{2},$$

because $0\le q_s(S_i=j)\le 1$; each row sums to one, and on $[0,t_1]$ its entries lie in $[\mathrm{c}_0,1-\mathrm{c}_0]$ with $\mathrm{c}_0=e^{-\epsilon t_1}/2$. Fix $t_1>0$ and let $\Gamma$ send a row path $\pi$ to the right-hand side of the first identity with $q_s=q_s^{\pi}$. It maps the closed set of continuous row paths on $[0,t_1]$ with $|\pi_{ij}(t)-1/2|\le(1-e^{-\epsilon t})/2$ into itself, and

$$\|\Gamma\pi(t)-\Gamma\pi'(t)\|_\infty \le \frac{2\epsilon t_1}{\mathrm{c}_0}\int_0^t\|\pi(s)-\pi'(s)\|_\infty ds, \qquad t\le t_1,$$

so the $n$th iterate of $\Gamma$ contracts the supremum distance by the factor $(2\epsilon t_1^2/\mathrm{c}_0)^n/n!$, and $\Gamma$ has exactly one fixed point in that set. Every solution lies in the set, solutions on nested intervals agree, and $s\mapsto q_s^{\pi}(S_i=j)$ is continuous because the clock has rate one, so the fixed point is differentiable and solves eq. (14.4). The rows are therefore a deterministic continuous function of time, and the row bound holds. This proves (a).

*Seed events.* The function $\epsilon t_*^{\circ}=\epsilon(2\log(1/\epsilon)+1)$ increases on $(0,1/100]$, with value $(1+4\log 10)/100<0.1022$ at $1/100$, so $t_*^{\circ}<2/\epsilon$ and, by (a), every row entry lies in $[0.4489,0.5511]$ on $[0,t_*^{\circ}]$, whence $\pi_{ik}(t)/\pi_{ij}(t)\le 0.5511/0.4489<1.23$ for the two entries of every row. For a sign $\sigma\in\{1,-1\}$ and a perfect matching $M\in\{S^{(0)},S^{(1)}\}$, let $\mathscr{B}_{\sigma,M}$ be the event that $|\mathbf{X}_0-\sigma\mathbf{x}^{\circ}|<\varrho_{\circ}$, $S_0=M$, and no proposal occurs on $[0,1/4]$. The initial configuration, the initial assignment, and the clock are independent, so

$$\Pr(\mathscr{B}_{\sigma,M}) = \frac{e^{-1/4}}{16}\mathcal{N}(\mathsf{m}_0,I_{12})\big\{|\mathbf{x}-\sigma\mathbf{x}^{\circ}|<\varrho_{\circ}\big\},$$

and the four events are disjoint because $|\mathbf{x}^{\circ}|=2>\varrho_{\circ}$.

*Episode.* Condition on $\mathscr{B}_{1,S^{(0)}}$. The configuration drift does not involve the rows, and the rows are deterministic, so the paragraph *Episode* of the proof above applies line by line from every starting point in the ball, whatever the law of $\mathbf{X}_0$ on it: the held Ornstein–Uhlenbeck process, whose noise bound involves only the Brownian motion; the contraction to the held path eq. (C.95); the residual bounds below $1/5$ and above $13/10$ on $[1/4, t_*^\circ]$, which make each receiver's exterior cost exceed its partner cost by more than one; and the relative-entropy bounds for $t \geq t_\epsilon$, all deterministic on the noise event. The rows enter only the exterior source probability: for a receiver $i$ with partner $j$ and exterior source $k$,

$$\beta_{ik}(\mathbf{x}; \pi(t)) \leq \frac{\pi_{ik}(t)}{\pi_{ij}(t)} \exp\Big[-\frac{\mathcal{D}_{e(i,k)}(\mathbf{x}) - \mathcal{D}_{e(i,j)}(\mathbf{x})}{\mathrm{T}}\Big] < 1.23 e^{-1/\mathrm{T}}$$

in place of $e^{-1/\mathrm{T}}$. A proposal therefore leaves $S^{(0)}$ with probability at most $5e^{-1/\mathrm{T}}$, the expected number of departing proposals up to $t_*^\circ$ is at most $5t_*^\circ e^{-1/\mathrm{T}}$, and the prefix event $\mathscr{A}_{t_*^\circ}$ of that paragraph fails with conditional probability at most

$$\tilde{p}_\epsilon = 73728\epsilon^4 t_*^\circ + 5t_*^\circ e^{-1/\mathrm{T}} < 0.008 + 10^{-14} < \tfrac{1}{2},$$

by the same monotonicity of $\epsilon^4 t_*^\circ$ and $e^{-1/\mathrm{T}} < \mathrm{T}$. On $\mathscr{B}_{1,S^{(0)}} \cap \mathscr{A}_{t_*^\circ}$, the block $\{0, 1\}$, of size two and joined to its complement by $e_1$ and $e_3$, carries on its internal edge $e_0$ a belief relative entropy above $(1 - \sqrt{3}\varrho_\circ)^2/2 > 1/4$ at time zero, because $|\mathbf{X}_0 - \mathbf{x}^\circ| < \varrho_\circ$ and every residual covector has norm $\sqrt{3}$; it has internal belief and model relative entropies at most $\epsilon^4$ on $[t_\epsilon, t_*^\circ]$, an interval of length one inside $[0, 2/\epsilon]$, and both boundary belief relative entropies above $1/2$ on $[1/4, t_*^\circ]$, so contrast formation occurs.

*The other seeds.* A cycle automorphism $g$, with link signs transformed under reflections, acts on configurations by an orthogonal signed permutation $\mathsf{O}_g$ with $U_{gS}(\mathsf{O}_g\mathbf{x}) = U_S(\mathbf{x})$ and $\mathcal{D}_{e(gi,gj)}(\mathsf{O}_g\mathbf{x}) = \mathcal{D}_{e(i,j)}(\mathbf{x})$. With $(g\pi)_{gi,gj} = \pi_{ij}$, it therefore maps the kinetics with row path $\pi$ from $(\mathbf{x}, S)$ to the kinetics with row path $g\pi$ from $(\mathsf{O}_g\mathbf{x}, gS)$, driven by the Brownian motion $\mathsf{O}_g W$ and by the receiver uniform variables relabeled by $g$, the variable of receiver $i$ serving receiver $gi$ and being replaced by one minus itself when $g$ is a reflection. Indeed a reflection sends the source $i+1$ of receiver $i$ to the source $gi-1$ of receiver $gi$, and the image receiver takes it exactly when the original variable is at most $\beta_{i,i+1}(\mathbf{x}; \pi) = \beta_{gi,gi-1}(\mathsf{O}_g\mathbf{x}; g\pi) = 1 - \beta_{gi,gi+1}(\mathsf{O}_g\mathbf{x}; g\pi)$, which is the rule of the image kinetics applied to one minus that variable, except on the null event that the variable equals its row; under the rotation the sources are carried along and the variable is kept. The relabeled family is again a family of independent uniform variables independent of the other inputs, so the image process has the law of the kinetics with rows $g\pi$, and $g\pi$ has the same ratio bound. The rotation $i \mapsto i+1$ fixes $\mathbf{x}^\circ$ and maps $S^{(0)}$ to $S^{(1)}$, the reflection $i \mapsto -i$ maps $(\mathbf{x}^\circ, S^{(0)})$ to $(-\mathbf{x}^\circ, S^{(1)})$, and their composition maps $(\mathbf{x}^\circ, S^{(0)})$ to $(-\mathbf{x}^\circ, S^{(0)})$, so the conditional bound holds on the other three seed events, with the images of the block. Summing over the four disjoint events,

$$\Pr(\text{contrast formation}) \geq (1 - \tilde{p}_\epsilon)\frac{2e^{-1/4}}{16} \sum_{\pm} \mathcal{N}(\mathsf{m}_0, I_{12})\big\{|\mathbf{x} \mp \mathbf{x}^\circ| < \varrho_\circ\big\},$$

which gives the first bound of (b) because $\tilde{p}_\epsilon < 1/2$. The Gaussian density on the ball about $\mathbf{x}^\circ$ is at least $(2\pi)^{-6}\exp[-(|\mathbf{x}^\circ - \mathrm{m}_0| + \varrho_\circ)^2/2]$, the ball has volume $\pi^6\varrho_\circ^{12}/6!$, and $(2\pi)^{-6}\pi^6/6! = 1/46080$, which gives the second bound after dropping the ball about $-\mathbf{x}^\circ$. Frozen rows have ratio one, so the argument and the bound hold without learning. The sub-event named in (b) is the event of part (a) of the proposition above with the block of the held matching. For fixed rows $\pi$ the proposal step is reversible for $q^\star_\pi$ by the same symmetry $q^\star_\pi(\mathbf{x}, S)\beta_{S'}(\mathbf{x};\pi) = q^\star_{\pi,\mathbf{X}}(\mathbf{x})\beta_S(\mathbf{x};\pi)\beta_{S'}(\mathbf{x};\pi)$, the sheet diffusions do not involve the rows, and the conditional law $\mathcal{N}(0, \mathrm{T}K_M^{-1})$ of $q^\star_\pi$ given $S_0 = M$ does not depend on $\pi$, so the paragraph *Stationary comparison* above gives the invariance and the bound $4e^{-1/(4\mathrm{T})}$ for every $\pi$. This proves (b).

*Invariant initial laws.* Since $\mathsf{O}_g$ is orthogonal, $g$ maps $\mathcal{N}(\mathrm{m}_0, I_{12})$ to $\mathcal{N}(\mathsf{O}_g\mathrm{m}_0, I_{12})$, so the initial law is invariant exactly when $\mathrm{m}_0$ is fixed by the rotation and the reflection. The rotation fixes exactly the means whose belief, model, and link blocks are constant, and the reflection reverses every edge and hence the sign of a constant link block, which leaves $(\mathrm{c}_1\mathbf{1}_4, \mathrm{c}_2\mathbf{1}_4, 0)$. For such a mean, since the uniform assignment law and the uniform rows are also invariant, $g$ maps the recognition law $q_t$ of the rows $\pi$ to the recognition law $g_\# q_t$ of the rows $g\pi$ from the same initial law, by the pathwise image of the paragraph *The other seeds*, whose relabeled driving inputs have the law of the original ones; and $g\pi$ solves eq. (14.4) with $g_\# q_t(S_{gi} = gj) = q_t(S_i = j)$, so $(g\pi, g_\# q)$ is a second solution from the same data and uniqueness in (a) gives $g\pi = \pi$. The reflection $j \mapsto 2i - j$ fixes receiver $i$ and exchanges its sources $i \pm 1$, so $\pi_{i,i+1} = \pi_{i,i-1} = 1/2$. This proves (c). □

## C.8 Depth one under exclusive source gating

*Proof of Proposition 13.24.* Throughout, $F = \mathrm{att}(S)$ and $W = W_F$ once a sheet $S$ is fixed, $P_W$ is the orthogonal projection onto $W$, and $\mathsf{Q}_W$, $\mathsf{Q}_{W^\perp}$ are matrices whose columns form orthonormal bases of $W$ and of $W^\perp$. Since $H_S$ is the sum of the positive semidefinite maps $\mathbf{a}_{e(i,S_i)}^\top \mathbf{a}_{e(i,S_i)}$, its kernel is $\bigcap_{e\in F} \ker \mathbf{a}_e = W_F$ and its range is $W_F^\perp$, the span of the rows of $\mathbf{a}_F$, so $\mathrm{k}_S = \mathrm{rank}\,\mathbf{a}_F$. For every edge the Gram matrix of the two rows of $\mathbf{a}_e$ is $\left(\begin{smallmatrix} 3 & 1 \\ 1 & 3 \end{smallmatrix}\right)$, as in eq. (C.94), so $\|\mathbf{a}_e\| = 2$.

*(a) Two bands.* Let $M = H_S + \mathrm{T}\Omega$. By Weyl's inequality [45, Sec. 4.3], with eigenvalues in increasing order, $\lambda_k(H_S) + \mathrm{T}\lambda_{\min}(\Omega) \le \lambda_k(M) \le \lambda_k(H_S) + \mathrm{T}\|\Omega\|$ for every $k$. The first $2N + |E| - \mathrm{k}_S$ eigenvalues of $H_S$ vanish, the remaining $\mathrm{k}_S$ lie in $[\mathrm{h}_S, \|H_S\|]$, and $\lambda_{\min}(\Omega) = \|\Omega^{-1}\|^{-1}$. The range was identified above, and $\pi$ does not enter $M$.

*(b) Gating wells.* For each receiver, the row sums to one and every weight is at least $\min_j \pi_{ij}$, so $\min_j \pi_{ij} e^{-\min_j \mathcal{D}_{e(i,j)}/\mathrm{T}} \le \sum_j \pi_{ij} e^{-\mathcal{D}_{e(i,j)}/\mathrm{T}} \le e^{-\min_j \mathcal{D}_{e(i,j)}/\mathrm{T}}$; taking $-\mathrm{T}\log$ and summing over $i$ gives eq. (13.56). Since $V_S$ is a sum over receivers and $\pi_S$ a product,

$$\sum_S \pi_S e^{-V_S/\mathrm{T}} = \prod_i \sum_{j\in\mathcal{N}_i} \pi_{ij} e^{-\mathcal{D}_{e(i,j)}/\mathrm{T}} = e^{-V_\pi/\mathrm{T}},$$

which is the configuration marginal up to the reference factor. If every edge of a covering set $F$ has zero residual at $\mathbf{x}$, every receiver has an option of zero cost, so $V_{\min}(\mathbf{x}) = 0$; a common translation of the belief and model means of the agents of one component of $(V, F)$, with the links

fixed, changes no residual of an edge of $F$, because the two endpoints of such an edge either both move or both stay, so $V_{\min}$ remains zero.

*(c) Rank formula.* In the coordinates $(\bar{k}, \mathfrak{d}, \omega)$ with $\mathfrak{d} = \bar{k} - \bar{m}$, the rows of $\mathbf{a}_e$ for $e : j \to i$ are $r_e^b = \bar{k}_i - \bar{k}_j - \omega_e$ and $r_e^m = r_e^b - (\mathfrak{d}_i - \mathfrak{d}_j)$; replacing the second by $r_e^m - r_e^b = -(\partial\mathfrak{d})_e$ is an invertible change within the rows of $e$. The rows $r_e^b$, $e \in F$, read only $(\bar{k}, \omega)$ and each has coefficient $-1$ on its own link coordinate and zero on every other, so they are linearly independent. The rows $(\partial\mathfrak{d})_e$, $e \in F$, read only $\mathfrak{d}$ and form the oriented incidence map of $(V, F)$ on $\mathbb{R}^V$, whose kernel consists of the functions constant on each component, so their rank is $N - \mathrm{cc}(F)$. The two families act on disjoint coordinate blocks, so $\operatorname{rank} \mathbf{a}_F = |F| + N - \mathrm{cc}(F)$.

*Minimum edge covers.* Let $F = \mathrm{att}(S)$. Every vertex selects an incident edge, so $F$ covers $V$ and every component of $(V, F)$ contains an edge; one edge from each component is a matching, so $\mathrm{cc}(F) \le \mathfrak{n}$, and a connected graph on $\mathrm{c}$ vertices has at least $\mathrm{c} - 1$ edges, so $|F| \ge N - \mathrm{cc}(F)$. Hence $\mathrm{k}_S \ge 2(N - \mathrm{cc}(F)) \ge 2(N - \mathfrak{n})$, with equality exactly when $F$ is a forest with $\mathfrak{n}$ components and therefore $N - \mathfrak{n}$ edges. The same two inequalities show that every edge cover has at least $N - \mathfrak{n}$ edges, and a maximum matching together with one incident edge at each unmatched vertex, which exists because no vertex is isolated, is an edge cover with exactly $N - \mathfrak{n}$ edges, since the other endpoint of each added edge is matched, or else the matching could be enlarged. An edge cover with $N - \mathfrak{n}$ edges has at least $\mathfrak{n}$ components and hence exactly $\mathfrak{n}$, and is a forest; so the covering forests with $\mathfrak{n}$ components are exactly the minimum edge covers. In a minimum edge cover every edge has an endpoint of degree one, since otherwise removing it would leave a smaller cover, and a tree in which every edge has a leaf endpoint is a star, because two vertices of degree at least two would be joined by a path whose first edge has no leaf endpoint. Conversely, a star with center $z$ is attended when every leaf selects $z$ and $z$ selects one leaf, which the source sets allow, so every minimum edge cover is some $\mathrm{att}(S)$. Adding to $F$ an edge that joins two components raises $|F|$ by one and lowers $\mathrm{cc}(F)$ by one, hence raises the rank by two.

*Sheet weights.* Since $H_S/\mathrm{T} + \Omega \succ 0$, $Z_S = (2\pi)^{(2N+|E|)/2} \det(H_S/\mathrm{T} + \Omega)^{-1/2}$. In the basis $(\mathsf{Q}_{W^\perp}, \mathsf{Q}_W)$ the matrix $H_S/\mathrm{T} + \Omega$ has the blocks $H_1/\mathrm{T} + \Omega_{11}$, $\Omega_{10}$, $\Omega_{01}$, $\Omega_{00}$, with $H_1 = \mathsf{Q}_{W^\perp}^\top H_S \mathsf{Q}_{W^\perp} \succeq \mathrm{h}_S I$ and $\Omega_{00} = \Omega|_W$. Eliminating the invertible block $\Omega_{00}$ by block row and column operations gives $\det(H_S/\mathrm{T} + \Omega) = \det \Omega_{00} \det(H_1/\mathrm{T} + \Sigma)$ with $\Sigma = \Omega_{11} - \Omega_{10}\Omega_{00}^{-1}\Omega_{01}$, and $0 \prec \Sigma \preceq \Omega_{11} \preceq \|\Omega\| I$. Then $\mathrm{T}\Sigma \preceq (\mathrm{T}\|\Omega\|/\mathrm{h}_S) H_1$, so $\det H_1 \le \det(H_1 + \mathrm{T}\Sigma) \le (1 + \mathrm{T}\|\Omega\|/\mathrm{h}_S)^{\mathrm{k}_S} \det H_1$, and $\det H_1 = \mathrm{pdet}\, H_S$. Since $\det(H_1/\mathrm{T} + \Sigma) = \mathrm{T}^{-\mathrm{k}_S} \det(H_1 + \mathrm{T}\Sigma)$, the determinant lies between $\mathrm{T}^{-\mathrm{k}_S}\mathrm{z}_S^{-2}$ and $(1 + \mathrm{T}\|\Omega\|/\mathrm{h}_S)^{\mathrm{k}_S}\mathrm{T}^{-\mathrm{k}_S}\mathrm{z}_S^{-2}$, which is eq. (13.57). The same block form shows that under $q_\Omega(\cdot \mid S)$, which is $\mathcal{N}(0, (H_S/\mathrm{T} + \Omega)^{-1})$, the coordinates $y = \mathsf{Q}_W^\top \mathbf{x}$ given $z = \mathsf{Q}_{W^\perp}^\top \mathbf{x}$ are Gaussian with precision $\Omega_{00}$, the $(W, W)$ block of the joint precision, which does not depend on $\mathrm{T}$.

*Merge cost.* $q_\Omega(S) = \pi_S Z_S / \sum_{S'} \pi_{S'} Z_{S'} \le \pi_S Z_S / \sum_{S' \in \mathcal{M}} \pi_{S'} Z_{S'}$. The upper bound of eq. (13.57) for $S$ and its lower bound for each $S' \in \mathcal{M}$, where $\mathrm{k}_{S'} = 2(N - \mathfrak{n})$ and $\mathrm{h}_{S'} \ge \mathrm{h}_*$, give $q_\Omega(S) \le \mathrm{T}^{(\mathrm{k}_S - 2(N-\mathfrak{n}))/2}(1 + \mathrm{T}\|\Omega\|/\mathrm{h}_*)^{N-\mathfrak{n}} \pi_S \mathrm{z}_S / \sum_{S' \in \mathcal{M}} \pi_{S'} \mathrm{z}_{S'}$. If $\mathrm{cc}(\mathrm{att}(S)) \le \mathfrak{n} - m$, then $\mathrm{k}_S \ge 2(N - \mathrm{cc}(\mathrm{att}(S))) \ge 2(N - \mathfrak{n}) + 2m$; with $\mathrm{T} \le 1$, summing over such $S$ gives eq. (13.58).

*(d) Block law.* The sheets with $\mathrm{cc}(\mathrm{att}(S)) < \mathfrak{n}$ carry mass at most $\mathrm{K}_1\mathrm{T}$ by eq. (13.58). Fix a sheet with $\mathrm{cc}(F) = \mathfrak{n}$. No component of $(V, F)$ contains two disjoint edges, since replacing that component's

edge in the matching of one edge per component would give $\mathfrak{n}+1$ disjoint edges. A connected graph in which no two edges are disjoint is a star or a triangle: fix an edge $uv$; every other edge meets $\{u,v\}$; if there are edges $vw$ and $uw'$ with $w\neq u$ and $w'\neq v$, then $w=w'$, and a further edge would have to meet $\{u,v\}$ and would then be disjoint from $vw$ or from $uw$, so the graph is the triangle $uvw$; otherwise all edges share one vertex. Replacing each triangle component by a two-edge path on the same three vertices gives a covering forest of stars with $\mathfrak{n}$ components, a minimum edge cover by (c), so the partition of $V$ into the components of $F$ is a minimum-edge-cover partition. The coherent-block partition $\mathfrak{P}_{\mathrm{a}}(\mathbf{x})$ differs from it only if some attended edge has $|\mathbf{a}_e\mathbf{x}|>\mathrm{a}$ or some edge joining two components of $F$ has $|\mathbf{a}_e\mathbf{x}|\le\mathrm{a}$. Under $q_\Omega(\cdot\mid S)$ the configuration is $\mathcal{N}(0,\mathrm{T}(H_S+\mathrm{T}\Omega)^{-1})$. For $e\in F$, $M=H_S+\mathrm{T}\Omega\succeq\mathbf{a}_e^\top\mathbf{a}_e$, so $\|\mathbf{a}_eM^{-1/2}\|\le1$ and the covariance $\mathrm{T}\mathbf{a}_eM^{-1}\mathbf{a}_e^\top$ of $\mathbf{a}_e\mathbf{x}$ is at most $\mathrm{T}I_2$; each of its two centered Gaussian coordinates has variance at most $\mathrm{T}$, and $|\mathbf{a}_e\mathbf{x}|>\mathrm{a}$ forces one of them beyond $\mathrm{a}/\sqrt{2}$, an event of probability at most $2\cdot2e^{-\mathrm{a}^2/(4\mathrm{T})}$. For an edge $e\notin F$ joining two components, $\operatorname{rank}\mathbf{a}_{F\cup\{e\}}=\operatorname{rank}\mathbf{a}_F+\operatorname{rank}(\mathbf{a}_e|_W)$, because $\ker\mathbf{a}_{F\cup\{e\}}=W\cap\ker\mathbf{a}_e$; by (c) the restriction has rank two, so $\mathbf{a}_e^W=\mathbf{a}_e\mathsf{Q}_W$ has two nonzero singular values, the least of which is at least $\mathrm{s}_*$. Writing $\mathbf{x}=\mathsf{Q}_{W^\perp}z+\mathsf{Q}_Wy$, the conditional law of $y$ given $z$ has precision $\Omega_{00}\preceq\|\Omega\|I$, so $\mathbf{a}_e^Wy$ is Gaussian in $\mathbb{R}^2$ with covariance $\mathbf{a}_e^W\Omega_{00}^{-1}\mathbf{a}_e^{W\top}\succeq\|\Omega\|^{-1}\mathbf{a}_e^W\mathbf{a}_e^{W\top}\succeq(\mathrm{s}_*^2/\|\Omega\|)I_2$ and density at most $(2\pi)^{-1}\|\Omega\|/\mathrm{s}_*^2$. The event $|\mathbf{a}_e\mathbf{x}|\le\mathrm{a}$ places $\mathbf{a}_e^Wy=\mathbf{a}_e\mathbf{x}-\mathbf{a}_e\mathsf{Q}_{W^\perp}z$ in a disc of radius $\mathrm{a}$ about a point determined by $z$, so its conditional probability is at most $\pi\mathrm{a}^2\cdot\|\Omega\|/(2\pi\mathrm{s}_*^2)$. A union bound over at most $|E|$ edges of each kind, integrated over the sheets, gives eq. (13.59). □

## C.9 Proof of Pairing Selection

*Proof of Theorem 14.5.* Write $p_i=\pi_{i,i+1}$, encode the assignment by $\sigma_i=+1$ when receiver $i$ selects $i+1$, and put $\varepsilon=\epsilon_{\mathrm{L}}$. The state remains $(\mathbf{x},\sigma)\in\mathbb{R}^{12}\times\{\pm1\}^4$, with the shared edge costs of eq. (13.50). The multiplicity of edge $e$ is $m_e(\sigma)=1_{\{\sigma_e=+1\}}+1_{\{\sigma_{e+1}=-1\}}$. In this proof define

$$\Omega=I+R_{\mathrm{ev}},\qquad \mathsf{A}_\sigma=H_\sigma+\mathrm{T}\Omega,\qquad H_\sigma=\nabla^2V_\sigma,\qquad \mathsf{A}_\sigma\succeq\mathrm{T}I.$$

Let $\mathsf{w}_\sigma=\det(\mathsf{A}_\sigma/\mathrm{T})^{-1/2}$, $\mathsf{Z}(p)=\sum_\sigma\pi_\sigma(p)\mathsf{w}_\sigma$, and

$$\mu_p(\mathbf{x},\sigma)=\frac{\pi_\sigma(p)}{(2\pi)^6\mathsf{Z}(p)}\exp\left[-\frac{\mathbf{x}^\top\mathsf{A}_\sigma\mathbf{x}}{2\mathrm{T}}\right],\qquad \pi_\sigma(p)=\prod_i p_i^{1_{\{\sigma_i=+1\}}}(1-p_i)^{1_{\{\sigma_i=-1\}}}.$$

Each held-sheet diffusion preserves its Gaussian factor, and a clock ring resamples $\sigma$ from its conditional law under $\mu_p$. Thus $\mu_p$ is invariant for the frozen rows, including boundary rows on their supported sheets. Set

$$\bar\beta_i(p)=\mu_p(\sigma_i=+1),\qquad \mathfrak{g}_i(p)=\partial_{p_i}\log\mathsf{Z}(p),\qquad \mathscr{V}_i(p)=\bar\beta_i(p)-p_i=p_i(1-p_i)\mathfrak{g}_i(p).\tag{C.96}$$

The last identity follows by differentiating the product prior. Since $\mathsf{Z}$ is a positive multilinear polynomial on the closed cube, $\mathfrak{g}$ and $\mathscr{V}$ extend analytically to a neighborhood of that cube and have bounded derivatives there. The well-posedness argument for proposition 14.4 applies at every

fixed $\mathrm{T} > 0$; it gives deterministic continuously differentiable rows and

$$\min_i\{p_i(t), 1-p_i(t)\} \geq \tfrac{1}{2}e^{-\varepsilon t}, \qquad \|p(t) - \tfrac{1}{2}\|_\infty \leq \tfrac{1}{2}\varepsilon t. \tag{C.97}$$

All constants below may depend on the fixed temperature. Constants in bounds on a fixed initial law may also depend on its mean, whereas the window constant in the conclusion will depend only on the temperature.

## Frozen convergence and the initial imprint

Here $\mathrm{Ent}_\nu(g) = \int g \log(g/\int g d\nu) d\nu$ for nonnegative $g$, and total variation is $\sup_A |q(A) - \nu(A)|$. For the time-one frozen kernel, the function $\mathsf{V}(\mathbf{x}) = |\mathbf{x}|^2$ satisfies

$$\mathbb{E}_{\mathbf{x}}|\mathbf{X}_t|^2 \leq e^{-2\mathrm{T}t}|\mathbf{x}|^2 + 12(1-e^{-2\mathrm{T}t}), \qquad P_1^p\mathsf{V} \leq e^{-2\mathrm{T}}\mathsf{V} + 12. \tag{C.98}$$

Indeed, jumps leave $\mathbf{X}$ unchanged, and the diffusion generator applied to its squared norm is at most $24\mathrm{T} - 2\mathrm{T}|\mathbf{X}|^2$. Choose $R > 24/(1-e^{-2\mathrm{T}})$. For any frozen $p$, some assignment $\sigma_*(p)$ has $\pi_{\sigma_*} \geq 1/16$. Restrict the rate-one clock to exactly one ring in $[0, 1/2]$ and none afterward. Its probability is $e^{-1}/2$. Starting with $|\mathbf{x}| \leq \sqrt{R}$, the first held Gaussian motion is in the ball of radius $r' = \sqrt{R} + \sqrt{24\mathrm{T}}$ at that ring with probability at least $1/2$. On the ball, the probability to redraw $\sigma_*$ is at least $16^{-1}e^{-v(r')/\mathrm{T}}$, where $v(r') = \max_{|\mathbf{z}|\leq r',\sigma} V_\sigma(\mathbf{z}) < \infty$. The subsequent held Gaussian transition lasts between $1/2$ and $1$. Its means are bounded and its covariance matrices have common positive lower and finite upper bounds. Its density consequently has a positive common lower bound on the unit ball. Hence, for some $\alpha_\mathrm{T} > 0$ independent of $p$,

$$P_1^p((\mathbf{x},\sigma),\cdot) \geq \alpha_\mathrm{T}\,\mathrm{Unif}(B_1) \otimes \delta_{\sigma_*(p)} \quad \text{when } |\mathbf{x}|^2 \leq R.$$

These are the drift and minorization hypotheses of Hairer and Mattingly [36, Assumptions 1–2 and Theorems 1.2–1.3]. The explicit dependence in their Theorem 1.3 makes the resulting constants uniform in the frozen rows, even though the minorizing measure depends on $p$. Thus $\mu_p$ is the unique invariant law and

$$|P_t^p\phi(\mathbf{x},\sigma) - \mu_p\phi| \leq \mathrm{K}_\mathrm{T}e^{-\lambda_\mathrm{T}t}(1+|\mathbf{x}|^2)\|\phi\|_\infty. \tag{C.99}$$

Intermediate times follow by applying the skeleton estimate to $P_r^p\phi$, $0 \leq r < 1$.

Put $\mathbf{v} = (1,-1,1,-1)$, $f_\mathbf{v}(\sigma) = \frac{1}{2}\sum_i \mathbf{v}_i\sigma_i$, and define

$$\mathsf{I}_i(\mathsf{m}_0) = \int_0^\infty [q_t^0(\sigma_i = +1) - \tfrac{1}{2}]dt, \qquad \mathsf{j}_\mathrm{T}(\mathsf{m}_0) = \mathbf{v}\cdot\mathsf{I}(\mathsf{m}_0). \tag{C.100}$$

At uniform rows, the stationary law is invariant under cycle rotations and reflections, so each stationary row is $1/2$ and $\mu_{1/2}f_\mathbf{v} = 0$. Estimate eq. (C.99) proves absolute convergence of these

integrals. It also gives a Poisson integral

$$\mathfrak{R}_{\mathrm{T}}(\mathbf{x},\sigma)=\int_0^\infty P_t^{1/2} f_{\mathbf{v}}(\mathbf{x},\sigma)dt, \qquad |\mathfrak{R}_{\mathrm{T}}(\mathbf{x},\sigma)|\le \mathrm{K}_{\mathrm{T}}(1+|\mathbf{x}|^2). \tag{C.101}$$

Convolution of this integral with the initial Gaussian expresses $\mathrm{j}_{\mathrm{T}}$ as an entire function of its complexified mean: for every compact set of complex means, all differentiated integrands are bounded by an integrable polynomial times a Gaussian. The same holds for I. A rotation by one changes the sign of $\mathrm{j}_{\mathrm{T}}$.

For $t\le 1/(10\varepsilon)$, all row entries are in $[\mathrm{c}_0,1-\mathrm{c}_0]$, $\mathrm{c}_0=0.45$. A posterior receiver probability has derivative at most $1/[4\mathrm{c}_0(1-\mathrm{c}_0)]\le 1/(2\mathrm{c}_0)$ with respect to its prior entry. Coupling the common clocks and their four posterior rows therefore gives

$$\|q_t^{\varepsilon}-q_t^{0}\|_{\mathrm{TV}}\le \frac{2}{\mathrm{c}_0}\int_0^t \|p(s)-\tfrac12\|_\infty ds\le \frac{\varepsilon t^2}{2\mathrm{c}_0}.$$

Variation of constants in the learning equation, followed by eq. (C.99), gives

$$\big|p(t)-\tfrac12-\varepsilon \mathsf{I}(\mathsf{m}_0)\big|_\infty \le \mathrm{K}_{\mathrm{T},\mathsf{m}_0}(\varepsilon e^{-\lambda_{\mathrm{T}} t}+\varepsilon^2 t)+\frac{\varepsilon^2 t^3}{6\mathrm{c}_0}. \tag{C.102}$$

In particular, at $t_1=\varepsilon^{-1/4}$, $p(t_1)-1/2=\varepsilon \mathsf{I}(\mathsf{m}_0)+o(\varepsilon)$.

## Nonvanishing of the matching imprint

Take $\mathbf{u}_\circ=(0_4,0_4,(0,1,0,-1))$, whose two residual vectors are $(0,-1,0,1)$. In the noiseless switched motion the subspace $\bar{k}=\bar{m}$ is invariant. After removing the scalar factor $e^{-\mathrm{T}t}$, its common residual vector satisfies

$$\dot\rho=-\mathsf{G}M_\sigma\rho, \qquad \mathsf{G}=4I-A_{C_4}, \qquad M_\sigma=\mathrm{diag}(m_e(\sigma)). \tag{C.103}$$

Here the cyclic difference matrix contributes $2I-A_{C_4}$ and the shared-link drift contributes $2I$. Of the sixteen assignments, six have $m_1=m_3$. Their first held motion is $e^{-ct}(0,-1,0,1)$, with $c=4m_1\in\{0,4,8\}$. At their first clock ring every receiver prefers the even-edge matching $S^{(0)}$. Holding that matching then leaves the odd residuals constant and the even residuals zero in the scalar-free coordinates.

Four further assignments have multiplicity patterns

$$(1,1,2,0),\quad (1,0,2,1),\quad (2,0,1,1),\quad (2,1,1,0).$$

The signed swap of coordinates $1,3$ and the rotation by two preserve the even-edge matching and carry the initial residual to itself or its negative. They reduce these four cases to $(1,1,2,0)$. Write its

scalar-free coordinates as $\rho = (-a, -b, -d, h)$. Then

$$\begin{pmatrix} \dot{a} \\ \dot{b} \\ \dot{d} \end{pmatrix} = \begin{pmatrix} -4 & 1 & 0 \\ 1 & -4 & 2 \\ 0 & 1 & -8 \end{pmatrix} \begin{pmatrix} a \\ b \\ d \end{pmatrix}, \qquad (a, b, d)(0) = (0, 1, 0), \qquad \dot{h} = -(a + 2d),\ h(0) = 1.$$

The nonnegative cone is invariant, and the negative row sums give $\max(a, b, d) \le 1$, hence $a \le 1/4$, $d \le 1/8$. The half-spaces $a - 3b/4 - 11d/25 \le 0$ and $d - b/4 \le 0$ are invariant: their outward derivatives on the corresponding boundary are $-b/400 - 7d/100$ and $-a/4 - b/8$, respectively. Consequently $a \le 43b/50$, $d \le b/4$, and $b \ge e^{-4t}$. The attended block is similar to a negative definite symmetric matrix because $\mathsf{G} \succ 0$. Solving $\rho(\infty) = \rho(0) - \mathsf{G}z$, with $z_3 = 0$ and the first three final coordinates zero, gives $h(\infty) = 6/7$. Thus $h \ge 6/7$.

If the first ring occurs at $t_{\mathrm{r}}$ and redraws $S^{(0)}$, the subsequent scalar-free held motion, in elapsed time $r$, is

$$\begin{aligned} a_r &= a_0 e^{-8r}, & d_r &= d_0 e^{-8r}, \\ b_r &= b_0 + \tfrac{1}{4}(a_0 + d_0)(1 - e^{-8r}), & h_r &= h_0 - \tfrac{1}{4}(a_0 + d_0)(1 - e^{-8r}) \ge \tfrac{171}{224}. \end{aligned}$$

Every adjacent odd-edge absolute residual exceeds its even-edge neighbor by at least $(7/50)e^{-4t_{\mathrm{r}}}$. Combining these four cases with the preceding six, the deterministic motion $\mathbf{Y}$ from $s\mathbf{u}_\circ$ has, throughout the hypothesized matching hold,

$$\sqrt{\mathcal{D}_{\mathrm{odd}}(\mathbf{Y}_t)} - \sqrt{\mathcal{D}_{\mathrm{even}}(\mathbf{Y}_t)} \ge \frac{7}{50} s e^{-\mathrm{T}t - 8t_{\mathrm{r}}}, \qquad t \ge t_{\mathrm{r}}. \tag{C.104}$$

There are therefore ten good initial assignments. For the remaining six only $f_{\mathbf{v}} \ge -2$ will be used.

Set $h_s = (\log s)/\mathrm{T}$, for $s > e$, and write the initial configuration as $s\mathbf{u}_\circ + \xi$, $\xi \sim \mathcal{N}(0, I)$. Along the actual assignment path decompose $\mathbf{X} = \mathbf{Y} + \mathbf{E} + \mathbf{Z}$, where the first two terms solve its switched deterministic equation from $s\mathbf{u}_\circ$ and $\xi$, respectively, while

$$d\mathbf{Z} = -\mathsf{A}_{\sigma_t}\mathbf{Z}dt + \sqrt{2\mathrm{T}}dW_t, \qquad \mathbf{Z}_0 = 0.$$

All three are adapted differential equations, and $|\mathbf{E}_t| \le e^{-\mathrm{T}t}|\xi|$. Itô's formula yields

$$\mathscr{L} e^{|\mathbf{Z}|^2/4} \le \mathrm{T}(6 - |\mathbf{Z}|^2/4) e^{|\mathbf{Z}|^2/4} \le \mathrm{T}e^5.$$

Stopping at radius $\log s$, and using $\mathbb{E}e^{|\xi|^2/4} = 64$, bounds the probability that either $|\xi| > \log s$ or $\sup_{t \le h_s} |\mathbf{Z}_t| > \log s$ by $(65 + \mathrm{T}e^5 h_s) e^{-(\log s)^2/4}$. On its complement each square-root edge cost changes from its $\mathbf{Y}$ value by at most $2\mathrm{C}_N \log s$, where $\mathrm{C}_N$ is the largest norm of a residual-pair map divided by $\sqrt{2}$.

Define

$$L_s(t_{\mathrm{r}}) = h_s - \frac{8}{\mathrm{T}} t_{\mathrm{r}} - \frac{2}{\mathrm{T}} \log\log s, \qquad B_s = (L_s(t_{\mathrm{r}}) - t_{\mathrm{r}})_+.$$

For a good initial assignment, eq. (C.104) gives a square-root gap at least $(7/50)(\log s)^2$ on $[t_{\mathrm{r}}, L_s]$. For sufficiently large $s$, while both noise bounds hold, the actual energy gap is at least $\kappa_0 (\log s)^4$,

$\kappa_0 = (7/50)^2/4$. Each receiver's wrong choice at a ring then has probability at most $e^{-\kappa_0(\log s)^4/\mathrm{T}}$. Apply this bound at the first ring, then use the rate-one compensator up to the first wrong redraw or exit from the noise bound, stopping at $L_s \wedge h_s$. The endpoint $L_s$ is known after the first ring, and the gap estimate holds until that first failure. Thus the matching holds throughout the nonempty interval $[t_{\mathrm{r}}, L_s]$ except on an event of probability at most

$$\delta_s = (65 + \mathrm{T}e^5 h_s)e^{-(\log s)^2/4} + 4(1+h_s)e^{-\kappa_0(\log s)^4/\mathrm{T}}, \qquad h_s\delta_s \longrightarrow 0. \tag{C.105}$$

Since $\mathbb{E}t_{\mathrm{r}} = 1$, a good initial assignment contributes at least

$$-2h_s + 4\mathbb{E}B_s - 4h_s\delta_s \geq 2h_s - 4(1 + 8/\mathrm{T}) - (8/\mathrm{T})\log\log s - 4h_s\delta_s$$

to $\mathbb{E}\int_0^{h_s} f_{\mathbf{v}}(\sigma_t)dt$. The other six contribute at least $-2h_s$. Averaging the sixteen uniform initial assignments gives

$$\mathbb{E}\int_0^{h_s} f_{\mathbf{v}}(\sigma_t)dt \geq \tfrac{1}{2}h_s - \frac{5}{\mathrm{T}}\log\log s - O_{\mathrm{T}}(1).$$

By eq. (C.98), $\mathbb{E}|\mathbf{X}_{h_s}|^2 \leq 14$. The Markov property and eq. (C.101) bound the integral after $h_s$ in absolute value by a constant independent of $s$. Therefore

$$\mathrm{j}_{\mathrm{T}}(s\mathbf{u}_\circ) \geq \frac{\log s}{2\mathrm{T}} - \frac{5}{\mathrm{T}}\log\log s - O_{\mathrm{T}}(1). \tag{C.106}$$

This proves that $\mathrm{j}_{\mathrm{T}}$ is not identically zero at any fixed $\mathrm{T} > 0$. Its zero set is Lebesgue-null. For completeness, this last assertion follows by induction on dimension: a nonzero analytic function in one variable has isolated zeros; in several variables at least one coefficient $\partial_n^k f(\mathbf{x}', 0)$ is nonzero as an analytic function, or analyticity would make $f$ vanish everywhere. Outside the null zero set of that coefficient, each one-dimensional slice has isolated zeros, and Fubini's theorem concludes the induction.

## Entropy balance for the evolving recognition law

Fix a finite horizon $h$. All rows are bounded away from zero on this interval by eq. (C.97). Write $q_\sigma(t, \mathbf{x})$ for the evolving density, $Q = \sum_\sigma q_\sigma$, and $b_\sigma(t, \mathbf{x}) = \mu_{p(t)}(\mathbf{x}, \sigma)/\sum_\eta \mu_{p(t)}(\mathbf{x}, \eta)$. For the forward Ornstein–Uhlenbeck semigroup $K_\sigma$, the last-proposal formula is

$$q_\sigma(t) = e^{-t}K_\sigma(t)q_\sigma(0) + \int_0^t e^{-(t-r)}K_\sigma(t-r)[b_\sigma(r)Q(r)]dr. \tag{C.107}$$

Iterating this positive formula, and bounding each proposal factor by one, reduces each term to a Gaussian motion along a deterministic sequence of sheets. Put $M = \max_\sigma \|\mathsf{A}_\sigma\|$. Its means have norm at most $|\mathsf{m}_0|$, and its covariances lie between $e^{-2Mh}I$ and $(1 + 2\mathrm{T}h)I$. Summing the at most $16^n$ sequences on each time simplex, of volume $t^n/n!$, gives a finite upper Gaussian envelope. The no-proposal term gives a lower Gaussian envelope. Consequently, uniformly on $[0, h]$,

$$0 < \mathrm{c}_h e^{-\mathrm{C}_h(1+|\mathbf{x}|^2)} \leq q_\sigma(t, \mathbf{x}) \leq \mathrm{C}_h e^{-\mathrm{c}_h|\mathbf{x}|^2}. \tag{C.108}$$

This bound does not assert Gaussianity conditional on a state-dependent assignment history.

Differentiating one Ornstein–Uhlenbeck kernel against a function bounded by a polynomial times a Gaussian gives a bound of the same type with factor $(t-r)^{-1/2}$. To check it, use the Gaussian kernel covariance bounds proportional to $t-r$, absorb its linear derivative into a Gaussian with twice the covariance, and complete the square in the convolution. This singularity is integrable in eq. (C.107), proving a first-derivative Gaussian bound. Moreover

$$\nabla b_\sigma = -\frac{b_\sigma}{\mathrm{T}}\left(H_\sigma \mathbf{x} - \sum_\eta b_\eta H_\eta \mathbf{x}\right), \qquad |\nabla b_\sigma| \le \mathrm{K_T}|\mathbf{x}|.$$

The source $b_\sigma Q$ therefore has a first derivative with the same type of bound. For the second derivative, transfer one derivative to this source using $\nabla K_\sigma(r) g = e^{r\mathsf{A}_\sigma} K_\sigma(r) \nabla g$, and differentiate the other kernel factor. The singularity is again only $r^{-1/2}$. The explicit initial Gaussian term supplies the bounds at zero. The forward equation obtained by taking a time difference quotient in eq. (C.107) then bounds the time derivative too:

$$|\nabla q_\sigma| + |\nabla^2 q_\sigma| + |\partial_t q_\sigma| \le \mathrm{C}_h(1+|\mathbf{x}|)^k e^{-\mathrm{c}_h|\mathbf{x}|^2}. \tag{C.109}$$

The derivative transfer follows directly by integration by parts in the Gaussian transition density; its inverse mean map is uniformly bounded on this finite interval. These estimates justify the differentiations, without assuming finite Fisher information.

Set $f_\sigma = q_\sigma/\mu_{p(t),\sigma}$ and $\mathsf{H}(t) = D_{\mathrm{KL}}(q_t \| \mu_{p(t)})$. The logarithms of both densities grow at most quadratically by eq. (C.108), and $\partial_t \log \mu_p$ is bounded in $\mathbf{x}$, since only its prior weight and normalizer vary. Together with eq. (C.109), this proves that $\mathsf{H}$ is absolutely continuous on $[0,h]$ and permits its integral to be differentiated. To integrate the diffusion by parts, first multiply by a smooth radial cutoff $\chi_R$, equal to one on $B_R$, zero outside $B_{2R}$, with $|\nabla \chi_R| \le \mathrm{C}/R$, increasing to one. On each compact cylinder the densities have positive lower bounds, so $\log f_\sigma$ is a legitimate test function. The boundary term is bounded by

$$\frac{\mathrm{C}_h}{R}\int_{R\le|\mathbf{x}|\le 2R} (1+|\mathbf{x}|)^{k+3} e^{-\mathrm{c}_h|\mathbf{x}|^2} d\mathbf{x} \longrightarrow 0.$$

Detailed balance $\mu_\sigma b_\eta = \mu_\eta b_\sigma$ symmetrizes the finite jump sum. Both resulting production integrands are nonnegative, so monotone convergence removes the cutoff and simultaneously proves finiteness and local time integrability of the production. The result is

$$\begin{aligned} \mathsf{H}'(t) &= -\mathscr{D}_{p(t)}(f) - \sum_\sigma \int q_\sigma \partial_t \log \mu_{p(t),\sigma}, \\ \mathscr{D}_p(f) &= \mathrm{T}\sum_\sigma \int \mu_{p,\sigma} \frac{|\nabla f_\sigma|^2}{f_\sigma} \\ &\quad + \frac{1}{2}\sum_{\sigma,\eta}\int \mu_{p,\sigma} b_\eta (f_\eta - f_\sigma)(\log f_\eta - \log f_\sigma). \end{aligned} \tag{C.110}$$

The equality holds almost everywhere and in integrated form on every finite interval. The jump term is integrable already from the Gaussian envelopes and quadratic logarithmic bounds; thus the cutoff argument does not presuppose a bound on the diffusion production. This also proves $\mathsf{H}(0) < \infty$.

## Entropy contraction at frozen rows

In the original twelve-dimensional coordinates put $\mathsf{B}_\sigma = \mathsf{A}_\sigma/\mathrm{T}$, the Gaussian sheet precision, and set

$$\kappa = \min_\sigma \lambda_{\min}(\mathsf{B}_\sigma) > 0, \qquad M = \max_\sigma \|\mathsf{B}_\sigma\|, \qquad q^*_\sigma = \mu_p(\sigma) = \frac{\pi_\sigma \mathsf{w}_\sigma}{\mathsf{Z}}, \qquad q^*_{\min} = \min_\sigma q^*_\sigma.$$

Let $\gamma_\sigma = \mathcal{N}(0, \mathsf{B}_\sigma^{-1})$, and, for a real function $u$, define

$$\mathscr{E}_c(u) = \mathrm{T} \sum_\sigma \int \mu_{p,\sigma} |\nabla u_\sigma|^2, \qquad \mathscr{E}_j(u) = \frac{1}{2} \sum_{\sigma,\eta} \int \mu_{p,\sigma} b_\eta (u_\sigma - u_\eta)^2.$$

The Gaussian inequalities needed below follow from the explicit Ornstein–Uhlenbeck semigroup with invariant law $\gamma_\sigma$. Its gradient commutes as $\nabla P_t g = e^{-t\mathsf{B}_\sigma} P_t \nabla g$. Differentiating entropy and applying Cauchy–Schwarz inside this kernel gives

$$-\frac{d}{dt} \operatorname{Ent}_{\gamma_\sigma}(P_t g) = \int \frac{|\nabla P_t g|^2}{P_t g} d\gamma_\sigma \le e^{-2\kappa t} \int \frac{|\nabla g|^2}{g} d\gamma_\sigma.$$

For bounded positive smooth $g$, the explicit Gaussian formula and dominated convergence make the entropy tend to zero. Integration, with $g = u^2$, gives

$$\operatorname{Ent}_{\gamma_\sigma}(u^2) \le \frac{2}{\kappa} \int |\nabla u|^2 d\gamma_\sigma, \qquad \operatorname{Var}_{\gamma_\sigma}(u) \le \frac{1}{\kappa} \int |\nabla u|^2 d\gamma_\sigma. \tag{C.111}$$

The second estimate follows by differentiating variance in the same semigroup. Positive regularization, truncation, and smooth cutoffs extend both estimates to their form domains: the cutoff errors tend to zero by integrability of the squared function and its gradient, and the entropy passes to the limit by its variational formula.

Replace $u$ by $|u|$, which decreases both energies, and normalize $\int u^2 d\mu_p = 1$. Put $z_\sigma = (\int u_\sigma^2 d\gamma_\sigma)^{1/2}$. The entropy chain rule gives

$$\operatorname{Ent}_{\mu_p}(u^2) = \sum_\sigma q^*_\sigma \operatorname{Ent}_{\gamma_\sigma}(u_\sigma^2) + \operatorname{Ent}_{q^*}(z^2).$$

Here $0 \le z_\sigma \le R = (q^*_{\min})^{-1/2}$. The function $\psi(z) = z^2 \log z^2 - z^2 + 1$ satisfies $\psi(1) = \psi'(1) = 0$ and $\psi''(z) = 4 + 4\log z$, whence $\psi(z) \le 2(1 + \log R)(z-1)^2$ on $[0, R]$. Since $\sum q^*_\sigma z_\sigma^2 = 1$,

$$\operatorname{Ent}_{q^*}(z^2) \le L_p \operatorname{Var}_{q^*}(z), \qquad L_p = 4 + 2\log(1/q^*_{\min}).$$

Indeed $\sum q^*_\sigma (z_\sigma - 1)^2 = 2(1 - \mathbb{E}_{q^*} z) \le 2 \operatorname{Var}_{q^*}(z)$.

Fix a ball $B = B_\rho$ in these coordinates, and define

$$m_\sigma = |B|^{-1}\int_B u_\sigma, \quad m_* = \sum_\sigma \pi_\sigma m_\sigma, \quad X_\sigma = \left[\frac{(2\pi)^6 \mathsf{w}_\sigma e^{\rho^2\|\mathsf{B}_\sigma\|/2}}{|B|} - 1\right]_+, \quad B_{\mathrm{T}} = \max_\sigma (1+\sqrt{X_\sigma})^2.$$

Variance about any center is an upper bound on the variance, so

$$\mathrm{Var}_{q^*}(z) \le 2\sum_\sigma q^*_\sigma (z_\sigma - m_\sigma)^2 + 2\sum_\sigma q^*_\sigma (m_\sigma - m_*)^2.$$

The first sum is at most $(B_{\mathrm{T}}/(\mathrm{T}\kappa))\mathscr{E}_c(u)$: $|z_\sigma - \gamma_\sigma u_\sigma| \le \sqrt{\mathrm{Var}_{\gamma_\sigma}(u_\sigma)}$, and Cauchy–Schwarz bounds $|m_\sigma - \gamma_\sigma u_\sigma|$ by the same standard deviation times $\sqrt{\chi^2(\mathrm{Unif}(B)\|\gamma_\sigma)} \le \sqrt{X_\sigma}$. For the second sum, Jensen's inequality gives

$$\sum_\sigma q^*_\sigma (m_\sigma - m_*)^2 \le \frac{\mathsf{w}_{\max}}{\mathsf{Z}|B|}\sum_{\sigma,\eta}\pi_\sigma\pi_\eta\int_B (u_\sigma - u_\eta)^2.$$

On $B$, the explicit Gaussian densities give

$$\mu_{p,\sigma} b_\eta \ge \frac{\pi_\sigma\pi_\eta e^{-\rho^2 M}}{(2\pi)^6 \mathsf{Z}}.$$

The normalizer therefore cancels, and the preceding sum is at most $2(2\pi)^6 \mathsf{w}_{\max} e^{\rho^2 M}\mathscr{E}_j(u)/|B|$. Combining these bounds proves

$$\mathrm{Ent}_{\mu_p}(u^2) \le \frac{2}{\mathrm{T}\kappa}(1 + L_p B_{\mathrm{T}})\mathscr{E}_c(u) + \frac{4L_p(2\pi)^6 \mathsf{w}_{\max} e^{\rho^2 M}}{|B|}\mathscr{E}_j(u).$$

For $u = \sqrt{f}$, the production is at least $4(\mathscr{E}_c(u) + \mathscr{E}_j(u))$. Its jump part uses $(a-b)(\log a - \log b) \ge 4(\sqrt{a} - \sqrt{b})^2$, which follows by Cauchy–Schwarz on the integral of $r^{-1/2}$ between $a$ and $b$. The inequality extends to the form domain by the preceding approximations. Finally,

$$q^*_{\min} \ge p^4_{\min}\frac{\mathsf{w}_{\min}}{\mathsf{w}_{\max}}, \qquad p_{\min} = \min_i\{p_i, 1-p_i\}.$$

Thus there are finite positive constants depending only on $\mathrm{T}$ with

$$\mathscr{D}_p(f) \ge \lambda(\mathrm{T}, p)\,\mathrm{Ent}_{\mu_p}(f), \qquad \lambda(\mathrm{T}, p) \ge \frac{\mathrm{c}_{\mathrm{T}}}{\mathrm{K}_{\mathrm{T}} + 8\log(1/p_{\min})}. \tag{C.112}$$

The factor $\mathrm{T}$ in $\mathscr{E}_c$ is the same diffusion coefficient as in eq. (C.110).

## The matching branch of the small-temperature landscape

The needed sheet weights can be computed using two four-dimensional determinants. Let $D$ be the cyclic difference matrix, and $C_\sigma = \mathrm{diag}(m_e(\sigma))$. In the orthogonal coordinates $(\bar{k} + \bar{m})/\sqrt{2}$, $(\bar{k} - \bar{m})/\sqrt{2}$, and $\omega$, the energy splits into $\frac{1}{2}(Du - \sqrt{2}\omega)^\top C_\sigma (Du - \sqrt{2}\omega) + \frac{1}{2}(Dv)^\top C_\sigma Dv$, while the

discrepancy coordinate has reference precision $5I$. The determinant identity for the first block then gives

$$\det \mathsf{A}_\sigma = \mathrm{T}^4 \det[\mathrm{T}I + C_\sigma(DD^\top + 2I)] \det[5\mathrm{T}I + D^\top C_\sigma D]. \tag{C.113}$$

The two matching assignments, eight end-path assignments, four middle-path assignments, and two oriented cycles have the respective determinants

$$\begin{aligned}
\mathsf{d}_M &= 25\mathrm{T}^8(\mathrm{T}+8)^2(5\mathrm{T}+4)^2,\\
\mathsf{d}_E &= 5\mathrm{T}^6(\mathrm{T}^3+16\mathrm{T}^2+77\mathrm{T}+112)(125\mathrm{T}^3+200\mathrm{T}^2+85\mathrm{T}+8),\\
\mathsf{d}_N &= 5\mathrm{T}^6(\mathrm{T}+4)(5\mathrm{T}+2)(\mathrm{T}^2+12\mathrm{T}+28)(25\mathrm{T}^2+30\mathrm{T}+4),\\
\mathsf{d}_C &= 5\mathrm{T}^5(\mathrm{T}+2)(\mathrm{T}+4)^2(\mathrm{T}+6)(5\mathrm{T}+2)^2(5\mathrm{T}+4).
\end{aligned}$$

These follow in eq. (C.113) from the multiplicity representatives $(2,0,2,0)$, $(2,1,1,0)$, $(1,2,1,0)$, and $(1,1,1,1)$, respectively. Normalize the matching weight to one and write

$$\mathsf{r}_E = \sqrt{\mathsf{d}_M/\mathsf{d}_E}, \qquad \mathsf{r}_N = \sqrt{\mathsf{d}_M/\mathsf{d}_N}, \qquad \mathsf{r}_C = \sqrt{\mathsf{d}_M/\mathsf{d}_C}.$$

Their leading terms are $\mathsf{r}_E, \mathsf{r}_N = \sqrt{40/7}\mathrm{T} + O(\mathrm{T}^2)$ and $\mathsf{r}_C = \sqrt{5/3}\mathrm{T}^{3/2}(1+O(\mathrm{T}))$. Choose $\mathrm{T}_0 > 0$ so that $0 < \mathsf{r}_E, \mathsf{r}_N, \mathsf{r}_C \le 1/8$ for $0 < \mathrm{T} \le \mathrm{T}_0$.

Use matching coordinates $x = (p_0, 1-p_1, p_2, 1-p_3)$. The normalized evidence is the multilinear polynomial whose vertex weights are one at $\mathbf{0}, \mathbf{1}$, $\mathsf{r}_E$ for subsets of cardinality one or three, $\mathsf{r}_N$ for the four adjacent pairs, and $\mathsf{r}_C$ for the two opposite pairs. Complementation preserves these weights, so the centered reduced field is odd. With $D_0 = 1 + 4\mathsf{r}_E + 2\mathsf{r}_N + \mathsf{r}_C$, its Jacobian at $x = \mathbf{1}/2$ is

$$\frac{1}{D_0}\begin{pmatrix}
0 & 1-\mathsf{r}_C & 1-2\mathsf{r}_N+\mathsf{r}_C & 1-\mathsf{r}_C\\
1-\mathsf{r}_C & 0 & 1-\mathsf{r}_C & 1-2\mathsf{r}_N+\mathsf{r}_C\\
1-2\mathsf{r}_N+\mathsf{r}_C & 1-\mathsf{r}_C & 0 & 1-\mathsf{r}_C\\
1-\mathsf{r}_C & 1-2\mathsf{r}_N+\mathsf{r}_C & 1-\mathsf{r}_C & 0
\end{pmatrix}.$$

For example, the adjacent Hessian entry of the evidence is $(1-\mathsf{r}_C)/2$, its uniform value is $D_0/8$, and the row mobility is $1/4$, giving the displayed entry. The eigenvalues are

$$\alpha = \frac{3-2\mathsf{r}_N-\mathsf{r}_C}{D_0} > 0, \qquad \beta = \frac{-1-2\mathsf{r}_N+3\mathsf{r}_C}{D_0} < 0, \qquad \eta = \frac{-1+2\mathsf{r}_N-\mathsf{r}_C}{D_0} < 0 \quad \text{twice.} \tag{C.114}$$

The positive eigendirection is $\mathbf{1}$ in matching coordinates, or $\mathbf{v}$ in the original rows. In fact the chosen bounds give $\alpha \ge 7/5$, $\beta \le -1/3$, and $\eta \le -2/5$.

Cyclic invariance makes the matching diagonal $x = z\mathbf{1}$ invariant. Writing $w = z(1-z)$, its normalized evidence is

$$Z_{\mathrm{diag}}(z) = 1 + 4(\mathsf{r}_E - 1)w + (2 - 8\mathsf{r}_E + 4\mathsf{r}_N + 2\mathsf{r}_C)w^2, \qquad \dot{z} = \frac{z(1-z)}{4Z_{\mathrm{diag}}(z)} Z'_{\mathrm{diag}}(z).$$

The derivative is $(1-2z)[4(\mathsf{r}_E-1)+2(2-8\mathsf{r}_E+4\mathsf{r}_N+2\mathsf{r}_C)w]$. The bracket is negative at both endpoints $w=0,1/4$, where its values are $4(\mathsf{r}_E-1)$ and $-3+2\mathsf{r}_N+\mathsf{r}_C$. It is affine in $w$, hence negative throughout. Thus the diagonal motion tends to $\mathbf{1}$ from $z>1/2$, and to $\mathbf{0}$ from $z<1/2$, reaching every prescribed endpoint neighborhood in finite slow time. At either vertex the Jacobian is $(\mathsf{r}_E-1)I$: a single receiver flip changes weight one to $\mathsf{r}_E$, the mobility vanishes, and its diagonal derivative supplies the factor $-(1-\mathsf{r}_E)$. Both vertices are therefore sinks with positive inward rate $\nu=1-\mathsf{r}_E$.

## Tracking, escape, and selection

Let $a_i(t)=q_t(\sigma_i=+1)$ and $\Delta_i(t)=a_i(t)-\bar\beta_i(p(t))$. Differentiating the normalized product prior gives

$$\partial_{p_i}\log\mu_p=\frac{1_{\{\sigma_i=+1\}}-\bar\beta_i(p)}{p_i(1-p_i)}.$$

Inserting the learning equation and eq. (C.96) in eq. (C.110) yields

$$\mathsf{H}'=-\mathscr{D}_p(f)-\varepsilon\sum_i\mathfrak{g}_i(p)\Delta_i-\varepsilon\sum_i\frac{\Delta_i^2}{p_i(1-p_i)}\le-\lambda(\mathrm{T},p)\mathsf{H}+\varepsilon\|\mathfrak{g}(p)\|_1\sqrt{\mathsf{H}/2}. \tag{C.115}$$

The last bound uses $|\Delta_i|\le\mathrm{TV}(q_t,\mu_p)\le\sqrt{\mathsf{H}/2}$. The nonpositive square is kept until this step; it removes any inverse-row factor from the forcing. Writing $\Lambda(t)=\int_0^t\lambda(\mathrm{T},p(r))dr$, comparison for $\sqrt{\mathsf{H}+\delta^2}$, followed by $\delta\downarrow0$, gives

$$\sqrt{\mathsf{H}(t)}\le e^{-\Lambda(t)/2}\sqrt{\mathsf{H}(0)}+\frac{\varepsilon}{2\sqrt{2}}\int_0^t e^{-(\Lambda(t)-\Lambda(r))/2}\|\mathfrak{g}(p(r))\|_1dr. \tag{C.116}$$

The bounded fitness, eq. (C.97), and eq. (C.112) show that, on $t\le2\mathrm{C}_\mathrm{T}\varepsilon^{-1}\log(1/\varepsilon)$,

$$\mathrm{TV}(q_t,\mu_{p(t)})\le e^{-\Lambda(t)/2}\sqrt{\mathsf{H}(0)/2}+O_\mathrm{T}(\varepsilon\log(1/\varepsilon)). \tag{C.117}$$

Indeed $\lambda$ is bounded below on this horizon by a positive constant divided by $1+\log(1/\varepsilon)$.

For the saddle a finer bound is needed. Let $\|\mathfrak{g}(p)-\mathfrak{g}(p')\|_1\le L_\mathrm{T}\|p-p'\|_\infty$, and suppose $\lambda\ge\lambda_\flat>0$ on $[s_0,t]$. The estimate $|\dot p_i|\le\varepsilon$ in eq. (C.116) gives

$$\mathrm{TV}(q_t,\mu_{p(t)})\le e^{-\lambda_\flat(t-s_0)/2}\sqrt{\mathsf{H}(s_0)/2}+\frac{\varepsilon}{2\lambda_\flat}\|\mathfrak{g}(p(t))\|_1+\frac{L_\mathrm{T}\varepsilon^2}{\lambda_\flat^2}. \tag{C.118}$$

The constants follow by integrating $1$ and $t-r$ against $e^{-\lambda_\flat(t-r)/2}$.

Fix $0<\mathrm{T}\le\mathrm{T}_0$ and a mean with $\mathsf{j}_\mathrm{T}(\mathsf{m}_0)\neq0$, and put $\mathbf{e}=\mathbf{v}/2$. In slow time $\mathfrak{s}=\varepsilon t$, write $y=p-1/2=a\mathbf{e}+z$, $z\perp\mathbf{e}$. By eq. (C.102), at $t_1=\varepsilon^{-1/4}$,

$$a(t_1)=\varepsilon(\mathsf{j}_\mathrm{T}(\mathsf{m}_0)/2+o(1)),\qquad|z(t_1)|=O(\varepsilon).$$

On $[0, t_1]$, symmetry gives $\mathfrak{g}(1/2) = 0$, so $\|\mathfrak{g}(p(t))\|_1 \le L_{\mathrm{T}}\varepsilon t/2$. The bound in eq. (C.116) gives $\sqrt{\mathsf{H}(t_1)} = O(\varepsilon^{7/4})$. Consequently, until exit from any sufficiently small fixed ball, eq. (C.118) implies

$$|\Delta|_2 \le \kappa_\varepsilon |y| + b_\varepsilon, \qquad \kappa_\varepsilon = O(\varepsilon), \quad b_\varepsilon = O(\varepsilon^{7/4}) = o(|a(t_1)|). \tag{C.119}$$

The constants here may depend on the fixed mean.

Let $\beta_\perp = \max(\beta, \eta) < 0$ and $\gamma = \alpha - \beta_\perp > 0$. Oddness and diagonal invariance of the analytic reduced field imply, for $|y| \le r$,

$$|\langle \mathscr{V}(1/2 + y), \mathbf{e}\rangle - \alpha a| \le \mathrm{C}|y|^3, \qquad |P_\perp \mathscr{V}(1/2 + y) - D\mathscr{V}(1/2)z| \le \mathrm{C}|y|^2|z|.$$

For the second inequality, subtract the zero perpendicular field on $a\mathbf{e}$ and integrate its derivative along the transverse segment. The derivative differs from its value at zero by $O(|y|^2)$ because the field is odd. Writing $w = |z|/|a|$, the upper right derivatives of norms give

$$\begin{aligned} \frac{d|a|}{d\mathfrak{s}} &\ge |a|[\alpha - (\mathrm{C}r^2 + \kappa_\varepsilon)(1 + w)] - b_\varepsilon, \\ D^+ w &\le -[\gamma - \mathrm{C}r^2 - (\mathrm{C}r^2 + \kappa_\varepsilon)(1 + w)]w + (1 + w)(\kappa_\varepsilon + b_\varepsilon/|a|). \end{aligned} \tag{C.120}$$

Choose $W$ larger than twice the initial limiting ratio, then choose $r > 0$ so small that the terms $\mathrm{C}r^2(1 + W)$ consume less than one quarter of both positive rates. For small enough $\varepsilon$, eq. (C.119) controls the remaining terms. A bootstrap then keeps $w \le W$, makes $|a|$ increase without changing sign, and gives

$$w(\mathfrak{s}) \le w(\mathfrak{s}_1)e^{-\gamma(\mathfrak{s}-\mathfrak{s}_1)/2} + \frac{2(1 + W)}{\gamma}(\kappa_\varepsilon + b_\varepsilon/|a(t_1)|).$$

The lower growth bound for $|a|$ yields an exit time no larger than

$$\frac{1 + O(r^2) + o(1)}{\alpha} \log(1/\varepsilon)$$

in slow time. The corresponding upper norm-growth bound forces the exit time to tend to infinity. Hence $w \to 0$ at exit, and the exit point tends to $\operatorname{sign}(\mathrm{j}_{\mathrm{T}})r\mathbf{e}$.

The reduced diagonal orbit from that point enters any fixed neighborhood of its matching vertex in finite slow time. The tracking error in eq. (C.117) tends to zero after $t_1$; finite-time Grönwall comparison therefore carries the evolving rows into the same neighborhood. At its selected vertex let $\zeta_i = |p_i - p_i^M|$. By the derivative computation above and continuity, a sufficiently small neighborhood has $(1 - \zeta_i)|\mathfrak{g}_i(p)| \ge \nu/2$, where $\nu = 1 - \mathsf{r}_E > 0$, and the fitness has the sign directed toward that vertex. Thus

$$D^+\zeta_i \le -\tfrac{\nu}{2}\zeta_i + \mathrm{TV}(q_t, \mu_{p(t)})$$

in slow time. The neighborhood remains invariant once the tracking floor is small enough, and $\zeta_i$ decays exponentially to a vanishing floor.

Choose any $\mathrm{C}_{\mathrm{T}} > 1/\alpha$, independently of the mean, and then choose the local radius so that its escape coefficient is less than $\mathrm{C}_{\mathrm{T}}$. The time remaining after the finite transit tends to infinity

uniformly on the slow-time window $[\mathrm{C_T}\log(1/\varepsilon), 2\mathrm{C_T}\log(1/\varepsilon)]$. The row distances therefore vanish uniformly on that window. Every normalized sheet weight is at most the matching weight one, so

$$\mu_p(M) \geq \pi_M \geq 1 - \sum_i \zeta_i, \qquad q_t(S \neq M) \leq \sum_i \zeta_i + \mathrm{TV}(q_t, \mu_{p(t)}).$$

The last total variation tends to zero by eq. (C.117); its integrated contraction exponent at the beginning of the window grows at least as a positive constant times $\varepsilon^{-1}$. This proves the first conclusion for every mean outside the null set in eq. (C.106). For the frozen arm, eq. (C.99) and rotation symmetry prove the second conclusion at every $\mathrm{T} > 0$ and every mean. The convergence is pointwise in the fixed mean; no uniform bound near the zero set of the imprint has been used. □

## C.10 Proof of Joint Selection Along a Schedule

*Proof of Corollary 14.6.* The saddle estimate eq. (C.114) gives $\alpha(\mathrm{T}) \geq 7/5$ throughout the stated temperature range. The proof of Theorem 14.5 permits any window constant greater than $1/\alpha(\mathrm{T})$, independently of the initial mean. Thus the common choice $2$ is valid. Write

$$\tau_\varepsilon = 2\varepsilon^{-1}\log(1/\varepsilon), \qquad W_\varepsilon = [\tau_\varepsilon, 2\tau_\varepsilon], \qquad 0 < \varepsilon < e^{-1}.$$

For a perfect matching $M$, define

$$\nu_{\mathrm{T}}^M = \mathcal{N}(0, \mathrm{T}\mathsf{A}_M^{-1}) \otimes \delta_M, \qquad \mu_{\mathrm{T}} = \mu_{\mathrm{T},1/2}.$$

Here $\mu_{\mathrm{T},p}$ is the frozen-row law used in the preceding proof, with its temperature now displayed explicitly. The inverse exists because $\mathsf{A}_M \succeq \mathrm{T}I$. Let $M_{\mathrm{T}}(\mathsf{m}_0)$ be the matching selected by the sign of $\mathsf{j}_{\mathrm{T}}(\mathsf{m}_0)$, choosing either matching on its null zero set.

First record the full one-time consequence of the fixed-temperature proof. Conditional on assignment $M$, the law $\mu_{\mathrm{T},p}$ has the Gaussian factor of $\nu_{\mathrm{T}}^M$, independently of $p$. The normalized sheet weights are at most the matching weight one. For $\zeta_i = |p_i - p_i^M|$, therefore,

$$\mu_{\mathrm{T},p}(M) \geq \pi_M \geq 1 - \sum_i \zeta_i, \qquad \mathrm{TV}(\mu_{\mathrm{T},p}, \nu_{\mathrm{T}}^M) = 1 - \mu_{\mathrm{T},p}(M) \leq \sum_i \zeta_i.$$

The row convergence and eq. (C.117) imply

$$\sup_{t \in W_\varepsilon} \mathrm{TV}(q_t^{\mathrm{T},\varepsilon,\mathsf{m}_0}, \nu_{\mathrm{T}}^{M_{\mathrm{T}}(\mathsf{m}_0)}) \longrightarrow 0$$

at fixed temperature outside the imprint's null zero set. Moreover eq. (C.116) itself gives convergence of the learned relative entropy to zero, uniformly on this window. For frozen rows, the entropy identity and contraction give

$$D_{\mathrm{KL}}(q_t^{\mathrm{T},0,\mathsf{m}_0} \| \mu_{\mathrm{T}}) \leq e^{-\lambda(\mathrm{T},1/2)t} D_{\mathrm{KL}}(q_0^{\mathrm{T},0,\mathsf{m}_0} \| \mu_{\mathrm{T}}).$$

The initial entropy is finite by the Gaussian density bounds, and $\lambda(\mathrm{T}, 1/2) > 0$. Frozen entropy and total variation therefore tend to zero on the window for every fixed mean.

Use the explicit entropy majorants as follows. Let

$$G_{\mathrm{T}} = \sup_{p\in[0,1]^4} \|\mathfrak{g}(p)\|_1, \qquad \ell_{\mathrm{T}}(\varepsilon) = \frac{\mathrm{c}_{\mathrm{T}}}{\mathrm{K}_{\mathrm{T}} + 8\log 2 + 32\log(1/\varepsilon)}, \qquad \lambda_u(\mathrm{T}) = \frac{\mathrm{c}_{\mathrm{T}}}{\mathrm{K}_{\mathrm{T}} + 8\log 2}.$$

The first constant is finite. The row lower bound and eq. (C.112) make $\ell_{\mathrm{T}}(\varepsilon)$ a rate lower bound throughout $t \le 2\tau_\varepsilon$, and $\lambda_u(\mathrm{T})$ a rate lower bound at uniform rows. The common initial entropy equals

$$H_0(\mathrm{T}, \mathsf{m}_0) = \log \mathsf{Z}(1/2) - 6 + \frac{\operatorname{tr} \overline{\mathsf{A}} + \mathsf{m}_0^\top \overline{\mathsf{A}} \mathsf{m}_0}{2\mathrm{T}}, \qquad \overline{\mathsf{A}} = \frac{1}{16}\sum_\sigma \mathsf{A}_\sigma.$$

Indeed the log density ratio is $\log \mathsf{Z}(1/2) - |\mathbf{x} - \mathsf{m}_0|^2/2 + \mathbf{x}^\top \mathsf{A}_\sigma \mathbf{x}/(2\mathrm{T})$; taking its initial expectation proves the formula. Thus $H_0$ is continuous and nonnegative. Uniformly on $W_\varepsilon$, the entropy tracking estimate gives

$$\begin{aligned} D_{\mathrm{KL}}(q_t^{\mathrm{T},\varepsilon,\mathsf{m}_0} \| \mu_{\mathrm{T},p(t)}) &\le B_L(\mathrm{T}, \varepsilon, \mathsf{m}_0) \\ &:= \left[ e^{-\ell_{\mathrm{T}}(\varepsilon)\tau_\varepsilon/2} \sqrt{H_0} + \frac{\varepsilon G_{\mathrm{T}}}{\sqrt{2}\ell_{\mathrm{T}}(\varepsilon)} \right]^2, \\ D_{\mathrm{KL}}(q_t^{\mathrm{T},0,\mathsf{m}_0} \| \mu_{\mathrm{T}}) &\le B_F(\mathrm{T}, \varepsilon, \mathsf{m}_0) := e^{-\lambda_u(\mathrm{T})\tau_\varepsilon} H_0. \end{aligned}$$

These majorants are continuous in positive mobility and mean and tend to zero at fixed temperature and mean as $\varepsilon \downarrow 0$.

Let $\mathcal{E}_{\mathrm{T}}(\varepsilon, \mathsf{m}_0)$ be the minimum of one and the maximum of the following five nonnegative quantities:

$$\begin{aligned} &\sup_{t\in W_\varepsilon} \max_i |p_i(t) - p_i^{M_{\mathrm{T}}(\mathsf{m}_0)}|, \\ &\sup_{t\in W_\varepsilon} \mathrm{TV}(q_t^{\mathrm{T},\varepsilon,\mathsf{m}_0}, \nu_{\mathrm{T}}^{M_{\mathrm{T}}(\mathsf{m}_0)}), \\ &B_L(\mathrm{T}, \varepsilon, \mathsf{m}_0), \\ &\sup_{t\in W_\varepsilon} \mathrm{TV}(q_t^{\mathrm{T},0,\mathsf{m}_0}, \mu_{\mathrm{T}}), \\ &B_F(\mathrm{T}, \varepsilon, \mathsf{m}_0). \end{aligned}$$

The preceding estimates prove $\mathcal{E}_{\mathrm{T}}(\varepsilon, \mathsf{m}_0) \to 0$ for almost every fixed mean at each fixed temperature.

We will take a supremum over mobility, so measurability must be checked. Fix a temperature and restrict the positive mobility and the mean to compact sets, with the mobility bounded away from zero. The corresponding horizons are bounded, and eq. (C.97) bounds every row away from zero and one. On this row set the posterior redraw probabilities are uniformly Lipschitz in the rows, independently of the configuration. Maximally couple the initial laws and then use common Brownian motions and proposal clocks until the first different redraw. The law difference is bounded by the initial total variation plus a constant times the integral of row differences. Subtracting the learning equations bounds the row difference by integrals of row and law differences and by the

mobility difference. Grönwall's inequality proves total-variation continuity in the initial mean and positive mobility on each bounded time interval. Time continuity follows from eq. (C.107).

Parametrizing the window by $t = u\tau_\varepsilon$, $1 \le u \le 2$, proves continuity in positive mobility of each row or total-variation window supremum. The explicit entropy majorants are continuous by their formulas. The matching as a function of the mean is Borel because its imprint is analytic. Thus the error is measurable in the mean, and its supremum over an interval of positive mobilities equals its supremum over rational mobilities. No uniform continuity at zero mobility or separate entropy-continuity assertion is needed.

Fix the sequence $\mathrm{T}_n \downarrow 0$ in the stated temperature interval and $0 < \eta_n < 1$ with $\eta_n \to 0$. For the prescribed exponent $a > 7/2$, put $h_n = \mathrm{T}_n^a$. The following argument in fact permits arbitrary positive caps $h_n$. On the mean space $\mathbb{R}^{12}$, introduce the auxiliary probability measure $\rho = \mathcal{N}(0, I_{12})$. Its density is strictly positive; it introduces no assumption that the physical mean is random. Set

$$F_n(r, \mathsf{m}_0) = \sup_{0<\varepsilon\le r} \mathcal{E}_{\mathrm{T}_n}(\varepsilon, \mathsf{m}_0).$$

The measurable functions $F_n$ are bounded by one. At fixed $n$, they tend to zero for $\rho$-almost every mean as $r \downarrow 0$. Dominated convergence yields

$$\int F_n(r, \mathsf{m}_0) d\rho(\mathsf{m}_0) \longrightarrow 0.$$

Choose $d_n = 2^{-k_n}$ recursively, where $k_n$ is the smallest positive integer satisfying

$$d_n \le \min(e^{-1}, 2^{-n}, h_n, d_{n-1}/2), \qquad \int F_n(d_n, \mathsf{m}_0) d\rho(\mathsf{m}_0) \le \eta_n 2^{-n},$$

with the $d_{n-1}/2$ condition omitted for $n = 1$. The preceding limit proves existence. This defines deterministic thresholds independent of the mean; it does not give numerical values or an explicit power law.

Let $B_n = \{\mathsf{m}_0 : F_n(d_n, \mathsf{m}_0) > \eta_n\}$. Markov's inequality and countable subadditivity give

$$\rho(B_n) \le 2^{-n}, \qquad \rho\left(\bigcup_{n\ge N} B_n\right) \le 2^{1-N}.$$

The decreasing intersection of these tail unions consequently has measure zero. Add the countable union of the imprint zero sets. The resulting exceptional set is Lebesgue-null because $\rho$ and Lebesgue measure have the same null sets. Outside it, for every sufficiently large $n$, all $0 < \varepsilon \le d_n$ satisfy $F_n(d_n, \mathsf{m}_0) \le \eta_n$. Since $\eta_n < 1$, clipping has no effect on this inequality: the row and total-variation errors and both entropy majorants are at most $\eta_n$. The majorants bound the actual entropies, giving all five stated conclusions.

The frozen matching probability is, exactly,

$$\mu_{\mathrm{T}}(M) = \frac{1}{2 + 8\mathsf{r}_E + 4\mathsf{r}_N + 2\mathsf{r}_C} = \frac{1}{2D_0(\mathrm{T})} = \frac{1}{2} + O(\mathrm{T}).$$

Indeed $\mathsf{r}_E, \mathsf{r}_N = O(\mathrm{T})$ and $\mathsf{r}_C = O(\mathrm{T}^{3/2})$, as proved from the sheet determinants. Pushforward cannot increase total variation. If $q^S$ denotes the assignment marginal, the two established total-variation estimates therefore imply, uniformly on $W_\varepsilon$,

$$\left| \mathrm{TV}(q_t^{\mathrm{T}_n,\varepsilon,\mathsf{m}_0,S}, q_t^{\mathrm{T}_n,0,\mathsf{m}_0,S}) - \left(1 - \frac{1}{2D_0(\mathrm{T}_n)}\right) \right| \le 2\eta_n.$$

Here the comparison distance is exactly $\mathrm{TV}(\delta_M, \mu^S_{\mathrm{T}_n}) = 1 - \mu_{\mathrm{T}_n}(M)$, for either matching. The assignment-law contrast thus tends to $1/2$, even if the selected matching changes with $n$.

Finally, defining $\mathrm{T}(\varepsilon) = \mathrm{T}_n$ on $d_{n+1} < \varepsilon \le d_n$ gives a single stepwise path with both parameters tending to zero. The proved bounds hold for all sufficiently small mobilities on this path. Arbitrary mobility caps are admissible; for example $h_n = \mathrm{T}_n^n$ forces $\varepsilon_n/\mathrm{T}_n^a \to 0$ for each fixed $a > 0$ whenever $\varepsilon_n \le d_n$. This imposes an upper bound and does not certify a prescribed equality between temperature and mobility.

Nothing in this argument makes the imprint sign eventually constant as temperature vanishes. Its fixed-temperature nonvanishing and null zero set give no quantitative lower bound at a fixed mean in that limit. Accordingly the matching is allowed to depend on $n$, and the result concerns complete one-time laws, not complete histories or conditional first-formation events. □

## An explicit temperature bound for tracking

**Proposition C.15 (Temperature-dependent entropy tracking).** In the four-cycle of theorem 14.5, let $0 < \mathrm{T} \le 1/100$, $p \in (0,1)^4$, and $p_{\min} = \min_i\{p_i, 1-p_i\}$. Define

$$\mathrm{K}_0 = 210\sqrt{2\pi/5}, \qquad \mathrm{c}_* = \min\left\{\frac{2}{1 + 4\mathrm{K}_0 e^{13/8}}, \frac{1}{\mathrm{K}_0 e^{13/4}}\right\}.$$

The entropy production of the full frozen law satisfies

$$\mathscr{D}_p(f) \ge \frac{\mathrm{c}_* \mathrm{T}^{5/2}}{4 + 3\log(1/\mathrm{T}) + 8\log(1/p_{\min})} \mathrm{Ent}_{\mu_p}(f).$$

Write $\varepsilon = \epsilon_{\mathrm{L}}$, $t_{\mathrm{L}} = 2\varepsilon^{-1}\log(1/\varepsilon)$, and

$$\mathrm{D}(\mathrm{T},\varepsilon) = 4 + 3\log(1/\mathrm{T}) + 8\log 2 + 32\log(1/\varepsilon), \qquad \ell = \frac{\mathrm{c}_* \mathrm{T}^{5/2}}{\mathrm{D}(\mathrm{T},\varepsilon)}.$$

For $0 < \varepsilon < e^{-1}$ and $0 \le t \le 2t_{\mathrm{L}}$,

$$\mathrm{TV}(q_t, \mu_{p(t)}) \le e^{-\ell t/2}\sqrt{\frac{13(12 + |\mathsf{m}_0|^2)}{4\mathrm{T}}} + \frac{2\varepsilon \mathrm{D}(\mathrm{T},\varepsilon)}{\mathrm{c}_* \mathrm{T}^{7/2}}. \tag{C.121}$$

In particular, for $\varepsilon = \mathrm{T}^{a}$ with any $a > 7/2$, this total variation tends to zero uniformly on $[t_{\mathrm{L}}, 2t_{\mathrm{L}}]$ and on bounded sets of initial means. This conclusion compares one-time laws; it supplies no imprint lower bound. [ESTABLISHED]

**Proof.** Use $u = (\bar{k} + \bar{m})/\sqrt{2}$, $v = (\bar{k} - \bar{m})/\sqrt{2}$, and $\omega$, and put $R = [D, -\sqrt{2}I]$. The common kernel $K = \ker R \oplus \ker D$ has dimension five. Every $\mathsf{H}_{\sigma}$ vanishes there, and the reference precision preserves $K$ and $N = K^{\perp}$. Its restrictions have eigenvalues one and five. The passive process on $K$ is therefore an independent Gaussian Ornstein–Uhlenbeck process with entropy contraction at least $2\mathrm{T}$. The active process on $N$, together with the assignment, is autonomous: all redraw costs depend only on its residuals. Rescaling its coordinate by $\mathrm{T}^{-1/2}$ gives unit diffusion and drift matrix $a_{\sigma} = (\mathsf{H}_{\sigma} + \mathrm{T}\Omega)|_{N}$. Since every edge multiplicity is at most two,

$$\mathrm{T}I \preceq a_{\sigma} \preceq 13I.$$

Indeed $RR^{\top} = DD^{\top} + 2I$ has eigenvalues $2, 4, 4, 6$, while $D^{\top}D$ has eigenvalues $0, 2, 2, 4$, so $\|\mathsf{H}_{\sigma}\| \le 12$.

Divide the four determinants in eq. (C.113) by the passive determinant $5\mathrm{T}^{5}$. Their positive coefficients, evaluated at $\mathrm{T} = 1/100$ for upper bounds, give the active bounds

$$5120\mathrm{T}^{3} \le d_{M} \le 5262\mathrm{T}^{3}, \quad 896\mathrm{T} \le d_{E} \le 1017\mathrm{T},$$
$$896\mathrm{T} \le d_{N} \le 2175\mathrm{T}, \qquad 3072 \le d_{C} < 5120.$$

Thus $w_{M}$ and $w_{C}$ are the largest and smallest active sheet weights, and $w_{M} \le \mathrm{T}^{-3/2}/\sqrt{5120}$, $w_{C}/w_{M} \ge \mathrm{T}^{3/2}$. The finite-label constant in the entropy proof consequently satisfies

$$L_{p} \le 4 + 3\log(1/\mathrm{T}) + 8\log(1/p_{\min}).$$

Apply the Gaussian-sheet and prior-weighted-ball argument of eqs. (C.111) and (C.112) in the seven active coordinates, using the ball of radius $1/2$, whose volume is $\pi^{3}/840$. Its Gaussian comparison factor is at most $4\mathrm{K}_{0}e^{13/8}\mathrm{T}^{-3/2}$, because

$$\frac{(2\pi)^{7/2}}{(\pi^{3}/840)\sqrt{5120}} = \mathrm{K}_{0}.$$

More explicitly, the two energy coefficients in that argument obey

$$C_{c} \le (2 + 8\mathrm{K}_{0}e^{13/8})L_{p}\mathrm{T}^{-5/2}, \qquad C_{j} \le 4\mathrm{K}_{0}e^{13/4}L_{p}\mathrm{T}^{-5/2}.$$

The continuous production equals four times its quadratic form and the jump production bounds four times its quadratic form. Hence the active rate is at least $\mathrm{c}_{*}\mathrm{T}^{5/2}/L_{p}$. Entropy decomposes over the independent passive and active factors. Cauchy–Schwarz bounds the energy of a marginal by the corresponding integrated energy, so the full rate is at least the smaller of the two rates. The

displayed active lower bound is less than $2\mathrm{T}$, proving the full-state inequality. The form-domain extension uses the positive regularization and cutoffs already used for eq. (C.112).

A one-receiver flip joins only the class pairs $M$–$E$, $N$–$E$, and $C$–$E$. The determinant bounds above make every ratio of such adjacent weights at most $1/\mathrm{T}$. Writing the evidence as $p_i Z_i^+ + (1-p_i)Z_i^-$, the ratio $Z_i^+/Z_i^-$ is between the extremal adjacent ratios. Therefore $|\partial_{p_i} \log Z| \le 1/\mathrm{T}$, and $\|\mathsf{g}\|_1 \le 4/\mathrm{T}$. The row lower bound gives the rate floor $\ell$ on the stated horizon. The initial entropy is at most $13(12+|\mathsf{m}_0|^2)/(2\mathrm{T})$: in the Gaussian density ratio, $\mathsf{B}_\sigma \preceq 13I/\mathrm{T}$, $\mathsf{Z} \le 1$, and the remaining quadratic term has nonpositive expectation. Insert these bounds in eq. (C.116) and use Pinsker's inequality to obtain eq. (C.121). For $\varepsilon = \mathrm{T}^a$, its second term is $O(\mathrm{T}^{a-7/2}\log(1/\mathrm{T}))$, while $\ell t_{\mathrm{L}}$ grows as a positive constant times $\mathrm{T}^{5/2-a}$. The exponential therefore dominates the initial $\mathrm{T}^{-1/2}$ factor, proving the last assertion. □

## C.11 Coherence and Persistence of the Selected Pairing

**Proposition C.16 (Coherence and persistence after pairing selection).** In the four-cycle of theorem 14.5, let $0 < \mathrm{T} \le 1/100$, let $M$ be one matching, and suppose $[t_0, t_0+\mathsf{h}] \subseteq [t_{\mathrm{L}}, 2t_{\mathrm{L}}]$, with $\mathsf{h} > 0$. Write $\nu_{\mathrm{T}}^M = \mathcal{N}(0, \mathrm{T}\mathsf{A}_M^{-1}) \otimes \delta_M$. Assume the learned and frozen recognition laws satisfy

$$\sup_{t_0 \le t \le t_0+\mathsf{h}} \mathrm{TV}(q_t^{\mathrm{L}}, \nu_{\mathrm{T}}^M) \le \delta, \qquad \mathrm{TV}(q_{t_0}^0, \mu_{\mathrm{T},1/2}) \le \delta.$$

Let $\mathsf{O}_{\mathrm{T},\mathsf{h}}^M$ be the stationary Ornstein–Uhlenbeck path law with drift $-\mathsf{A}_M\mathbf{x}$, diffusion covariance $2\mathrm{T}I$, and constant assignment $M$. For the configuration and assignment path laws on this interval,

$$\mathrm{TV}(\mathsf{P}_{t_0,\mathsf{h}}^{\mathrm{L}}, \mathsf{O}_{\mathrm{T},\mathsf{h}}^M) \le (\mathsf{h}+2)\delta, \tag{C.122}$$

$$\mathrm{TV}\left(\mathsf{P}_{t_0,\mathsf{h}}^0, \tfrac{1}{2}\mathsf{O}_{\mathrm{T},\mathsf{h}}^{M_0} + \tfrac{1}{2}\mathsf{O}_{\mathrm{T},\mathsf{h}}^{M_1}\right) \le \delta + 16\mathrm{T} + 31\mathrm{T}\mathsf{h}. \tag{C.123}$$

The frozen arm keeps its original uniform rows.

For resolution $\mathsf{a} > 0$, let $\mathscr{E}_M(\mathsf{a},\mathsf{h})$ require throughout the interval that the assignment is $M$, the stacked residual norm $\sqrt{(r^b)^2+(r^m)^2}$ on its two edges is at most $\mathsf{a}$, and that norm on each exterior edge exceeds $\mathsf{a}$. Define

$$\mathsf{b}(\mathrm{T},\mathsf{a},\mathsf{h}) = 8\sqrt{2/\pi}\,\mathsf{a} + (8+3400\mathsf{h})\mathrm{T}/\mathsf{a}^2.$$

Then

$$\Pr_{\mathrm{L}}(\mathscr{E}_M^c) \le (\mathsf{h}+2)\delta + \mathsf{b}(\mathrm{T},\mathsf{a},\mathsf{h}). \tag{C.124}$$

In particular, the choice $\mathsf{a} = \mathrm{T}^{1/4}$, $\mathsf{h} = \mathrm{T}^{-1/4}$, and $\delta \le \mathrm{T}$ makes this probability tend to zero. The matching-window label, equal to $M_j$ when $\mathscr{E}_{M_j}$ holds and to an additional symbol otherwise, converges to $\delta_M$ in the learned arm and to the equal matching mixture in the frozen arm.

At one time, define $\mathfrak{P}_{\mathsf{a}}(\mathbf{x})$ as the connected-component partition of edges whose stacked residual norm $\sqrt{(r^b)^2+(r^m)^2}$ is at most $\mathsf{a}$. For times when both one-time hypotheses above hold, put

$$\mathsf{e}(\mathrm{T},\mathsf{a}) = 4\sqrt{2/\pi}\mathsf{a} + 4\mathrm{T}/\mathsf{a}^2.$$

The learned and frozen partition laws satisfy

$$\left|\mathrm{TV}(q_t^{\mathrm{L}}\mathfrak{P}_{\mathsf{a}}, q_t^{0}\mathfrak{P}_{\mathsf{a}}) - \tfrac{1}{2}\right| \le 2\delta + 16\mathrm{T} + 2\mathsf{e}(\mathrm{T},\mathsf{a}). \tag{C.125}$$

Thus the contrast persists for the coherent partition, not only the source label.

These conclusions apply along corollary 14.6 with tolerance $\delta \le \mathrm{T}$, mobility cap $\epsilon_{\mathrm{L}} \le \mathrm{T}^2$, and $t_0 = t_{\mathrm{L}}$. The matching can vary along that schedule. The internal restoring rates are $8+\mathrm{T}$ and $4+5\mathrm{T}$, on an event whose probability tends to one. The resolution therefore has restoring ratio at least $4\mathsf{a}^2/\mathrm{T} \to \infty$. If the rows are also within $\delta$ of their matching vertex, the stationary law with those rows frozen realizes the same persistence event with probability at least $1-4(\mathsf{h}+2)\delta-\mathsf{b}$. This late coherence is restored equilibrium coherence; it supplies no entrance or first-formation assertion. [ESTABLISHED]

*Proof.* First retain the proposal mechanism in the comparison. Put $r(t) = q_t^{\mathrm{L}}(S \neq M)$, and let $b_{\neq M}(t,\mathbf{x})$ be the probability of proposing a different assignment at a clock ring. Since the clock has rate one,

$$r'(t) = \mathbb{E}b_{\neq M}(t,\mathbf{X}_t) - r(t).$$

If $N_{\neq M}$ counts proposals to an assignment other than $M$ on the interval, its compensator gives

$$\mathbb{E}N_{\neq M} = r(t_0+\mathsf{h}) - r(t_0) + \int_{t_0}^{t_0+\mathsf{h}} r(t)dt. \tag{C.126}$$

Consequently $\Pr(\exists t \in [t_0, t_0+\mathsf{h}] : S_t \neq M) \le (\mathsf{h}+1)\delta$. Couple the initial law maximally with $\nu_{\mathrm{T}}^M$, and use the same future Brownian motion for the native process and the stationary matching process. Their paths agree until an initial mismatch or a proposal to another assignment. The union bound and eq. (C.126) prove eq. (C.122). This argument counts proposals of the native process and does not condition its diffusion on an assignment history.

For the frozen comparison, use the normalized weights $1, \mathsf{r}_E, \mathsf{r}_N, \mathsf{r}_C$ from appendix C.9. The determinant bounds in the proof of proposition C.15 give

$$\mathsf{r}_E, \mathsf{r}_N \le \tfrac{5}{2}\mathrm{T}, \qquad \mathsf{r}_C \le \tfrac{3}{2}\mathrm{T}^{3/2}.$$

The invariant mass $\kappa_{\mathrm{T}}$ of nonmatching assignments is therefore at most $16\mathrm{T}$. Conditional on either matching, its configuration law is the Gaussian factor of $\nu_{\mathrm{T}}^M$. The invariant law is consequently within $\kappa_{\mathrm{T}}$ of the equal mixture of these two laws.

Detailed balance bounds the stationary flux from $M$ to all nonmatching assignments by $\kappa_{\mathrm{T}}$. For the other matching $M'$, write $\mu_M(\mathbf{x})$ for the joint invariant density on the matching sheet. Pointwise,

$$\mu_M b_{M'} = \frac{\mu_M \mu_{M'}}{\sum_S \mu_S} \le \tfrac{1}{2}\sqrt{\mu_M \mu_{M'}}.$$

Since $(\mathsf{A}_M + \mathsf{A}_{M'})/2$ is the oriented-cycle matrix, Gaussian integration gives $\int \sqrt{\mu_M \mu_{M'}} = \mu_{\mathrm{T},1/2}(M)\mathsf{r}_C$. Dividing the two flux bounds by the matching mass yields

$$\mathbb{E}_{\nu_{\mathrm{T}}^M} b_{\neq M} \le 8\mathsf{r}_E + 4\mathsf{r}_N + \tfrac{5}{2}\mathsf{r}_C \le 31\mathrm{T}.$$

Thus the stationary flux out of the two matching sheets together is at most $31\mathrm{T}$. Couple the stationary native process at its initial time with the equal matching mixture and use common Brownian increments until the first departure from its coupled matching. Its failure probability is at most $16\mathrm{T} + 31\mathrm{Th}$. The frozen native process started at $q_{t_0}^0$ differs from its stationary path law by at most its initial total variation $\delta$, because the frozen transition kernel is the same. This proves eq. (C.123).

It remains to bound residuals on the comparison path. The four normal coordinates $z_g^{\pm} = n_g^{\pm} \cdot \mathbf{X}$, $g \in M$, from eq. (C.94) have rates $8 + \mathrm{T}$ and $4 + 5\mathrm{T}$. Their stationary variances are at most $\mathrm{T}/4$, and

$$r_g^b = \sqrt{2} z_g^+ + z_g^-, \qquad r_g^m = \sqrt{2} z_g^+ - z_g^-.$$

For a scalar Ornstein–Uhlenbeck noise starting at zero, integration by parts in its stochastic convolution gives

$$Z_t = \sqrt{2\mathrm{T}}\left[W_t - \lambda \int_0^t e^{-\lambda(t-u)} W_u du\right], \qquad \mathbb{E}\sup_{t \le \mathsf{h}} |Z_t|^2 \le 32\mathrm{Th}.$$

The last inequality uses $\sup|Z_t| \le 2\sqrt{2\mathrm{T}}\sup|W_t|$ and the Brownian maximal estimate proved in the matching-episode argument of appendix C.7. For the four-dimensional stationary normal vector $z$, it follows that

$$\mathbb{E}\sup_{t \le \mathsf{h}} |z_t|^2 \le 2\mathrm{T} + 256\mathrm{Th}.$$

Each selected stacked residual norm is at most $2|z_t|$, since $(r_g^b)^2 + (r_g^m)^2 = 4(z_g^+)^2 + 2(z_g^-)^2$. Thus the probability that any internal stacked norm exceeds $\mathsf{a}$ at any time is at most $(8 + 1024\mathsf{h})\mathrm{T}/\mathsf{a}^2$.

Every exterior link coordinate is absent from $H_M$, and its block in $\mathsf{A}_M$ is the scalar $\mathrm{T}$. It is therefore an independent $\mathcal{N}(0,1)$ coordinate in the stationary Gaussian. Each exterior belief or model residual has variance at least one. The probability that any of these four residuals initially has absolute value at most $2\mathsf{a}$ is at most $8\sqrt{2/\pi}\,\mathsf{a}$.

To control their increments, diagonalize the symmetric positive matrix $\mathsf{A}_M$. In a scalar eigencoordinate with rate $\lambda > 0$, stationarity and the preceding noise estimate imply

$$\mathbb{E}\sup_{t \le \mathsf{h}} |X_t - X_0|^2 \le \frac{2\mathrm{T}}{\lambda}(1 - e^{-\lambda \mathsf{h}})^2 + 64\mathrm{Th} \le 66\mathrm{Th},$$

since $(1-e^{-v})^2 \le v$ for $v \ge 0$. Summing twelve coordinates gives

$$\mathbb{E}\sup_{t\le \mathsf{h}} |\mathbf{X}_t - \mathbf{X}_0|^2 \le 792\mathrm{Th}. \tag{C.127}$$

Every residual covector has norm $\sqrt{3}$. Except with probability $2376\mathrm{Th}/\mathsf{a}^2$, all exterior increments are therefore at most $\mathsf{a}$ in absolute value. The initial small-ball estimate and the two maximal estimates prove $\mathsf{O}^{M}_{\mathrm{T},\mathsf{h}}(\mathscr{E}^{c}_{M}) \le \mathsf{b}$. Together with eq. (C.122), this proves eq. (C.124). The frozen path comparison proves the stated matching-window label limit whenever $\mathsf{b} \to 0$, $\mathsf{h}\delta \to 0$, and $\mathrm{Th} \to 0$, in particular for the displayed choices.

For a single time no increment estimate is needed. The normal second moment is at most $\mathrm{T}$, so the internal failure probability is at most $4\mathrm{T}/\mathsf{a}^2$. The exterior small-ball bound at radius $\mathsf{a}$ is $4\sqrt{2/\pi}\mathsf{a}$. Thus the partition induced by $\nu_{\mathrm{T}}^{M_j}$ differs from the matching partition $\mathfrak{P}_j$ with probability at most $\mathsf{e}$. The learned partition law is within $\delta + \mathsf{e}$ of $\delta_{\mathfrak{P}_M}$; the frozen one is within $\delta + 16\mathrm{T} + \mathsf{e}$ of $(\delta_{\mathfrak{P}_0} + \delta_{\mathfrak{P}_1})/2$. These matching partitions are distinct, so the triangle inequality for total variation proves eq. (C.125).

The schedule in corollary 14.6 permits tolerance $\mathrm{T}$ and mobility cap $\mathrm{T}^2$. Then $\mathrm{T}^{-1/4} \le t_{\mathrm{L}}$ for sufficiently small $\mathrm{T}$, so the following persistence window lies within the selection window. All errors just bounded tend to zero. The normal rates establish the restoring ratio on the native event. Finally, if $\max_i |p_i - p_i^M| \le \delta$, $\mathrm{TV}(\mu_{\mathrm{T},p}, \nu_{\mathrm{T}}^{M}) = 1 - \mu_{\mathrm{T},p}(M) \le 4\delta$. Applying the same occupation and coupling argument to this stationary frozen-row process gives the stationary lower bound in the statement. The persistence event is therefore already typical at equilibrium for those rows and carries no first-entrance claim. □

### C.11.1 The first contrast on the original episode window

**Proposition C.17 (A first-block null under slower source learning).** In the model of theorem 14.5, let $0 < \mathrm{T} \le 10^{-16}$, and let the initial law be $\mathcal{N}(\mathsf{m}_0, I_{12})$ times uniform assignments. The learned and frozen arms have the same initial mean and differ only in the source-learning mobility $\epsilon_{\mathrm{L}}$. Set

$$H_{\mathrm{T}} = 2\mathrm{T}^{-1/8}.$$

Keep the contrast event of proposition 14.4 with its regime parameter equal to $\mathrm{T}^{1/8}$: an eligible proper connected block has an internal belief relative entropy exceeding $1/4$ initially, and for a duration-one interval inside $[0, H_{\mathrm{T}}]$ all its internal belief and model relative entropies are at most $\sqrt{\mathrm{T}}$, while its boundary belief relative entropies exceed $1/2$. For each block take the infimum of its qualifying completion times, with value infinity when there is no such interval. Let $K_{\mathrm{T}}$ be the block with least such infimum, using one fixed tie rule in both arms, and give it a cemetery value when no block qualifies. Write $E_{\mathrm{T}}$ for the event that its value is not the cemetery value. The infimum need not itself be a qualifying completion time.

Define the positive, temperature-independent constant

$$\mathsf{c}(\mathsf{m}_0) = \frac{e^{-1/4}\varrho_\circ^{12}}{737280}\exp\left[-\frac{1}{2}(|\mathbf{x}^\circ - \mathsf{m}_0| + \varrho_\circ)^2\right], \qquad \mathsf{d}_{\mathrm{T}} = \min\{1, 8\epsilon_{\mathrm{L}}\mathrm{T}^{-1/4}\}.$$

Then

$$\Pr_0(E_{\mathrm{T}}) \geq \mathsf{c}(\mathsf{m}_0), \qquad \Pr_{\epsilon_{\mathrm{L}}}(E_{\mathrm{T}}) \geq \mathsf{c}(\mathsf{m}_0) - \mathsf{d}_{\mathrm{T}}.$$

If $\mathsf{d}_{\mathrm{T}} < \mathsf{c}(\mathsf{m}_0)$, the two conditional block laws satisfy

$$\mathrm{TV}\left(\mathrm{Law}_{\epsilon_{\mathrm{L}}}(K_{\mathrm{T}} \mid E_{\mathrm{T}}), \mathrm{Law}_0(K_{\mathrm{T}} \mid E_{\mathrm{T}})\right) \leq \frac{\mathsf{d}_{\mathrm{T}}}{\mathsf{c}(\mathsf{m}_0)}.$$

In particular their difference vanishes for every fixed Gaussian mean whenever $\mathrm{T} \downarrow 0$ and $\epsilon_{\mathrm{L}}\mathrm{T}^{-1/4} \to 0$. This includes the caps $\epsilon_{\mathrm{L}} \leq \mathrm{T}_n^a$, $a > 7/2$, in corollary 14.6. [ESTABLISHED]

*Proof.* First prove a uniform comparison of the complete native path laws on any finite horizon $H$. The learned rows are deterministic and satisfy

$$|p_i(t) - 1/2| \leq \epsilon_{\mathrm{L}} t/2.$$

At a common configuration, write $A, B > 0$ for a receiver's likelihood factors for its two sources. Its posterior positive-source probabilities in the two arms are

$$b(p) = \frac{pA}{pA + (1-p)B}, \qquad b(1/2) = \frac{A}{A+B}.$$

Direct subtraction gives

$$|b(p) - b(1/2)| = \frac{AB|2p-1|}{(pA + (1-p)B)(A+B)} \leq |2p-1|.$$

Indeed the denominator is $pA^2 + AB + (1-p)B^2 \geq AB$. This bound is independent of temperature and the likelihood ratio, without an interior-row restriction.

Use identical initial samples, a common Brownian motion, and the same rate-one proposal clock in the two arms. Until the first different redraw, their configurations and assignments coincide: a held-sheet drift and its noise do not depend on the prior rows. At each ring, couple each receiver's two Bernoulli laws maximally. The union bound and the row estimate bound the probability of any different receiver choice at time $t$ by

$$\sum_i |b_i(p_i(t)) - b_i(1/2)| \leq 4\epsilon_{\mathrm{L}} t.$$

The compensator of the rate-one clock, stopped at the first mismatch, bounds its probability before $H$ by $\int_0^H 4\epsilon_{\mathrm{L}} t dt = 2\epsilon_{\mathrm{L}} H^2$. Consequently,

$$\mathrm{TV}\left(\mathrm{Law}_{\epsilon_{\mathrm{L}}}(\mathbf{X}_{[0,H]}, S_{[0,H]}), \mathrm{Law}_0(\mathbf{X}_{[0,H]}, S_{[0,H]})\right) \leq \min\{1, 2\epsilon_{\mathrm{L}} H^2\}. \tag{C.128}$$

The bound follows because the complete paths are identical unless the coupling has had a mismatch. Taking $H = H_{\mathrm{T}}$ gives $\mathsf{d}_{\mathrm{T}}$.

The first-block mark is a measurable function of these paths. To see this directly, for a fixed block let $A_I(t)$ be the maximum of its internal belief and model relative entropies on $[t-1, t]$, and let $B_I(t)$ be the minimum boundary belief relative entropy there. Configuration paths are continuous, so these two window functions are continuous. Replace the strict condition $B_I(t) > 1/2$ by the countable union of conditions $B_I(t) \geq 1/2 + 1/k$, $k \geq 1$. For each $k$ the qualifying end times form a closed subset of the compact interval $[1, H_{\mathrm{T}}]$. Its first time is measurable, as may be checked by minimizing $\max(A_I(t) - \sqrt{\mathrm{T}}, 1/2 + 1/k - B_I(t))$ on each compact initial subinterval. The infimum over $k$, the initial eligibility condition, and the finite comparison over blocks preserve measurability. This construction implements the strict boundary convention without asserting that its infimum is attained.

A measurable map of paths does not increase total variation. Let $a_I, b_I$ be the probabilities of the non-cemetery marks in the learned and frozen arms, $p = \sum_I a_I$, $q = \sum_I b_I$, and $d$ the total-variation distance of the two cemetery-augmented mark laws. Then

$$2d = \sum_I |a_I - b_I| + |p - q|, \qquad d \leq \mathsf{d}_{\mathrm{T}}.$$

If $p \geq q > 0$, the triangle inequality gives

$$\frac{1}{2} \sum_I \left| \frac{a_I}{p} - \frac{b_I}{q} \right| \leq \frac{\sum_I |a_I - b_I| + (p - q)}{2p} = \frac{d}{p}.$$

Interchanging the arms proves the bound $d/\max(p, q)$. The frozen matching-seed sub-event in proposition 14.4 belongs to $E_{\mathrm{T}}$ and has probability at least $\mathsf{c}(\mathsf{m}_0)$. Thus $q \geq \mathsf{c}(\mathsf{m}_0)$ and $p \geq q - d \geq \mathsf{c}(\mathsf{m}_0) - \mathsf{d}_{\mathrm{T}}$. The claimed conditional bound follows, and its limiting conclusion is immediate. □

The proposition keeps the original contrast event and uses its seeded sub-event only as a probability lower bound. That sub-event already has the restoring interaction and stationary comparison of proposition 13.22; the proof does not certify every history in the broader contrast event under the coherence-episode rule. Any common measurable admissibility restriction retaining the seeded sub-event has the same comparison.

The result does not cover the original coupled choice $\epsilon_{\mathrm{L}} = \mathrm{T}^{1/8}$, for which its bound is vacuous. It also does not compare different initial means or establish a conditional parent error. On the slower schedules, late assignment selection can coexist with a vanishing difference between the conditional laws of the first contrast on $[0, H_{\mathrm{T}}]$. These are different observables on different horizons.

## C.12 Activation from the Actual Conditioned Recognition Law

**Proposition C.18 (Conditional information after a completed coherence prefix).** In the learned four-cycle, fix $0 < \mathrm{T} \leq 1/100$, a selected matching $M$, and $s = t_{\mathrm{L}} \geq 1$. Suppose, throughout

$[s, 2s]$,

$$D_{\mathrm{KL}}(q_t \| \mu_{\mathrm{T},p(t)}) \le \eta, \qquad \mathrm{TV}(q_t, \nu_{\mathrm{T}}^M) \le \delta.$$

Set $\mathsf{a} = \mathrm{T}^{1/4}$, $\theta = s + 1$, and let $E = \mathscr{E}_M(\mathsf{a}, 1)$ be the actual matching and coherence event of proposition C.16 on $[s, s+1]$. Assume

$$\rho = 3\delta + \mathsf{b}(\mathrm{T}, \mathsf{a}, 1) < \tfrac{1}{2}, \qquad h_b(r) = -r \log r - (1-r)\log(1-r), \qquad \kappa_E = \frac{\eta + h_b(\rho)}{1-\rho}.$$

For all $t \ge \theta$, write $q_t^E = \mathrm{Law}(\mathbf{X}_t, S_t \mid E)$, always conditioning on this same completed prefix. Keep the single marked channel $C$ of proposition 13.22(b), including its identity action off the matching strata and its identity extension on evaluator samples. The parent activation law is exactly

$$\mu_{\mathrm{act}} = C_{\#} q_\theta^E.$$

Let $\mathcal{H}_{\mathrm{opt}}(q)$ be the mean discarded conditional relative entropy of $q$ against the restricted optimizer under $C$. Let $\mathcal{H}_{\mathrm{phys}}(q)$ be the corresponding mean discarded conditional relative entropy against the full physical posterior, after lifting $q$ with its recognition evaluator and retaining every evaluator sample. Then, for every $t \in [\theta, 2s]$,

$$D_{\mathrm{KL}}(q_t^E \| \mu_{\mathrm{T},p(t)}) \le \kappa_E, \tag{C.129}$$

$$\mathcal{H}_{\mathrm{opt}}(q_t^E) \le \kappa_E, \tag{C.130}$$

$$\mathcal{H}_{\mathrm{phys}}(q_t^E) \le \kappa_E + \log \frac{4 + 5\mathrm{T}}{4 + \mathrm{T}} \le \kappa_E + \mathrm{T}. \tag{C.131}$$

The two hidden-conditional bounds hold with any strictly positive alternative source rows in their reference laws; the full-law bound eq. (C.129) refers to the native rows $p(t)$.

These are unscaled, dimensionless relative entropies. The corresponding scaled optimizer free-energy loss is at most $\mathrm{T}\kappa_E$; the unscaled physical negative-ELBO loss is at most $\kappa_E + \mathrm{T}$, and its $\mathrm{T}$-scaled value is at most $\mathrm{T}\kappa_E + \mathrm{T}^2$. Along the joint selection schedule with $\eta, \delta \le \mathrm{T}$,

$$\rho = O(\mathrm{T}^{1/4}), \qquad \kappa_E = O(\mathrm{T}^{1/4} \log(1/\mathrm{T})).$$

The assumption $\rho < 1/2$ then holds for all sufficiently small $\mathrm{T}$, so these bounds are asymptotic. [ESTABLISHED]

*Proof.* The learned-arm persistence bound eq. (C.124), applied with $t_0 = s$ and $\mathsf{h} = 1$, gives $r := \Pr(E^c) \le \rho$; its proof uses only the learned-arm hypothesis. Thus $E$ is a $\theta$-measurable event of probability at least $1 - \rho$. The law $q_\theta^E$ is the actual conditional law and may have retained–hidden correlations and a non-Gaussian hidden conditional. Its marked pushforward is normalized by construction. At later times no further event is conditioned upon. In particular, the deterministic native row path is not recomputed from $q_t^E$.

For a fixed $t \in [\theta, 2s]$, put $B = \mathbf{1}_E$ and $R_t = (\mathbf{X}_t, S_t)$. Compute the relative entropy of the joint law of $(B, R_t)$ against $\mathrm{Bernoulli}(1-r) \otimes \mu_{\mathrm{T},p(t)}$ in two orders:

$$\begin{aligned}(1-r)D_{\mathrm{KL}}(q_t^E \| \mu_{\mathrm{T},p(t)}) + rD_{\mathrm{KL}}(q_t^{E^c} \| \mu_{\mathrm{T},p(t)}) \\ = D_{\mathrm{KL}}(q_t \| \mu_{\mathrm{T},p(t)}) + \mathbb{I}(B; R_t) \le \eta + h_b(r).\end{aligned} \tag{C.132}$$

When $r = 0$, use just the first term. The second term is nonnegative, $h_b$ is increasing on $[0, 1/2]$, and $1 - r \ge 1 - \rho$. This proves eq. (C.129). It also establishes the needed absolute continuity and finiteness; directly, $q_t^E \le q_t/(1-r)$.

The channel's fibers can be checked explicitly. On a matching with edges $e, f$, let $F_M$ have the orthonormal columns

$$(n_e^+ - n_f^+)/\sqrt{2}, \qquad n_e^-, \qquad n_f^-.$$

Write $Z_M = \operatorname{im} F_M$, $W_M = Z_M^{\perp}$, and $\mathbf{x} = y + F_M z$. The normal identities in eq. (C.94) give

$$\nu_{\mathrm{T},Z} = \mathcal{N}(0, \mathrm{T}K_Z^{-1}), \qquad K_Z = \mathrm{diag}(8 + \mathrm{T}, 4 + 5\mathrm{T}, 4 + 5\mathrm{T})$$

as the hidden optimizer conditional. Both the interaction Hessian and the evaluator precision preserve $W_M$ and $Z_M$. Consequently this conditional is independent of the retained coordinate and of the prior rows. It is a reference conditional, not an assertion that the actual conditioned recognition law has that form. The same statement holds on the other matching by rotation; off the two matchings the channel is the identity and has no discarded conditional.

The target of $C$ is a finite disjoint union of Euclidean strata, so disintegration gives the exact entropy chain rule

$$D_{\mathrm{KL}}(q_t^E \| \mu_{\mathrm{T},p(t)}) = D_{\mathrm{KL}}(C_{\#} q_t^E \| C_{\#} \mu_{\mathrm{T},p(t)}) + \mathcal{H}_{\mathrm{opt}}(q_t^E). \tag{C.133}$$

Dropping its retained term proves eq. (C.130). This argument uses the complete actual law; it neither resets the hidden state nor assumes Gaussianity conditional on a source history.

For the physical bound, work first on one matching fiber. Its physical outer hidden Gaussian is

$$\Pi_{\mathrm{T},Z} = \mathcal{N}(0, \mathrm{T}(K_Z^{\mathrm{phys}})^{-1}), \qquad K_Z^{\mathrm{phys}} = \mathrm{diag}(8 + \mathrm{T}, 4 + \mathrm{T}, 4 + \mathrm{T}).$$

All displayed matrices are positive definite. Write $v = 1 - \sigma_{\mathrm{ev}}^2 = 1/2$, $d_y = \mathfrak{D}y$, and $d_z = \mathfrak{D}F_M z$. Since $R_{\mathrm{ev}} = \mathfrak{D}^{\top}\mathfrak{D}/v$ preserves the split, $d_y \cdot d_z = 0$; its hidden eigenvalues are $0, 4, 4$.

Let $U = (m, k)$ denote the full evaluator sample vectors. At configuration $\mathbf{x}$, the physical evaluator has $m \sim \mathcal{N}(\bar{m}, \sigma_{\mathrm{ev}}^2 I)$ and independent increment $k - m \sim \mathcal{N}(0, vI)$. Its recognition counterpart, eq. (13.6), has the same first law and increment mean $d_y + d_z$. At a fixed retained $y$, form the comparison probability law

$$\overline{P}_y(dz, dU) = \Pi_{\mathrm{T},Z}(dz) P_{\mathrm{ev}}(dU \mid y + F_M z) \exp\left[\frac{d_y \cdot (k - m) - |d_y|^2/2}{v}\right]. \tag{C.134}$$

Completing the Gaussian square in $k - m$ proves that the multiplier integrates to one for every $z$. It changes the increment mean to $d_y$ and keeps $\Pi_{\mathrm{T},Z}$ as its $z$ marginal. The multiplier is retained-measurable in $(y, U)$, so $\overline{P}_y$ has the physical hidden conditional given these retained variables. At a fixed $z$, the recognition evaluator differs by the additional increment mean $d_z$, with relative entropy $|d_z|^2/(2v)$.

Compute the joint relative entropy first through $z$, then through $U$, and drop the nonnegative evaluator-marginal term. For the actual hidden law $q(dz \mid M, y)$, this gives

$$\mathcal{H}_{\mathrm{phys}}(q \mid M, y) \le D_{\mathrm{KL}}(q(dz \mid M, y) \| \Pi_{\mathrm{T},Z}) + \mathbb{E}_q[|d_z|^2/(2v) \mid M, y].$$

All terms are well-defined: the native process has finite second moments, conditioning increases the second-moment bound by a factor at most $(1 - \rho)^{-1}$, and the Gaussian reference logarithms are quadratic. The optimizer conditional is the exact tilt

$$\nu_{\mathrm{T},Z}(dz) = \frac{\Pi_{\mathrm{T},Z}(dz) e^{-|d_z|^2/(2v)}}{Z_{\mathrm{tilt}}}, \qquad Z_{\mathrm{tilt}} = \frac{4 + \mathrm{T}}{4 + 5\mathrm{T}}.$$

For its normalizer, the common coordinate is unchanged and each of the two discrepancy coordinates changes rate from $4 + \mathrm{T}$ to $4 + 5\mathrm{T}$; their two square-root Gaussian determinant factors multiply to $Z_{\mathrm{tilt}}$. Therefore

$$D_{\mathrm{KL}}(q(dz \mid M, y) \| \Pi_{\mathrm{T},Z}) + \mathbb{E}_q[|d_z|^2/(2v) \mid M, y] = D_{\mathrm{KL}}(q(dz \mid M, y) \| \nu_{\mathrm{T},Z}) + \log \frac{4 + 5\mathrm{T}}{4 + \mathrm{T}}.$$

Integrate over the retained law. Matching strata have total probability at most one, and all other strata have zero discarded loss. Finally,

$$0 \le \log \frac{4 + 5\mathrm{T}}{4 + \mathrm{T}} = \log \left( 1 + \frac{4\mathrm{T}}{4 + \mathrm{T}} \right) \le \frac{4\mathrm{T}}{4 + \mathrm{T}} \le \mathrm{T}.$$

This proves eq. (C.131). Only recognition has been conditioned on $E$; the physical target still uses the original observed records, with all evaluator samples retained.

A source-prior factor $\pi_S$ cancels in any conditional law on a fiber because $S$ is retained. The hidden optimizer and physical reference conditionals are consequently the same for every strictly positive choice of prior rows. The full-law entropy need not remain small for different rows, because its retained term may change. An imposed zero prior on an occupied stratum would remove reference support and is excluded; ordinary finite-time learning from positive rows preserves positivity.

The exact free-energy chain rules give the stated scaled and unscaled losses. With $\eta, \delta \le \mathrm{T}$ and $\mathsf{a} = \mathrm{T}^{1/4}$,

$$\rho \le 3\mathrm{T} + 8\sqrt{2/\pi}\mathrm{T}^{1/4} + 3408\mathrm{T}^{1/2}.$$

Use $h_b(\rho) \le \rho \log(e/\rho)$ and $1 - \rho \ge 1/2$ eventually to obtain the displayed order for $\kappa_E$. □

**Remark C.19 (A high-probability prefix is needed).** The high probability of the declared prefix has mathematical content. Under the exact matching Gaussian $\nu_{\mathrm{T}}^{M}$, a hidden coordinate $z_1$ is independent of the retained state and has zero unconditioned hidden relative entropy. Conditioning on $\{z_1 > 0\}$, of probability $1/2$, instead gives a density ratio two on that half-space and hidden relative entropy $\log 2$ at every temperature. Thus a probability merely bounded below does not suffice for a vanishing conditional loss. Conditioning on the entire realized fine history is different again: it fixes the hidden endpoint and gives a Dirac conditional, singular relative to the nondegenerate Gaussian fiber. The activation here uses the full joint law conditional on $E$, not a realized fine history. [ESTABLISHED]

The proposition supplies conditional information at activation and for the native future marginals under that same completed prefix. It does not yet compare an independently evolving parent path. Its prefix is the late restored coherence event, and carries no first-formation or recursive-coalescence conclusion.

## Inherited dynamics from the conditioned activation law

**Proposition C.20 (Dynamics of the conditioned parent).** Use the selected matching, channel, and conditioned activation of proposition C.18. More explicitly, let $0 < \mathrm{T} \le 1/100$, $s = t_{\mathrm{L}}$, $\theta = s + 1$, and let $\mathscr{E}$ be the selected coherence event on $[s, \theta]$ at resolution $\mathsf{a}$ from proposition C.16. Assume $\Pr(\mathscr{E}^c) \le \rho < 1$, and, throughout $[s, \theta + \mathsf{h}] \subseteq [t_{\mathrm{L}}, 2t_{\mathrm{L}}]$,

$$\mathrm{TV}(q_t, \nu_{\mathrm{T}}^{M}) \le \delta, \qquad \max_i |p_i(t) - p_i^{M}| \le \delta.$$

Write $\varepsilon = \epsilon_{\mathrm{L}}$, $d = \delta + \varepsilon\mathsf{h} < 1/4$, and use relative future time $u = t - \theta$. Initialize the parent with

$$q_0^{C} = C_{\#}\,\mathrm{Law}(\mathbf{X}_\theta, S_\theta \mid \mathscr{E}), \qquad p^{C}(0) = p(\theta).$$

On the stratum $\{S\} \times W_S$, its continuous generator is $-\mathrm{A}_S y \cdot \nabla_{W_S} + \mathrm{T}\Delta_{W_S}$. Its rate-one proposal kernel is

$$J_p((S, y), \cdot) = \int \kappa_S(d\mathbf{x} \mid y) \sum_{S'} b_p(S' \mid \mathbf{x}) \delta_{(S', P_{W_{S'}}\mathbf{x})}, \tag{C.135}$$

where $b_p(S' \mid \mathbf{x}) = \mu_{\mathrm{T},p}(S' \mid \mathbf{x})$ is the product of receiver posterior rows at rows $p$, and $\kappa_S$ is the optimizing hidden conditional: the three independent Gaussian coordinates on a matching stratum, and a point mass off the matchings. The parent rows evolve from its own law:

$$\frac{d}{du} p_i^{C}(u) = \varepsilon\big(q_u^{C}(\sigma_i = +1) - p_i^{C}(u)\big). \tag{C.136}$$

This coupled parent system has a unique global law with interior rows.

Let $\mathrm{P}_{\theta,\mathsf{h}}^{\mathrm{fine},C}$ be the law of $(C(\mathbf{X}_{\theta+u}, S_{\theta+u}))_{0 \le u \le \mathsf{h}}$ conditioned only on $\mathscr{E}$, with the fine process continuing its original row path. Let $\mathrm{P}_{\mathsf{h}}^{C}$ be the parent path law from the displayed activation and its own

learning equation. Then

$$\mathrm{TV}(\mathsf{P}^{\mathrm{fine},C}_{\theta,\mathsf{h}}, \mathsf{P}^{C}_{\mathsf{h}}) \le 2\rho + (\mathsf{h}+3)\delta + \frac{4\mathsf{h}d}{1-4d}. \tag{C.137}$$

The comparison includes all retained coordinates and the assignment history, hence also the loop history. Self-proposals remain part of the rate-one kernel.

Along corollary 14.6, choose $\delta \le \mathrm{T}$, $\varepsilon \le \mathrm{T}^2$, $\mathsf{a} = \mathrm{T}^{1/4}$, and $\mathsf{h} = \mathrm{T}^{-1/4}$. One may take $\rho = 3\delta + \mathsf{b}(\mathrm{T}, \mathsf{a}, 1)$, where $\mathsf{b}$ is from proposition C.16. The right side of eq. (C.137) is then $O(\mathrm{T}^{1/4})$. This is a conditional, finite-window approximation after the selected coherence event; the activation law is not replaced by a Gaussian law. [ESTABLISHED]

*Proof.* For a matching $M$, the decomposition $\mathbb{R}^{12} = W_M \oplus Z_M$ of proposition 13.22 is preserved by $H_M$, $\Omega$, and $\mathsf{A}_M$. The hidden drift matrix is $\mathrm{diag}(8+\mathrm{T}, 4+5\mathrm{T}, 4+5\mathrm{T})$. The prior weight $\pi_S(p)$ cancels from the conditional law at fixed $(S, y)$, so $\kappa_S$ is independent of $p$. Off the matchings no coordinate is discarded. Thus the displayed continuous generator is the projection of every held-sheet diffusion, and eq. (C.135) is a probability kernel. A self-proposal has $S' = S$ and $P_{W_S}\mathbf{x} = y$, so it leaves the parent state unchanged. At a different assignment it uses that assignment's projection, including its twelve-dimensional stratum when appropriate.

For fixed rows, the pushed optimizing law $C_{\#}\mu_{\mathrm{T},p}$ is invariant. Indeed each diffusion is reversible on its Gaussian stratum, and $\mu_{\mathrm{T},p}(d\mathbf{x}, S)b_p(S' \mid \mathbf{x})$ is symmetric in $S, S'$. Pushing both endpoints through $C$ gives detailed balance for $J_p$. The loop covector annihilates each $Z_M$, so this construction retains the loop through the stratum changes as well.

We next check the existence assertion with evolving parent rows. For any continuous interior row path, concatenate the linear diffusions with the rate-one proposals, drawing $\kappa_S$ at each proposal. On every finite interval the clock has finitely many rings and every inserted Gaussian is finite almost surely, so this construction is nonexplosive. Its second moments remain finite: projection cannot increase the norm, an inserted hidden draw adds expected squared norm at most $5\mathrm{T}/8$, and the diffusion generator applied to $|y|^2$ is at most $24\mathrm{T} - 2\mathrm{T}|y|^2$. The full generator is consequently bounded above by $25\mathrm{T} - 2\mathrm{T}|y|^2$. The activation law has finite second moment because it conditions a finite-moment fine law on an event of positive probability.

If two row paths have entries in $[\mathrm{c}_0, 1-\mathrm{c}_0]$, couple their parent processes by the same initial state, Brownian increments, clocks, hidden Gaussian draws, and receiver uniforms until their first different proposal. The derivative of a posterior receiver probability with respect to its prior entry is at most $1/[4\mathrm{c}_0(1-\mathrm{c}_0)]$. Therefore

$$\mathrm{TV}(q^{C,p}_u, q^{C,\widetilde{p}}_u) \le \frac{1}{\mathrm{c}_0(1-\mathrm{c}_0)} \int_0^u \|p(v) - \widetilde{p}(v)\|_\infty dv.$$

The variation-of-constants map for eq. (C.136) consequently has Lipschitz constant at most $\varepsilon u_*^2/[2\mathrm{c}_0(1-\mathrm{c}_0)]$ on an interval of length $u_*$. On a short enough interval it is a contraction on continuous row paths in an interior cube. Its image stays in that cube after further shortening the

interval, since the row velocity has absolute value at most $\varepsilon$. Moreover its solution satisfies

$$p_i^C(u) \geq p_i^C(0)e^{-\varepsilon u}, \qquad 1 - p_i^C(u) \geq (1 - p_i^C(0))e^{-\varepsilon u}.$$

On any prescribed finite horizon these lower bounds give a common interior cube and a positive contraction step length. Successive steps give existence and uniqueness on that horizon, and then on all finite horizons. The source marginal is continuous because its jump rate is bounded, so the resulting rows are continuously differentiable. This constructs learning from the parent's own law, without imposing the fine row path after activation.

Let $\overline{\nu}_{\mathrm{T}}^M$ be the retained Gaussian on $W_M$. Write $\mathsf{O}_{\mathrm{T},\mathsf{h}}^C$ for its stationary held-$M$ path law. The projection of $\mathsf{O}_{\mathrm{T},\mathsf{h}}^M$ from proposition C.16 is exactly this law. Only the learned-arm conclusions of that proposition are used, and their proof does not use its frozen-arm hypothesis. Conditioning any probability law on an event of probability $1 - r$ changes its total variation by at most $r$ after any measurable projection. Thus the conditioned native retained future obeys

$$\mathrm{TV}(\mathsf{P}_{\theta,\mathsf{h}}^{\mathrm{fine},C}, \mathsf{O}_{\mathrm{T},\mathsf{h}}^C) \leq \rho + (\mathsf{h}+2)\delta. \tag{C.138}$$

This applies to a prefix event through $\theta$. No condition is placed on future coherence or future proposals. The same argument at the activation time gives $\mathrm{TV}(q_0^C, \overline{\nu}_{\mathrm{T}}^M \otimes \delta_M) \leq \rho + \delta$.

For the parent's own rows let $\zeta_i^C$ be the probability in row $i$ assigned to the nonpartner of $M$. Its equation and the bound on any probability give

$$0 \leq \zeta_i^C(u) \leq \zeta_i^C(0) + \varepsilon u \leq \delta + \varepsilon u \leq d.$$

The row-independent normalized sheet weights are at most the matching weight one for $0 < \mathrm{T} \leq 1/100$ by proposition C.15; hence, for any fixed row vector with this bound,

$$\mu_{\mathrm{T},p}(M) \geq \pi_M(p) \geq 1 - \sum_i \zeta_i \geq 1 - 4d.$$

Define the parent departure probability at a proposal by $k_p(y) = J_p((M,y), \{S' \neq M\})$. Detailed balance, and the fact that the optimizing conditional retained law on $M$ is $\overline{\nu}_{\mathrm{T}}^M$ independently of the rows, give

$$\int k_p(y)\overline{\nu}_{\mathrm{T}}^M(dy) \leq \frac{1 - \mu_{\mathrm{T},p}(M)}{\mu_{\mathrm{T},p}(M)} \leq \frac{4d}{1-4d}. \tag{C.139}$$

The first inequality bounds stationary flux from $M$ to its complement by the complement's stationary mass.

Maximally couple the parent activation with the stationary retained matching law and use common Brownian increments, clocks, and parent proposal draws up to the first parent departure. Before departure, the continuous paths agree because the matching drift splits, and self-proposals do not move the retained coordinate. The row path $p^C(u)$ constructed above is deterministic, even though it was computed from the parent's own law. Under the stationary comparison path, its marginal

at every time is $\overline{v}_{\mathrm{T}}^{M}$. Bound the first-departure compensator with its stopped indicator by the compensator evaluated on that whole stationary path. Applying eq. (C.139) at each deterministic $p^{C}(u)$ gives

$$\Pr(\text{parent departure before } \mathsf{h} \text{ while initially coupled}) \leq \frac{4\mathrm{h}d}{1-4d}.$$

No Gaussian assertion is made about the parent's law conditioned on survival or about the native hidden variables. We have proved

$$\mathrm{TV}(\mathsf{P}_{\mathsf{h}}^{C}, \mathsf{O}_{\mathrm{T},\mathsf{h}}^{C}) \leq \rho + \delta + \frac{4\mathrm{h}d}{1-4d}.$$

Combining this with eq. (C.138) proves eq. (C.137).

For the displayed joint choices, $d \leq \mathrm{T} + \mathrm{T}^{7/4} \leq 2\mathrm{T}$, $\mathrm{h}\delta \leq \mathrm{T}^{3/4}$, and $\rho = O(\mathrm{T}^{1/4})$. The activation and following window fit inside the selection window because $1 + \mathrm{T}^{-1/4} \leq t_{\mathrm{L}}$ for sufficiently small temperature under the mobility cap. The departure term is $O(\mathrm{T}^{3/4})$, giving the stated total error. The Gaussian comparison laws served only to bound that error; the parent's activation and subsequent learning remain the laws in the statement. □

# Appendix D

# Notation

Table D.1 lists the symbols that recur across chapters, each with the one meaning it has in this manuscript and the statement or section that introduces it. A superscript $x \in \{b, m\}$ names the belief or the model channel. Symbols used inside a single chapter are declared at its opening, and those of the formation family are collected in Table 13.1. Metrics are sans-serif, statement-local constants are upright, and $\mathcal{N}(\mu, \Sigma)$ is the normal law with mean $\mu$ and covariance $\Sigma$.

**Table D.1:** Notation of the main manuscript.

| Symbol | Meaning | Introduced |
|---|---|---|
| **Base point, bundles, and frames** | | |
| $\mathcal{C}$, $c$, $c_*$ | contextual base, a context, the fixed base point | definition 2.1 |
| $\pi : \mathscr{P}_G \to \mathcal{C}$, $G$, $e_G$ | principal right $G$-bundle, its structure group, the identity of $G$ | definition 2.2 |
| $u_i$, $\sigma$, $U_i$ | common frame of agent $i$ at $c_*$ for both channels; a reference point of $(\mathscr{P}_G)_{c_*}$; the frame relative to it, $u_i = \sigma \cdot U_i^{-1}$ | definition 2.2 |
| $g_i$, $R_i^x = \rho_x(g_i)$ | passive re-expression $u_i' = u_i \cdot g_i^{-1}$ (vertex gauge); the represented sample-coordinate map | section 2.2 |
| $\varphi_i$ | local function $F(u_i) = u_i \cdot \varphi_i$ of a fiber automorphism $F$ | proposition 2.6 |
| $T_{ij}$ | pointwise frame comparison, $u_j = u_i \cdot T_{ij}$, $T_{ij} = U_i U_j^{-1}$ | proposition 2.5 |
| $(\mathsf{K}, \mathscr{K})$, $(\mathsf{M}, \mathscr{M})$ | belief and model sample fibers | definition 2.2 |
| $\rho_b, \rho_m$; $\widehat{\rho}_x$ | representations of $G$ on the sample fibers; the pushforward action $\widehat{\rho}_x(g)q = (\rho_x(g))_{\#}q$ on laws | definition 2.2 |
| $\mathcal{B}_x$, $\mathcal{E}_x$, $\varpi_x$ | law fiber, associated bundle $\mathscr{P}_G \times_{\widehat{\rho}_x} \mathcal{B}_x$, and its projection | definition 2.4 |
| $q_i^b$, $q_i^m$ | belief-law and model-law sections of agent $i$, and their coordinate laws at $c_*$ | definition 2.8 |
| $\Phi$, $\widetilde{\Phi}$; $\Phi_i$, $\widetilde{\Phi}_i$ | cross morphisms $\mathcal{E}_b \to \mathcal{E}_m$ and $\mathcal{E}_m \to \mathcal{E}_b$; their representatives in the frame $u_i$ | definition 2.9 |
| **Interaction multigraph and links** | | |
| $V$, $N$, $i, j$ | agent set, number of agents (also the extensive size of a finite labeled family), agent indices | definition 2.8 |
| $\mathfrak{I} = (V, E, {}^{-})$, $e : j \to i$, $\bar{e}$ | interaction multigraph at $c_*$, the support of the source priors; oriented edge copy; reversal | definition 2.11 |
| $\Theta_e^x$, $\mathsf{t}_e^x$ | independent link on $e$, with $\Theta_{\bar{e}}^x = (\Theta_e^x)^{-1}$; its link map $\mathsf{t}_e^x(u_j \cdot g) = u_i \cdot \Theta_e^x g$ | definition 2.11 |

*continued on the next page*

| Symbol | Meaning | Introduced |
|---|---|---|
| $\hat{\Theta}_e^x$ | per-edge twist $U_i^{-1}\Theta_e^x U_j$, whose conjugacy class is invariant | proposition 2.12 |
| $\Theta_p^x$, $\mathrm{Hol}^x(p)$ | path product along a walk $p$; holonomy of a closed walk | eq. (2.30) |
| $\Theta_{I\leftarrow i}^x$, $\mathrm{Hol}_I^x$, $\widehat{\Theta}_e^x$ | tree transport to the root of a block $I$, based holonomy group of $I$, block-dressed boundary edge | definition 10.4 |
| **Samples, records, and the joint** | | |
| $k_i$, $m_i$; $\mathsf{Y}_i$ | belief sample and model presentation of agent $i$; its latent block $\mathsf{K}_i \times \mathsf{M}_i$ | definition 3.1 |
| $y \in \mathsf{Y}$, $o \in \mathsf{O}$ | population latent sample and latent space; observation and observation space | definition 3.1 |
| $X \in \mathsf{X}$, $\theta$ | structural data and their configuration space; generative parameters | definitions 3.2 and 4.1 |
| $a \in \mathfrak{A}$, $\partial a$, $K_a$, $E_{a,o}$ | record, its scope, its normalized record kernel, and its negative log density at the record $o$ | definition 3.1 and hypothesis 5.5 |
| $\mathrm{ev}_i$, $K^X_{i,m_i}$ | evaluator of agent $i$; the evaluated mechanism of the presentation $m_i$ | definition 3.3 |
| $\mathrm{Acc}_i$, $\mathbb{Q}_{i,o,X}$, $\mathrm{Cpl}(\mathbb{Q}_{\bullet,o,X})$ | access map, local recognition law, and the class of couplings of the local laws | definitions 3.2 and 3.13 |
| $\mathfrak{q}_i^x$, $\mathsf{Y}_i^{\mathrm{law}}$; $\mathsf{Y}^{\mathrm{cfg}}$, $\mathrm{pres}$, $\kappa^{\mathrm{cfg}}$ | law points and the law-point block of agent $i$; configuration space, presentation map, presentation kernel | definition 3.19 |
| $\mathbb{P}_{\theta,V}$, $\mathbb{Q}_{V,o,X}$, $\mathbf{\Pi}_{\theta,V,o,X}$ | constructed joint, selected recognition coupling, posterior slice | definitions 3.15 and 4.2 |
| $\widetilde{p}_X$ | unnormalized potential of an undirected model | definition 4.5 |
| $\mathcal{L}_V$, $\mathcal{F}_V = -\mathcal{L}_V$; $\mathcal{L}^{\mathrm{ext}}$, $\mathcal{F}^{\mathrm{ext}}$ | evidence bound and variational free energy; their extended versions, defined on every recognition law | definition 5.6 |
| $D_{\mathrm{KL}}$, $\mathrm{TC}$ | relative entropy with values in $[0,+\infty]$; total correlation | definition 3.10 and corollary 5.1 |
| **Attention and selective communication** | | |
| $\pi_{ij}$, $\mathcal{N}_i$, $J_i$ | source prior of receiver $i$, its source set (the support of $\pi_i$), the source label | definition 6.1 |
| $(i,h)$ | receiving row of head $h$ of agent $i$ | definition 6.1 |
| $\beta_{ij}^x$, $\alpha_i^x$, $\eta_{ij}^x$ | attention row of receiver $i$, receiver occupancy, and event law $\eta = \alpha\beta$, with the role mark $P$, $\Pi$, or $Q$ of the law they are computed from ($\beta^m$ names the model channel); the event law carries no head index | definition 6.4 |
| $\tau_i$, $\mathsf{c}_{ij}$, $\mathsf{r}_i$, $\mathcal{F}_i^{\mathrm{att},x}$ | attention temperature, replicator cost, row relaxation rate, row free energy | corollary 6.7 and definition 12.8 |
| $S_e^Q$, $S_e^K$, $H_e^{\mathrm{cmp}}$; $\mu_i^Q$, $\mu_j^K$ | query and key comparison maps into the comparison space; projected query and key means | definition 6.15 |
| $K_e^v$; $\mathsf{L}_e^{\mathrm{enc}}$, $\mathsf{L}_e^{\mathrm{tr}}$, $\mathsf{L}_e^{\mathrm{dec}}$, $\mathsf{L}_e$, $\mathsf{L}_e^{\mathrm{val}}$ | composed value kernel of an edge copy; encoder, transmission, decoder, composite, and value linear maps | definitions 6.10 and 6.12 |
| $\Sigma_i^+$, $\Sigma_i^{\mathrm{pred}}$, $\widetilde{W}_i = W_i/\tau_i$ | receiver posterior and predictive covariances; tempered receiver precision | proposition 6.19 |
| $S_j$, $\Upsilon$, $v$ | source sample; random receiver side information and its values | proposition 6.24 |

*continued on the next page*

| Symbol | Meaning | Introduced |
|---|---|---|
| **Fiber geometry, histories, and time scales** | | |
| $g^F$, $T^F$, $\mathbf{s}$; $\lambda$ | Fisher tensor, Amari–Chentsov tensor, score; coordinates of a recognition family | definition 7.2 |
| $\mathfrak{k}_q$, $\mathfrak{g}_q$ | orbit Fisher form and isotropy algebra at a law $q$ | corollary 7.7 |
| $\mathcal{Q}_i$, $\mathcal{Q}_V$, $\mathsf{G}$ | configuration manifolds of an agent and of the population at $c_*$; a declared configuration metric | definition 7.13 |
| $\mathcal{Q}_\ell$, $\mathsf{F}_\ell$, $\mathscr{V}_\ell$, $\mathrm{Fl}_t^{\mathscr{V}_\ell}$ | configuration manifold at depth $\ell$, coarse configuration map, natural-gradient field, and its flow | definition 7.44 |
| $\mathfrak{t}$ | Fisher duration, the information clock of a selected history | definition 7.35 |
| $\varkappa$; $\mathfrak{K}$, $\mathfrak{C}$; $\mathrm{ret}$ | parity amplitude; residual second-moment tensor and signed cross tensor of a joint score; retained-law map | definition 7.22, proposition 7.18, and theorem 7.14 |
| $f$, $s = (s_\mathrm{m}, \theta)$ | fast and slow coordinate blocks | definition 8.2 |
| $\mathsf{G}_\mathrm{f}$, $\mathsf{G}_\mathrm{s}$, $\varepsilon$, $\mathsf{G}_\varepsilon$, $F_\mathrm{red}$ | block metrics, rate ratio, rate-separated metric $\varepsilon\mathsf{G}_\mathrm{f} \oplus \mathsf{G}_\mathrm{s}$, reduced functional | lemma 8.3 and proposition 8.4 |
| **Coarse channels and the network state** | | |
| $\ell$ | depth of coarse-graining, and the level index of the formation family | definition 9.20 |
| $I$, $J$; $A$ | child blocks; the parent label of a block | corollary 9.10 |
| $C_A$, $\mathsf{Z}_A$, $\Delta_A$ | parent coarse kernel, parent state space, conditional-information defect of the coarse chain rule | corollary 9.10 and theorem 9.6 |
| $\mathfrak{X}_\ell$, $C_{\ell k}$ | objects and arrows of a scale diagram | definition 9.20 |
| $\gamma \in \mathfrak{G}$, $W_\gamma$ | label of a finite labeled family (a mixture or geometry label) and its weight | definition 9.22 |
| $\mathsf{s}_\gamma$, $\mathsf{v}_\gamma$, $\mathcal{I}_\gamma$, $\bar{\mathsf{p}}$, $\mathsf{w}$ | bulk log-mass density, tie slope $\nabla\mathsf{s}_\gamma$, bulk Fisher density, limiting label law, window coordinate $\mathsf{w} = N(\theta - \theta_0)$ | section 9.8 |
| $L_I$, $\mathbb{V}_{I,\mathrm{fix}}$, $f_I$ | connection Laplacian of a block, its fixed fiber $\ker L_I$, and the retained rank $\dim \mathbb{V}_{I,\mathrm{fix}}$ | definition 10.16 |
| **Exact coarse-graining and the multiscale free energy** | | |
| $V_\ell$, $\mathsf{Y}_\ell$, $C_\ell$, $C^\dagger$ | agents and state space at depth $\ell$, the depth arrow, the reverse kernel of a channel | definition 11.1 |
| $(\varrho_\ell, \mathsf{m}_\ell)$, $M_\ell$, $\bar{\mathsf{m}}_\ell$, $H_\ell$ | measure pair of a reference probability and a finite evidence measure; its mass, normalized law, and action | definition 11.1 |
| $\mathsf{U}$, $\mathcal{A}$ | reverse conditional expectation and the exact action increment of a channel | theorem 11.4 |
| $\mathcal{G}_\ell$, $\mathsf{g}$, $\mathcal{T}_\ell$ | space of Hoeffding interaction coordinates, an element, and the exact interaction step | definitions 11.9 and 11.12 |
| $D_\ell$, $D_{n\leftarrow\ell}$ | finite-step multipliers of the interaction step | definition 11.18 |
| $\mathsf{P}_\ell^\mathrm{ret}$, $r_{\ell+1}^{\mathcal{G}}$ | retained projection and closure residual | definition 11.21 |
| $Y_\ell = (Z_\ell, \mathfrak{N}_\ell)$, $\mathbf{Y}$ | state at depth $\ell$, made of coordinates and the graph datum; the whole tower | definition 12.1 |
| $R_\ell \in \mathsf{R}_\ell$; $C^\mathrm{struct}_{\ell+1\leftarrow\ell}$, $\widehat{\mathsf{Y}}_{\ell+1\leftarrow\ell}$ | structural mark sampled by the structural channel, such as a random partition, and its space; the structural channel and its augmented transition state | definition 12.3 |
| $\mathfrak{m}_\ell$, $\mathfrak{M}_\ell$ | membership kernel and joint endpoint membership kernel | definition 12.14 |

# Appendix E

# Open Problems

Table E.1 collects the open problems of this manuscript in the order of the chapters. Each row names the item, the numbered statement where it arises, which also states what holds, and a condition that would settle it. Established results, refutations among them, stay at their statements and are not listed. Open problems of the base and Gaussian settings are collected in the two supplements.

**Table E.1:** Open problems of the main manuscript.

| Item | Locator | Settling condition |
|---|---|---|
| **Part I: single-point foundations** | | |
| Nondominated likelihood families | open problem 3.9 | A treatment of the evidence bound, its local disintegrations, and the coarse channels at the level of joint measures and regular conditional laws, with a rule for record-null versions, in which the exact identities of chapters 5 and 9 are proved, or hypotheses that restore a jointly measurable likelihood; negatively, a record family for which those identities provably have no measure-level counterpart. |
| A normalized law-point prior | open problem 3.23 | For a class of law points, a sigma-finite reference on the law-point or presentation space, a finite normalizer for the exponentiated law-point energies, the transformation law of the prior under the active gauge action, record kernels that read law points, and a compatibility relation between the law-point posterior and the sample-level recognition marginals. |
| An intrinsic criterion for recognition structure | open problem 5.14 | A functional on admissible channels or partitions, expressible relative to one fixed joint and not monotone toward an endpoint, whose optimizer is nondegenerate; negatively, a proof of endpoint degeneracy for a precisely stated class of such functionals. |

*continued on the next page*

| Item | Locator | Settling condition |
|---|---|---|
| Convergence of the alternating scheme | open problem 5.32 | For models of hypothesis 5.5 with finitely many latent states, parameters in an open cone, and joint densities twice continuously differentiable in the parameters, under a full exact E phase alternating with the exact M-coordinate with a unique maximizer: conditions checkable on the model, including compactness of the upper level sets of the evidence, met by some model of the class with parameter dimension at least two whose stationary set at the limit value of an orbit is not discrete, under which every orbit converges to one stationary point of the extended bound; or a model of that class with compact upper level sets whose orbit has more than one limit point. |
| Discrete creation and deletion of interaction edges | open problem 6.9 | For the threshold rule of the statement (multigraphs of bounded multiplicity; deletion at a rate that is a continuous nonincreasing function of the recognition row weight and creation at a rate that is a continuous nonincreasing function of the averaged comparison energy, each positive below a threshold; rows on the current support following the replicator flow): for some pair of rate functions, a normalized joint at $c_*$ with the edge set as a structural coordinate, declared with a named law and coupled to the link, interface, and label variables, in which the collective VFE decomposition, the event law, and its pushforward extend to the random support and the collective VFE is nonincreasing under the master equation of the rule from every initial support and recognition law; or a proof that for every pair of rate functions no joint of that class has this property. |
| Label-level multi-head attention | open problem 6.22 | For heads of one receiver with exclusive labels and comparison records whose projected queries are correlated under the receiver's conditional Gaussian law and whose integrated positive label table does not factor (the factorizing subclass, including untied channel-split heads under a channel-product receiver law, is excluded): an exact recursion with bounded head count, a family of comparison-record models closed under the meta-attention pushforward composed with a declared head-merging channel, in which the parent's per-head rows are the disintegrations of the merged event law and its posterior mean is the coarse coordinate of its children's posterior means; or a member and a partition for which no comparison-record model at the block with fewer heads than the block has rows reproduces the posterior mean of the block's coarse coordinate. |

*continued on the next page*

| Item | Locator | Settling condition |
|---|---|---|
| Formation with evolving selective interfaces | open problem 6.27 | For a connected vector-tier block with Gaussian belief means as fast coordinates, factored linear communication through surjective interfaces of one fixed rank below the fiber dimension, comparison records read through the receiver's interface, the kinetics of definition 13.21 at temperature $\mathrm{T}$, and interfaces learned as generative parameters by a natural-gradient flow at mobility $\epsilon_{\mathrm{M}}$: a parent that is a retained mechanism of fixed dimension below the block's, independent of $\mathrm{T}$, $\epsilon_{\mathrm{M}}$, and the horizon, whose path law of the communicated coordinates is within a relative entropy tending to zero as $\epsilon_{\mathrm{M}} \to 0$ on horizons of order $1/\epsilon_{\mathrm{M}}$, uniformly over $\mathrm{T} \in (0, \mathrm{T}_0]$; or a member whose learned interface path has a rest point at which no such parent is projectable, or in which every such parent keeps that relative entropy bounded below as $\epsilon_{\mathrm{M}}$ and $\mathrm{T}$ tend to zero along a stated relation. |
| Global Fisher quotients and canonical block data | open problem 7.15 | A criterion, checkable on the configuration manifolds at $c_*$ and proved, under which the null foliation of a pulled-back Fisher tensor has a Hausdorff manifold leaf space, for instance through a free and proper group action, together with block data on the quotient constructed from a stated typing of agents and channels, or a counterexample showing that the typing does not determine them. |
| Scope beyond the finite parity witness | open problem 7.29 | For a stated regular family of agent laws and a stated class of morphisms containing kernels that do not preserve products, a smooth right inverse of the marginal map that is natural under the class, and correlated lifts equivariant under a connected structure group acting on continuous law fibers, each by a construction with proof or a no-go theorem. |
| Quantitative meta-agent clock ratio | open problem 7.43 | A theorem computing or bounding the ratio of a meta-agent's Fisher duration to its constituents' collective duration, for a stated coarse map and a stated law of constituent velocities, including its dependence on $N$ for coherent and for incoherent histories. |
| Slow-manifold reduction of the collective free energy | open problem 8.7 | For the sampled four-cycle dynamics of definition 13.21 with model means and links at the common slow rate $\varepsilon_{\mathrm{m}}$ of proposition 8.6, no parameter learning, and the split by mobility: a bound on the distance between the slow coordinates of the sampled process and the reduced slow flow of its assignment path over a bounded slow-time interval, explicit in $\varepsilon_{\mathrm{m}}$, $\mathrm{T}$, and the number of assignment changes, tending to zero with $\varepsilon_{\mathrm{m}}$ and $\mathrm{T}$, for every positive integer $\mathrm{n}$ with probability tending to one on paths with at most $\mathrm{n}$ changes per unit slow time; or a positive integer $\mathrm{n}$, mobility ratios and temperatures tending to zero, and initial states from which, with probability bounded below, the path has at most $\mathrm{n}$ changes per unit slow time and that distance stays bounded below. |
| **Part II: exact coarse-graining at one base point** | | |

*continued on the next page*

| Item | Locator | Settling condition |
|---|---|---|
| Equivariant Bayes versions beyond the finite and compact cases | open problem 9.18 | For a locally compact second countable group acting continuously on Polish spaces, an equivariant channel $K$, and laws $Q$ with noncompact stabilizer whose posterior has no weakly continuous version on the support of $QK$: an equivariant measurable selection theorem giving, for every member, an everywhere-defined version equivariant under the stabilizer, jointly measurable across every measurable family of such laws meeting several orbits; or a pair $(Q, K)$ of the class with no everywhere-defined version equivariant under the stabilizer. |
| A natural selector of partitions and membership laws | open problem 9.19 | A relabeling-natural, gauge-compatible functional of the laws, not monotone under refinement, with a nondegenerate optimizer; or a theorem that no such selector exists in a stated class. |
| Invariant retained subspaces | open problem 11.25 | For finite agent spaces with fixed product references of full support, at least two blocks, and deterministic block-product channels whose block maps depend on every coordinate of their block, with base laws of full support in the class of product laws or of all laws: a proper invariant pair with fixed retained projections, invariant under the step of every base law in a nonempty open subset of the class, whose depth-$\ell$ range is not contained in the sum of the pullback sector and a grouping sector of at least two groups; or a proof that off a closed nowhere dense subset of the class every proper invariant pair has its depth-$\ell$ range in such a sum. |
| Dynamic closure of the pointwise parent | open problem 11.38 | For switching diffusions of the type of definition 13.21 indexed by $\epsilon$, channels as in part (b) of proposition 13.22 that hide coordinates read by the receiver rows on an assignment of weight bounded below, and parents that are time-homogeneous Markov on the retained assignment and coordinates or carry in addition a filter in the explicitly declared regular class of definition 11.37, at two levels (windows of fixed length uniformly up to the switching time, and the complete history over horizons on which the expected number of assignment changes diverges): a proof that every member admits such a parent whose retained and assignment path law is within a total variation tending to zero with $\epsilon$ of the native one; or a member for which every such parent keeps that distance bounded below. |
| Accounting across episodes | open problem 12.26 | For label-sufficient latent routes of definition 12.23, whose episode model is that of the parameter route with a finite label coordinate, the posterior slice as recognition coupling, and the aggregate update of constant step $\epsilon_{\mathrm{e}}$, whose read-out reads the label and an auxiliary preference coordinate only, and whose records are silent and informative or have likelihood vectors spanning a cone on which no nonzero polynomial vanishes: a route whose per-episode discrepancy from the aggregate update applied to its own filtered label law and label posterior is $o(\epsilon_{\mathrm{e}})$ uniformly over record histories and over episodes up to any fixed multiple of $1/\epsilon_{\mathrm{e}}$; or a lower bound of order $\epsilon_{\mathrm{e}}$ on that discrepancy for every such route. |

*continued on the next page*

| Item | Locator | Settling condition |
|---|---|---|
| Persistent multiscale hierarchy | open problem 12.27 | For models at $c_*$ with structural marks from a finite partition catalog inside both laws, descending the multiscale free energy by a reversible diffusion at temperature $\mathrm{T}$ with rate-one redraws of the marks, with energies, precisions, prior ratios, hyperparameters, and mobilities independent of $\mathrm{T}$, as $\mathrm{T} \to 0$: three results together, a lineage of at least two levels, formed with probability bounded below from every initial law with a density and a finite second moment, whose second level depends on the first (a declared handoff comparison with a margin bounded below over counterfactual handoffs, and no occurrence with probability tending to one uniformly over the handoff class) and is a coherence episode (a restoring clause at a declared resolution and a horizon clause), a nondegenerate intrinsic partition selector, and a separated slow manifold on which depth-$(\ell+1)$ variables feed back consistently into depth-$\ell$ kernels; or a proof that in this class no second level produced by descent, apart from a supplied-scale sequence whose event occurs with probability tending to one uniformly over the handoff class, meets the restoring and horizon conditions and exceeds its counterfactual statistic by such a margin. |
| **Part III: formation and interpretation** | | |
| Sharpness of the doubly exponential regime | open problem 13.16 | In the family of theorem 13.17 under (H1), (H3), and (H4), with power-law strengths satisfying $3\mathrm{Q}_\ell \geq 4\mathrm{Q}_{\ell+1} + 20$ and $\mathrm{Q}_\ell \geq 14$ below the top, the regime formed from $\mathrm{Q}_1$, and an entrance constant above the thresholds of the theorem and of proposition 13.15: conclusions (a) to (e) for every such schedule with constants and preparation neighborhoods independent of $L$, through an early estimate of the retained Fisher information by the future-layer energy, a history-environment bound without the top-interface recursion, and depth-uniform flag and jump constants; or an admitted preparation and a schedule of the class under which a conclusion fails, or a lower bound showing that a conclusion cannot hold with constants independent of $L$. |
| A family in which learning or holonomy selects the components | open problem 13.20 | On the four-cycle of proposition 14.4, with its temperature, learning mobility, and horizon, Gaussian initial laws $\mathcal{N}(\mathrm{m}_0, I_{12})$ with $\mathrm{m}_0$ outside a Lebesgue-null set, and parents of the type of part (b) of proposition 13.22: such parents with separately vanishing conditional and dynamical errors, together with a total-variation difference bounded below as $\epsilon \to 0$ between the conditional laws, given contrast formation, of the first formed block in the learned and the frozen-row arm, or between those of two initial means that differ only in their loop coordinate; or a proof that for almost every $\mathrm{m}_0$ both conditional differences vanish as $\epsilon \to 0$. |

*continued on the next page*

| Item | Locator | Settling condition |
|---|---|---|
| Repeated endogenous formation | open problem 13.25 | For models at $c_*$ with source-gated attention on a finite connected graph, outside the family of definition 13.23, whose energies, precisions, prior ratios, hyperparameters, and mobilities do not depend on $\mathrm{T}$, as $\mathrm{T} \to 0$, and a second coalescence declared by a handoff time, first- and second-stage events with their blocks, one statistic of the handoff state, and a handoff class fixed before the comparison: such a coalescence, with learned generative parameters and from initially nonmatching assignments, that does not occur with probability tending to one uniformly over the handoff class, whose statistic exceeds its counterfactual value by a margin bounded below, whose block is held at a declared resolution by a restoring interaction on horizons short compared with the stationary entrance time of the event, and whose parents have separately vanishing conditional and dynamical errors; or a proof that in this class every coalescence of first-stage blocks on such horizons fails the margin or the restoring clause or has a second-stage event occurring with probability tending to one uniformly over the handoff class. |
| Physical-law identification | open problem 14.7 | For a named target system whose data the model does not generate, observed through a declared record kernel that under the model's own law separates two objects of the coarse-graining differing in a gauge-invariant datum, with an observable, an estimator, and a validation and uncertainty protocol set before the data: held-out evidence, robust across admissible blocking schemes, that an object of the coarse-graining meets a baseline margin with probability-level closure; or held-out evidence under the same protocol that the margin is not met. |
| Single-point testability | open problem 14.8 | For a named target system at $c_*$ whose data the model does not generate, observed with sampling units and a noise model through a declared record kernel that reads the links of a closed walk or the coordinates of a candidate block: a test fixed before the data, with an estimator of gauge-typed parameters, a held-out statistic, an acceptance tolerance, and baselines without link holonomy or without formation, whose failure would falsify nontrivial holonomy on that walk or formation of that block, with a proof that under the kernel the law of the tested statistic depends on that hypothesis; negatively, a proof that under the kernel every statistic of the records has the same law for two link configurations differing only in the conjugacy class of the walk's holonomy, or whether or not the block is coherent at the declared resolution. |

# Appendix F

# Numerical Index

Table F.1 indexes the statements whose status is NUMERICAL in the sense of Table 1.1, each with its check, what the check tests, its control, and its seed. No statement of the main manuscript has that status: every lemma, proposition, theorem, corollary, and claim-bearing remark is proved here or cited to a primary source, and computations that corroborate its proved statements are not indexed here.

**Table F.1:** Numerical index of the main manuscript.

| Claim | Check | What is tested | Control | Seed |
|---|---|---|---|---|
| No statement of the main manuscript has the status NUMERICAL. | | | | |

# Bibliography


[1] P.-A. Absil, R. Mahony, and B. Andrews. Convergence of the iterates of descent methods for analytic cost functions. *SIAM Journal on Optimization*, 16(2):531–547, 2005. doi: 10.1137/040605266.

[2] Shun-ichi Amari and Hiroshi Nagaoka. *Methods of Information Geometry*, volume 191 of *Translations of Mathematical Monographs*. American Mathematical Society and Oxford University Press, 2000.

[3] N. Ay, J. Jost, H.V. Lê, and L. Schwachhöfer. Information geometry and sufficient statistics. *Probability Theory and Related Fields*, 162:327–364, 2015.

[4] N. Ay, J. Jost, H. V. Lê, and L. Schwachhöfer. *Information Geometry*, volume 64 of *Ergebnisse der Mathematik und ihrer Grenzgebiete*. Springer, 2017.

[5] Julio Backhoff-Veraguas and Clara Unterberger. On the specific relative entropy between martingale diffusions on the line, 2023. URL https://arxiv.org/abs/2207.03312.

[6] Dominique Bakry and Michel Émery. Diffusions hypercontractives. *Séminaire de Probabilités*, 19:177–206, 1985. URL https://numdam.org/item/SPS_1985__19__177_0/.

[7] Martin Bauer, Martins Bruveris, and Peter W. Michor. Uniqueness of the Fisher–Rao metric on the space of smooth densities. *Bulletin of the London Mathematical Society*, 48(3):499–506, 2016. doi: 10.1112/blms/bdw020.

[8] Cédric Bény and Tobias J. Osborne. Information-geometric approach to the renormalization group. *Physical Review A*, 92(2):022330, 2015. doi: 10.1103/PhysRevA.92.022330.

[9] David S. Berman and Marc S. Klinger. The inverse of exact renormalization group flows as statistical inference. *Entropy*, 26(5):389, 2024. doi: 10.3390/e26050389.

[10] David S. Berman, Jonathan J. Heckman, and Marc Klinger. On the dynamics of inference and learning, 2022.

[11] David S. Berman, Marc S. Klinger, and Alexander G. Stapleton. Bayesian renormalization. *Machine Learning: Science and Technology*, 4(4):045011, 2023. doi: 10.1088/2632-2153/ad0102.

[12] C.M. Bishop. *Pattern Recognition and Machine Learning*. Springer, 2006.

[13] Christian Borgs and Roman Kotecký. A rigorous theory of finite-size scaling at first-order phase transitions. *Journal of Statistical Physics*, 61(1–2):79–119, 1990. doi: 10.1007/BF01013955.

[14] Stéphane Boucheron, Gábor Lugosi, and Pascal Massart. *Concentration Inequalities: A Nonasymptotic Theory of Independence*. Oxford University Press, Oxford, 2013. ISBN 978-0-19-953525-5. doi: 10.1093/acprof:oso/9780199535255.001.0001.

[15] L. L. Campbell. An extended čencov characterization of the information metric. *Proceedings of the American Mathematical Society*, 98(1):135–141, 1986. doi: 10.1090/S0002-9939-1986-0848890-5.

[16] Eric A. Carlen. Superadditivity of Fisher's information and logarithmic Sobolev inequalities. *Journal of Functional Analysis*, 101(1):194–211, 1991. doi: 10.1016/0022-1236(91)90155-X.

[17] Ernst Cassirer. *Substance and Function and Einstein's Theory of Relativity*. Open Court, Chicago, 1923. Translated by W. C. Swabey and M. C. Swabey; originally published as Substanzbegriff und Funktionsbegriff, 1910.

[18] Alessio Catanzaro, Diego Garlaschelli, and Subodh P. Patil. Renormalization of interacting random graph models. *Physical Review E*, 113(2):024314, 2026. doi: 10.1103/34n8-pw8x.

[19] T. Cohen and M. Welling. Group equivariant convolutional networks. In *International Conference on Machine Learning*, pages 2990–2999, 2016.

[20] Taco S. Cohen, Maurice Weiler, Berkay Kicanaoglu, and Max Welling. Gauge equivariant convolutional networks and the icosahedral CNN. In *Proceedings of the 36th International Conference on Machine Learning (ICML)*, volume 97 of *Proceedings of Machine Learning Research*, pages 1321–1330. PMLR, 2019. URL https://proceedings.mlr.press/v97/cohen19d.html.

[21] T.M. Cover and J.A. Thomas. *Elements of Information Theory*. Wiley-Interscience, Hoboken, 2nd edition, 2006.

[22] Michael Creutz. Gauge fixing, the transfer matrix, and confinement on a lattice. *Physical Review D*, 15(4):1128–1136, 1977. doi: 10.1103/PhysRevD.15.1128.

[23] James P. Crutchfield and David P. Feldman. Regularities unseen, randomness observed: Levels of entropy convergence. *Chaos*, 13(1):25–54, 2003. doi: 10.1063/1.1530990.

[24] I. Csiszár. Information-type measures of difference of probability distributions and indirect observations. *Studia Scientiarum Mathematicarum Hungarica*, 2:299–318, 1967.

[25] John M. Danskin. *The Theory of Max-Min and its Application to Weapons Allocation Problems*, volume 5 of *Ökonometrie und Unternehmensforschung / Econometrics and Operations Research*. Springer-Verlag, Berlin, Heidelberg, 1967. doi: 10.1007/978-3-642-46092-0. Chapter III, "The directional derivative", pp. 19–32.

[26] A.P. Dempster, N.M. Laird, and D.B. Rubin. Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society Series B*, 39:1–38, 1977.

[27] R. L. Dobrushin. Central limit theorem for nonstationary Markov chains. I. *Theory of Probability and Its Applications*, 1(1):65–80, 1956. doi: 10.1137/1101006.

[28] James G. Dowty. Chentsov's theorem for exponential families. *Information Geometry*, 1(1): 117–135, 2018. doi: 10.1007/s41884-018-0006-4.

[29] Michael Esfeld and Vincent Lam. Moderate structural realism about space-time. *Synthese*, 160(1):27–46, 2008. doi: 10.1007/s11229-006-9076-2.

[30] N. Fenichel. Geometric singular perturbation theory for ordinary differential equations. *Journal of Differential Equations*, 31(1):53–98, 1979.

[31] Michael Friedman. *Dynamics of Reason*. CSLI Publications, Stanford, CA, 2001. ISBN 978-1575863467.

[32] Karl J. Friston, Thomas Parr, Conor Heins, Axel Constant, Daniel Friedman, Takuya Isomura, Chris Fields, Tim Verbelen, Maxwell J. D. Ramstead, John Clippinger, and Christopher D. Frith. Federated inference and belief sharing. *Neuroscience & Biobehavioral Reviews*, 156:105500, 2024. doi: 10.1016/j.neubiorev.2023.105500.

[33] Andrea Gabrielli, Diego Garlaschelli, Subodh P. Patil, and M. Ángeles Serrano. Network renormalization. *Nature Reviews Physics*, 7(4):203–219, 2025. doi: 10.1038/s42254-025-00817-5.

[34] Elena Garuccio, Margherita Lalli, and Diego Garlaschelli. Multiscale network renormalization: scale-invariance without geometry. *Physical Review Research*, 5(4):043101, 2023. doi: 10.1103/PhysRevResearch.5.043101.

[35] Michèle Giry. A categorical approach to probability theory. In *Categorical Aspects of Topology and Analysis*, volume 915 of *Lecture Notes in Mathematics*, pages 68–85. Springer, Berlin, Heidelberg, 1982. doi: 10.1007/BFb0092872.

[36] Martin Hairer and Jonathan C. Mattingly. Yet another look at Harris' ergodic theorem for Markov chains. In *Seminar on Stochastic Analysis, Random Fields and Applications VI*, volume 63 of *Progress in Probability*, pages 109–117. Birkhäuser, 2011. doi: 10.1007/978-3-0348-0021-1_7. URL https://arxiv.org/abs/0810.2777.

[37] Jakob Hansen and Robert Ghrist. Toward a spectral theory of cellular sheaves. *Journal of Applied and Computational Topology*, 3:315–358, 2019. doi: 10.1007/s41468-019-00038-7.

[38] Jakob Hansen and Robert Ghrist. Opinion dynamics on discourse sheaves. *SIAM Journal on Applied Mathematics*, 2021. doi: 10.1137/20M1341088.

[39] Carsten Hartmann, Lara Neureither, and Upanshu Sharma. Coarse-graining of non-reversible stochastic differential equations: quantitative results and connections to averaging, 2020. URL https://arxiv.org/abs/1911.06081.

[40] Lingshen He, Yiming Dong, Yisen Wang, Dacheng Tao, and Zhouchen Lin. Gauge equivariant transformer. In *Advances in Neural Information Processing Systems*, volume 34, 2021. URL https://proceedings.neurips.cc/paper/2021/hash/e57c6b956a6521b28495f2886ca0977a-Abstract.html.

[41] Conor Heins, Beren Millidge, Lancelot Da Costa, Richard P. Mann, Karl J. Friston, and Iain D. Couzin. Collective behavior from surprise minimization. *Proceedings of the National Academy of Sciences*, 121(17):e2320239121, 2024. doi: 10.1073/pnas.2320239121.

[42] D.D. Hoffman, M. Singh, and C. Prakash. The interface theory of perception. *Psychonomic Bulletin & Review*, 22:1480–1506, 2015.

[43] Donald D. Hoffman. *The Case Against Reality: Why Evolution Hid the Truth from Our Eyes*. W. W. Norton & Company, New York, 2019. ISBN 978-0393254693.

[44] Matthew D. Hoffman and Matthew J. Johnson. ELBO surgery: Yet another way to carve up the variational evidence lower bound. In *Workshop in Advances in Approximate Bayesian Inference, 30th Conference on Neural Information Processing Systems (NIPS 2016)*, Barcelona, Spain, 2016. Checked against the author PDF; the average-KL decomposition is Eqs. (11)–(12).

[45] R.A. Horn and C.R. Johnson. *Matrix Analysis*. Cambridge University Press, 2nd edition, 2013.

[46] Jessica N. Howard, Marc S. Klinger, Anindita Maiti, and Alexander G. Stapleton. Bayesian RG flow in neural network field theories. *SciPost Physics Core*, 8(1):027, 2025. doi: 10.21468/SciPostPhysCore.8.1.027.

[47] Antonin Joly and Nicolas Keriven. Graph coarsening with message-passing guarantees. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi: 10.52202/079017-3648.

[48] Abram Kagan and Lawrence A. Shepp. A sufficiency paradox: An insufficient statistic preserving the Fisher information. *The American Statistician*, 59(1):54–56, 2005. doi: 10.1198/000313005X21041.

[49] Olav Kallenberg. *Foundations of Modern Probability*. Springer, Cham, 3 edition, 2021.

[50] Immanuel Kant. *Critique of Pure Reason*. Cambridge University Press, Cambridge, 1998. ISBN 978-0521657297. Originally published 1781/1787.

[51] John G. Kemeny and J. Laurie Snell. *Finite Markov Chains*. Undergraduate Texts in Mathematics. Springer, New York, 1976. Reprint of the 1960 Van Nostrand edition.

[52] S. Kobayashi and K. Nomizu. *Foundations of Differential Geometry*, volume 1. Wiley, 1963.

[53] S. Kullback and R.A. Leibler. On information and sufficiency. *Annals of Mathematical Statistics*, 22:79–86, 1951.

[54] James Ladyman. Structural realism. In Edward N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, winter 2014 edition, 2014. URL https://plato.stanford.edu/archives/win2014/entries/structural-realism/.

[55] James Ladyman, Don Ross, David Spurrett, and John Collier. *Every Thing Must Go: Metaphysics Naturalized*. Oxford University Press, Oxford, 2007. ISBN 978-0199276196. doi: 10.1093/acprof:oso/9780199276196.001.0001.

[56] Hông Vân Lê. The uniqueness of the Fisher metric as information metric. *Annals of the Institute of Statistical Mathematics*, 69:879–896, 2017.

[57] John M. Lee. *Introduction to Smooth Manifolds*, volume 218 of *Graduate Texts in Mathematics*. Springer, New York, second edition, 2012. doi: 10.1007/978-1-4419-9982-5.

[58] Oren E. Livne and Achi Brandt. Lean algebraic multigrid (LAMG): Fast graph Laplacian linear solver. *SIAM Journal on Scientific Computing*, 34(4):B499–B522, 2012. doi: 10.1137/110843563.

[59] Thomas A. Louis. Finding the observed information matrix when using the EM algorithm. *Journal of the Royal Statistical Society Series B*, 44(2):226–233, 1982. doi: 10.1111/j.2517-6161.1982.tb01203.x.

[60] Andreas Loukas. Graph reduction with spectral and cut guarantees. *Journal of Machine Learning Research*, 20(116):1–42, 2019.

[61] Pankaj Mehta and David J. Schwab. An exact mapping between the variational renormalization group and deep learning, 2014.

[62] Dov Monderer and Lloyd S. Shapley. Potential games. *Games and Economic Behavior*, 14(1): 124–143, 1996. doi: 10.1006/game.1996.0044.

[63] Sadao Nakajima. On quantum theory of transport phenomena. *Progress of Theoretical Physics*, 20(6):948–959, 1958. doi: 10.1143/PTP.20.948.

[64] Radford M. Neal and Geoffrey E. Hinton. A view of the EM algorithm that justifies incremental, sparse, and other variants. In Michael I. Jordan, editor, *Learning in Graphical Models*, volume 89 of *NATO ASI Series*, pages 355–368. Springer, Dordrecht, 1998. doi: 10.1007/978-94-011-5014-9_12. URL https://link.springer.com/chapter/10.1007/978-94-011-5014-9_12.

[65] Yvan Notay. An aggregation-based algebraic multigrid method. *Electronic Transactions on Numerical Analysis*, 37:123–146, 2010.

[66] Kenta Oono and Taiji Suzuki. Graph neural networks exponentially lose expressive power for node classification. In *International Conference on Learning Representations*, 2020.

[67] Dénes Petz. Sufficiency of channels over von Neumann algebras. *The Quarterly Journal of Mathematics*, 39(1):97–108, 1988. doi: 10.1093/qmath/39.1.97.

[68] David Pollard. A note on insufficiency and the preservation of fisher information. In *From Probability to Statistics and Back: High-Dimensional Models and Processes*, volume 9 of *IMS Collections*, pages 266–275. Institute of Mathematical Statistics, Beachwood, OH, 2013. doi: 10.1214/12-IMSCOLL919. URL https://doi.org/10.1214/12-IMSCOLL919.

[69] Hans Reichenbach. *The Theory of Relativity and A Priori Knowledge*. University of California Press, Berkeley, 1965.

[70] Herbert Robbins. An empirical Bayes approach to statistics. In Jerzy Neyman, editor, *Proceedings of the Third Berkeley Symposium on Mathematical Statistics and Probability, Volume 1: Contributions to the Theory of Statistics*, pages 157–163. University of California Press, Berkeley, 1956. Checked against the Project Euclid chapter record.

[71] Biswa Sengupta, Arturo Tozzi, Gerald K. Cooray, Pamela K. Douglas, and Karl J. Friston. Towards a neuronal gauge theory. *PLoS Biology*, 14(3):e1002400, 2016. doi: 10.1371/journal.pbio.1002400.

[72] C.E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27: 379–423, 623–656, 1948.

[73] Amit Singer and Hau-Tieng Wu. Vector diffusion maps and the connection Laplacian. *Communications on Pure and Applied Mathematics*, 65(8):1067–1144, 2012. doi: 10.1002/cpa.21395.

[74] A. J. Stam. Some inequalities satisfied by the quantities of information of Fisher and Shannon. *Information and Control*, 2(2):101–112, 1959. doi: 10.1016/S0019-9958(59)90348-1.

[75] Peter Thestrup Waade, Christoffer Lundbak Olesen, Jonathan Ehrenreich Laursen, Samuel William Nehrer, Conor Heins, Karl Friston, and Christoph Mathys. As one and many: Relating individual and emergent group-level generative models in active inference. *Entropy*, 27(2):143, 2025. doi: 10.3390/e27020143.

[76] A. N. Tikhonov. Systems of differential equations containing small parameters in the derivatives. *Matematicheskii Sbornik (N.S.)*, 31(73)(3):575–586, 1952. In Russian.

[77] Jakub Tomczak and Max Welling. VAE with a VampPrior. In Amos Storkey and Fernando Perez-Cruz, editors, *Proceedings of the Twenty-First International Conference on Artificial Intelligence and Statistics*, volume 84 of *Proceedings of Machine Learning Research*, pages 1214–1223. PMLR, 2018. URL https://proceedings.mlr.press/v84/tomczak18a.html. Checked against the PMLR page and PDF; the aggregated posterior as ELBO-optimal prior is Eqs. (7)–(8).

[78] Bas C. van Fraassen. *The Scientific Image*. Oxford University Press, Oxford, 1980. ISBN 978-0198244271. doi: 10.1093/0198244274.001.0001.

[79] Petr Vaněk, Jan Mandel, and Marian Brezina. Algebraic multigrid by smoothed aggregation for second and fourth order elliptic problems. *Computing*, 56(3):179–196, 1996. doi: 10.1007/BF02238511.

[80] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, pages 5998–6008, 2017.

[81] Nikolai Nikolaevich Čencov. *Statistical Decision Rules and Optimal Inference*, volume 53 of *Translations of Mathematical Monographs*. American Mathematical Society, Providence, RI, 1982. ISBN 978-0821830512. Translated from Russian by Lev J. Leifman; originally published as *Statisticheskie reshayushchie pravila i optimalnye resheniya*, Moscow, Nauka, 1972.

[82] Pablo Villegas, Tommaso Gili, Guido Caldarelli, and Andrea Gabrielli. Laplacian renormalization group for heterogeneous networks. *Nature Physics*, 19(3):445–450, 2023. doi: 10.1038/s41567-022-01866-8.

[83] Christof Wetterich. Linear lattice gauge theory. *Nuclear Physics B*, 884:44–65, 2014. doi: 10.1016/j.nuclphysb.2014.04.002.

[84] John Archibald Wheeler. Information, physics, quantum: The search for links. In Wojciech H. Zurek, editor, *Complexity, Entropy, and the Physics of Information*, pages 3–28. Addison-Wesley, Redwood City, CA, 1990.

[85] K.G. Wilson. Confinement of quarks. *Physical Review D*, 10(8):2445–2459, 1974.

[86] J. Worrall. Structural realism: The best of both worlds? *Dialectica*, 43(1-2):99–124, 1989.

[87] C. F. Jeff Wu. On the convergence properties of the EM algorithm. *The Annals of Statistics*, 11(1):95–103, 1983.

[88] Zhitao Ying, Jiaxuan You, Christopher Morris, Xiang Ren, William L. Hamilton, and Jure Leskovec. Hierarchical graph representation learning with differentiable pooling. In *Advances in Neural Information Processing Systems*, volume 31, 2018.

[89] Robert Zwanzig. Ensemble method in the theory of irreversibility. *The Journal of Chemical Physics*, 33(5):1338–1341, 1960. doi: 10.1063/1.1731409.